Beyond the distributed monolith

...

Blanca Garcia Gil
Principal Systems Engineer, BBC
@blanquish

Table of contents

BBC NEWS

BBC SPORT

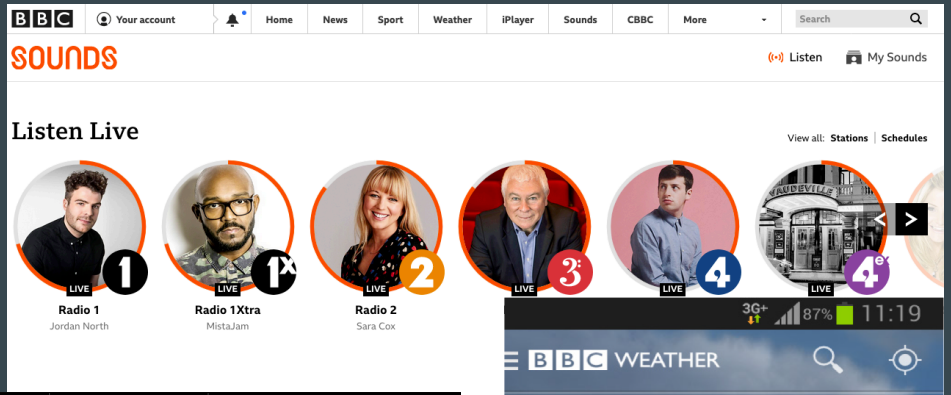

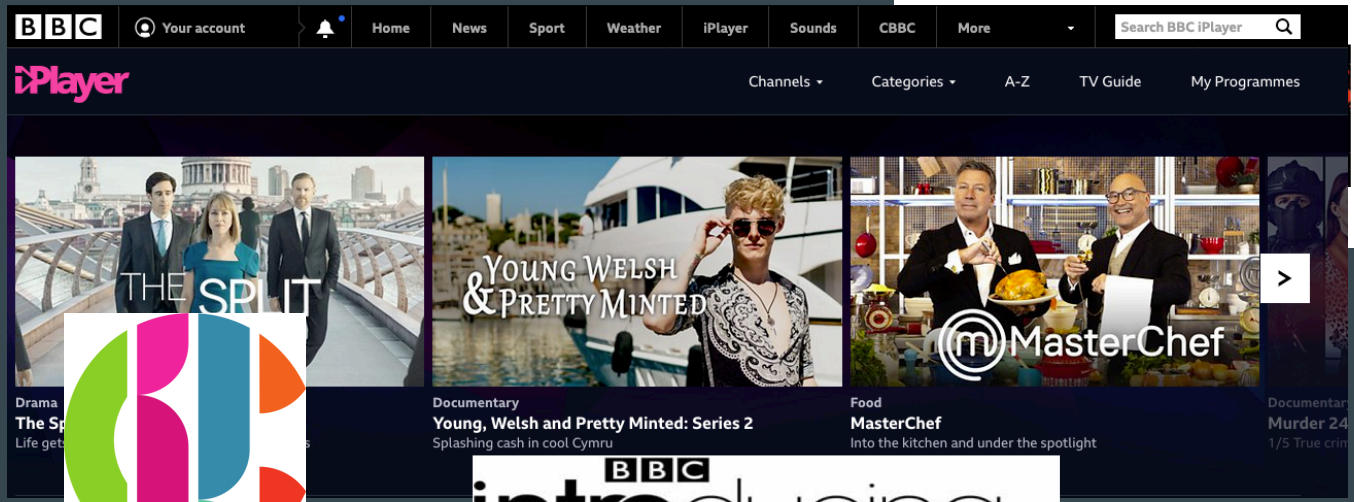

BBC introducing...

A more personal BBC


Why?

Privacy promise

A more *personal* BBC for you

A better BBC for *everyone*

https://www.bbc.co.uk/usingthebbc/privacy-promise

The acceptable face of personalisation

- Dr Who

Daleks

2. How does analytics processing work?

Typical data analytics end to end pipeline

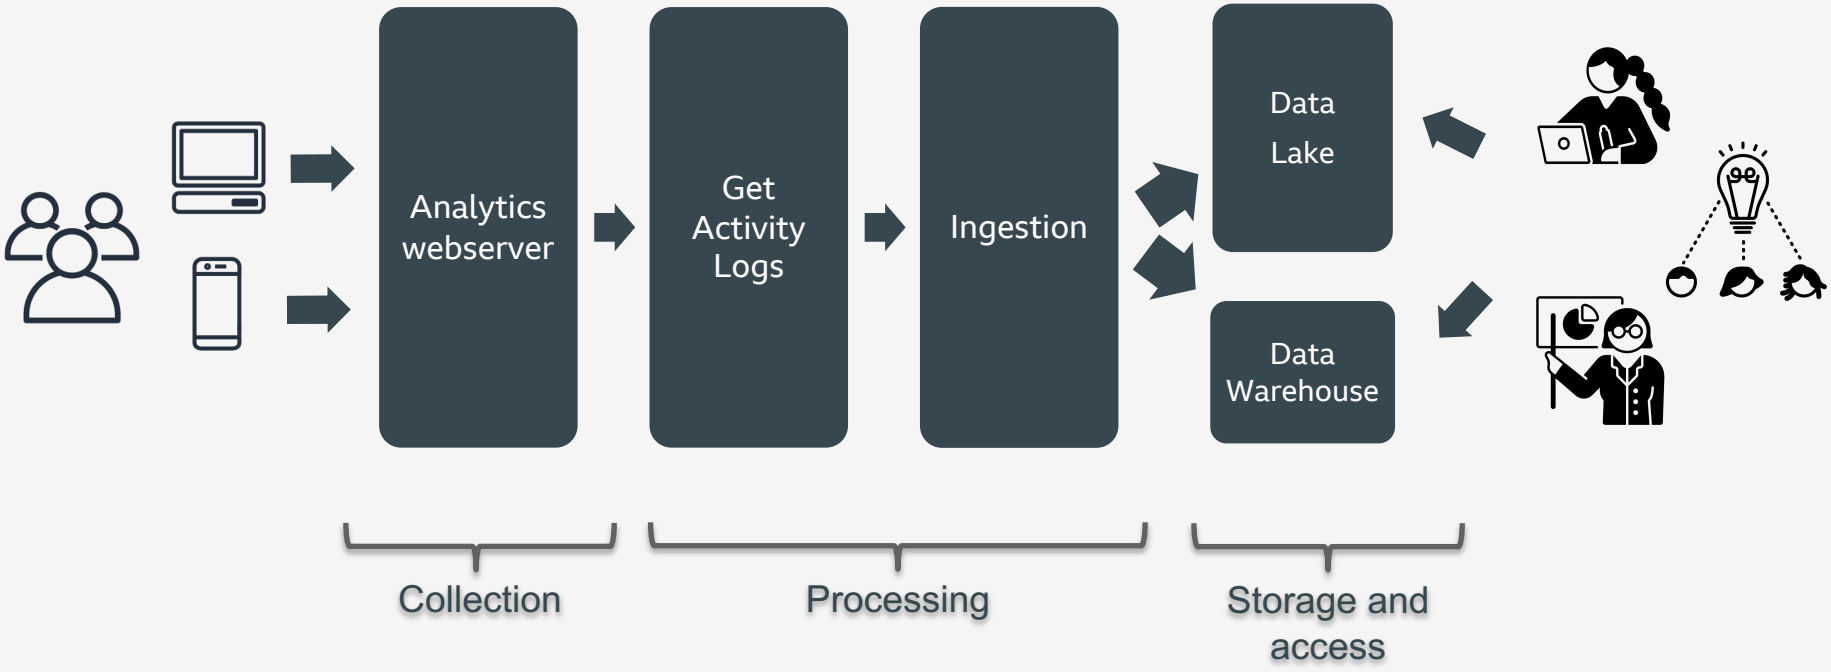

Typical data analytics ingestion flow

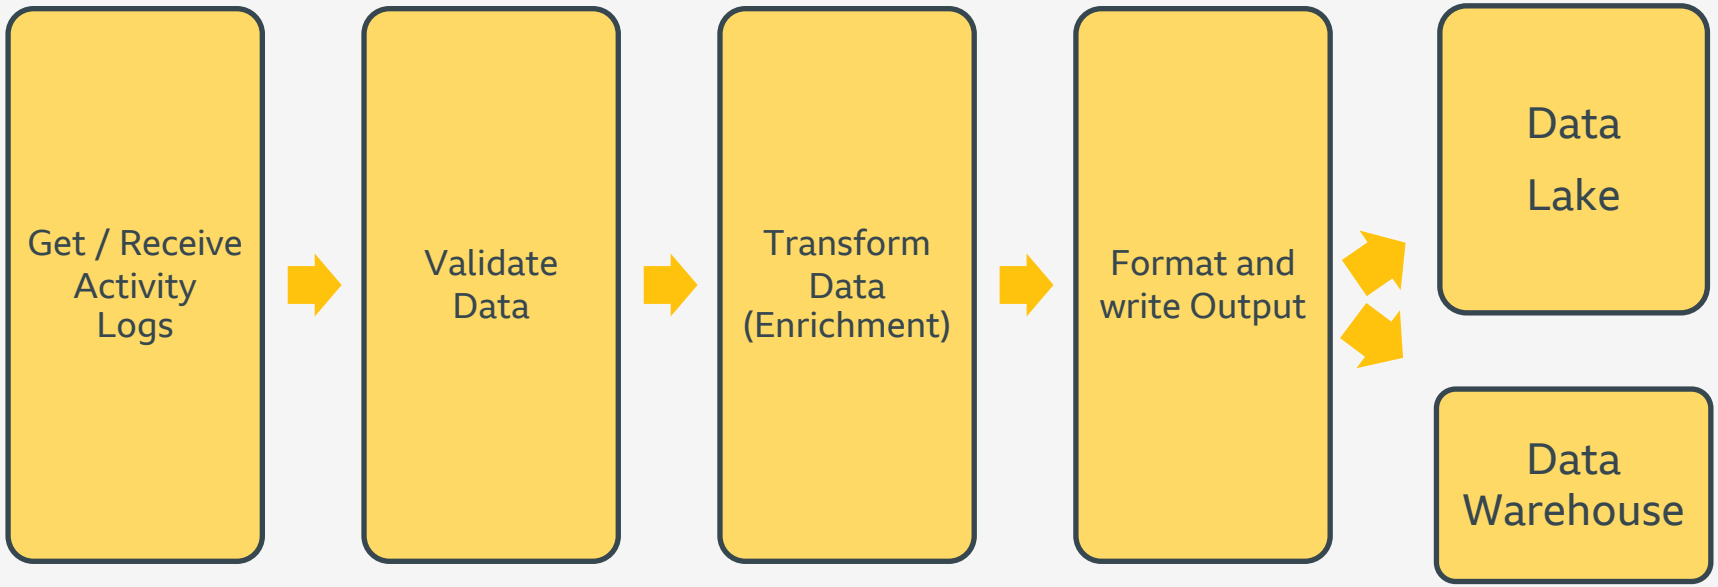

3. The distributed monolith and lessons learnt

How our data analytics pipeline architecture was

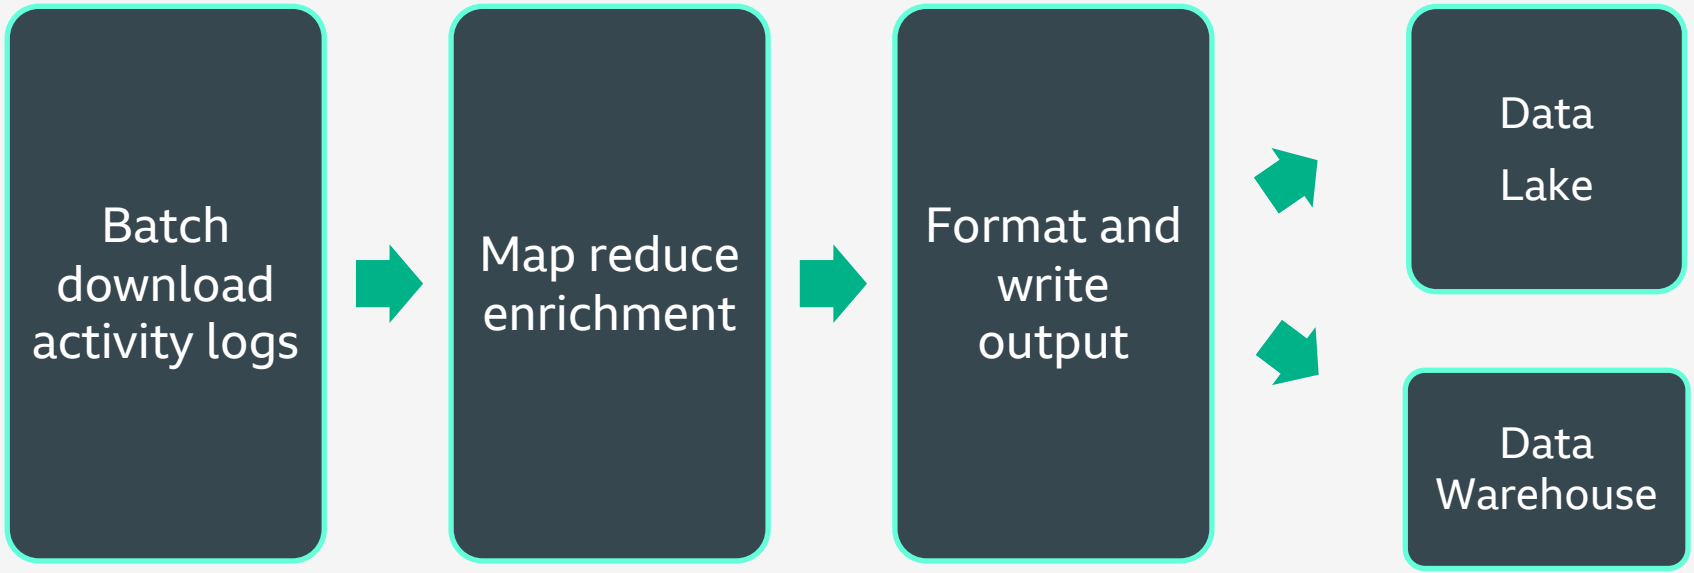

How our data analytics pipeline architecture was

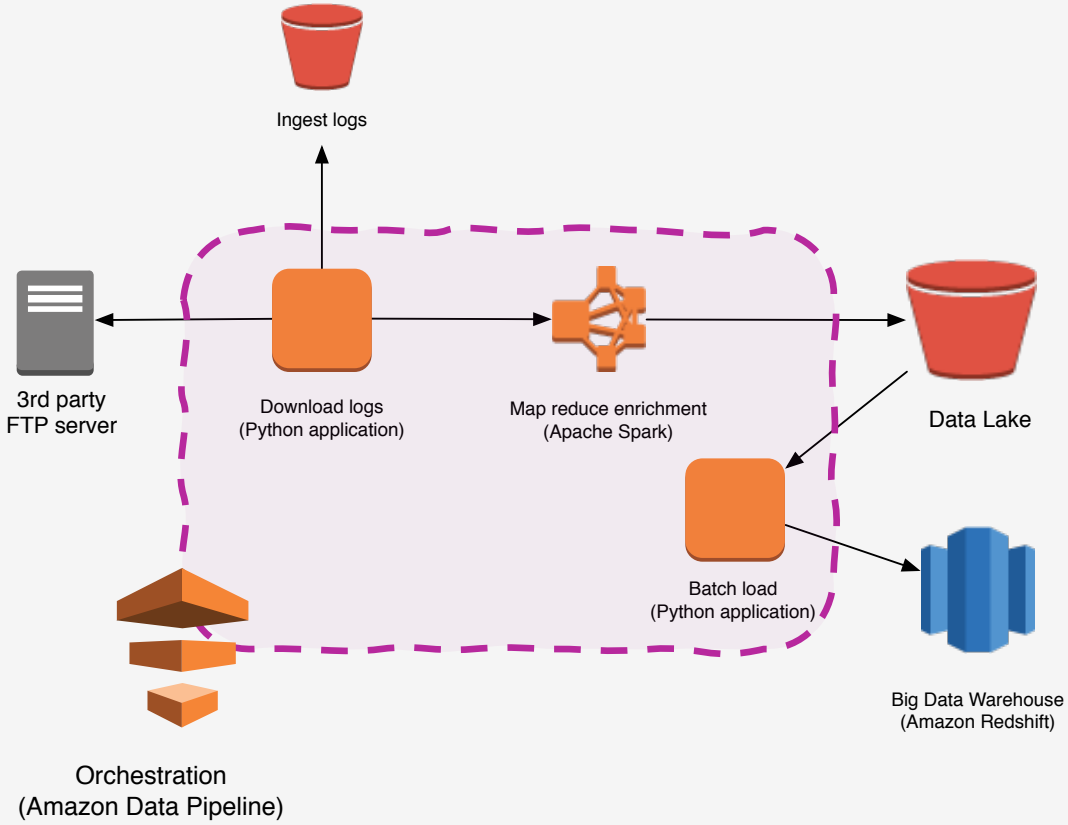

Microservices
+
Big Data

Highest scale data pipeline within our team

Billions of messages per day

Peta byte scale data lake

Lessons learnt

Lesson #1: batch processing

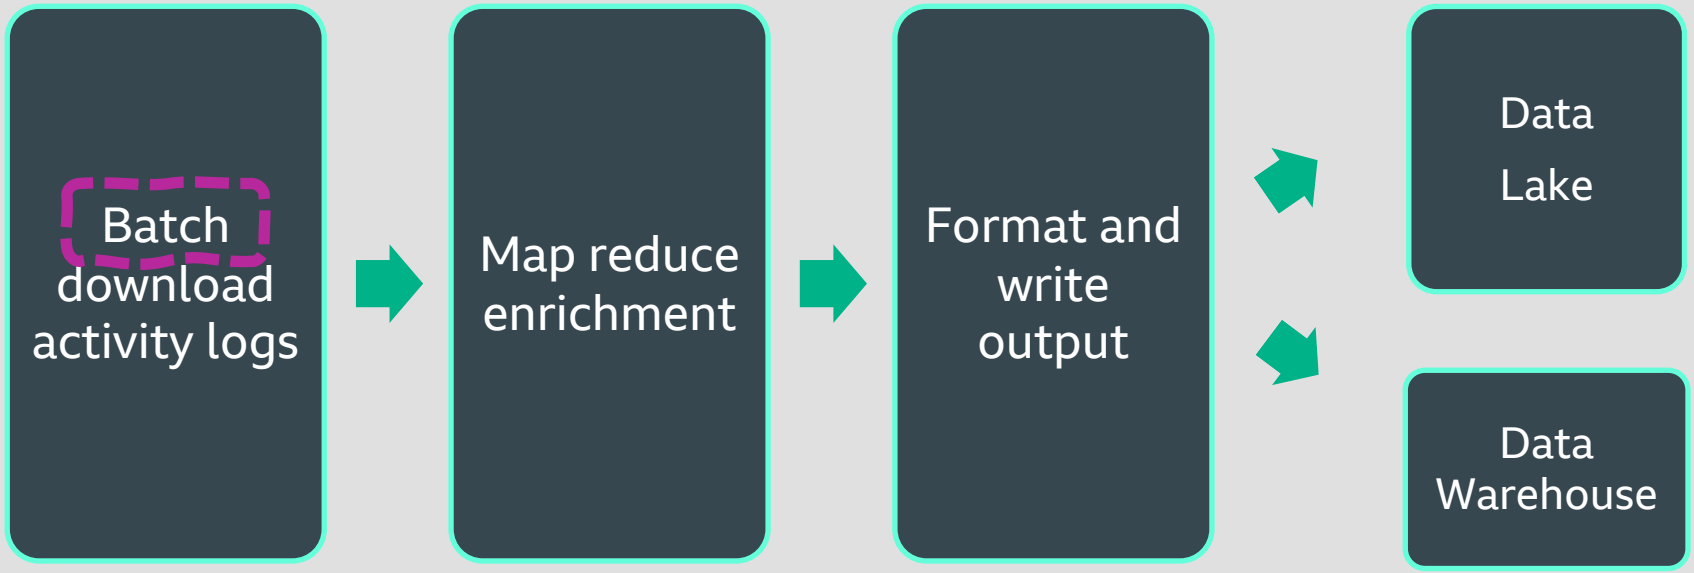


Finding needles

Lesson #2: validating input and testing

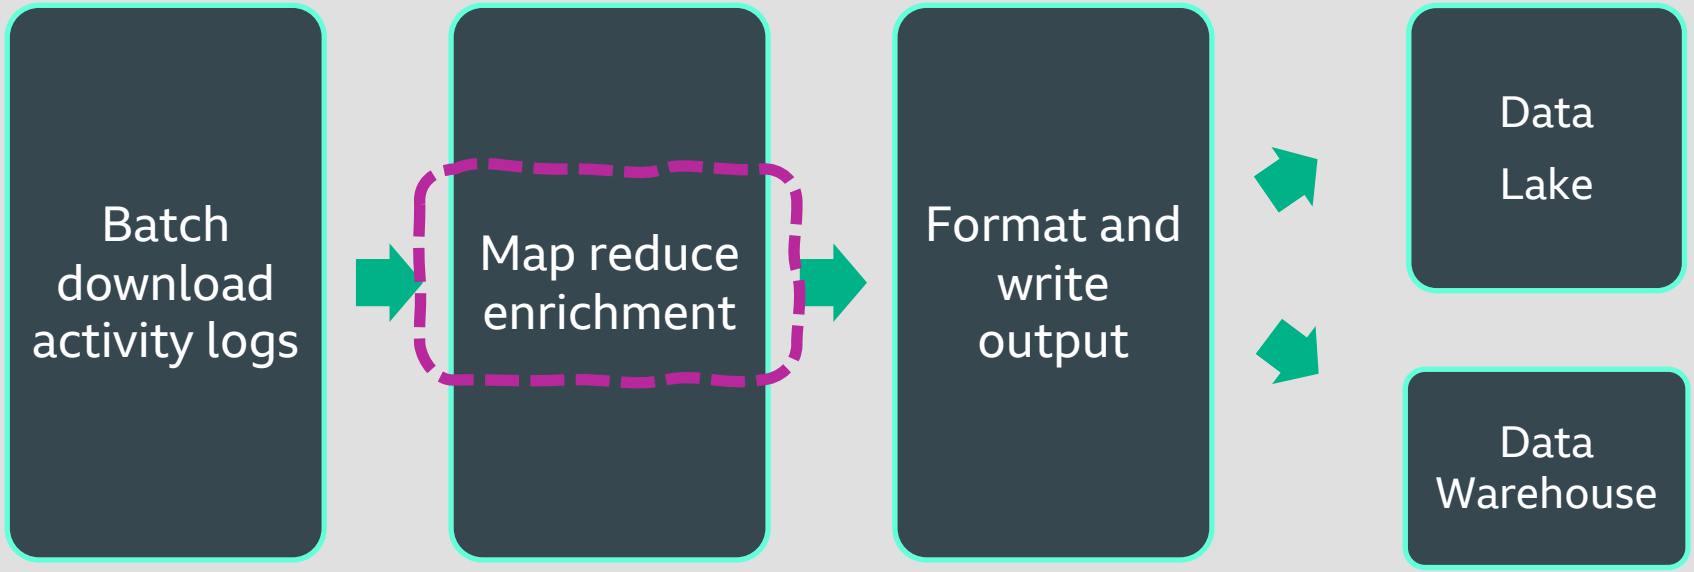

Testing

Very few tests made it hard to make changes to the code with confidence

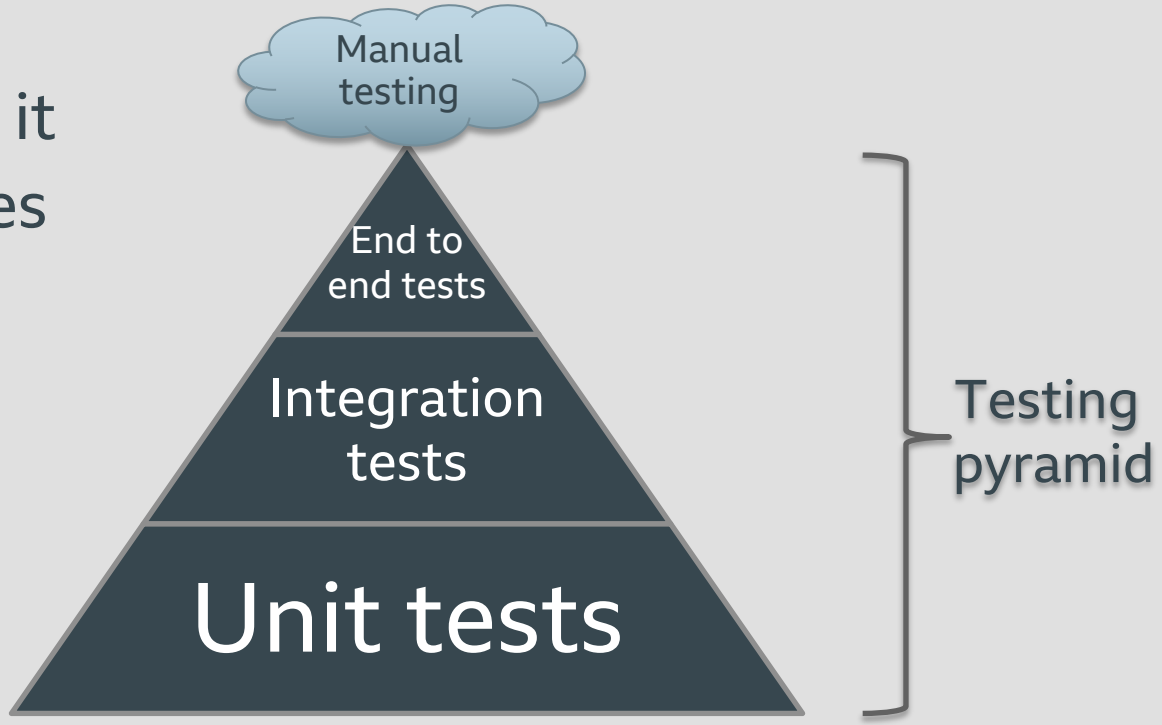

Lesson #3: tight coupling

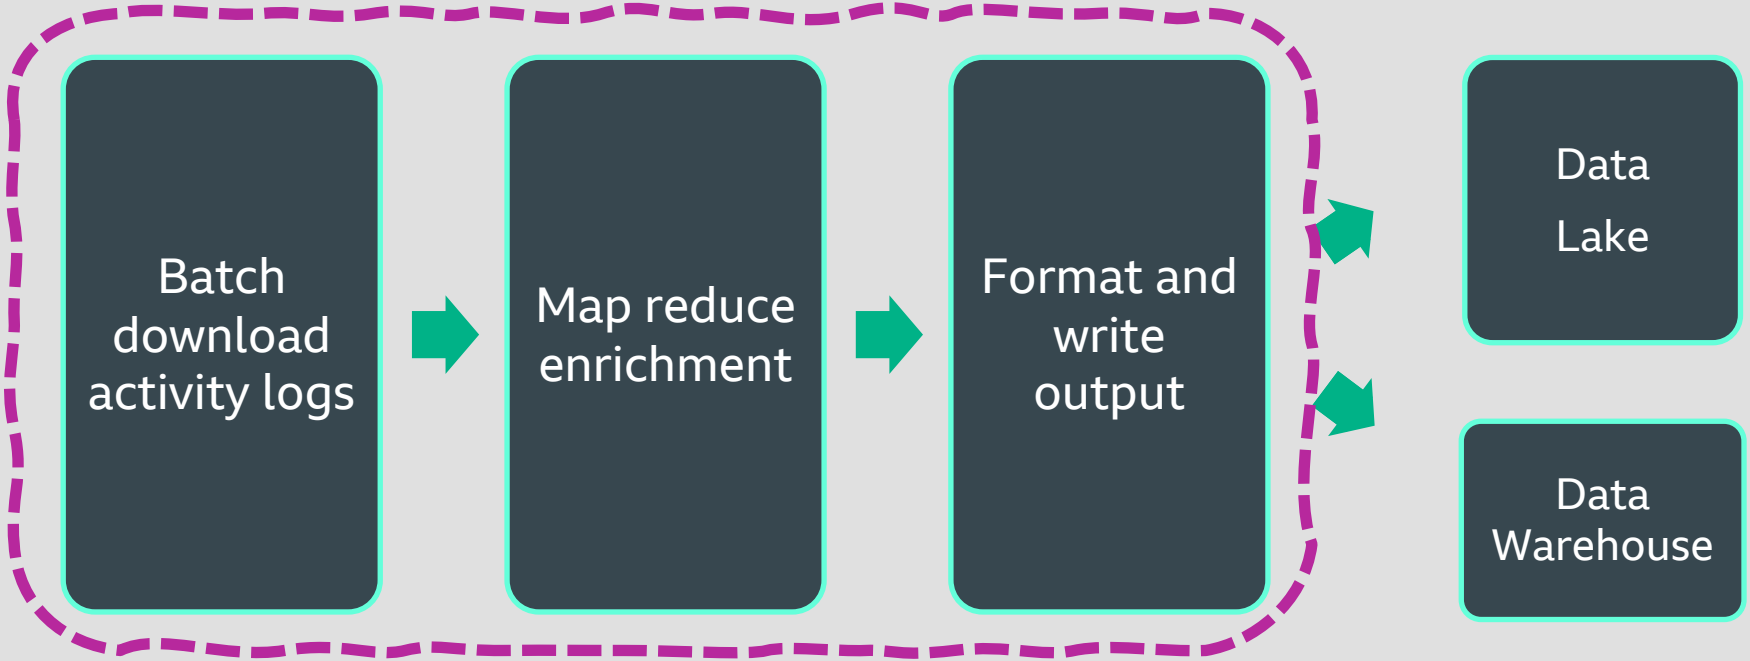

Lesson #3: tight coupling

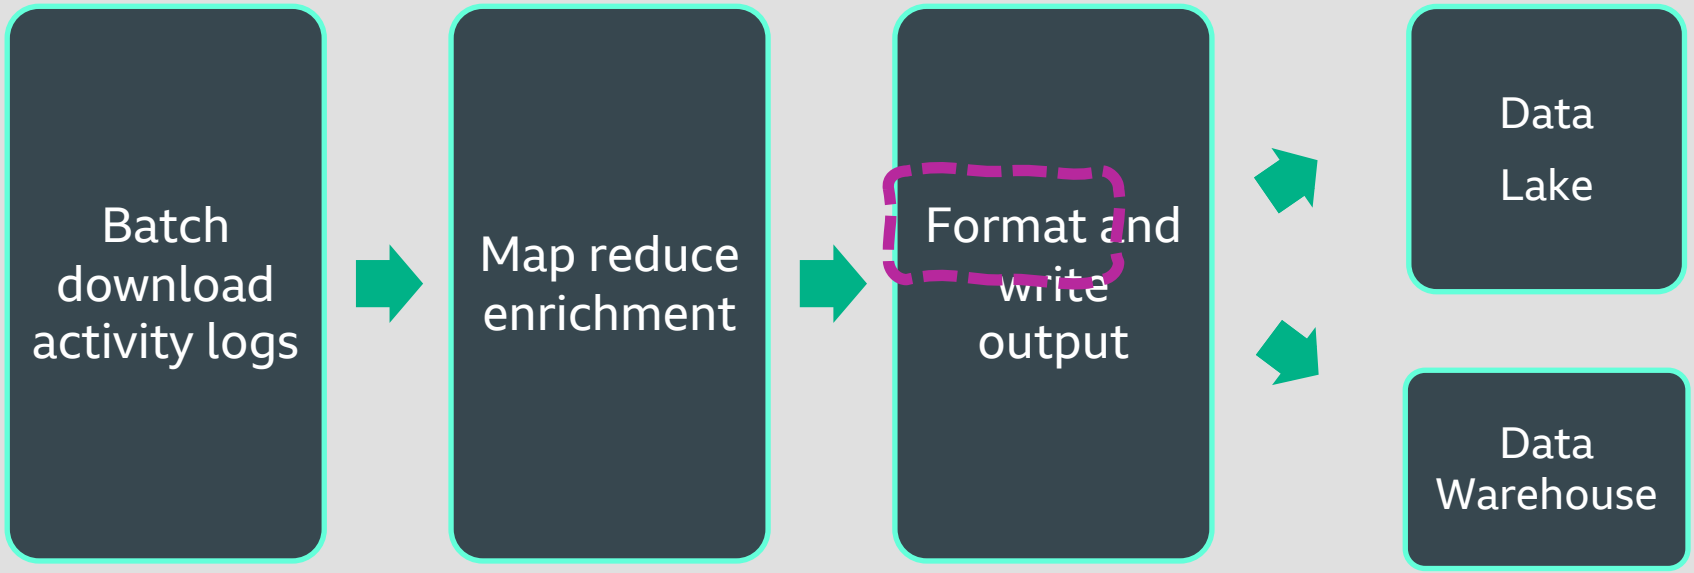

Lesson #4: Monitoring

Is every alert worth your team being interrupted?

Lesson #5: understanding our traffic patterns

Big News Days

Data volume is ever increasing

Lesson #6: Cost effective solution

Lesson #7: Getting feedback from our internal users

4. Designing a new analytics ingestion architecture

“Small, autonomous services that work together, modelled around a business domain”

Definition of microservices from “Building microservices” book by Sam Newman

Revisiting the typical data analytics ingestion flow

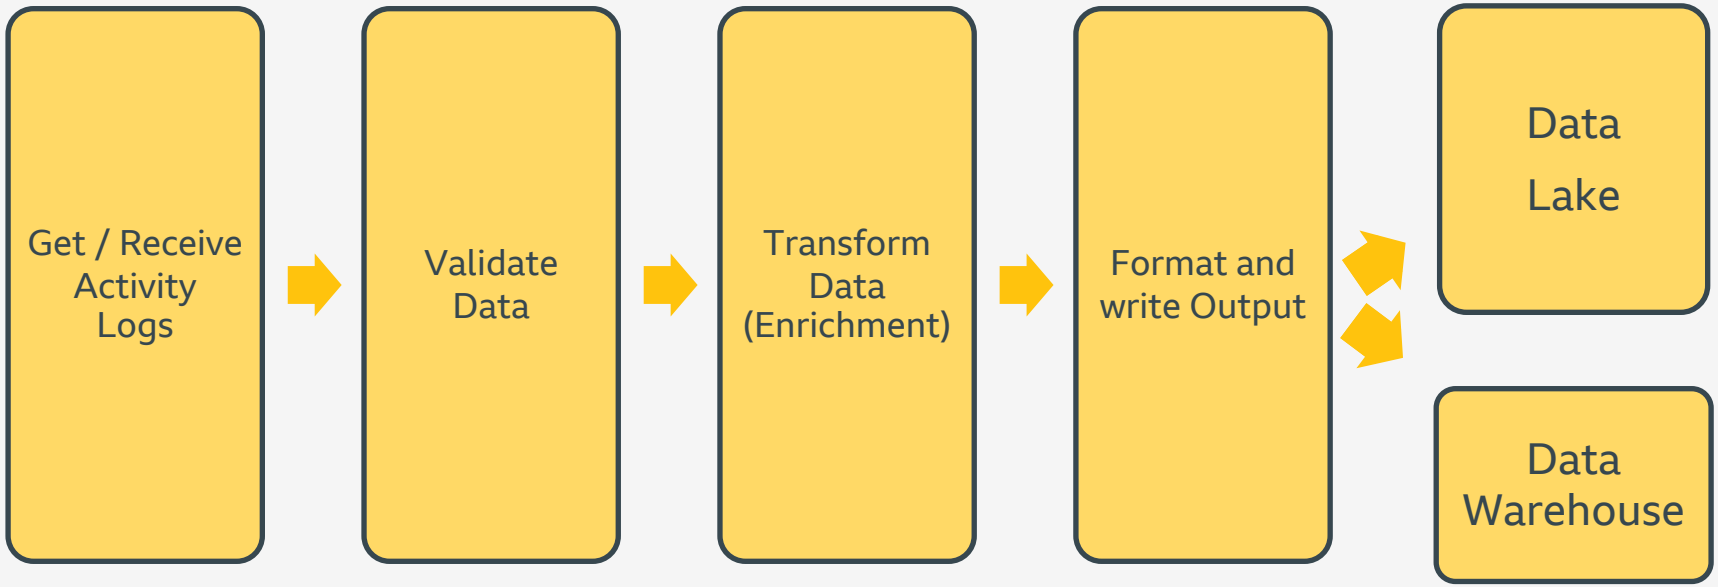

Smaller problems to solve

Squads to focus on each bounded context:

1. Receiving activity logs and keeping track of the data received
2. Model driven validation, enrichment and transformation to columnar format
3. Data aggregations (summaries) based on users' needs to make easier getting value from the data

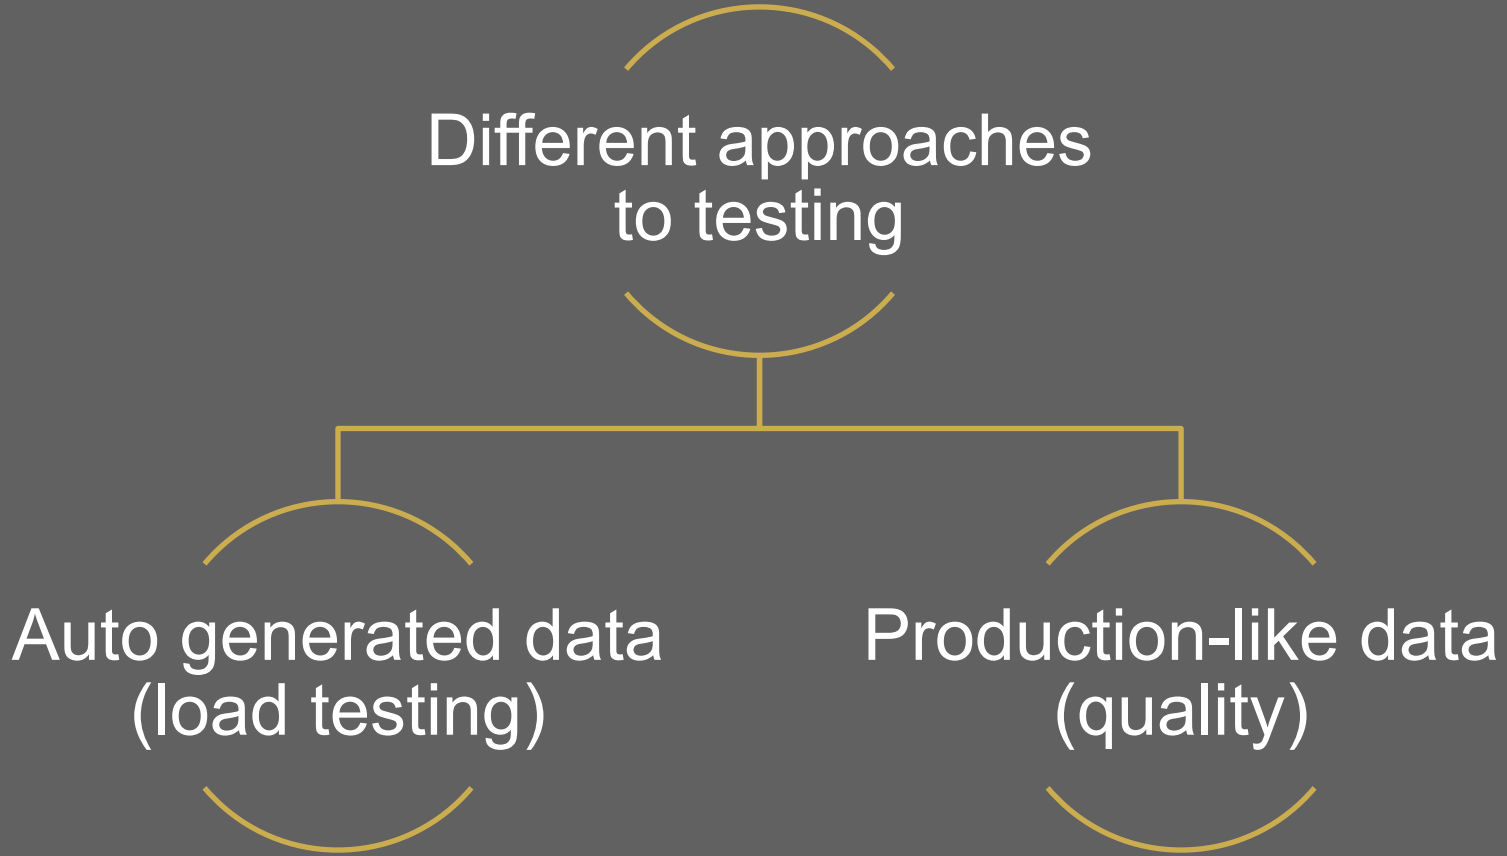
Different approaches to testing
Auto generated data (load testing)
Production-like data (quality)

New pipeline architecture

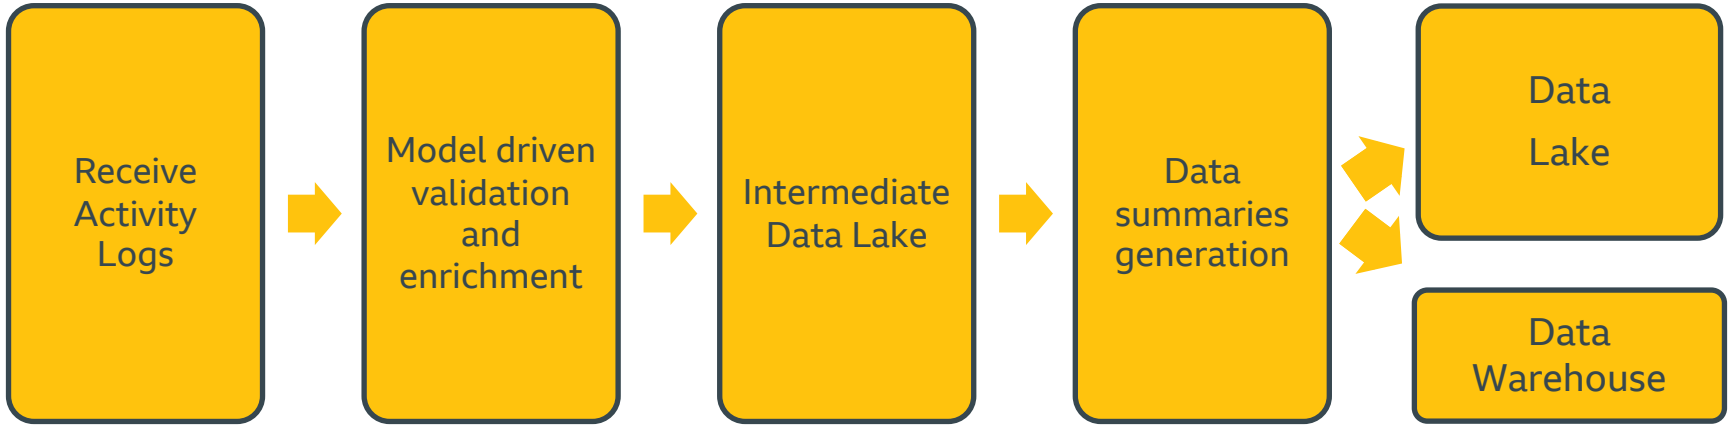

Keeping track of the data

Introduced a lineage store

New pipeline architecture

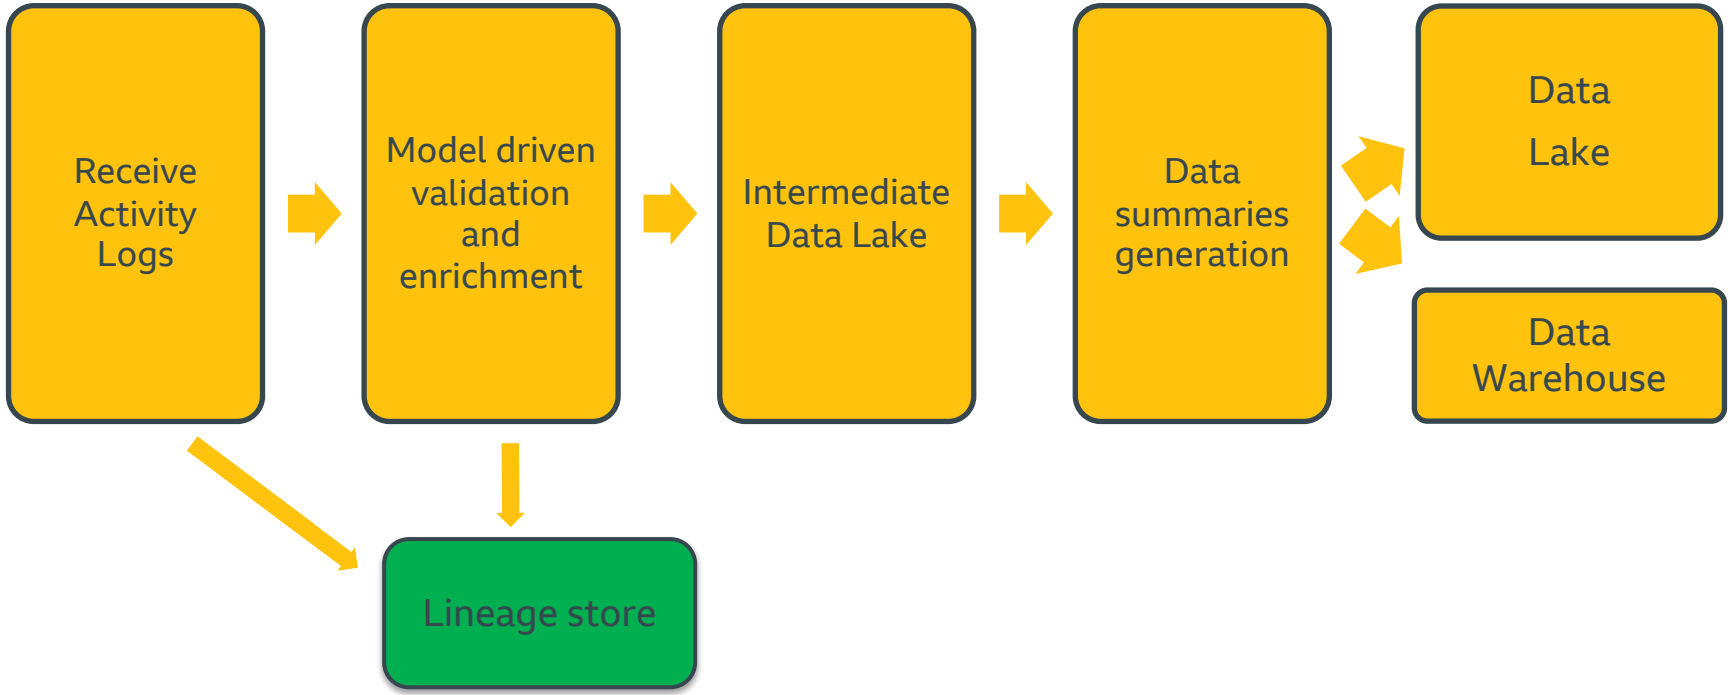

Operating the new architecture

Minimize the knowledge needed to be able to share the operational load within the team

💗 command line tools

New pipeline architecture

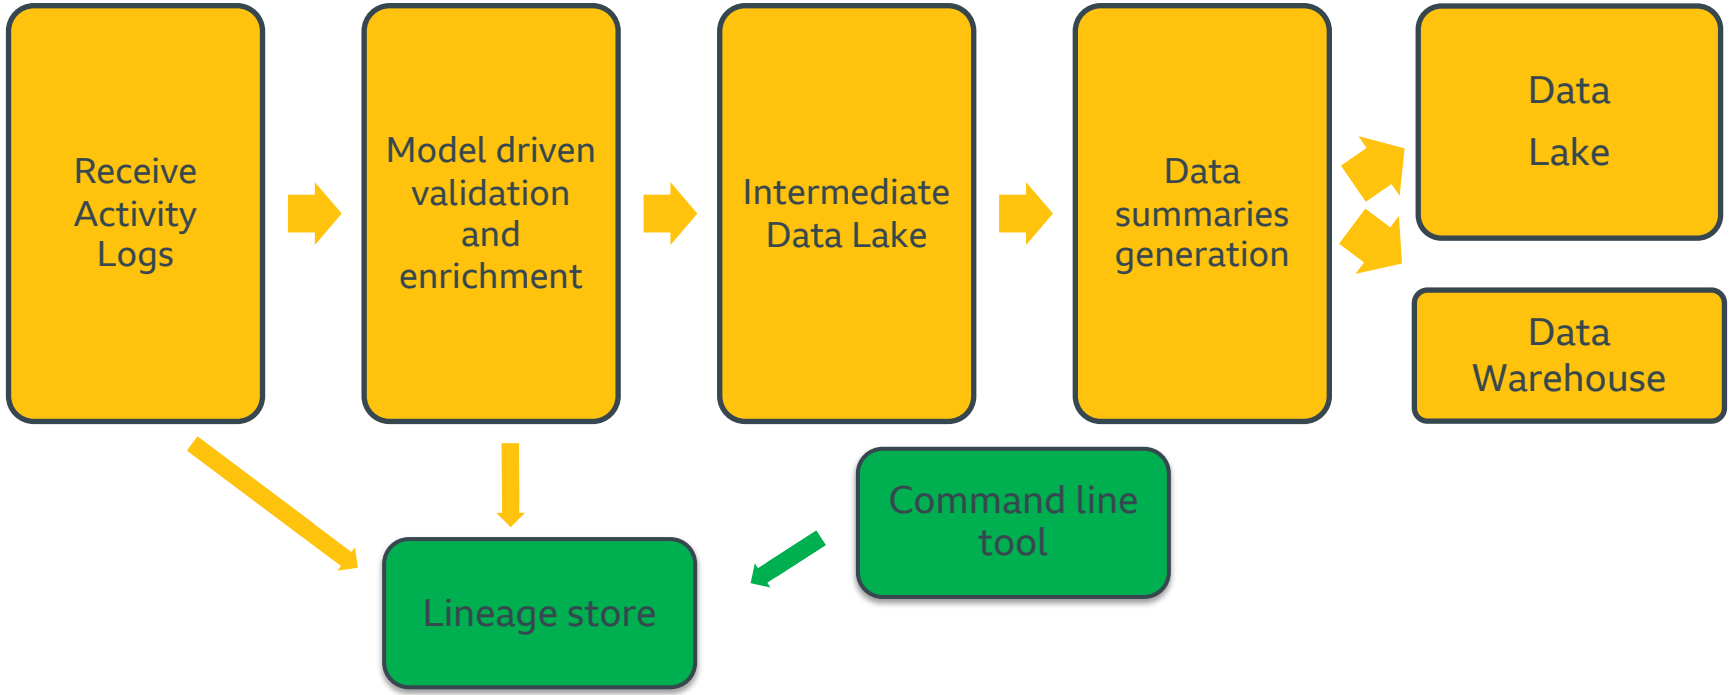

Operating the new architecture

Alerting on missing data

New pipeline architecture

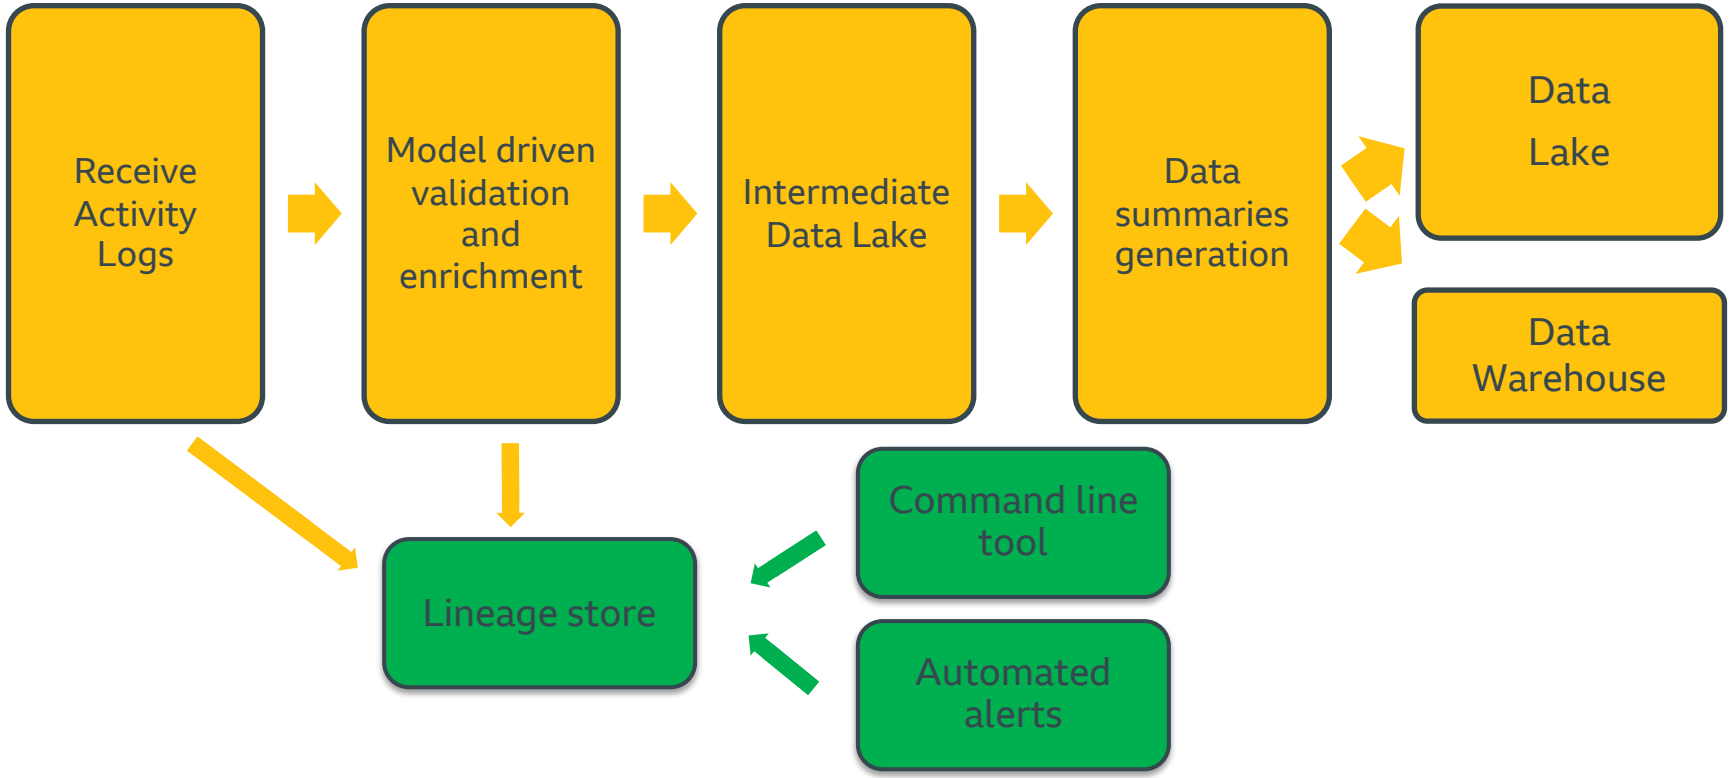

Operating the new architecture

Replaying data without stopping everything else

New pipeline architecture

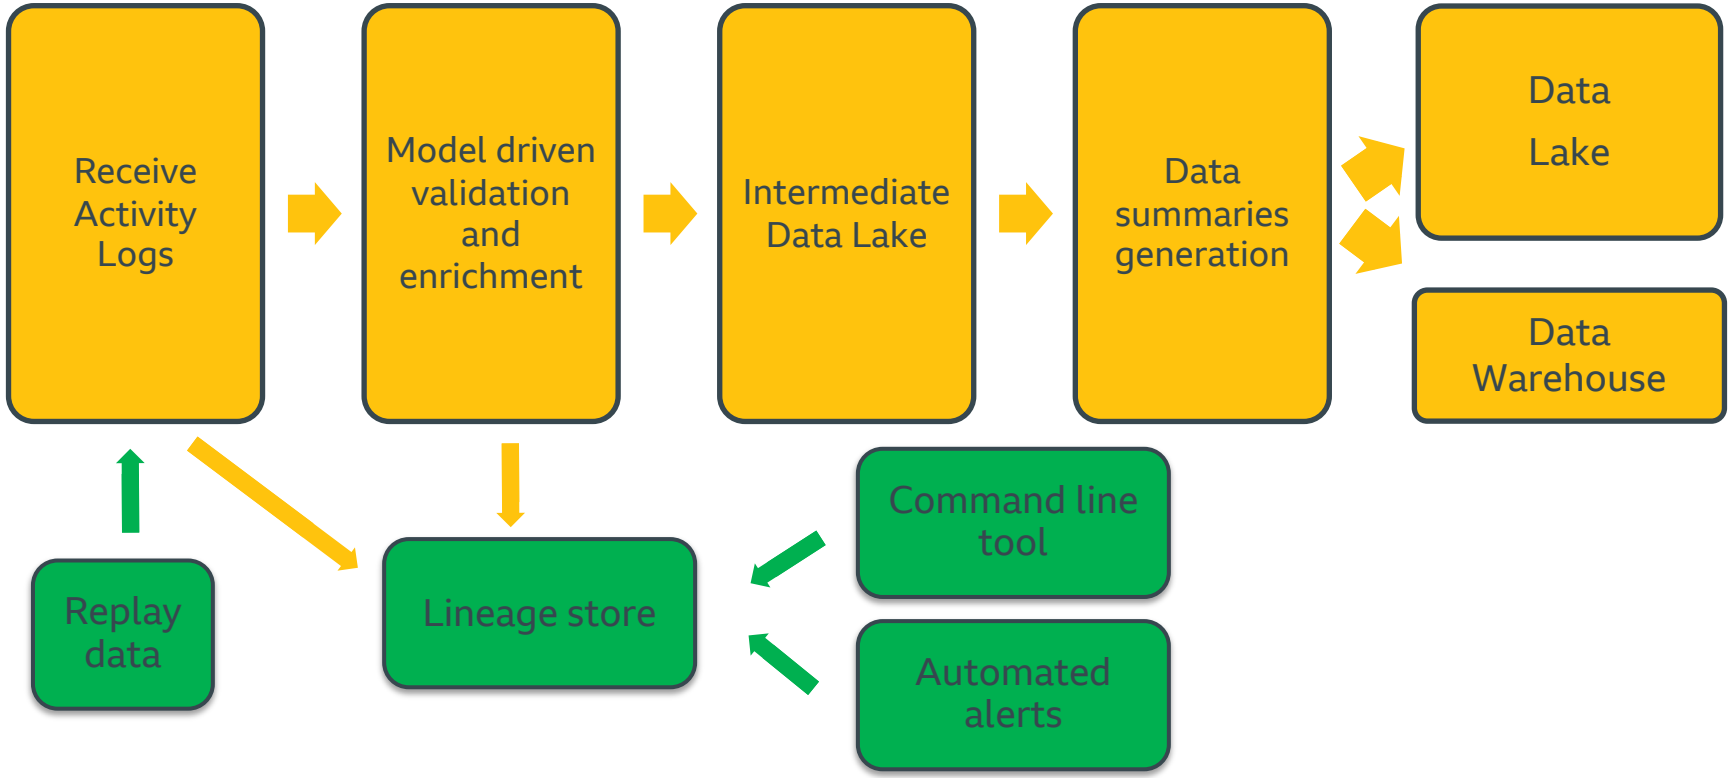

New pipeline architecture

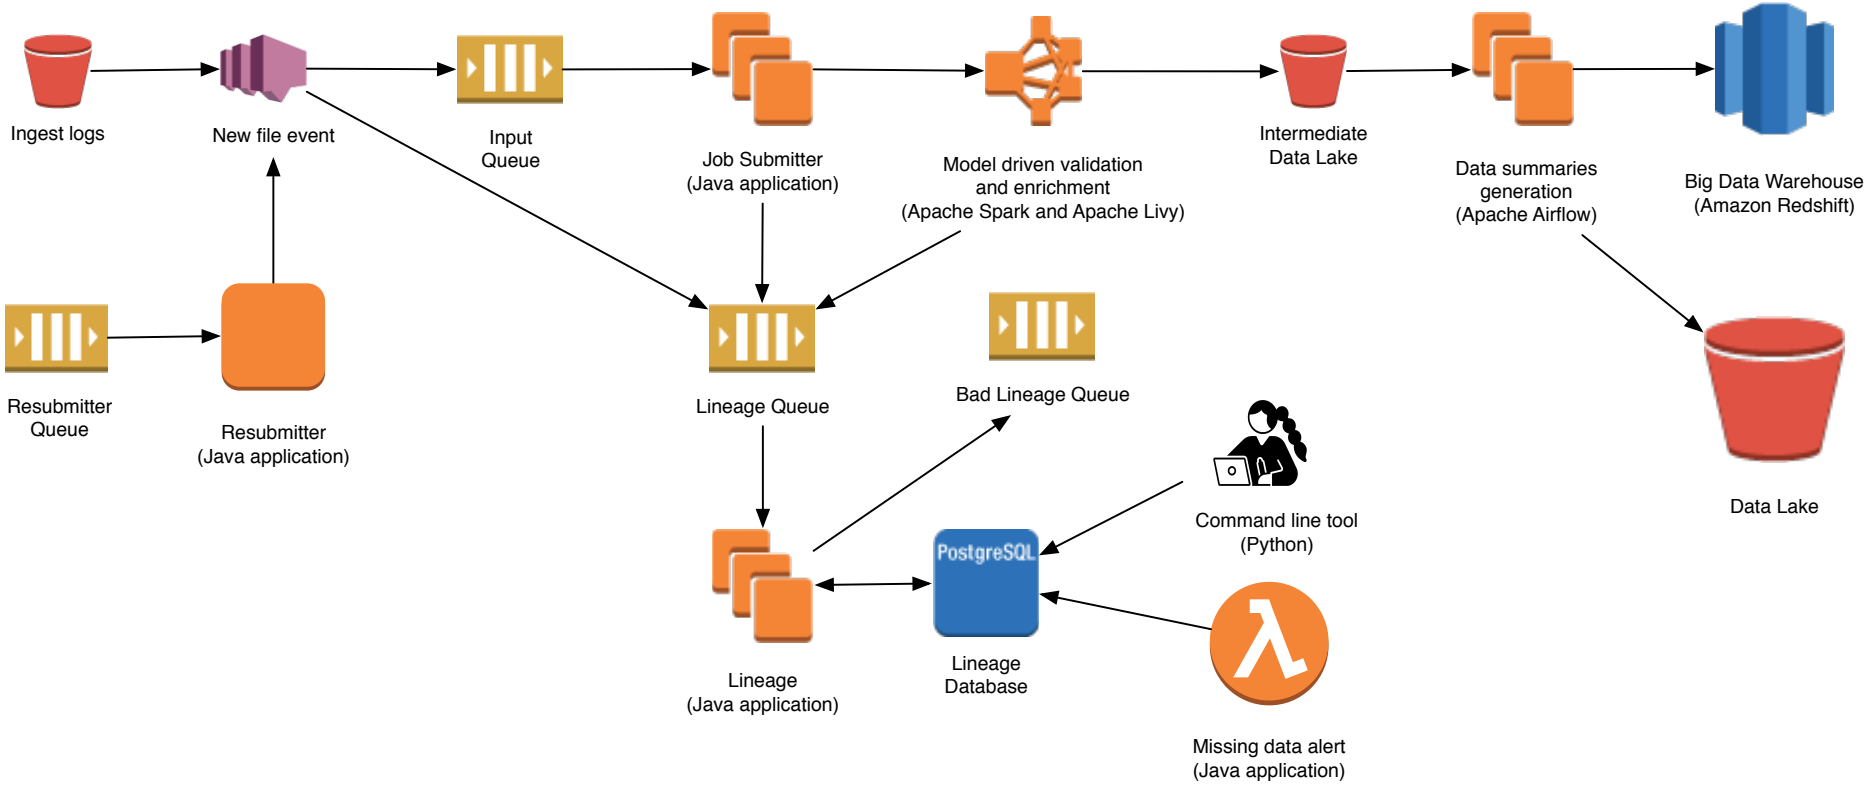

Benefit #1

Enabling team to choose fit for purpose tech and architecture

Benefit #1: choose fit for purpose tech

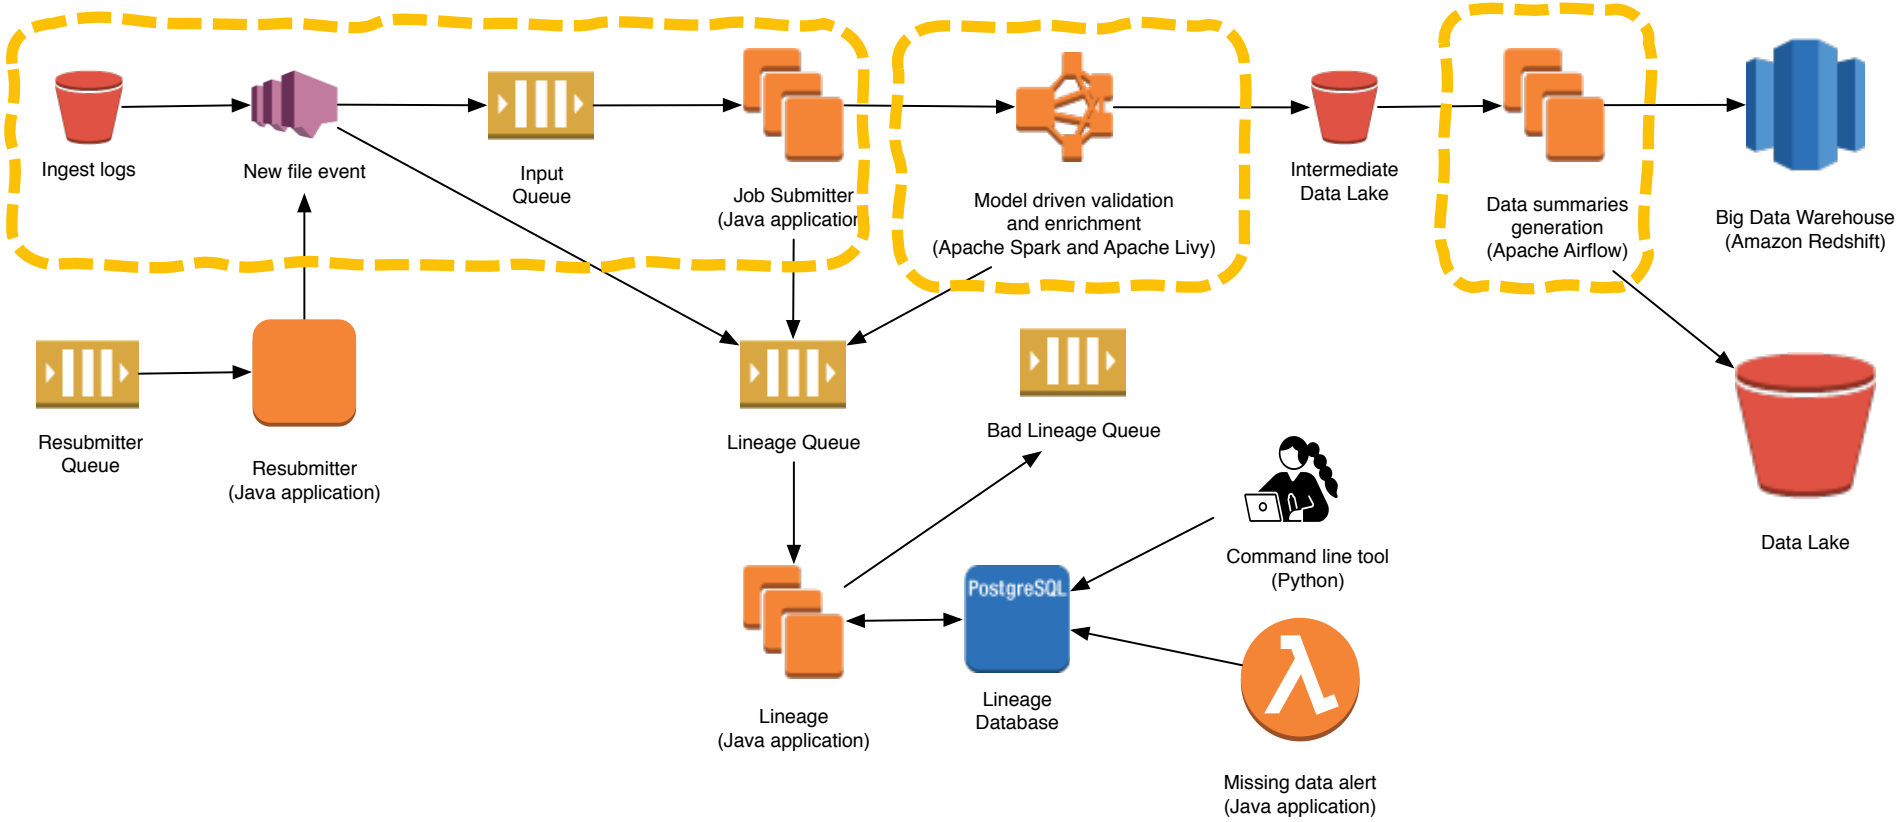

Benefit #2

Make it easier to change or replace microservices

Benefit #2: make it easier to change or replace services

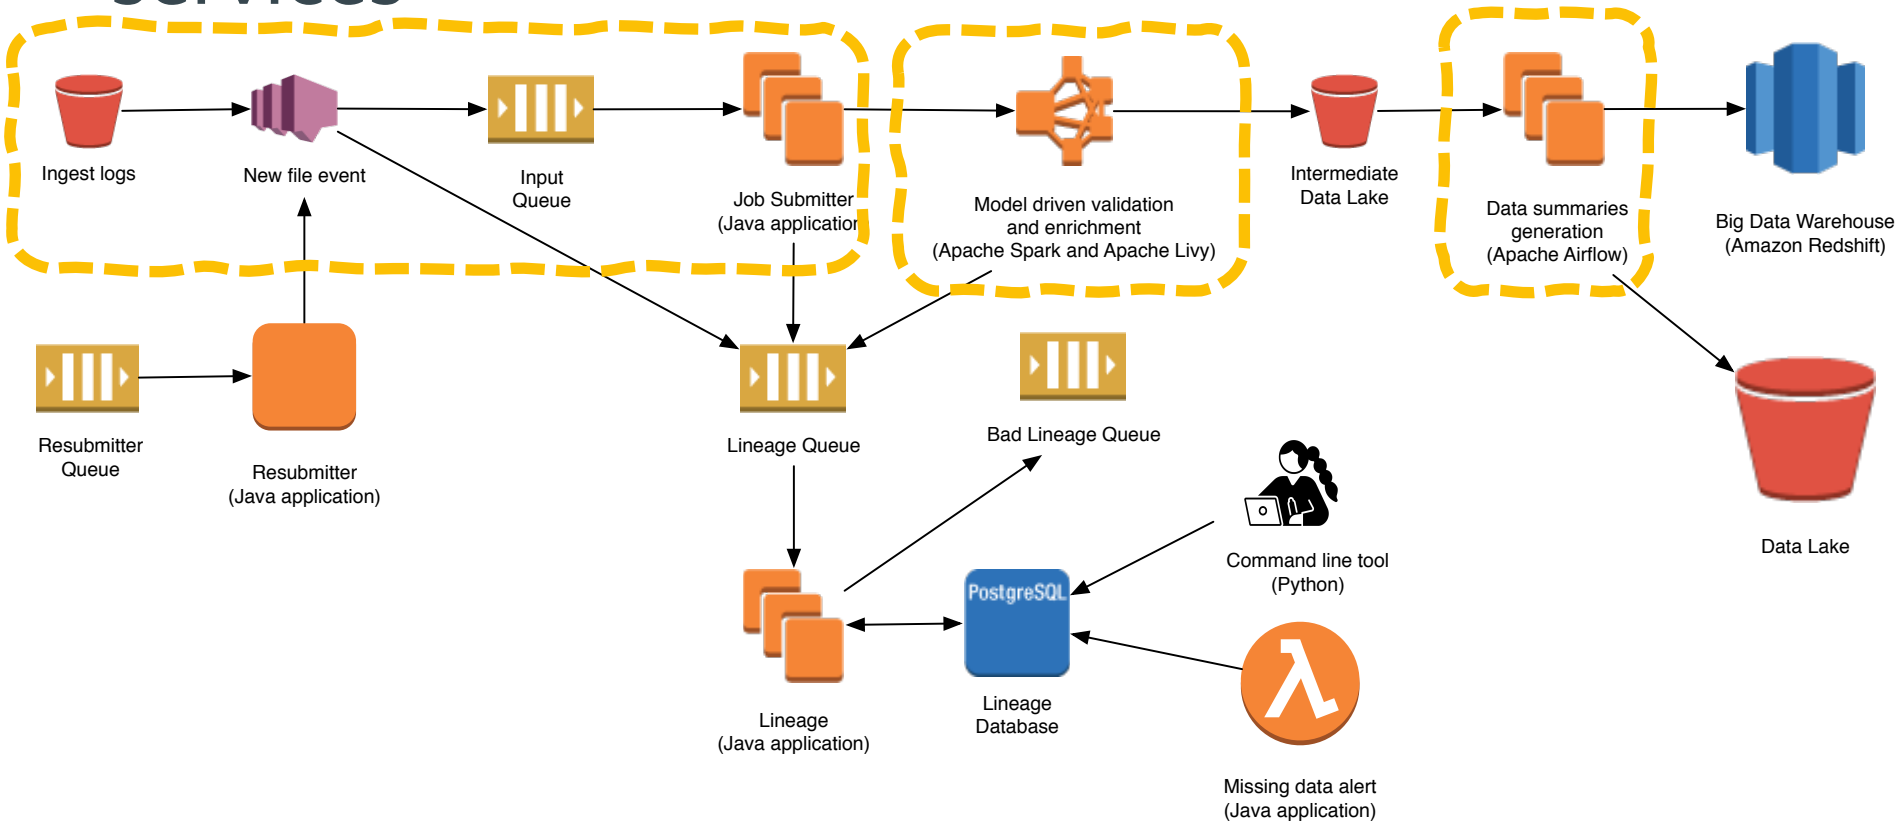

Benefit #3

Isolate failure

Benefit #3: isolate failure

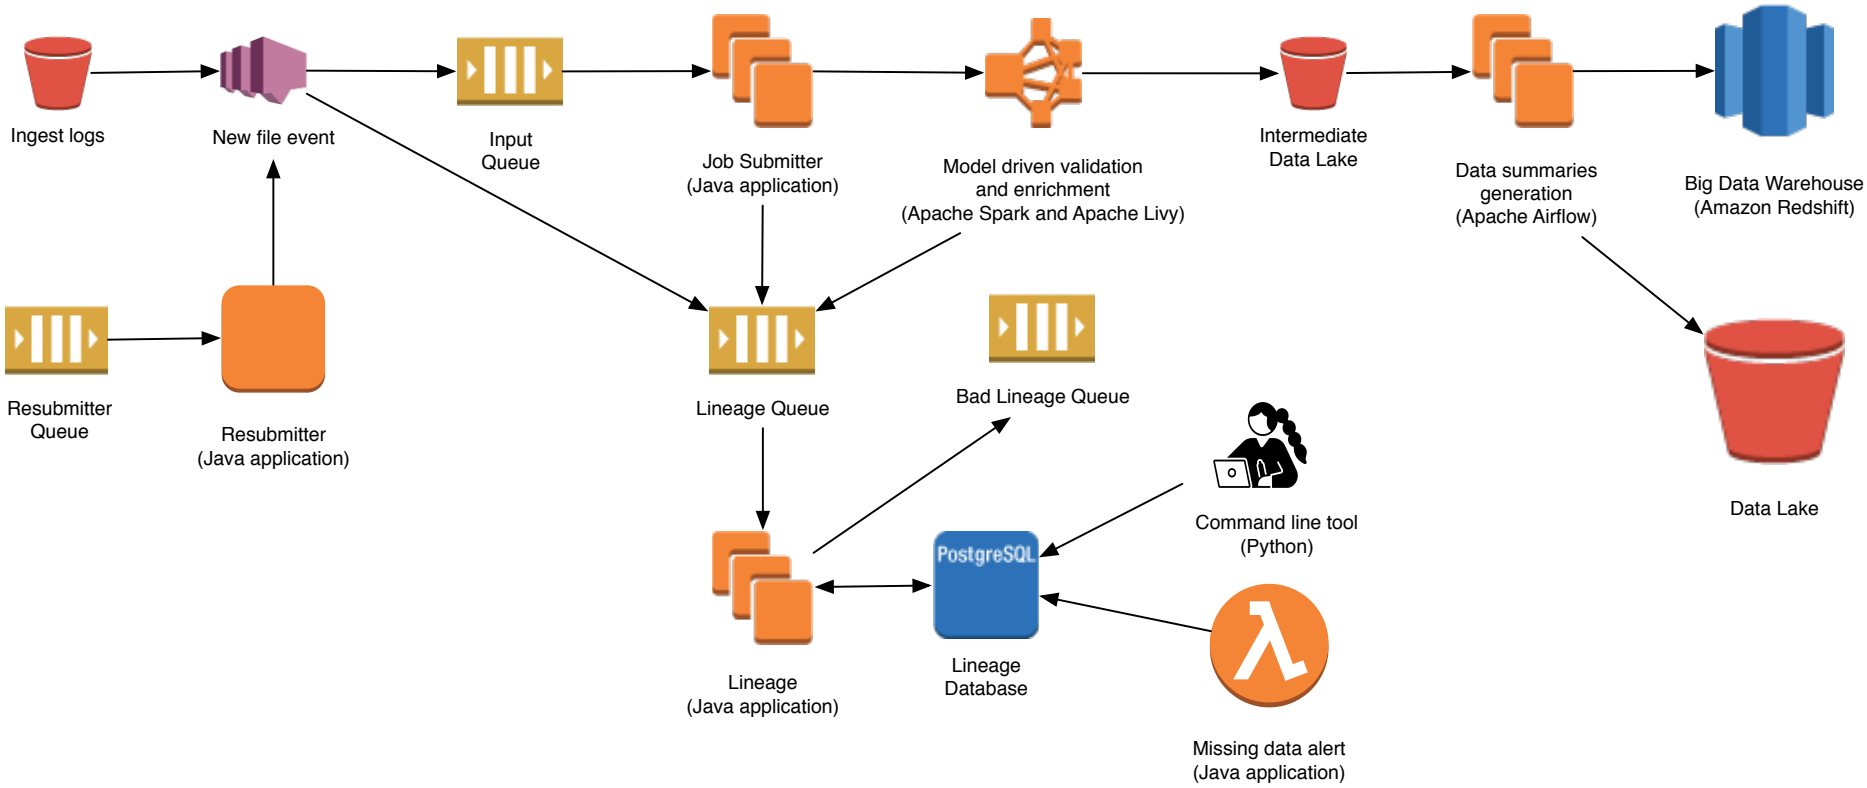

Benefit #4

Organic system growth as we operate it and learn

Benefit #4: organic system growth

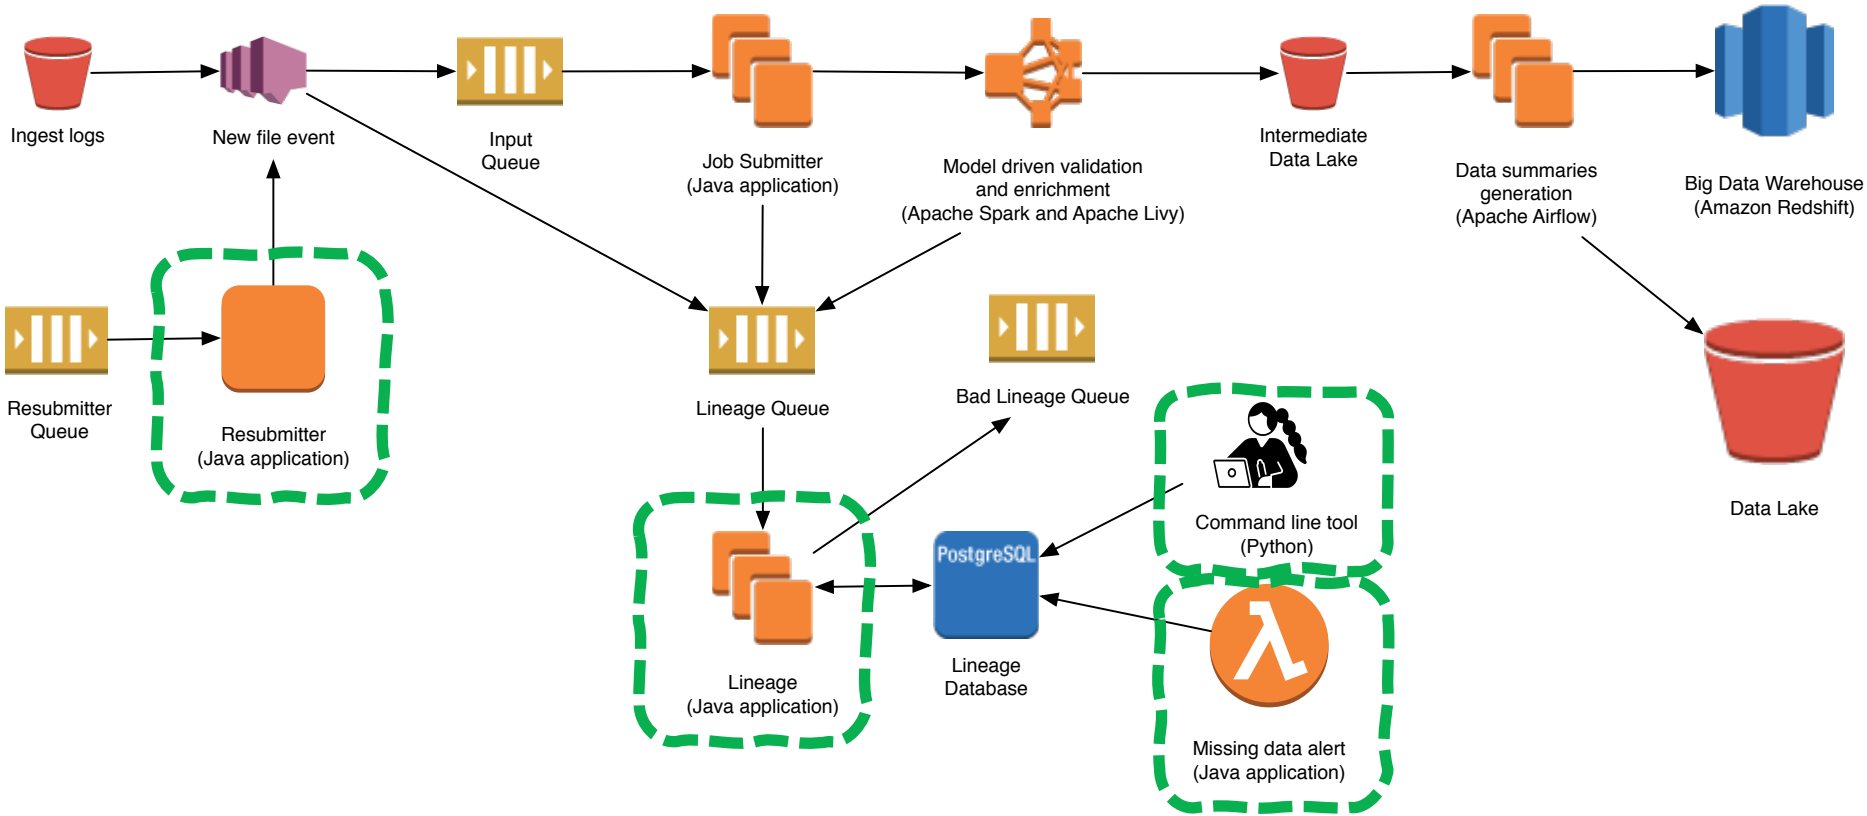

How our data analytics pipeline architecture was

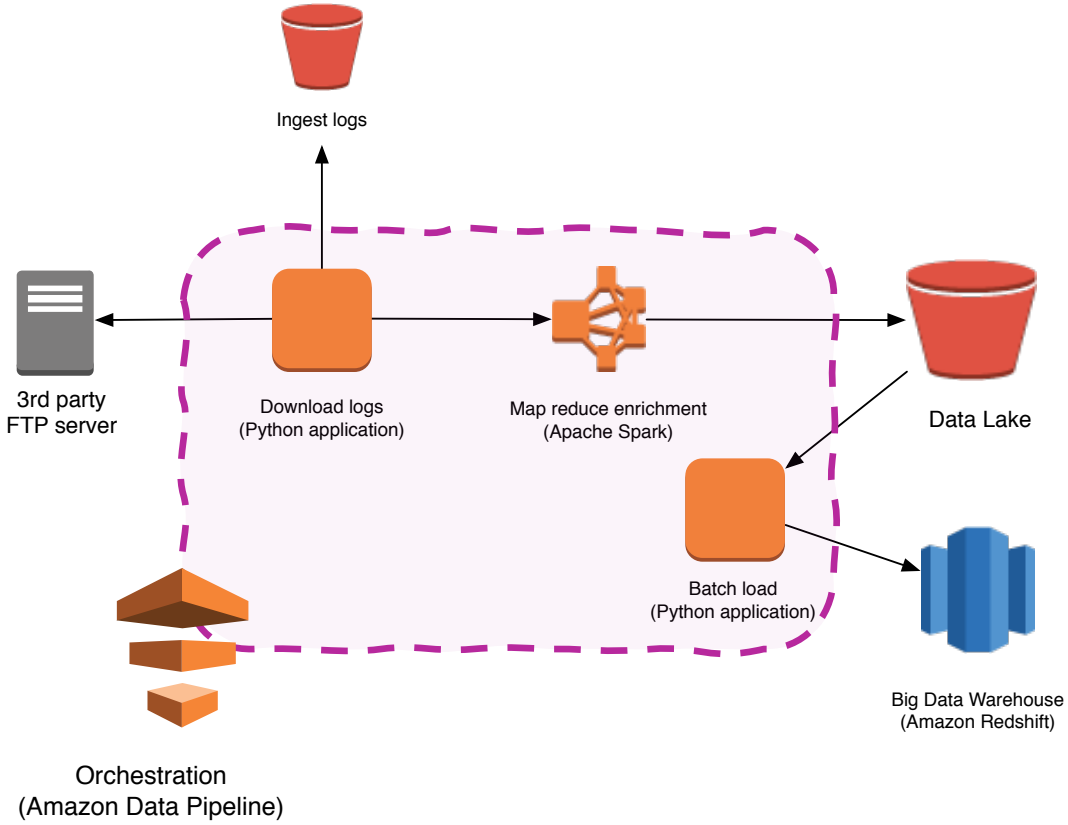

New pipeline architecture

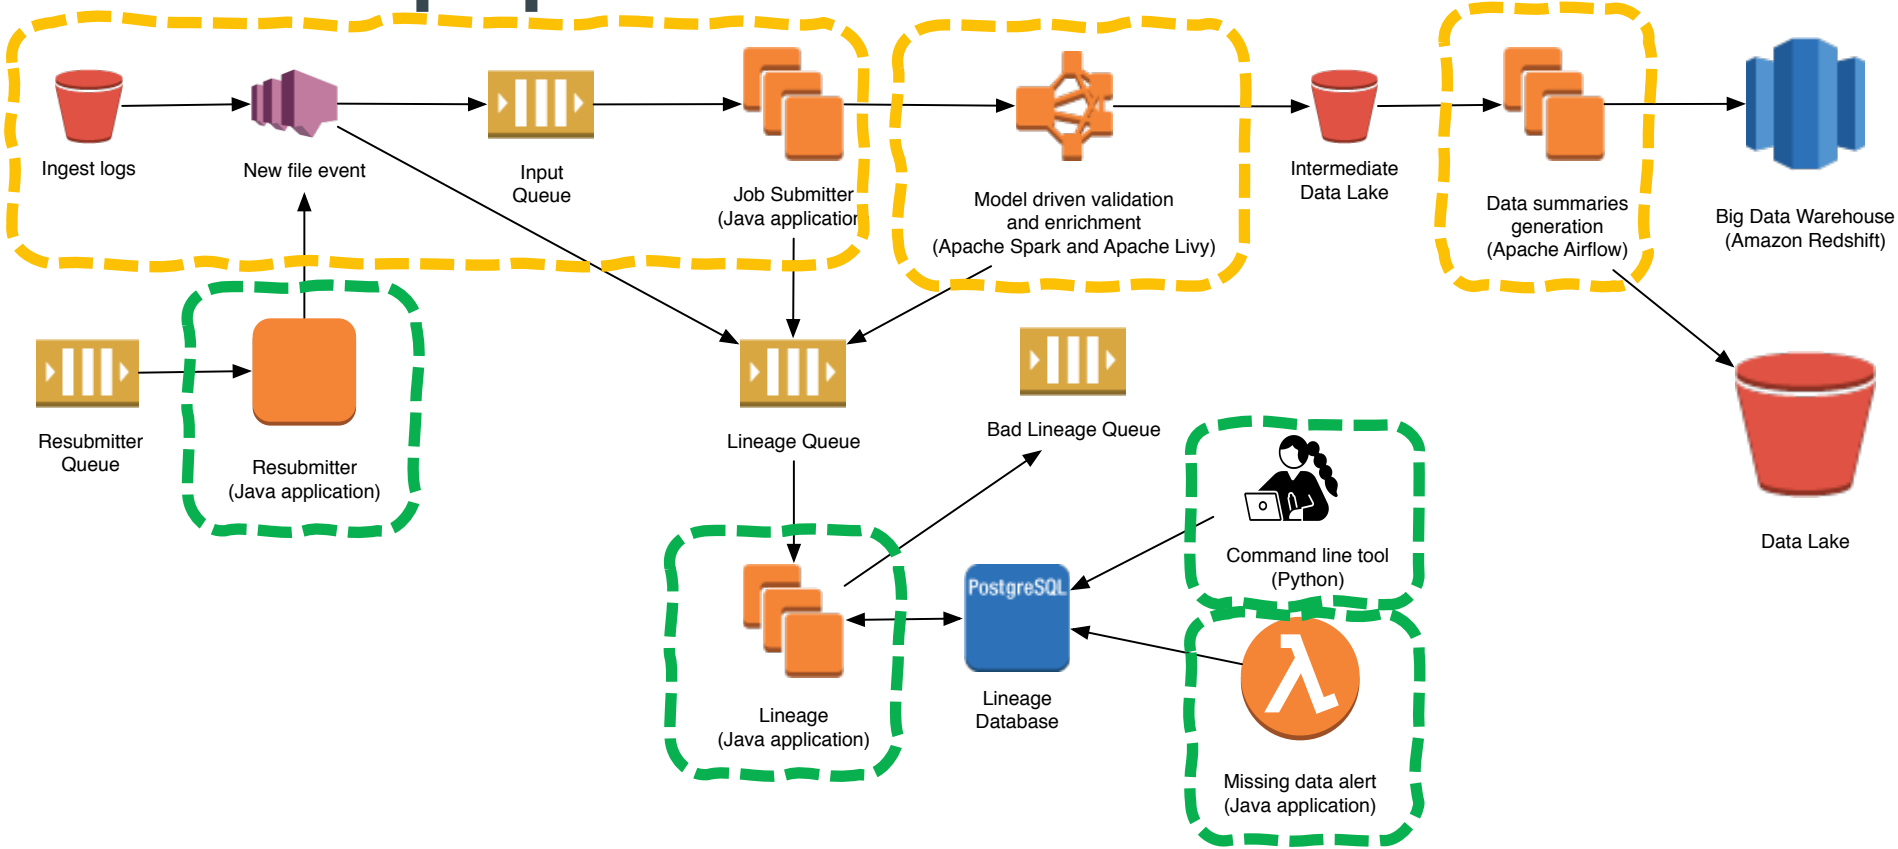

Main takeaways

Design with change in mind, you can’t predict how your traffic will evolve over time

Make sure everyone in the team can triage live issues

Choosing languages and tools which the whole team owns

5. The future of the Data Platform

Challenges as we look into the future

1. How will this architecture evolve as our data load increases?

Challenges as we look into the future

2. What are the future usage requirements for our data platform?

Challenges as we look into the future

3. How can we make it easier for our users to self serve while keeping the data secure?


A data story
KILLING EVE

Thank you!
@blanquish

BBC HD 1080 lines