

The background features a vibrant, multi-colored gradient with diagonal stripes in shades of blue, purple, orange, and yellow. A solid blue diagonal band runs from the bottom-left to the top-right.

AWS
re:Invent

DAT403 - R

Amazon DynamoDB deep dive: Advanced design patterns

Rick Houlihan
Principal Technologist, NoSQL
Amazon Web Services



© 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved.



Agenda

- Brief history of data processing (Why NoSQL?)
- Overview of Amazon DynamoDB
- NoSQL data modeling
 - Normalized versus de-normalized schema
- Common NoSQL design patterns
 - Composite keys, hierarchical data, relational data
- Modeling real applications








History of data processing

“History repeats itself because nobody was listening the first time.”


– Anonymous

Timeline of database technology

Data Pressure





Technology adoption and the hype curve



Why NoSQL?

SQL	NoSQL
Optimized for storage	Optimized for compute
Normalized/relational	De-normalized/hierarchical
Ad hoc queries	Instantiated views
Scale vertically	Scale horizontally
Good for OLAP	Built for OLTP at scale



Amazon DynamoDB



Fully managed NoSQL

Document or Wide Column

Scales to any workload




Fast and consistent

Access control

Event-driven programming

Table



Mandatory
Key-value access pattern
Determines data distribution

Optional
Model 1:N relationships
Enables rich query capabilities

All items for key
==, <, >, >=, <=
“begins with”
“between”
“contains”
“in”
sorted results
counts
top/bottom N values

Partition overloading

Use generic keys to facilitate heterogeneous partitions

Primary Key		Attributes			
PK	SK	Source	Location	URL	CustomerType
Customer_1	2019-11-29T08:31:28Z#O1	Online	US	www.amazon.com	Regular
	2019-11-29T08:31:28Z#O1#I1	ASIN	Status	Product	FCCID
		B07G6CQQYQ	PROCESSING	BOOM 3	JNZS00170
	Customer_1	Login	Email	Name	Address
		jdoe	john@example.com	John Doe	123 5th Street, New York, NY

SELECT * WHERE PK=Customer_1 AND SK > 2019-10-29


Secondary indexes

Online indexing

Support secondary access patterns


Index across all partition keys

Use composite sort keys for compound indexes



Partition/shard keys in NoSQL


Partition/shard key is used for building an unordered hash index
Allows table to be partitioned for scale



Write sharding

Salt indexed keys to support high-density aggregations on GSIs

Primary Key		Attributes			
PK	SK	Source	Location	Store	CustomerType
Customer_1	2019-11-29T08:31:28Z#O1	Online#(0-N)	US	www.amazon.com	Regular
	2019-11-29T08:31:28Z#O1#I1	ASIN	Status	Product	FCCID
		B07G6CQQY#(0-N)	PROCESSING	BOOM 3	JNZS00170
Customer_1		Login	Email	Name	Address
		jdoe	john@example.com	John Doe	123 5th Street, New York, NY



- Abstract partitioning from clients behind an API
- Write across many partitions
- Use parallel processes to increase read throughput

Index overloading

Use generic keys once more to use indexes for multiple access patterns

Primary Key		Attributes			
PK	SK	GSI1PK	GSI1SK	Store	CustomerType
Customer_1	2019-11-29T08:31:28Z#O1	Online#(0-N)	US	www.amazon.com	Regular
		GSI1PK	GSI1SK	Product	FCCID
	2019-11-29T08:31:28Z#O1#I1	B07G6CQQY#(0-N)	PROCESSING	BOOM 3	JNZS00170
		Login	Email	Name	Address
	Customer_1	jdoe	john@example.com	John Doe	123 5th Street, New York, NY

Index overloading

SELECT * WHERE PK=ONLINE#0 AND SK=US

...

SELECT * WHERE PK=ONLINE#N AND SK=US

Primary Key		Attributes				
GSI1PK	GSI1SK	PK	SK	Store	CustomerType	
Online#(0-N)	US	Customer_1	2019-11-29T08:31:28Z#O1	www.amazon.com	Regular	
B07G6CQQY#(0-N)		Customer_1	2019-11-29T08:31:28Z#O1#I1	Product	FCCID	
	PROCESSING	PK	SK	BOOM 3	JNZS00170	
		Customer_1	2019-11-29T08:31:28Z#O1#I1			

SELECT * WHERE PK=B07G6CQQY#0 AND SK=PROCESSING

...


SELECT * WHERE PK=B07G6CQQY#N AND SK=PROCESSING

Scaling NoSQL

“We are stuck with technology when what we really want is just stuff that works.”

– Douglas Adams

What bad NoSQL looks like



Getting the most out of DynamoDB throughput


“To get the most out of DynamoDB throughput, create tables where the partition key element has a large number of distinct values, and values are requested fairly uniformly, as randomly as possible.”

—*DynamoDB Developer Guide*

Space: Access is evenly spread over the key space


Time: Requests arrive evenly spaced in time

Much better picture




Auto scaling

Throughput automatically adapts to your actual traffic




Without auto scaling




With auto scaling

Performance at any scale

High request volume




Consistent low latency



Many **millions** of requests per second per table

Millisecond variance

Global-scale events: Elastic is the new normal




NoSQL data modeling


“If we have data, let's look at data. If all we have are opinions, let's go with mine.”

– Jim Barksdale

It's all about relationships




Social network



Document management



Process control




IT monitoring



Data trees

SQL vs. NoSQL design pattern



Ad hoc “joins” in SQL

```
SELECT * FROM PRODUCTS
INNER JOIN BOOKS ON
productId = productId
WHERE name = "Book Title"
```

```
SELECT * FROM PRODUCTS
INNER JOIN ALBUMS ON
productId = productId
INNER JOIN TRACKS ON
albumId = albumId
WHERE name = "Album Title"
```

```
SELECT * FROM PRODUCTS
INNER JOIN VIDEOS ON
productId = productId
INNER JOIN ACTORVIDEO ON
videoid = videoid
INNER JOIN ACTORS ON
actorId = actorId
WHERE name = "Movie Title"
```

productId	name	type	price
1	Frankenstein	Book	11.99
2	Dire Straits	Album	17.49
3	Big	Video	14.99
4	Jane Eyre	Book	10.99
5	The Dark Side of the Moon	Album	17.49
6	Saving Private Ryan	Video	18.99

bookId	productId	author	publisher	ISBN-10
1	1	Mary Shelley	Bantam	553212478
2	4	Charlotte Brontë	Wordsworth	1853260207

albumId	productId	artist	producer	releaseDate
1	2	Dire Straits	Muff Winwood	10/7/78
2	5	Pink Floyd	Pink Floyd	3/1/73

videoid	productId	writer	director	releaseDate
1	3	Ann Spielberg	Penny Marshall	6/5/88
2	6	Robert Rodat	Steven Spielberg	7/21/98

actorVideoid	videoid	actorId	character
1	1	1	Josh
2	2	1	Captain Miller
3	1	2	Susan
4	1	3	MacMillan

actorId	gender	name	birthDate
1	M	Tom Hanks	7/9/56
2	F	Elizabeth Perki	11/18/60
3	M	Robert Loggia	1/3/30

Time Complexity

$$O(\log(N) \log(D) \log(M) \log(N) \log(M) \log(M))$$

trackId	albumId	song	duration
1	1	Down to the Waterline	3:55
2	1	Water of Love	5:23
3	1	Setting Me Up	3:18
4	1	Six Blade Knife	4:10
5	1	Southbound Again	2:58
6	1	Sultans of Swing	5:47
7	1	In the Gallery	6:16
8	1	Wild West End	4:42
9	1	Lions	5:05
10	2	Speak to Me	1:13
11	2	Breathe	2:43
12	2	On the Run	3:36
13	2	Time	4:36
14	2	The Great Gig in the Sky	19:27
15	2	Money	6:23
16	2	Us and Them	7:49
17	2	Any Colour You Like	3:26
18	2	Brain Damage	3:49
19	2	Eclipse	2:03

Modeled “joins” in NoSQL

`SELECT * WHERE PK="Book Title"`

`SELECT * WHERE PK="Album Title"`

`SELECT * WHERE PK="Movie Title"`

Time Complexity

$O(1)$

Primary Key		Attributes			
PK	SK	Type	Price	Publisher	ISBN-10
Frankenstein	Mary Shelley	book	11.99	Bantam	553212478
	Dire Straits	Type	Price	Producer	ReleaseDate
		album	17.49	Muff Winwood	10/7/78
	Down to the Waterline	Duration	TrackNo		
		3:55	1		
	Water of Love	Duration	TrackNo		
		5:23	2		
	Setting Me Up	Duration	TrackNo		
		3:18	3		
	Six Blade Knife	Duration	TrackNo		
		4:10	4		
Dire Straits	Southbound Again	Duration	TrackNo		
		2:58	5		
	Sultans of Swing	Duration	TrackNo		
		5:47	6		
	In the Gallery	Duration	TrackNo		
		6:16	7		
	Wild West End	Duration	TrackNo		
		4:42	8		
	Lions	Duration	TrackNo		
		5:05	9		
Big	Penny Marshall	Type	Price	Writer	ReleaseDate
		video	14.99	Ann Spielberg	6/5/88
	Tom Hanks	Character	Gender	BirthDate	
		Josh	Male	7/9/56	
	Elizabeth Perkins	Character	Gender	BirthDate	
		Susan	Female	11/18/60	
	Robert Loggia	Character	Gender	BirthDate	
		MacMillan	Male	1/3/30	
	Tom Hanks	Gender	BirthDate	Bio	
		Male	7/9/56	{...}	

Modeled “joins” in NoSQL

SELECT * WHERE SK=“Author Name”

SELECT * WHERE SK=“Song Title”

SELECT * WHERE SK=“Actor Name”

SELECT * WHERE SK=“Director Name”

SELECT * WHERE SK=“Musician”

Swap PK and SK on index

PK	SK	Attributes			
		Type	Price	Publisher	ISBN-10
Mary Shelley	Frankenstein	book	11.99	Bantam	553212478
		Duration	TrackNo		
Sultans of Swing	Dire Straits	5:47	6		
		Sultans of Swing: The Very Best of Dire Straits		Duration	TrackNo
Tom Hanks	Big	5:50	1		
		Type	Gender	BirthDate	
Tom Hanks	Saving Private Ryan	Josh	Male	7/9/56	
		Character	Gender	BirthDate	
Penny Marshall	Big	Captain Miller	Male	7/9/56	
		Gender	BirthDate	Bio	
Dire Straits	Tom Hanks	Male	7/9/56	{...}	
		Type	Price	Writer	ReleaseDate
Dire Straits	Dire Straits	video	14.99	Ann Spielberg	6/5/88
		Type	Price	Producer	ReleaseDate
Dire Straits	Sultans of Swing: The Very Best of Dire Straits	album	17.49	Muff Winwood	10/7/78
		Type	Price	Producer	ReleaseDate
Dire Straits	Sultans of Swing: The Very Best of Dire Straits	album	25.99	Various	10/19/98

Document vs. wide column data modeling

```
{  
  _id: "john@example.com",      ← Default "_id" index supports K/V access  
  firstName: "John",  
  lastName: "Doe",  
  address: "123 A Street",  
  city: "Seattle",  
  state: "WA",  
  building: "SEA58",  
  floor: "07.650.O1"           ← Compound index on "building.floor" supports  
                             subtree aggregations for employees by  
                             location: SELECT * WHERE building ==  
                             "SEA58" AND floor startsWith("07")  
}
```

Document vs. wide column

```
{  
  _id: "john@example.com",  
  firstName: "John",  
  lastName: "Doe",  
  address: "123 A Street",  
  city: "Seattle",  
  state: "WA",  
  building: "SEA58",  
  floor: "07.650.01"
```

PK (_id)	firstName	lastName	address	city	state	GSIPK	GSISK
	}						

Indexing efficiently in NoSQL


Document	Wide column
Default index on <code>_id</code>	Partition Key defines default index
Query planner selects the index	User specifies the index
Include Shard Key or suffer	Partition Key value always required
Optimize with Compound Indexes	Use Projections to “pre-load” the index

Complex queries

“Computers are useless. They can only give you answers.”

– Pablo Picasso

Serverless & event driven architecture



Composite keys

“Hierarchies are celestial. In hell all are equal.”

– Nicolás Gómez Dávila

Multi-value sorts and filters



Partition key

Sort key

Secondary index

Opponent	Date	Gameld	Status	Host
Alice	2014-10-02	d9bl3	DONE	David
Carol	2014-10-08	o2pnB	IN_PROGRESS	Bob
Bob	2014-09-30	72f49	PENDING	Alice
Bob	2014-10-03	b932s	PENDING	Carol
Bob	2014-10-03	ef9ca	IN_PROGRESS	David

Approach 1: Query filter

```
SELECT * FROM Game  
WHERE Opponent='Bob'  
ORDER BY Date DESC  
FILTER ON Status='PENDING'
```



Secondary index

Opponent	Date	GameId	Status	Host
Alice	2014-10-02	d9bl3	DONE	David
Carol	2014-10-08	o2pnb	IN_PROGRESS	Bob
Bob	2014-09-30	72f49	PENDING	Alice
Bob	2014-10-03	b932s	PENDING	Carol
Bob	2014-10-03	ef9ca	IN_PROGRESS	David

(Filtered out)


Approach 2: Composite key

Status	Date	StatusDate
DONE	2014-10-02	DONE_2014-10-02
IN_PROGRESS	2014-10-08	IN_PROGRESS_2014-10-08
IN_PROGRESS	2014-10-03	IN_PROGRESS_2014-10-03
PENDING	2014-10-03	PENDING_2014-09-30
PENDING	2014-09-30	PENDING_2014-10-03

Approach 2: Composite key

Partition key

Sort key



Secondary index

<u>Opponent</u>	<u>StatusDate</u>		<u>Gameld</u>	<u>Host</u>
Alice	DONE_2014-10-02		d9bl3	David
Carol	IN_PROGRESS_2014-10-08		o2pnb	Bob
Bob	IN_PROGRESS_2014-10-03		ef9ca	David
Bob	PENDING_2014-09-30		72f49	Alice
Bob	PENDING_2014-10-03		b932s	Carol

Approach 2: Composite key

```
SELECT * FROM Game  
WHERE Opponent='Bob'  
    AND StatusDate BEGINS_WITH 'PENDING'
```



Secondary index


Opponent	StatusDate	GameId	Host
Alice	DONE_2014-10-02	d9bl3	David
Carol	IN_PROGRESS_2014-10-08	o2pnb	Bob
Bob	IN_PROGRESS_2014-10-03	ef9ca	David
Bob	PENDING_2014-09-30	72f49	Alice
Bob	PENDING_2014-10-03	b932s	Carol

Modeling relational data

“Dude, where’s my lookup table?”

- Anonymous Amazon SDE

Modeling complex relationships



Access Patterns	
	Get meetings
1	by date and email
2	by date and employeeID
3	by date and building/floor/room
	Load employee dashboard by email
4	Get employee data
5	Get meetings
6	Get tickets
7	Get reservations
8	Get time cards
	Get employee info
9	by employeeID
10	by email
	Get Ticket history
11	by Ticket ID
12	by employee email
13	by assignee email
	Get employees
14	by city, building, floor, aisle, desk
15	by manager
	Get assigned tickets
16	by email
	Get Tickets
17	by last touched > 24 hours
	Get project(s)
18	by status, start and target date
19	by name
	Get project history
20	by date range
21	by role
	Get Rooms
22	by buildingID
23	by availability and time range

The table

Access Patterns			Key Condition		Filter Condition
2	Get meetings	by date and employeeId	PK = employeeId, SK between(date1, date2)		duration > 0
3	Get meetings	by date and building/floor/room	PK = buildingId, SK between(date1, date2)		SK contains(building/floor/room)
9	Get employee info	by employeeID	PK = employeeId, SK startsWith("E")		
11	Get Ticket History	by Ticket ID	PK = ticketId		
19	Get project	by name	PK = projectName, SK = projectName		
20	Get project history	by date range	PK = projectName, SK between(date1, date2)		
21	Get project history	by role	PK = projectName		role = roleName
22	Get rooms	by buildingId	PK = buildingID		
23	Get rooms	by Availability and Time Range	PK = buildingId, SK between(date1, date2)		

Primary key		Attributes					
Partition key: pk	Sort key: sk						
SEA58	2019-08-20T10:00:00Z 07.106	GSI1pk	GSIsk	Duration	Attendees	Subject	
		john@example.com	2019-08-20T10:00:00Z 07.106	30	[...]	Discuss ProjectX	
	2019-08-20T10:15:00Z 07.106	Attendees	Subject	Organizer			
		[...]	Discuss ProjectX	john@example.com			
EMPLOYEE_1	2019-08-20T10:00:00Z 07.106	GSI1pk	GSIsk	Duration	Attendees	Subject	
		richard@example.com	2019-08-20T10:00:00Z 07.106	30	[...]	Discuss ProjectX	
	E#999	GSI1pk	GSIsk	GSI3pk	GSI3sk	Name	Title
		richard@example.com	E#999	SEA	58.07.105.B2	Richard Roe	IT Support
EMPLOYEE_2	E#777	GSI1pk	GSIsk	GSI3pk	GSI3sk	Name	Title
		john@example.com	E#777	SEA	58.09.203.A1	John Doe	CEO
ProjectX	2019-09-06 john@example.com	GSI1pk	GSIsk	Hours	Role		
		john@example.com	2019-09-06	12	TPM		
	2019-09-06 richard@example.com	GSI1pk	GSIsk	Hours	Role		
		richard@example.com	2019-09-06	24	SDE2		
	ProjectX	GSI1pk	GSIsk	Description	TargetDelivery		
		Active	2019-08-30	Some project	2020-01-30		
Ticket_1	2019-08-15T12:35:00Z	GSI1pk	GSIsk	Subject	GSI3pk	GSI3sk	GSI2pk
		john@example.com	2019-08-15T12:35:00Z	Badge replacement	7	2019-08-16T12:35:00Z	richard@example.com
	2019-08-15T12:35:05Z	GSI1pk	GSIsk	GSI2pk	Message		Dog ate my badge.
		john@example.com	2019-08-15T12:35:05Z	richard@example.com	Request received.		

The index schema (GSI1)

Access Patterns			Key Condition	Filter Condition
1	Get Meetings	by date and email	GSI1PK = email, GSISK between(date1, date2)	duration > 0
4	Load dashboard by email	Get employee data	GSI1PK = email, GSISK > 30 days ago	None
5		Get meetings		
6		Get tickets		
7		Get reservations		
8		Get time cards		
10	Get Employee info	by email	GSI1PK = email, GSISK startsWith("E")	
12	Get Ticket History	by employee email	GSI1PK = email	PK = ticketId
18	Get Projects	by status, start and target date	GSI2PK = status, GSISK > startDate	targetDelivery < targetDate

Primary key		Attributes								
Partition key: GSI1pk	Sort key: GSIsk	pk	sk	Subject	GSI3pk	GSI3sk	GSI2pk	Message		
staylor@abc.com	2019-08-15T12:35:00Z	pk	sk	Subject	GSI3pk	GSI3sk	GSI2pk	Message		
		Ticket_1	2019-08-15T12:35:00Z	Badge replacement	7	2019-08-16T12:35:00Z	bhana@abc.com	Dog ate my badge.		
	2019-08-15T12:35:05Z	pk	sk	GSI2pk	Message	Request received.				
		Ticket_1	2019-08-15T12:35:05Z	bhana@abc.com						
	2019-08-20T10:00:00Z 07.106	pk	sk	Duration	Attendees	Subject	Discuss ProjectX			
		SEA58	2019-08-20T10:00:00Z 07.106	30	[...]					
bhana@abc.com	2019-09-06	pk	sk	Hours	Role	Discuss ProjectX				
		ProjectX	2019-09-06 staylor@abc.com	12	TPM					
	E#777	pk	sk	GSI3pk	GSI3sk	Name	Title	GSI2pk		
		EMPLOYEE_2	E#777	SEA	58.09.203.A1	Steven Taylor	CEO	staylor@abc.com		
	2019-08-20T10:00:00Z 07.106	pk	sk	Duration	Attendees	Subject	Discuss ProjectX			
		EMPLOYEE_1	2019-08-20T10:00:00Z 07.106	30	[...]					
Active	2019-09-06	pk	sk	Hours	Role	SDE2				
		ProjectX	2019-09-06 bhana@abc.com	24						
	E#999	pk	sk	GSI3pk	GSI3sk	Name	Title	GSI2pk		
		EMPLOYEE_1	E#999	SEA	58.07.105.B2	Benny Hana	IT Support	staylor@abc.com		
Active	2019-08-30	pk	sk	Description	TargetDelivery	Some project				
		ProjectX	ProjectX	Some project	2020-01-30					

The index schema (GSI2)

Access Patterns			Key Condition	Filter Condition
13	Get Ticket History		by assignee email	GSI2PK = email
15	Get employees		by manager	GSI2PK = email, SK > 3

Primary key		Attributes							
Partition key: GSI2pk	Sort key: GSIsk	pk	sk	GSI1pk	GSI3pk	GSI3sk	Name	Title	
john@example.com	E#777	pk	sk	GSI1pk	GSI3pk	GSI3sk	Name	Title	
		EMPLOYEE_2	E#777	john@example.com	SEA	58.09.203.A1	John Doe	CEO	
	E#999	pk	sk	GSI1pk	GSI3pk	GSI3sk	Name	Title	
		EMPLOYEE_1	E#999	richard@example.com	SEA	58.07.105.B2	Richard Roe	IT Support	
richard@example.com	2019-08-15T12:35:00Z	pk	sk	GSI1pk	Subject	GSI3pk	GSI3sk	Message	
		Ticket_1	2019-08-15T12:35:00Z	john@example.com	Badge replacement	7	2019-08-16T12:35:00Z	Dog ate my badge.	
	2019-08-15T12:35:05Z	pk	sk	GSI1pk	Message				
		Ticket_1	2019-08-15T12:35:05Z	john@example.com	Request received.				

The index schema (GSI3)

Access Patterns			Key Condition	Filter Condition
14	Get employees	by city, building, floor, aisle, desk	GSI3PK = city, GSI3SK startsWith(building/floor/aisle/desk)	
17	Get tickets	by last touched > 24 hours	GSI3PK = (0-N), GSI3SK < yesterday	

Primary key		Attributes						
Partition key: GSI3pk	Sort key: GSI3sk	pk	sk	GSI1pk	GSIsk	Name	Title	GSI2pk
SEA	58.07.105.B2	pk	sk	GSI1pk	GSIsk	Name	Title	GSI2pk
		EMPLOYEE_1	E#999	richard@example.com	E#999	Richard Roe	IT Support	john@example.com
	58.09.203.A1	pk	sk	GSI1pk	GSIsk	Name	Title	GSI2pk
		EMPLOYEE_2	E#777	john@example.com	E#777	John Doe	CEO	john@example.com
7	2019-08-16T12:35:00Z	pk	sk	GSI1pk	GSIsk	Subject	GSI2pk	Message
		Ticket_1	2019-08-15T12:35:00Z	john@example.com	2019-08-15T12:35:00Z	Badge replacement	richard@example.com	Dog ate my badge.

The final result

Access Patterns		Table/Index	Key Condition	Filter Condition
Get meetings				
1	by date and email	GSI1	GSI1PK = email, GSISK between(date1, date2)	duration > 0
2	by date and employeeld	Table	PK = employeeld, SK between(date1, date2)	duration > 0
3	by date and building/floor/room	Table	PK = buildingId, SK between(date1, date2)	SK contains(building/floor/room)
Load employee dashboard by email				
4	Get employee data	GSI1	GSI1PK = email, GSISK > 30 days ago	None
5	Get meetings			
6	Get tickets			
7	Get reservations			
8	Get time cards			
Get employee info				
9	by employeeID	Table	PK = employeeld, SK startsWith("E")	
10	by email	GSI1	GSI1PK = email, GSISK startsWith("E")	
Get Ticket history				
11	by Ticket ID	Table	PK = ticketId	
12	by employee email	GSI1	GSI1PK = email	PK = ticketId
13	by assignee email	GSI2	GSI2PK = email	PK = ticketId
Get employees				
14	by city, building, floor, aisle, desk	GSI3	GSI3PK = city, GSI3SK startsWith(building/floor/aisle/desk)	
15	by manager	GSI2	GSI2PK = email, SK > 3	
Get assigned tickets				
16	by email	GSI2	GSI1PK = email	PK = ticketId
Get Tickets				
17	by last touched > 24 hours	GSI3	GSI3PK = (0-N), GSI3SK < yesterday	
Get project(s)				
18	by status, start and target date	GSI1	GSI2PK = status, GSISK > startDate	targetDelivery < targetDate
19	by name	Table	PK = projectName, SK = projectName	
Get project history				
20	by date range	Table	PK = projectName, SK between(date1, date2)	
21	by role	Table	PK = projectName	role = roleName
Get Rooms				
22	by buildingId	Table	PK = buildingID	
23	by Availability and Time Range	Table	PK = buildingId, SK between(date1, date2)	

Design for common patterns

“To understand is to perceive patterns.”

– Isaiah Berlin

Access patterns matter

Primary Key		Attributes				
PK	SK
Client1	Quote1_v1	200+ Attributes (50KB avg)				
	Quote1_v2
	Quote1_v3	200+ Attributes (50KB avg)				
	Quote1_v4	200+ Attributes (50KB avg)				
	Quote1_v5	200+ Attributes (50KB avg)				


- Insurance quote service
- Store all versions
- 200+ attributes per quote
- 50KB average record size
- 800 quotes-per-minute peak
- 1K WCU provisioned

Optimized for writes

- Version items as categories are updated
- Send all versions when queried
- Process with client-side logic
- 50 WCU provisioned

Primary Key		Attributes			
PK	SK	type	date	status	price
Client1	Quote1_v1_Toplevel	type	date	status	price
	
	Quote1_v1_Mileage	mileage	carType	...	priceAdj
		5000	1
	Quote1_v2_Mileage	mileage	carType	...	priceAdj
		25000	1.25

NoSQL Workbench for DynamoDB



- Use the tool designed by and for the AWS specialist SA team
- Model your data, visualize your designs, generate your code
- <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/workbench.html>

Conclusions

- NoSQL does not mean non-relational
- The ERD still matters
- RDBMS is not deprecated by NoSQL
- Use NoSQL for OLTP or DSS at scale
- Use RDBMS for OLAP

Learn databases with AWS Training and Certification

Resources created by the experts at AWS to help you build and validate database skills



25+ free digital training courses cover topics and services related to databases, including:

- Amazon Aurora
- Amazon Neptune
- Amazon DocumentDB
- Amazon DynamoDB
- Amazon ElastiCache
- Amazon Redshift
- Amazon RDS



Validate expertise with the new **AWS Certified Database - Specialty** beta exam

[Visit aws.training](https://aws.training)

Thank you!



Please complete the session
survey in the mobile app.