


# Transaction Properties: ACID vs. BASE



# Client/Server with Transaction Processing

- ◆ *Transactions are a way to make ACID operations a general commodity*  
*[Transaction Processing Concepts and Techniques, Jim Gray and Andreas Reuter, 1993]*

## ☐ **Atomicity**

- ☐ a transaction is an indivisible unit of work
- ☐ an all-or-nothing proposition
- ☐ all updates to a database, displays on the clients' screens, message queues
- ☐ e.g., salary increase for all 1 million employees or none

## ☐ **Consistency**

- ☐ a transaction is an indivisible unit of work
- ☐  $S \rightarrow [T \mid \text{abort}] \rightarrow S$
- ☐ integrity constraints (e.g.,  $\text{mgr.salaray} > \text{salary}$ )

## ☐ **Isolation**

- ☐ a transaction's behavior not affected by other transactions running concurrently
- ☐ e.g., reserve a seat
- ☐ serialization techniques

## ☐ **Durability**

- ☐ persistence
- ☐ a transaction's effects are permanent after it commits.

*ACID is like motherhood and apple pie. It's necessary- and you can't have too much of it. OK, enough preaching.*

# ACID Properties: Failures

- Isolation Failure:

$T_1$  transfers 10 from A to B.  $T_2$  transfers 10 from B to A. Combined, there are four actions:

- $T_1$  subtracts 10 from A.
- $T_1$  adds 10 to B.
- $T_2$  subtracts 10 from B.
- $T_2$  adds 10 to A.

By interleaving the transactions, the actual order of actions might be:

- $T_1$  subtracts 10 from A.
- $T_2$  subtracts 10 from B.
- $T_2$  adds 10 to A.
- $T_1$  adds 10 to B.

Now, consider what happens if  $T_1$  fails halfway through.

# ACID Properties:

## Failures

- Durability Failure:

$T_1$  transfers 10 from A to B.

- It removes 10 from A. It then adds 10 to B.
- At this point, a "success" message is sent to the user.
- However, the changes are still queued in the [disk buffer](#) waiting to be committed to the disk.
- Power fails and the changes are lost.
- The user assumes (understandably) that the changes have been made.

# CAP Theorem (Brewer's Theorem)

- it is impossible for a [distributed computer system](#) to simultaneously provide all three of the following guarantees:
  - [Consistency](#): all nodes see the same data at the same time
  - [Availability](#): Node failures do not prevent other survivors from continuing to operate (a guarantee that every request receives a response about whether it succeeded or failed)
  - [Partition tolerance](#): the system continues to operate despite arbitrary partitioning due to network failures (e.g., message loss)
- A distributed system can satisfy any two of these guarantees at the same time but not all three.

# CAP Theorem (Brewer's Theorem):

## So, what can be done?

- In a distributed system, a network (of networks) is inevitable (by definition).
- Failures can, and will, occur to a networked system -> partitioned tolerance should be accommodated.
- Then, the only option left is choosing between Consistency and Availability- i.e., CA doesn't make any sense (except when we have, e.g., a single-site databases; 2-phase commit, cache validation protocols)
- Not necessarily in a mutually exclusive manner, but possibly by partial accommodation of both -> trade-off analyses important

**[In terms of architecture, what ontological concept is this about?]**

- AP: A partitioned node returns
  - a correct value, if in a consistent state;
  - a timeout error or an error, otherwise
  - e.g., DynamoDB, CouchDB, and Cassandra
- CP: A partitioned node returns the most recent version of the data, which could be stale.
  - e.g., MongoDB, Redis, AppFabric Caching, and MemcacheDB

# BASE

(Basically Available, Soft-State, Eventually Consistent)

- **SQL databases:**

- Structured query language for
- Traditional relational databases (unique keys, single valued, no update/insertion/deletion anomalies)
- Well structured data
- ACID properties should hold

- **NoSQL (Not only SQL) databases:**

- triggered by the storage needs of [Web 2.0](#) companies such as [Facebook](#), [Google](#) and [Amazon.com](#)
- Not necessarily well structured – e.g., pictures, documents, web page description, video clips, etc.
- Lately of increasing importance due to big data
- ACID properties may not hold -> no properties at all then???
- focuses on availability of data even in the presence of multiple failures
- spread data across many storage systems with a high degree of replication.

# BASE

(Basically Available, Soft-State, Eventually Consistent)

- **Rationale:**

- It's ok to use stale data (Accounting systems do this all the time. It's called "closing out the books.") ; it's ok to give approximate answers
- Use resource versioning -> say what the data really is about – no more, no less.

The value of x is 5, at time T and date D

- So, shift the PH from 0-6 (acidic) to 8-14 (basic) – pure water's PH is 7 and neutral

- **Can some compromise be made between C and A?:**

- instead of completely giving up on C, for A
- Instead of completely giving up on A, instead of C




# BASE

(Basically Available, Soft-State, Eventually Consistent)

- **Basic Availability:** fulfill request, even in partial consistency.
- **Soft State:** abandon the consistency requirements of the ACID model pretty much completely
- **Eventual Consistency:** at some point in the future, data will converge to a consistent state; delayed consistency, as opposed to immediate consistency of the ACID properties.
  - purely a [liveness](#) guarantee (reads eventually return the requested value); but
  - does not make [safety](#) guarantees, i.e.,
  - an eventually consistent system can return any value before it converges

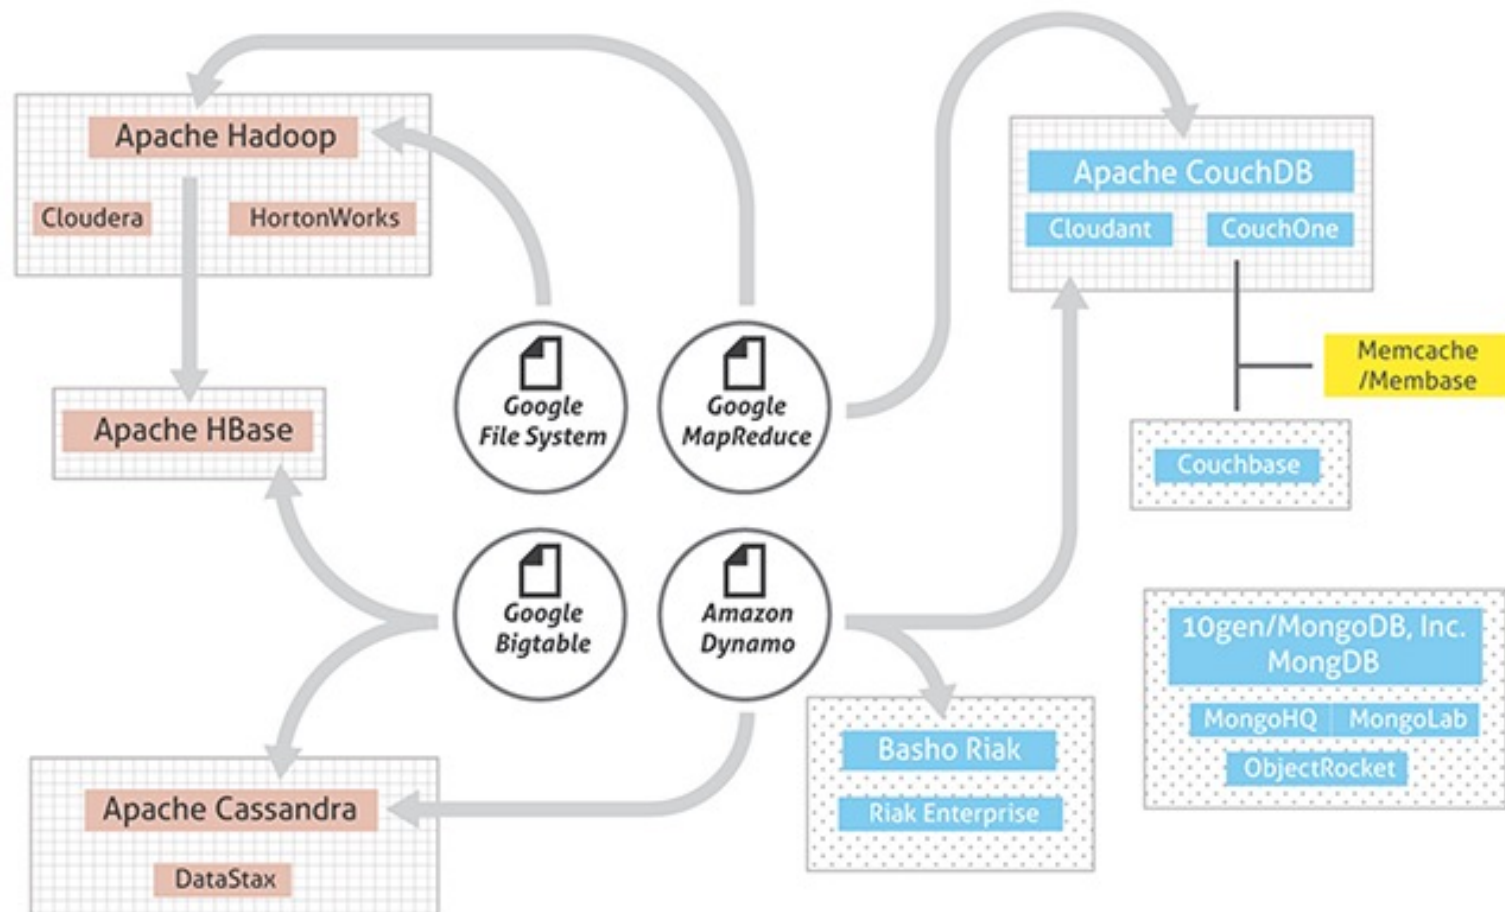
# HOW TO WRITE A CV

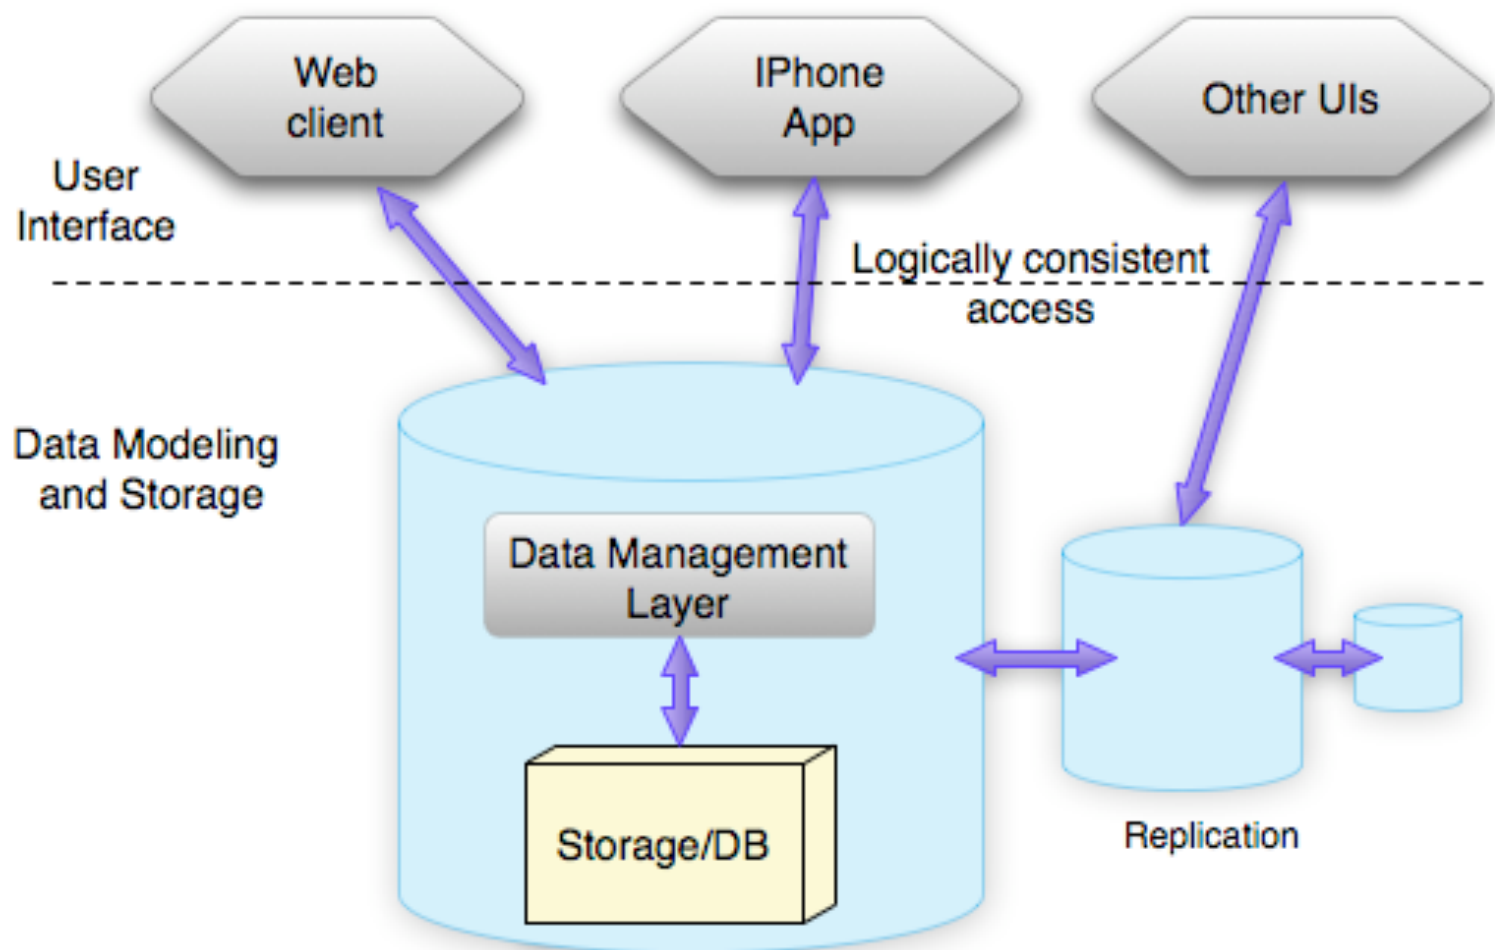


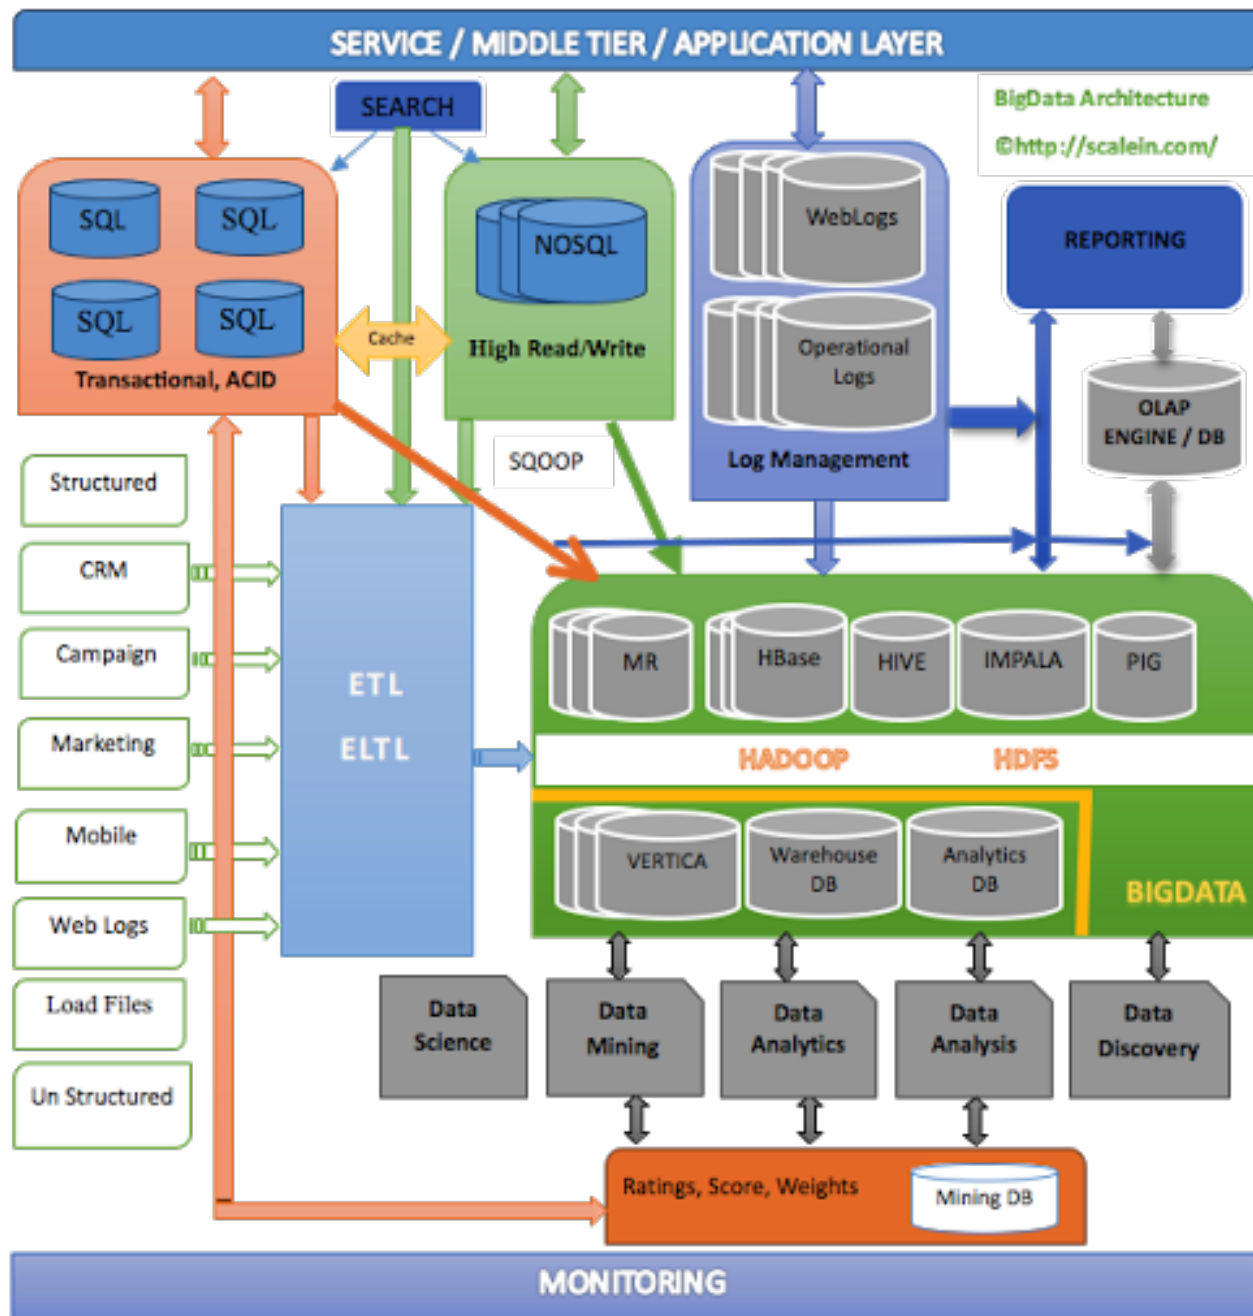
Leverage the NoSQL boom

# NoSQL FAMILY TREE

Understanding the Architecture that Runs Tomorrow's Web







S  
a