

Apache Iceberg



Ryan Blue
2019 Big Data Orchestration Summit

NETFLIX

Netflix's Data Warehouse



5-year Challenges

- Smarter processing engines
 - CBO, better join implementations
 - Result set caching, materialized views
- Reduce manual data maintenance
 - Data librarian services
 - Declarative instead of imperative

Problem Whack-a-mole

- Unsafe operations are everywhere
 - Writing to multiple partitions
 - Renaming a column
- Interaction with object stores causes major headaches
 - Eventual consistency to performance problems
 - Output committers can't fix it
- Endless scale challenges

What is Iceberg?



**Iceberg is a scalable format for
tables with a lot of best
practices built in.**

A format?

**We already have Parquet, Avro
and ORC . . .**

A table format.

A table format

- File formats help you modify or skip data in a single file
- Table formats do the same thing for a collection of files
- To demonstrate this, consider Hive tables . . .

Hive Tables

- Key idea: **organize data in a directory tree**

```
date=20180513/  
  |- hour=18/  
    |- ...  
  |- hour=19/  
    |- part-000.parquet  
    |- ...  
    |- part-031.parquet  
  |- hour=20/  
    |- ...
```

Hive Tables

- Filter: WHERE **date** = '20180513' AND **hour** = 19

date=20180513/

| - **hour=18/**

| | - ...

| - **hour=19/**

| | - **part-000.parquet**

| | - ...

| | - **part-031.parquet**

| - **hour=20/**

| | - ...

Hive Metastore

- Problem: **too much directory listing** for large tables
- Solution: use HMS to track partitions

`date=20180513/hour=19 -> hdfs:/.../date=20180513/hour=19`

`date=20180513/hour=20 -> hdfs:/.../date=20180513/hour=20`

- The file system still tracks the files in each partition . . .

Hive Tables: Problems

- State is kept in both the **metastore** and in a **file system**
- Changes are not atomic without locking
- Requires directory listing
 - $O(n)$ listing calls, $n = \#$ matching partitions
 - Eventual consistency breaks correctness

Hive Tables: Benefits

- Everything supports Hive tables*
 - Engines: Hive, Spark, Presto, Flink, Pig
 - Tools: Hudi, NiFi, Flume, Sqoop
- **Simplicity** and **ubiquity** have made Hive tables indispensable
- The whole ecosystem uses the same at-rest data!

Iceberg



Iceberg's Goals

- **An open spec and community for at-rest data interchange**
 - Maintain a clear spec for the format
 - Design for multiple implementations across languages
 - Support needs across projects to avoid fragmentation

Iceberg's Goals

- **Improve scale and reliability**
 - Work on a single node, scale to a cluster
 - All changes are atomic, with serializable isolation
 - Native support for cloud object stores
 - Support many concurrent writers

Iceberg's Goals

- **Fix persistent usability problems**
 - In-place evolution for schema and layout (no side-effects)
 - Hide partitioning: insulate queries from physical layout
 - Support time-travel, rollback, and metadata inspection
 - Configure tables, not jobs
- Tables should have **no unpleasant surprises**


Iceberg's Design

- Key idea: **track all files in a table** over time
 - A **snapshot** is a complete list of files in a table
 - Each write produces and commits a new snapshot



Iceberg's Design

- Readers use the current snapshot
- Writers optimistically create new snapshots, then commit



In reality, it's a bit more complicated.

Iceberg Design Benefits

- All changes are atomic
- No expensive (or inconsistent) file system operations
- Snapshots are indexed for scan planning on a single node
- CBO metrics are reliable
- Versions for incremental updates and **materialized views**

Iceberg at Netflix



Scale

- Production tables: **tens of petabytes, millions of partitions**
 - Scan planning fits on a single node
 - Advanced filtering enables more use cases
 - Overall performance is better
- Low latency queries are faster for large tables

Concurrency

- Production Flink pipeline writing in 3 AWS regions
- Lift service moving data into a single region
- Merge service compacting small files

Usability

- Rollback is popular
- Metadata tables
 - Track down the version a job read
 - Find the process that wrote a bad version

Future Work

- Spark vectorization for faster bulk reads
 - Presto vectorization already done
- Row-level delete encodings
 - MERGE INTO
 - ID equality predicates

Thank you!

Questions?

NETFLIX

Ryan Blue
rblue@netflix.com