


Complex Database Queries with PostgreSQL



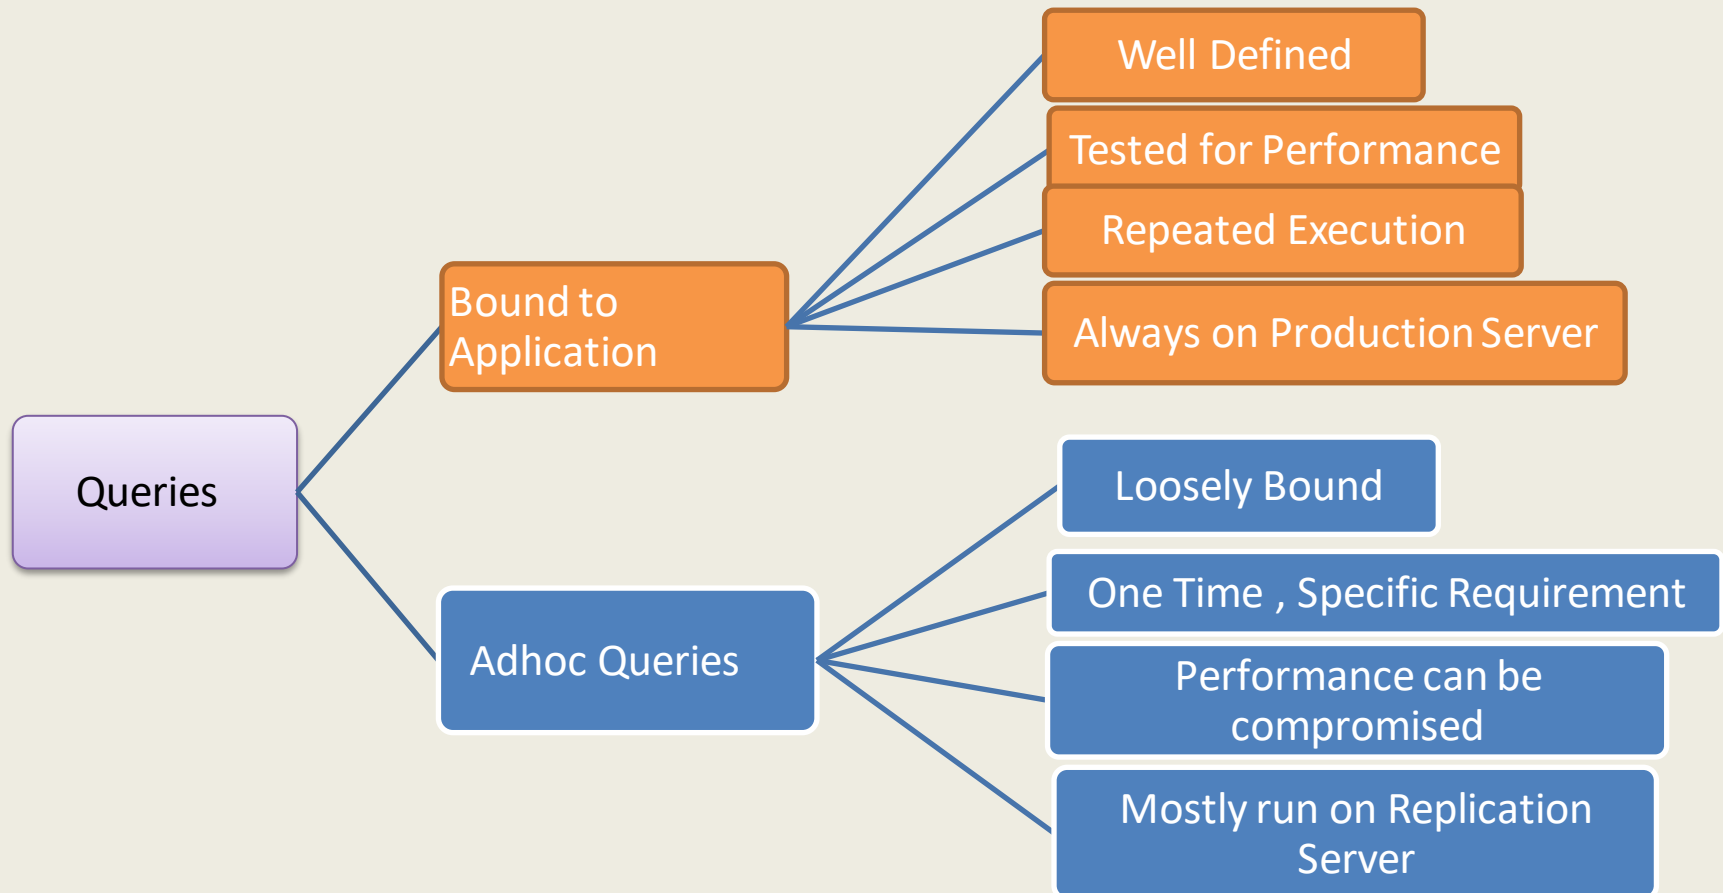
Dr.B.Hemalatha
Deputy Chief System Manager
IIT Kharagpur-INDIA

SQL Queries

- Vital part of all applications
- Creating, retrieving and manipulating data in an efficient way.—CRUD



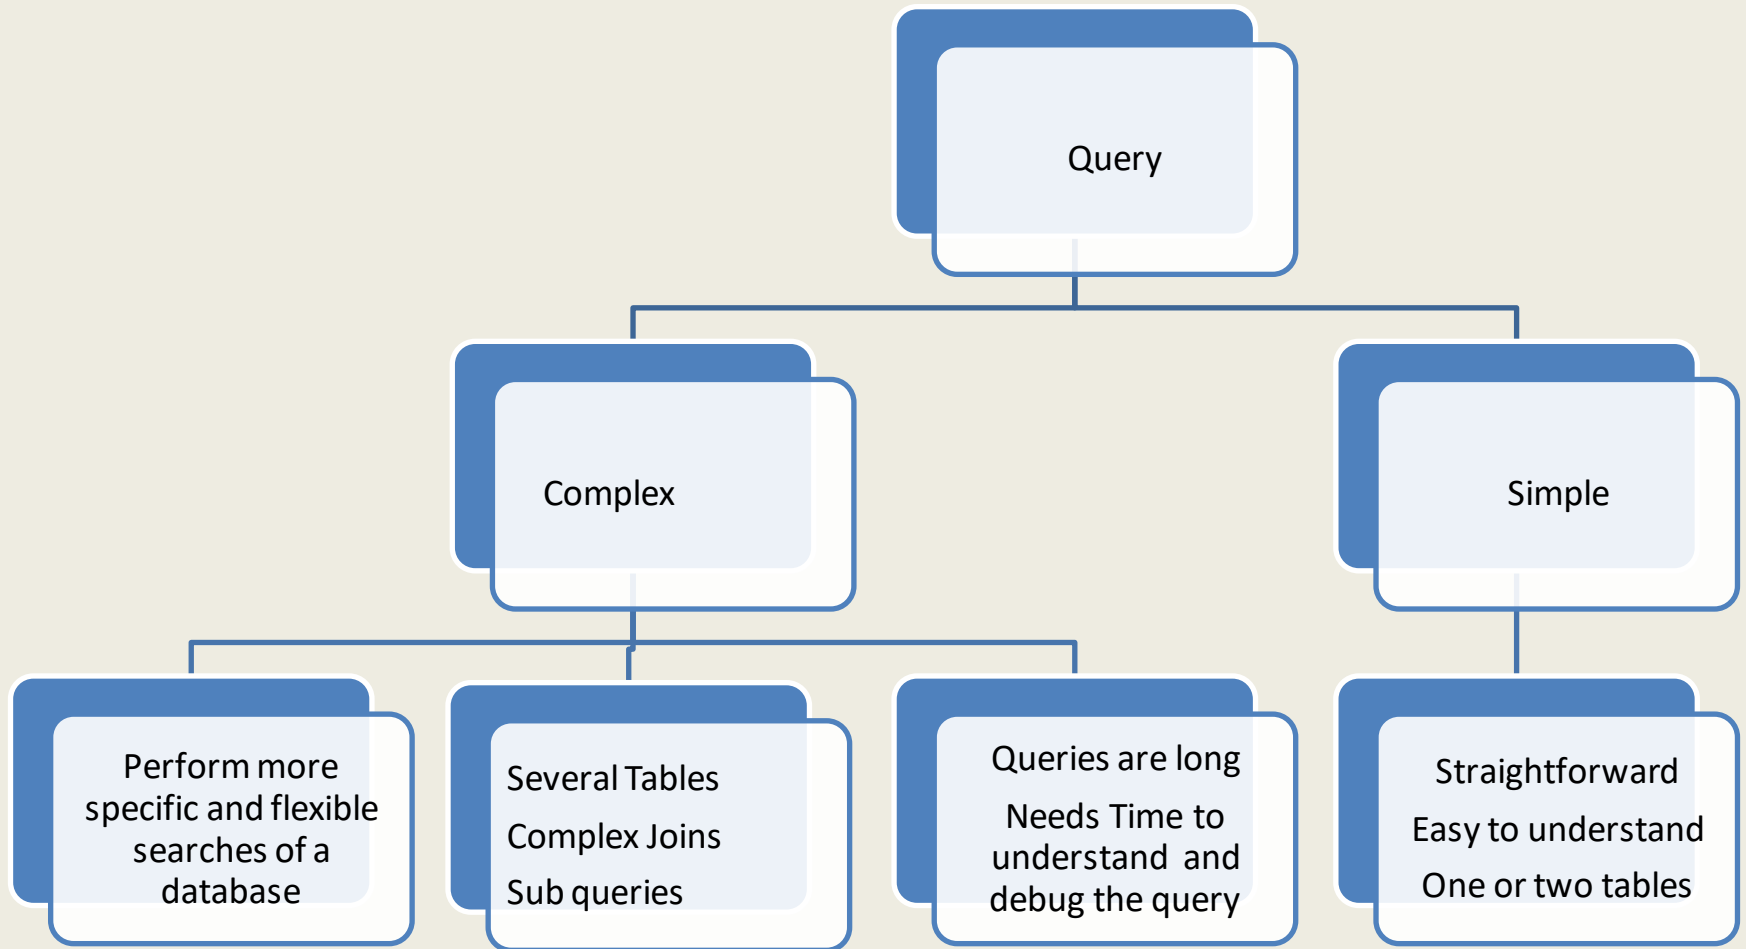
SQL Queries—Broad Classification




Adhoc Queries

- Adhoc queries are needed
 - To cater to the Data Requirements for Analytics
 - RTI requirements that are diverse and different
 - Specific requirements from National and International bodies
 - To test the output of a batch program
- More often the data requirement is in a particular format (pivot/un pivot)
 - In most cases these need complex queries to generate the data
- Adhoc queries work on data that will not normally change in that time span. Not currently transactional.
- Adhoc queries are run on a Streaming Replication Server using
 - pgAdmin Client
 - Microsoft Excel through ODBC drivers
 - Customized Application Interface
- Features available in PostGre 9.x and upwards make adhoc data retrieval easier

SQL Queries—Degree of Difficulty



Complex Queries--Construction



PostgreSQL and SQL Standard Conformance

- PostgreSQL tries to conform with the [SQL standard](#) where such conformance does not contradict traditional features or could lead to poor architectural decisions. Many of the features required by the SQL standard are supported, though sometimes with slightly differing syntax or function. Further moves towards conformance can be expected over time. As of the version 11 release in October 2018, PostgreSQL conforms to at least 160 of the 179 mandatory features for SQL:2011 Core conformance, where as of this writing, no relational database meets full conformance with this standard.

(Source: <https://www.postgresql.org/about/>)

A few PostgreSQL features for Complex Queries

- **JSON Support**
- **Filter Clause** –to subsets of data meeting certain conditions, used with aggregate and window functions
- **Array Aggregates** -Powerful way of converting several rows into one column
- **Window Functions** —Operates on a set of rows(OVER (PARTITION))
- **Statistical Functions** with the WITHIN GROUP clause
- **Grouping Sets** —Complex Grouping operations on all possible groups, rollup & cube (Multiple Group by in a single query)
- **Common Table Expressions** (CTE –WITH)
- **Use of Temporary Functions** (pg_temp)
- All the above features provide a elegant way of writing a complex query and delegating data fetching to the database

Platform & Examples

- All queries run & checked on pgAdmin4 version 3.2 on 64 bit Windows 10 Pro.
- Database "PostgreSQL 10.5 on x86_64-pc-linux-gnu, compiled by gcc (GCC) 4.8.5 20150623 (Red Hat 4.8.5-28), 64-bit"
- Representative Results only, real time object names and data have been masked.

JSON--Utility

- JSON Support -PostgreSQL provided JSON support since 9.2 and JSONB from 9.4 onwards. Postgre added the JSON feature well before the SQL standard added it in 2016 release.
- JSON objects particularly useful in complex queries to pivot and unpivot data.

Quick Examples-Unpivoting using JSON

| Choices |
|---------------|
| Regno |
| Choice1 |
| Choice2 |
| Choice3 |
| -- |
| Choicen |
| Other columns |

Straight query: `select regno, choice1, choice2,...choicen from choices`

Unpivoting

```
select choices."regno" As regno, j.key , j.val
from choices, LATERAL jsonb_each_text
(to_jsonb(choices)) AS j(key,val)
where adm_year = 'current' and j.key ilike '%cho%' and
j.key <> 'choice_allotted' and coalesce(j.value, '') not in ('')
order by regno, char_length(j.key), j.key
```

| Regno | Key | Value |
|---------|---------|-------|
| 17XXX02 | choice1 | EE |
| 17XXX02 | choice2 | ME |
| 17XXX02 | choice3 | IE |
| | | |

Columns are
converted to rows

Example-GROUPING SETS

```
select insti_scode, branch, count(distinct(regno))  
from allocations where allotted = 1 and session  
= 'current ' group by rollup(1,2)
```

| Insti_scode | branch | count |
|-------------|--------|-------|
| IITKGP | AE | 31 |
| IITKGP | AG | 28 |
| IITKGP | AR | 40 |
| : | | |
| IITKGP | | 1324 |

GRAND TOTAL

Branch
wise count
and grand
total, 2
group by
clauses

Filter Clause—Elegantly replaces CASE WHEN (pivots the output)

```
select branch, count(*) as tot, count(*) filter (where  
gender = 'F') as female, count(*) filter (where gender= 'M')  
as male, from allotted where insti_scode = 'IITKGP' and  
allotted = 1 group by rollup(1) order by 3 desc nulls first
```

| branch | tot | female | male |
|--------|------|--------|------|
| | 1587 | 263 | 1324 |
| ME | 181 | 27 | 154 |
| GG | 110 | 21 | 89 |
| EE | 118 | 18 | 100 |
| EC | 114 | 17 | 97 |

Overall Count

Branch wise
count with
gender
segregation and
overall count

Window Functions –Example using rank()

```
select * from (select a.roll,b.name,b.course,a.cgpa , rank() over
(partition by b.course order by a.cgpa desc) as rr from
performance a, studentmaster b where a.roll = b.roll and a.semno
= 'final' and a.session = 2016 and a.cgpa > 8.5 ) a where rr <=2
```

| Roll | Name | Course | CGPA | Rank |
|-----------|---------|--------|------|------|
| 13AEXXX1 | XXXXXXX | AE | 9.02 | 1 |
| 13CHXXX33 | XXXXXXX | CH | 9.21 | 1 |
| 13CHXXX95 | XXXXXXX | CH | 9.06 | 2 |
| 13CHXXX01 | XXXXXXX | CH | 9.06 | 2 |
| 13CSXX110 | XXXXXXX | CS | 9.26 | 1 |

Course wise
ranks. Neat
way of
picking ties.

Window Functions –Cumulative Totals

```
select semno, sem_crd_taken, sem_crd_cleared,  
sum(sem_crd_taken) over ( order by semno rows between  
unbounded preceding and current row ), sum(sem_crd_cleared)  
over ( order by semno rows between unbounded preceding and  
current row )from performance where roll = 'xxxxxxx' group by  
semno, sem_crd_taken,sem_crd_cleared
```

| Semno | Taken | Cleared | Cum Taken | Cum Cleared |
|-------|-------|---------|-----------|-------------|
| 1 | 23 | 23 | 23 | 23 |
| 2 | 22 | 22 | 45 | 45 |
| 3 | 21 | 21 | 66 | 66 |
| 4 | 21 | 21 | 87 | 87 |
| 5 | 22 | 22 | 109 | 109 |
| 6 | 22 | 22 | 131 | 131 |
| 7 | 23 | 23 | 154 | 154 |
| 8 | 22 | 22 | 176 | 176 |

Cumulative
Credits
Taken and
Cleared in a
single query

Statistical Functions(Mean, Median, Mode in one go)

- **select deptcode, round(avg(cgpa),2) , mode() within group (order by cgpa), percentile_cont(0.5) within group (order by cgpa), percentile_disc(0.5) within group (order by cgpa) from performance where semno = 'Final' and session = 'current' group by 1**

| Dept | Avg | Mode | MedianC | MedianD |
|------|------|------|---------|---------|
| AE | 7.34 | 8.48 | 7.73 | 7.73 |
| AG | 6.92 | 8.18 | 7.36 | 7.36 |
| AR | 7.21 | 6.52 | 7.52 | 7.52 |
| BT | 7.2 | 6.22 | 7.3 | 7.3 |

Deptwise
Mean,
Median &
Mode of
CGPA

Statistical Functions –Comparison with Local and Global Average

```
select deptcode,roll, semno, cgpa,round(avg(cgpa) over (partition  
by deptcode order by deptcode) ,2) as depavg,
```

```
cgpa,round(avg(cgpa) over (),2) as overall_avg from performance  
where year = 2018 and semno =1 group by 1,2,3,4 order by 1,2
```

| Dept | Regno | Semno | CGPA | Depavg | Overall avg |
|------|-----------|-------|------|--------|-------------|
| AE | 18AEXXXX1 | 1 | 7.73 | 7.34 | 7.74 |
| AE | 18AEXXXX2 | 1 | 9.18 | 7.34 | 7.74 |
| : | : | : | : | : | : |
| AR | 18ARXXXX4 | 1 | 8.09 | 7.21 | 7.74 |
| AR | 18ARXXXX5 | 1 | 6.55 | 7.21 | 7.74 |

Comparing
individual
scores with
local and
global
averages

Array_agg functions—Subject Overlaps

```
select a.subno, count(distinct a.roll), array_agg(distinct b.subno order by b.subno) as
associations, count(distinct b.subno) as intersections from registration a, registration
b where a.subno in (select subno from subjects_master where subject_type = 'T') and
b.subno in (select subno from subjects_master where subject_type = 'T') and a.roll =
b.roll group by 1 order by 4 desc
```

| Subno | Strength | Associations | Intersections |
|---------|----------|---------------------------------------|---------------|
| EP60020 | 300 | {AE40008,AE40034,A E51007 ...} | 259 |
| MA20104 | 753 | {AE21002,AE21004,A E21008 ...} | 203 |
| BM40002 | 104 | {AE31002,AE31004,A E31006, | 161 |
| AI61002 | 142 | {AE40008,AE40010,A E51007, | 161 |
| EP60002 | 137 | {AE21002,AE21004,A E21008,AE40008, | 156 |

Array of subjects
having common
registrations
through array_agg
function. Rows
are converted to a
single column

Temp Functions

```
create or replace function pg_temp.curr ( r VARCHAR)
```

```
//Temporary function to fetch the curricular details of a student
```

```
RETURNS TABLE ( r1 bpchar, cc varchar, sem smallint, patt1 text[], tot1 numeric)
```

```
AS $$
```

```
BEGIN
```

```
RETURN QUERY
```

```
(select roll, coursecode, maxsem, array_agg(subno order by subno) as patt, tot  
from z group by 1, 2, 3 5 order by 1, 2, 3) ;
```

```
END; $$ LANGUAGE 'plpgsql';
```

- Gives a dynamic character
- Can be used like a relation
- Lasts as long as session exists

Temporary Functions in Query

- `select * from pg_temp.curr('13MEXXXX')`
- `select a.*, b.name from studentmas b, pg_temp.curr(b.rollno) a where b.roll ilike '13ME%'` --Function used like a relation

| Roll | Coursecode | Max Semno | Pattern | Tot Sub |
|----------|------------|-----------|------------------------------|---------|
| 13MEXXX1 | ME | | 8{Br-1,De-50,El-6,HS-1,IE-1} | 59 |
| 13MEXXX2 | ME | | 8{Br-1,De-50,El-6,HS-1,IE-1} | 59 |

Common Table Expressions

- A common table expression (CTE) is as a temporary result set that is defined within the execution scope of a single SQL statement.
- A CTE is similar to a derived table and lasts only for the duration of the query.
- Unlike a derived table, a CTE can be self-referencing and can be referenced multiple times in the same query.

CTE--Explained

- Complex Query with nesting of subqueries

-*Select ... from* (*select ... from* (*select...from*) x) y
Outer↑Most Next↑Level Inner↑Most

-Reading the statement inside out

- CTE provides a compact way to break queries into manageable and comprehensible blocks

```
– WITH innermost as(select...), nextlevel as (select..
  From t1, innermost), outermost as(select * from
  innermost, nextlevel where <> )
select * from outermost ----Main Query
```

CTE--Advantages

- Using a CTE offers the advantages of improved readability and ease in maintenance of complex queries. The query can be divided into separate, simple, logical building blocks. These simple blocks can then be used to build more complex, interim CTEs until the final result set is generated.
- CTEs are statement scoped views
- Non Recursive or Recursive
- Read Top down

CTE—Non recursive example

Grade Distribution and Subwise Grades through CTE

- **WITH** **toppers** as (**select** * from(select roll, dept,cgpa rank() over (partition by dept order by cgpa desc) as rr) x where rr<=2) , **grade_distri** as (**select** roll, dept,cgpa,count(*), count(*) filter (where grade = 'EX') ... from **toppers**, performance group by roll,dept,cgpa), **subwise_grades**(**select** rollno, array_agg(subno || '-' grade order by grade) as patt from grades, **toppers** group by subno)
- **select** a.*, b.patt from grade_distri, subwise_grades

| Roll | Dept | Cgpa | TOT | EX | A | B | C | D | P | F | Patt |
|------|------|------|-----|----|----|----|---|---|---|---|-------------------|
| 13AE | AE | 8.85 | 64 | 17 | 23 | 18 | 4 | 1 | 1 | 0 | {AE21-A,...EV2-P} |

CTE—Repeated Condition

- CTEs are useful when a condition has to be repeatedly applied in the where clause.
- WITH `current_exam_session` as (select subno, exam_session from `exam_timetable` where session = 'current' and exam_type = 'mid'), `student_exam_timetable` as (select roll, a.subno, exam_session from `current_exam_session` a, registration b where a.subno = b.subno), `sub_wise_slot_occupancy` as (select a.subno,a.exam_session,array_agg(distinct TRIM(b.exam_session) order by TRIM(b.exam_session)) as occupancy from `student_exam_timetable` a, `student_exam_timetable` b where a.roll = b.roll group by a.subno, a.exam_session)
- Main Query
select * from `sub_wise_slot_occupancy`

| Subno | Exam Session | Occupancy |
|---------|--------------|------------------|
| AI60001 | S2 | {S1,S2,S4,S5,S7} |

CTE—Other Uses

- Assign column names to unnamed columns using CTE

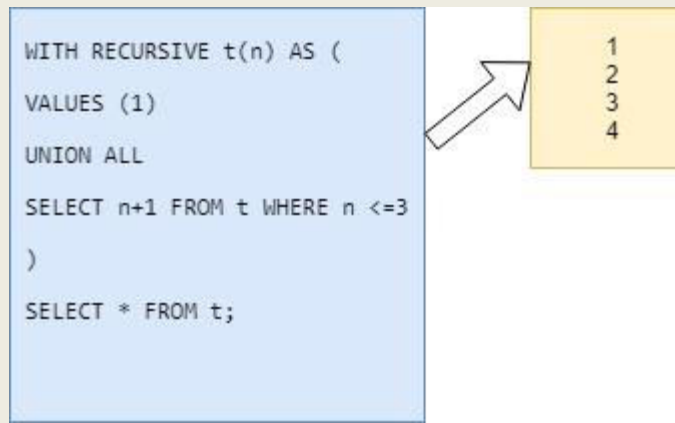
```
WITH grademst(grade, points) as  
( values ('A', 9), ('B',8),('C', 7),('D', 6), ('P',5),('F',0), (' ',0))  
select grade, points from grademst
```

- Overload a schema table with CTE

| Grade | Points |
|-------|--------|
| A | 9 |
| B | 8 |
| : | |
| | 0 |

CTE-Recursive

- Refer to themselves in the second leg of UNION ALL
- Example:



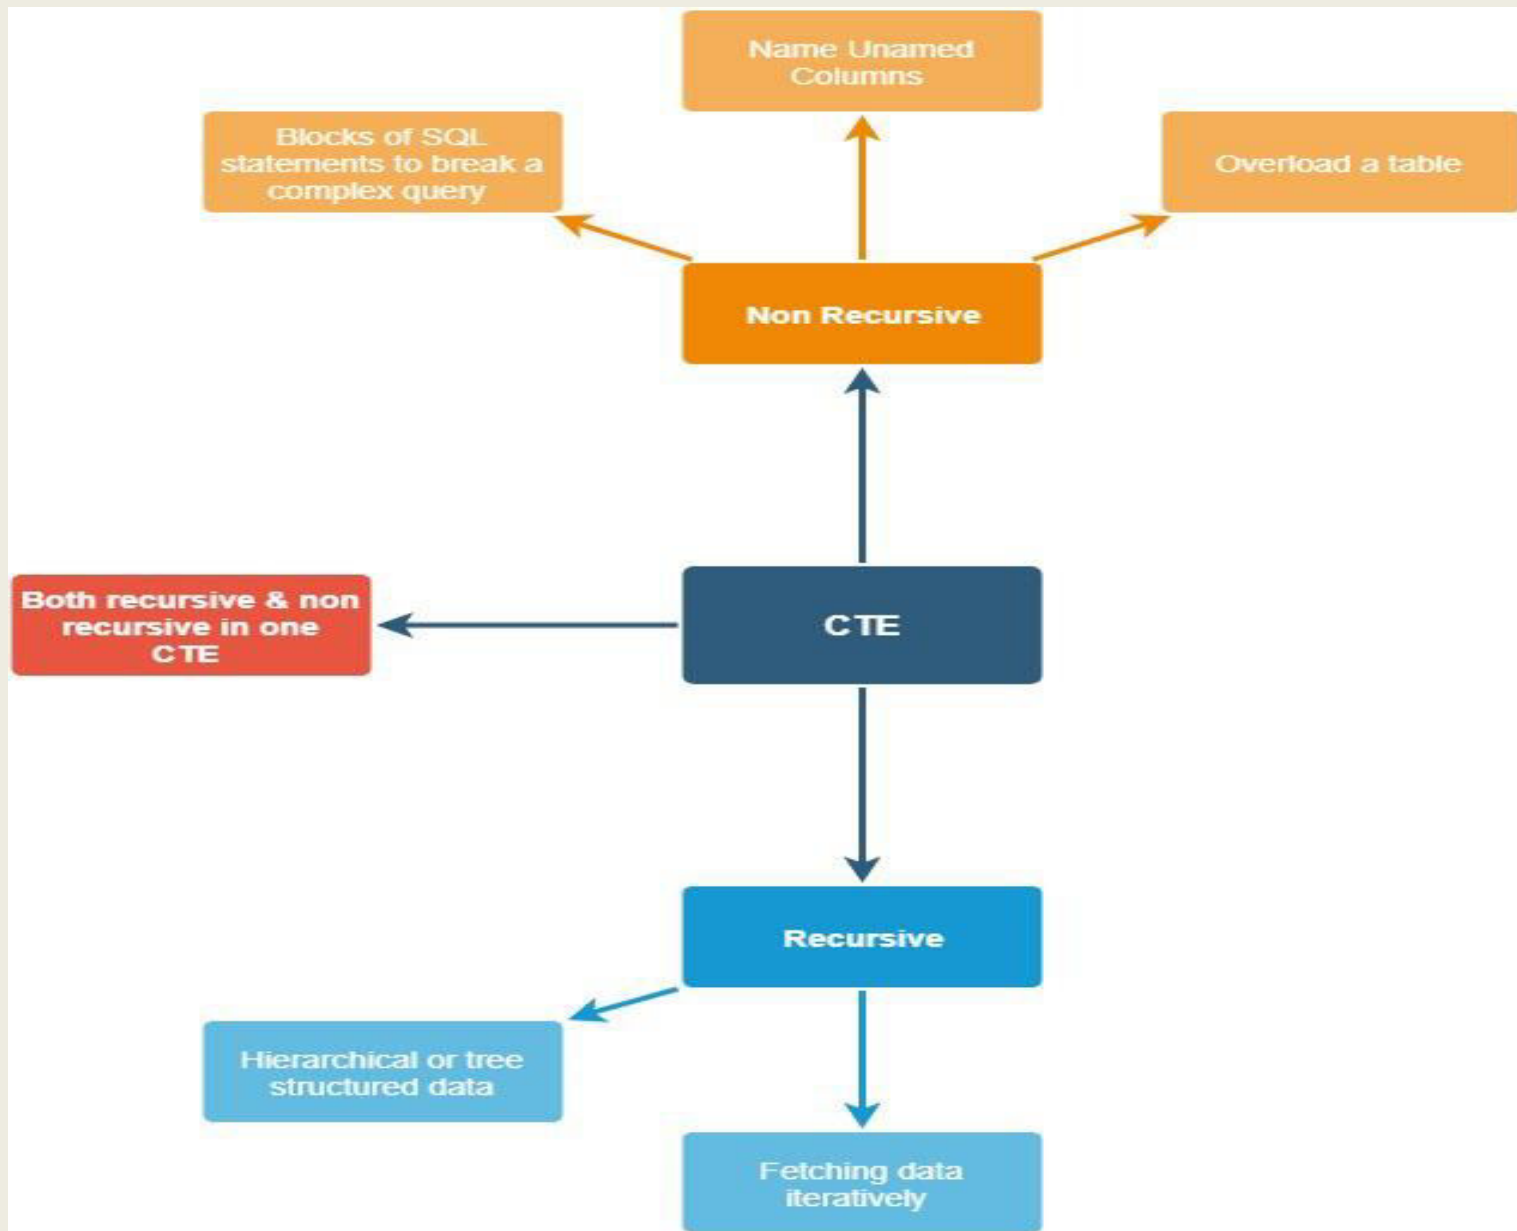
- Actually an iteration—similar to while <cond> loop
- Execution terminates when no tuples are returned.
- To prevent an infinite loop, limit clause can be used in the outer select

Recursive CTE -Example

- Semester wise toppers for a particular batch
- First with picks the profile of the toppers for semester 1 and is recursive over the semesters upto final

```
-WITH recursive toppers_eachsem as (select roll,name, semno,course, cgpa,
rank()over (partition by course order by cgpa desc) as rr from performance where
adm_year='2013' and semno = 1 group by roll,name, semno,course,cgpa ) UNION
ALL
(select roll,name, a.semno,course, ,cgpa,rank()over (partition by course order by cgpa
desc) as rr from performance a, (select distinct semno+1 as semc from
toppers_eachsem) where adm_year='2013' and a.semno = c.semno group by
roll,name, semno,course,cgpa)
select course,semno, rollno,name,cgpa,rr from toppers_eachsem where rr =1 order
by course,semno
```

| Course | Semno | Roll | Name | Cgpa | rank |
|--------|-------|---------|--------|------|------|
| AE | 1 | 13AEXXX | Xyxyxy | 9.1 | 1 |
| AE | 2 | 13AEZZZ | ZYZYZY | 9.4 | 1 |



WITH CLAUSE—PERFORMANCE ISSUES IN POSTGRESQL

- While most other databases process WITH in the same way as they process views and derived tables and optimise the overall query, there is big difference in PostgreSQL.
- The PostgreSQL query planner considers each with query and the main statement separately.
- The usual query optimization methods are yet to be implemented for CTE operations
 - Projection Pushdown(eliminating unnecessary columns)
 - Predicate Pushdown (apply the predicate as early as possible)
 - Sort Elimination (redundant sort operations)
- This makes the with clause an **optimization fence** meta data, such as the order in which a with query returns rows, is not available during optimization of the main query.(ref: @MarkusWinad)

WITH CLAUSE—Performance Example

- `explain with current_exam_session as (select subno, exam_session from exam_timetable where session = 'current' and exam_type = 'Mid') , student_exam_timetable as (select roll, a.subno, exam_session from current_exam_session a, registration b where a.subno = b.subno)`

`select * from student_exam_timetable`

Output

"CTE Scan on student_exam_timetable (cost=7046.72..8797.34 rows=87531 width=80)"

--In the absence of WITH

- `explain select roll, a.subno, exam_session from exam_timetable a, registration b where session = 'current' and exam_type = 'Mid' and a.subno = b.subno`

"Hash Join (cost=114.99..3594.84 rows=45293 width=21)"

Complex Queries –Application in Academic Information System

- Mostly Adhoc queries for checking batch operations, reports and Analytics
- Activities used include
 - Exam Scheduling (using **WITH and aggregate functions**) (*B Hemalatha. A Multistage Technique for Examination Timetabling. International Journal of Computer Applications 181(33):5-11, December 2018*)
 - Curriculum Compliance by generating patterns using **array_agg** functions
 - Class Scheduling (using **WITH and array_agg**)
 - Performance Ranking for various criteria (**WITH** in conjunction with **window functions** ,**aggegrate functions** along with the **filter** clause)
 - Reports needing Unpivoting/Pivoting through **JSON**

CONCLUSIONS

- Complex queries can be compacted through WITH clause
- The JSON and filter clause enable us to unpivot and pivot data
- Rich features available through Array_agg and statistical functions enable more compaction of queries
- Combining all the above, the data manipulation task can be delegated at the database level itself, especially for Adhoc queries

References

- <https://www.postgresql.org/docs/10/index.html>
- <https://modern-sql.com/> --@MarkusWinand
- <https://momjian.us/main/presentations/sql.html>
- <https://www.tutorialspoint.com/postgresql/>
- <http://www.postgresqltutorial.com/>
- <https://www.linuxjournal.com/content/postgresql-10-great-new-version-great-database>

ACKNOWLEDGEMENTS

- IIT Kharagpur
- The entire student community of IIT KGP
- Organizers of PGConf India 2019.