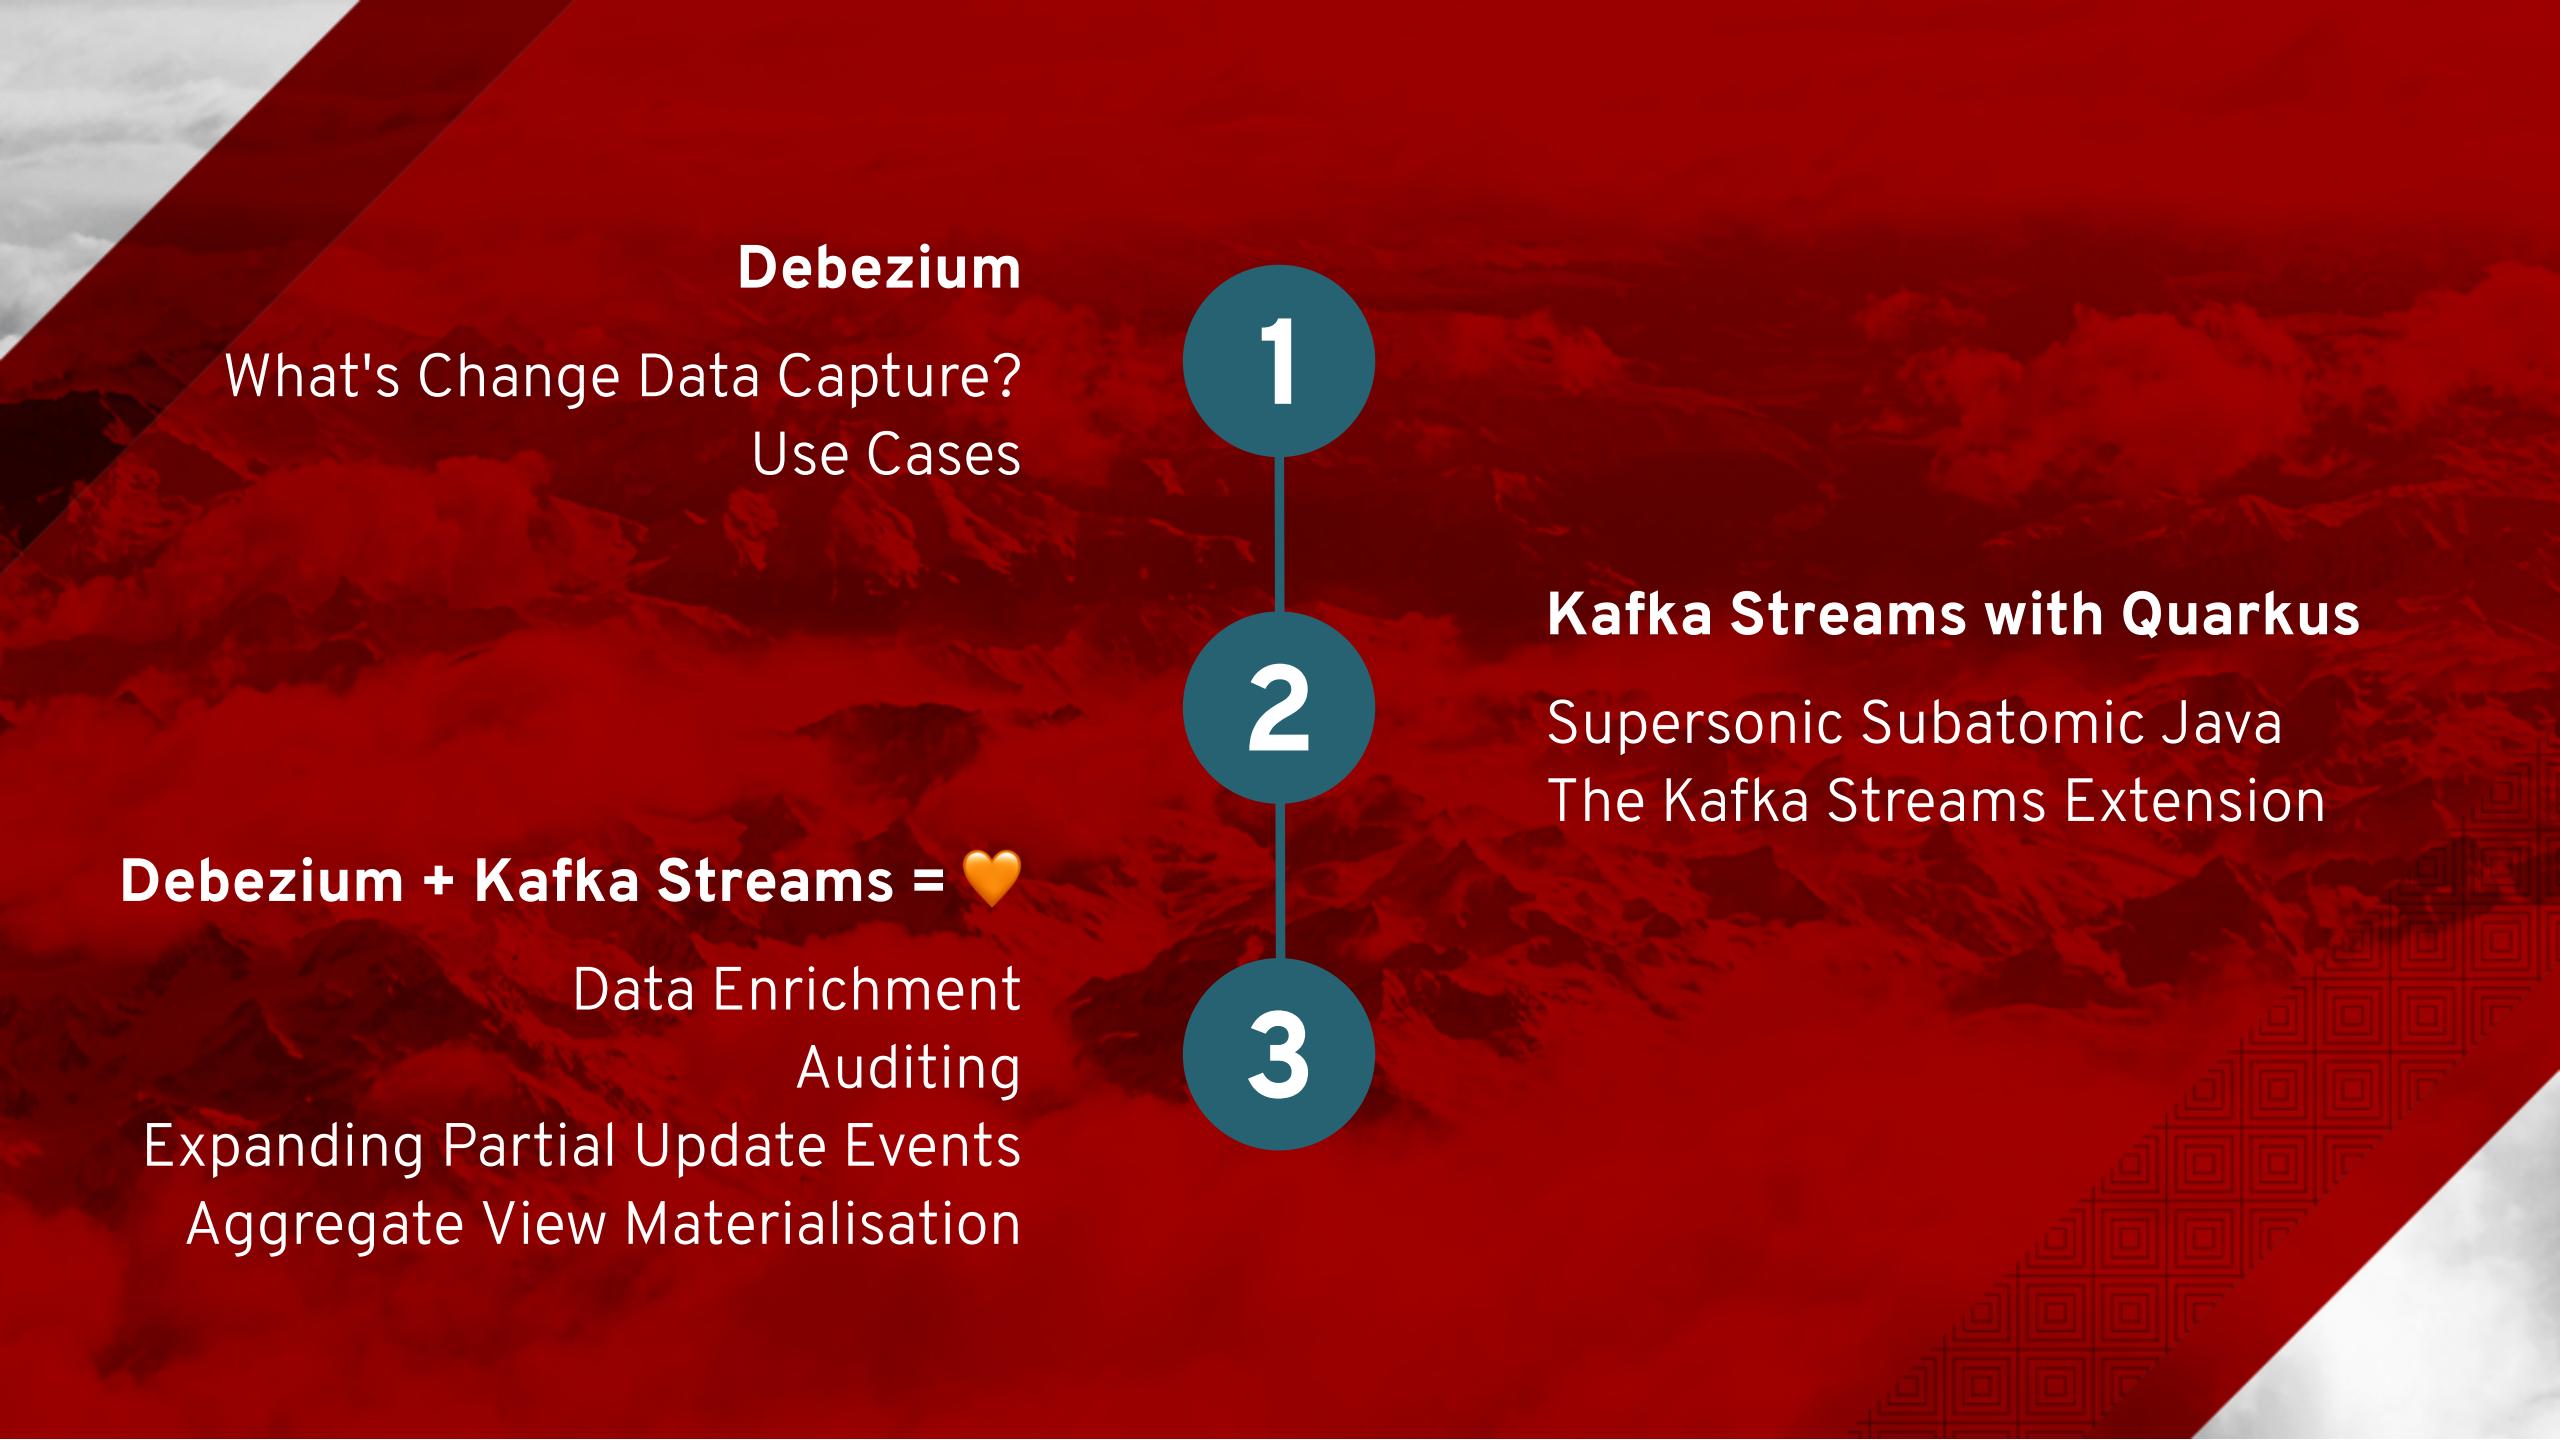




Change Data Capture Pipelines With Debezium and Kafka Streams

Gunnar Morling
Software Engineer
@gunnarmorling





What's Change Data Capture?
Use Cases

Debezium

Debezium + Kafka Streams = ❤️

Data Enrichment
Auditing
Expanding Partial Update Events
Aggregate View Materialisation



Kafka Streams with Quarkus
Supersonic Subatomic Java
The Kafka Streams Extension

Gunnar Morling

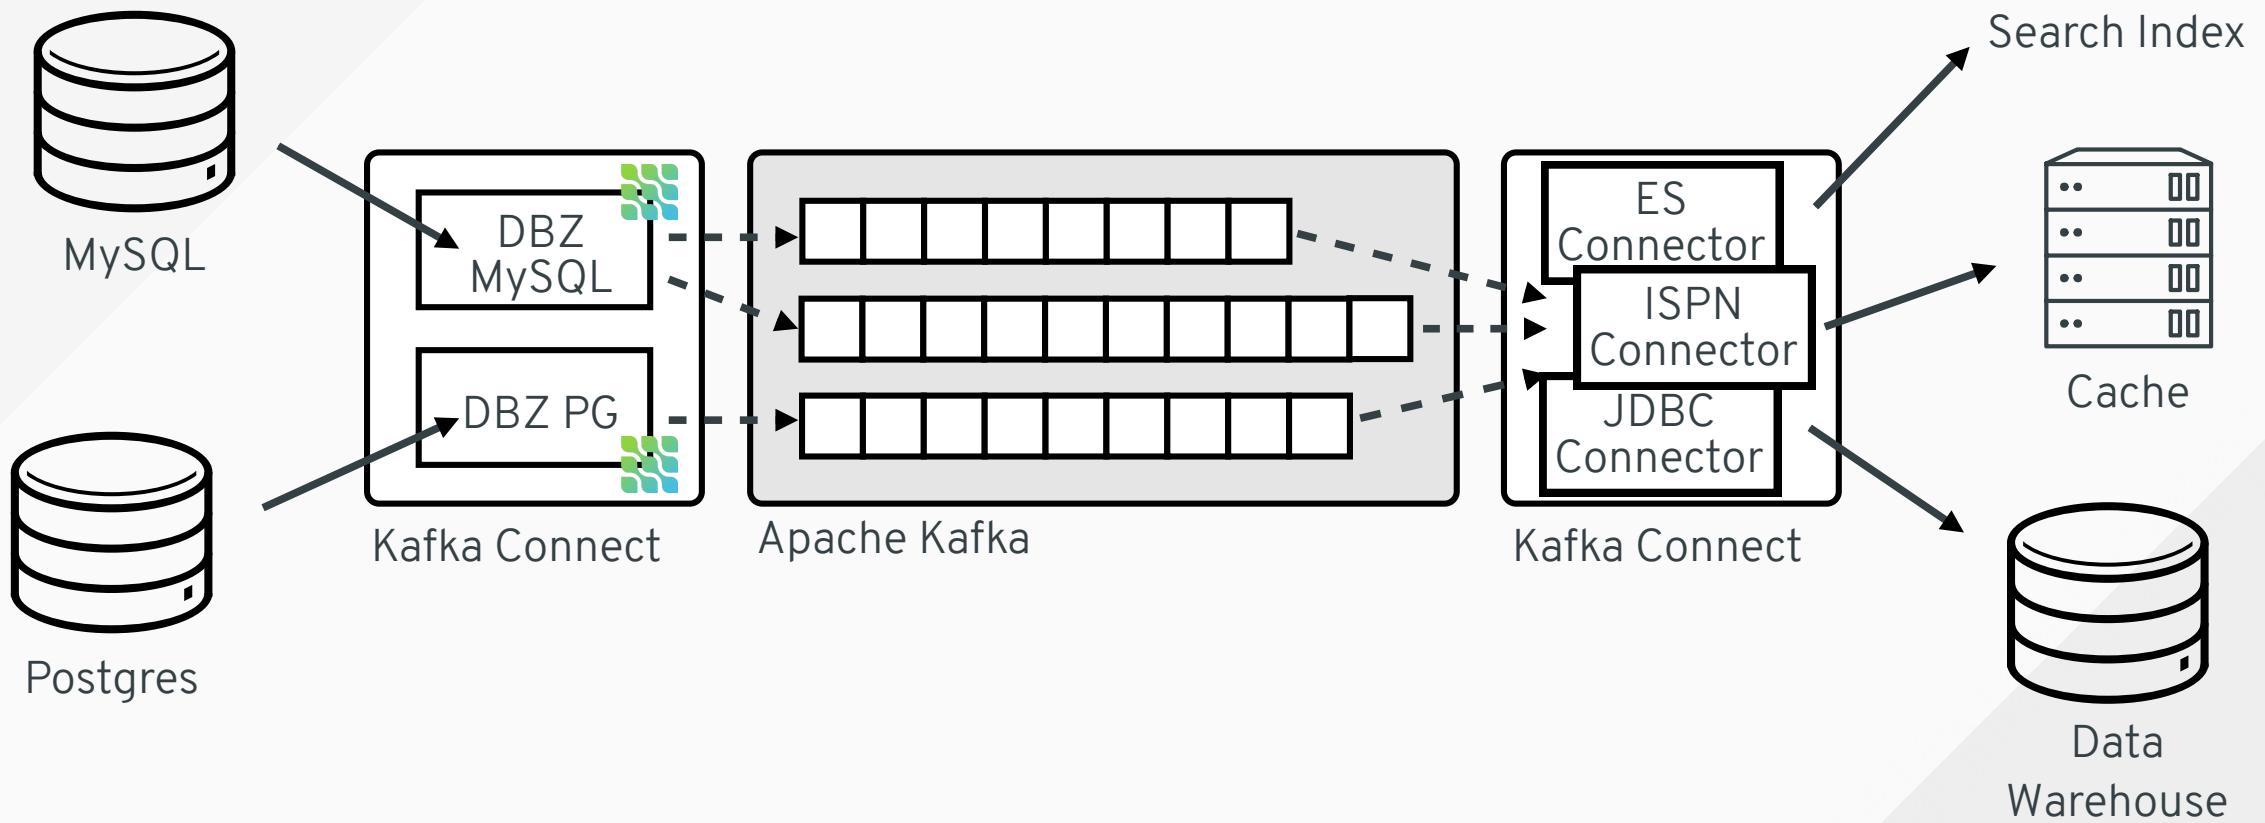
- Open source software engineer at Red Hat
 - **Debezium**
 - Quarkus
 - Hibernate
- **Spec Lead** for Bean Validation 2.0
- Other projects: **Deptective**, MapStruct
- Java Champion





Debezium

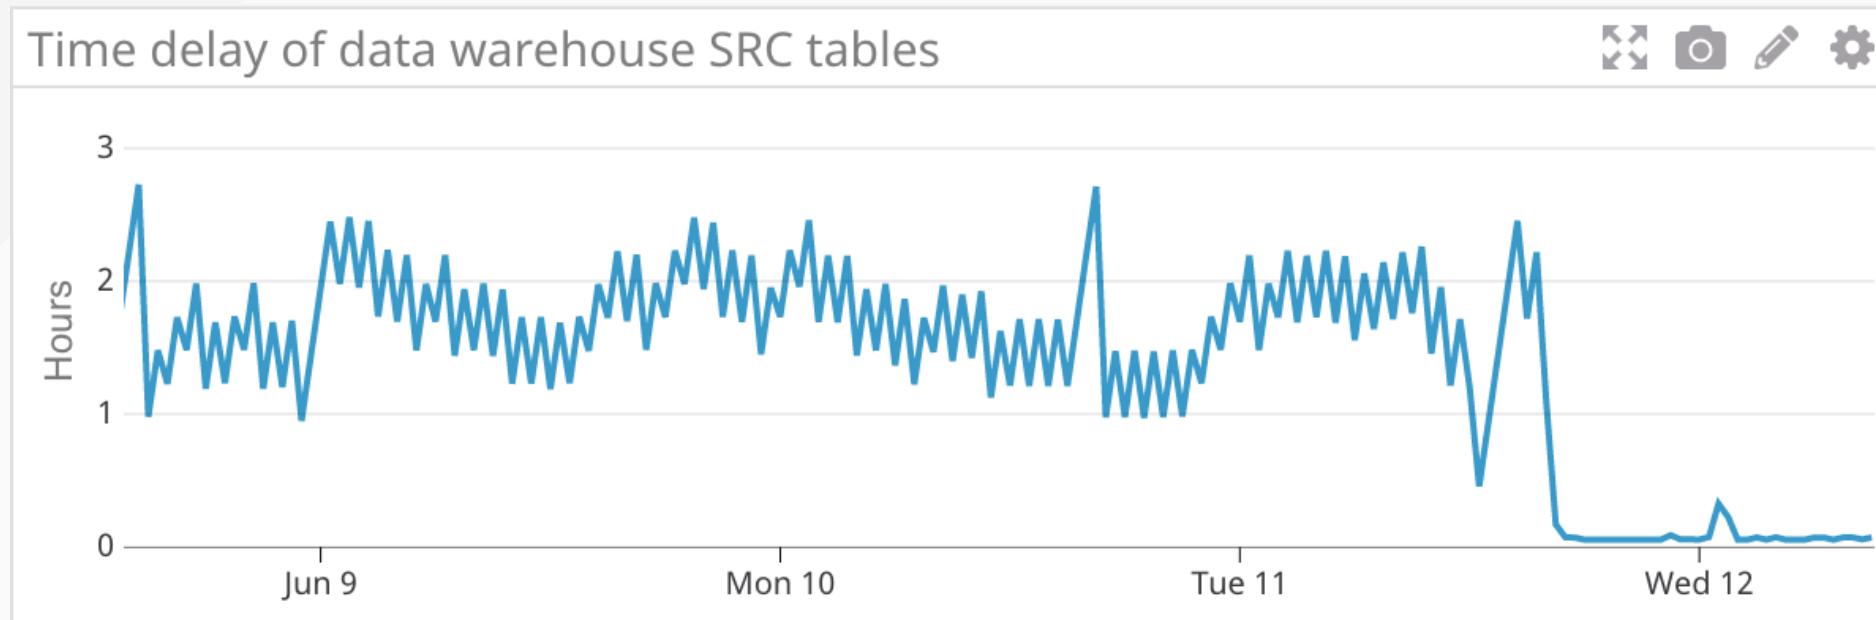
Enabling Zero-Code Data Streaming Pipelines





Debezium

Low-Latency Change Data Streaming

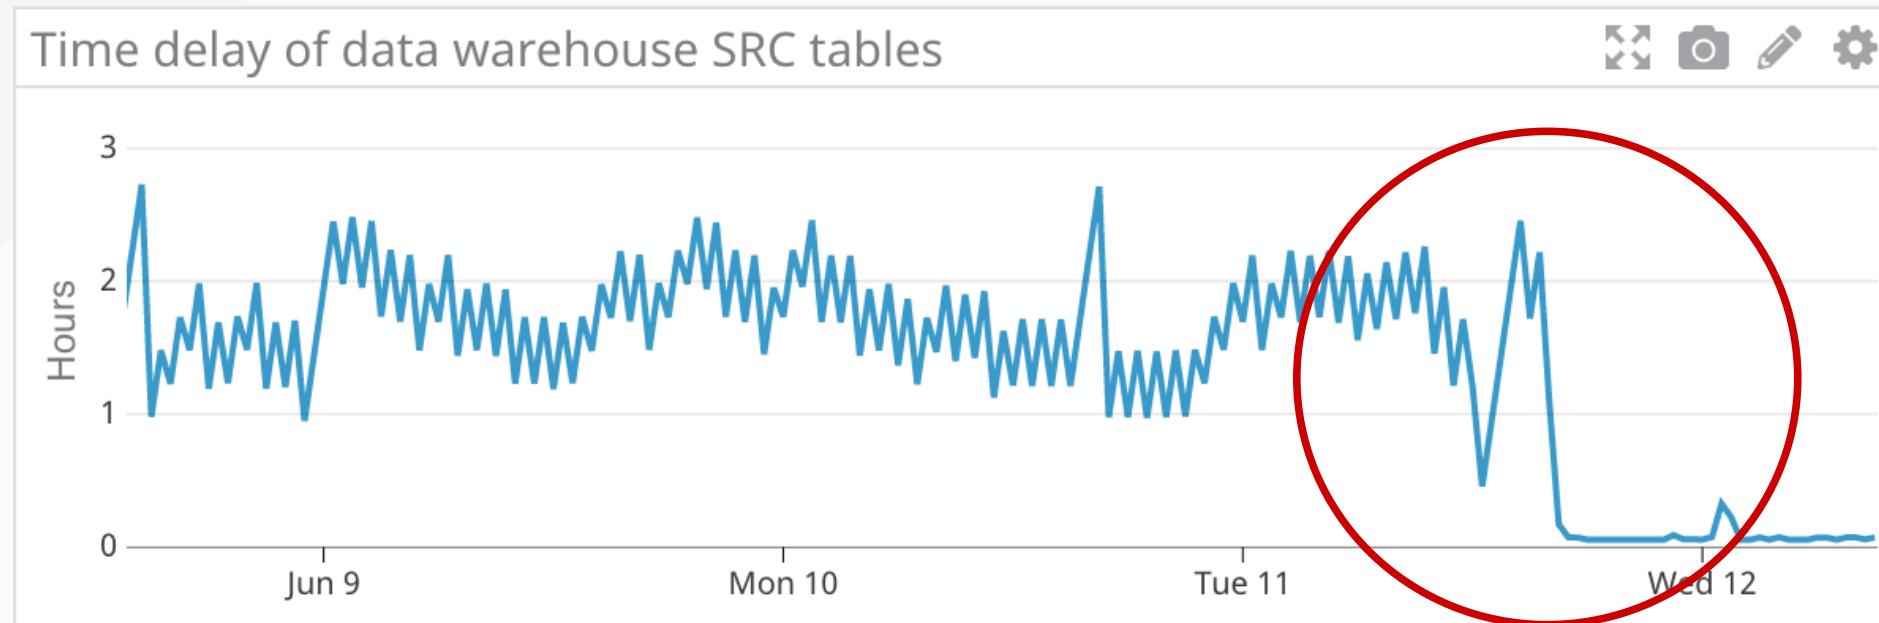


<https://medium.com/convoy-tech/>



Debezium

Low-Latency Change Data Streaming



<https://medium.com/convoy-tech/>



Debezium Connectors

- MySQL
- Postgres
- MongoDB
- SQL Server
- Cassandra (Incubating)
- Oracle (Incubating)
- Db2 (Incubating)
- Future additions: Vitess, MariaDB





CDC - "Liberation for Your Data"

Gunnar Morling 
@gunnarmorling

Change data capture is one giant enabler; it lets you

- * replicate data
- * feed search indexes
- * update caches
- * run streaming queries
- * sync data between microservices
- * maintain denormalized views
- * create audit logs and so much more.

Ultimately, it's liberation for your data.

 Tweet übersetzen
13:46 - 30. Apr. 2019

27 Retweets 67 „Gefällt mir“-Angaben

Q 2 T 27 67



CDC - "Liberation for Your Data"

Gunnar Morling (@gunnarmorling)

Change data capture is one giant enabler; it lets you

- * replicate data
- * feed search indexes
- * update caches
- * run streaming queries
- * sync data between microservices
- * maintain denormalized views
- * create audit logs and so much more.

Ultimately, it's liberation for your data.

Tweet übersetzen
13:46 - 30. Apr. 2019

27 Retweets 67 „Gefällt mir“-Angaben

2 27 67



Change Event Structure

- Key: Primary key of table
- Value: Describing the change event
 - Old row state
 - New row state
 - Metadata

```
{  
  "before": null,  
  "after": {  
    "id": 1004,  
    "first_name": "Anne",  
    "last_name": "Kretchmar",  
    "email": "annek@noanswer.org"  
  },  
  "source": {  
    "name": "dbserver1",  
    "server_id": 0,  
    "ts_sec": 0,  
    "file": "mysql-bin.000003",  
    "pos": 154,  
    "row": 0,  
    "snapshot": true,  
    "db": "inventory",  
    "table": "customers"  
  },  
  "op": "c",  
  "ts_ms": 1486500577691  
}
```





Log- vs. Query-Based CDC

	Query-Based	Log-Based
All data changes are captured	-	+
No polling delay or overhead	-	+
Transparent to writing applications and models	-	+
Can capture deletes and old record state	-	+
Installation/Configuration	+	-

Debezium

What's Change Data Capture?
Use Cases

Debezium + Kafka Streams = ❤

Data Enrichment
Auditing
Expanding Partial Update Events
Aggregate View Materialisation



Kafka Streams with Quarkus

Supersonic Subatomic Java
The Kafka Streams Extension



Quarkus

Supersonic Subatomic Java



“A Kubernetes Native Java stack tailored for OpenJDK HotSpot and GraalVM, crafted from the best of breed Java libraries and standards.



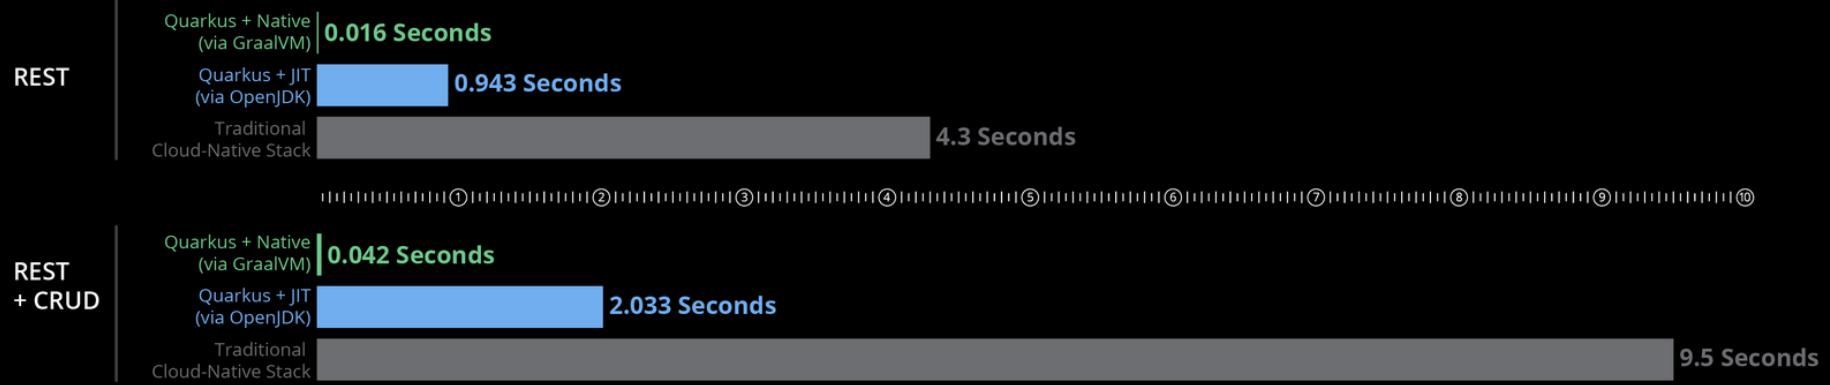
```
$ ./my-native-java-rest-app  
Quarkus started in 0.008s
```

Memory (RSS) in Megabytes*

*Tested on a single-core machine



BOOT + First Response Time

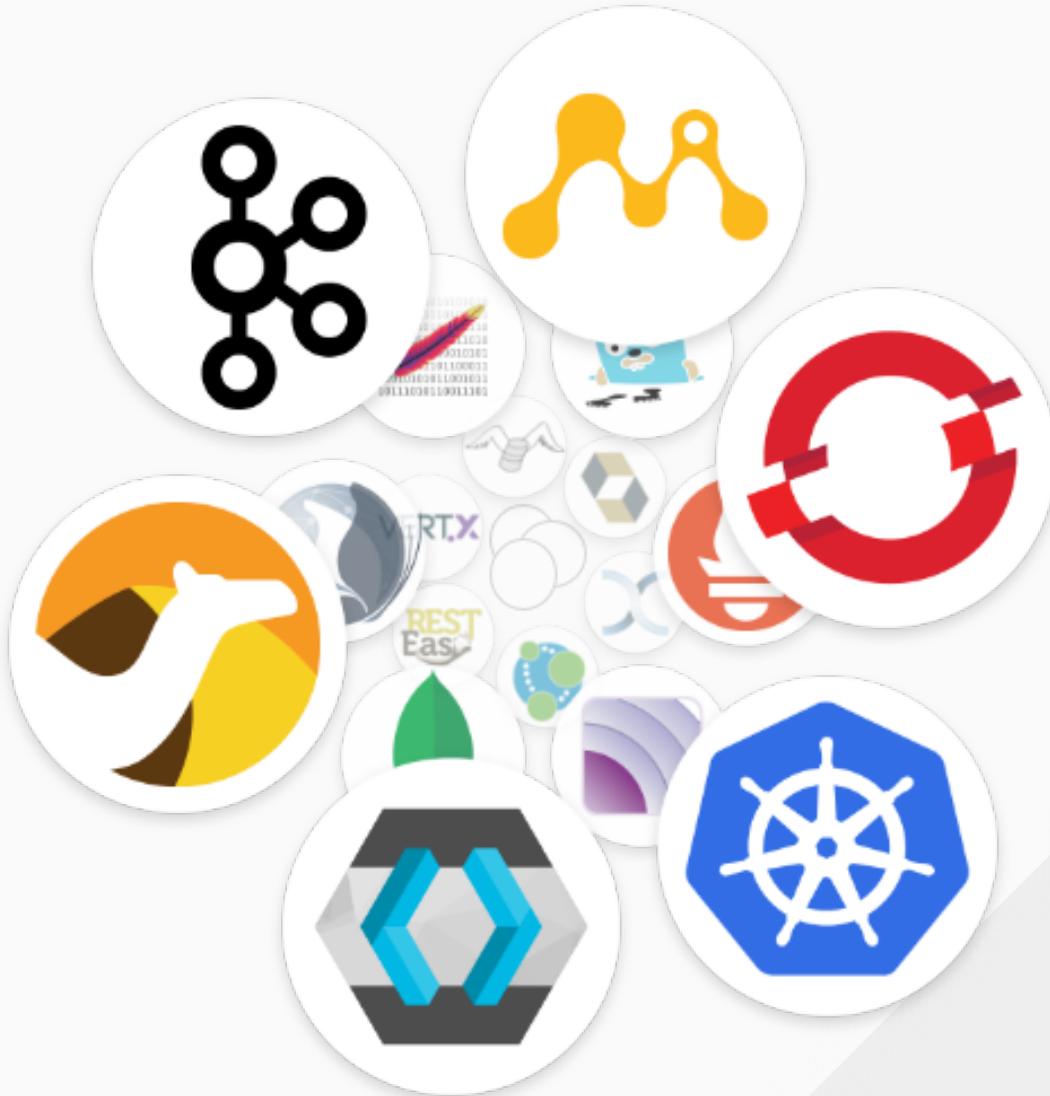




Quarkus

Supersonic Subatomic Java

- **Developer joy**
 - **Imperative and Reactive**
 - **Best-of-breed libraries**

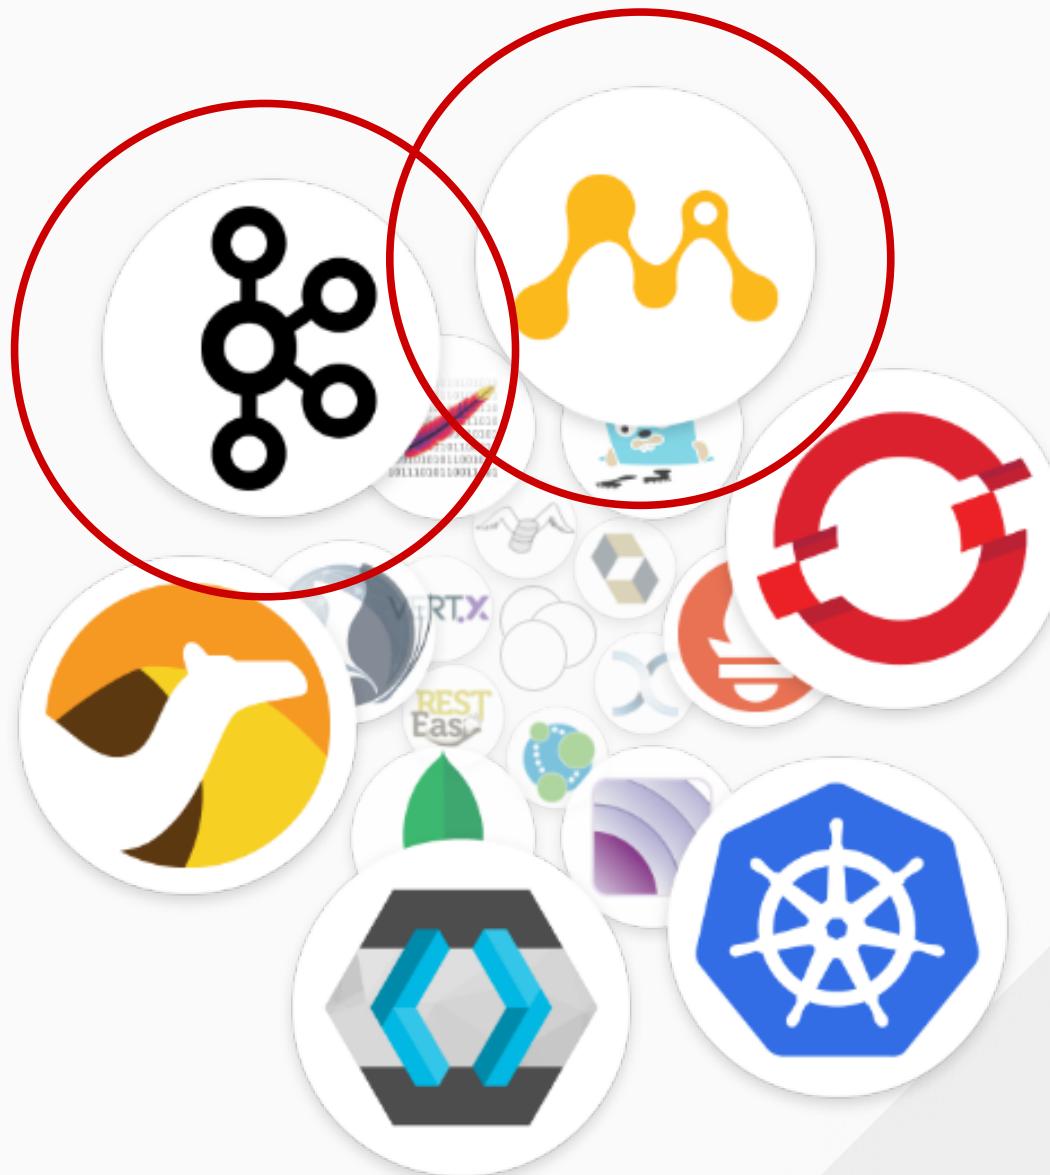




Quarkus

Supersonic Subatomic Java

- **Developer joy**
- **Imperative and Reactive**
- Best-of-breed **libraries**

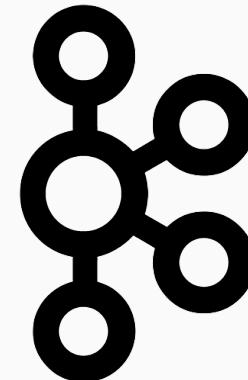




Quarkus

The Kafka Streams Extension

- **Management** of topology
- **Health** checks
- **Dev Mode**
- Support for **native binaries** via GraalVM



GraalVM™

Debezium

What's Change Data Capture?
Use Cases

Debezium + Kafka Streams = ❤

Data Enrichment
Auditing
Expanding Partial Update Events
Aggregate View Materialisation



Kafka Streams with Quarkus
Supersonic Subatomic Java
The Kafka Streams Extension

A photograph of a group of people swimming in the ocean at sunset. The water is a warm, golden color, and the sky above is a bright yellow. The silhouettes of the swimmers are visible against the water. In the background, there are distant hills or mountains.

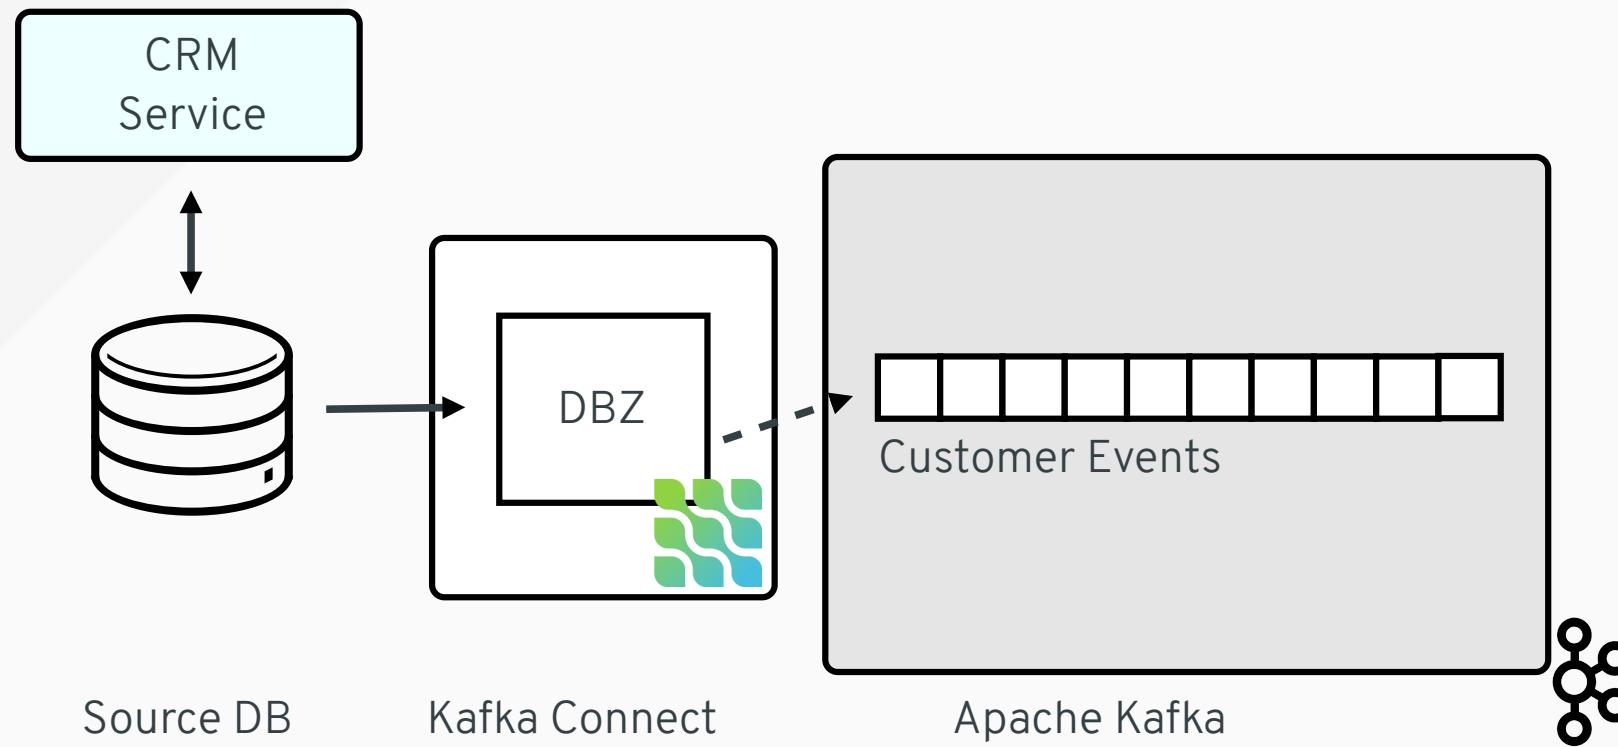
Data Enrichment - Demo

A nighttime photograph of a city street scene. In the foreground, a chain-link fence runs along a curved concrete barrier. Above the fence, a bridge with a metal railing spans the street. The background features a tall building with many lit windows and palm trees. The overall image has a motion blur effect, particularly in the lights from the building and the bridge.

Auditing



Auditing

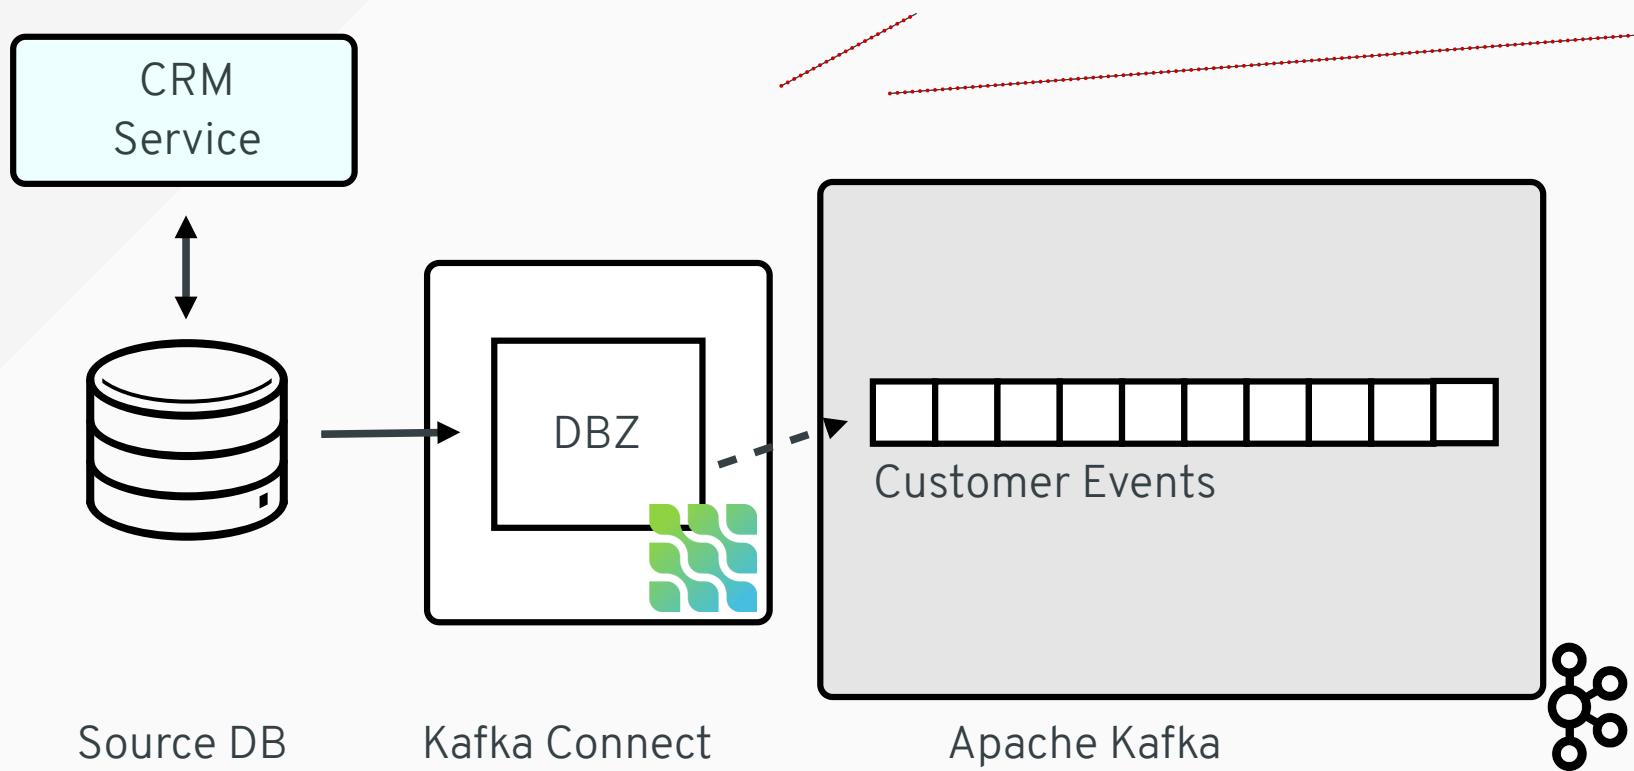




Auditing

"Transactions" table

Id	User	Use Case
tx-1	Bob	Create Customer
tx-2	Sarah	Delete Customer
tx-3	Rebecca	Update Customer

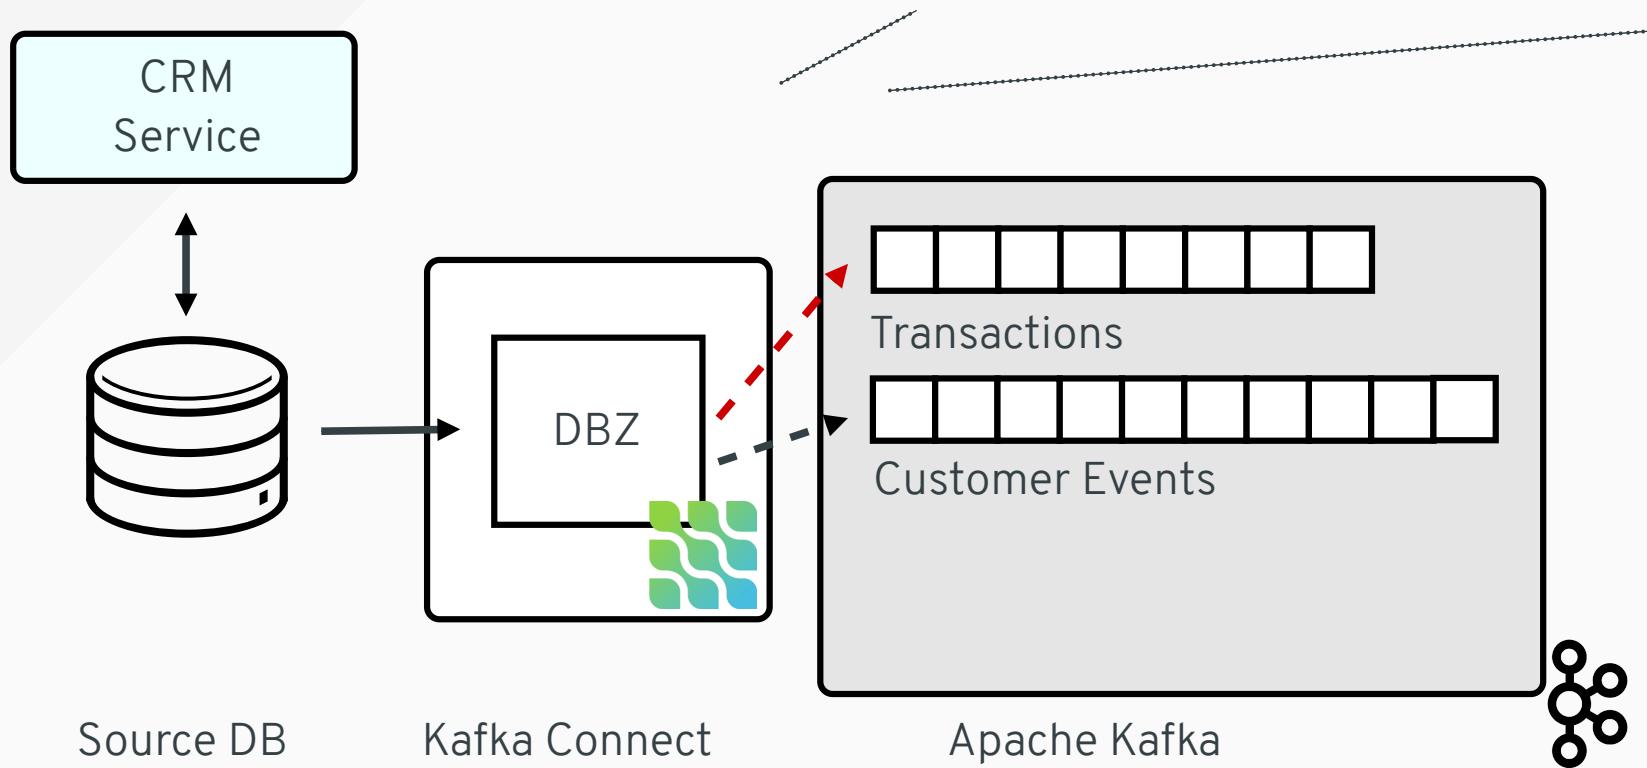




Auditing

"Transactions" table

Id	User	Use Case
tx-1	Bob	Create Customer
tx-2	Sarah	Delete Customer
tx-3	Rebecca	Update Customer

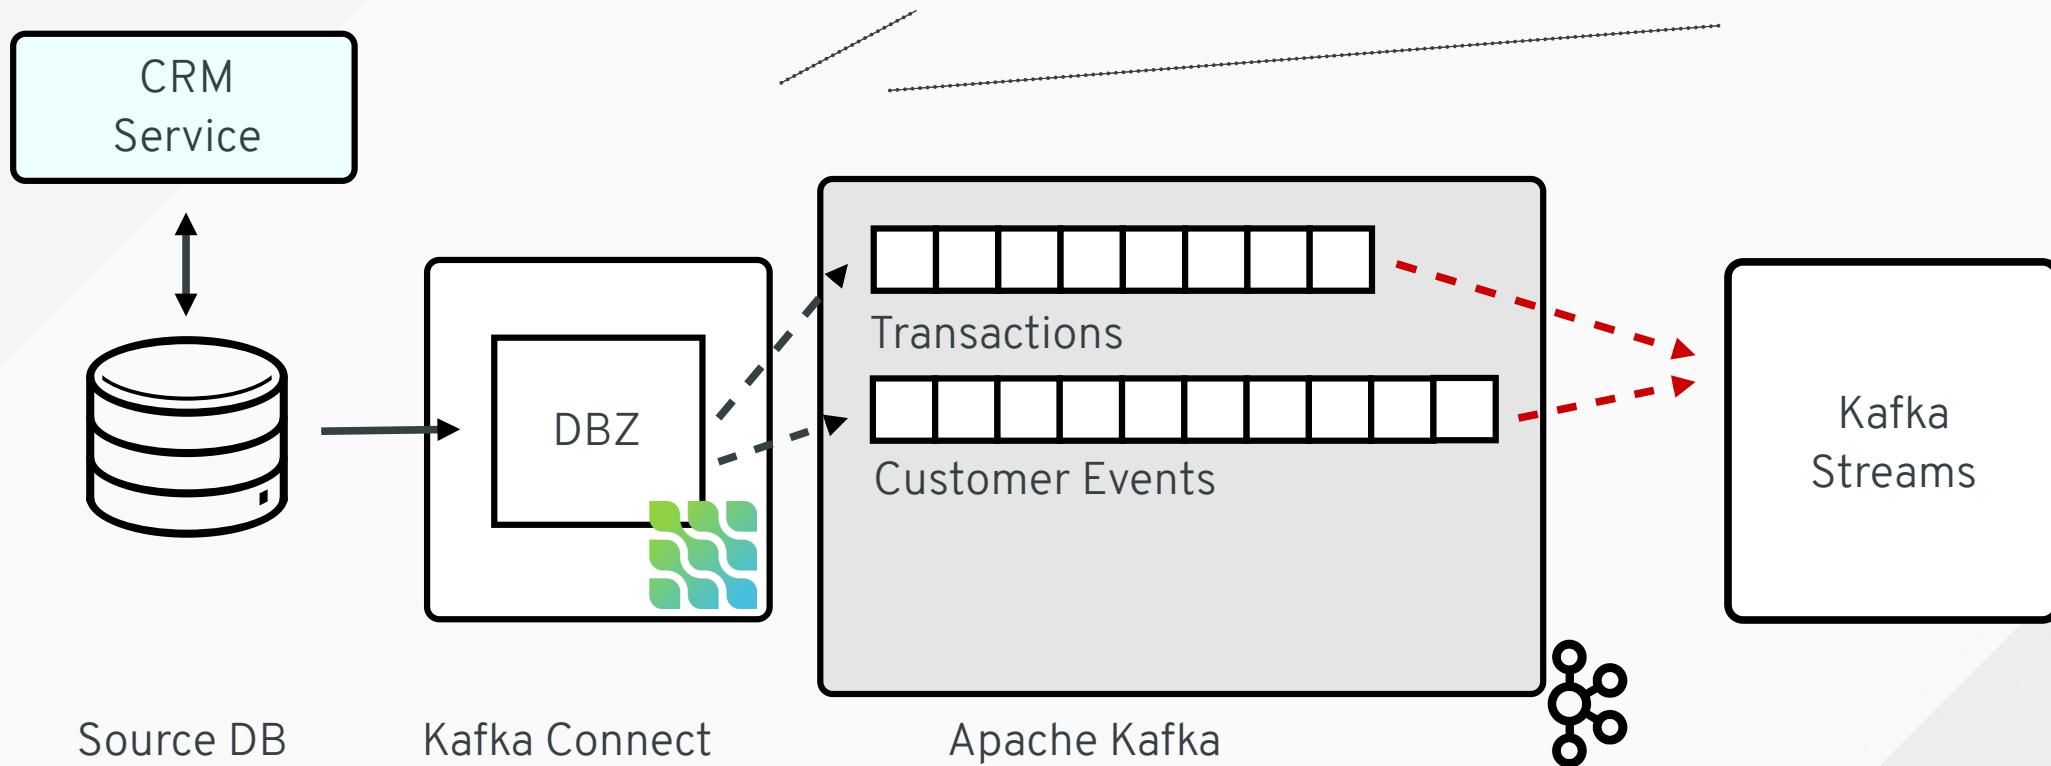




Auditing

"Transactions" table

Id	User	Use Case
tx-1	Bob	Create Customer
tx-2	Sarah	Delete Customer
tx-3	Rebecca	Update Customer

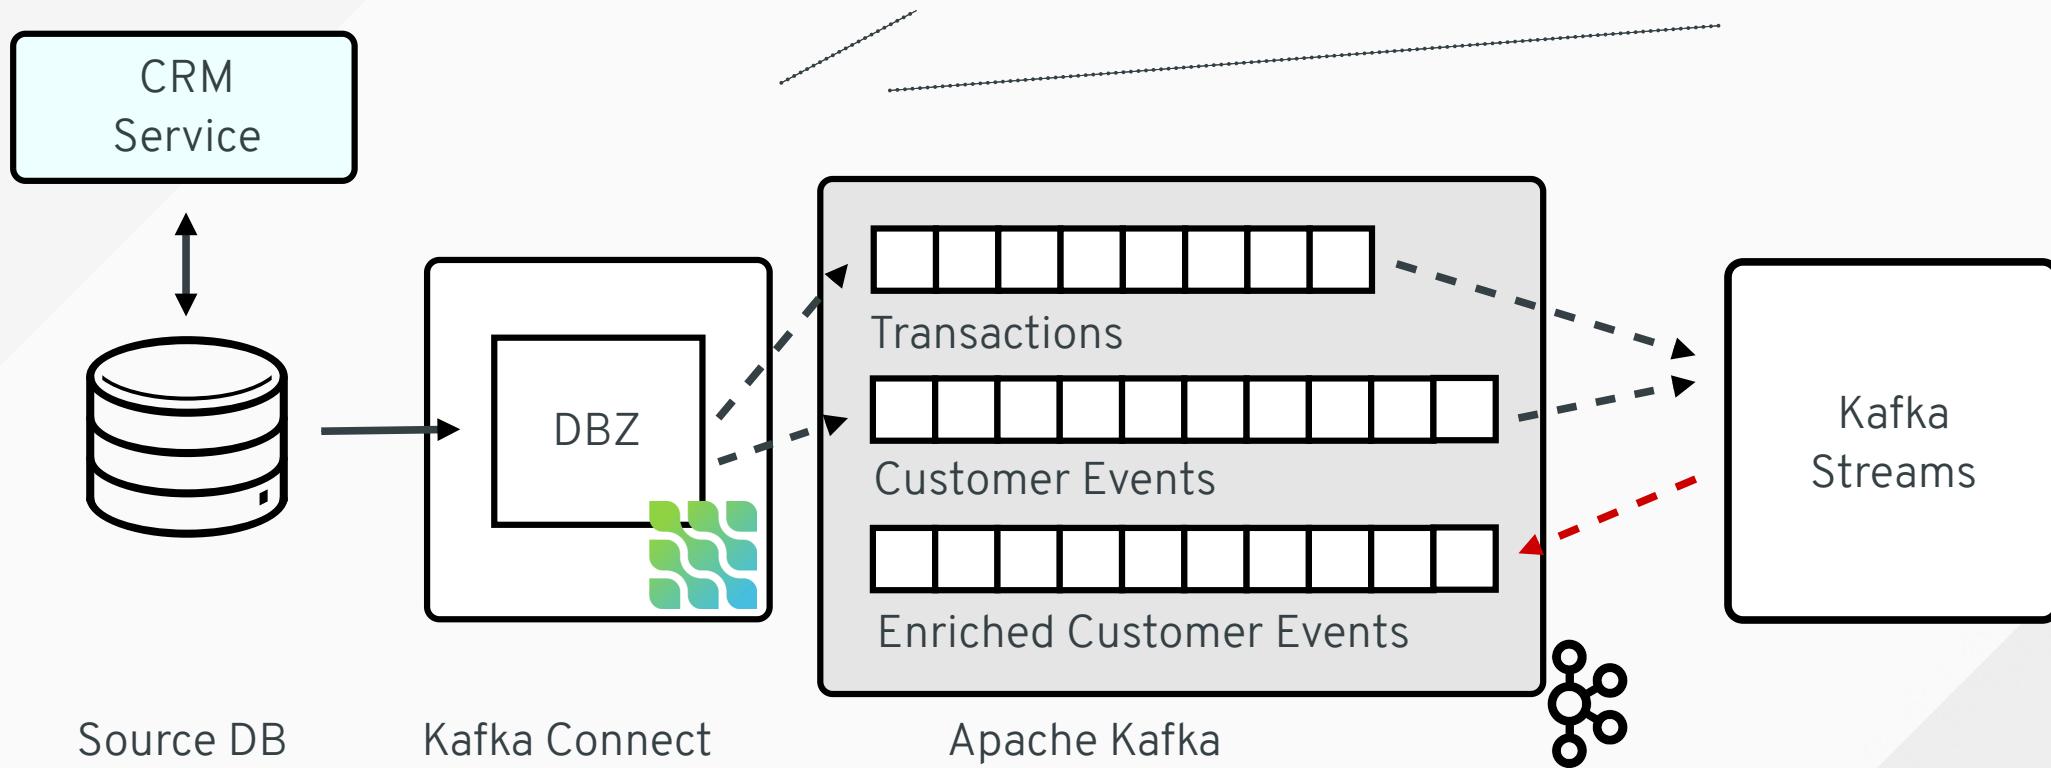




Auditing

"Transactions" table

Id	User	Use Case
tx-1	Bob	Create Customer
tx-2	Sarah	Delete Customer
tx-3	Rebecca	Update Customer





Auditing

```
{  
  "before": {  
    "id": 1004,  
    "last_name": "Kretchmar",  
    "email": "annek@example.com"  
  },  
  "after": {  
    "id": 1004,  
    "last_name": "Kretchmar",  
    "email": "annek@noanswer.org"  
  },  
  "source": {  
    "name": "dbserver1",  
    "table": "customers",  
    "txId": "tx-3"  
  },  
  "op": "u",  
  "ts_ms": 1486500577691  
}
```

Customers



```
{  
    "id": "tx-3"  
}  
  
{  
    "before": null,  
    "after": {  
        "id": "tx-3",  
        "user": "Rebecca",  
        "use_case": "Update customer"  
    },  
    "source": {  
        "name": "dbserver1",  
        "table": "transactions",  
        "txId": "tx-3"  
    },  
    "op": "c",  
    "ts_ms": 1486500577691  
}
```

Transactions

```
{  
    "before": {  
        "id": 1004,  
        "last_name": "Kretchmar",  
        "email": "annek@example.com"  
    },  
    "after": {  
        "id": 1004,  
        "last_name": "Kretchmar",  
        "email": "annek@noanswer.org"  
    },  
    "source": {  
        "name": "dbserver1",  
        "table": "customers",  
        "txId": "tx-3"  
    },  
    "op": "u",  
    "ts_ms": 1486500577691  
}
```

Customers



```
{  
  "id": "tx-3"  
}  
  
{  
  "before": null,  
  "after": {  
    "id": "tx-3",  
    "user": "Rebecca",  
    "use_case": "Update customer"  
  },  
  "source": {  
    "name": "dbserver1",  
    "table": "transactions",  
    "txId": "tx-3"  
  },  
  "op": "c",  
  "ts_ms": 1486500577691  
}
```

Transactions

```
{  
  "before": {  
    "id": 1004,  
    "last_name": "Kretchmar",  
    "email": "annek@example.com"  
  },  
  "after": {  
    "id": 1004,  
    "last_name": "Kretchmar",  
    "email": "annek@noanswer.org"  
  },  
  "source": {  
    "name": "dbserver1",  
    "table": "customers",  
    "txId": "tx-3"  
  },  
  "op": "u",  
  "ts_ms": 1486500577691  
}
```

Customers



Auditing

```
{  
  "before": {  
    "id": 1004,  
    "last_name": "Kretchmar",  
    "email": "annek@example.com"  
  },  
  "after": {  
    "id": 1004,  
    "last_name": "Kretchmar",  
    "email": "annek@noanswer.org"  
  },  
  "source": {  
    "name": "dbserver1",  
    "table": "customers",  
    "txId": "tx-3",  
    "user": "Rebecca",  
    "use_case": "Update customer"  
  },  
  "op": "u",  
  "ts_ms": 1486500577691  
}
```

Enriched Customers



Auditing

- Non-trivial join implementation
 - **no ordering** across topics
 - need to **buffer change events** until TX data available
- bit.ly/debezium-auditlogs

```
@Override
public KeyValue<JsonObject, JsonObject>
    transform(JsonObject key, JsonObject value) {

    boolean enrichedAllBufferedEvents =
        enrichAndEmitBufferedEvents();

    if (!enrichedAllBufferedEvents) {
        bufferChangeEvent(key, value);
        return null;
    }

    KeyValue<JsonObject, JsonObject> enriched =
        enrichWithTxMetaData(key, value);
    if (enriched == null) {
        bufferChangeEvent(key, value);
    }

    return enriched;
}
```

Expanding Partial Update Events





Expanding Partial Update Events

Examples

- MongoDB update events ("patch")
- Postgres
 - Replica identity **not FULL**
 - **TOAST-ed** columns
- Cassandra update events
- MySQL with row image minimal

```
{  
  "before": { ... },  
  "after": {  
    "id": 1004,  
    "first_name": "Dana",  
    "last_name": "Kretchmar",  
    "email": "annek@noanswer.org",  
    "biography":  
      "__debezium_unavailable_value"  
  },  
  "source": { ... },  
  "op": "u",  
  "ts_ms": 1570448151611  
}
```



Expanding Partial Update Events

Examples

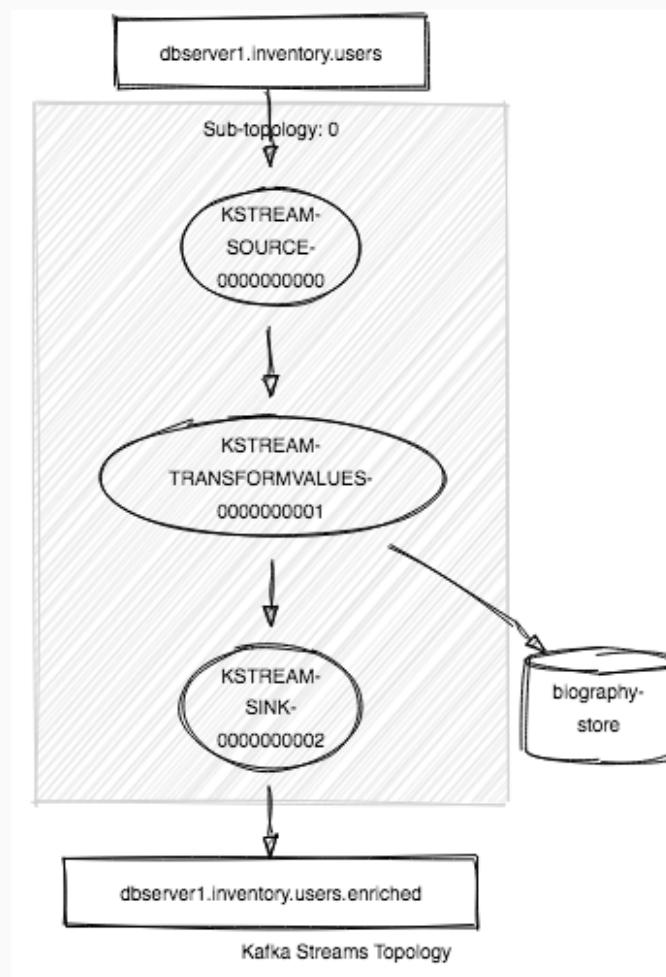
- MongoDB update events ("patch")
- Postgres
 - Replica identity **not FULL**
 - **TOAST-ed** columns
- Cassandra update events
- MySQL with row image minimal

```
{  
  "before": { ... },  
  "after": {  
    "id": 1004,  
    "first_name": "Dana",  
    "last_name": "Kretchmar",  
    "email": "annek@noanswer.org",  
    "biography":  
      "__debezium_unavailable_value"  
  },  
  "source": { ... },  
  "op": "u",  
  "ts_ms": 1570448151611  
}
```



Expanding Partial Update Events

Topology



[https://zz85.github.io/
kafka-streams-viz/](https://zz85.github.io/kafka-streams-viz/)



Expanding Partial Update Events

Obtaining missing values from a state store

```
class ToastColumnValueProvider implements ValueTransformerWithKey<JsonObject, JsonObject, JsonObject>

    private KeyValueStore<JsonObject, String> biographyStore;

    @Override
    public void init(ProcessorContext context) {
        biographyStore = (KeyValueStore<JsonObject, String>) context.getStateStore(
            TopologyProducer.BIOGRAPHY_STORE);
    }

    @Override
    public JsonObject transform(JsonObject key, JsonObject value) {
        // ...
    }
}
```



Expanding Partial Update Events

Obtaining missing values from a state store

```
class ToastColumnValueProvider implements ValueTransformerWithKey<JsonObject, JsonObject, JsonObject>

    private KeyValueStore<JsonObject, String> biographyStore

    @Override
    public void init(ProcessorContext context) {
        biographyStore = (KeyValueStore<JsonObject, String>) context.getStateStore(
            TopologyProducer.BIOGRAPHY_STORE);
    }

    @Override
    public JsonObject transform(JsonObject key, JsonObject value) {
        // ...
    }
}
```



Expanding Partial Update Events

Obtaining missing values from a state store

```
JsonObject payload = value.getJsonObject("payload");
JsonObject newRowState = payload.getJsonObject("after");
String biography = newRowState.getString("biography");

if (isUnavailableValueMarker(biography)) {
    String currentValue = biographyStore.get(key);
    newRowState = Json.createObjectBuilder(newRowState)
        .add("biography", currentValue)
        .build();
    // ...
}
else {
    biographyStore.put(key, biography);
}
return value;
```

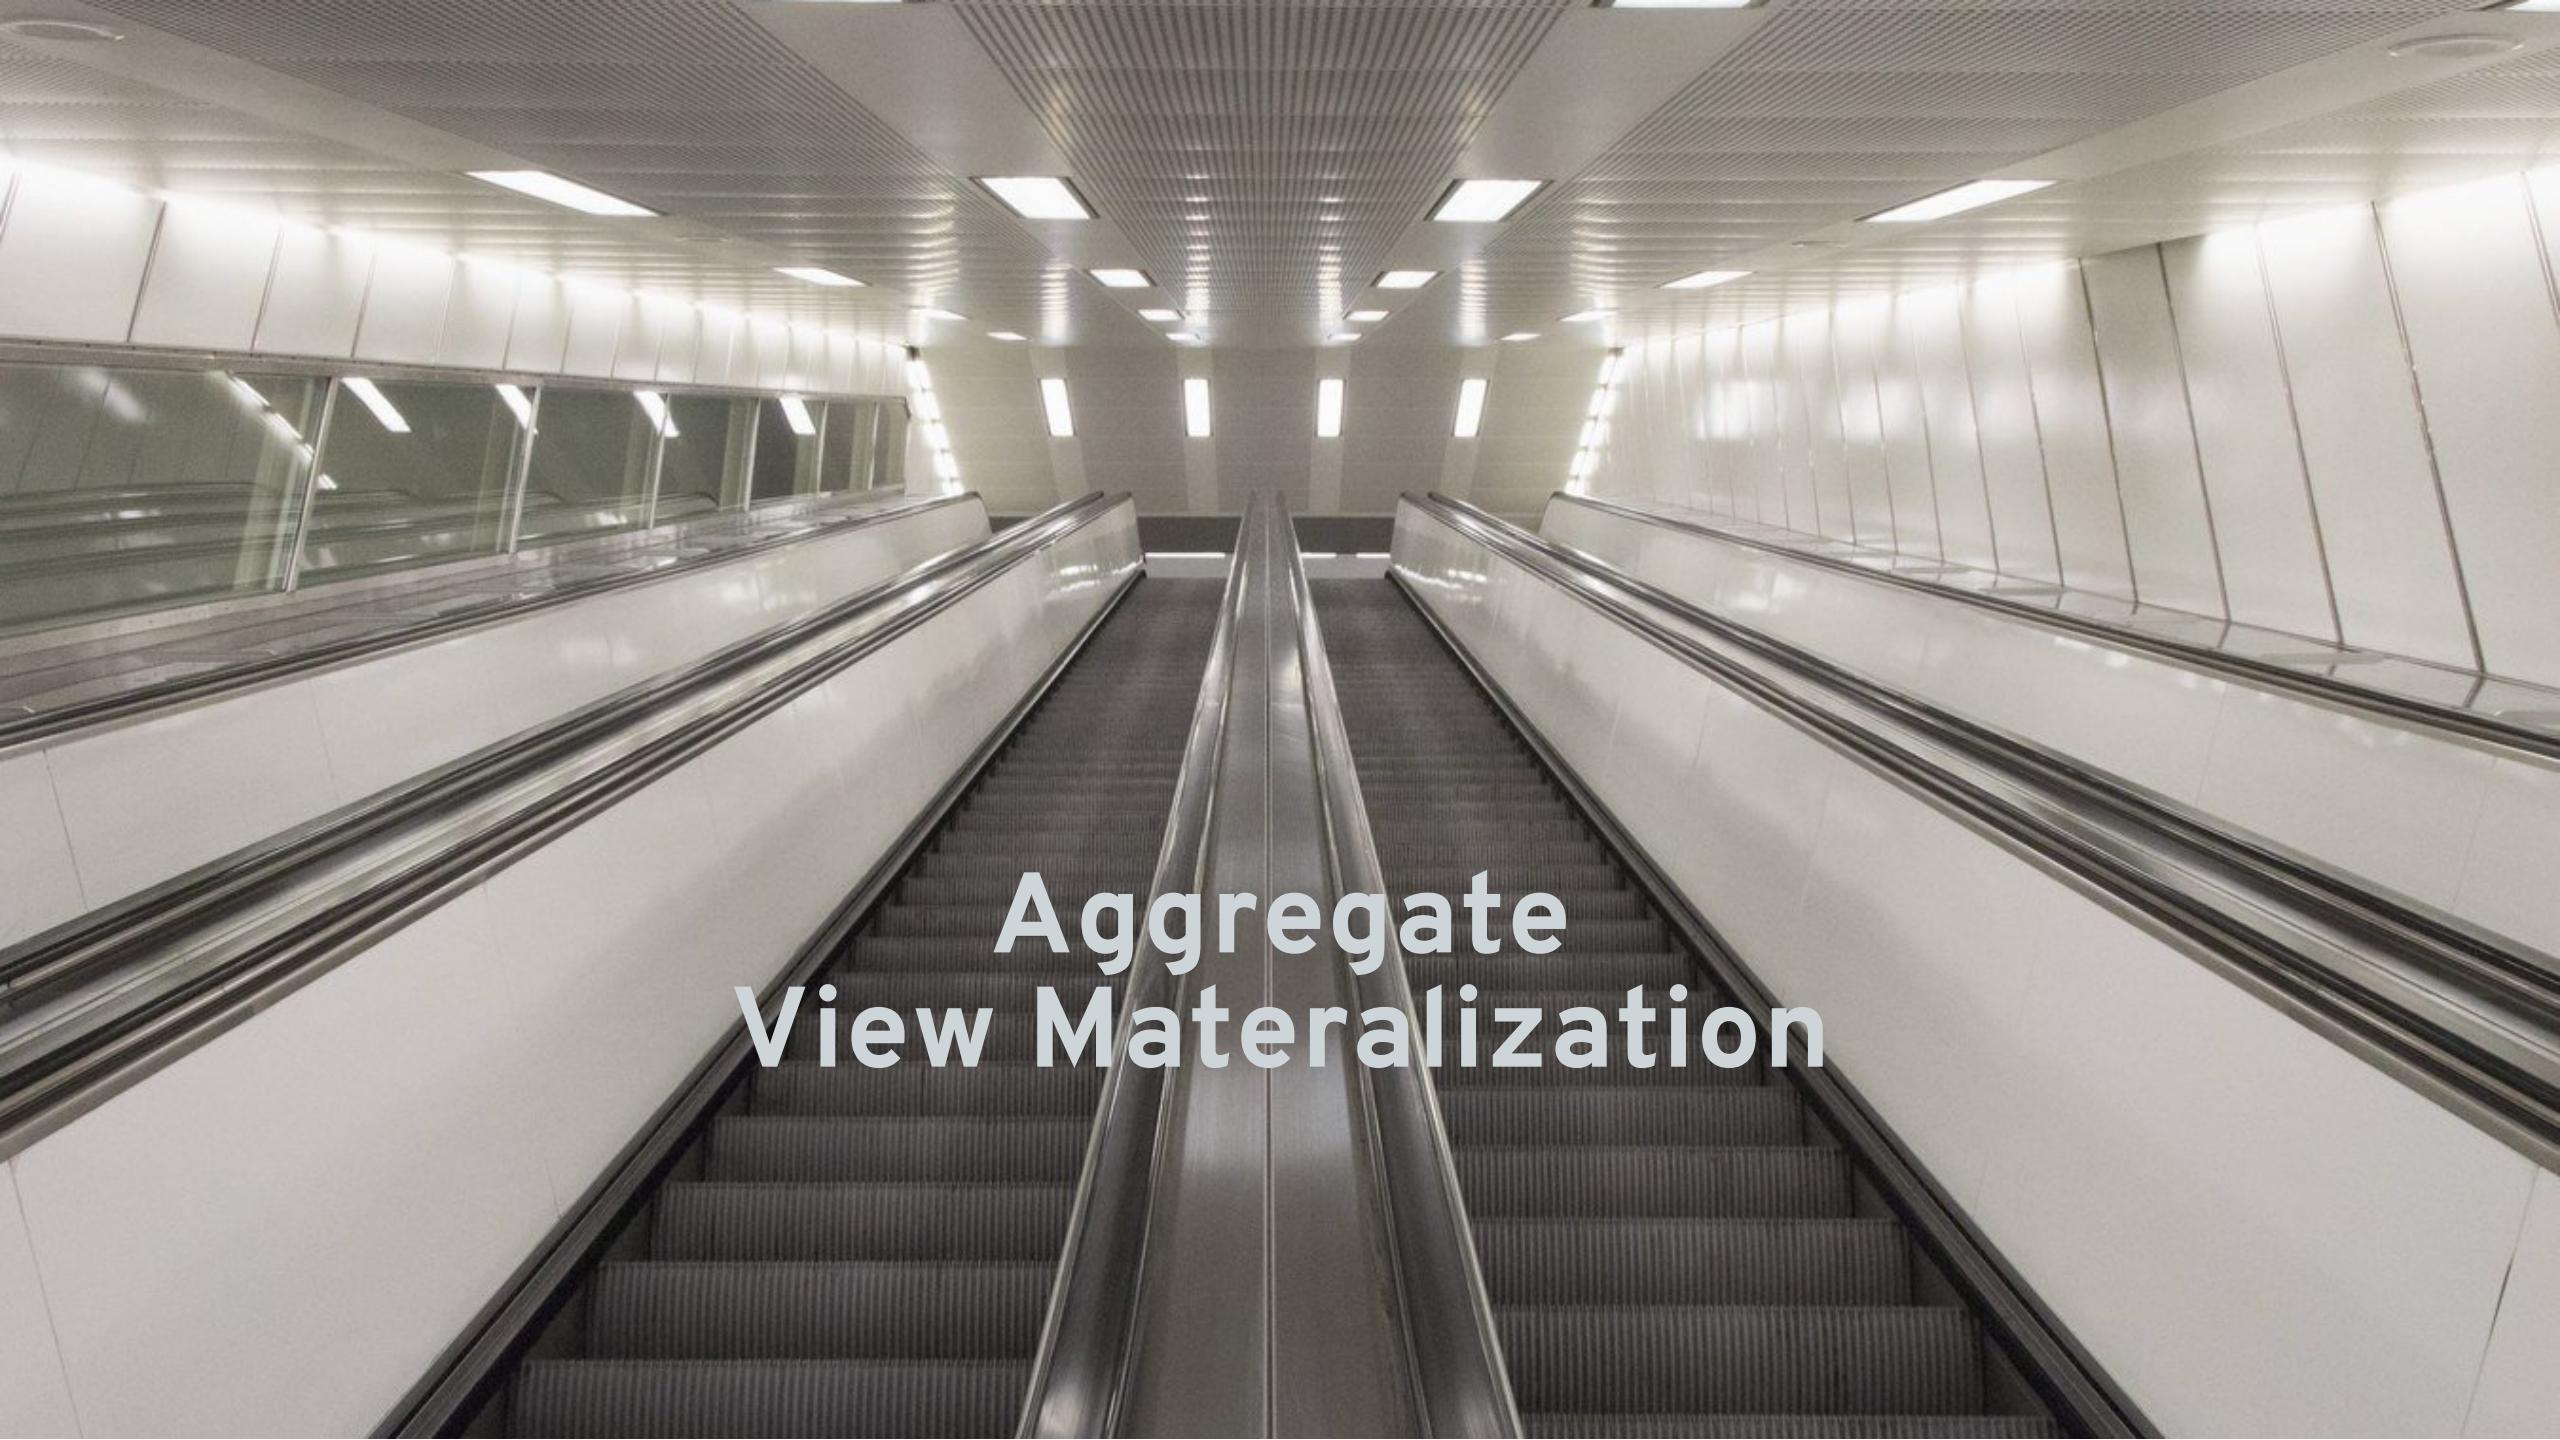


Expanding Partial Update Events

Obtaining missing values from a state store

```
JsonObject payload = value.getJsonObject("payload");
JsonObject newRowState = payload.getJsonObject("after");
String biography = newRowState.getString("biography");

if (isUnavailableValueMarker(biography)) {
    String currentValue = biographyStore.get(key);
    newRowState = Json.createObjectBuilder(newRowState)
        .add("biography", currentValue)
        .build();
    // ...
}
else {
    biographyStore.put(key, biography)
}
return value;
```

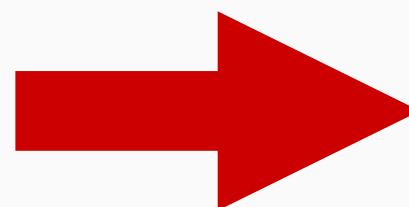
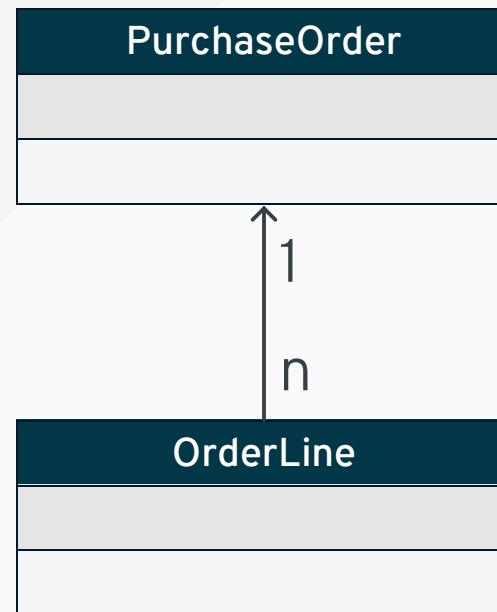
A perspective view of a modern subway station platform. Two parallel escalators lead down from the platform level. The ceiling is made of a grid of recessed lights. The walls are light-colored and feature vertical panels. The overall atmosphere is clean and minimalist.

Aggregate View Materialization



Aggregate View Materialization

From Multiple Topics to One View



```
{  
  "purchaseOrderId" : "order-123",  
  "orderDate" : "2020-08-24",  
  "customerId": "customer-123",  
  "orderLines" : [  
    {  
      "orderLineId" : "orderLine-456",  
      "productId" : "product-789",  
      "quantity" : 2,  
      "price" : 59.99  
    },  
    {  
      "orderLineId" : "orderLine-234",  
      "productId" : "product-567",  
      "quantity" : 1,  
      "price" : 14.99  
    }  
  ]  
}
```



Aggregate View Materialization

Non-Key Joins (KIP-213)

```
KTable<Long, OrderLine> orderLines = ...;
KTable<Integer, PurchaseOrder> purchaseOrders = ...;

KTable<Integer, PurchaseOrderWithLines> purchaseOrdersWithOrderLines = orderLines
    .join(
        purchaseOrders,
        orderLine -> orderLine.purchaseOrderId,
        (orderLine, purchaseOrder) -> new OrderLineAndPurchaseOrder(orderLine, purchaseOrder))
    .groupBy(
        (orderLineId, lineAndOrder) -> KeyValue.pair(lineAndOrder.purchaseOrder.id, lineAndOrder))
    .aggregate(
        PurchaseOrderWithLines::new,
        (Integer key, OrderLineAndPurchaseOrder value, PurchaseOrderWithLines agg) -> agg.addLine(value),
        (Integer key, OrderLineAndPurchaseOrder value, PurchaseOrderWithLines agg) -> agg.removeLine(value)
    );
```



Aggregate View Materialization

Non-Key Joins (KIP-213)

```
KTable<Long, OrderLine> orderLines = ...;  
KTable<Integer, PurchaseOrder> purchaseOrders = ...;  
  
KTable<Integer, PurchaseOrderWithLines> purchaseOrdersWithOrderLines = orderLines  
.join(  
    purchaseOrders,  
    orderLine -> orderLine.purchaseOrderId,  
    (orderLine, purchaseOrder) -> new OrderLineAndPurchaseOrder(orderLine, purchaseOrder))  
.groupBy(  
    (orderLineId, lineAndOrder) -> KeyValue.pair(lineAndOrder.purchaseOrder.id, lineAndOrder))  
.aggregate(  
    PurchaseOrderWithLines::new,  
    (Integer key, OrderLineAndPurchaseOrder value, PurchaseOrderWithLines agg) -> agg.addLine(value),  
    (Integer key, OrderLineAndPurchaseOrder value, PurchaseOrderWithLines agg) -> agg.removeLine(value)  
);
```



Aggregate View Materialization

Non-Key Joins (KIP-213)

```
KTable<Long, OrderLine> orderLines = ...;  
KTable<Integer, PurchaseOrder> purchaseOrders = ...;  
  
KTable<Integer, PurchaseOrderWithLines> purchaseOrdersWithOrderLines = orderLines  
.join(  
    purchaseOrders,  
    orderLine -> orderLine.purchaseOrderId,  
    (orderLine, purchaseOrder) -> new OrderLineAndPurchaseOrder(orderLine, purchaseOrder))  
.groupBy(  
    (orderLineId, lineAndOrder) -> KeyValue.pair(lineAndOrder.purchaseOrder.id, lineAndOrder))  
.aggregate(  
    PurchaseOrderWithLines::new,  
    (Integer key, OrderLineAndPurchaseOrder value, PurchaseOrderWithLines agg) -> agg.addLine(value),  
    (Integer key, OrderLineAndPurchaseOrder value, PurchaseOrderWithLines agg) -> agg.removeLine(value)  
);
```



Aggregate View Materialization

Non-Key Joins (KIP-213)

```
KTable<Long, OrderLine> orderLines = ...;  
KTable<Integer, PurchaseOrder> purchaseOrders = ...;  
  
KTable<Integer, PurchaseOrderWithLines> purchaseOrdersWithOrderLines = orderLines  
.join(  
    purchaseOrders,  
    orderLine -> orderLine.purchaseOrderId,  
    (orderLine, purchaseOrder) -> new OrderLineAndPurchaseOrder(orderLine, purchaseOrder))  
.groupBy(  
    (orderLineId, lineAndOrder) -> KeyValue.pair(lineAndOrder.purchaseOrder.id, lineAndOrder))  
.aggregate(  
    PurchaseOrderWithLines::new,  
    (Integer key, OrderLineAndPurchaseOrder value, PurchaseOrderWithLines agg) -> agg.addLine(value),  
    (Integer key, OrderLineAndPurchaseOrder value, PurchaseOrderWithLines agg) -> agg.removeLine(value)  
);
```



Aggregate View Materialization

Non-Key Joins (KIP-213)

```
KTable<Long, OrderLine> orderLines = ...;  
KTable<Integer, PurchaseOrder> purchaseOrders = ...;  
  
KTable<Integer, PurchaseOrderWithLines> purchaseOrdersWithOrderLines = orderLines  
.join(  
    purchaseOrders,  
    orderLine -> orderLine.purchaseOrderId,  
    (orderLine, purchaseOrder) -> new OrderLineAndPurchaseOrder(orderLine, purchaseOrder))  
.groupBy(  
    (orderLineId, lineAndOrder) -> KeyValue.pair(lineAndOrder.purchaseOrder.id, lineAndOrder))  
.aggregate(  
    PurchaseOrderWithLines::new,  
    (Integer key, OrderLineAndPurchaseOrder value, PurchaseOrderWithLines agg) -> agg.addLine(value),  
    (Integer key, OrderLineAndPurchaseOrder value, PurchaseOrderWithLines agg) -> agg.removeLine(value)  
>;
```



Aggregate View Materialization

Awareness of Transaction Boundaries

- TX metadata in change events (e.g. dbserver1.inventory.orderline)

```
{  
    "before": null,  
    "after": { ... },  
    "source": { ... },  
    "op": "c",  
    "ts_ms": "1580390884335",  
    "transaction": {  
        "id": "571",  
        "total_order": "1",  
        "data_collection_order": "1"  
    }  
}
```



Aggregate View Materialization

Awareness of Transaction Boundaries

- Topic with **BEGIN/END markers**
- Enable consumers to buffer all events of one transaction

```
{  
  "transactionId" : "571",  
  "eventType" : "begin transaction",  
  "ts_ms" : 1486500577125  
}
```

BEGIN

```
{  
  "transactionId" : "571",  
  "ts_ms" : 1486500577691,  
  "eventType" : "end transaction",  
  "eventCount" : [  
    {  
      "name" : "dbserver1.inventory.order",  
      "count" : 1  
    },  
    {  
      "name" : "dbserver1.inventory.orderLine",  
      "count" : 5  
    }  
  ]  
}
```

END



Takeaways

Many Use Cases for Debezium and Kafka Streams

- Data **enrichment**
- Creating **aggregated events**
- **Stream queries**
- **Interactive query** services for legacy databases



Takeaways

Many Use Cases for Debezium and Kafka Streams

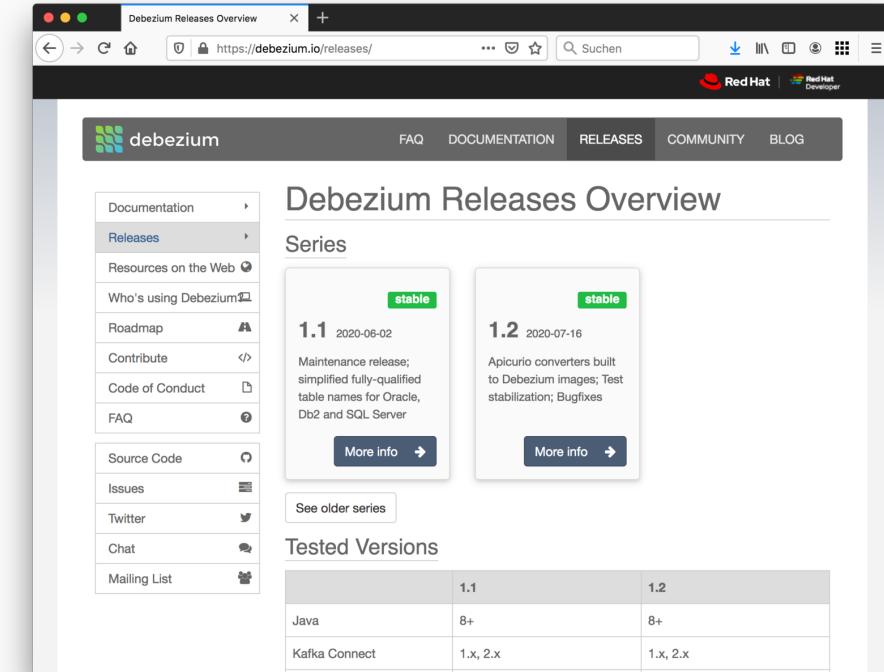
- Data **enrichment**
- Creating **aggregated events**
- **Stream queries**
- **Interactive query** services for legacy databases

Debezium + Kafka
Streams = ❤️



Resources

- **Website:** debezium.io
- **Examples:** github.com/debezium/debezium-examples/
- **Latest news:**  [@debezium](https://twitter.com/debezium)



The screenshot shows the Debezium Releases Overview page. At the top, there's a navigation bar with links for Documentation, Releases (which is currently selected), Resources on the Web, Who's using Debezium, Roadmap, Contribute, Code of Conduct, FAQ, Source Code, Issues, Twitter, Chat, and Mailing List. Below the navigation is a search bar and a "Suchen" button. The main content area is titled "Debezium Releases Overview". It features a section for "Series" with two cards: "1.1" (stable, 2020-06-02) and "1.2" (stable, 2020-07-16). Each card includes a "More info" button. Below this is a "See older series" link. A "Tested Versions" table follows, comparing Java and Kafka Connect versions 1.1 and 1.2. The table has columns for Java (8+), 1.1, 1.2, Kafka Connect (1.x, 2.x), and 1.1, 1.2 respectively.

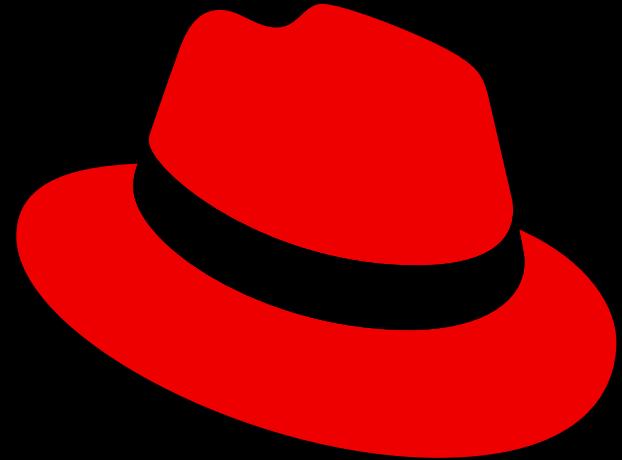
	1.1	1.2
Java	8+	8+
Kafka Connect	1.x, 2.x	1.x, 2.x



Q & A

✉ gunnar@hibernate.org

🐦 [@gunnarmorling](https://twitter.com/gunnarmorling)



Red Hat