
A decorative background graphic consisting of a network of nodes and edges. The nodes are represented by circles of varying sizes and colors (blue, grey, white), and the edges are thin grey lines connecting them. The network is distributed across the slide, with a denser cluster on the left and a more sparse one on the bottom right.

Distributed consensus and fault tolerance

Lecture 1 / 2

Georgios Bitzes, CERN
iCSC 2017



Let's talk about distributed systems

Lecture 1

- Introduction to distributed systems
- Replication and split-brain
- Strong vs eventual consistency
- The raft consensus algorithm

Lecture 2


- Two Generals, Byzantine Generals
- Byzantine fault tolerance
- Bitcoin and blockchain consensus

The background of the slide is a light gray network graph. It consists of numerous nodes, represented by small circles, some of which are solid gray and others are hollow with a gray outline. These nodes are interconnected by a web of thin, light gray lines representing edges. The overall pattern is dense and non-uniform, suggesting a complex, distributed system.

Introduction to distributed systems


Distributed systems are all around us

- Infrastructure:
 - Networking, routing algorithms
 - Flight control systems
 - Banking systems, ATMs
- Internet services: running these on a single machine would be unthinkable
 - Facebook, Twitter
 - Google search, gmail
 - Github
- But the need for distributed systems appears long before we have to scale to millions of users



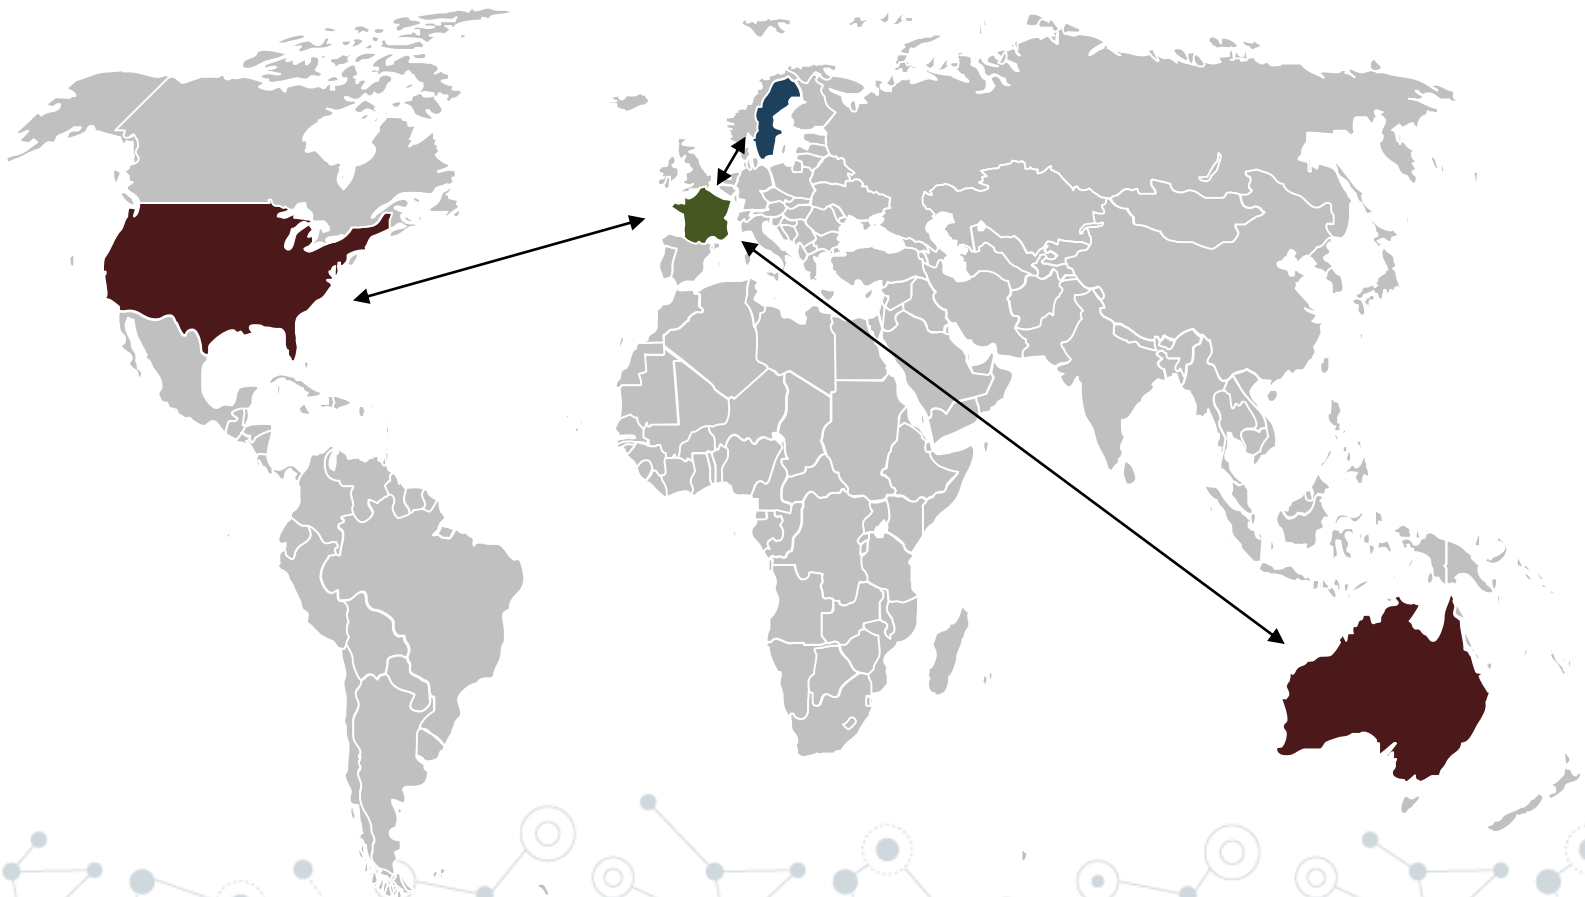
Scalability

- As the load increases, one server no longer enough to serve all clients
- We're bound to hit some bottleneck
 - CPU power
 - Memory size
 - Memory bandwidth
 - IO operations per second
 - Network bandwidth
 - Software, more often than not cannot fully exploit our powerful hardware



Latency

With only a single server, some clients will suffer from high latency



Fault tolerance

- Our system has a single point of failure
 - one faulty hard drive and the service goes down
- Downtime of certain critical services can cause great disruption: credit card processing, air traffic control, or... stackoverflow.com



Distributed systems to the rescue

- A distributed system solves the above problems nicely
 - **Scalable performance:** add more machines as needed
 - **Lower latencies:** clients connect to the server nearest to them
 - **Fault tolerance:** if a machine goes down, the others can detect it and take over its responsibilities
- **Caveat:** distributed systems add *a lot* of complexity



What is a distributed system?

- Distributed system: A collection of independent computers that appears to its users as a **single, coherent** system
- Nodes coordinate by **exchanging messages** through the network



Distributed system: added complexity

Making a system distributed adds some complexity:

- **Unreliable network:** messages get lost, delayed, re-ordered, corrupted
- **Node failures:** the more machines we have, the higher the probability of some failing
- **Latency:** communication far slower than using local shared memory or IPC
- **Limited bandwidth**
- **Security:** “is this message really coming from who I think it is” ?



Multiple single points of failure?


- What if every node depended on every other for its correct operation?
- A single failure will bring them all down
- Just because a system is distributed, doesn't mean it's fault tolerant



“

A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable.

Leslie Lamport



The background of the slide is a light gray network graph. It consists of numerous nodes, represented by small circles, some of which are solid gray and others are hollow with a gray outline. These nodes are interconnected by a web of thin, light gray lines representing edges. The overall pattern is dense and non-uniform, filling the entire background.

Replication and split-brain

Replication as a means to fault tolerance

- A database machine holding 25% of all our data crashes and burns – hard drive is **unrecoverable**
- What happens next? Some possibilities...
 1. No backups, *data lost forever* – oops. Not acceptable, will create bad publicity, erode user trust
 2. A backup taken the previous day is manually restored. Much better, but new user data in the last 24h is lost, and *long downtime* during manual restore
 3. A replicated database machine takes over immediately after the failure – no data loss, no downtime, *users don't even notice*



Key-value stores (used in future examples)

- Simple database, clients can perform 2 operations
 - **write** a value into a key
 - **read** it back.

```
127.0.0.1:6379> SET favorite_food pickles
OK
127.0.0.1:6379> GET favorite_food
"pickles"
127.0.0.1:6379> SET favorite_language c++
OK
127.0.0.1:6379> GET favorite_language
"C++"
```

A naive replication protocol

- Assume we want our key-value store replicated on 3 nodes. 1 2
- Let's invent our own simple replication protocol: 3
 1. A client sends a write request: propagate the change to all other sibling nodes. Ignore errors.
 2. A client does a read: give back the local value stored on the contacted node
 3. On receipt of a propagated change from a sibling node: simply apply it by updating the local value.




A naive replication protocol (2)

- In simple cases and good network conditions, this protocol might appear to work.
 - <https://gbitzes.github.io/icsc/animation/#naiverep>
- What if a node **goes down** for maintenance for 5 minutes?
 - All writes within that window are not replicated onto it
- What if certain nodes receive the updates in a **different order**?



The problem: Split brain



The problem: Split brain (2)

- Our naive protocol will **inevitably** lead to split brain
- What we want: replicated nodes to agree on the state of the key-value store, that they all come to a **consensus** about each update
- We'll talk about a correct algorithm later



Failure model

The failure model we'll concern ourselves for now:

1. Fail-recover faults

- A node goes down, stops responding to messages.
- The other nodes can detect this through timeouts – “if 127.0.0.10 doesn't respond in 100ms, it's down”
- ... but failed nodes can recover


2. Delayed / lost messages: Messages between nodes can be arbitrarily delayed or lost, but **not** corrupted

○ Later on: **byzantine faults**





Network partitions

- Tolerance to delayed and lost messages implies tolerance to network partitions
- **Network partition:** parts of the system become disconnected from each other



Network partitions (2)

- Problematic because each part is making decisions independently
- What if clients try to write different values to each partitioned part? Which value is “correct”?





The background of the slide features a complex, repeating pattern of a network graph. It consists of numerous small, light-blue circular nodes connected by thin, light-blue lines representing edges. Some nodes are solid circles, while others are circles with a smaller concentric circle inside, creating a layered or hierarchical appearance. The connections between nodes form a dense, interconnected web that fills the entire background.

Strong vs eventual consistency

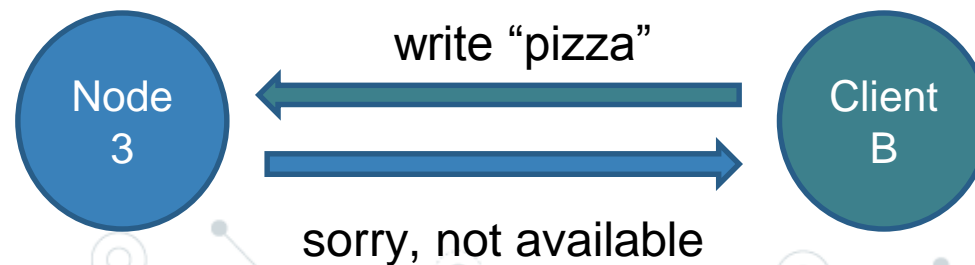

CAP theorem

- When faced with Network Partitions (P), a distributed system can be either Consistent (C) or Available (A)



Strong consistency

- System appears **externally consistent** to clients
- Requests are **refused** if consistency cannot be guaranteed





Strong consistency (2)

- Internal, temporary inconsistencies: **inevitable** in distributed systems
- Strong consistency: internal inconsistencies are resolved and **not exposed** to clients
- The cost: **sacrifice** availability in favor of consistency



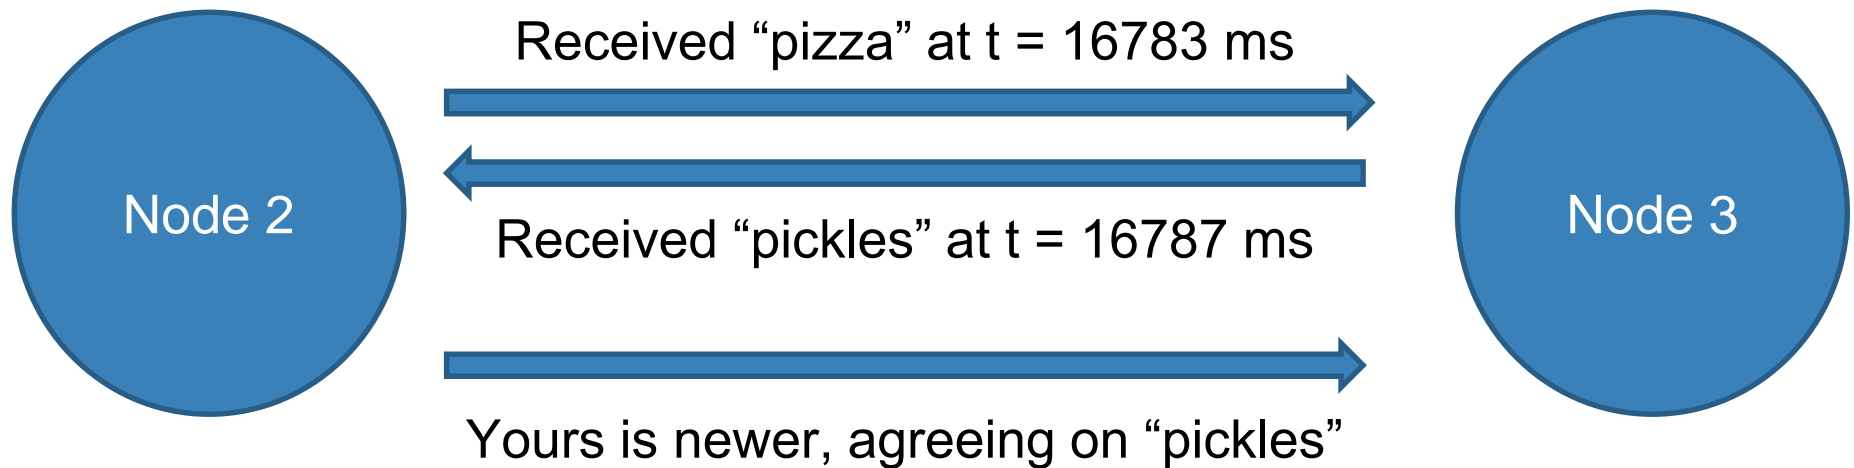
Eventual consistency

- Best-effort: try to be consistent, but without guarantees
- Internal inconsistencies **may be exposed** to clients



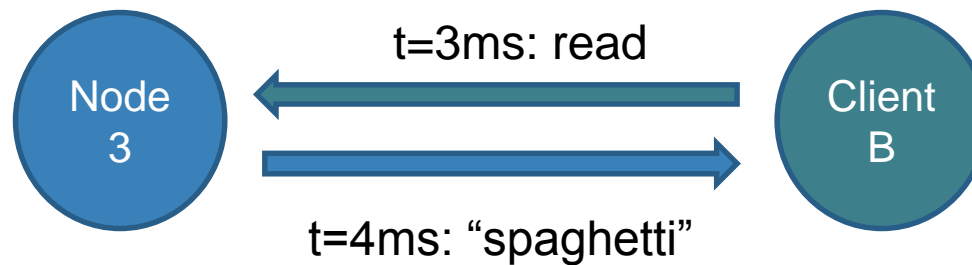

Eventual consistency (2)

- Conflicts are resolved after partition heals
- Common approach: *“last writer wins”*



Eventual consistency (3)

- Clients may receive stale values



Which one is better?

- Depends on the **application**, both are useful
- Eventual consistency: generally more **performant** and **scalable**
 - example: DNS, Amazon S3
- Strong consistency: **safety guarantee**, every read receives the most recent write or error
 - example: certain SQL databases



Question

Which form of consistency does our naive replication protocol provide?

1. Strong consistency
 2. Eventual consistency
 3. Neither
- Correct answer: neither.
 - In our protocol, inconsistencies are never resolved – not even “eventually”




The background of the slide features a complex, repeating pattern of a network graph. It consists of numerous small circular nodes, some of which are solid grey and others are hollow with a grey outline. These nodes are interconnected by a web of thin, light-grey lines representing edges. The overall effect is a dense, textured background that suggests a distributed system or a complex network topology.

The raft consensus algorithm


Distributed consensus

- A number of nodes come to an **agreement** about a **value**
- Any node can **crash** or **recover** at any time



The raft consensus algorithm

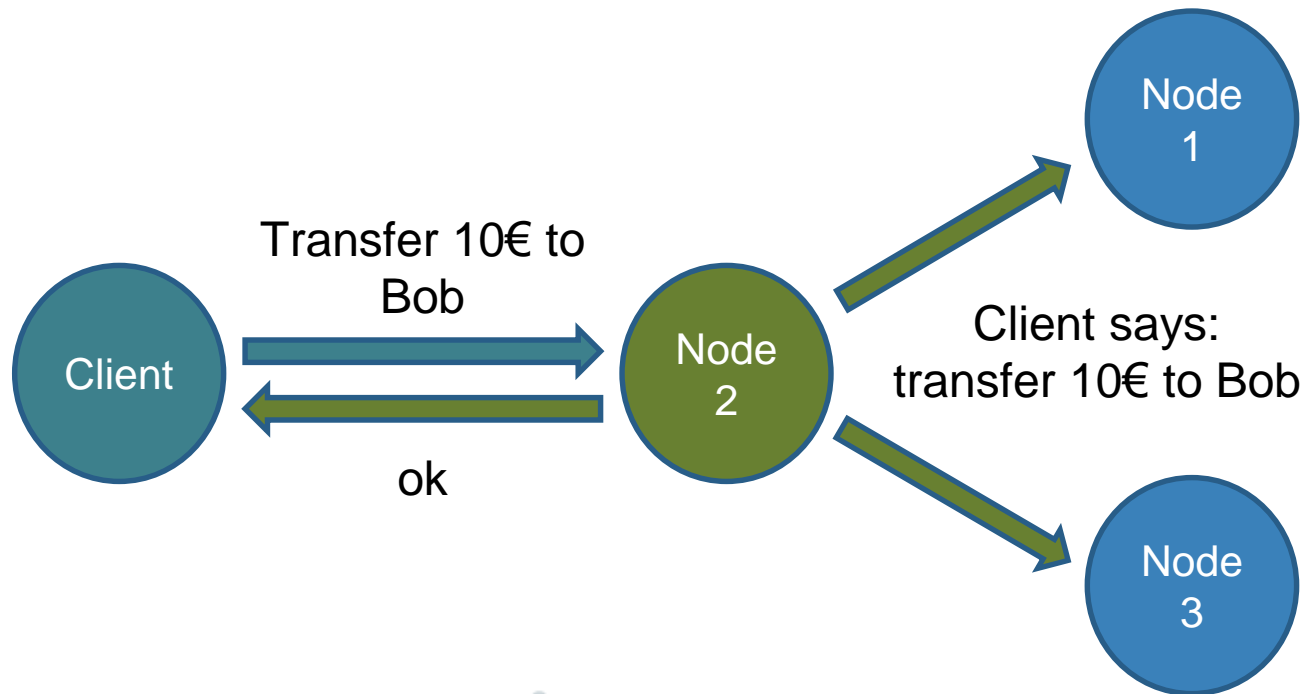
- Raft is a proven, correct consensus algorithm, most suitable for strongly consistent system



- Other consensus algorithms also exist – most notably paxos


Master – slave replication

- One of the nodes is elected to become the *master* (or *leader*)



Master – slave replication (2)

- Easy way to give a **global ordering** to writes from clients – can reconcile simultaneous updates
- Clients contacting followers are **redirected**




Leader failover

- The leader sends regular heartbeats to all followers
- If a follower stops receiving heartbeats, it assumes leader failure and triggers an **election**
- An election is won if a node receives positive votes from at least a **majority** of the cluster



Heartbeats

Haven't heard from
the leader for
2 sec... Something
is wrong



Follower




Election terms

- A leader is elected for a specific **election term** based on **majority vote**
- Only **one** leader can be elected per term
- Possible to have multiple leaders at a time, due to network partitions – all will have **different** terms
- The leader with the **highest** term “wins” and can override decisions made by the others



Leader election

A successful election: 2 out of 3 nodes agree on the new leader

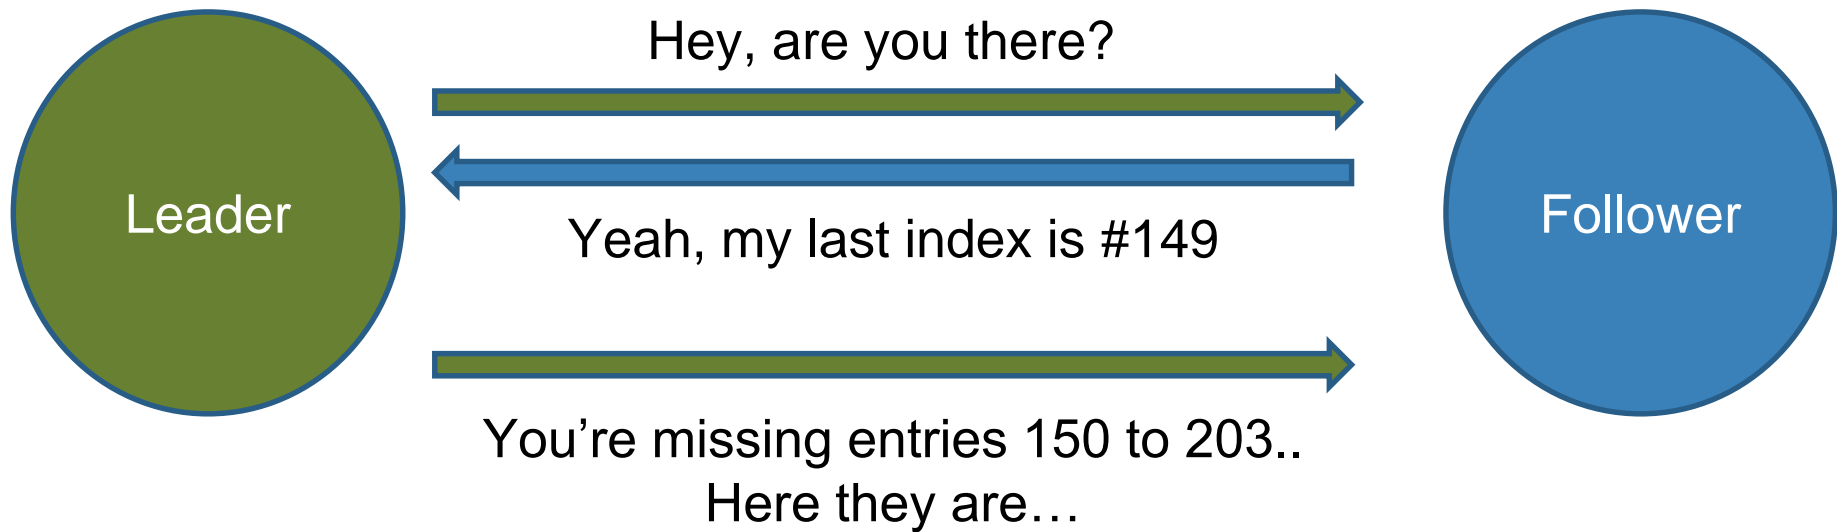


Bringing outdated nodes up-to-date

- One of the followers goes offline for 10 minutes – how to bring it up-to-date?
- Record all writes into an indexed log, and replicate it
- Contains also the term during which the entry was recorded

Index	Term	Contents
0	1	SET food pizza
1	1	SET language c++
2	1	SET food pickles
3	5	SET answer_to_life 42
...	

Log replication



Application of log entries

- Log entries represent changes to the database (also called the state machine)
- At some point, these changes have to be applied to the state machine

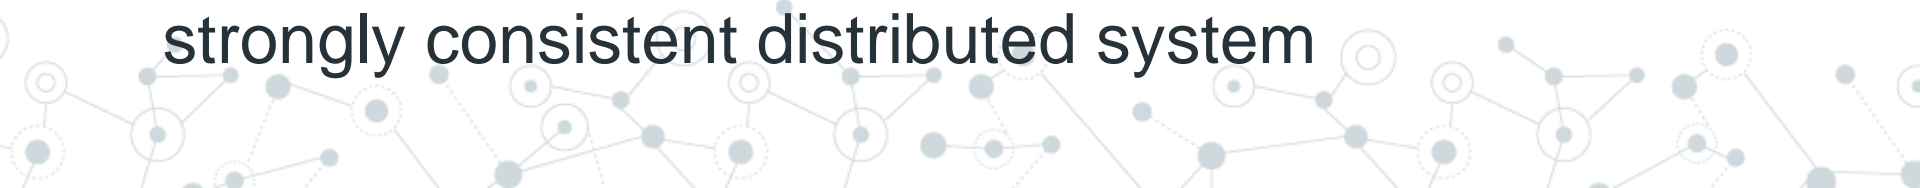


State machine replication

- Not safe to apply an entry immediately, since in rare cases it might get **rolled back** by a subsequent leader
- Log entries are considered **committed** once a majority of nodes have them
- Once an entry is committed, it's guaranteed that it won't be rolled back, and can be safely applied to the state machine
- Only nodes that have all committed entries can ever succeed during a leader election



Summary

- Distributed systems are necessary to provide scalability, low latency, and fault tolerance
 - Replicated nodes can offer transparent failover, but be careful of split brain
 - Several approaches to replication, need to compromise between consistency, availability, scalability, and more
 - The Raft consensus algorithm offers a good basis for a strongly consistent distributed system
- 
- A decorative network diagram at the bottom of the slide, featuring a series of interconnected nodes and lines, representing a distributed system topology. The nodes are represented by circles of varying sizes, some solid and some outlined, connected by thin lines. The overall pattern is a complex, interconnected web.

Bedtime reading

- PACELC theorem, extension to CAP theorem
https://en.wikipedia.org/wiki/PACELC_theorem
- The raft paper
<https://raft.github.io/raft.pdf>
- Paxos algorithm
[https://en.wikipedia.org/wiki/Paxos_\(computer_science\)](https://en.wikipedia.org/wiki/Paxos_(computer_science))
- Strong consistency models
<https://aphyr.com/posts/313-strong-consistency-models>



Thanks

Any questions?

See you in Lecture 2

Georgios.bitzes@cern.ch

