



Everything You Wanted to Know About Distributed Tracing.

Strange Loop Conference

September 2019

AJUA

Who Am I



- Hungai Amuhinda
- Twitter: [@Hungai](https://twitter.com/@Hungai)
- Nairobi, Kenya
- Work: Ajua, Infra Team
- Website: hungaikev.in

How to use this talk

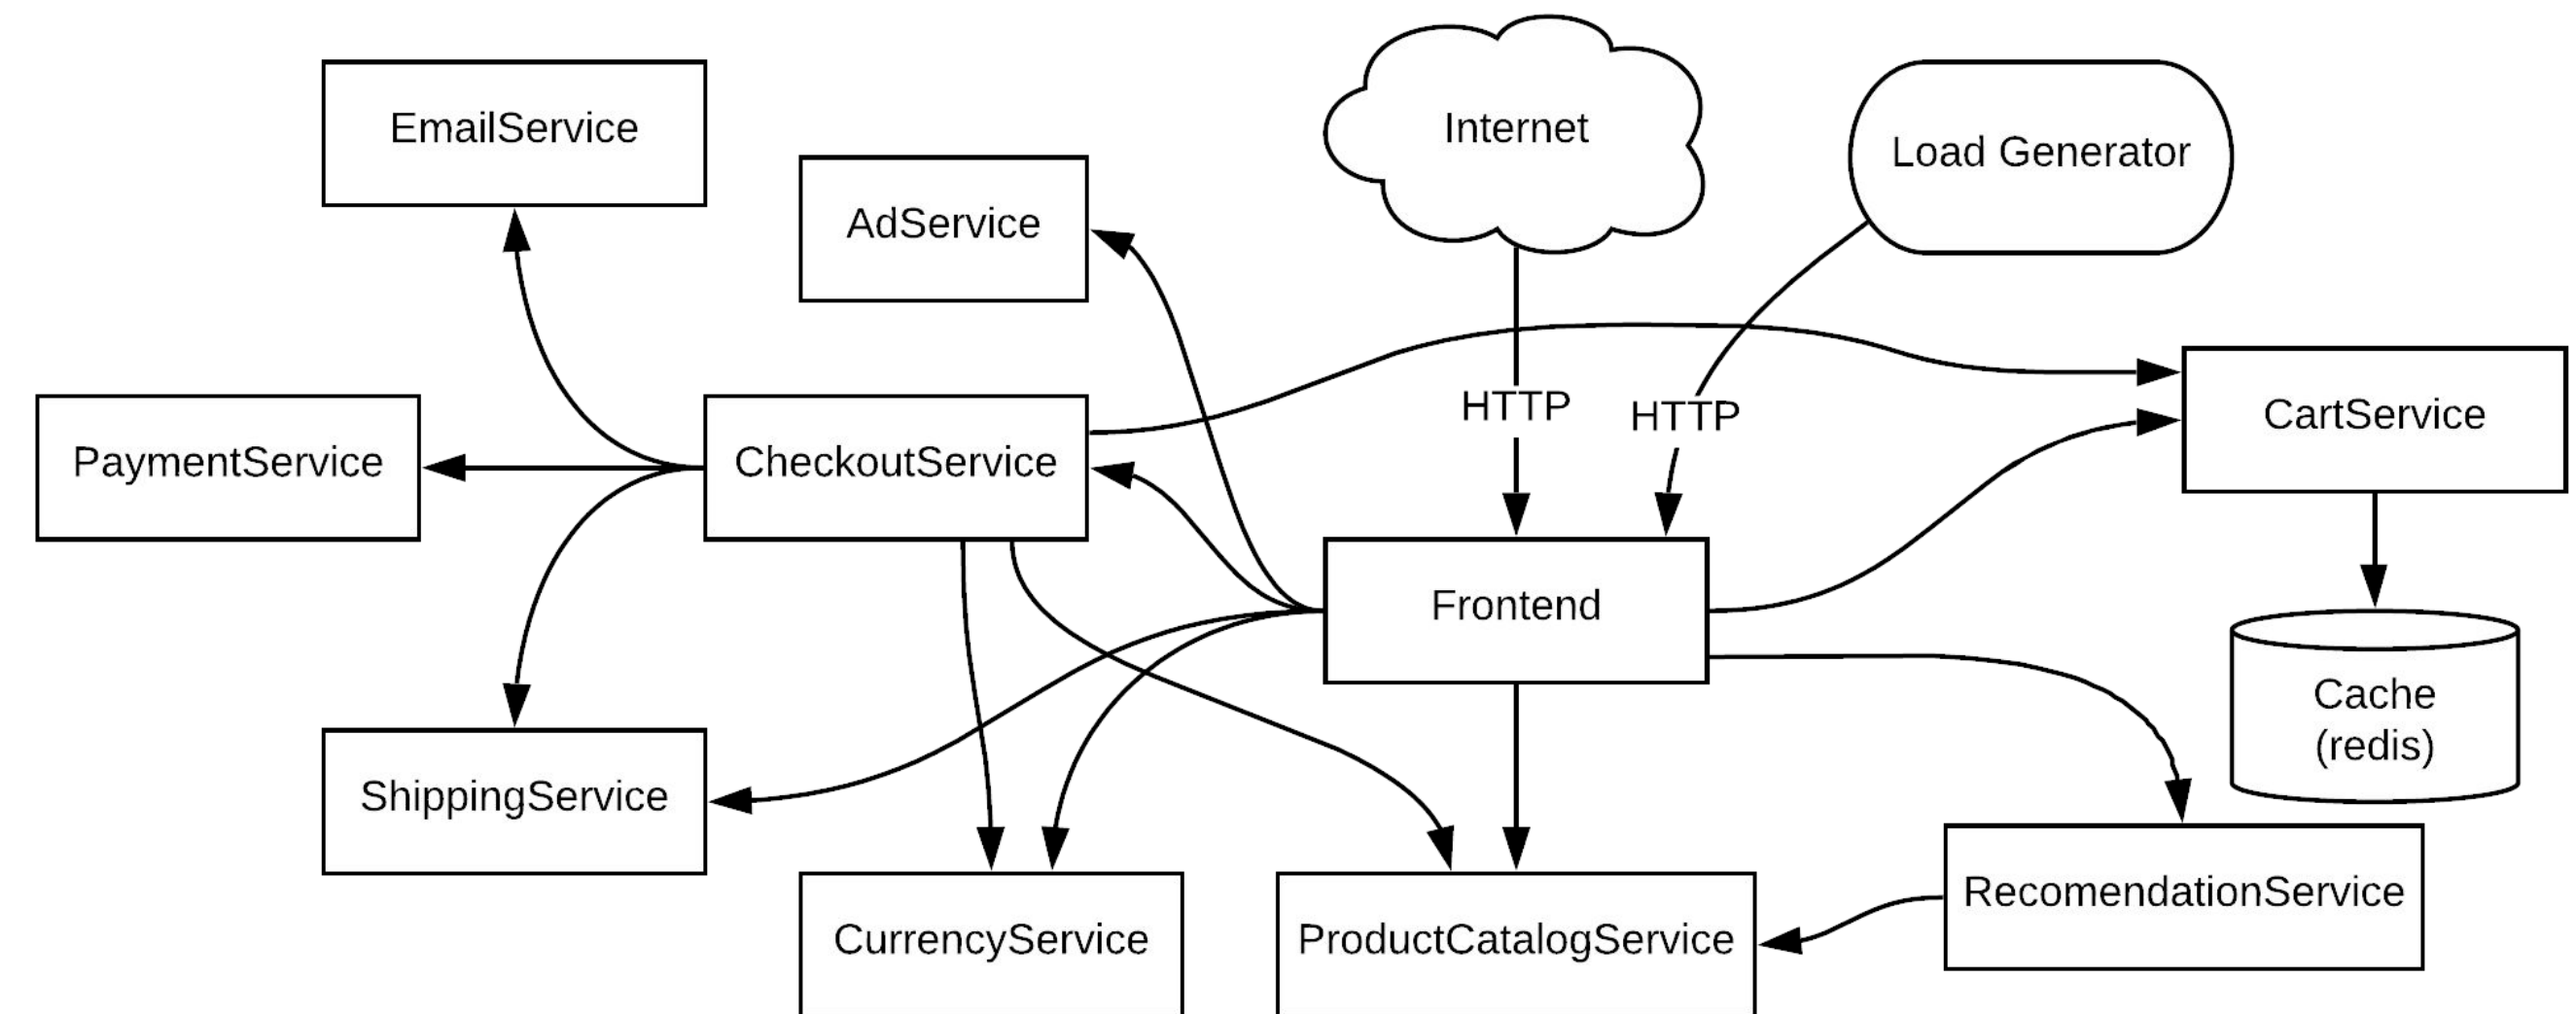
- **Please don't just follow it blindly:** There are often times when you will need to do things differently.
- **Software is all about trade-offs:** Very few decisions are about right and wrong.
- **Try things for yourself:** Why not make Friday afternoons a time to play and experiment?

Software Evolution.

- Monolith
- On Prem
- Single Language
- Single Stack
- Virtual Machines

Software Evolution.

- Microservices
- Containers
- Multi Cloud/ Hybrid
- Polyglot
- Containers
- Serverless/ Cloud Functions



New Architectures/ New Challenges.

- **Observability**
- Deployment / Packaging
- Configuration Management
- Debugging
- Secrets Management

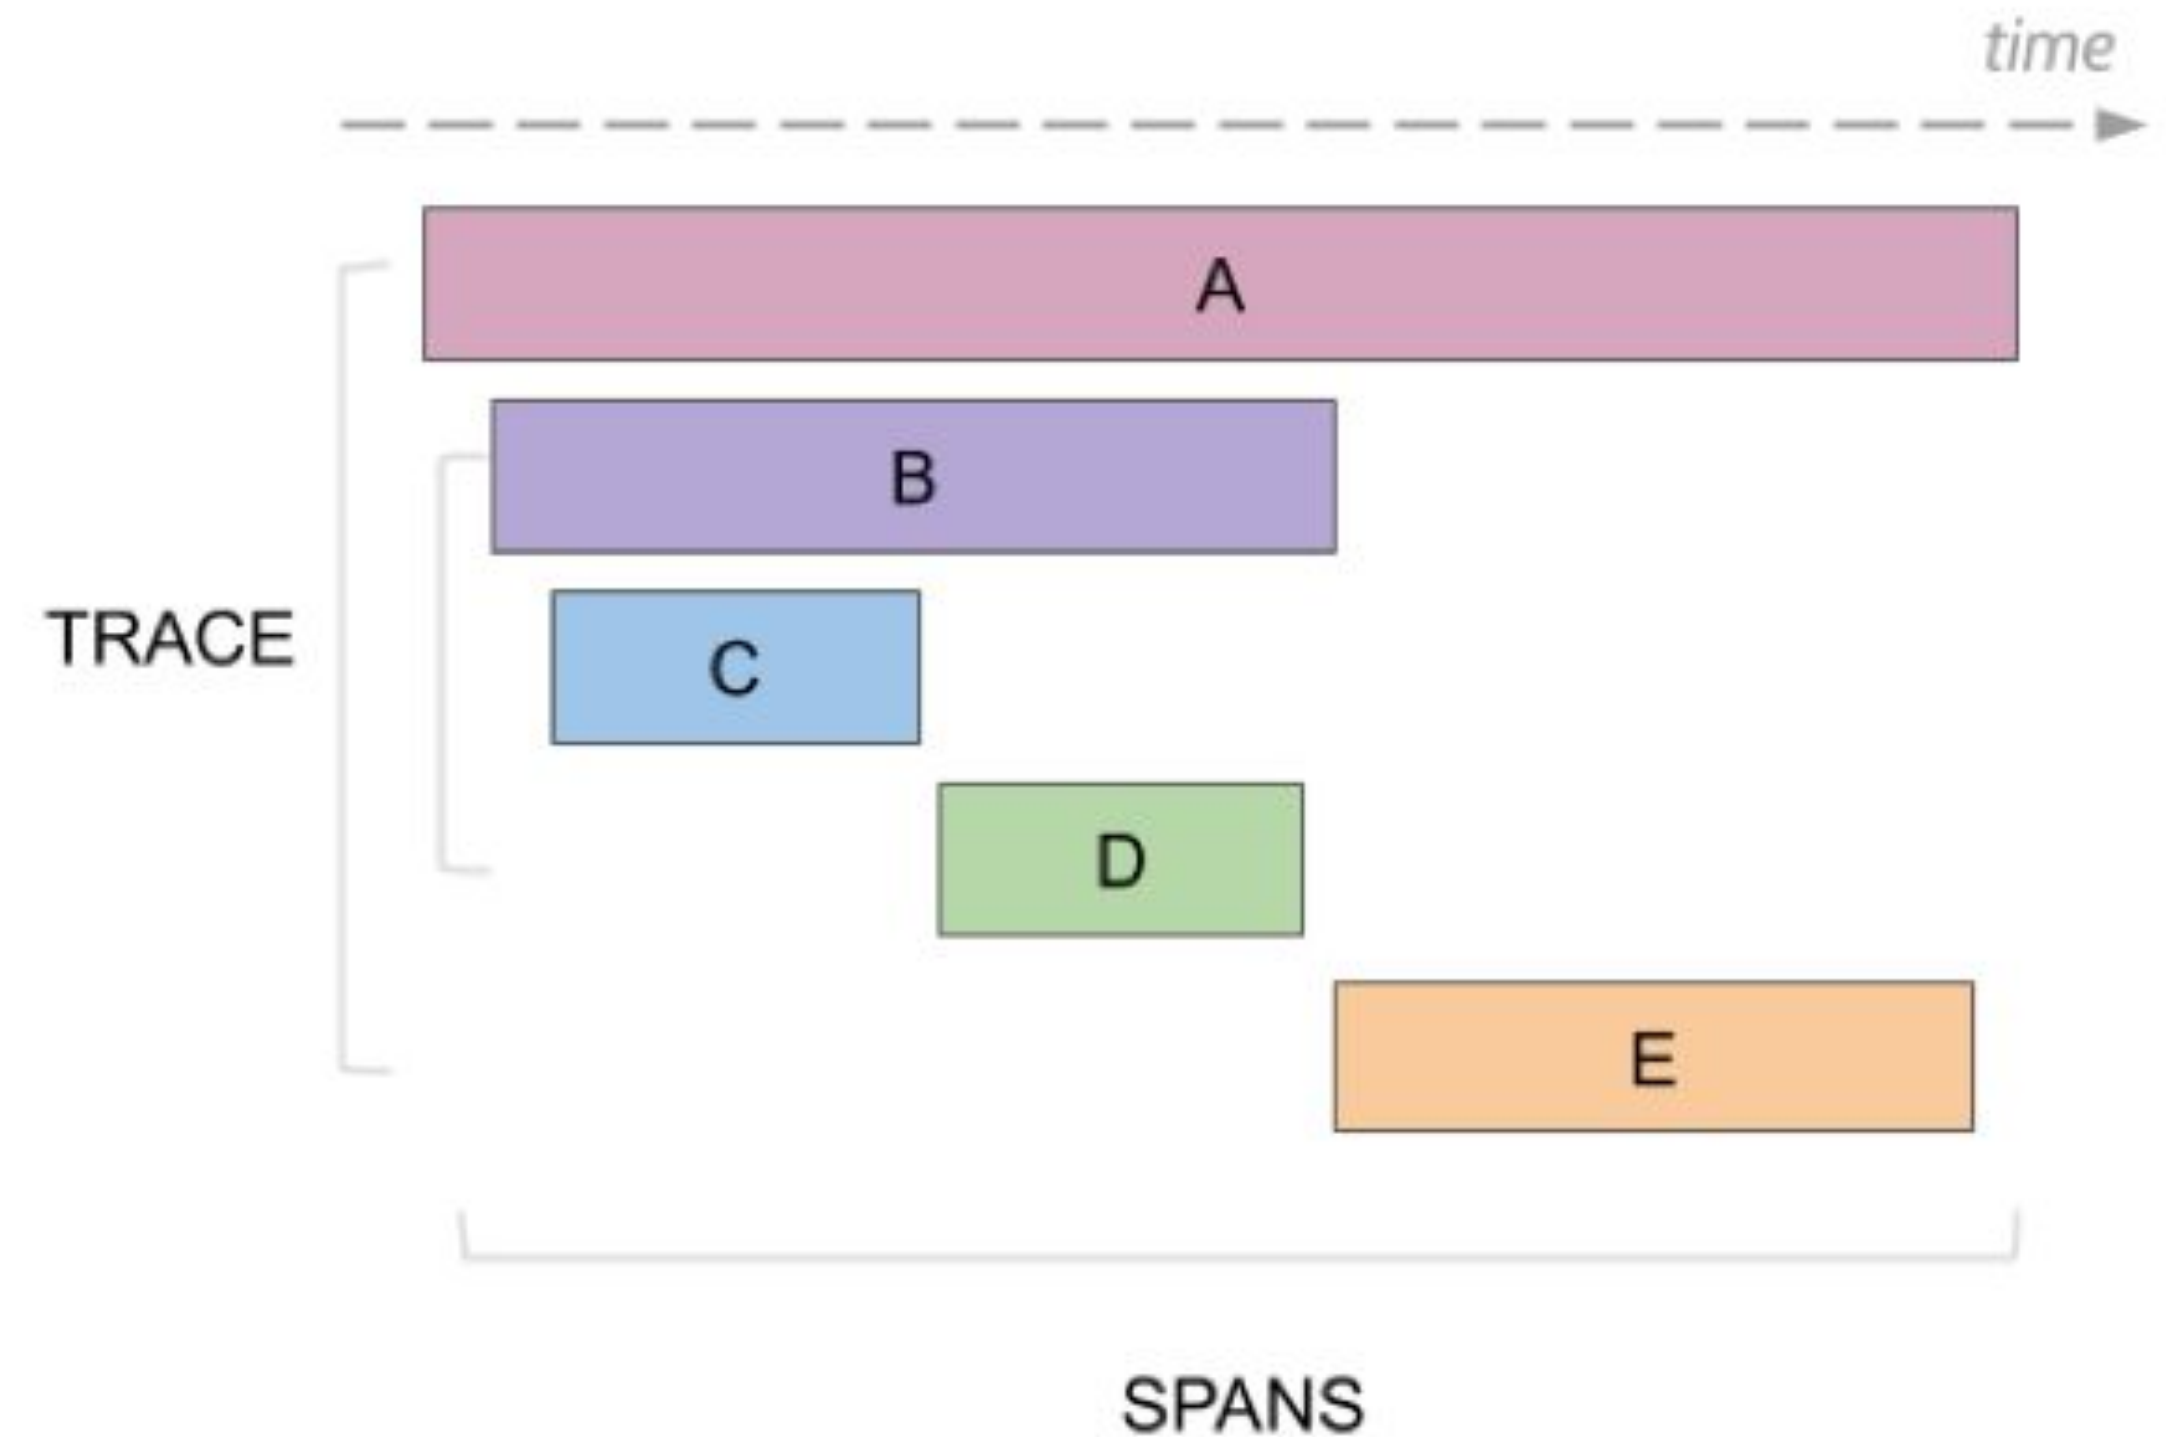
Meet: Distributed Tracing

“**Distributed Tracing**, also called distributed request tracing, is a method used to profile and monitor applications, especially those built using a microservices architecture. Distributed tracing helps pinpoint where failures occur and what causes poor performance.”

Distributed Tracing - Terminology

Trace - a trace is a tree of spans that follows the course of a request or system from its source to its ultimate destination.

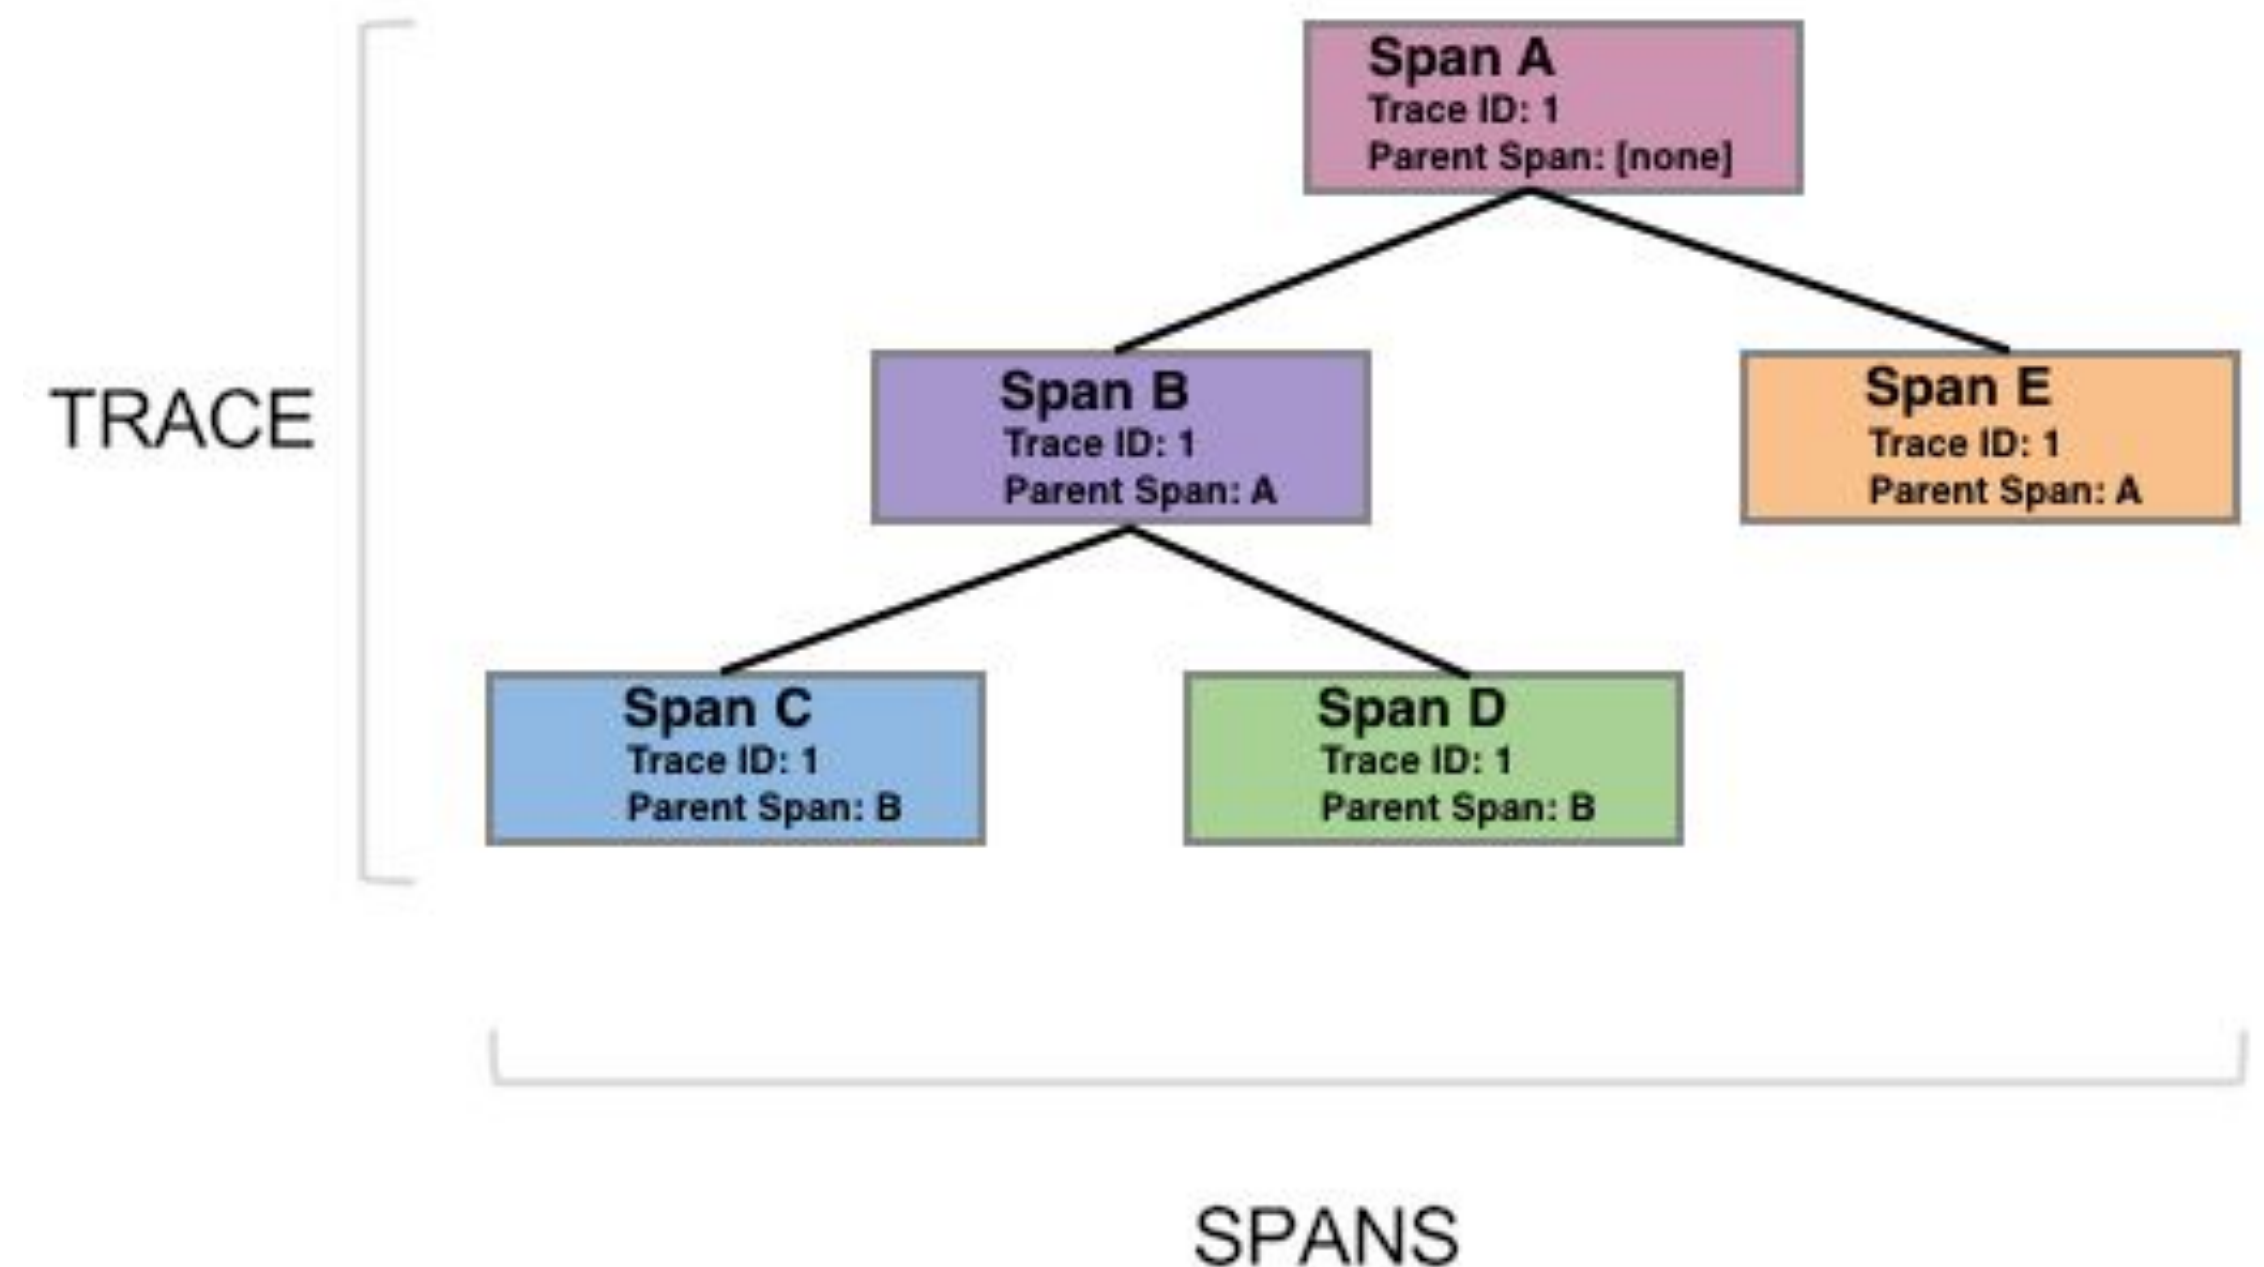
Each **trace** is a narrative that tells the requests story as it travels through the system.



Distributed Tracing - Terminology

Span - are logical units of work in a distributed system. They all have a name, a start time, and a duration.

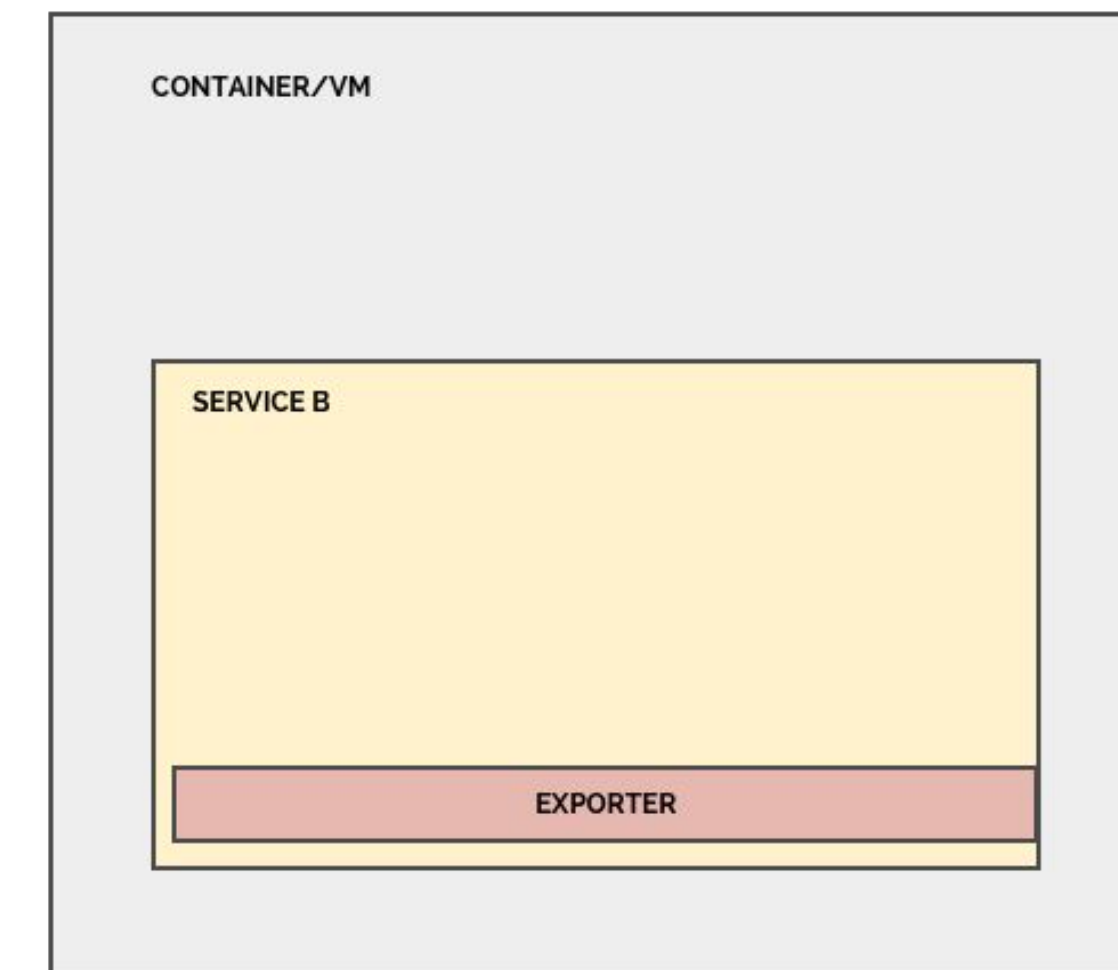
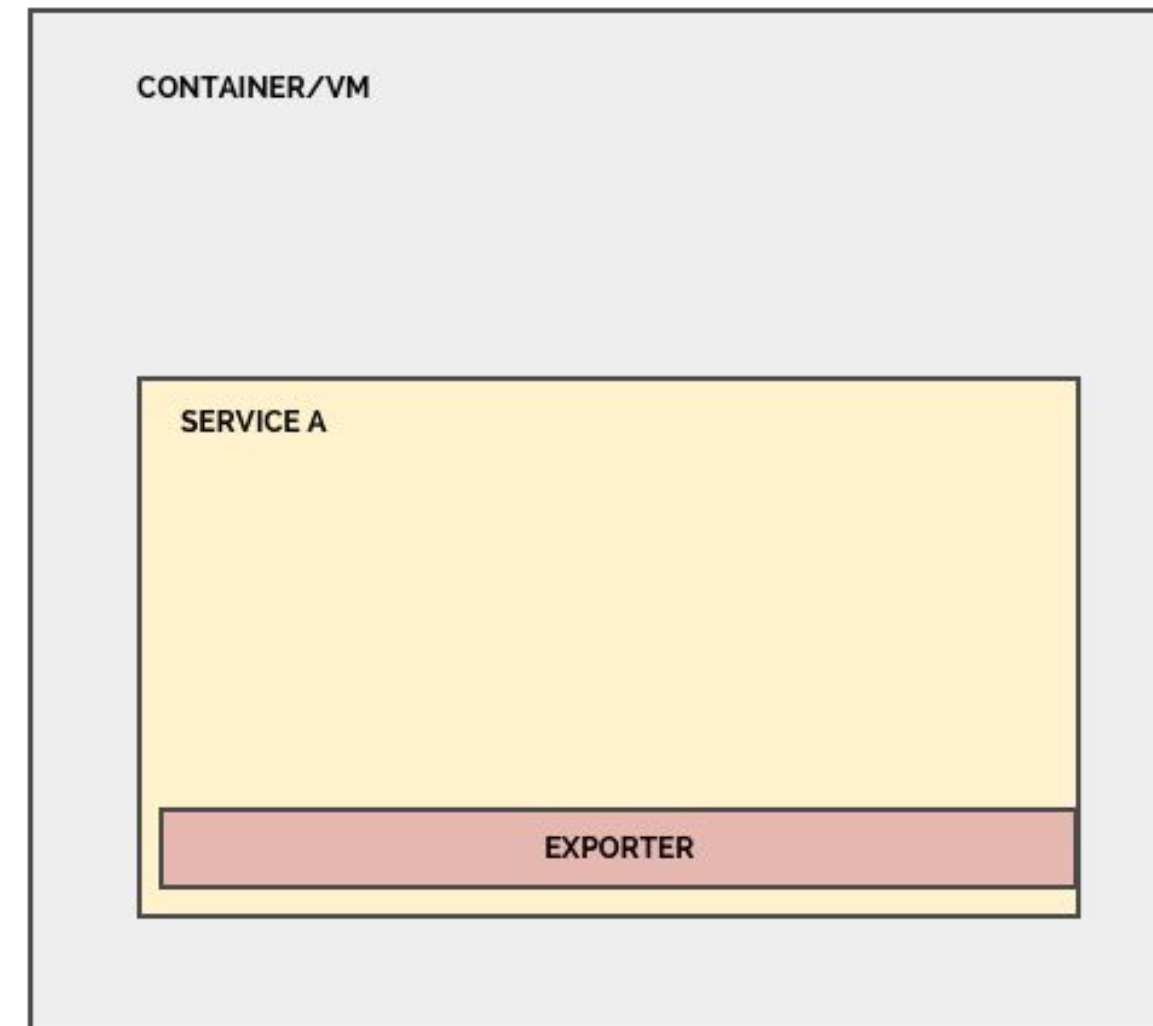
Each **Span** captures important data points specific to the current process handling the request.



Distributed Tracing - Terminology

Context Propagation:

Incoming
Request

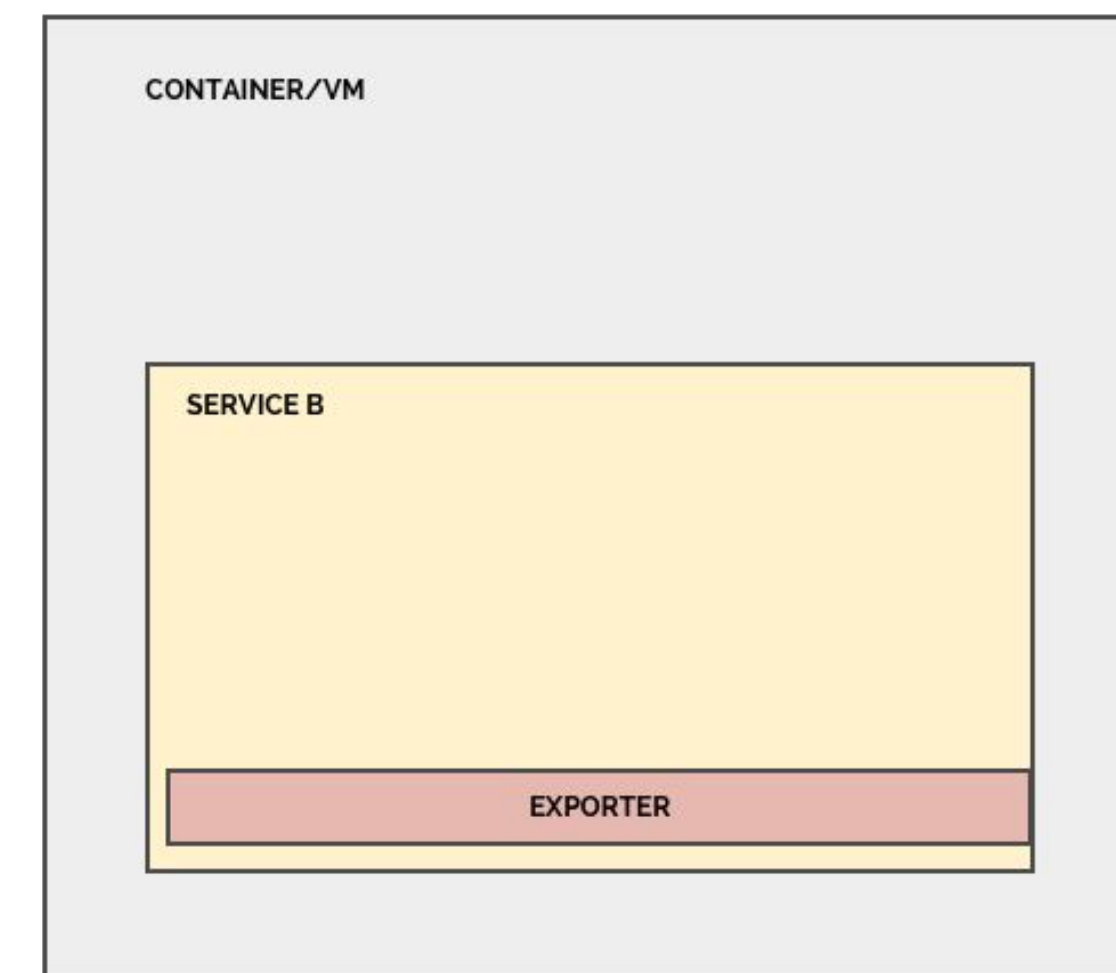
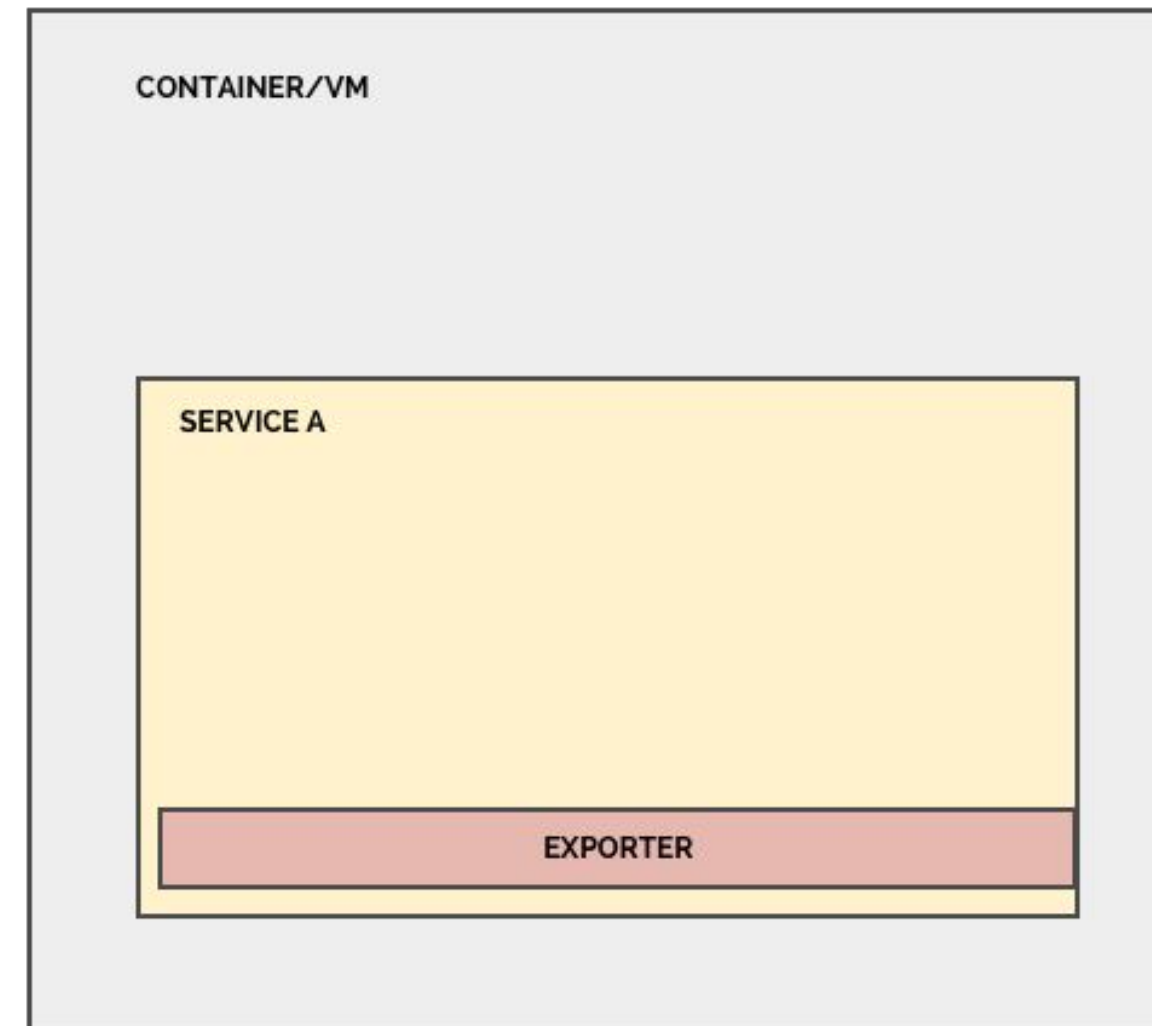



Distributed Tracing - Terminology

Context Propagation:

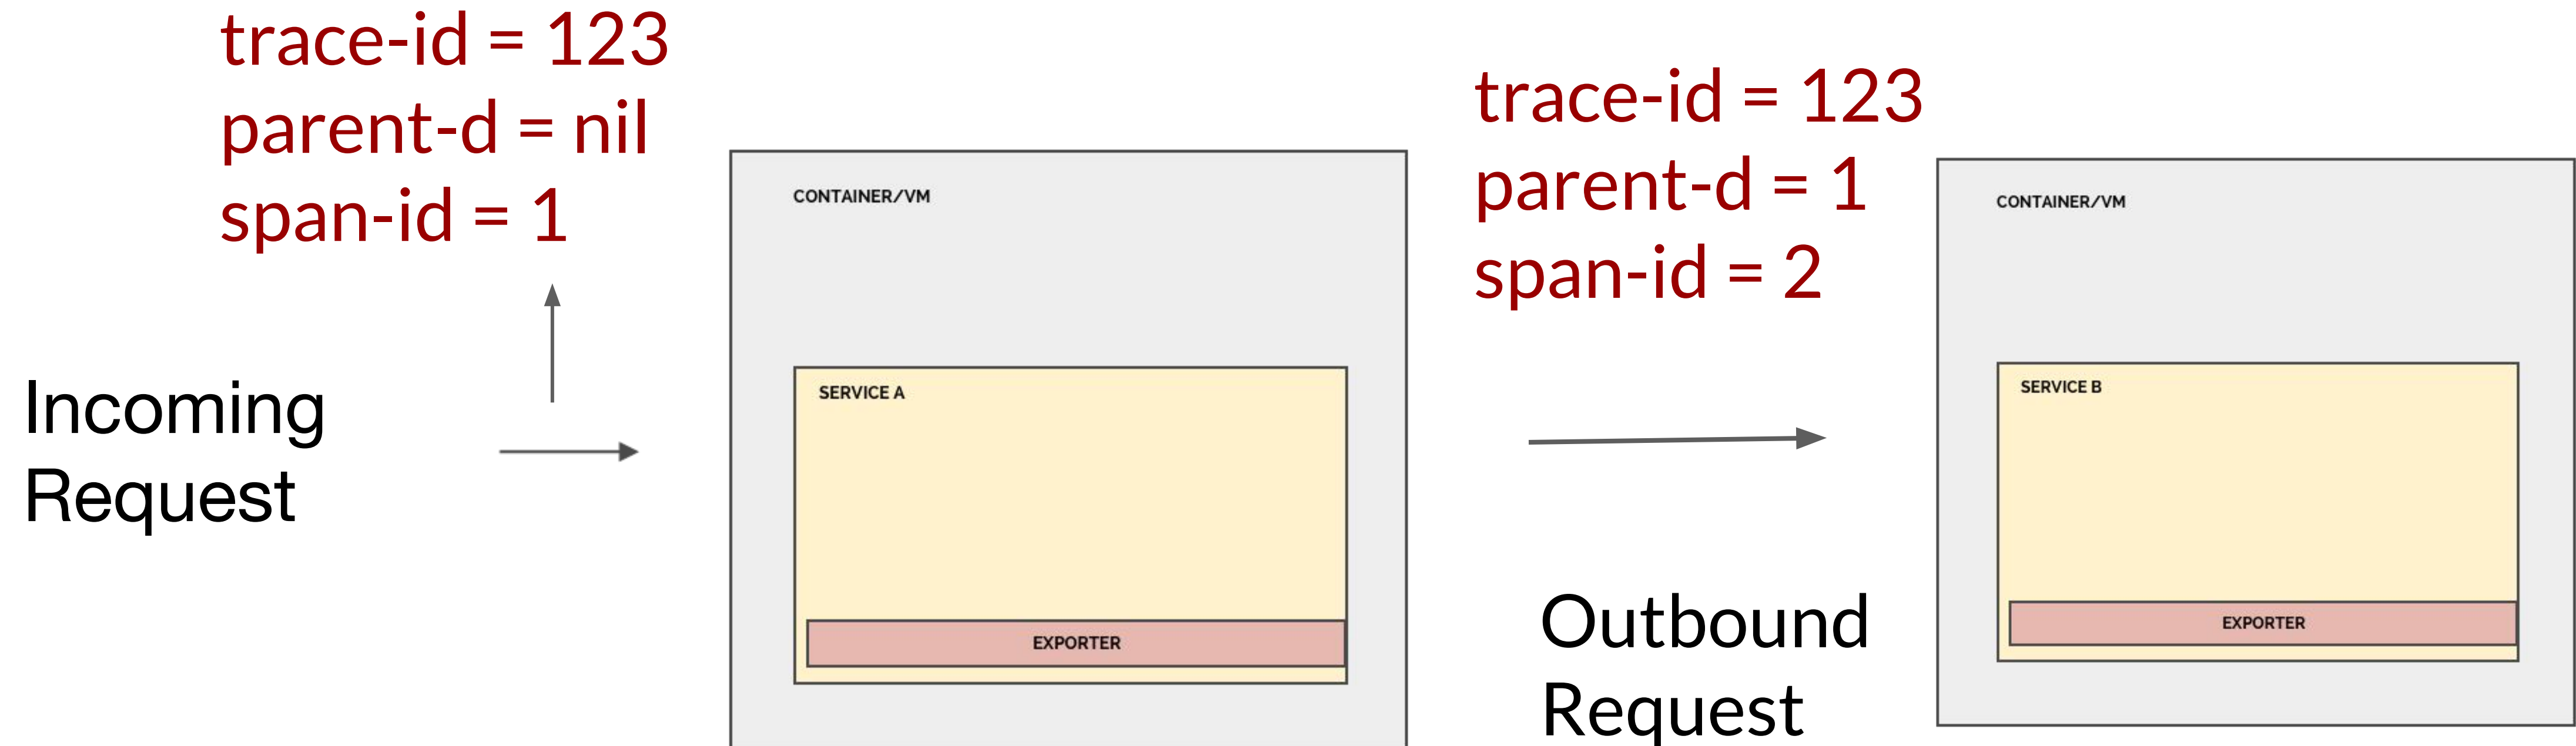
trace-id = 123
parent-d = nil
span-id = 1

Incoming
Request



Distributed Tracing - Terminology

Context Propagation:



Distributed Tracing - Terminology

Tags & Logs: both annotate the span with some contextual information.

- Tags typically apply to the whole span, while logs represent some events that happened during the span execution.
- A log always has a timestamp that falls within the span's start-end time interval.
- The tracing system does not explicitly track causality between logged events the way it keeps track of causality relationships between spans, because it can be inferred from the timestamps.

What questions can tracing help us answer?

Distributed Tracing:

- What services did a **request pass through**?

Distributed Tracing:

- What services did a **request pass through**?
- What occurred in **each service** for a given request?

Distributed Tracing:

- What services did a **request pass through**?
- What occurred in **each service** for a given request?
- Where did the **error** happen?

Distributed Tracing:

- What services did a **request pass through**?
- What occurred in **each service** for a given request?
- Where did the **error** happen?
- Where are the **bottlenecks**?

Distributed Tracing:

- What services did a **request pass through**?
- What occurred in **each service** for a given request?
- Where did the **error** happen?
- Where are the **bottlenecks**?
- What is the **critical path** for a request?

Distributed Tracing:

- What services did a **request pass through**?
- What occurred in **each service** for a given request?
- Where did the **error** happen?
- Where are the **bottlenecks**?
- What is the **critical path** for a request?
- Who should I **page**?

If tracing is so good why isn't everyone using it?

If tracing is so good why isn't everyone using it?

- Not much education or not many publicized case studies on the benefits.

If tracing is so good why isn't everyone using it?

- Not much education or not many publicized case studies on the benefits.
- Vendor Lock in is unacceptable: Instrumentation must be decoupled from vendors

If tracing is so good why isn't everyone using it?

- Not much education or not many publicized case studies on the benefits.
- Vendor Lock in is unacceptable: Instrumentation must be decoupled from vendors .
- Inconsistent APIs: Tracing semantics must not be language dependent.

If tracing is so good why isn't everyone using it?

- Not much education or not many publicized case studies on the benefits.
- Vendor Lock in is unacceptable: Instrumentation must be decoupled from vendors .
- Inconsistent APIs: Tracing semantics must not be language dependent.
- Handoff woes: Tracing libs in Project X do not handoff to tracing libs in Project Y.

Meet OpenTelemetry

OpenTelemetry



Open Telemetry is made up of an integrated set of APIs and libraries as well as a collection mechanism via a agent and collector. These components are used to generate, collect, and describe telemetry about distributed systems.

Problems OpenTelemetry solves:

- Vendor neutrality for tracing, monitoring and logging
- Context Propagation.

OpenTelemetry (opentelemetry.io) Is:


- **Single set of APIs** for tracing and metrics collection.
- **Standardized Context Propagation.**
- **Exporters** for sending data to backend of choice.
- **Collector** for smart traces & metrics aggregation.
- **Integrations** with popular web, RPC and storage frameworks.

OpenTelemetry (opentelemetry.io) Is:

Next major version of the [OpenTracing](#) and [OpenCensus](#) projects.



OpenTelemetry Roadmap:



Merging OpenTracing and OpenCensus: Goals and Non-Goals



Ben Sigelman [Follow](#)

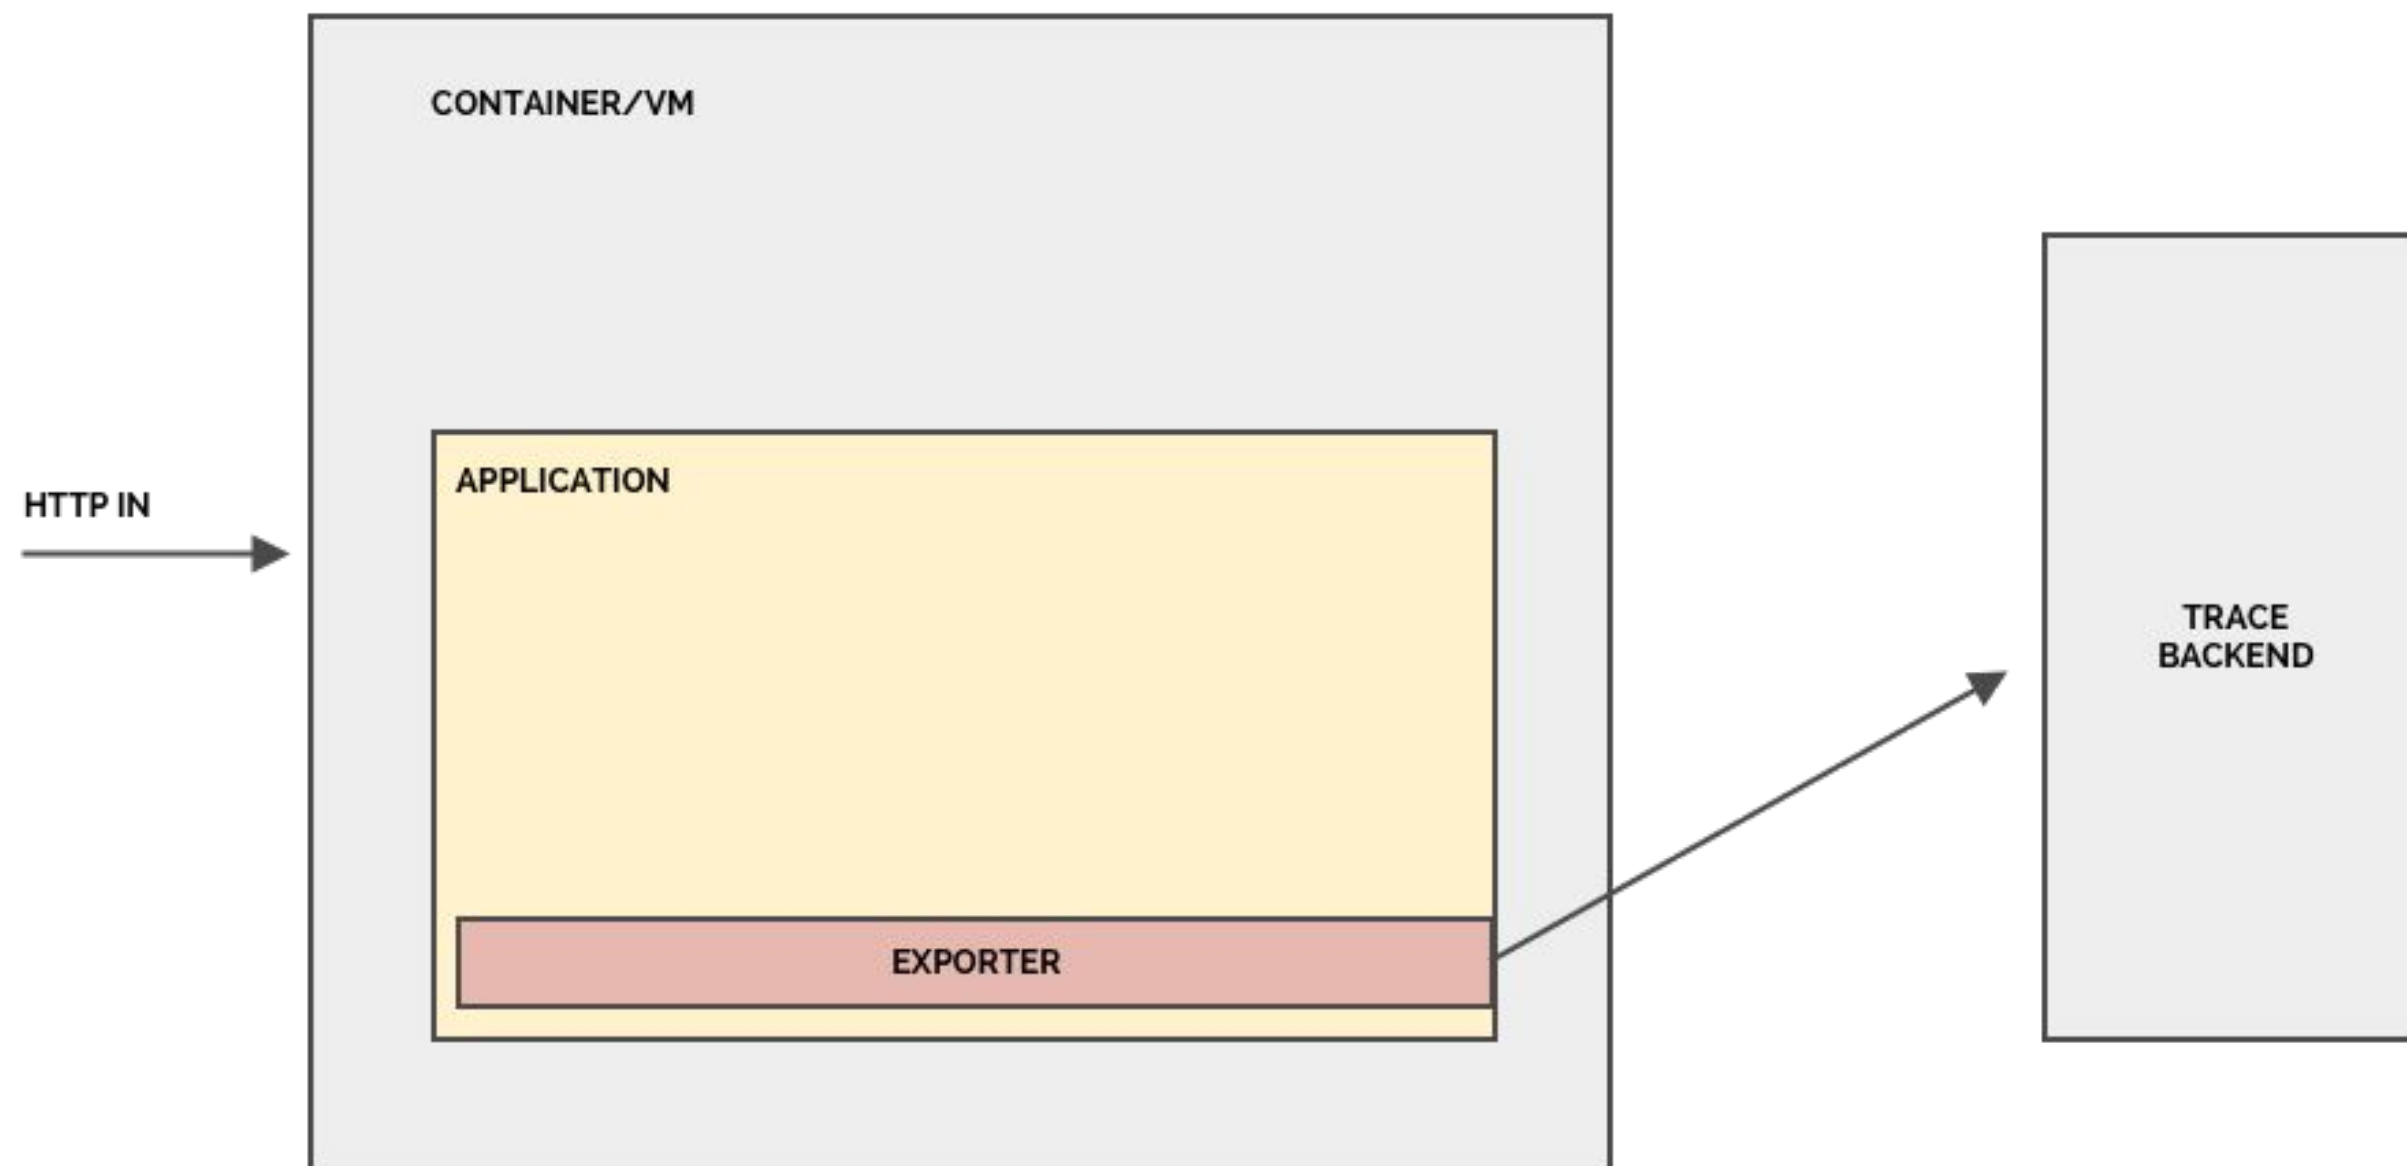
Mar 29 · 3 min read

Announcement: <https://medium.com/opentracing/merging-opentracing-and-opencensus-f0fe9c7ca6f0>

Roadmap: <https://medium.com/opentracing/a-roadmap-to-convergence-b074e5815289>

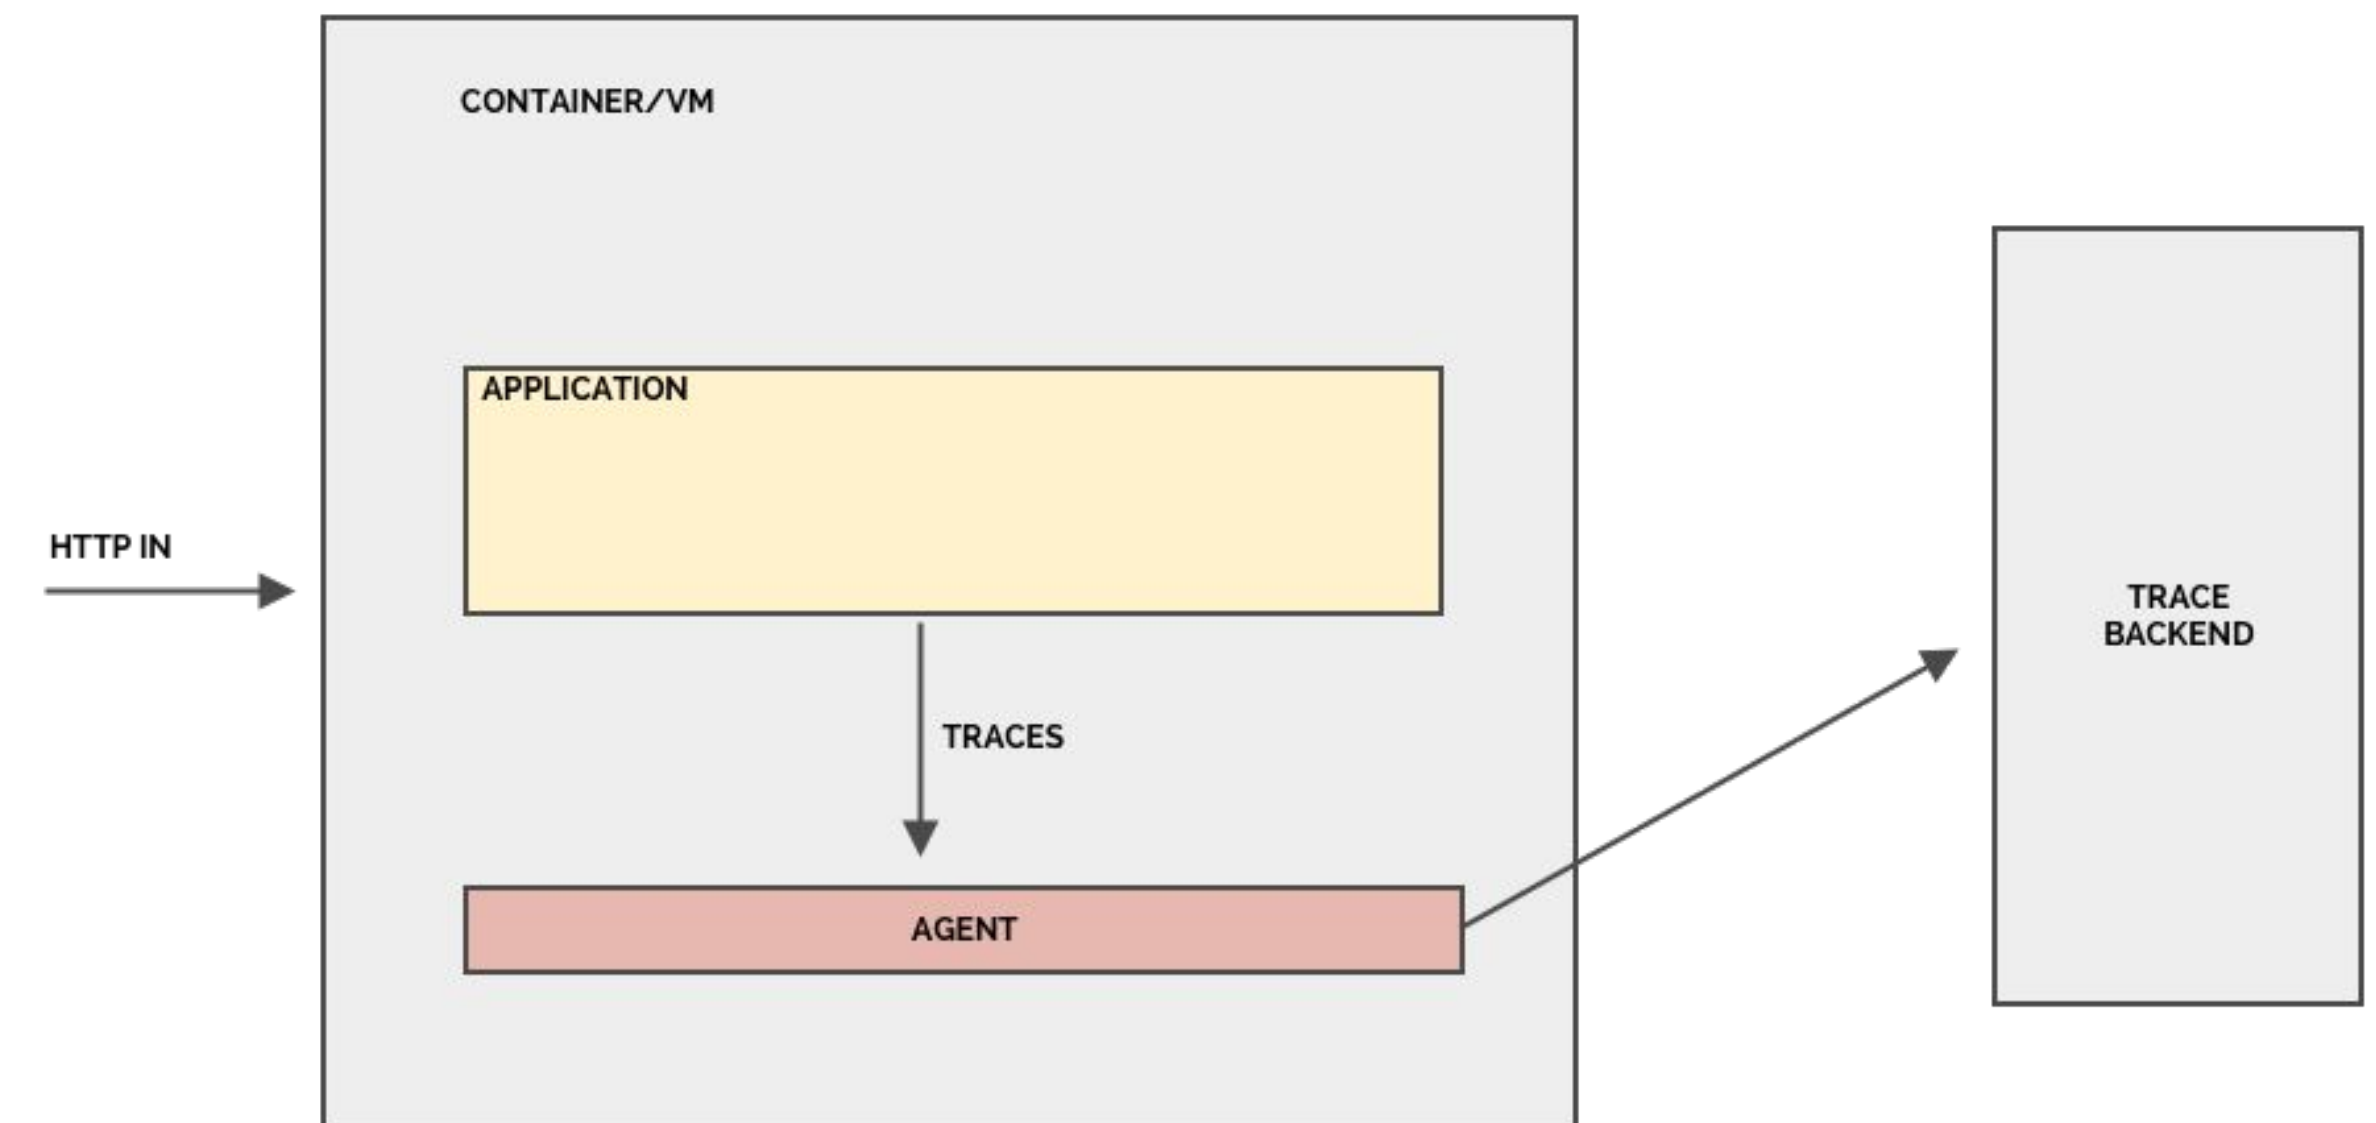
Tracing with OpenTelemetry - The Options

Agentless



Initialize exporter in app code

Using an Agent



Install the agent alongside the app

OpenTelemetry: How to get Involved

Github: <https://github.com/open-telemetry>

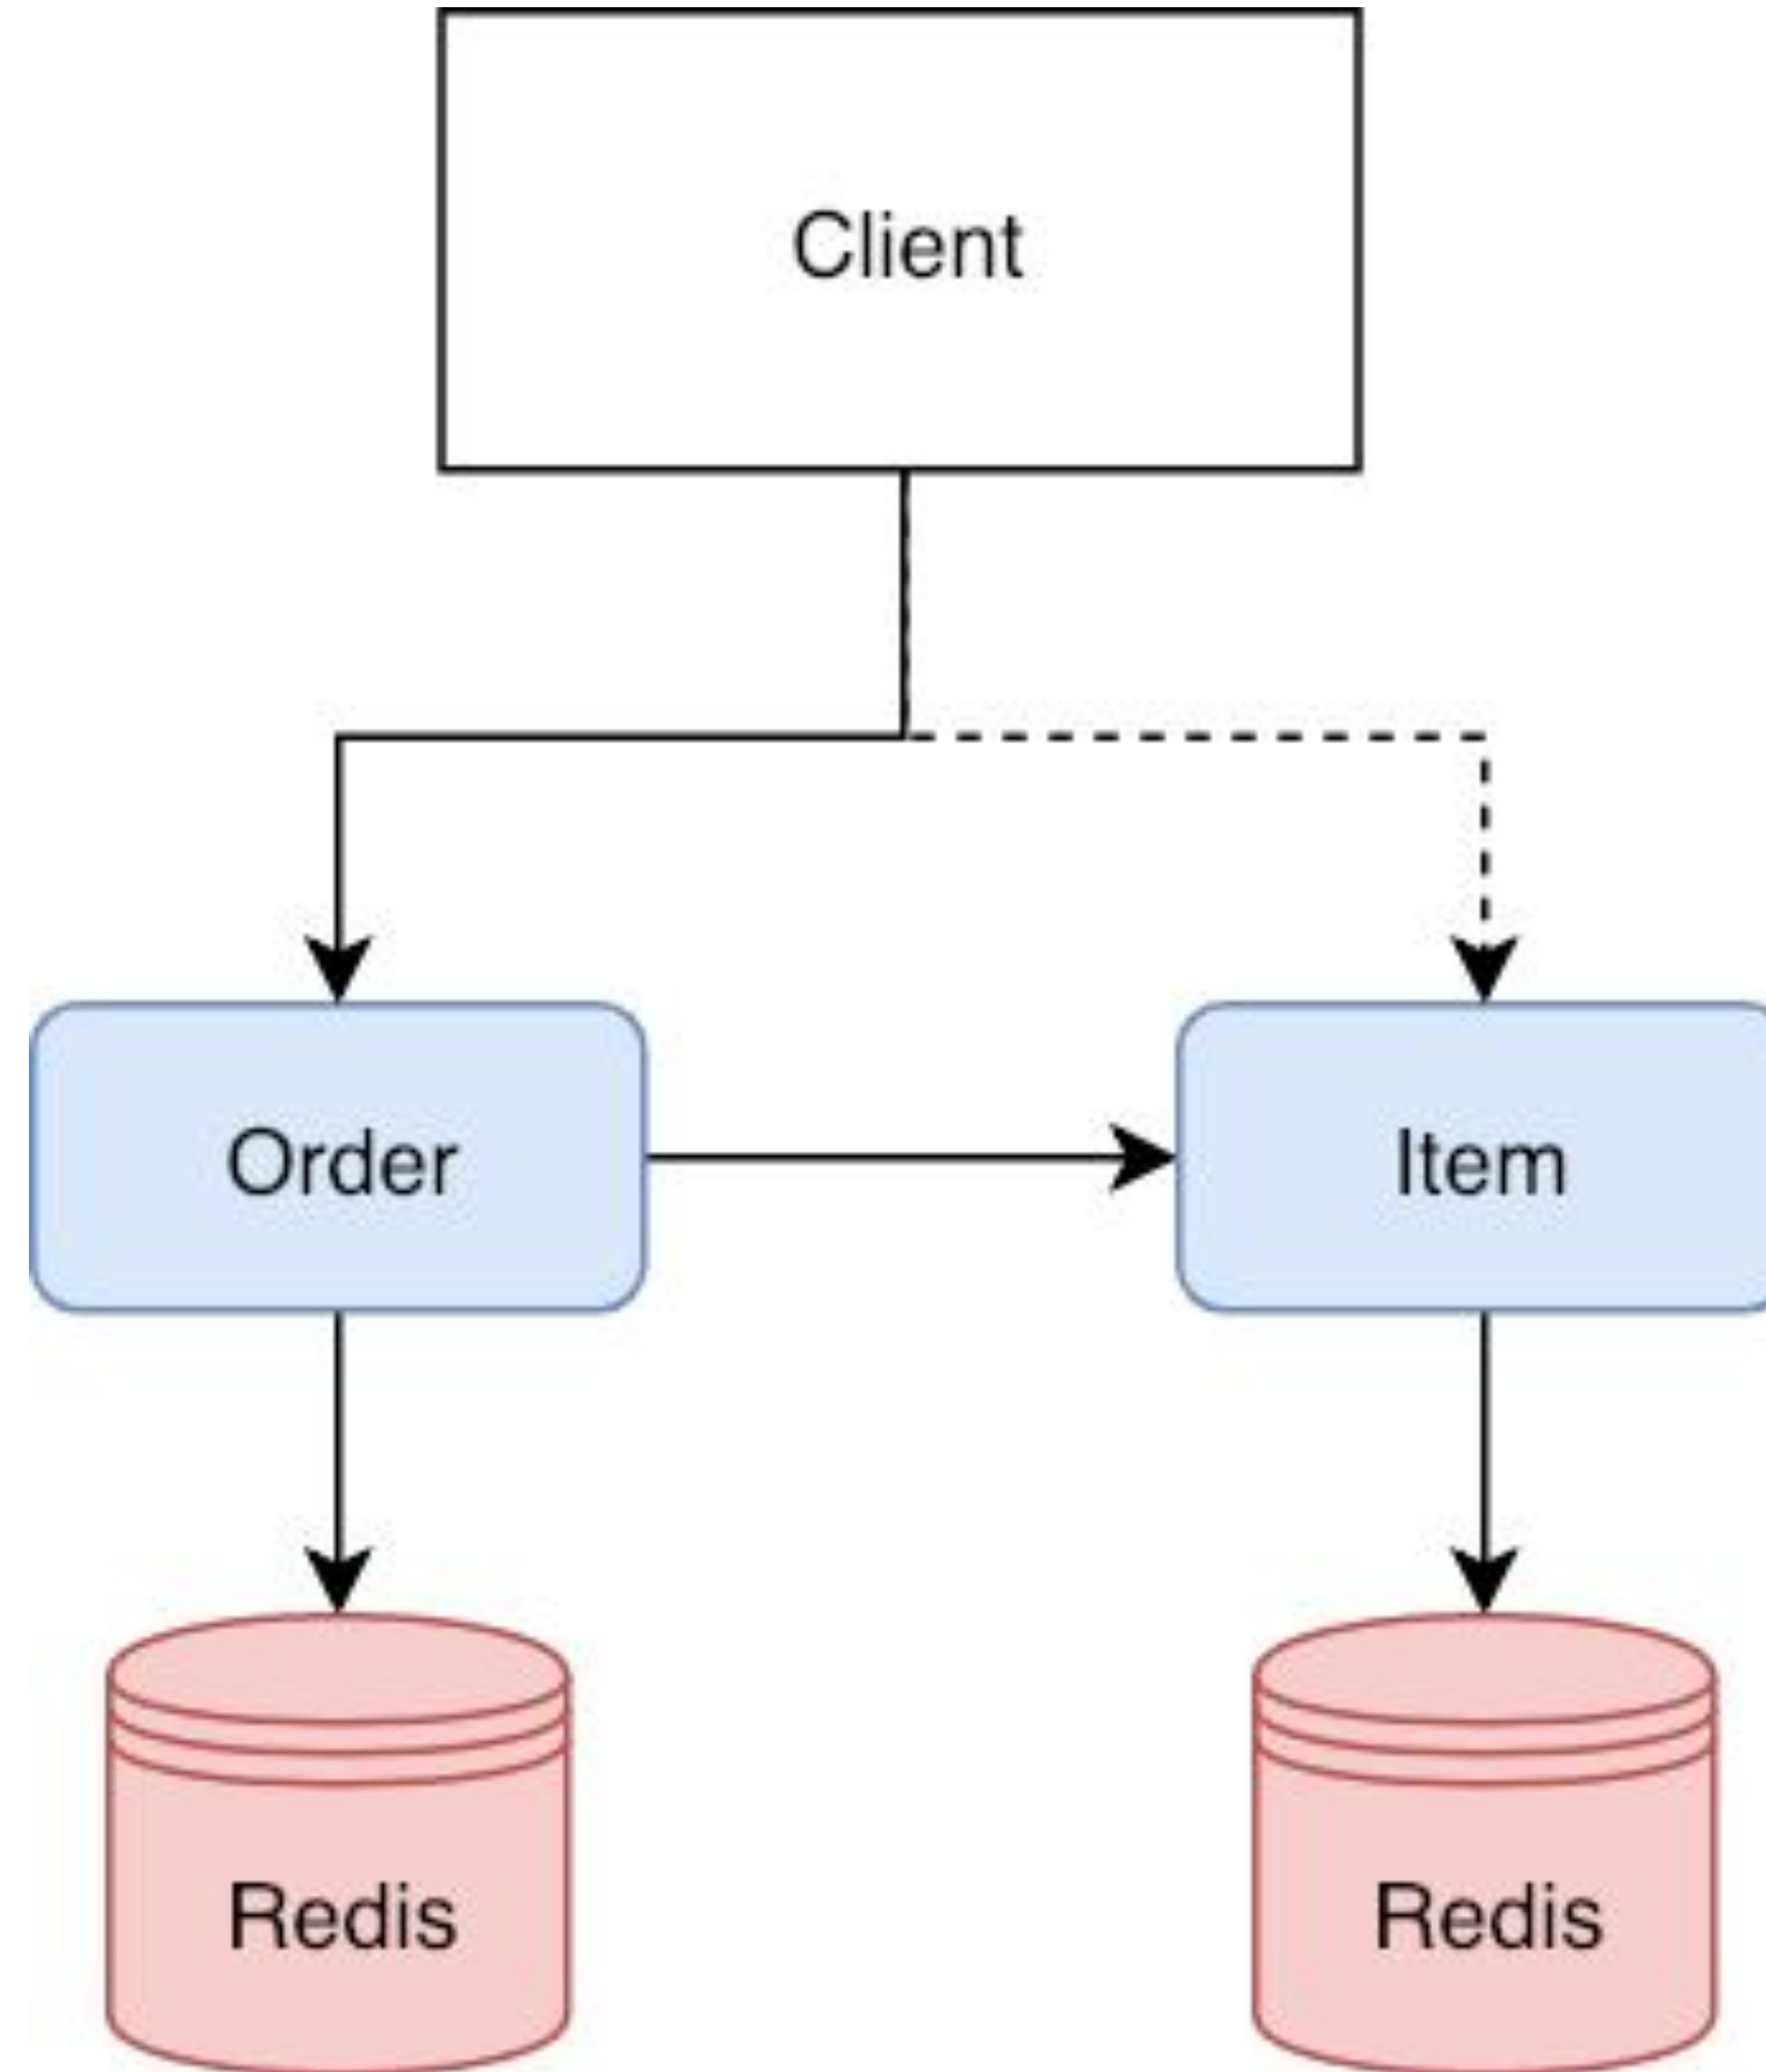
Gitter: <https://gitter.im/open-telemetry>

Languages:

- .NET SDK
- GoLang SDK
- Java SDK
- JavaScript SDK
- Python SDK
- Ruby SIG
- Erlang/Elixir SDK



EXAMPLE



Create a Tracer

```
125 // Creating tracer
126 var tracer ot.Tracer
127 var closer io.Closer
128 tracer, closer, err = util.InitTracer("item", s.logger)
129 if err != nil {
130     s.logger.Warnw("unable to initialize tracer",
131         "error", err,
132     )
133 } else {
134     defer closer.Close()
135     ot.SetGlobalTracer(tracer)
136 }
137
```

Create a Tracer

```
14
15 // InitTracer returns an instance of Jaeger Tracer that samples 100% of traces and logs all spans
16 func InitTracer(serviceName string, logger *Logger) (ot.Tracer, io.Closer, error) {
17     cfg, err := config.FromEnv()
18     if err != nil {
19         return nil, nil, err
20     }
21
22     cfg.Sampler.Type = "const"
23     cfg.Sampler.Param = 1
24     cfg.Reporter.LogSpans = false
25
26     tracer, closer, err := cfg.New(
27         serviceName,
28         config.Logger(logger),
29         config.Metrics(prometheus.New()),
30     )
31     if err != nil {
32         return nil, nil, err
33     }
34     return tracer, closer, nil
35 }
```


Tracers



JAEGER



honeycomb.io



Skywalking



ZIPKIN



LIGHTSTEP



DATADOG



AJUA

Instrument

```
37 // TracerMiddleware adds a Span to the request Context ready for other handlers to use it.
38 func TracerMiddleware(inner http.Handler, route Route) http.HandlerFunc {
39     return http.HandlerFunc(func(w http.ResponseWriter, r *http.Request) {
40         var ctx context.Context
41         var span ot.Span
42         tracer := ot.GlobalTracer()
43
44         // If possible, extract span context from headers
45         spanCtx, _ := tracer.Extract(ot.HTTPHeaders, ot.HTTPHeadersCarrier(r.Header))
46         if spanCtx == nil {
47             span = tracer.StartSpan("request")
48             defer span.Finish()
49             ctx = ot.ContextWithSpan(r.Context(), span)
50         } else {
51             span = tracer.StartSpan("request", ext.RPCServerOption(spanCtx))
52             defer span.Finish()
53             ctx = ot.ContextWithSpan(r.Context(), span)
54         }
55         for k, v := range r.Header {
56             span.SetTag(fmt.Sprintf("header.%s", k), v)
57         }
58
59         // TODO: capture return code as tag in root trace
60         span.SetTag("method", r.Method)
61         span.SetTag("url", r.URL.Path)
62         span.SetTag("handler", route.Name)
63
64         r = r.WithContext(ctx)
65         inner.ServeHTTP(w, r)
66     })
67 }
```


Instrument

```
204 // getItem retrieves a single Item by ID from Redis.
205 func (s *Server) getItem() http.HandlerFunc {
206     return func(w http.ResponseWriter, r *http.Request) {
207         span, ctx := ot.StartSpanFromContext(r.Context(), "getItem")
208         defer span.Finish()
209
210         pr := mux.Vars(r)
211         key := pr["id"]
212
213         item, err := s.RedisGetItem(ctx, key)
214         if err != nil {
215             s.logger.Errorw("unable to get key from redis",
216                 "key", key,
217                 "error", err,
218             )
219             s.Respond(ctx, http.StatusInternalServerError, "unable to retrieve item", 0, nil, w)
220             return
221         }
222         if item == nil {
223             s.Respond(ctx, http.StatusNotFound, fmt.Sprintf("item with ID %s doesn't exist", key), 0, nil, w)
224             return
225         }
226         s.Respond(ctx, http.StatusOK, "item retrieved", 1, []*Item{item}, w)
227     }
228 }
```

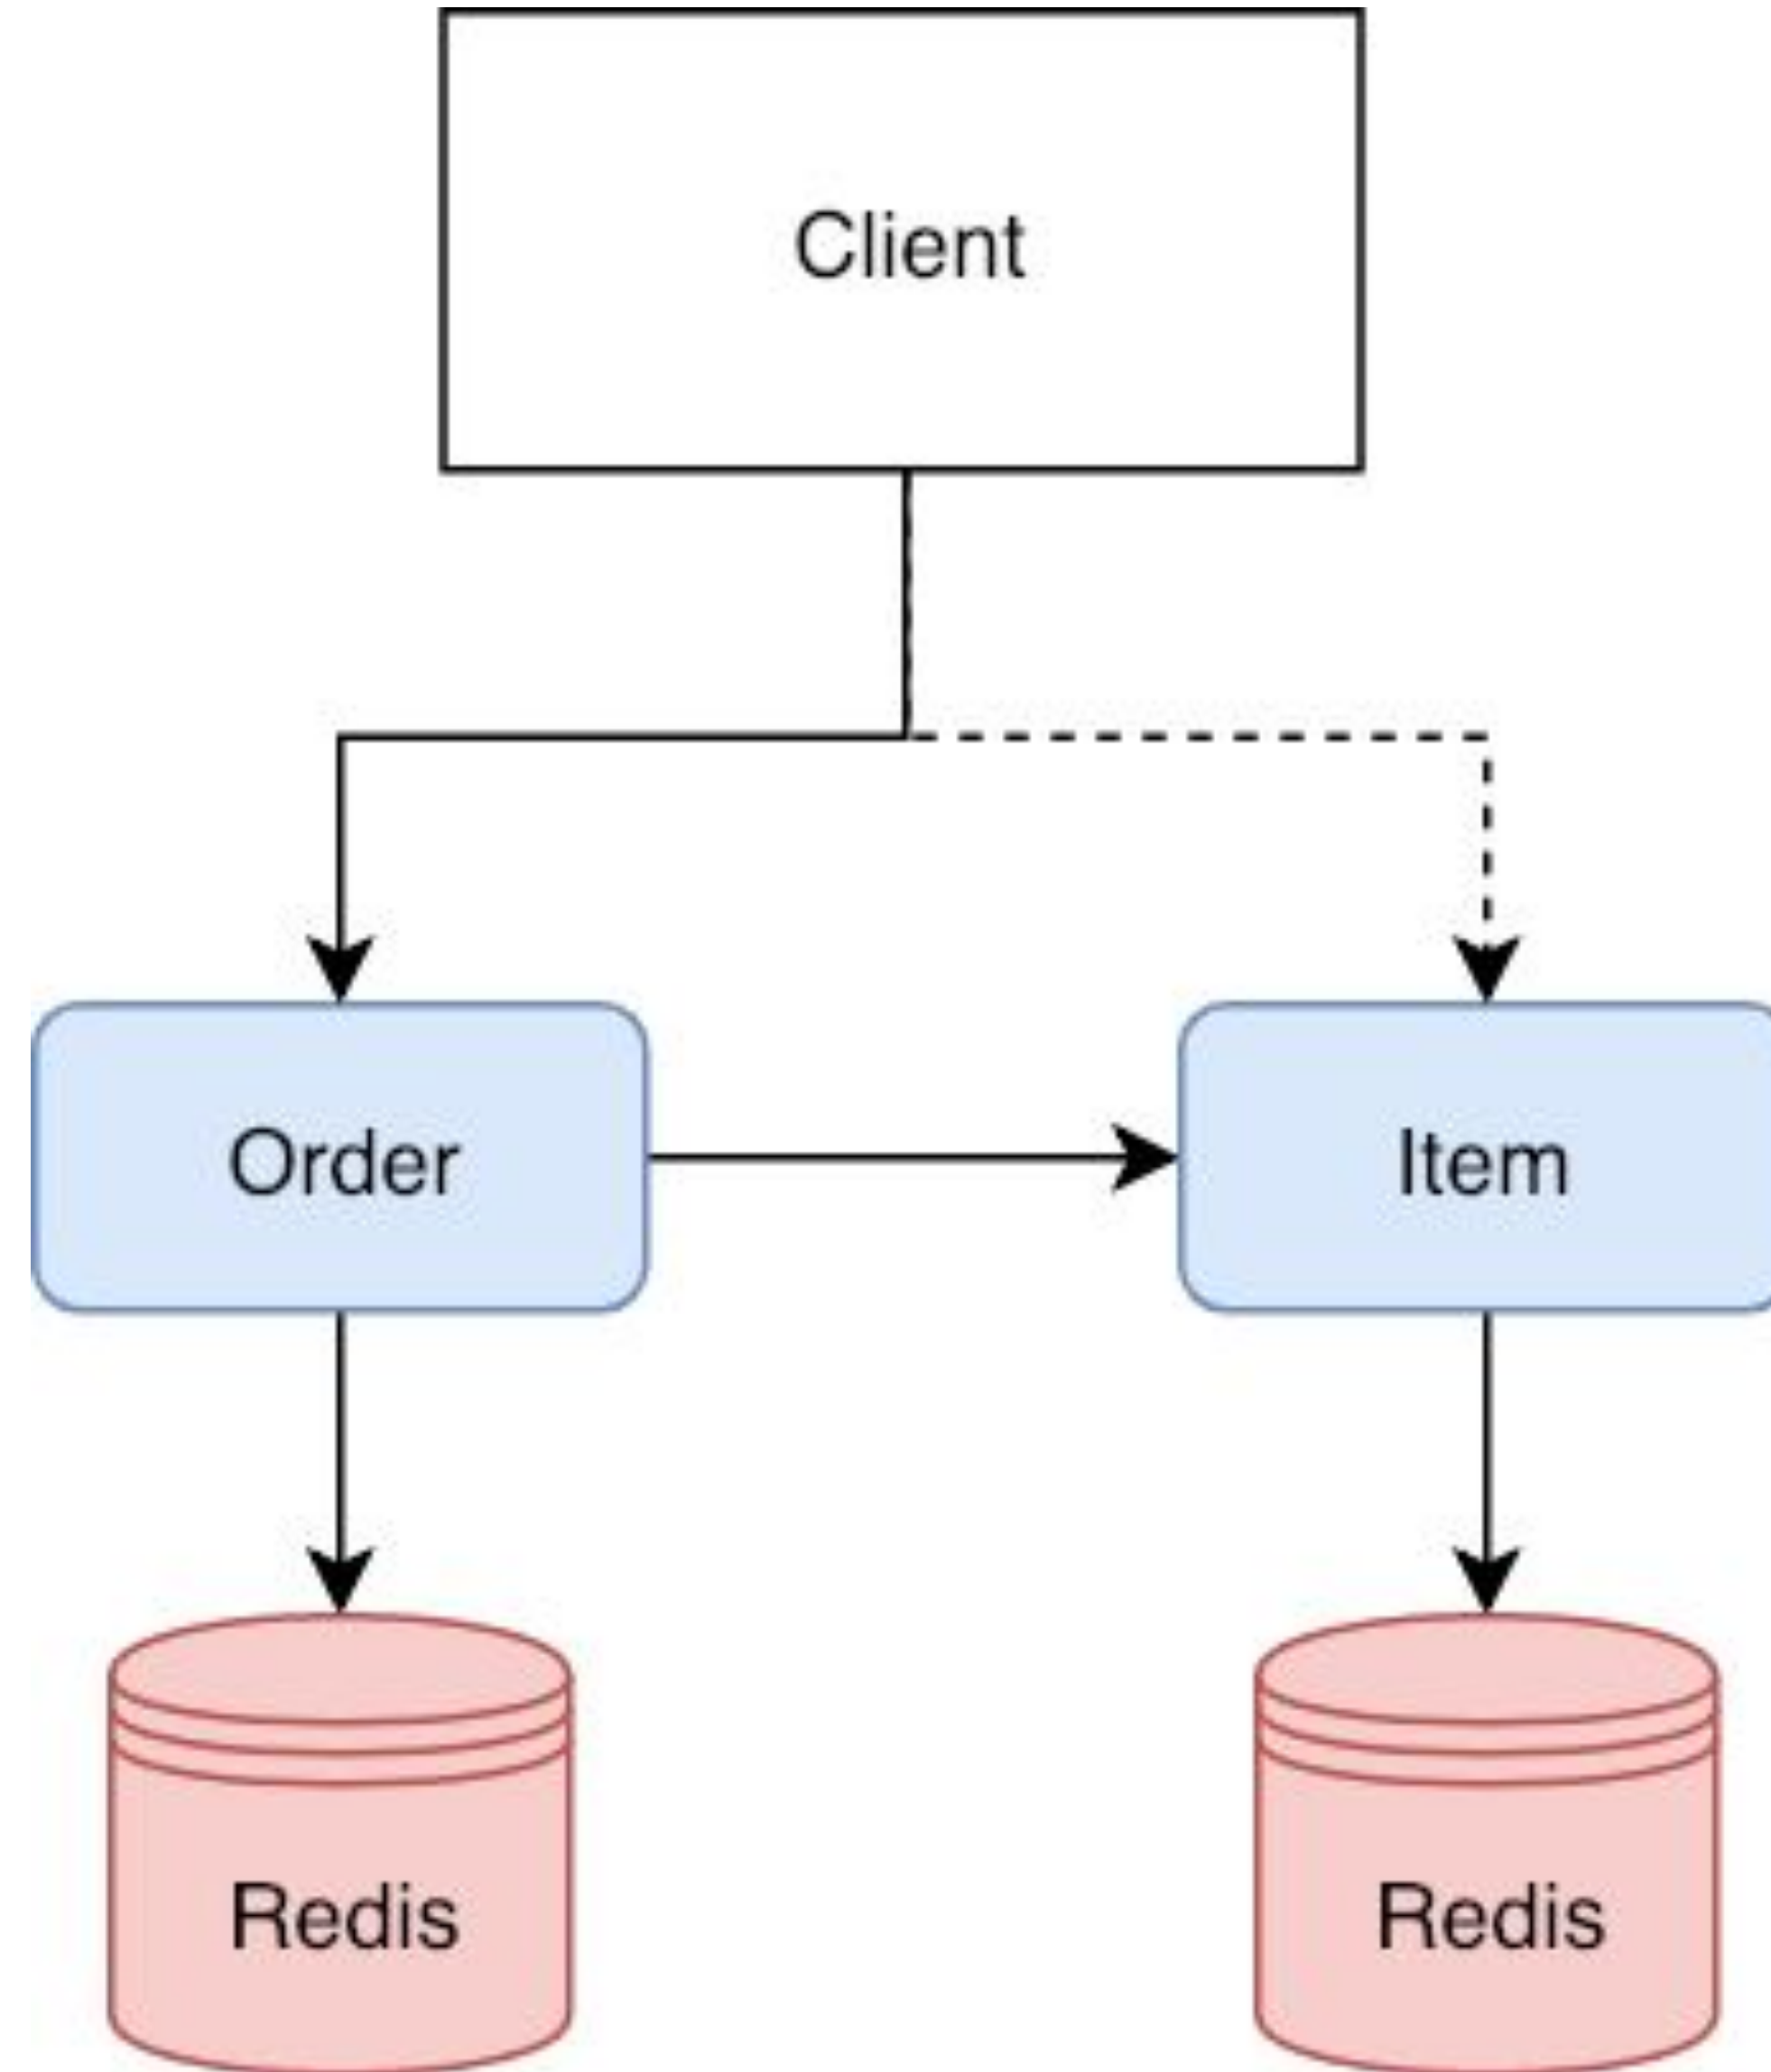
Introducing Jaeger:

- Open source distributed tracing platform.
- Inspired by Google Dapper and OpenZipkin
- Created by Uber in 2015 and donated to CNCF in 2017.
- Compliant with both OpenTracing and OpenCensus.
- Supports multiple storage options (Cassandra, ElasticSearch, In-Memory)
- Compatible with Apache Kafka for backpressure management.



DEMO

[obitech/micro-obs](https://github.com/obitech/micro-obs)



Conclusion

- Tracing is crucial for understanding complex, microservices applications.
- Distributed tracing provides a base view of the system that can drastically shorten feedback loops and the number of people involved incidents.
- Tracing provides much more context, allowing an on call responder to better understand the system and get further on their own before involving more people.

Thank You

Twitter: @Hungai

Email: hungaikevin@gmail.com



AJUA
