




Kubernetes: What is “reconciliation”?




Tim Hockin <thockin@google.com>
@thockin

Assume there's a cloud API to make shapes.


Why shapes? It's just concrete enough to reason about,
while not getting stuck in the details.







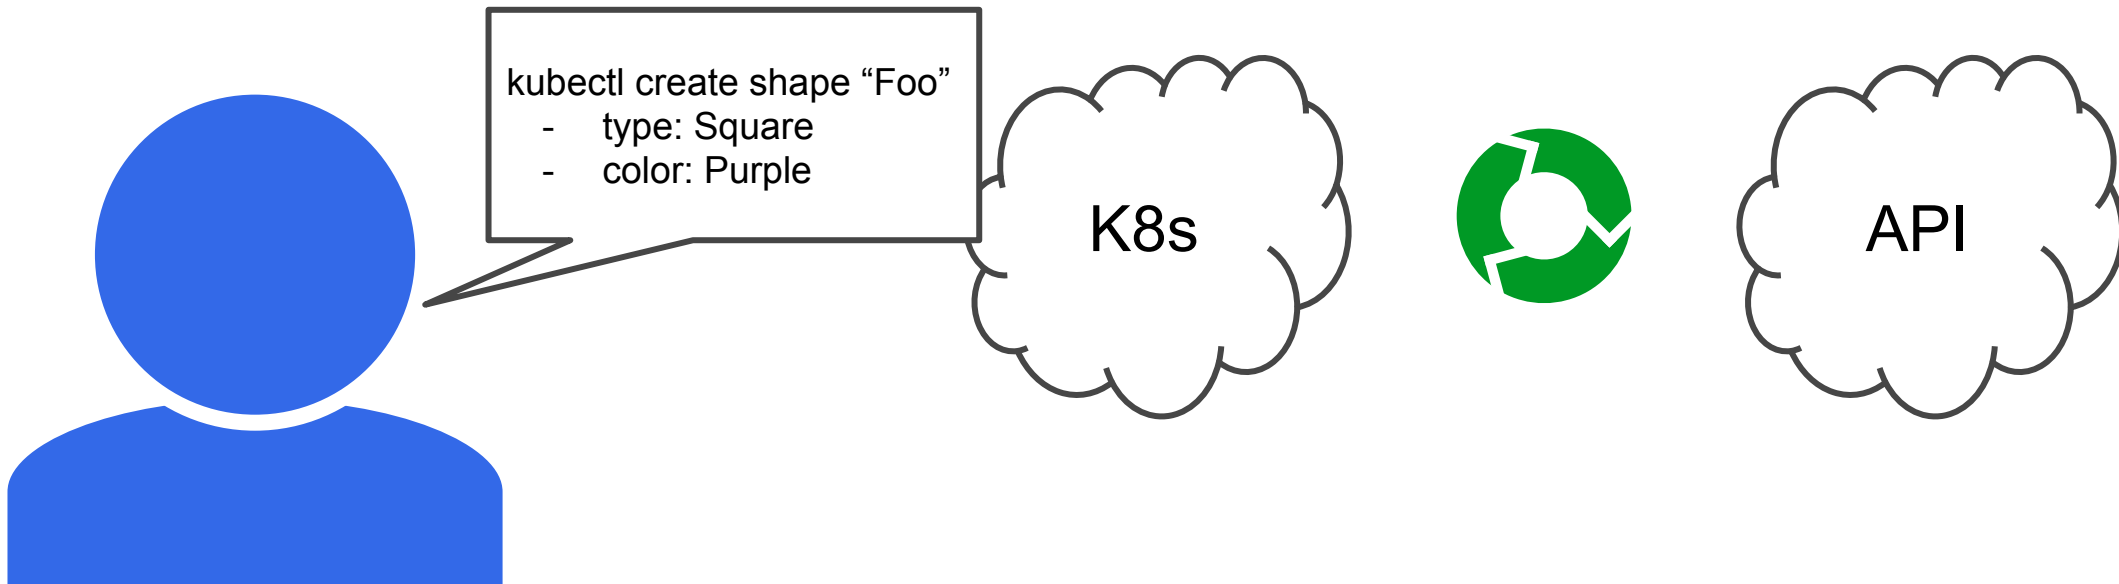
Foo

API





This API is fine, but I want to wrap it into a declarative system (e.g. Kubernetes)





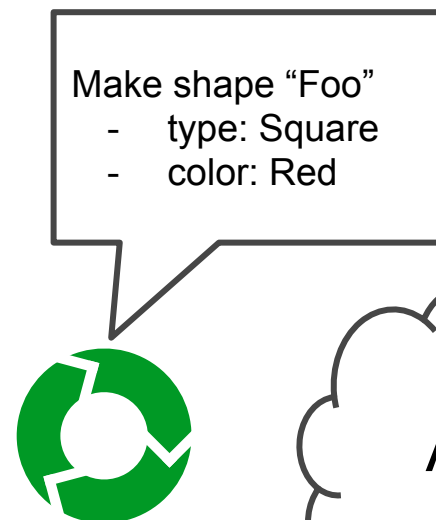



kind: Shape
name: Foo
type: Square
color: Red







kind: Shape
name: Foo
type: Square
color: Red







kind: Shape
name: Foo
type: Square
color: Red

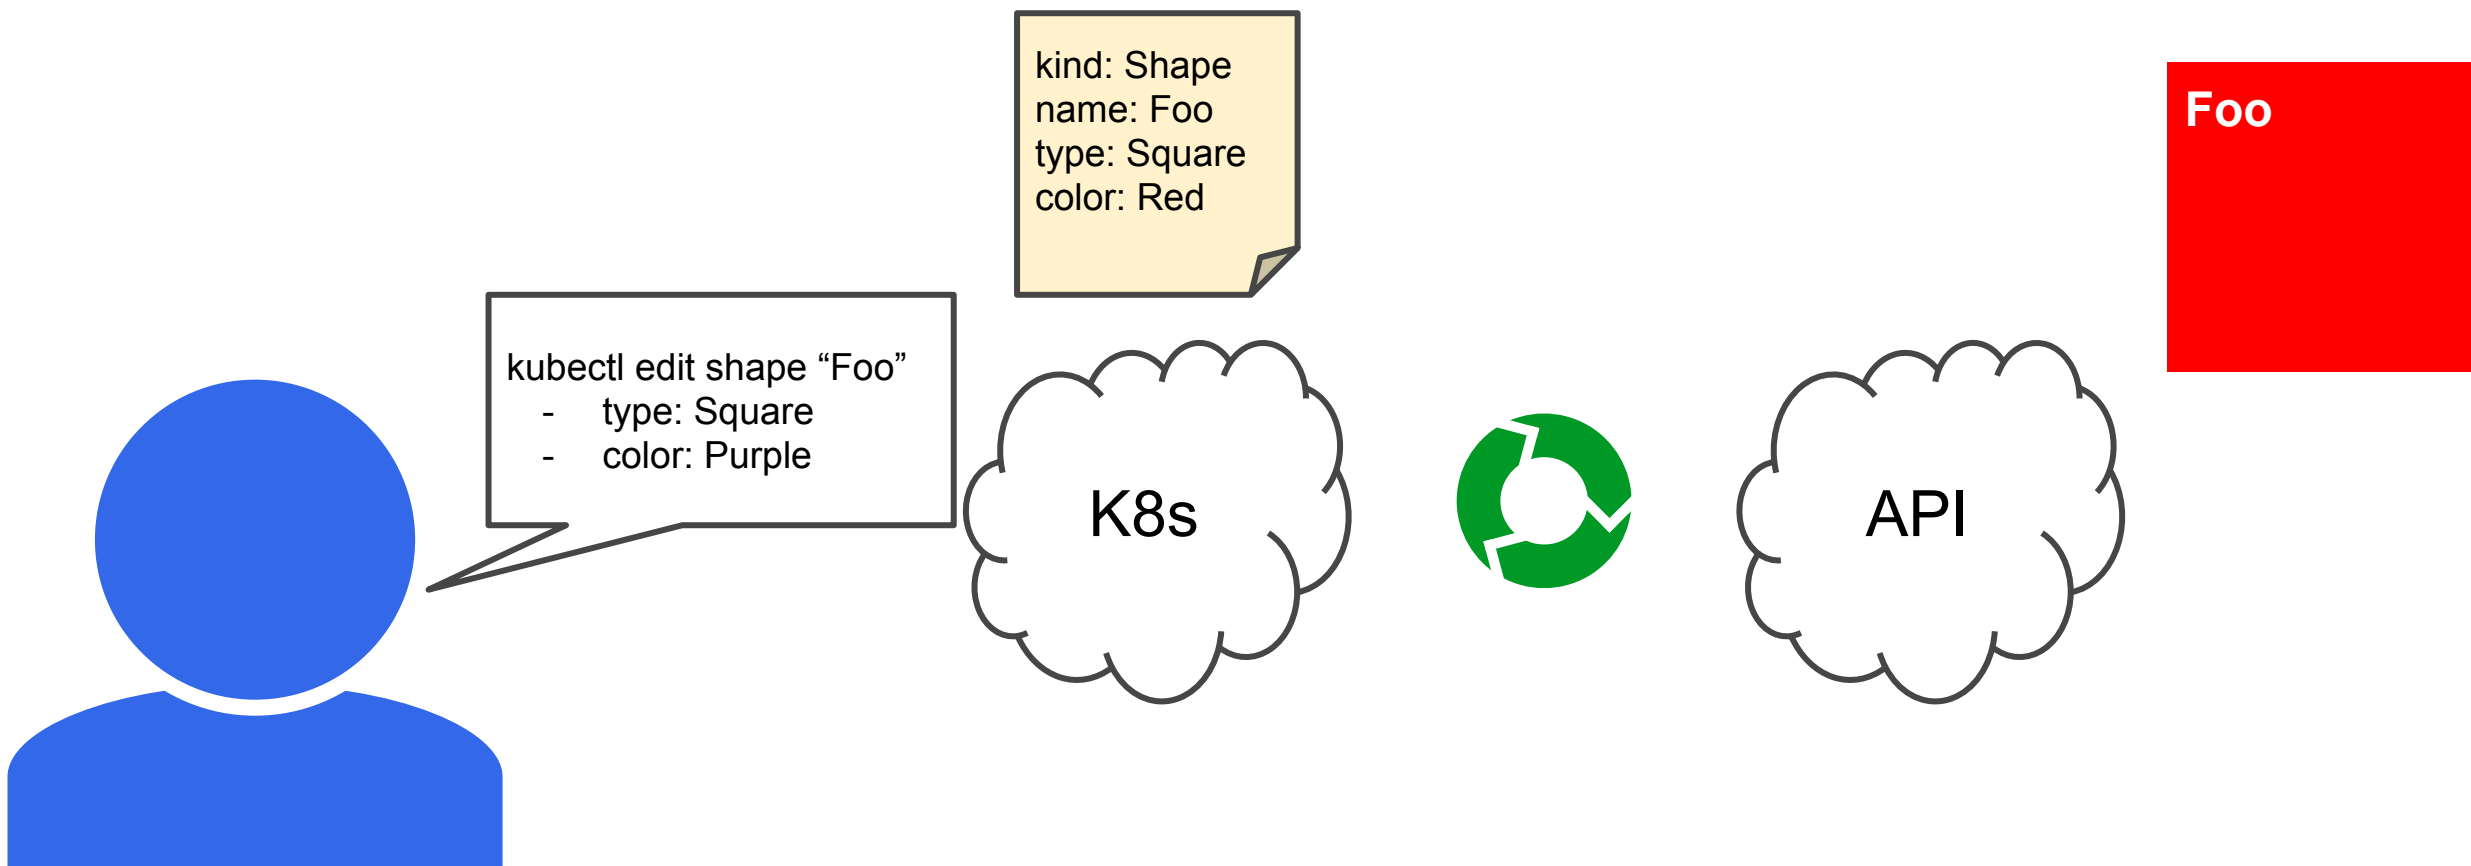


The controller will keep my Kubernetes object in sync with the underlying API





kind: Shape
name: Foo
type: Square
color: Red










```
kind: Shape
name: Foo
type: Square
color: Purple
```








kind: Shape
name: Foo
type: Square
color: Purple





```
kind: Shape
name: Foo
type: Square
color: Purple
```






This is what we call “reconciliation”. Specifically, this is **uni-directional reconciliation**.

What happens if...





kind: Shape
name: Foo
type: Square
color: Purple





kind: Shape
name: Foo
type: Square
color: Purple





Delete shape Foo





kind: Shape
name: Foo
type: Square
color: Purple




The shape I wanted has inadvertently been removed by a human or other system.

If we only reconcile in one direction, we will never fix it!


We need to observe that the underlying state has changed and re-assert the state we want.



kind: Shape
name: Foo
type: Square
color: Purple





Make shape "Foo"
- type: Square
- color: Purple





kind: Shape
name: Foo
type: Square
color: Purple





We usually call this **bi-directional reconciliation**.

But it gets funkier. What if...



kind: Shape
name: Foo
type: Square
color: Purple







Create shape Bar

- type: Circle
- color: Red




kind: Shape
name: Foo
type: Square
color: Purple

Foo





kind: Shape
name: Foo
type: Square
color: Purple



What should the controller do?

Does it expect to have exclusive use of all shapes? If so,
clean up!




Does it expect to share the shapes API with other users?
If so, leave it alone!

Right?

I said it gets **funkier**. What if...





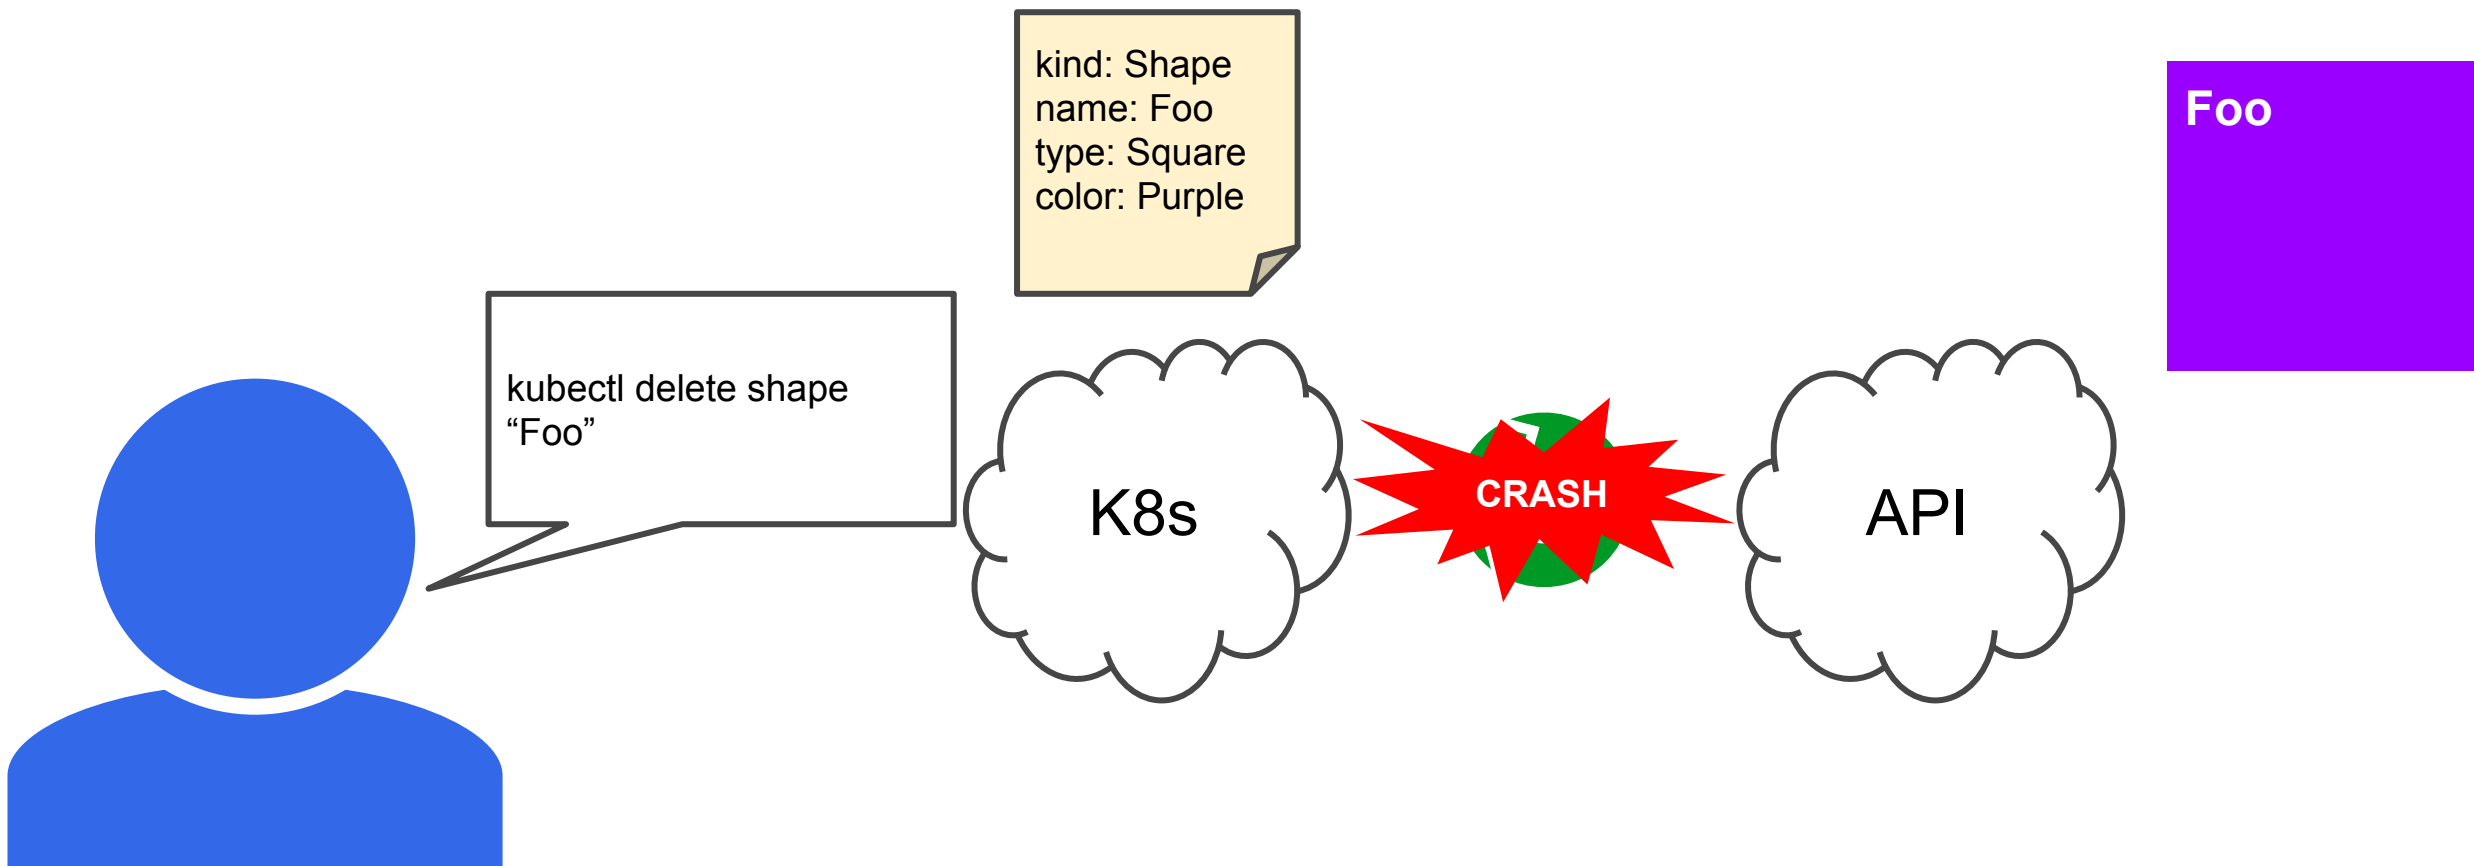
```
kind: Shape
name: Foo
type: Square
color: Purple
```











kind: Shape
name: Foo
type: Square
color: Purple











The controller **missed the deletion** of the shape, but as we saw earlier, it ignored things it doesn't know.

This is a **LEAK!**

The controller has to know which shapes it **owns** and which it doesn't.


HOW to do that depends on the API. Examples:

- Special name prefixes
- Metadata (labels, tags, description)
- Controller-specific checkpoints




Foo

labels:
owner: k8s




Foo

labels:
owner: k8s




Foo
labels:
owner: k8s





Foo

labels:
owner: k8s



Foo

labels:
owner: k8s



This is sometimes called the “list-watch” pattern.

Now the controller will keep things it owns in sync and ignores other things.



What if...



kind: Shape
name: Foo
type: Square
color: Purple




Foo

labels:
owner: k8s





kind: Shape
name: Foo
type: Square
color: Purple



Create shape Bar

- type: Circle
- color: Red
- labels:
 - owner: k8s

Foo



labels:
owner: k8s



kind: Shape
name: Foo
type: Square
color: Purple

Foo

labels:
owner: k8s




Bar


labels:
owner: k8s



kind: Shape
name: Foo
type: Square
color: Purple



List shapes where
owner=k8s



Foo


labels:
owner: k8s

Bar


labels:
owner: k8s



kind: Shape
name: Foo
type: Square
color: Purple



Foo
Bar



Foo

labels:
owner: k8s

Bar

labels:
owner: k8s



kind: Shape
name: Foo
type: Square
color: Purple

I don't have a
"Bar"

Foo

labels:
owner: k8s

K8s




API

Bar


labels:
owner: k8s



kind: Shape
name: Foo
type: Square
color: Purple




Delete shape "Bar"



Foo

labels:
owner: k8s







kind: Shape
name: Foo
type: Square
color: Purple

Foo

labels:
owner: k8s



Note that while doing a full reconciliation at startup is necessary, it is not sufficient.

Good controllers will reconcile against underlying APIs continuously or at least periodically.

How does this apply to real life?

This pattern is found in almost every case where Kubernetes layers on top of some other API. Examples:

- Cloud load-balancers for Services & Ingress
- Cloud disks for PersistentVolumes
- iptables rules for Services
- Running containers for Pods



Sadly, not every controller gets this right.

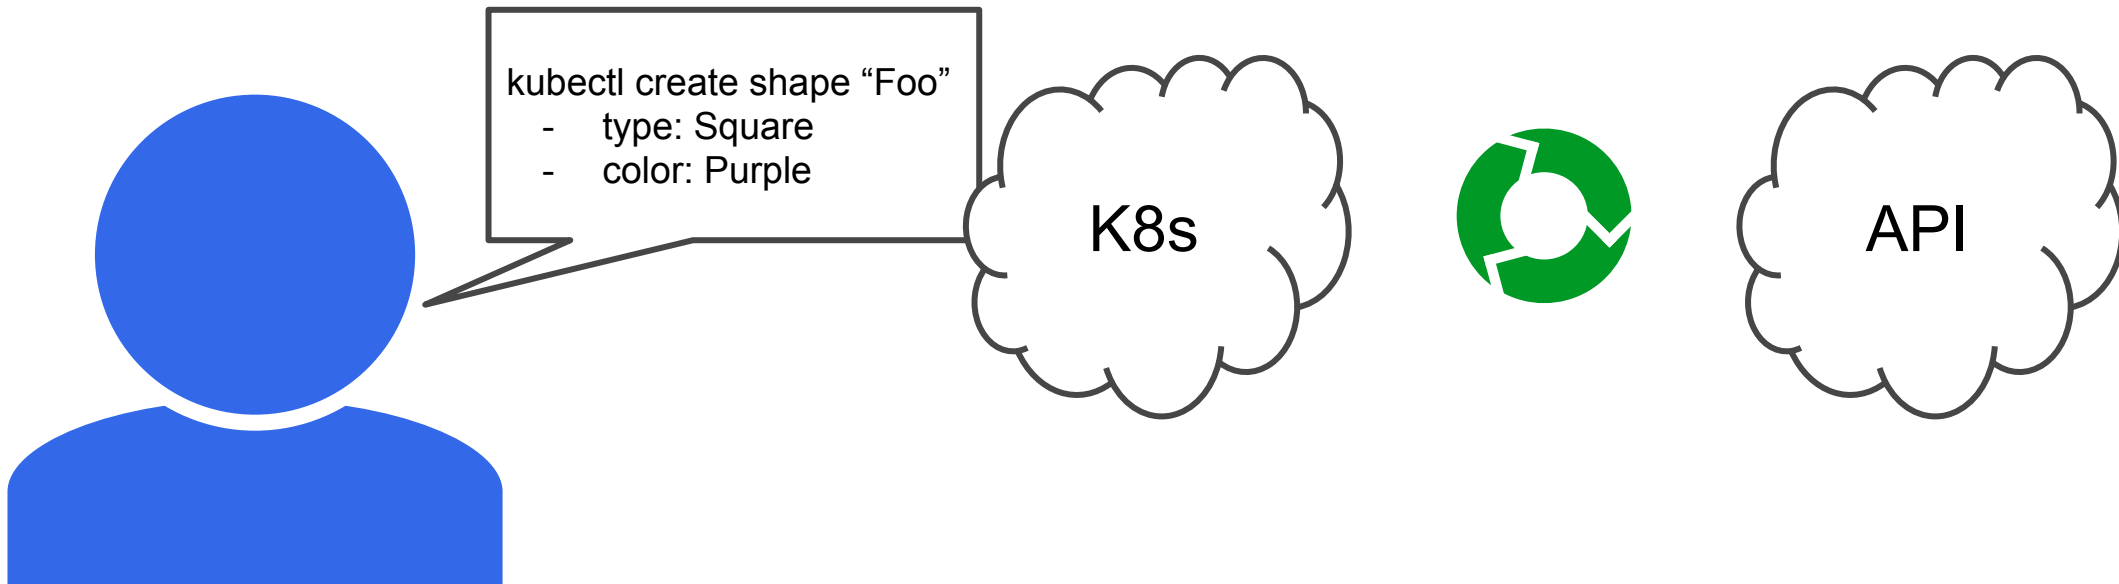
While every controller should strive for complete reconciliation, sometimes the underlying API makes it very hard or expensive or even just impossible.

:(

There are some techniques that can mitigate the lack of mechanisms to denote ownership (or augment them).



Finalizers

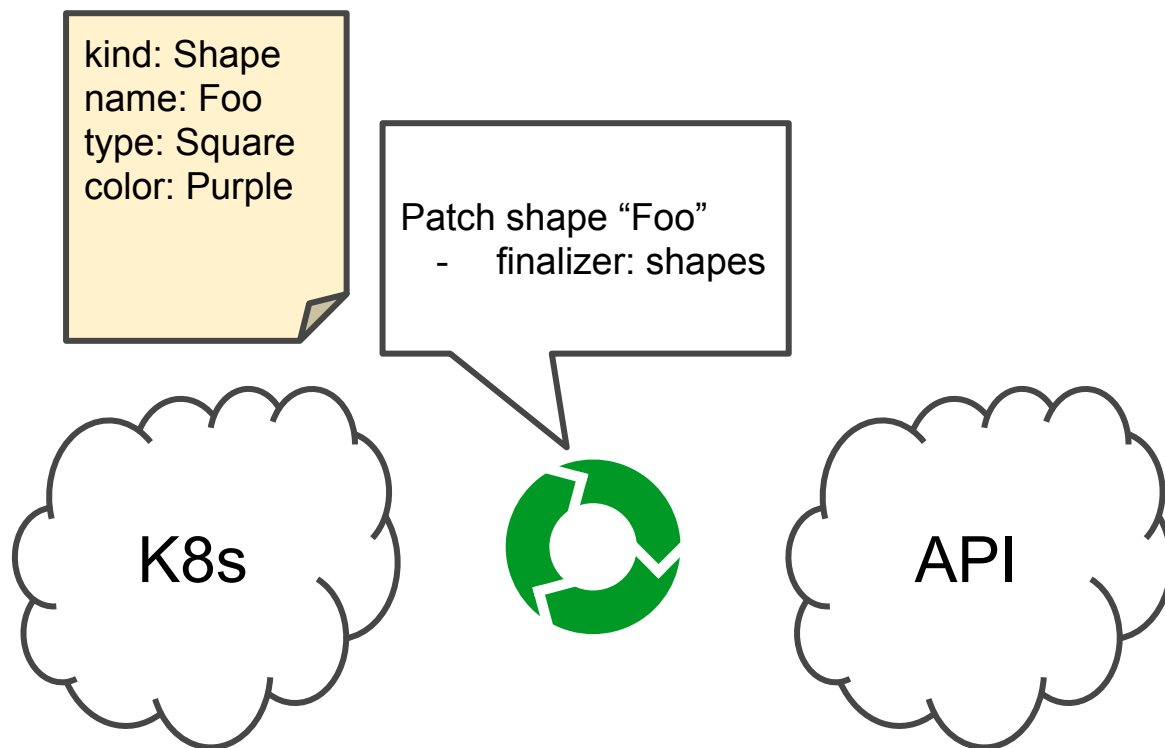








kind: Shape
name: Foo
type: Square
color: Purple





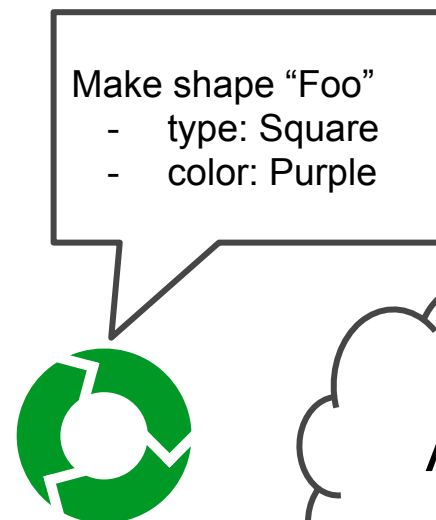



kind: Shape
name: Foo
type: Square
color: Purple
finalizers:
- shapes








kind: Shape
name: Foo
type: Square
color: Purple
finalizers:
- shapes







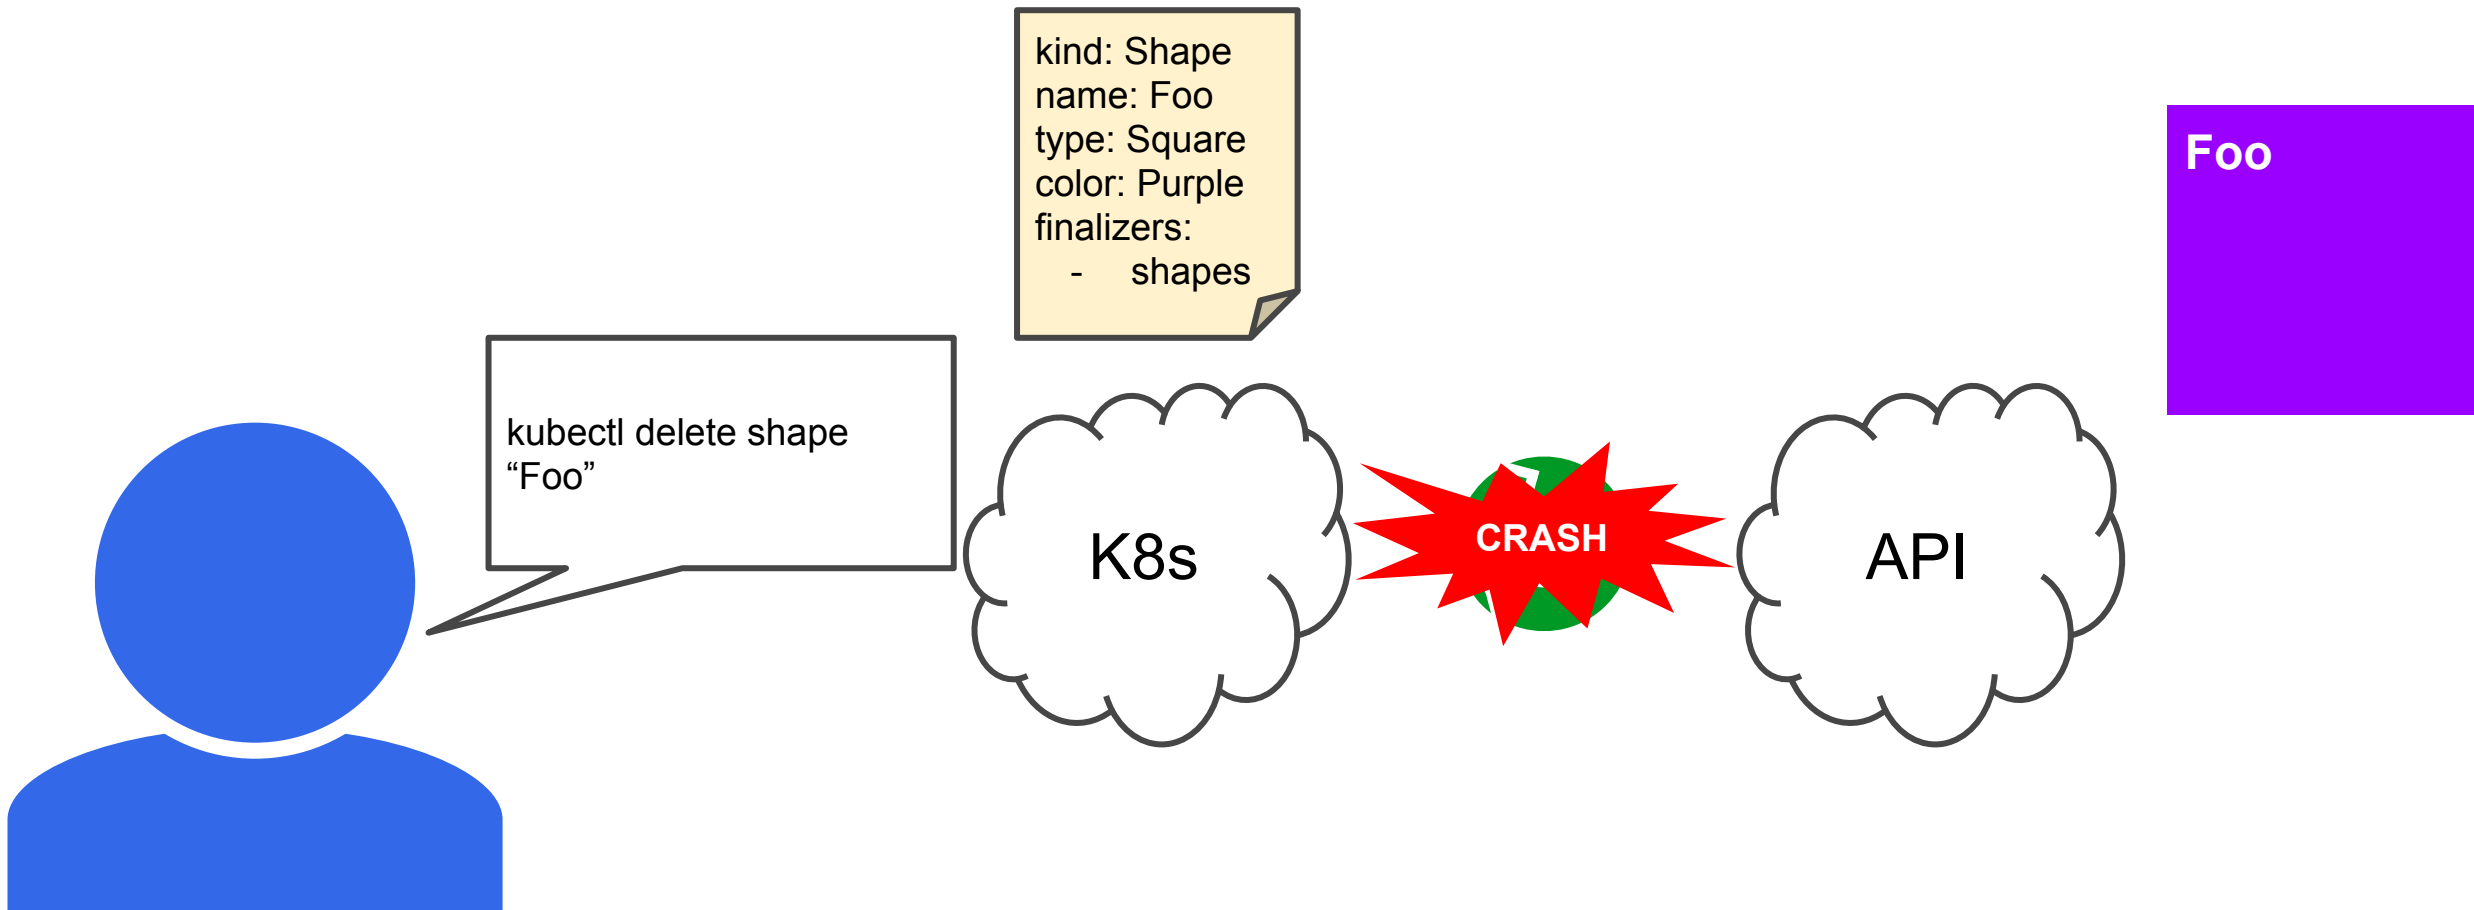
```
kind: Shape
name: Foo
type: Square
color: Purple
finalizers:
- shapes
```



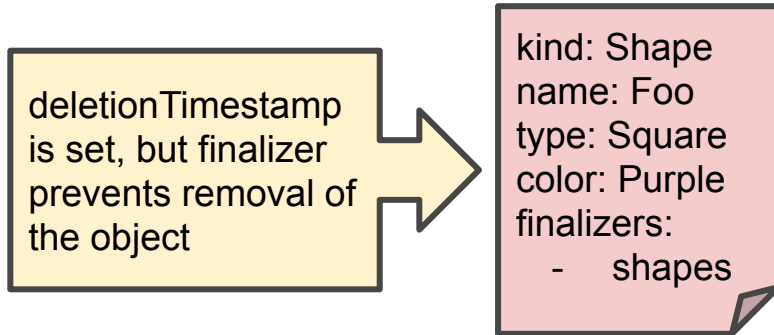




kind: Shape
name: Foo
type: Square
color: Purple
finalizers:
- shapes












kind: Shape
name: Foo
type: Square
color: Purple
finalizers:
- shapes






```
kind: Shape
name: Foo
type: Square
color: Purple
finalizers:
- shapes
```




The controller observes the pending deletion of the shape.



```
kind: Shape
name: Foo
type: Square
color: Purple
finalizers:
- shapes
```



Delete shape "Foo"






Foo



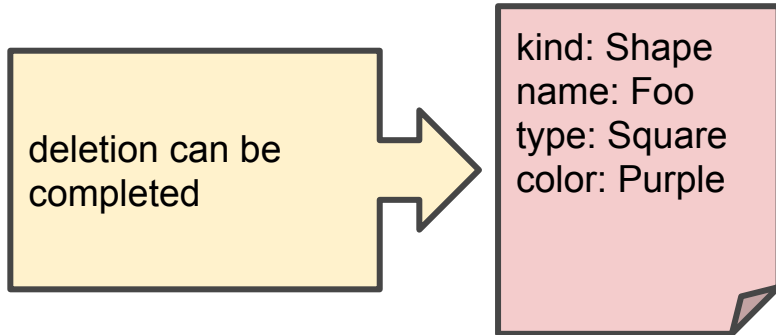
labels:
owner: k8s





kind: Shape
name: Foo
type: Square
color: Purple
finalizers:
- shapes














CustomResources

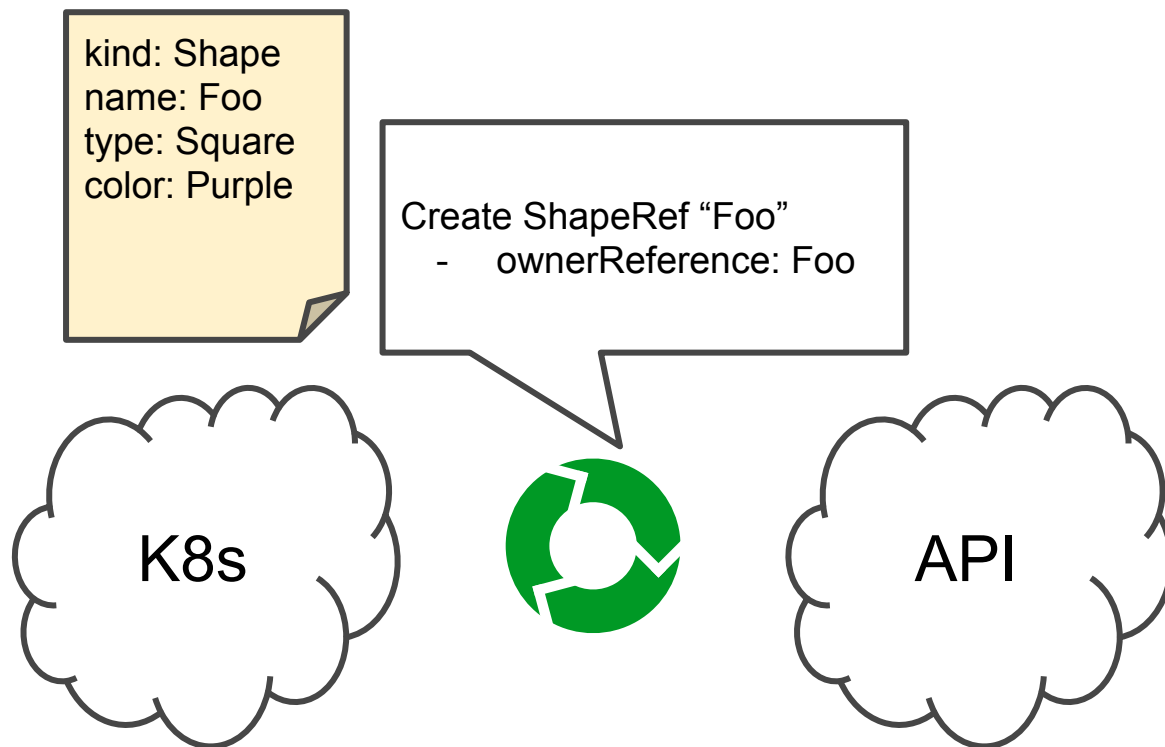




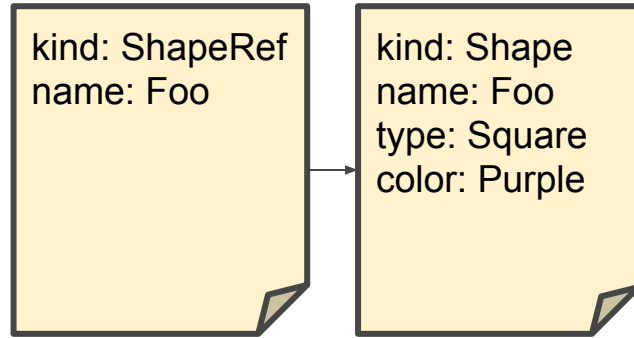




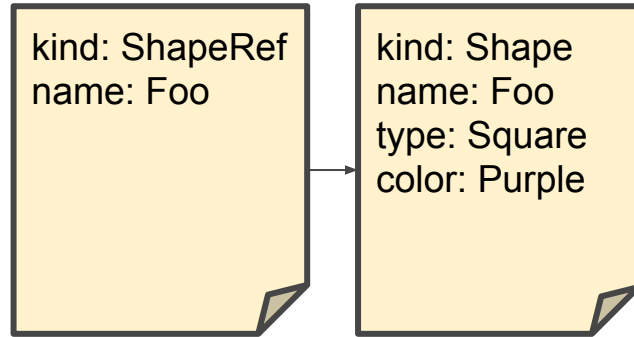




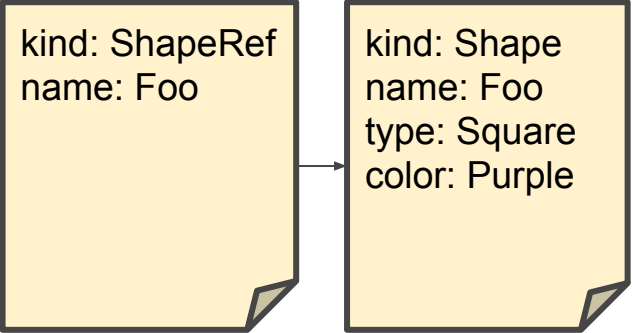
kind: Shape
name: Foo
type: Square
color: Purple


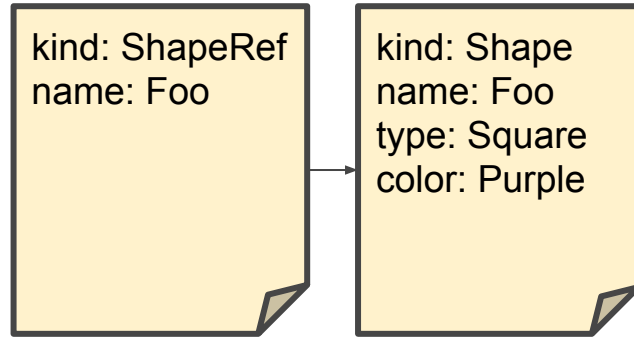


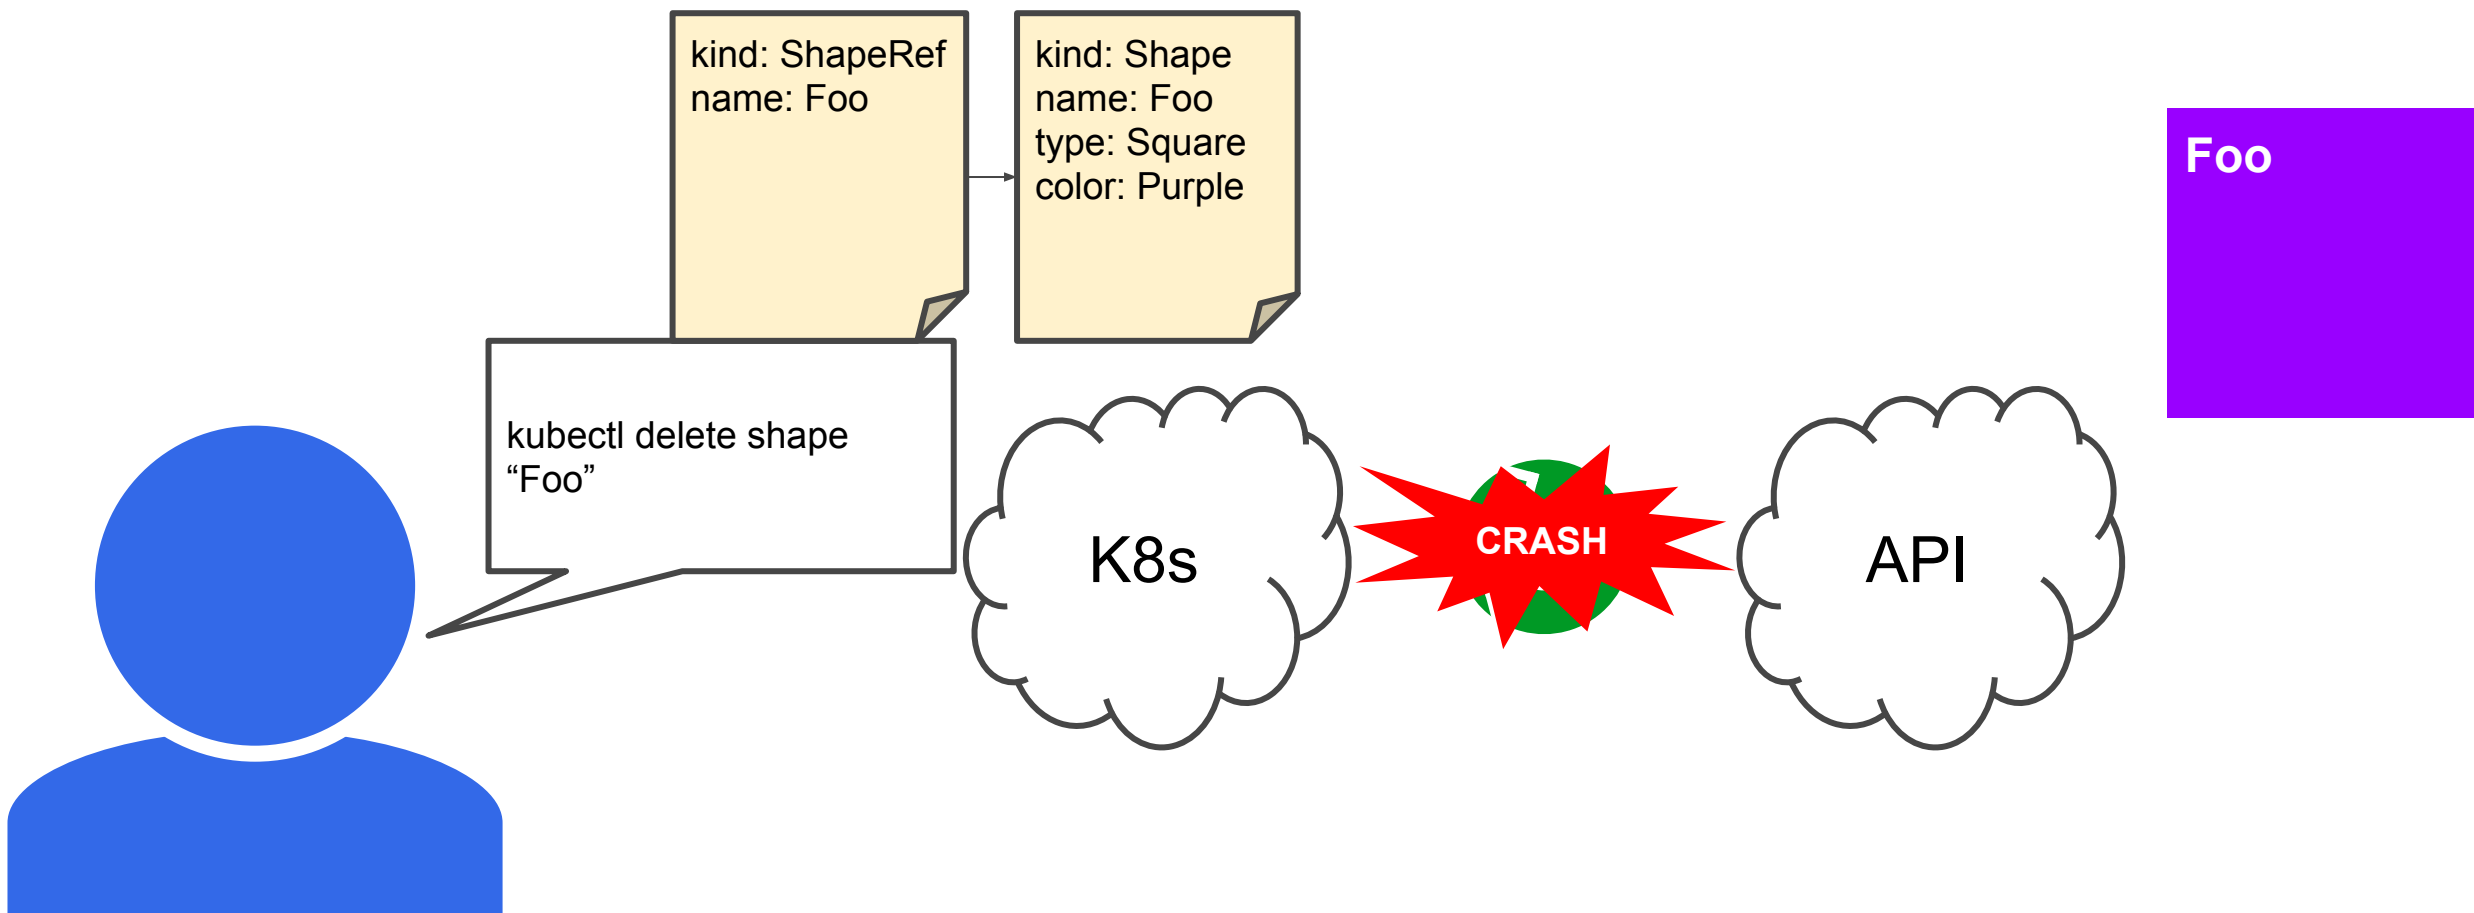


















kind: ShapeRef
name: Foo













The controller did not observe the deletion of the Shape,
but it does observe the dangling ShapeRef.






kind: ShapeRef
name: Foo



Foo
labels:
owner: k8s







In most of these mechanisms, there's some amount of
“you broke it, you bought it”.

If a user deletes the ShapeRef or removes the finalizer or
edits the underlying metadata, the linkage can be broken.

You broke it, you get to keep the pieces.