


Materialized Views (MV)



Why Materialized Views ?

- Relieve the pain of manual denormalization

```
CREATE TABLE user(  
  id int PRIMARY KEY,  
  country text,  
  ...  
);  
CREATE TABLE user_by_country(  
  country text,  
  id int,  
  ...,  
  PRIMARY KEY(country, id)  
);
```

Materialized View In Action

```
CREATE MATERIALIZED VIEW user_by_country  
AS SELECT country, id, firstname, lastname  
FROM user  
WHERE country IS NOT NULL AND id IS NOT NULL  
PRIMARY KEY(country, id)
```



```
CREATE TABLE user_by_country (  
  country text,  
  id int,  
  firstname text,  
  lastname text,  
  PRIMARY KEY(country, id));
```

Materialized View Syntax

CREATE MATERIALIZED VIEW [IF NOT EXISTS]

keyspace_name.view_name

Must select **all primary key columns** of base table

AS SELECT column₁, column₂, ...

FROM keyspace_name.table_name

WHERE column₁ IS NOT NULL AND column₂ IS NOT NULL ...

PRIMARY KEY(column₁, column₂, ...)

- **IS NOT NULL** condition for now
- more complex conditions in future


- at least **all primary key columns** of base table (ordering can be different)
- **maximum 1 column NOT pk** from base table



Materialized Views Demo


Materialized View Impl

UPDATE user
SET country='FR'
WHERE id=1



Materialized View Impl


UPDATE user
SET country='FR'
WHERE id=1



②
Acquire **local lock** on
base table partition

Materialized View Impl

UPDATE user
SET country='FR'
WHERE id=1




③

Local read to fetch current values
`SELECT * FROM user WHERE id=1`

Materialized View Impl

UPDATE user
SET country='FR'
WHERE id=1




④

Create **BatchLog** with

- DELETE FROM user_by_country
WHERE country = 'old_value'
- INSERT INTO
user_by_country(country, id, ...)
VALUES('FR', 1, ...)

Materialized View Impl


UPDATE user
SET country='FR'
WHERE id=1



⑤
Execute **async BatchLog**
to **paired view replica**
with CL = **ONE**

Materialized View Impl


UPDATE user
SET country='FR'
WHERE id=1



⑥

Apply base table update locally
SET COUNTRY='FR'

Materialized View Impl




UPDATE user
SET country='FR'
WHERE id=1

⑦
Release local lock


Materialized View Impl

UPDATE user
SET country='FR'
WHERE id=1



Materialized View Impl


UPDATE user
SET country='FR'
WHERE id=1



Materialized Views impl explained

What is paired replica ?

- Base primary replica for **id=1** is paired with MV primary replica for **country='FR'**
- Base secondary replica for **id=1** is paired with MV secondary replica for **country='FR'**
- etc ...



Materialized Views impl explained

Why **local lock** on base replica ?

- to provide **serializability** on concurrent updates

Why **BatchLog** on base replica ?

- necessary for **eventual durability**
- replay the MV delete + insert until successful

Why **BatchLog** on base replica uses **CL = ONE** ?

- each base replica is responsible for update of its **paired** MV replica
- $CL > ONE$ will create un-necessary **duplicated mutations**


MV Failure Cases: concurrent updates *without local lock*

1) UPDATE ... SET country='US'

- Read base row (country='UK')
- DELETE FROM mv WHERE country='UK'
- INSERT INTO mv ...(country) VALUES('US')
- Send async **BatchLog**
- Apply update country='US'

2) UPDATE ... SET country='FR'

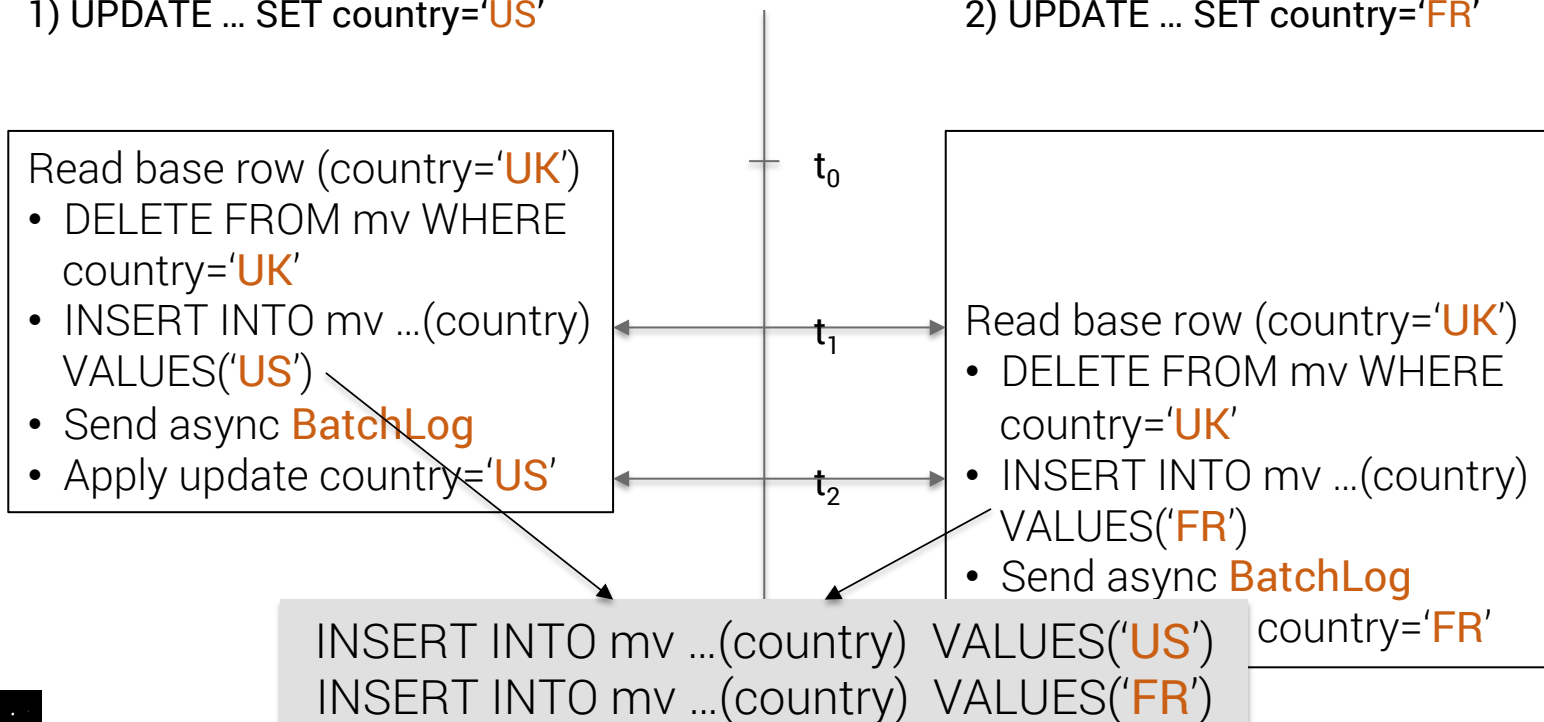
- Read base row (country='UK')
- DELETE FROM mv WHERE country='UK'
- INSERT INTO mv ...(country) VALUES('FR')
- Send async **BatchLog**
- Apply update country='FR'



MV Failure Cases: concurrent updates *without local lock*

1) UPDATE ... SET country='US'

2) UPDATE ... SET country='FR'



MV Failure Cases: concurrent updates *with local lock*

1) UPDATE ... SET country='US'



Read base row (country='UK')

- DELETE FROM mv WHERE country='UK'
- INSERT INTO mv ...(country) VALUES('US')
- Send async **BatchLog**
- Apply update country='US'



2) UPDATE ... SET country='FR'



Read base row (country='US')

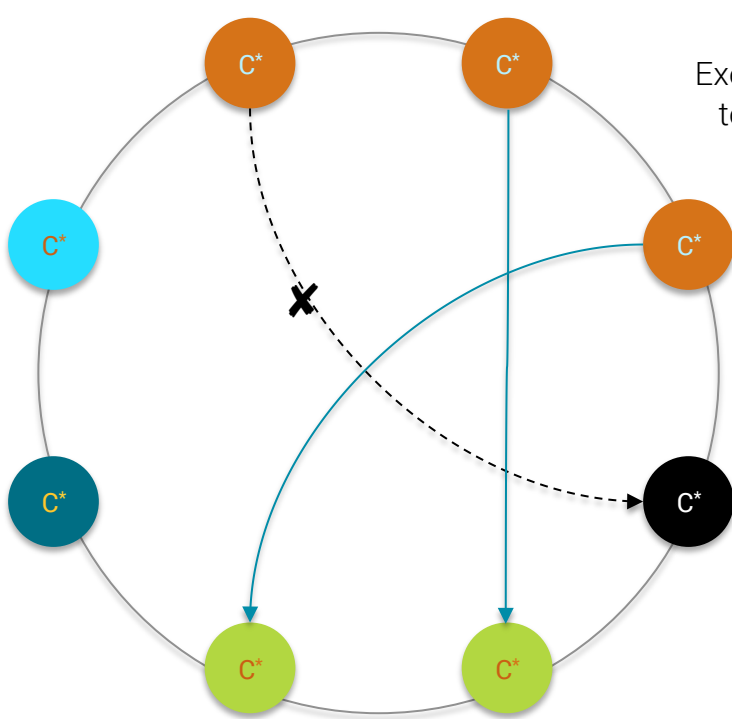
- DELETE FROM mv WHERE country='US'
- INSERT INTO mv ...(country) VALUES('FR')
- Send async **BatchLog**
- Apply update country='FR'



21

@doanduyhai

MV Failure Cases: failed updates to MV




⑤ Execute **async BatchLog** to **paired view replica** with CL = **ONE**

MV replica down

```
UPDATE user
SET country='FR'
WHERE id=1
```

MV Failure Cases: failed updates to MV

UPDATE user
SET country='FR'
WHERE id=1



MV replica up

Materialized View Performance

- Write performance
 - local lock
 - local read-before-write for MV → update contention on partition (*most of perf hits*)
 - local batchlog for MV
 - 📖 you only pay this price **once** whatever number of MV
 - for each base table update: **mv_count x 2** (*DELETE + INSERT*) extra mutations

Materialized View Performance

- Write performance vs **manual denormalization**
 - MV better because *no client-server network traffic for read-before-write*
 - MV better because *less network traffic for multiple views (client-side BATCH)*
- Makes developer life easier → **priceless**

Materialized View Performance

- Read performance vs **secondary index**
 - MV better because **single node read** (secondary index can hit many nodes)
 - MV better because **single read path** (secondary index = *read index + read data*)

Materialized Views Consistency

- Consistency level
 - CL honoured for base table, **ONE** for MV + local batchlog
- **Weaker consistency guarantees** for MV than for base table.
- Exemple, write at **QUORUM**
 - guarantee that **QUORUM** replicas of base tables have received write
 - guarantee that **QUORUM** of MV replicas will **eventually** receive *DELETE + INSERT*

Materialized Views Gotchas

- Beware of hot spots !!! MV **user_by_gender** 🤯
- Only 1 non-pk column for MV
- No **static column** for MV view
 - *1:1 relationship* between static column & partition key
 - if MV supports static column → MV has same partition key as base table → useless ...

Materialized Views Operations

- Repair, read-repair, index rebuild, decommission ...
 - repair on base replica (*mutation-based* repair) → update on MV paired replica
 - possible to repair a MV *independently from base table*
 - read-repair on MV behaves as normal read-repair
 - read-repair on base table updates MV
 - hints replay on base table updates MV

Materialized Views Operations

- MV build status ?
 - system.views_builds_in_progress
 - system.built_views
 - data are **node-local** !!!

```
cqlsh:system> select * from system.views_builds_in_progress;
```

keyspace_name	view_name	generation_number	last_token
---------------	-----------	-------------------	------------

(0 rows)

```
cqlsh:system> select * from system.built_views ;
```

keyspace_name	view_name
---------------	-----------

music	albums_by_year
-------	----------------

music	artists_by_country
-------	--------------------

Materialized Views Schema Ops

- Schema
 - MV can be tuned as normal table (**ALTER MATERIALIZED VIEW ...**)
 - cannot drop column from base table used by MV
 - can add column to base table, initial value = null from MV
 - cannot drop base table, drop all MVs first

Single non-pk column limitation

- Because of Cassandra consistency model
- Because null value forbidden for primary key column

```
CREATE MATERIALIZED VIEW user_by_gender_and_age
AS SELECT country, id, firstname, lastname
FROM user
WHERE gender IS NOT NULL AND age IS NOT NULL AND id IS NOT NULL
PRIMARY KEY((gender, age) id)
```

Single non-pk column limitation

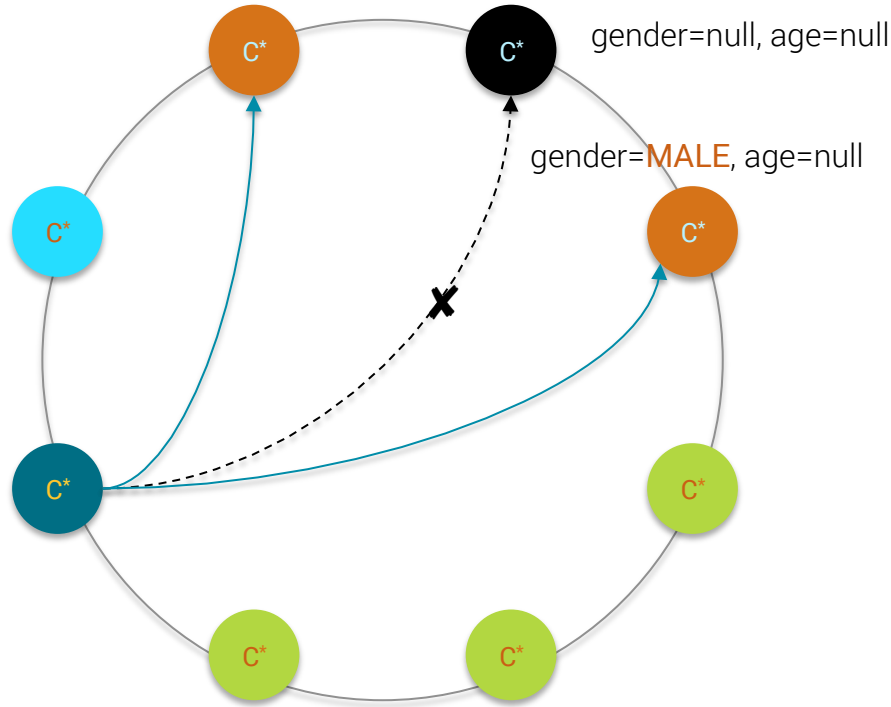
- Possible RULE : **UPDATE MV ONLY IF ALL COLUMNS IN PK NOT NULL**

Multiple non-PK columns in MV PK

gender=MALE, age=null

CL = QUORUM

```
UPDATE user  
SET gender='MALE'  
WHERE id=1
```

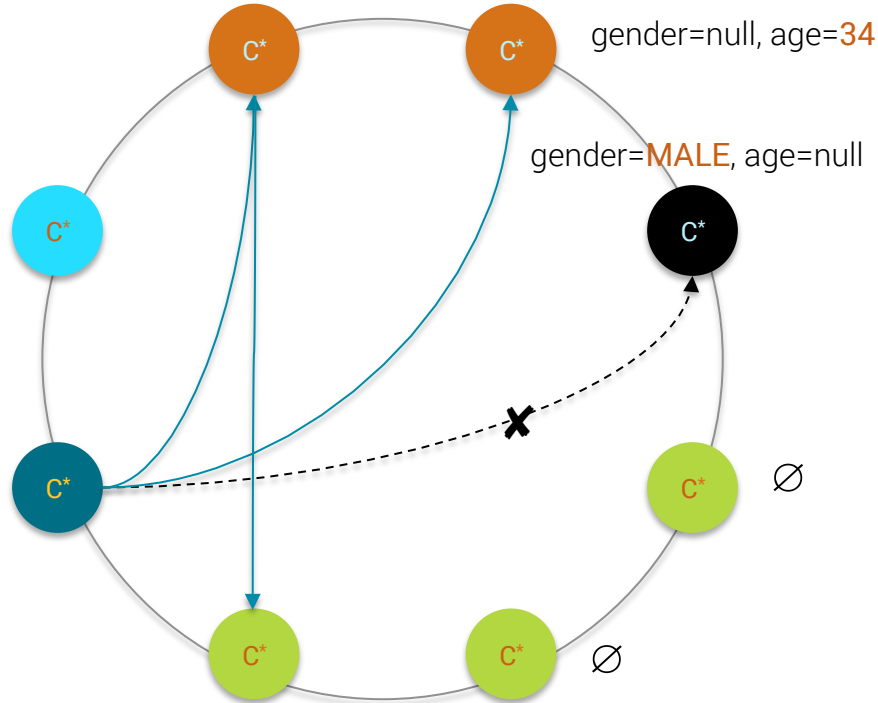


Multiple non-PK columns in MV PK

gender=MALE, age=34

CL = QUORUM

```
UPDATE user  
SET age=34  
WHERE id=1
```



gender=MALE, age=34

Multiple non-PK columns in MV PK


CL = QUORUM

```
SELECT age,gender  
FROM user  
WHERE id=1
```

1

34

MALE



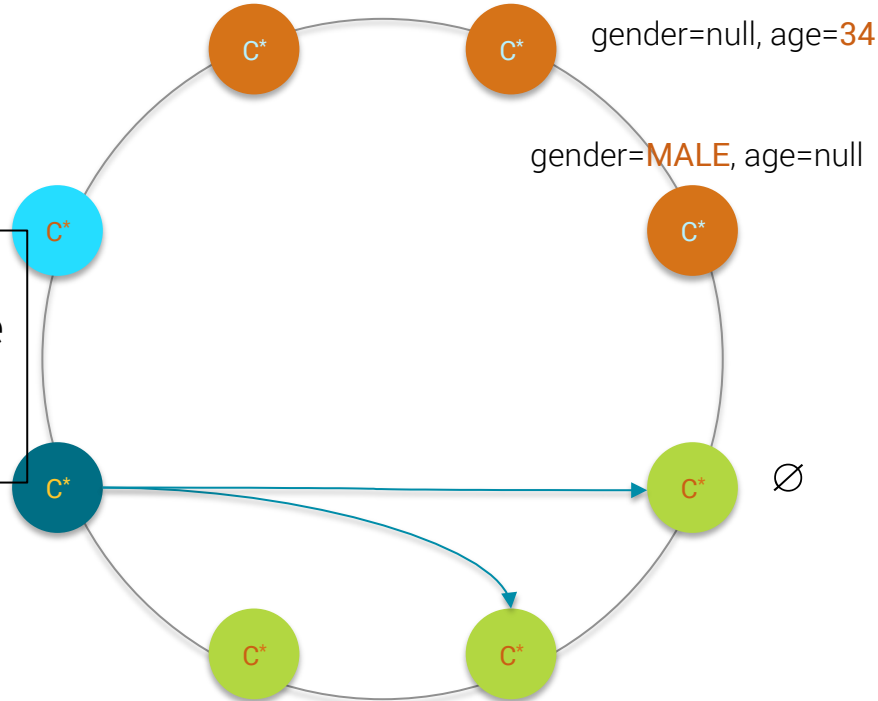
gender=MALE, age=34

Multiple non-PK columns in MV PK

gender=MALE, age=34

CL = QUORUM

```
SELECT * FROM  
user_by_gender_and_age  
WHERE gender='MALE'  
AND age=34
```



gender=MALE, age=34

Single non-pk column limitation

- Possible RULE 2: **ALLOW NULL VALUE FOR COLUMN IN PK**

```
INSERT INTO user(id, name) VALUES(1, 'John DOE');
```

```
...
```

```
...
```

```
...
```

```
INSERT INTO user(id, name) VALUES(1000_000, 'Helen SUE');
```

No age, No gender

Single non-pk column limitation

- Possible RULE 2: **ALLOW NULL VALUE FOR COLUMN IN PK**

(*null*, *null*) single partition with **10^6 users**



Living in Danger