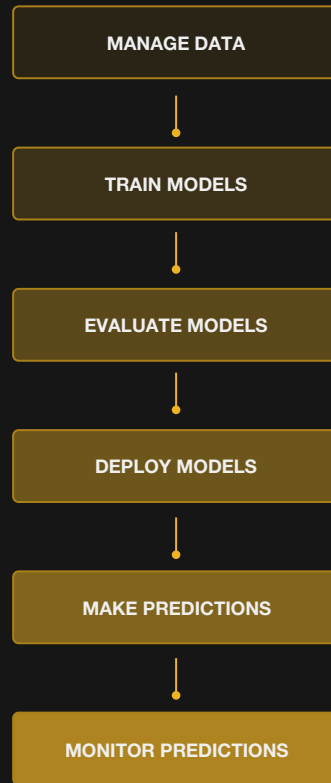


Michelangelo @ Uber

Enable engineers and data scientists across the company to easily build and deploy machine learning solutions at scale.



ML-as-a-service

- Managing Data/Features
- Tools for managing, end-to-end, heterogeneous training workflows
- Batch, online & mobile serving
- Feature and Model drift monitoring



Feature Engineering @ Uber

- Example: ETA for EATS order
- Key ML *features*
 - How large is the order?
 - How busy is the restaurant?
 - How quick is the restaurant?
 - How busy is the traffic?



Managing Features

One of the hardest problems in ML

- Finding good Features & labels
- Data in production: reliability, scale, low latency
- Data parity: training/serving skew
- Real-time features: traditional tools don't work

Palette Feature Store

Uber-specific *curated* and *crowd-sourced* feature database that is easy to use with machine learning projects.

One stop shop

- Search for features in single catalog/spec: *rider, driver, restaurant, trip, eaters, etc.*
- Define new features + create production pipelines from spec
- Share features across Uber: cut redundancy, use consistent data
- Enable tooling: Data Drift Detection, Auto Feature Selection, etc.

Feature Store Organization

Organized as <entity>:<feature-group>:<feature-name>:<join-key>

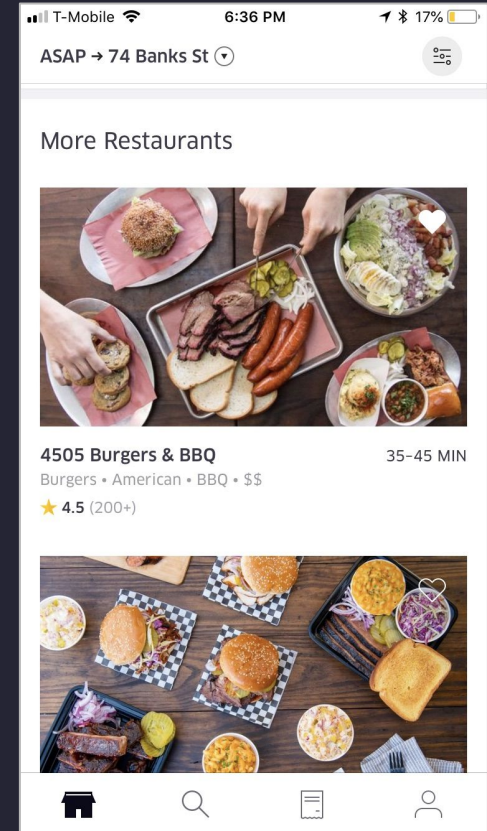
Eg. @palette:restaurant:realtime_group:orders_last_30min:restaurant_uuid

Backed by a dual datastore system: similarities to lambda

- Offline
 - Offline (Hive based) store for bulk access of features
 - Bulk retrieval of features across time
- Online
 - KV store (Cassandra) for serving latest known value
 - Supports lookup/join of latest feature values in real time
- Data synced between online & offline
 - Key to avoiding training/serving skew

EATS Features revisited

- How large is the order? ← Input
- How busy is the restaurant?
- How quick is the restaurant?
- How busy is the traffic?



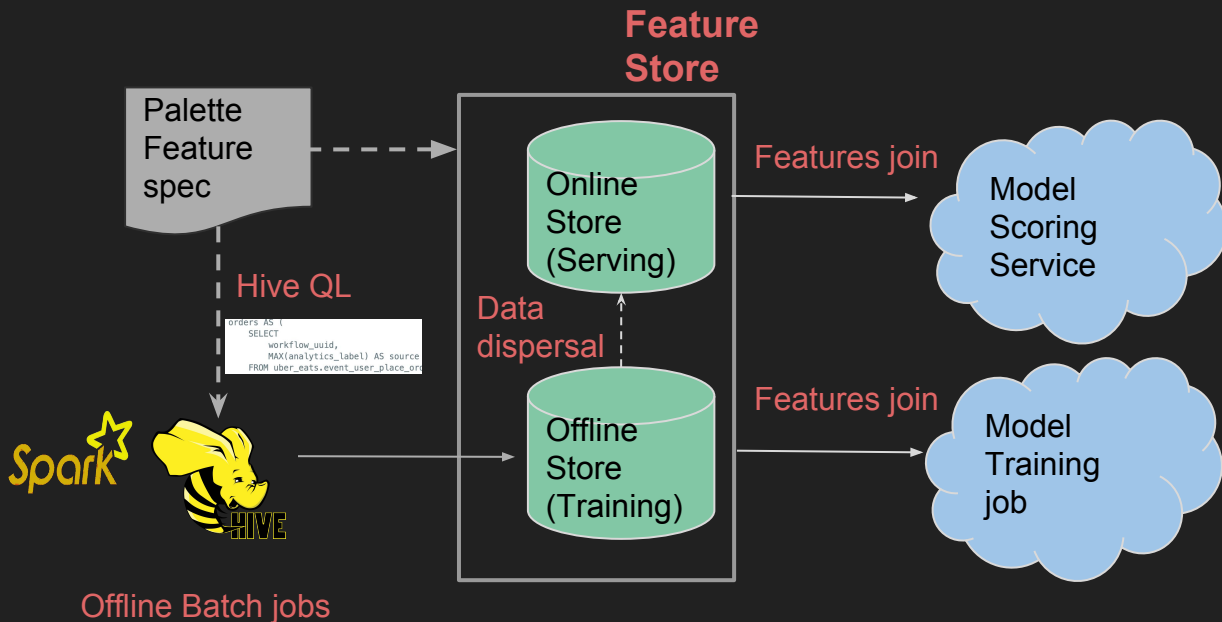
Creating Batch Features

General trends, not sensitive to exact time of event

Ingested from Hive queries or Spark jobs

How quick is the restaurant ?

- Aggregate trends
- Use Hive QL from warehouse
- @palette:restaurant:batch_aggr:
prepTime:rlid



Apache Hive and Apache Spark are either registered trademarks or trademarks of the Apache Software Foundation in the United States and/or other countries. No endorsement by The Apache Software Foundation is implied by the use of this mark.

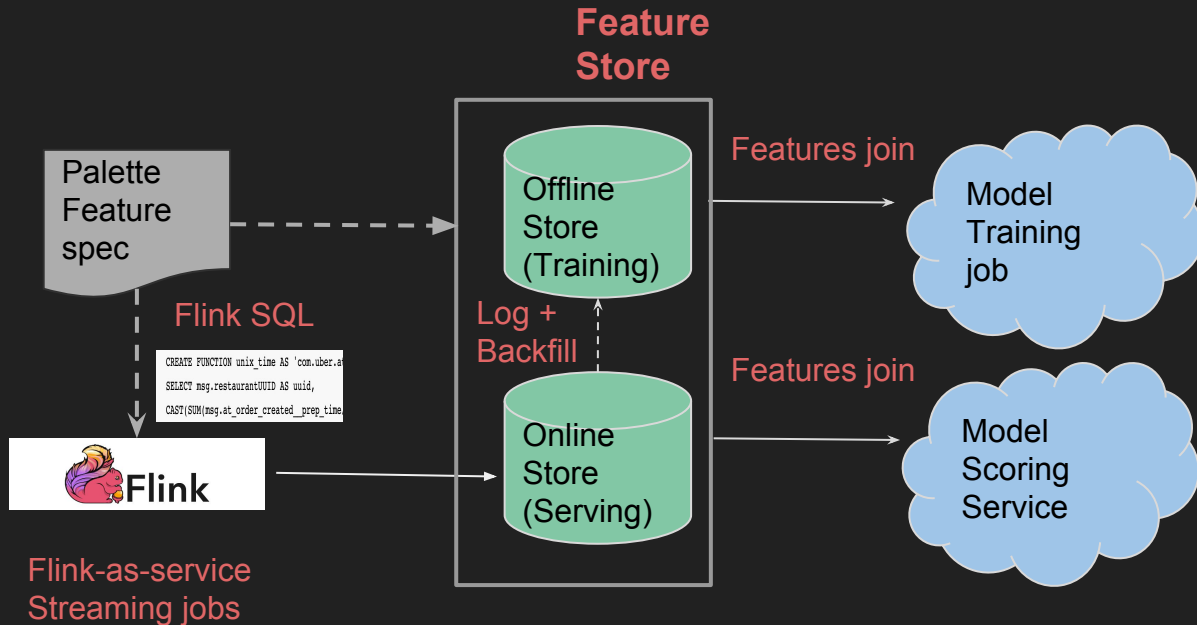
Creating Real-time Features

Features reflecting the latest state of the world

Ingest from streaming jobs

How busy is the restaurant ?

- kafka topic with events
- perform realtime aggregations
- @palette:restaurant:rt_aggr:**nMeal**:rld



Apache Flink is either a registered trademark or trademark of the Apache Software Foundation in the United States and/or other countries. No endorsement by The Apache Software Foundation is implied by the use of this mark.

Bring Your Own Features

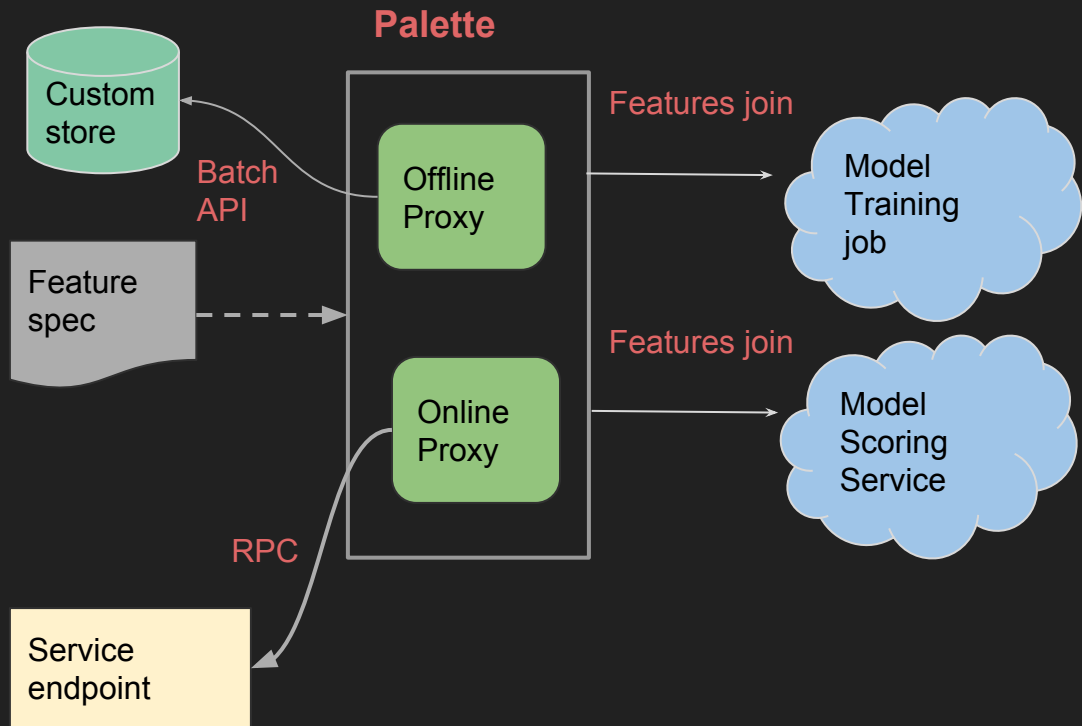
Feature maintained by customers

Mechanisms for hooking
serving/training endpoints

Users maintain data parity

How busy is the region ?

- RPC: external traffic feed
- Log RPCs for training
- `@palette:region:traffic:nBusy:regionId`

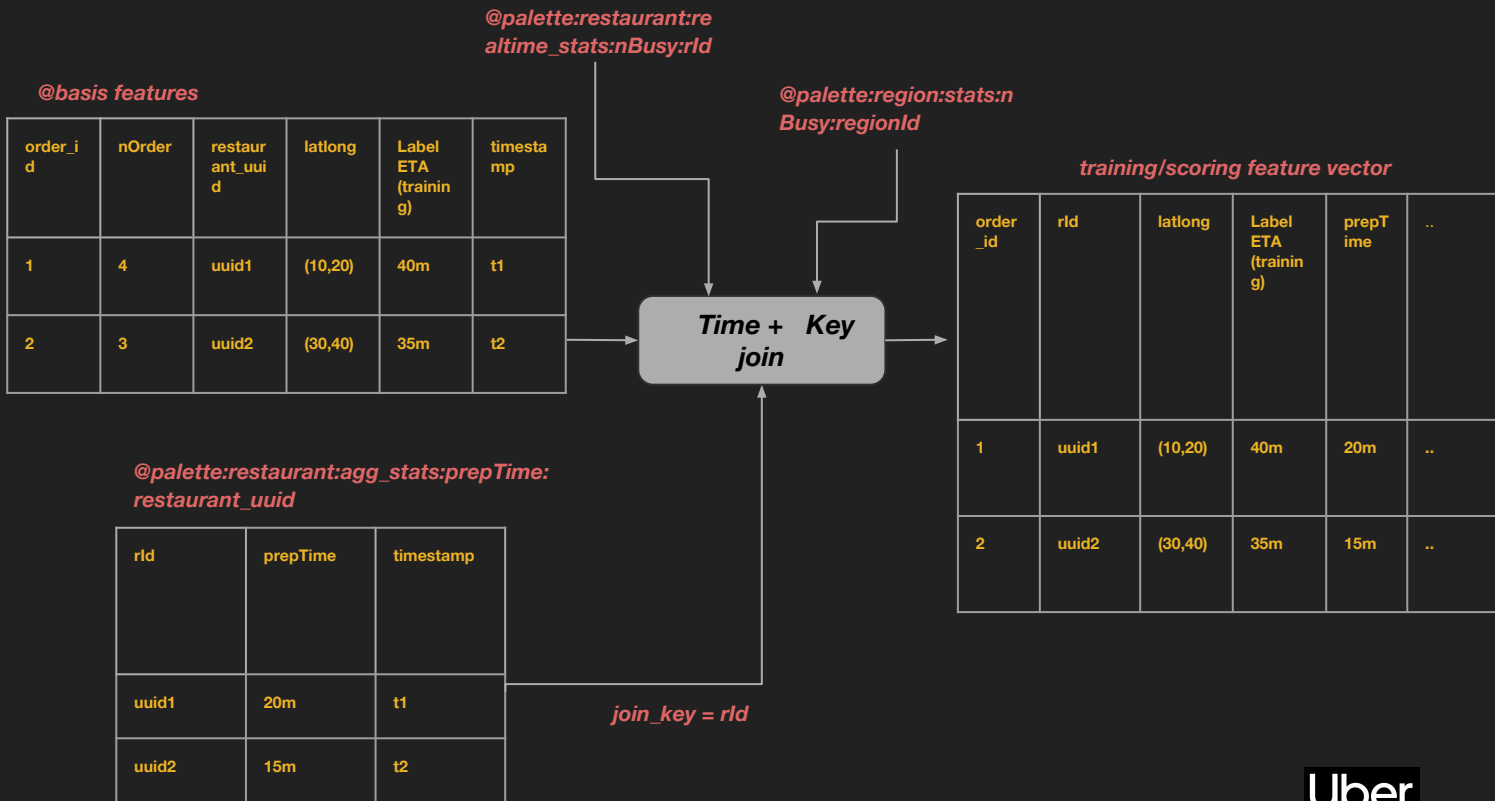


Palette Feature Joins

Join *@basis* features
with supplied
@palette features into
single feature vector

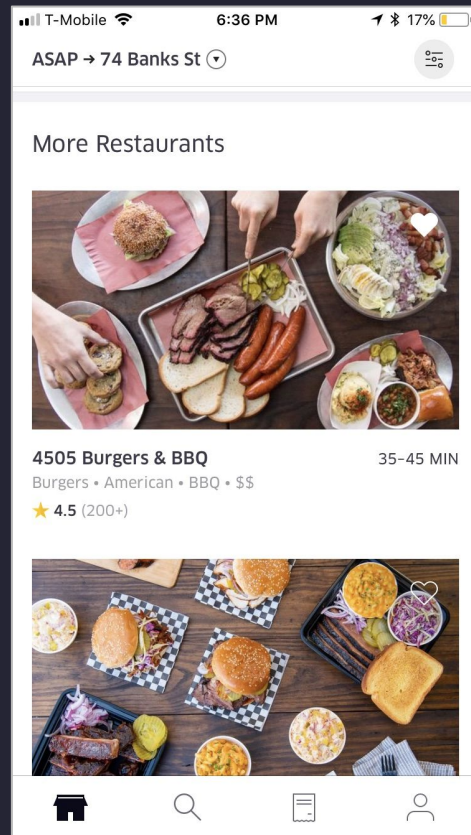
Join billion+ rows at
points-in-time:
dominates overhead

Join/Lookup 10s of
tables at serving time
at low latency



Done with Feature Engineering ?

- Feature Store Features
 - *nOrder*: How large is the order? (*basis*)
 - *nMeal*: How busy is the restaurant? (*near real-time*)
 - *prepTime*: How quick is the restaurant? (*batch feature*)
 - *nBusy*: How busy is the traffic? (*external feature*)
- Ready to use ?
 - Model specific feature transformations
 - Chaining of features
- **Feature Transformers**



Feature Consumption

Feature Store Features

- *nOrder*: input feature
- *nMeal*: consume directly
- *prepTime*: needs transformation before use
- *nBusy*: input latlong but need regionId

Setting up consumption pipelines

- *nMeal*: *r_id* -> *nMeal*
- *prepTime*: *r_id* -> *prepTime* -> *featureImpute*
- *nBusy*: *r_id* -> *lat, log* -> *regionId(lat, log)* -> *nBusy*

In arbitrary order

Michelangelo Transformers

Transformer: Given a record defined a set of fields, add/modify/remove fields in the record

PipelineModel: A sequence of transformers

Spark ML: Transformer/PipelineModel on DataFrames

Michelangelo Transformers: extended transformer framework for both Apache Spark and Apache Spark-less environments

Estimator: Analyze the data and produce a transformer

Defining a Pipeline Model

Feature consumption

- Feature extraction: Palette feature retrieval expressed as a transform
- Feature mutation: Scala-like DSL for simple transforms
- Model-centric Feature Engineering: string indexer, one-hot encoder, threshold decision
- Result retrieval

Modeling

- Model inferencing (also Michelangelo Transformer)

Join Palette Features

Apply Feature Eng Rules

String Indexing

One-Hot Encoding

DL Inferencing

Result Retrieval

Michelangelo Transformers Example

```
class MyEstimator(override val uid: String) extends  
  Estimator[MyEstimator] with Params with DefaultParamsWritable {  
  
  ...  
  
  override def fit(dataset: Dataset[_]): MyModel = ...  
  
}
```

```
class MyModel (override val uid: String) extends Model[MyModel] with  
  MyModelParam with MLWritable with MATransformer {
```

...

```
override def transform(dataset: Dataset[_]): DataFrame = ...
```

```
override def scoreInstance(instance: util.Map[String, Object]): util.Map[String,  
Object] = ...
```

```
}
```

Palette retrieval as a Transformer

```
tx_p1 = PaletteTransformer([
    "@palette:restaurant:realtime_feature:nMeal:r_id",
    "@palette:restaurant:batch_feature:prepTime:r_id",
    "@palette:restaurant:property:lat:r_id",
    "@palette:restaurant:property:log:r_id"
])

tx_p2 = PaletteTransformer([
    "@palette:region:service_feature:nBusy:region_id"
])
```

Palette Feature Transformer

Hive Access

Cassandra Access

RPC Feature Proxy

Feature Meta Store

DSL Estimator / Transformer

```
es_dsl1 = DSLEstimator(lambdas = [  
    ["region_id", "regionId(@palette:restaurant:property:lat:r_id,  
        @palette:restaurant:property:r_id")  
])  
  
es_dsl2 = DSLEstimator(lambdas = [  
    ["prepTime": nFill(nVal("@palette:restaurant:batch_feature:prepTime:r_id"),  
        avg("@palette:restaurant:batch_feature:prepTime:r_id"))],  
    ["nMeal": nVal("@palette:restaurant:realtime_feature:nMean:r_id")],  
    ["nOrder": nVal("@basis:nOrder")],  
    ["nBusy": nVal("@palette:region:service_feature:nBusy:region_id")]  
])
```

DSL Estimator

Code Gen / Compiler

DSL Transformer

Online classloader

Offline classloader

Uber Eats Example Cont.

Computation order

- *nMeal: rld -> nMeal*
- *prepTime: rld -> prepTime -> featureImpute*
- *busyScale: rld -> lat, log -> regionId(lat, log) -> busyScale*

Palette Transformer
id -> nMeal
id -> prepTime
id -> lat, log

DSL Transformer
lag, log -> regionID

Palette Transformer
regionID -> nBusy

DSL Transformer
impute(nMeal)
impute(preptime)

Dev Tools: Authoring and Debugging a Pipeline

Palette feature generation

- Apache Hive QL, Apache Flink SQL

Interactive authoring

- PySpark + iPython Jupyter notebook

Centralized model store

- Serialization / Deserialization (Spark ML, MLReadable/Writeable)
- Online and offline accessibility

```
basis_feature_sql = "..."  
df = spark.sql(basis_feature_sql)
```

```
pipeline = Pipeline(stages=[tx_p1, es_dsl1, tx_p2, es_dsl2t, vec_asm, l_r)  
pipeline_model = pipeline.fit(df)  
scored_def = pipeline_model.transform(df)
```

```
model_id = MA_store.save_model(pipeline_model)
```

```
draft_id = MA_store.save_pipeline(basis_feature_sql, pipeline)  
retrain_job = MA_API.train(draft_id, new_basis_feature_sql)
```

Takeaways

Feature Store: Batch, Realtime and External Features with online and offline parity

Offline scalability: Joins across billions of rows

Online serving latency: Parallel IO, fast storage with caching

Feature Transformers: Setup chains of transformations at training/serving time

Pipeline reliability and monitoring out-of-the-box