



# SQL Tuning Tips You Can't Do Without

---

**Maria Colgan**

Distinguished Product Manager  
Database Server Technologies  
May 2021

 @SQLMaria



# Disclaimer

---

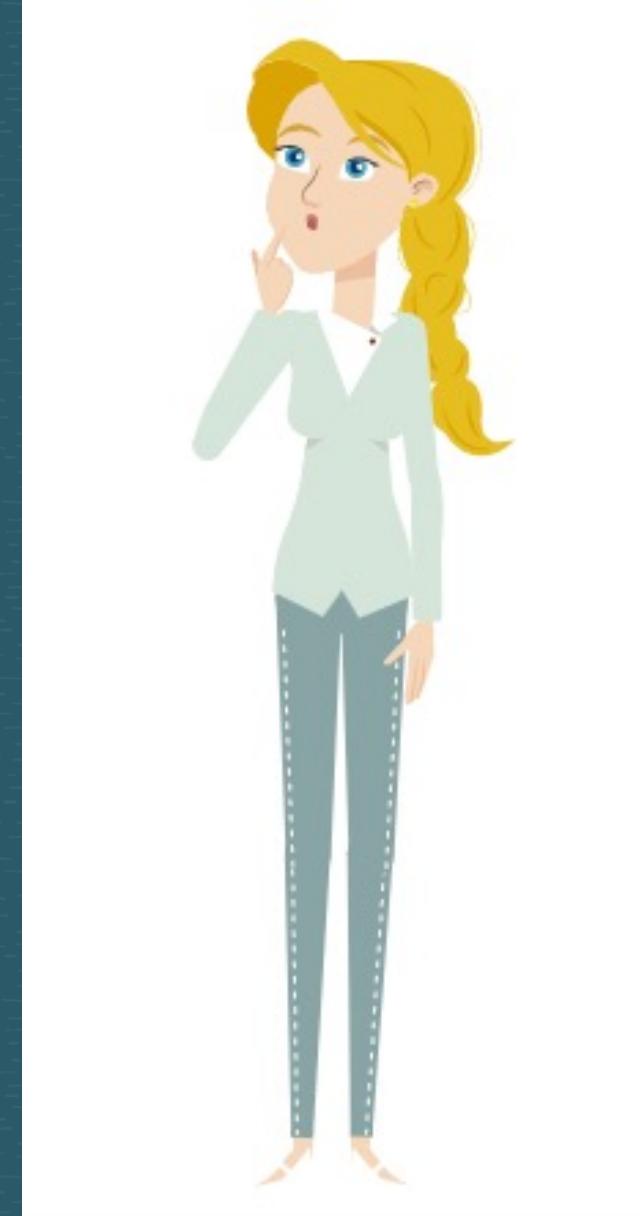
The goal of this session to provide you tips on how to correct common problems identified from a SQL execution plans

This session will not make you an Optimizer expert instantly or give you the power to tune SQL statements with the flick of your wrist!

# Problem Statement 1

---

**Why Didn't The Optimizer Pick  
The Join Method I Expected?**



# What Join Method Should The Optimizer Pick?

## The Query

```
SELECT      p.prod_name, SUM(s.quantity_sold)
FROM        sales s, products p
WHERE       s.prod_id = p.prod_id
AND         s.time_id = '03-MAY-21'
GROUP BY    p.prod_name;
```

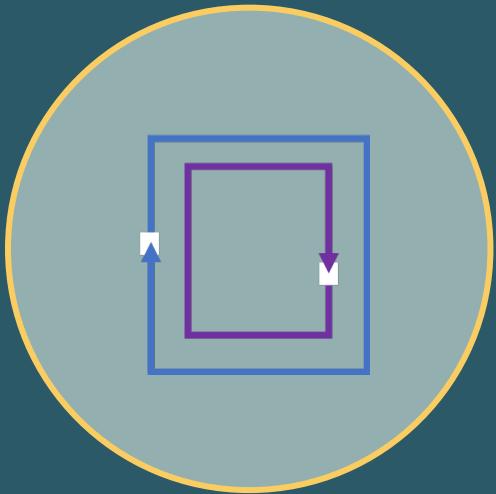
SALES 

SALES table has  
10 million rows

PRODUCTS

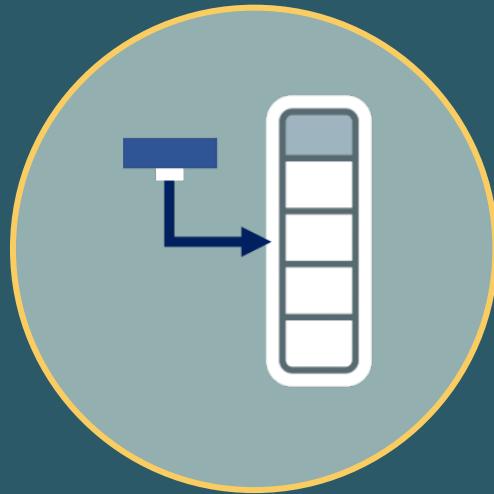
PRODUCTS table  
has 100,000 rows

# Oracle Optimizer Has 3 Join Methods



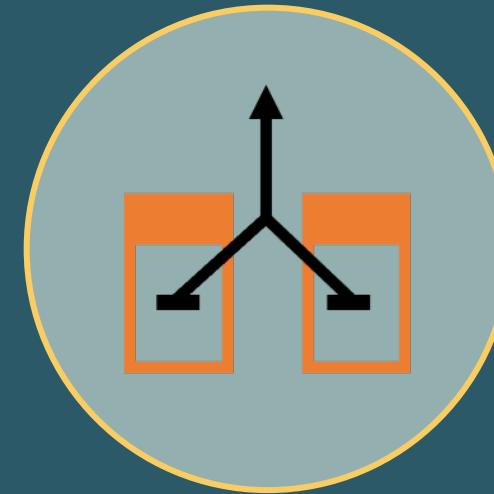
## NESTED LOOPS

For every row in the outer table, Oracle accesses all the rows in the inner table



## HASH JOIN

Smaller of two tables is scan and resulting rows used to build a hash table on the join key in memory. Larger table is then scan, join column of the resulting rows are hashed and the values used to probe the hash table to find matching rows



## SORT MERGE JOIN

Consists of two steps:  
1. Sort both inputs on the join key  
2. Merge the sorted lists together

# Optimizer Chose A NESTED LOOP Instead of HASH JOIN Join

```
SELECT      p.prod_name, SUM(s.quantity_sold)
FROM        sales s, products p
WHERE       s.prod_id = p.prod_id
AND         s.time_id = '03-MAY-21';
GROUP BY    p.prod_name;
```

Id	Operation	Name	Starts	Rows
0	SELECT STATEMENT		1	
1	HASH GROUP BY		1	1
2	NESTED LOOPS		1	1
3	NESTED LOOPS		1	1
4	PARTITION RANGE SINGLE		1	1
*	5	TABLE ACCESS STORAGE FULL	0	1
*	6	INDEX UNIQUE SCAN	0	1
7	TABLE ACCESS BY INDEX ROWID	PRODUCTS	0	1

SALES table has  
10 million rows

PRODUCTS table  
has 100,000 rows

# First Thing To Check Is the Cardinality Estimates

```
SELECT      p.prod_name, SUM(s.quantity_sold)
FROM        sales s, products p
WHERE       s.prod_id = p.prod_id
AND         s.time_id = '03-MAY-21';
GROUP BY    p.prod_name;
```

1. Always start by looking at the Cardinality Estimate

Id	Operation	Name	Starts	Rows
0	SELECT STATEMENT		1	1
1	HASH GROUP BY		1	1
2	NESTED LOOPS		1	1
3	NESTED LOOPS		1	1
4	PARTITION RANGE SINGLE		1	1
*	5	TABLE ACCESS STORAGE FULL	0	1
*	6	INDEX UNIQUE SCAN	0	1
7	TABLE ACCESS BY INDEX ROWID	PRODUCTS	0	1

# First Thing To Check Is the Cardinality Estimates

```
SELECT      p.prod_name, SUM(s.quantity_sold)
FROM        sales s, products p
WHERE       s.prod_id = p.prod_id
AND         s.time_id = '03-MAY-21';
GROUP BY    p.prod_name;
```

1. Always start by looking at the Cardinality Estimate
2. The zero's in the Starts column also look suspicious

Id	Operation	Name	Starts	Rows
0	SELECT STATEMENT		1	1
1	HASH GROUP BY		1	1
2	NESTED LOOPS		1	1
3	NESTED LOOPS		1	1
4	PARTITION RANGE SINGLE		1	1
*	TABLE ACCESS STORAGE FULL	SALES	0	1
*	INDEX UNIQUE SCAN	PRODUCTS_PK	0	1
7	TABLE ACCESS BY INDEX ROWID	PRODUCTS	0	1

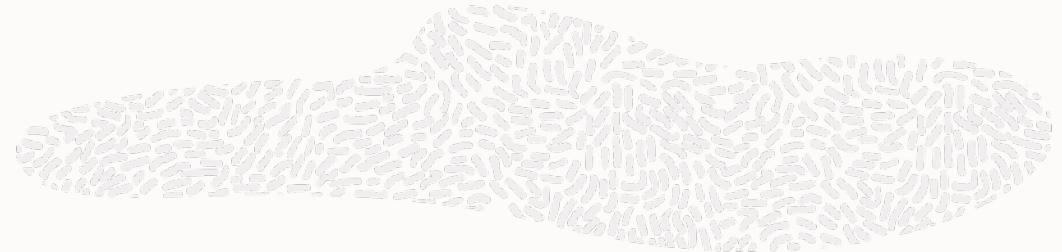
# We are going to need more information



Let's see what additional  
information we can find in  
the database



# Check Table Metadata

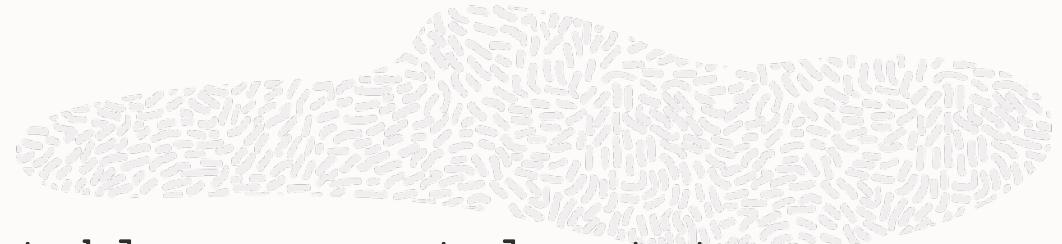


```
SELECT table_name, partitioned  
FROM user_tables  
WHERE table_name in  
( 'SALES' , 'PRODUCTS' );
```

TABLE_NAME	PARTITIONED
------------	-------------

PRODUCTS	NO
SALES	YES

# Check Table Metadata and Statistics



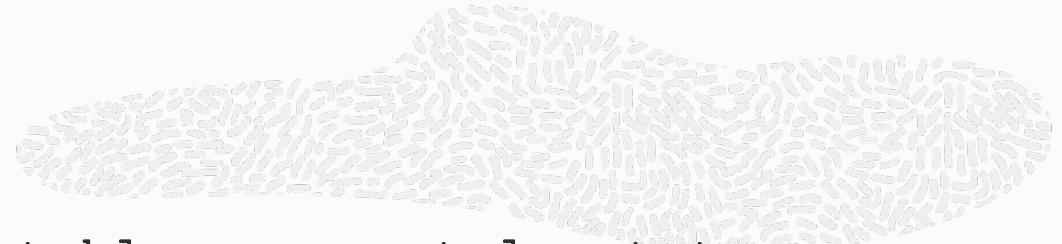
```
SELECT table_name, partitioned  
FROM user_tables  
WHERE table_name in  
( 'SALES' , 'PRODUCTS' );
```

TABLE_NAME	PARTITIONED
PRODUCTS	NO
SALES	YES

```
SELECT table_name, stale_stats  
FROM user_tab_statistics  
WHERE table_name in ( 'SALES' , 'PRODUCTS' );
```

TABLE_NAME	STALE_STATS
PRODUCTS	NO
SALES	NO
SALES	

# Check Table Metadata and Statistics



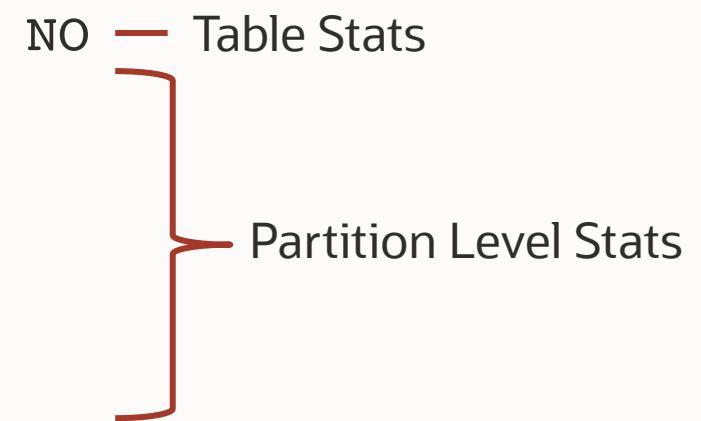
```
SELECT table_name, partitioned  
FROM user_tables  
WHERE table_name in  
( 'SALES' , 'PRODUCTS' );
```

TABLE_NAME	PARTITIONED
PRODUCTS	NO
SALES	YES

```
SELECT table_name, stale_stats  
FROM user_tab_statistics  
WHERE table_name in ( 'SALES' , 'PRODUCTS' );
```

TABLE_NAME	STALE_STATS
PRODUCTS	NO
SALES	NO

SALES



# Next Check the WHERE Clause Predicate Versus The Data



## Where Clause Predicate

```
SELECT      p.prod_name,  
            SUM(s.quantity_sold)  
FROM        sales s, products p  
WHERE       s.prod_id = p.prod_id  
AND         s.time_id = '03-MAY-21'  
GROUP BY    p.prod_name;
```

## The Data

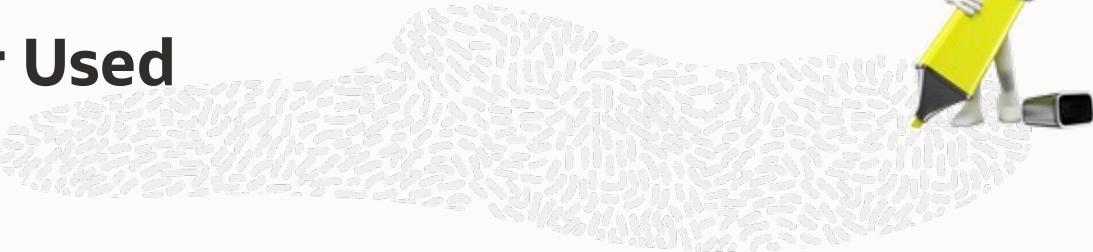
```
SELECT min(time_id),  
       max(time_id)  
FROM   sales;
```

MIN(DATE\_ID)      MAX(DATE\_ID)

01-JAN-04

03-MAY-21

# Compare Data To The Statistics Optimizer Used



```
SELECT max(time_id)  
FROM sales;
```

**MAX(DATE\_ID)**

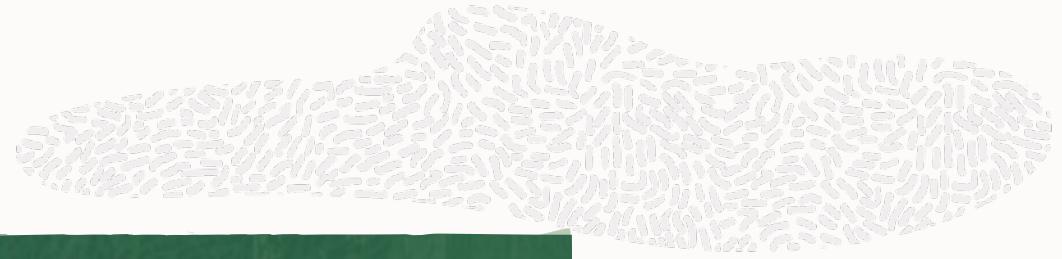
03-MAY-21

```
DECLARE  
    rv RAW(32) ;  
    dt DATE;  
BEGIN  
    SELECT high_value INTO rv  
    FROM user_tab_col_statistics  
    WHERE table_name='SALES'  
    AND column_name='TIME_ID';  
    dbms_stats.convert_raw_value(rv, dt);  
    dbms_output.put_line( TO_CHAR(dt,'dd-MON-yy'));  
END;  
/
```

31-MAR-2021

PL/SQL procedure successfully completed.

# What Do We Know So Far



**01**

Statistics for TIME\_ID column show  
max value as '31-MAR-2021'

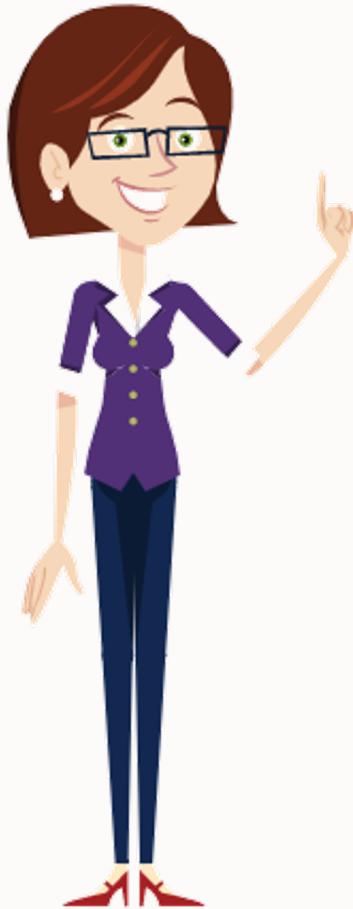
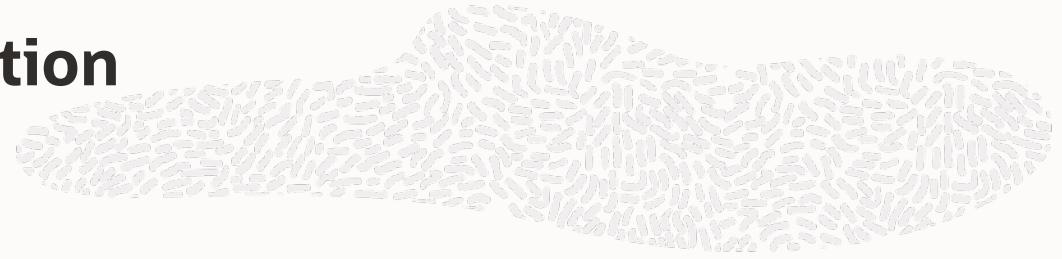
**02**

Our query is looking for sales  
on '03-MAY-2021'

**03**

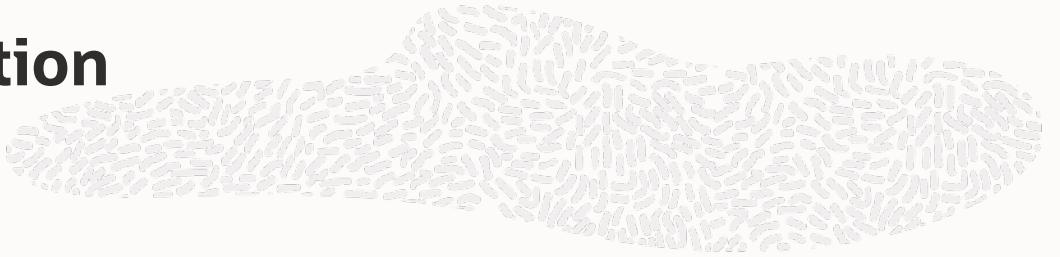
Optimizer checks if predicate falls  
between the min, max value of column

# What Will The Optimizer Do In This Situation



If the Optimizer doesn't find the value in the WHERE clause predicate between the MIN and MAX value of the column, it considers it **OUT\_OF\_RANGE**

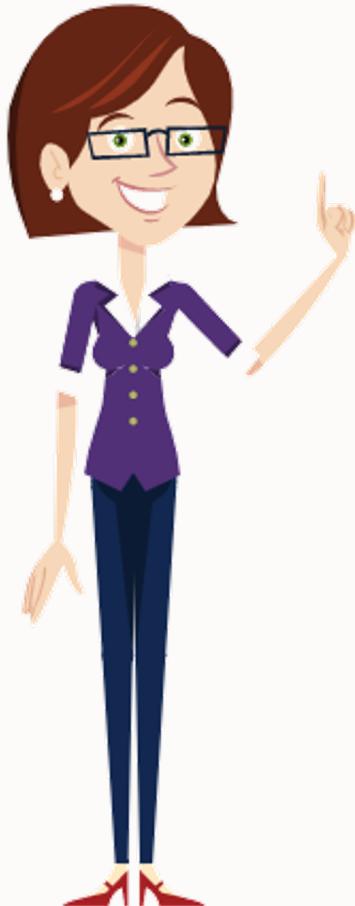
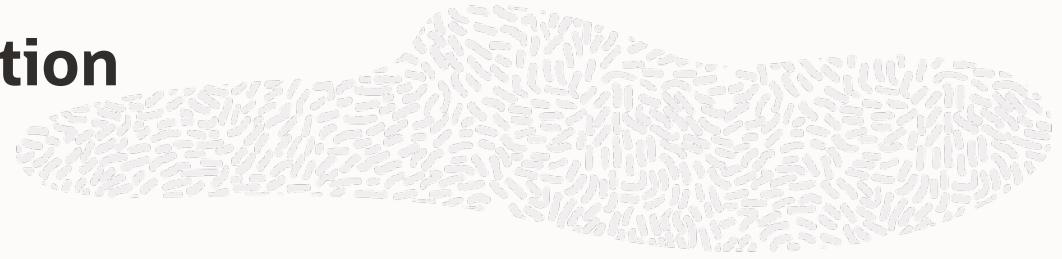
# What Will The Optimizer Do In This Situation



If the Optimizer doesn't find the value in the WHERE clause predicate between the MIN and MAX value of the column, it considers it **OUT\_OF\_RANGE**

Optimizer **prorates cardinality** based on the distance between the predicate value and the maximum value

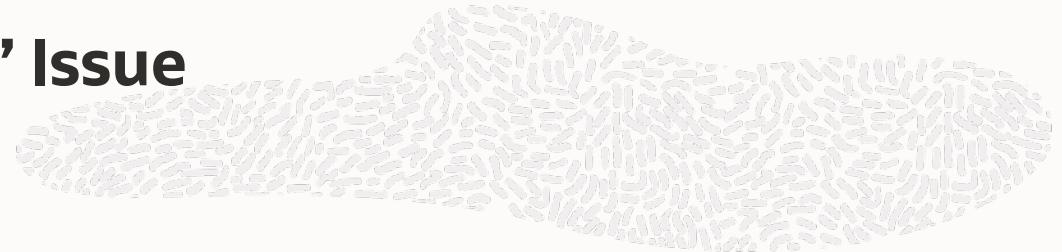
# What Will The Optimizer Do In This Situation



Alternatively we could have looked in the **Optimizer trace** file (10053 trace), where we would have found:

*“Using prorated density: 0.000000 of col #1 as selectivity of out-of-range/non-existent value pred”*

# Stale Statistics Caused an “Out of Range” Issue



```
SELECT      p.prod_name, SUM(s.quantity_sold)
FROM        sales s, products p
WHERE       s.prod_id = p.prod_id
AND         s.time_id = '03-MAY-21';
GROUP BY    p.prod_name;
```

Optimizer assumes there are no rows for time\_id='03-MAY-21' because it is out side of the [MIN,MAX] range in the statistics

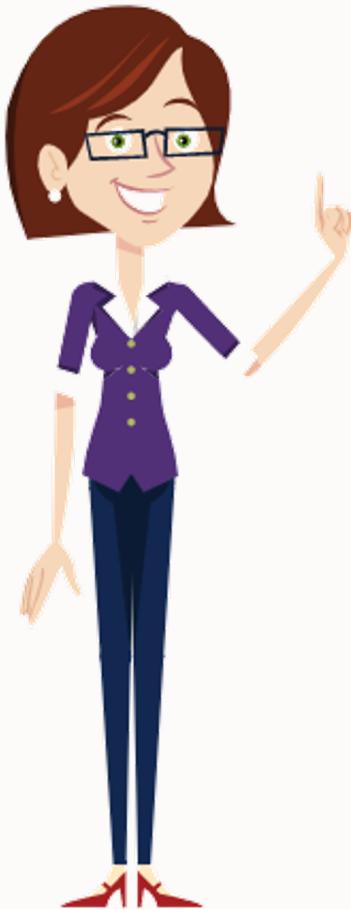
Id	Operation	Name	Starts	Rows
0	SELECT STATEMENT		1	1
1	HASH GROUP BY		1	1
2	NESTED LOOPS		1	1
3	NESTED LOOPS		1	1
4	PARTITION RANGE SINGLE		1	1
*	5	TABLE ACCESS STORAGE FULL	0	1
*	6	INDEX UNIQUE SCAN	0	1
7	TABLE ACCESS BY INDEX ROWID	PRODUCTS	0	1

Always be wary of a cardinality estimates of 1 !



*But my statistics aren't stale?*

# What Will The Optimizer Do In This Situation



By default stats are only considered stale after 10% of the rows have changed (inserted, updated, deleted)

Our tables SALES has 10 million rows in it, therefore **1 million rows** need to be changed before the statistics are considered stale



## Fix Out-of-Range Cardinality Misestimate

### Option 1 - USE DBMS\_STATS.SET\_TABLE\_PREFS To Change Staleness Threshold

To avoid out-of-range problems lower the staleness threshold for larger tables

```
BEGIN  
    dbms_stats.set_table_prefs('SH', 'SALES', 'STALE_PERCENT', '1');  
    dbms_stats.gather_table_stats('SH', 'SALES');  
END;  
/
```



## Fix Out-of-Range Cardinality Misestimate

### Option 2 - USE DBMS\_STATS.COPY\_TABLE\_STATS()

Copies stats from source partition to destination partition

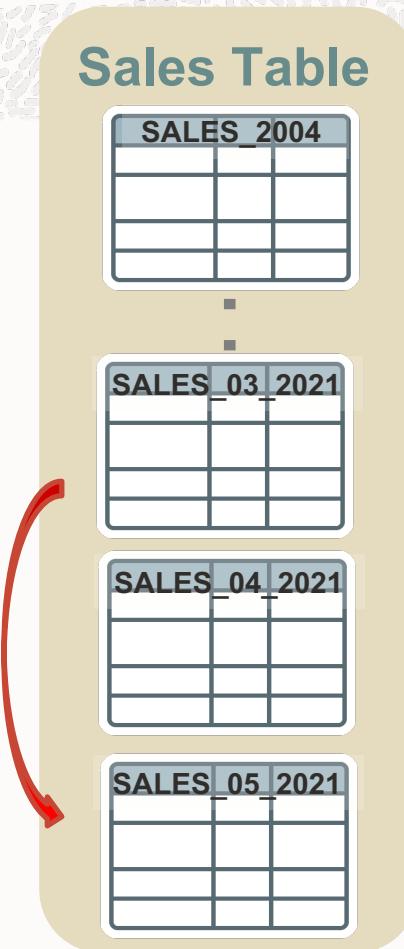
Adjusts min & max values for partition column at both partition & global level

Copies statistics of the dependent objects

- Columns, local indexes etc.
- Does not update global indexes

BEGIN

```
  dbms_stats.copy_table_stats( 'SH','SALES','SALES_03_2021','SALES_05_2021');  
END;  
/
```



# Fixing Out-of-Range Error Means Hash Join Plan Chosen Automatically

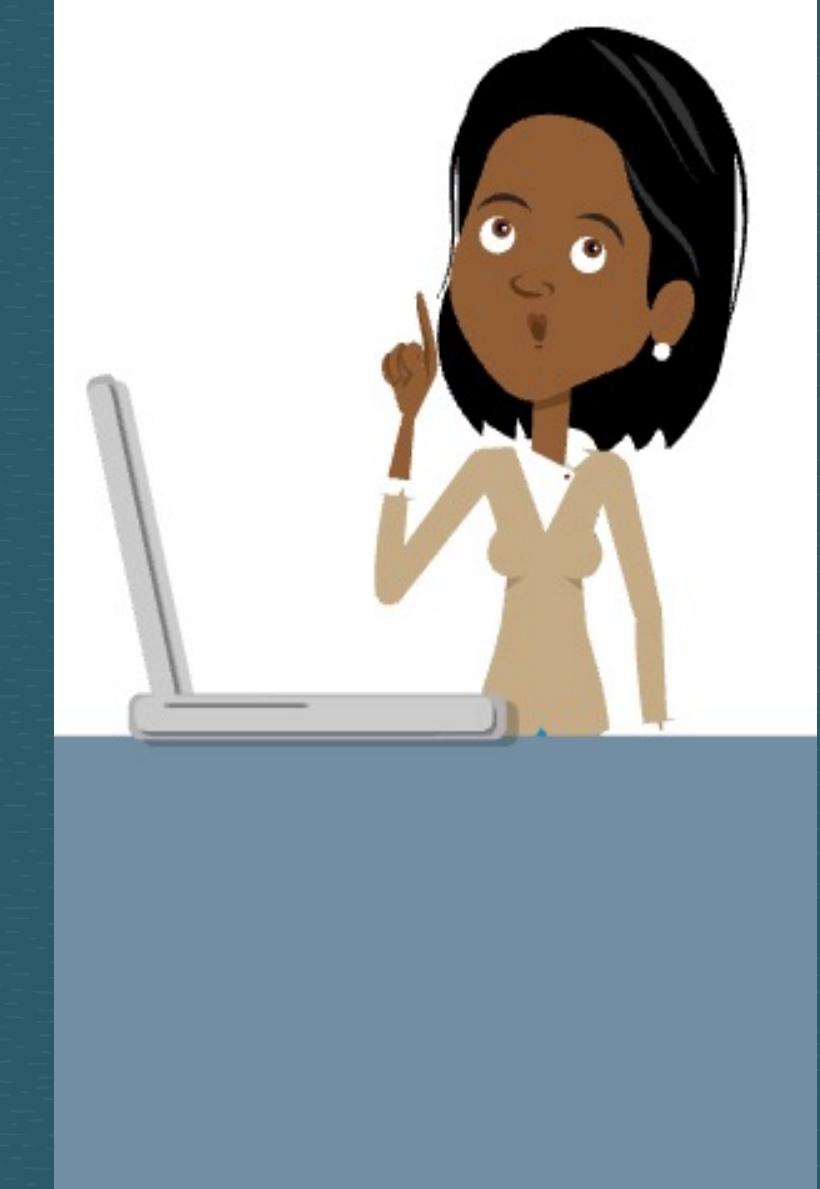
```
SELECT      p.prod_name, SUM(s.quantity_sold)
FROM        sales s, products p
WHERE       s.prod_id = p.prod_id
AND         s.time_id = '03-MAY-21'
GROUP BY    p.prod_name;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Pstart	Pstop
0	SELECT STATEMENT		1	32	81 (20)		
1	HASH GROUP BY		1	32	81 (20)		
* 2	<b>HASH JOIN</b>		3671	114K	80 (19)		
3	JOIN FILTER CREATE	:BF0000	3671	114K	80 (19)		
4	PARTITION RANGE SINGLE		3596	61132	17 (12)	KEY	KEY
* 5	TABLE ACCESS STORAGE FULL	SALES	3596	61132	17 (12)	KEY	KEY
6	JOIN FILTER USE	:BF0000	1398K	20M	58 (14)		
* 7	TABLE ACCESS STORAGE FULL	PRODUCTS	1398K	20M	58 (14)		

## Problem Statement 2

---

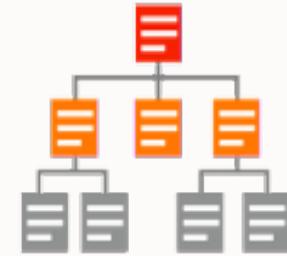
Why Didn't The Optimizer Pick  
The Index I Expected?



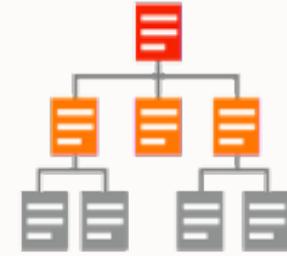
# What Index Should The Optimizer Pick?

## The Query

```
SELECT comment  
FROM Sales  
WHERE prod_id = 141  
AND cust_id < 8938;
```



PROD\_CUST\_IND  
(prod\_id, cust\_id)



PROD\_CUST\_COM\_IND  
(prod\_id, comment, cust\_id)

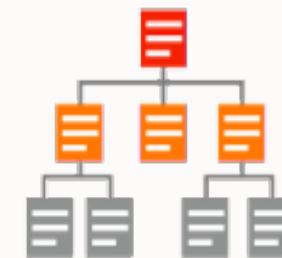


SALES table has  
2 million rows

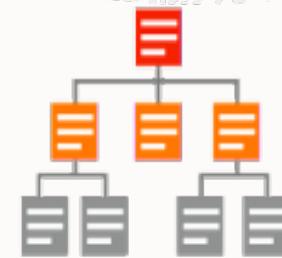
# Why Didn't The Optimizer Use The Index With All The Columns Needed?

## The Query

```
SELECT comment  
FROM Sales  
WHERE prod_id = 141  
AND cust_id < 8938;
```



PROD\_CUST\_IND  
(prod\_id, cust\_id)



PROD\_CUST\_COM\_IND  
(prod\_id, comment, cust\_id)

Id   Operation	Name	Rows	COST ( %CPU )
0   SELECT STATEMENT		2579	2580 (1)
1   TABLE ACCESS BY INDEX ROWID	SALES	2579	2580 (1)
2   INDEX RANGE SCAN	PROD_CUST_IND	2579	9 (1)

# Let's Confirm Our Preferred Plan is Possible

## The Query

```
SELECT /*+ INDEX(PROD_CUST_COM_IND) */  
      comment  
  FROM Sales  
 WHERE prod_id = 141  
   AND cust_id < 8938;
```

The Plan is possible but why is the cost so high?

Id	Operation	Name	Rows	COST (%CPU)	
0	SELECT STATEMENT		2579	9728	(1)
1	INDEX RANGE SCAN	PROD_CUST_COM_IND	2579	9728	(1)

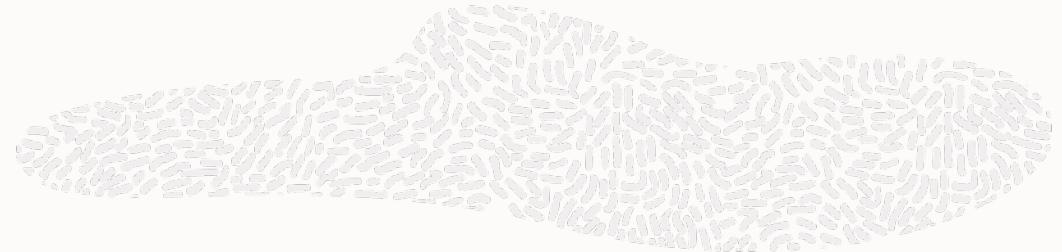
# We are going to need more information



Let's see what additional  
information we can find in  
the database



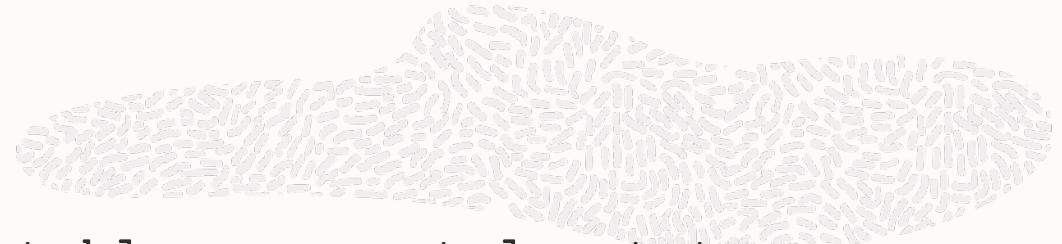
# Check Table Metadata



```
SELECT table_name, partitioned  
FROM user_tables  
WHERE table_name = 'SALES';
```

TABLE_NAME	PARTITIONED
SALES	NO

## Check Table Metadata and Statistics



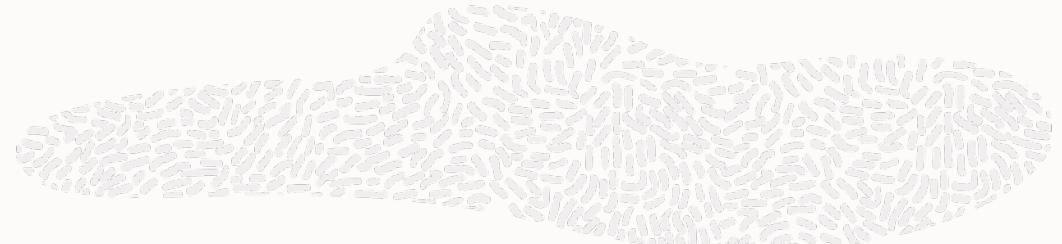
```
SELECT table_name, partitioned  
FROM user_tables  
WHERE table_name = 'SALES';
```

TABLE_NAME	PARTITIONED
SALES	NO

```
SELECT table_name, stale_stats  
FROM user_tab_statistics  
WHERE table_name = 'SALES';
```

TABLE_NAME	STALE_STATS
SALES	NO

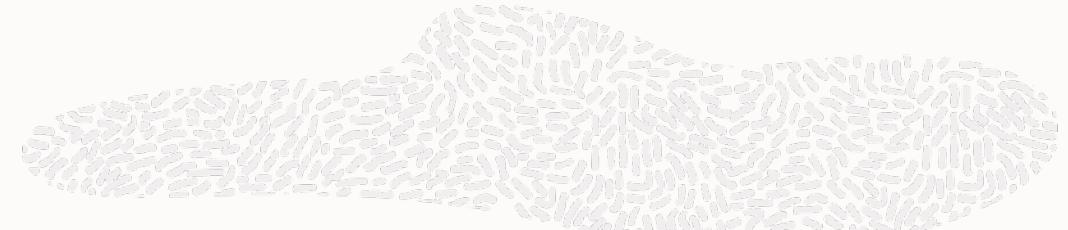
## Check Index Metadata and Statistics



```
SELECT index_name, stale_stats  
FROM   user_ind_statistics  
WHERE  table_name = 'SALES';
```

<b>INDEX_NAME</b>	<b>STALE_STATS</b>
PROD_CUST_COM_IND	NO
PROD_CUST	NO

## Check Index Metadata and Statistics



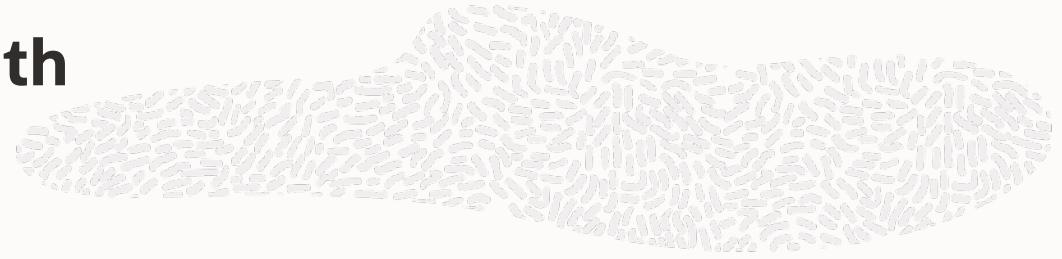
```
SELECT index_name, stale_stats  
FROM user_ind_statistics  
WHERE table_name = 'SALES';
```

```
SELECT index_name, leaf_blocks, blevel  
FROM user_indexes  
WHERE table_name = 'SALES';
```

INDEX_NAME	STALE_STATS	INDEX_NAME	LEAF_BLOCKS	BLEVEL
PROD_CUST_COM_IND	NO	PROD_CUST_COM_IND	699467	9
PROD_CUST	NO	PROD_CUST	5468	2

Why Does PROD\_CUST\_COM\_IND have so many leaf blocks?

# Understand The Data You Are Dealing With



DESC sales

Name	Null?	Type
ORDER_ID	NOT NULL	NUMBER
CUST_ID	NOT NULL	NUMBER
PRODUCT_ID	NOT NULL	NUMBER
SUPPLIER_ID	NOT NULL	NUMBER
DATE_ID		DATE
AMOUNT_SOLD		NUMBER
PRICE		NUMBER
REVENUE		NUMBER
SHIPMODE		CHAR(10)
COMMENTS		VARCHAR2(2000)

# Need To Look At How The Indexes Are Being Used



```
SELECT /*+ INDEX(PROD_CUST_COM_IND) */ comment  
FROM   Sales  
WHERE prod_id = 141  
AND   cust_id < 8938;
```

Id	Operation	Name	Rows	COST (%CPU)
0	SELECT STATEMENT		2579	9728 (1)
1	INDEX RANGE SCAN	PROD_CUST_COM_IND	2579	9728 (1)

Predicate Information (identified by operation id):

```
1 - access("PROD_ID"=141 AND "CUST_ID"<8938)  
      filter("CUST_ID"<8938)
```

Remember the order of the column  
in the PROD\_CUST\_COM\_IND index  
(prod\_id, comment, cust\_id)



# Access Predicates Versus Filter Predicates

```
SELECT /*+ gather_plan_statistics */ count(*) FROM sales2 WHERE prod_id=to_number('139')
```

Plan hash value: 1631620387

Id	Operation	Name	Starts	E-Rows	Cost (%CPU)	A-Rows
0	SELECT STATEMENT		1	1	35 (100)	1
1	SORT AGGREGATE		1	1		1
* 2	INDEX RANGE SCAN	MY_PROD_IND	1	12762	35 (0)	11574

Predicate Information (identified by operation id):

2 - access("PROD\_ID"=139)

## Access predicate

- Where clause predicate used for data retrieval
  - The start and stop keys for an index
  - If rowids are passed to a table scan

# Access Predicates Versus Filter Predicates



```
SQL> SELECT username  
  2   FROM my_users  
  3 WHERE username LIKE 'MAR%';
```

USERNAME

MARIA

Plan hash value: 2982854235

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT				2 (100)	
1	TABLE ACCESS FULL	MY_USERS	1	66	2 (0)	00:00:01

Predicate Information (identified by operation id):

```
1 - filter("USERNAME" LIKE 'MAR%')
```

## Filter predicate

- Where clause predicate that is not used for data retrieval but to eliminate uninteresting row once the data is found
- Requires additional CPU resources to apply filters

# Need To Look At How The Indexes Are Being Used



```
SELECT comment  
FROM   Sales  
WHERE prod_id = 141  
AND   cust_id < 8938;
```

Id	Operation	Name	Rows	COST (%CPU)
0	SELECT STATEMENT		2579	2580 (1)
1	TABLE ACCESS BY INDEX ROWID	MY_SALES	2579	2580 (1)
2	INDEX RANGE SCAN	PROD_CUST_IND	2579	9 (1)

Predicate Information (identified by operation id):

```
1 - access("PRODUCT_ID"=517538 AND "CUST_ID"><8938)
```

Both columns are used  
as Access Predicates

# Costing of Each Index Access is Different



Id	Operation	Name	Rows	COST (%CPU)	
0	SELECT STATEMENT		2579	2580	(1)
1	TABLE ACCESS BY INDEX ROWID	MY_SALES	2579	2580	(1)
2	INDEX RANGE SCAN	PROD_CUST_IND	2579	9	(1)

Predicate Information (identified by operation id):

1 - access("PRODUCT\_ID"=517538 AND "CUST\_ID"<8938)

Both columns are used as **Access Predicates**

Cost calculation for this index:

blevel + (cardinality X leaf blocks)  
#rows

Id	Operation	Name	Rows	COST (%CPU)	
0	SELECT STATEMENT		2579	9728	(1)
1	INDEX RANGE SCAN	PROD_CUST_COM_IND	2579	9728	(1)

Predicate Information (identified by operation id):

1 - access("PRODUCT\_ID"=517538 AND "CUST\_ID"<8938)  
- filter("CUST\_ID"<8938)

Only the PROD\_ID is used as **Access Predicates**

Cost calculation for this index:

blevel + (1 X Leaf Block)  
NDV of col1

## Costing of Each Index Access is Different

Cost calculation for PROD\_CUST\_IND index:

$$2 + (2579 \div 2098400)(5468) = \underline{9}$$

Cost calculation for PROD\_CUST\_COM\_IND index:

$$9 + (0.013 \times 699467) = \underline{9724}$$

# Common Misconceptions On What Index Will Be Chosen

A common misconception is the Optimizer will always pick the index with all the necessary columns in it



# Common Misconceptions On What Index Will Be Chosen



A common misconception is the Optimizer will always pick the index with all the necessary columns in it

Fact is the Optimizer picks the index based on the cost of the access



# Common Misconceptions On What Index Will Be Chosen



A common misconception is the Optimizer will always pick the index with all the necessary columns in it

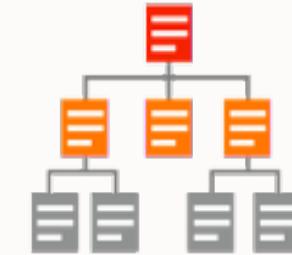
Fact is the Optimizer picks the index based on the cost of the access

How the WHERE clause predicates are used and the order of the columns in the index have a massive impact on the Optimizers choice



## Solution: Recreate PROD\_CUST\_COM\_IND With Different Column Order

```
CREATE INDEX Prod_cust_com_ind2  
ON Sales(prod_id, cust_id, comment);
```

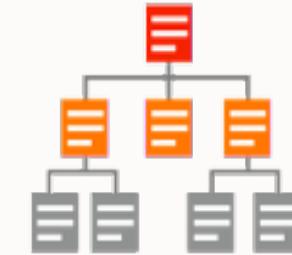


PROD\_CUST\_COM\_IND2  
(prod\_id, cust\_id, comment)

# Solution: Recreate PROD\_CUST\_COM\_IND With Different Column Order

## The Query

```
SELECT comment  
FROM Sales  
WHERE prod_id = 141  
AND cust_id < 8938;
```



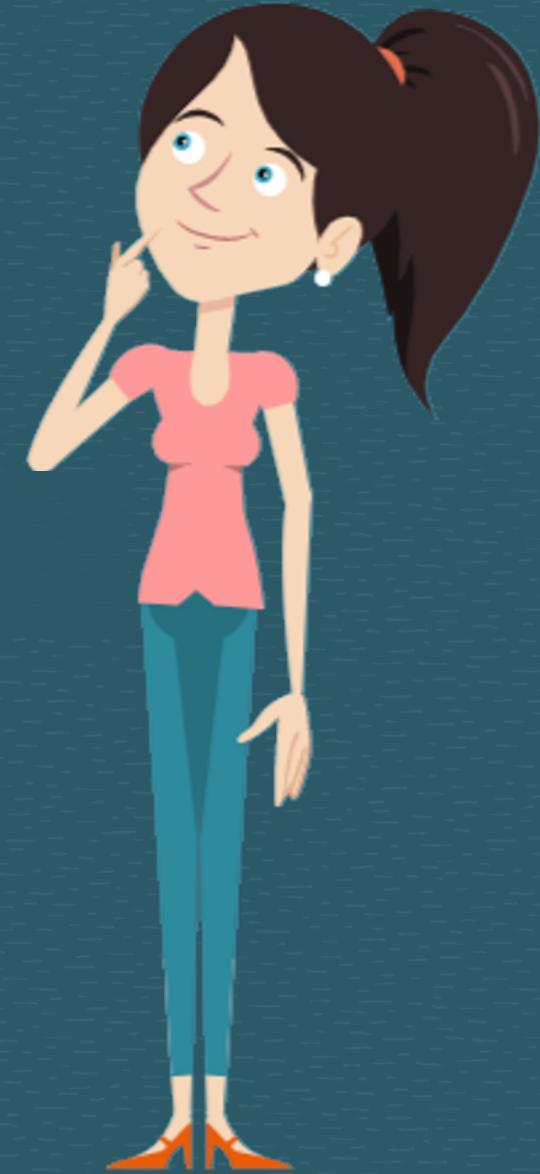
PROD\_CUST\_COM\_IND2  
(prod\_id, cust\_id, comment)

Id	Operation	Name	Rows	COST (%CPU)
0	SELECT STATEMENT		2579	1935 (100)
1	INDEX RANGE SCAN	PROD_CUST_COM_IND2	2579	1935 (0)

## Problem Statement 3

---

**Why Didn't The Optimizer Pick  
The Join Method I Expected?**

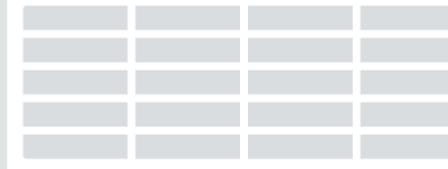


# What Join Method Should The Optimizer Pick?

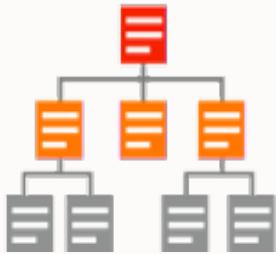
## The Query

```
SELECT COUNT(*)  
FROM   sales s, customers c  
WHERE  s.cust_id = c.cust_id  
AND    substr(c.cust_state_province,1,2)='CA';
```

SALES 

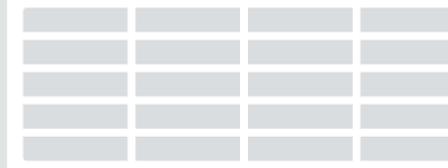


SALES table has  
10 million rows



SALES\_CUST\_BIX  
(cust\_id)

CUSTOMERS



CUSTOMERS table  
has 10,000 rows

# Got a Nested Loops Join Instead of Hash Join

## The Query

```
SELECT COUNT(*)  
FROM   sales s, customers c  
WHERE  s.cust_id = c.cust_id  
AND    substr(c.cust_state_province,1,2)='CA';
```

Id	Operation	Name	Rows	Cost (%CPU)
0	SELECT STATEMENT			26 (100)
1	SORT AGGREGATE		1	
2	NESTED LOOPS		41	26 (0)
3	TABLE ACCESS FULL	CUSTOMER	100	2 (0)
* 4	BITMAP CONVERSION COUNT		4	26 (0)
* 5	BITMAP INDEX SINGLE VALUE	SALES_CUST_BIX		

# Let's Confirm Our Preferred Plan is Possible

## The Query

```
SELECT /*+ USE_HASH(s) */ COUNT(*)  
FROM   sales s, customers c  
WHERE  s.cust_id = c.cust_id  
AND    substr(c.cust_state_province,1,2)='CA';
```

The Plan is possible but why is the cost so high?

Id	Operation	Name	Rows	Cost (%CPU)
0	SELECT STATEMENT			1114 (100)
1	SORT AGGREGATE		1	
2	HASH JOIN		14732	1114 (2)
* 4	TABLE ACCESS FULL	CUSTOMER	100	406 (1)
* 5	TABLE ACCESS FULL	SALES	10M	703 (2)

# We are going to need more information



**Let's see what actually happened  
during the statement execution**

# Need to Gather More Information



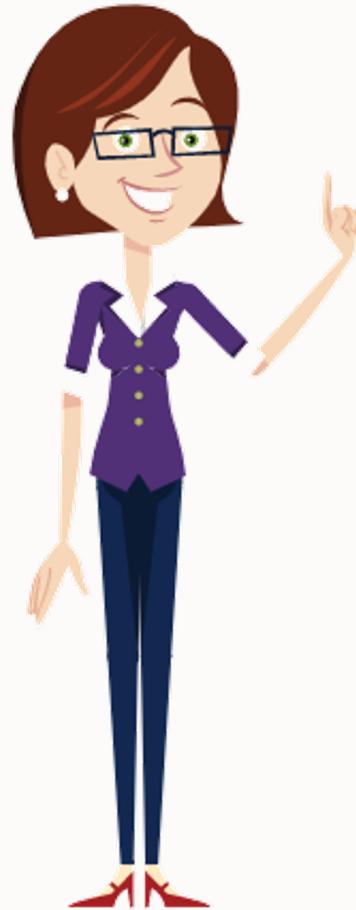
```
SELECT /*+ GATHER_PLAN_STATISTICS */ COUNT(*)  
FROM sales s, customers c  
WHERE s.cust_id = c.cust_id  
AND substr(c.cust_state_province,1,2)='CA';
```

1. Always start by looking at the Cardinality Estimate

```
SELECT * FROM table(dbms_xplan.display_cursor(format=>'ALLSTATS LAST'));
```

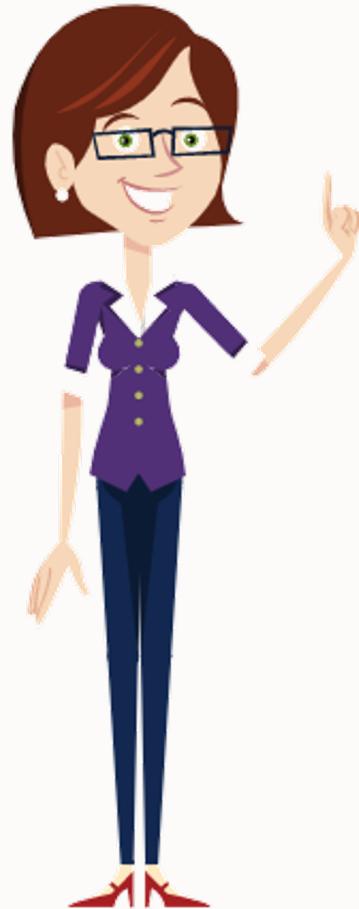
Id	Operation	Name	E-Rows	A-Rows	Buffer
0	SELECT STATEMENT		1	1	15984
1	SORT AGGREGATE		1	1	15984
2	NESTED LOOPS		1471	154K	15984
3	TABLE ACCESS FULL	CUSTOMER	100	6656	15984
* 4	BITMAP CONVERSION COUNT		4	154K	2914
* 5	INDEX RANGE SCAN	SALES_CUST_BIX			13034

# Why is the Cardinality Estimate So Wrong?



Always start by looking at the cardinality estimate for the table on the left-hand side of the join

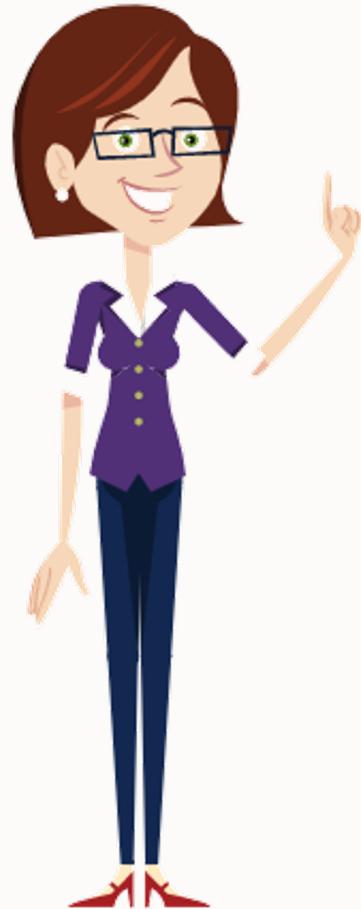
# Why is the Cardinality Estimate So Wrong?



Always start by looking at the cardinality estimate for the table on the left-hand side of the join

Customers is the table on the left hand side of our join  
It has one WHERE clause predicates is  
`substr(c.cust_state_province, 1, 2) = 'CA'`

# Why is the Cardinality Estimate So Wrong?



Always start by looking at the cardinality estimate for the table on the left-hand side of the join

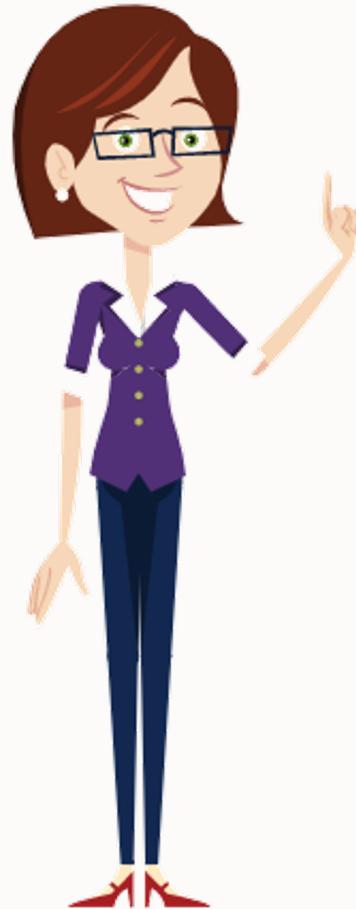
Customers is the table on the left hand side of our join

It has one WHERE clause predicates is

`substr(c.cust_state_province, 1, 2) = 'CA'`

What is the first thing you noticed about this predicate?

# Why is the Cardinality Estimate So Wrong?



Always start by looking at the cardinality estimate for the table on the left-hand side of the join

Customers is the table on the left hand side of our join  
It has one WHERE clause predicate:

`substr(c.cust_state_province, 1, 2) = 'CA'`

There is a function being applied to the column in the WHERE clause predicate

# Why is the Cardinality Estimate So Wrong?



Always start by looking at the cardinality estimate for the table on the left-hand side of the join

Customers is the table on the left hand side of our join

It has one WHERE clause predicate:

`substr(c.cust_state_province, 1, 2) = 'CA'`

There is a function being applied to the column in the WHERE clause predicate

The Optimizer has no idea how the function affects the values in the column

# Function Wrapped Column Blinds the Optimizer

Optimizer Doesn't Know How Function Affects Values In The Column

## The Query

```
SELECT COUNT(*)  
FROM   customers c  
WHERE  substr(c.cust_state_province,1,2)='CA';
```

Remember  
CUSTOMERS table  
has 10,000 rows

Id	Operation	Name	Rows	Cost (%CPU)
0	SELECT STATEMENT			2 (100)
1	SORT AGGREGATE		1	
* 4	TABLE ACCESS FULL	CUSTOMER	100	2 (0)

Optimizer guesses the cardinality to be 1% of rows

# Solution: Tell Optimizer How Function Affects Column Values

## Option 1 Create Extended Statistics

```
SELECT
```

```
dbms_stats.create_extended_stats(null, 'CUSTOMERS',  
                                '(SUBSTR(CUST_STATE_PROVINCE,1,2))')  
FROM dual;
```

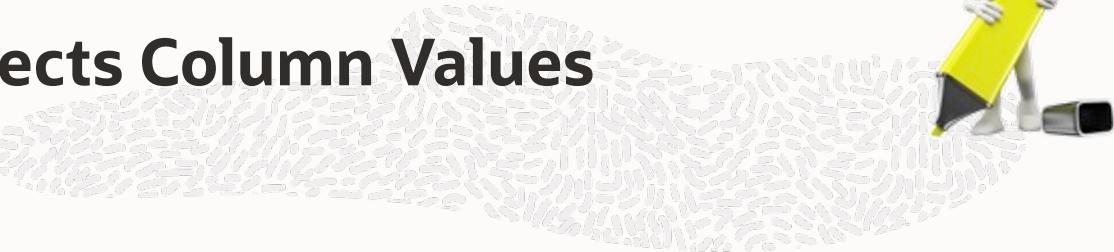
```
DBMS_STATS.CREATE_EXTENDED_STATS(NULL, 'CUSTOMERS', '(SUBSTR(CUST_STATE_PROVINCE,1,2))')
```

---

```
SYS_STUAJYEDA$07W8CRW1A18K4Q_G
```

# **Solution: Tell Optimizer How Function Affects Column Values**

## **How To Identify Extended Statistics**



```
SELECT column_name, num_distinct, histogram  
FROM   user_tab_col_statistics  
WHERE  table_name='CUSTOMERS';
```

COLUMN_NAME	NUM_DISTINCT	HISTOGRAM
SYS_STUAJYEDA\$07W8CRW1A18K4Q_G	84	FREQUENCY
C_CUST_ID	335	HYBRID
CUST_YEAR_OF_BIRTH	69	FREQUENCY
CUST_VALID	2	FREQUENCY
CUST_TOTAL_ID	1	FREQUENCY
CUST_TOTAL	1	FREQUENCY
CUST_STREET_ADDRESS	340	HYBRID
:		

# Solution: Tell Optimizer How Function Affects Column Values

## The Query

```
SELECT COUNT(*)  
FROM   sales s, customers c  
WHERE  s.cust_id = c.cust_id  
AND    substr(c.cust_state_province,1,2)='CA';
```

Id	Operation	Name	Rows	Cost (%CPU)
0	SELECT STATEMENT			1114 (100)
1	SORT AGGREGATE		1	
2	HASH JOIN		14732	1114 (2)
* 4	TABLE ACCESS FULL	CUSTOMER	6656	406 (1)
* 5	TABLE ACCESS FULL	SALES	10M	703 (2)

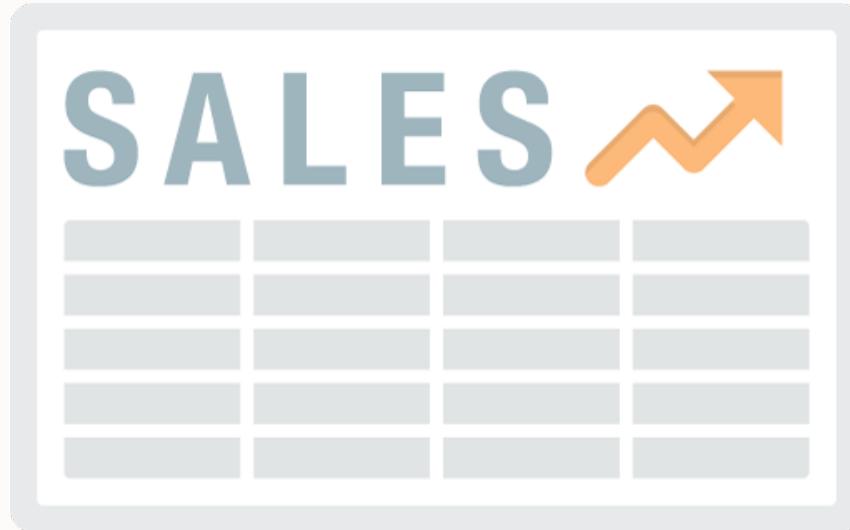
## Problem Statement 4

---

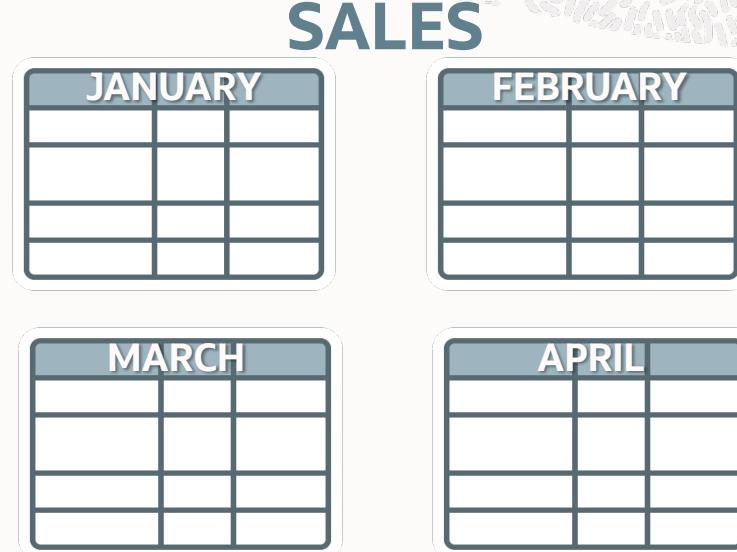
Why Didn't The Optimizer Use  
Partition Pruning?



# Partitioning Provides Flexibility & Efficiency at Scale



Large Table  
Difficult to  
Manage



Partitions  
Divide and Conquer  
Easier to Manage  
Improve Performance

Transparent to applications

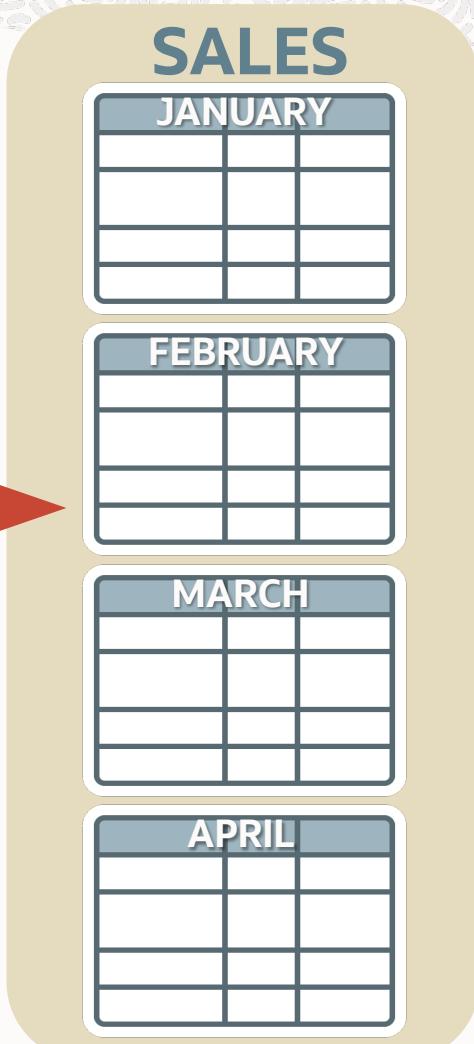
# Partitioning Provides Performance Acceleration



## The Query

```
SELECT SUM(s.sales_amount)
FROM sales s
WHERE s.sales_date = 'FEBRUARY';
```

Only the relevant partition is accessed

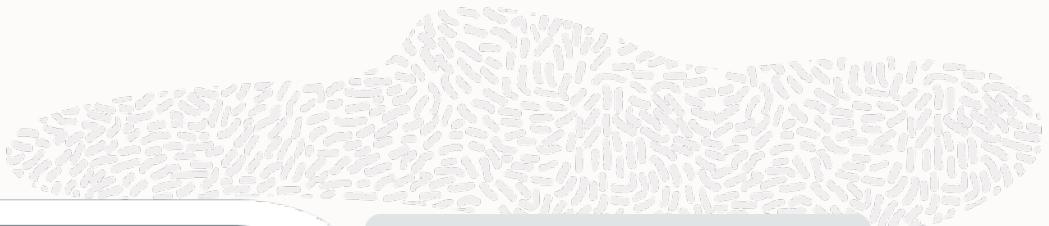


Id	Operation	Name	Rows	Cost (%CPU)	Pstart	Pstop
0	SELECT STATEMENT			16 (100)		
1	SORT AGGREGATE		1			
2	PARTITION RANGE SINGLE		4018	16 (7)	12	12
*	TABLE ACCESS STORAGE FULL	SALES	4018	16 (7)	12	12

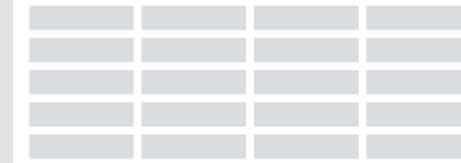
# Expected Partition Pruning

## The Query

```
SELECT      SUM(s.amount_sold)
FROM        sales s,
WHERE       TO_CHAR(s.time_id,'YYYYMMDD')
BETWEEN    '20200101' AND '20201231';
```



SALES 



SALES table has  
10 million rows

SALES is partitioned on the TIME\_ID column on a quarterly basis and has 28 partitions

# Expected Partition Pruning

## The Query

```
SELECT      SUM(s.amount_sold)
FROM        sales s,
WHERE       TO_CHAR(s.time_id,'YYYYMMDD')
BETWEEN    '20200101' AND '20201231';
```

SALES 

SALES table has  
10 million rows

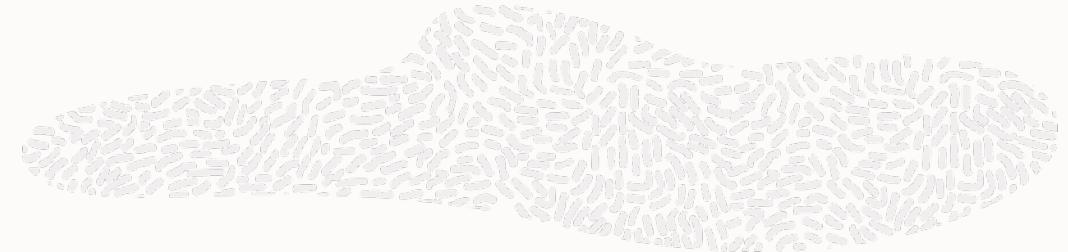
Id	Operation	Name	ROWS	Cost (%CPU)	Pstart	Pstop
0	SELECT STATEMENT			552 (100)		2
1	SORT AGGREGATE		13			2
2	PARTITION RANGE ALL		10M	552 (7)	1	28
* 3	TABLE ACCESS FULL	SALES	10M	552 (7)	1	28

# We are going to need more information



**Let's look at how the WHERE clause predicates are being used in the plan**

# Expected Partition Pruning



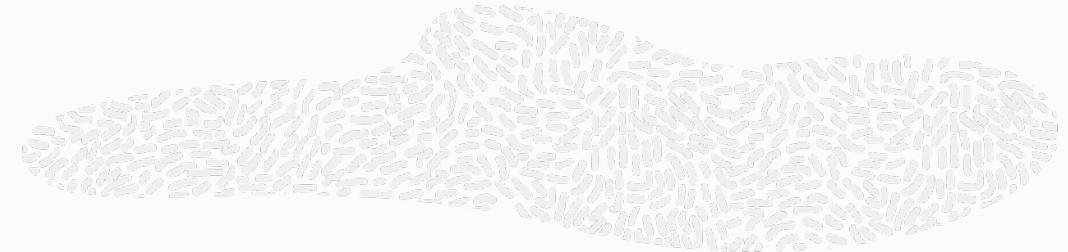
Id	Operation	Name	ROWS	Cost (%CPU)	Pstart	Pstop
0	SELECT STATEMENT			552 (100)		2
1	SORT AGGREGATE		13			2
2	<b>PARTITION RANGE ALL</b>		10M	552 (7)	1	28
* 3	TABLE ACCESS FULL	SALES	10M	552 (7)	1	28

Predicate Information (identified by operation id):

3 - **filter**((TO\_CHAR(INTERNAL\_FUNCTION("S"."TIME\_ID"),'YYYYMMDD')>='20190101' AND TO\_CHAR(INTERNAL\_FUNCTION("S"."TIME\_ID"),'YYYYMMDD')<='20191231'))

**For Partition Pruning To Occur We Need An Access Predicate**

# Expected Partition Pruning



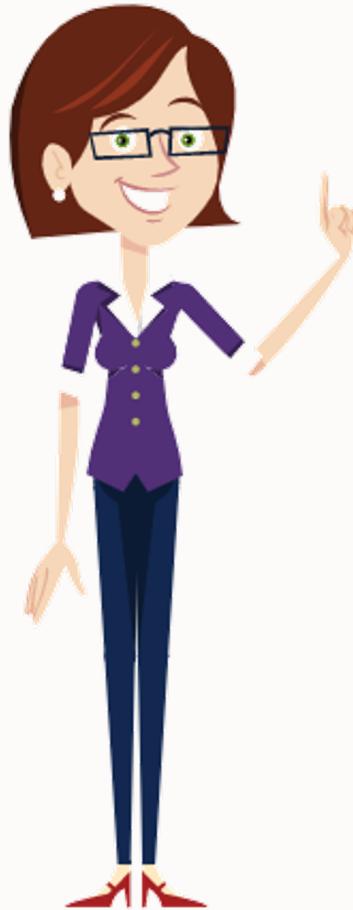
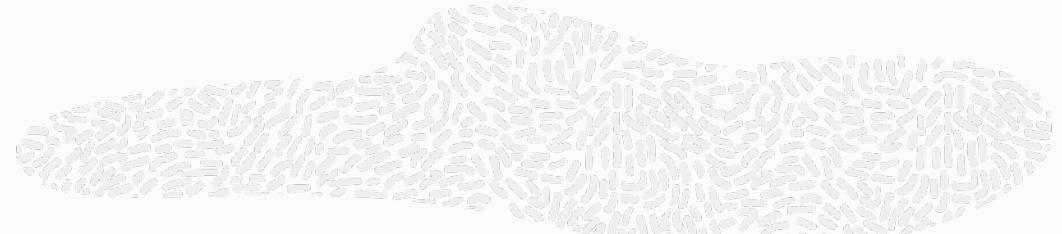
Id	Operation	Name	ROWS	Cost (%CPU)	Pstart	Pstop
0	SELECT STATEMENT			552 (100)		2
1	SORT AGGREGATE		13			2
2	<b>PARTITION RANGE ALL</b>		10M	552 (7)	1	28
* 3	TABLE ACCESS FULL	SALES	10M	552 (7)	1	28

Predicate Information (identified by operation id):

```
3 - filter((TO_CHAR(INTERNAL_FUNCTION("S"."TIME_ID"),'YYYYMMDD')>='20190101' AND  
           TO_CHAR(INTERNAL_FUNCTION("S"."TIME_ID"),'YYYYMMDD')<='20191231'))
```

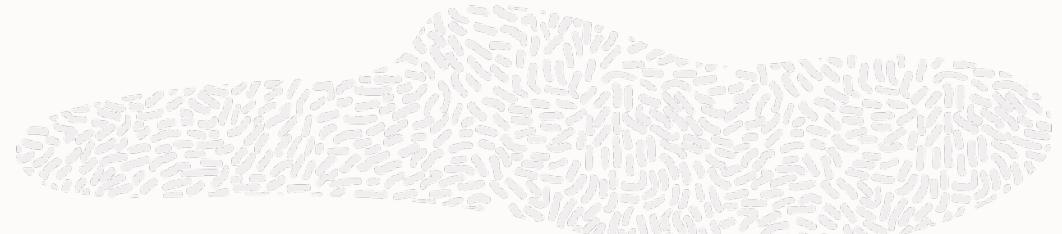
Why Has An Additional INTERNAL\_FUNCTION Been Added To Our Predicate?

# Why We Didn't Get Partition Pruning



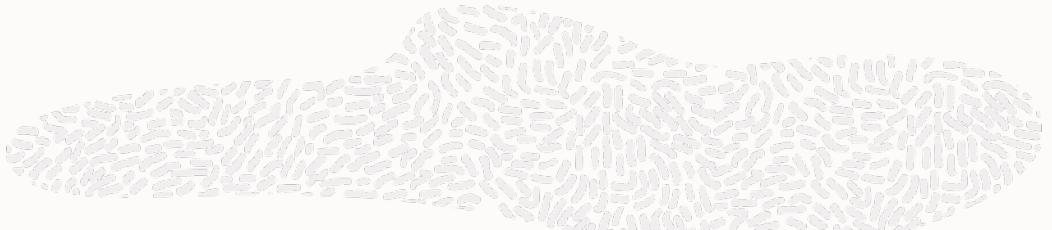
- The presences of an INTERNAL\_FUNCTION typically means a data type conversion has occurred

# Why We Didn't Get Partition Pruning



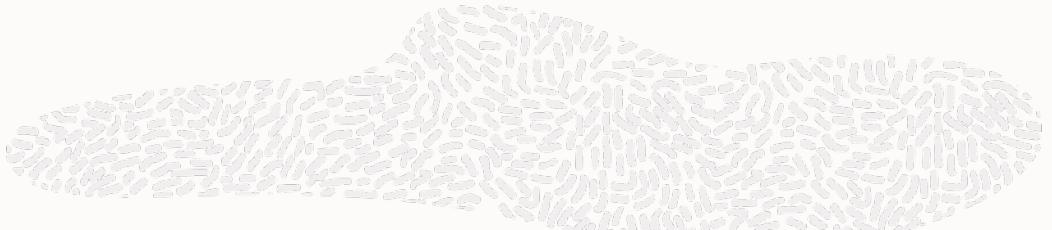
- The presences of an INTERNAL\_FUNCTION typically means a data type conversion has occurred
- A data type conversion is needed when column type & predicate type don't match
- Predicate is `TO_CHAR(s.time_id, 'YYYYMMDD')`

# Why We Didn't Get Partition Pruning



- The presences of an INTERNAL\_FUNCTION typically means a data type conversion has occurred
- A data type conversion is needed when column type & predicate type don't match
- Predicate is `TO_CHAR(s.time_id, 'YYYYMMDD')`
- `TIME_ID` is actually a date column
- Optimizer has no idea how function will affect the values in `TIME_ID` column

# Why We Didn't Get Partition Pruning



- The presences of an INTERNAL\_FUNCTION typically means a data type conversion has occurred
- A data type conversion is needed when column type & predicate type don't match
- Predicate is `TO_CHAR(s.time_id, 'YYYYMMDD')`
- `TIME_ID` is actually a date column
- Optimizer has no idea how function will affect the values in `TIME_ID` column
- Optimizer can't determine which partitions will be accessed now

# Solution: Using Inverse Function On Other Side Of Predicate

## The Query

```
SELECT SUM(s.amount_sold)
FROM sales s,
WHERE s.time_id BETWEEN TO_DATE('20200101', 'YYYYMMDD') AND
                      TO_DATE(' 20201231', 'YYYYMMDD');
```

Id	Operation	Name	ROWS	Cost (%CPU)	Pstart	Pstop
0	SELECT STATEMENT			552 (100)		2
1	SORT AGGREGATE		13			2
2	PARTITION RANGE ALL		10M	552 (7)	22	24
*	TABLE ACCESS FULL	SALES	10M	552 (7)	22	24

Predicate Information (identified by operation id):

3 - Access("S"."TIME\_ID">>='20190101' TO\_DATE('20190101','YYYYMMDD') AND  
TO\_DATE(' 20191231','YYYYMMDD'))

## Important Tip on Using Functions

### Using Inverse Function On Other Side Of Predicate

Keep the following in mind when deciding where to place the function

- Try to place functions on top of constants (literals, binds) rather than on columns

## Important Tip on Using Functions

### Using Inverse Function On Other Side Of Predicate

Keep the following in mind when deciding where to place the function

- Try to place functions on top of constants (literals, binds) rather than on columns
- Avoid using a function on index columns or partition keys as it prevents pruning or use of index

## Important Tip on Using Functions

### Using Inverse Function On Other Side Of Predicate

Keep the following in mind when deciding where to place the function

- Try to place functions on top of constants (literals, binds) rather than on columns
- Avoid using a function on index columns or partition keys as it prevents pruning or use of index
- For function-based index to be considered, use that exact function as specified in index

## Important Tip on Using Functions

### Using Inverse Function On Other Side Of Predicate

Keep the following in mind when deciding where to place the function

- Try to place functions on top of constants (literals, binds) rather than on columns
- Avoid using a function on index columns or partition keys as it prevents pruning or use of index
- For function-based index to be considered, use that exact function as specified in index
- If multiple predicates involve the same columns, write predicates such that they share common expressions For example,

WHERE  $f(a) = b$   
AND       $a = c$

Should be  
rewritten as

WHERE  $a = \text{inv\_f}(b)$   
AND       $a = c$

This will allow transitive predicate  $c = \text{inv\_f}(b)$  to be added by the optimizer

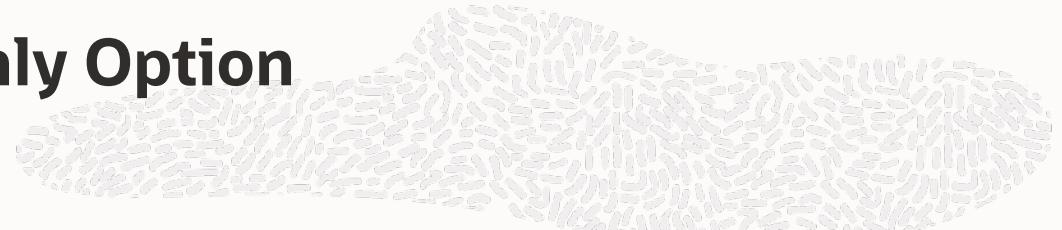
## Question Number 5

---

Why Is The Optimizer Ignoring My Hints?

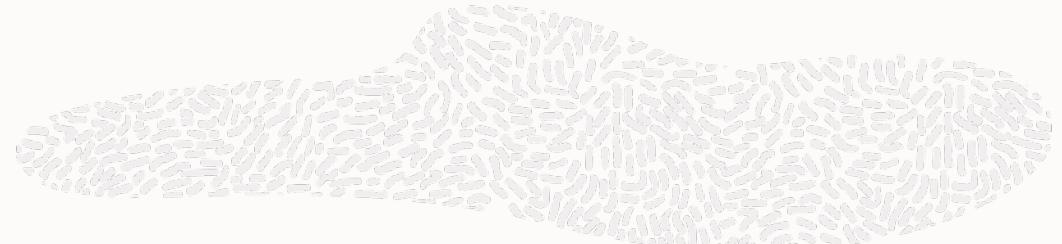


# Some Times an Optimizer Hint is Your Only Option



- Hints allow you to influence the Optimizer when it has to choose between several possibilities
- A hint is a directive that will be followed when applicable
- Can influence everything from the Optimizer mode used to each operation in the execution
- Automatically means the Cost Based Optimizer will be used
- Only exception is the RULE hint but it must be used alone

# Why Are Optimizer Hints Ignored?

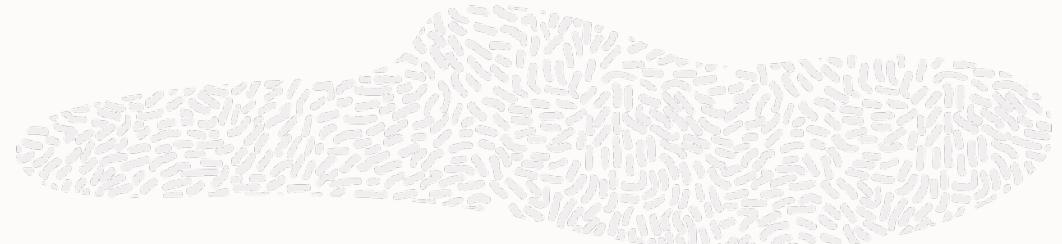


- An Optimizer hint is never ignored
- Hints only evaluated when they apply to a decision that has to be made
- Often times hint that's aren't used because they are irrelevant or not legal

**How Can I tell What's Happening With My Hint?**

# Why Are Optimizer Hints Ignored?

How to find out what happened to you hint



- An Optimizer hint is never ignored
- Hints only evaluated when they apply to a decision that has to be made
- Often times hint that's aren't used because they are irrelevant or not legal

In **10053 trace file** you will find:

*“Dumping Hints*

*=====*

```
Atom_hint=(@=0X124360178 err=0 resol=0 used=1 token=454 org=1 lvl=1 txt=ALL_ROWS)
Atom_hint=(@=0X2af785e0c260 err=0 resol=1 used=1 token=448 org=1 lvl=3 txt=FULL ("E"))
===== -- END SQL Statement Dump =====
```

ERR indicates if there is an error with hint

USED indicates the hint was used during the evaluation of the part of the plan it pertains to But Doesn't mean the final plan will reflect it

# Why Are Optimizer Hints Ignored?

How to find out what happened to your hint

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT				9 (100)	
1	HASH GROUP BY		1	52	9 (23)	00:00:01
* 2	HASH JOIN		1	52	8 (13)	00:00:01
3	PARTITION RANGE SINGLE		2	34	2 (0)	00:00:01
* 4	TABLE ACCESS FULL	SALES	2	34	2 (0)	00:00:01
* 5	TABLE ACCESS FULL	CUSTOMERS	13	455	5 (0)	00:00:01

Predicate Information (identified by operation id):

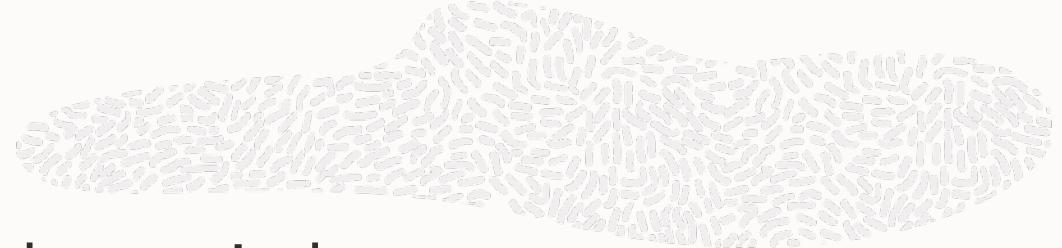
```
2 - access("C"."CUST_ID"="S"."CUST_ID")
4 - filter("S"."TIME_ID"='30-DEC-99')
5 - filter(("C"."CUST_CITY"='Los Angeles' AND "CUST_STATE_PROVINCE"='CA'))
```

Hint Report (identified by operation id / Query Block Name / Object Alias):  
Total hints for statement: 1 (U - Unused (1))

```
4 - SEL$1 / S@SEL$1
U - USE NL(s)
```

New hint info under the plan  
in 19c with  
DBMS\_XPLAN.DISPLAY\_CURSOR

# Why Are Optimizer Hints Ignored?



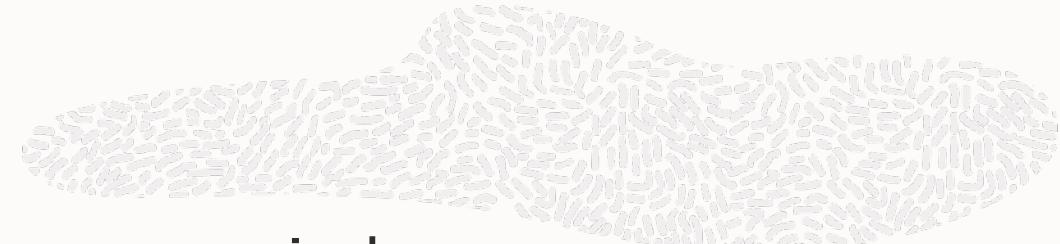
Employees table has a unique index called pk\_emp index

```
SELECT /*+ INDEX(e emp_pk) */ * FROM employees e;
```

Id	Operation	Name	Rows	COST (%CPU)	
0	SELECT STATEMENT		2579	8743	(1)
1	TABLE ACCESS FULL	EMPLOYEES	2579	8743	(1)

# Why Are Optimizer Hints Ignored?

## Syntax and Spelling Mistakes



EMPLOYEES table has a unique index called **PK\_EMP** index

```
SELECT /*+ INDEX(e emp_pk) */ * FROM employees e;
```

Id	Operation	Name	Rows	COST (%CPU)	
0	SELECT STATEMENT		2579	8743	(1)
1	TABLE ACCESS FULL	EMPLOYEES	2579	8743	(1)

Hint Report (identified by operation id / Query Block Name / Object Alias):

Total hints for statement: 1 (U - Unused (1))

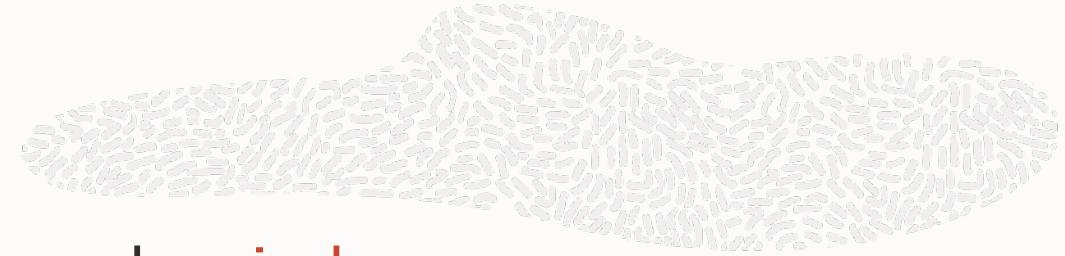
---

```
1 - SEL$1 /S@SEL$1
```

U - index (e emp\_pk) / index specified in the hint doesn't exist

19<sup>c</sup>

# Why Are Optimizer Hints Ignored?



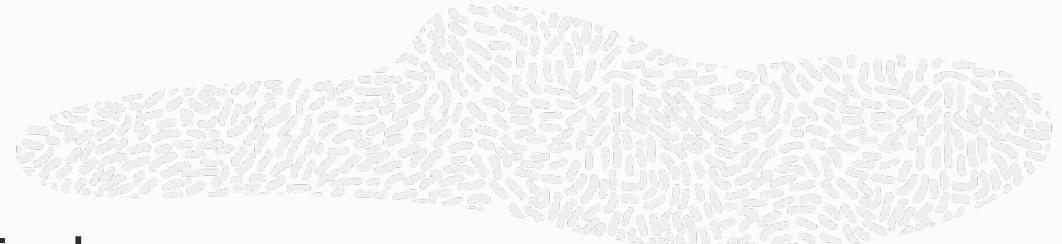
The MY\_PROMOTIONS table has 5,000 rows and no indexes

```
SELECT /*+ INDEX(p) */ Count(*)  
FROM   my_promotions p  
WHERE  promo_category = 'TV'  
AND    promo_begin_date = '05-OCT-20';
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	15	3 (100)	
1	SORT AGGREGATE		1	15	3 (0)	
* 2	TABLE ACCESS FULL	MY_PROMOTIONS	1	15	3 (0)	00:00:01

# Why Are Optimizer Hints Ignored?

Invalid Hint as No Indexes exist on



## Specifying an index hint on a table with no indexes

```
SELECT /*+ INDEX(p) */ Count(*)  
FROM   my_promotions p  
WHERE  promo_category = 'TV'  
AND    promo_begin_date = '05-OCT-20';
```

Invalid hint because no indexes exist on the table

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT				3 (100)	
1	SORT AGGREGATE		1	15		
* 2	TABLE ACCESS FULL	MY_PROMOTIONS	1	15	3 (0)	00:00:01

Hint Report (identified by operation id / Query Block Name / Object Alias):

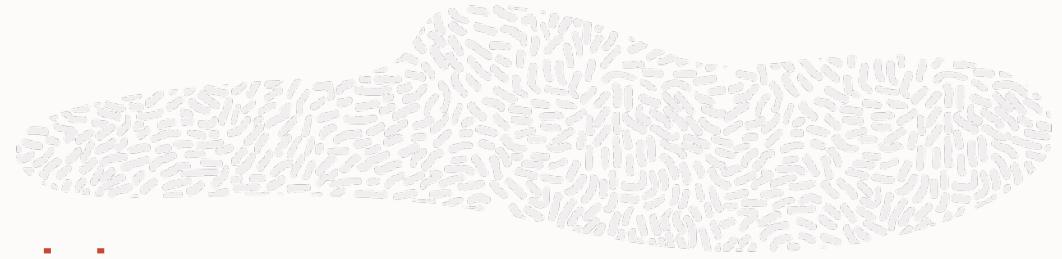
Total hints for statement: 1 (U - Unused (1))

19<sup>c</sup>

---

```
2 - SEL$1 /P@SEL$1  
      U - INDEX(p)
```

# Why Are Optimizer Hints Ignored?



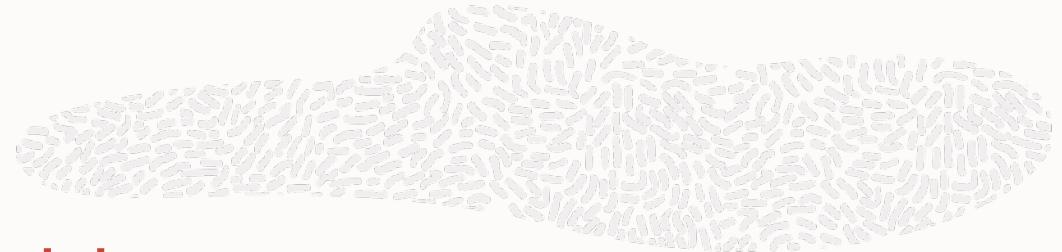
Requesting a hash join hint for **non-equality** join

```
SELECT /*+ USE_HASH(e s) */ e.first_name, e.last_name  
FROM   employees e, salary_grade s  
WHERE  e.salary BETWEEN s.low_sal AND s.high_sal;
```

	Id	Operation		Name		Rows		Bytes		Cost	(%CPU)	
	0	SELECT STATEMENT				1		45		4	(0)	
	1	NESTED LOOPS				1		45		4	(0)	
*	2	TABLE ACCESS STORAGE FULL	I	SALARY_GRADE	I	1		26		2	(0)	
*	3	TABLE ACCESS STORAGE FULL	I	EMPLOYEES	I	1		19		2	(0)	

# Why Are Optimizer Hints Ignored?

Illegal hint



Requesting a hash join hint for **non-equality** join

```
SELECT /*+ USE_HASH(e s) */ e.first_name, e.last_name  
FROM   employees e, salary_grade s  
WHERE  e.salary BETWEEN s.low_sal AND s.high_sal;
```

Id   Operation	Name	Rows	Bytes	Cost (CPU)
0   SELECT STATEMENT		1	45	4 (0)
1   NESTED LOOPS		1	45	4 (0)
* 2   TABLE ACCESS STORAGE FULL	SALARY_GRADE	1	26	2 (0)
* 3   TABLE ACCESS STORAGE FULL	EMPLOYEES	1	19	2 (0)

Illegal hint because a hash join can't be used for a non-equality join predicate

Hint Report (identified by operation id / Query Block Name / Object Alias):  
Total hints for statement: 1 (U - **Unused** (1))

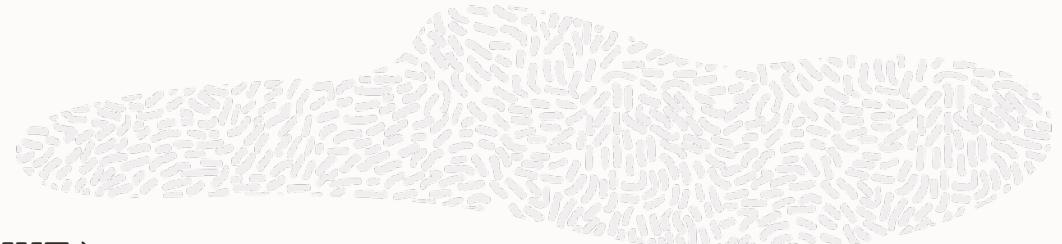
---

```
2 - SEL$1 /P@SEL$1  
      U - USE_HASH(e s)
```

19c

# Why Are Optimizer Hints Ignored?

Requesting a Bloom Filter not to be used



```
SELECT /*+ PX_no_px_join_filter */ SUM(REVENUE)
FROM   Sales S, Customers c
WHERE  s.cust_id=c.c_cust_id;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)
0	SELECT STATEMENT				3792 (100)
1	SORT AGGREGATE		1	17	
* 2	HASH JOIN		4126	70142	3792 (5)
3	JOIN FILTER CREATE	:BF0000	1159	5795	2 (0)
4	TABLE ACCESS STORAGE FULL	CUSTOMERS	1159	5795	2 (0)
5	JOIN FILTER USE	:BF0000	10M	114M	3759 (5)
6	PARTITION RANGE ALL		10M	114M	3759 (5)
* 7	TABLE ACCESS STORAGE FULL	SALES	10M	114M	3759 (5)

# Why Are Optimizer Hints Ignored?

Requesting a Bloom Filter not to be used

```
SELECT /*+ PX_no_px_join_filter */ SUM(REVENUE)
  FROM Sales S, Customers c
 WHERE s.cust_id=c.c_cust_id;
```

Id	Operation	Name	Rows	Bytes	Cost	(%CPU)
0	SELECT STATEMENT				3792	(100)
1	SORT AGGREGATE		1	17		
* 2	HASH JOIN		4126	70142	3792	(5)
3	JOIN FILTER CREATE	:BF0000	1159	5795	2	(0)
4	TABLE ACCESS STORAGE FULL	CUSTOMERS	1159	5795	2	(0)
5	JOIN FILTER USE	:BF0000	10M	114M	3759	(5)
6	PARTITION RANGE ALL		10M	114M	3759	(5)
* 7	TABLE ACCESS STORAGE FULL	SALES	10M	114M	3759	(5)

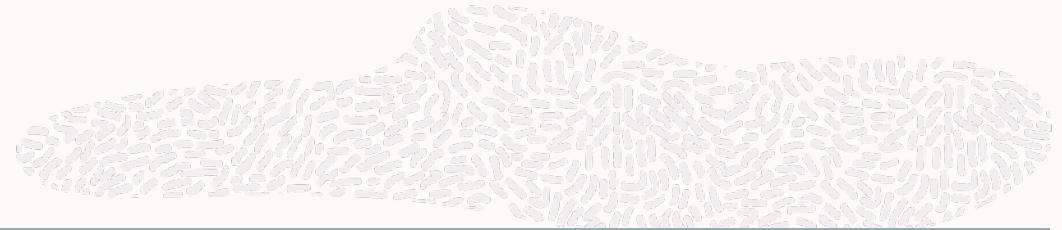
The hint is misspelt, which caused a syntax error.  
The correct hint is  
**NO\_PX\_JOIN\_FILTER**

Hint Report (identified by operation id / Query Block Name / Object Alias): **19c**  
Total hints for statement: 1 (E - Syntax error(1))

1 - SEL\$1

E - PX\_no\_px\_join\_filter

# Tips To Remember

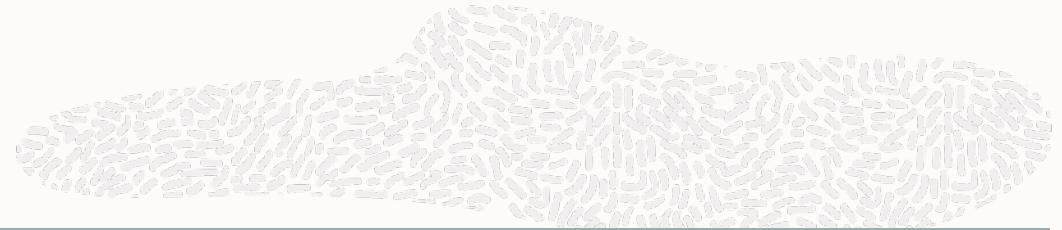


1

## Statistics

Always check and correct  
your cardinality estimates  
first

# Tips To Remember



1

## Statistics

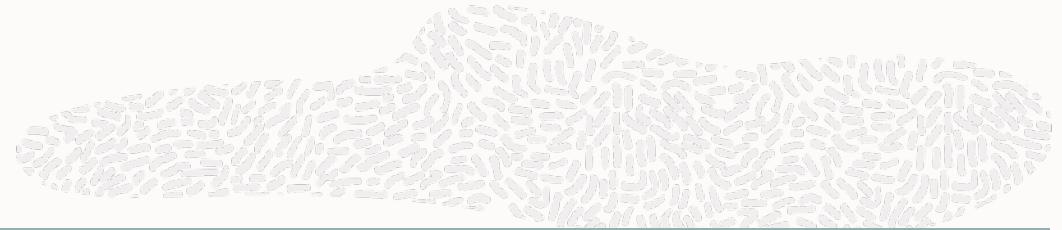
Always check and correct  
your cardinality estimates  
first

2

## SQL Statement

Look at how the SQL  
statement is written and how  
the predicates are being used

# Tips To Remember



1

## Statistics

Always check and correct  
your cardinality estimates  
first

2

## SQL Statement

Look at how the SQL  
statement is written and how  
the predicates are being used

3

## Hints

Only add hints as a last  
resort and confirm your hint  
is really being used

# For More Information Please Visit



<https://sqlmaria.com>



<https://blogs.oracle.com/Optimizer>



<https://twitter.com/@SQLMaria>



<https://www.facebook.com/SQLMaria>

---

## Questions?



# ORACLE

