
Ray - Scalability from a Laptop to a Cluster

Dean Wampler - April 8, 2020

dean@anyscale.com

[@deanwampler](https://twitter.com/deanwampler)

<https://ray.io>

<https://anyscale.com>

Checkout our online events
this Summer:

<https://anyscale.com/events>



© 2019-2020, Anyscale.io



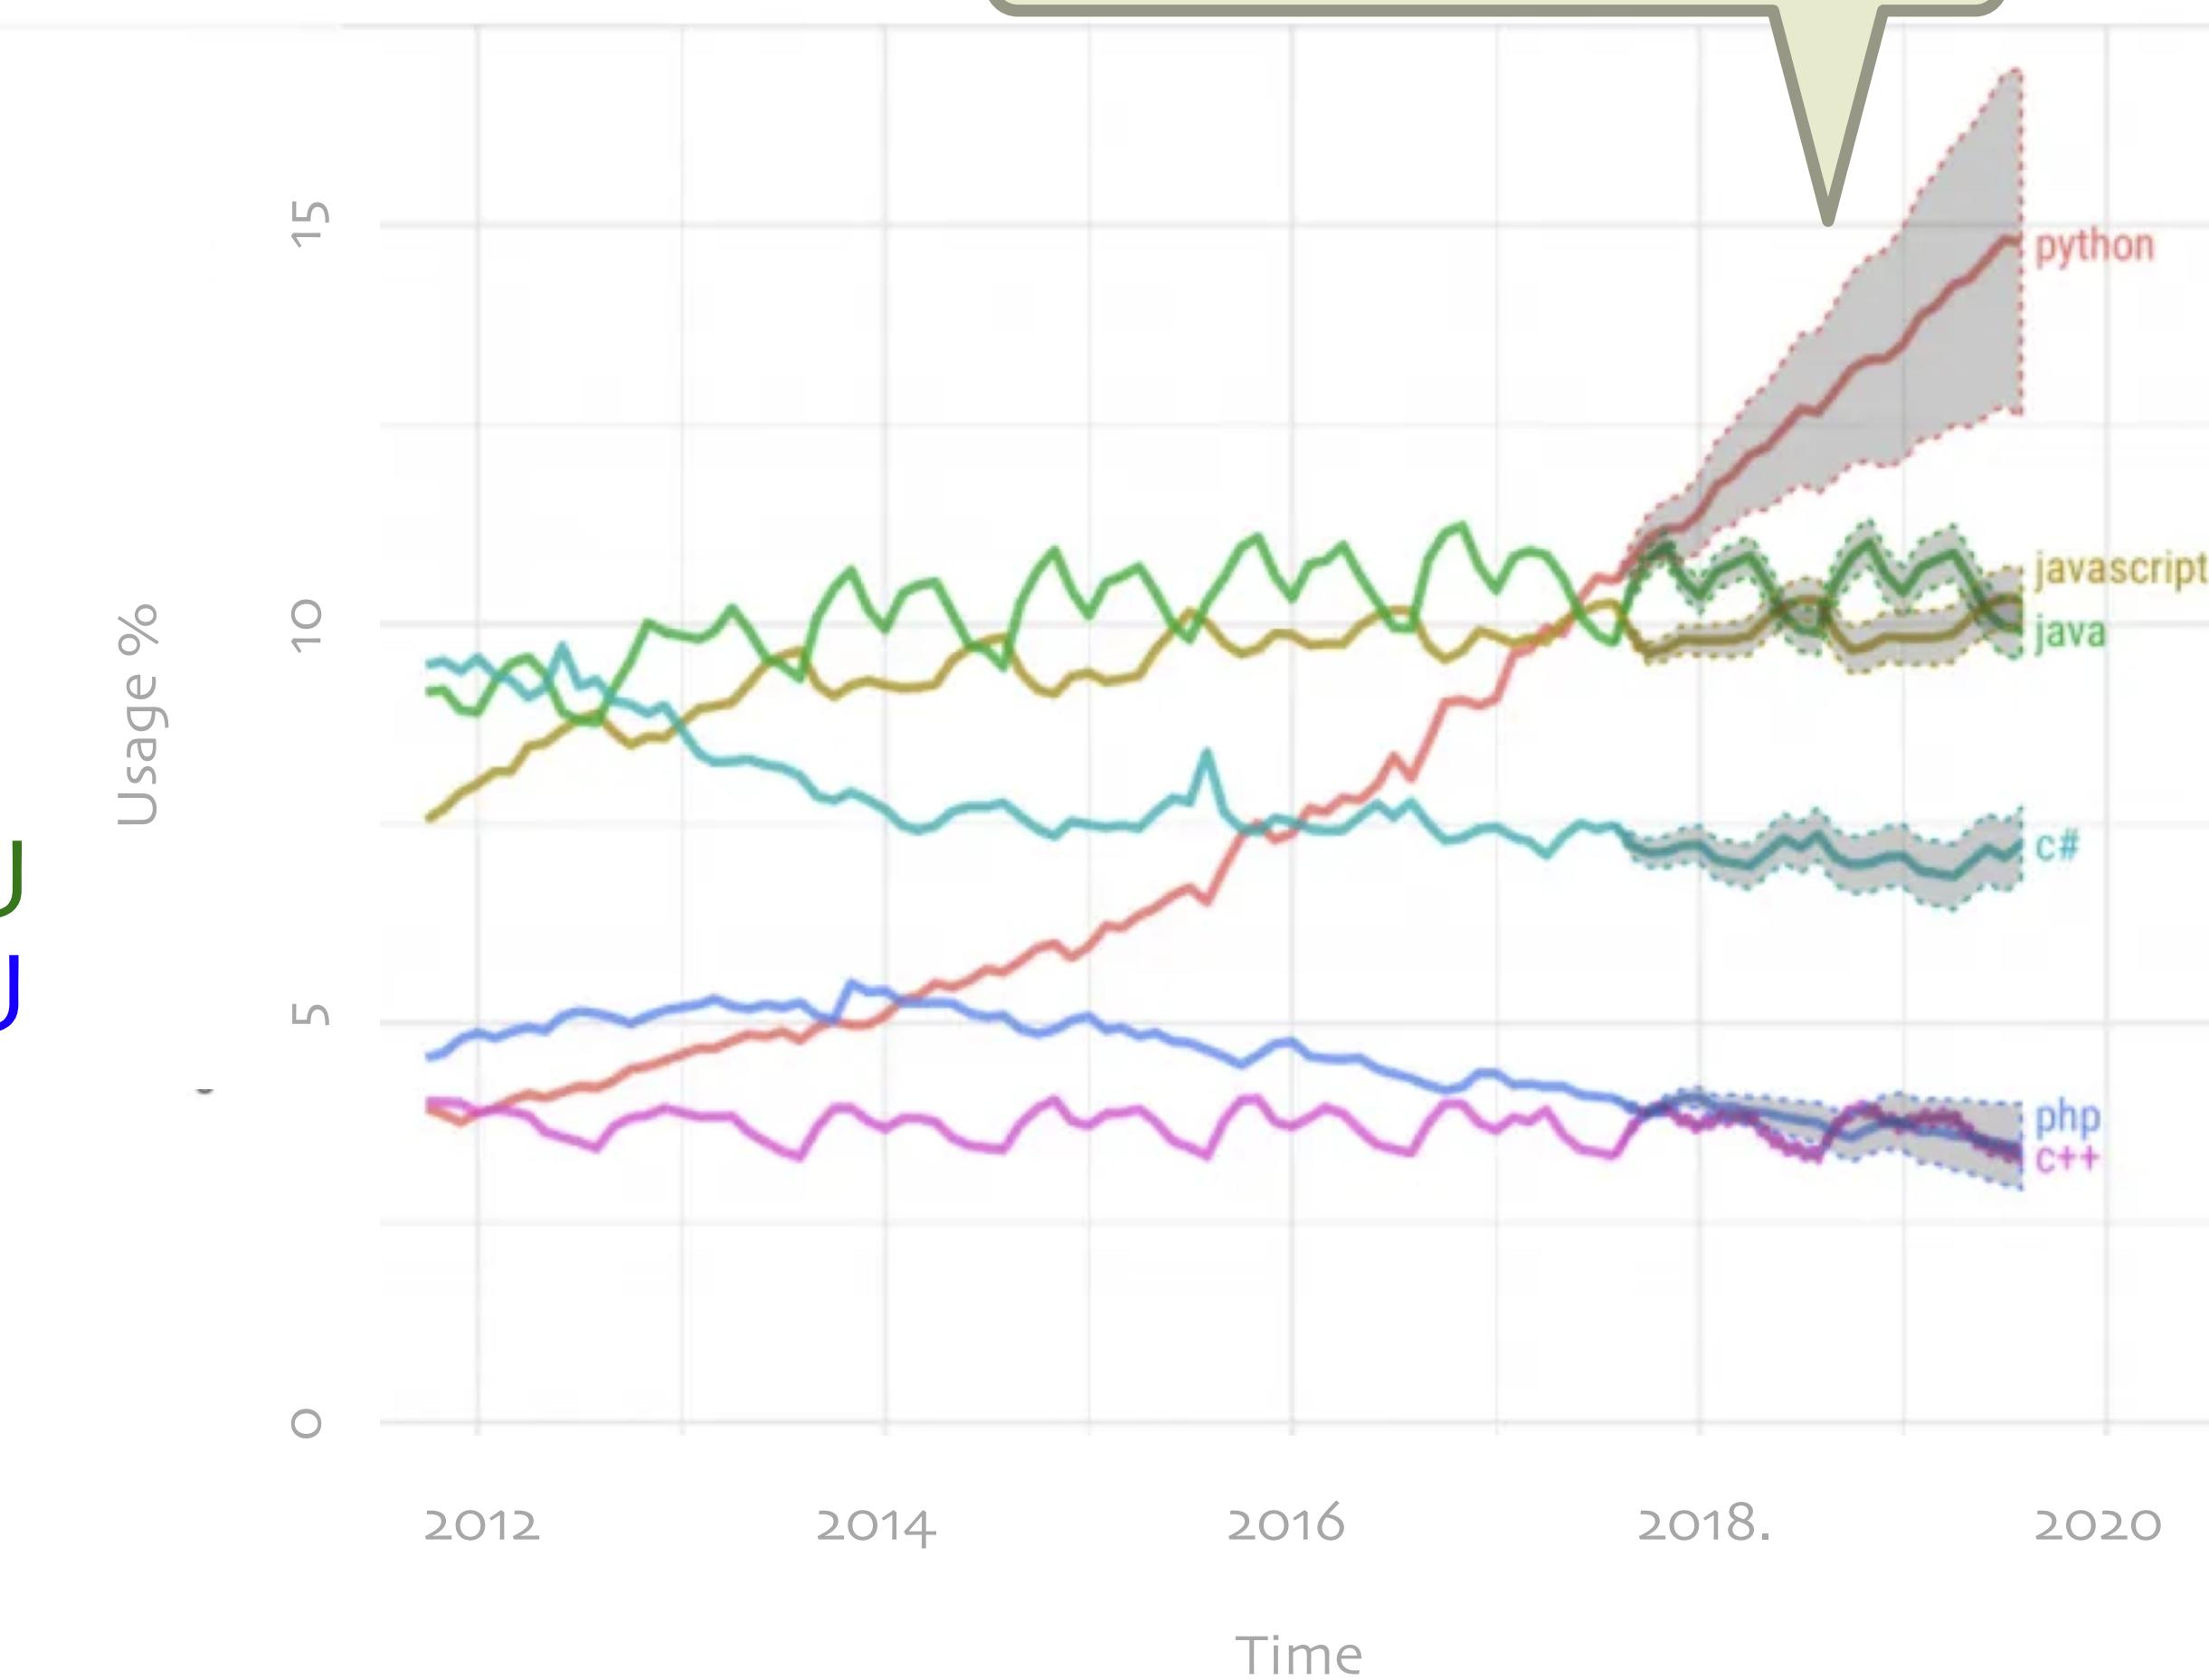
Two Major Trends

Model sizes and therefore compute requirements outstripping Moore's Law



Hence, there is a pressing need for robust, easy to use solutions for distributed Python

Python growth driven by ML/AI and other data science workloads



The ML Landscape Today

All require distributed implementations to scale

Featurization



Streaming



Hyperparam Tuning



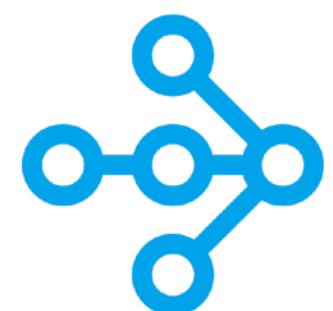
Training



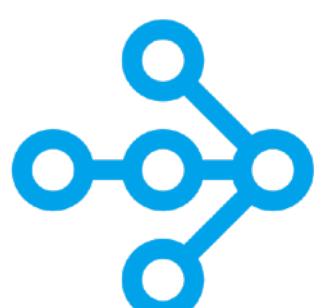
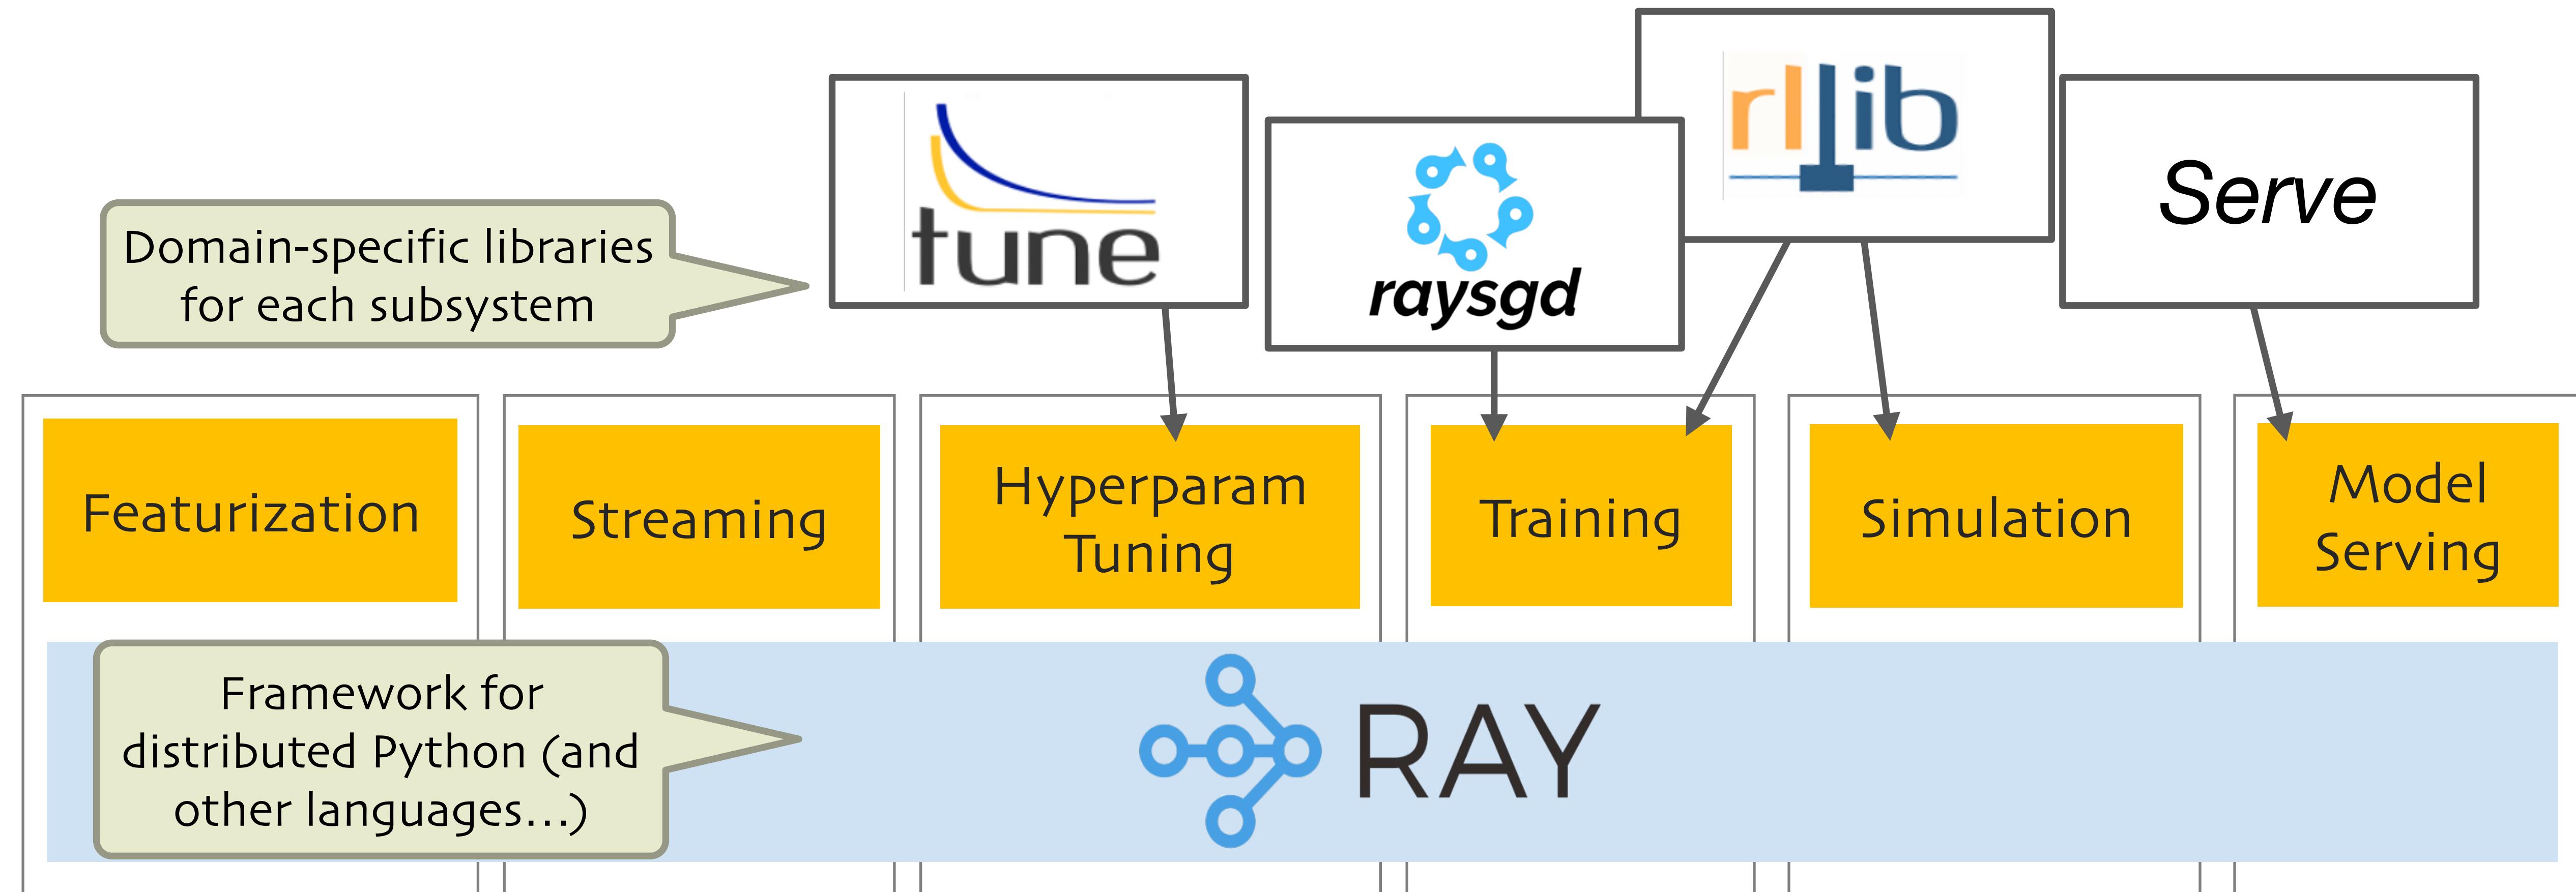
Simulation



Model Serving



The Ray Vision: Sharing a Common Framework



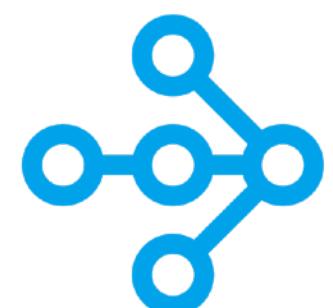
API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
def add_arrays(a, b):  
    return np.add(a, b)
```

The Python you
already know...



API - Designed to Be Intuitive and Concise

Functions -> Tasks

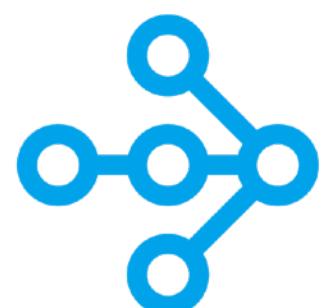
```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

For completeness, add these first:

```
import ray  
import numpy as np  
ray.init()
```

Now these functions
are remote "tasks"



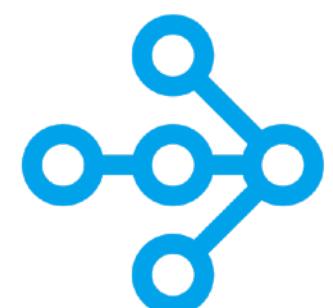
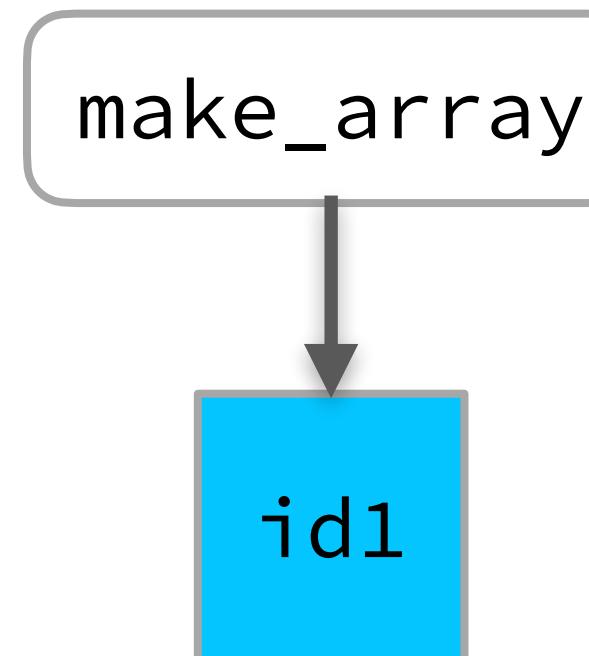
API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

```
id1 = make_array.remote(...)
```



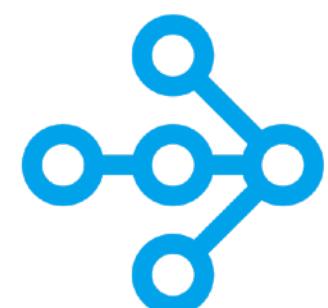
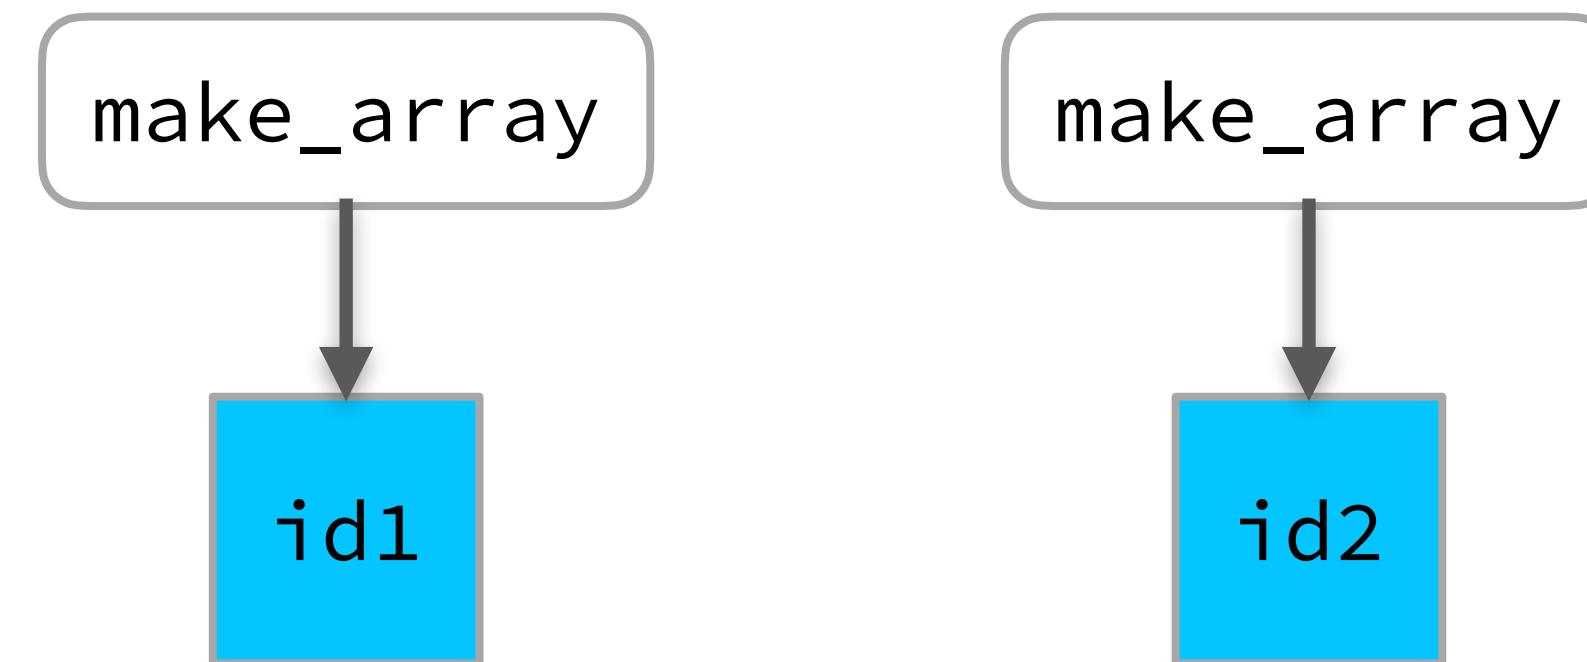
API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

```
id1 = make_array.remote(...)  
id2 = make_array.remote(...)
```



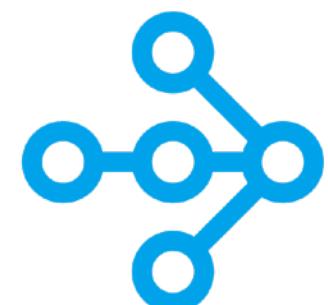
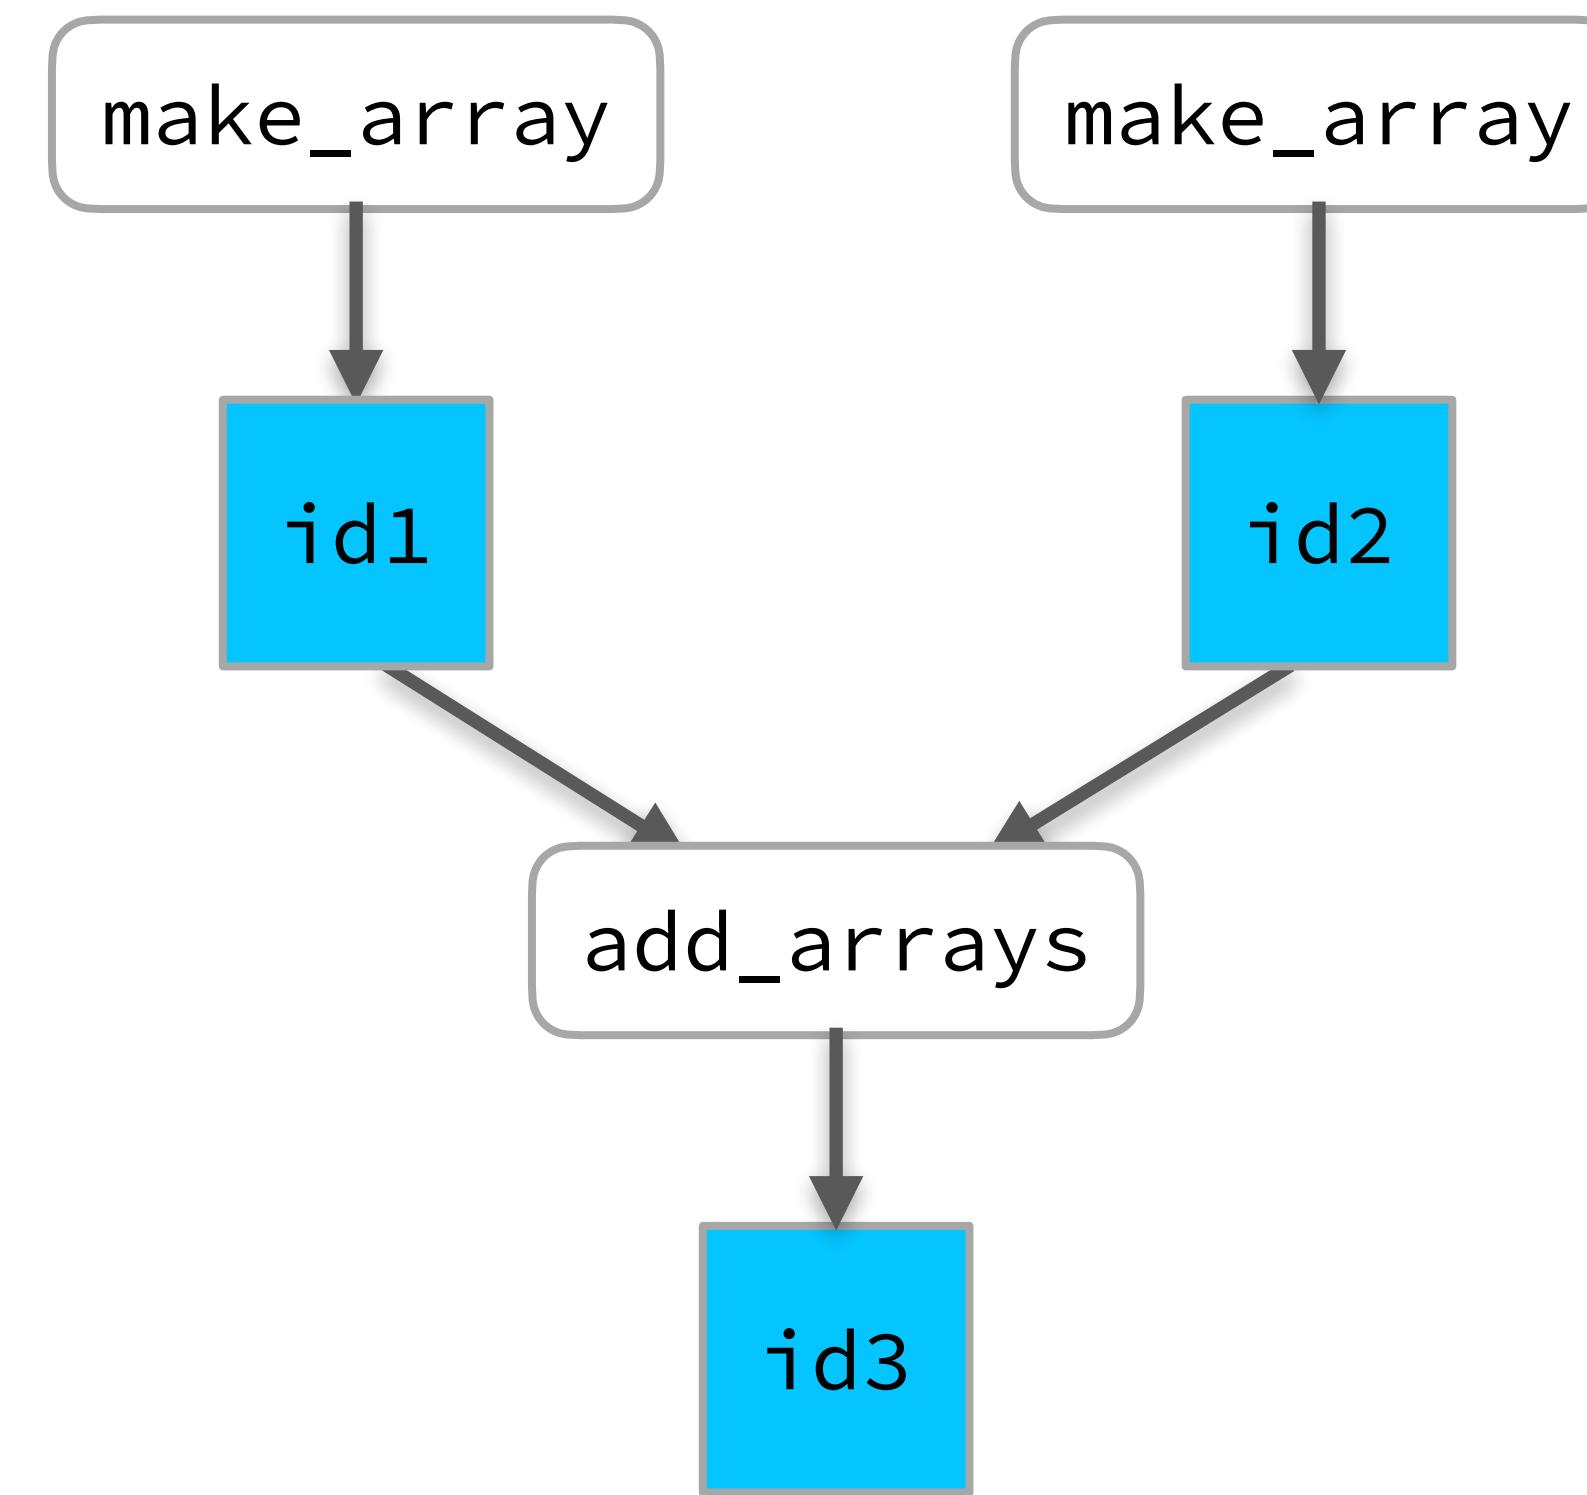
API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

```
id1 = make_array.remote(...)  
id2 = make_array.remote(...)  
id3 = add_arrays.remote(id1, id2)
```



API - Designed to Be Intuitive and Concise

Functions -> Tasks

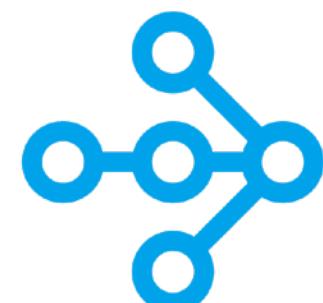
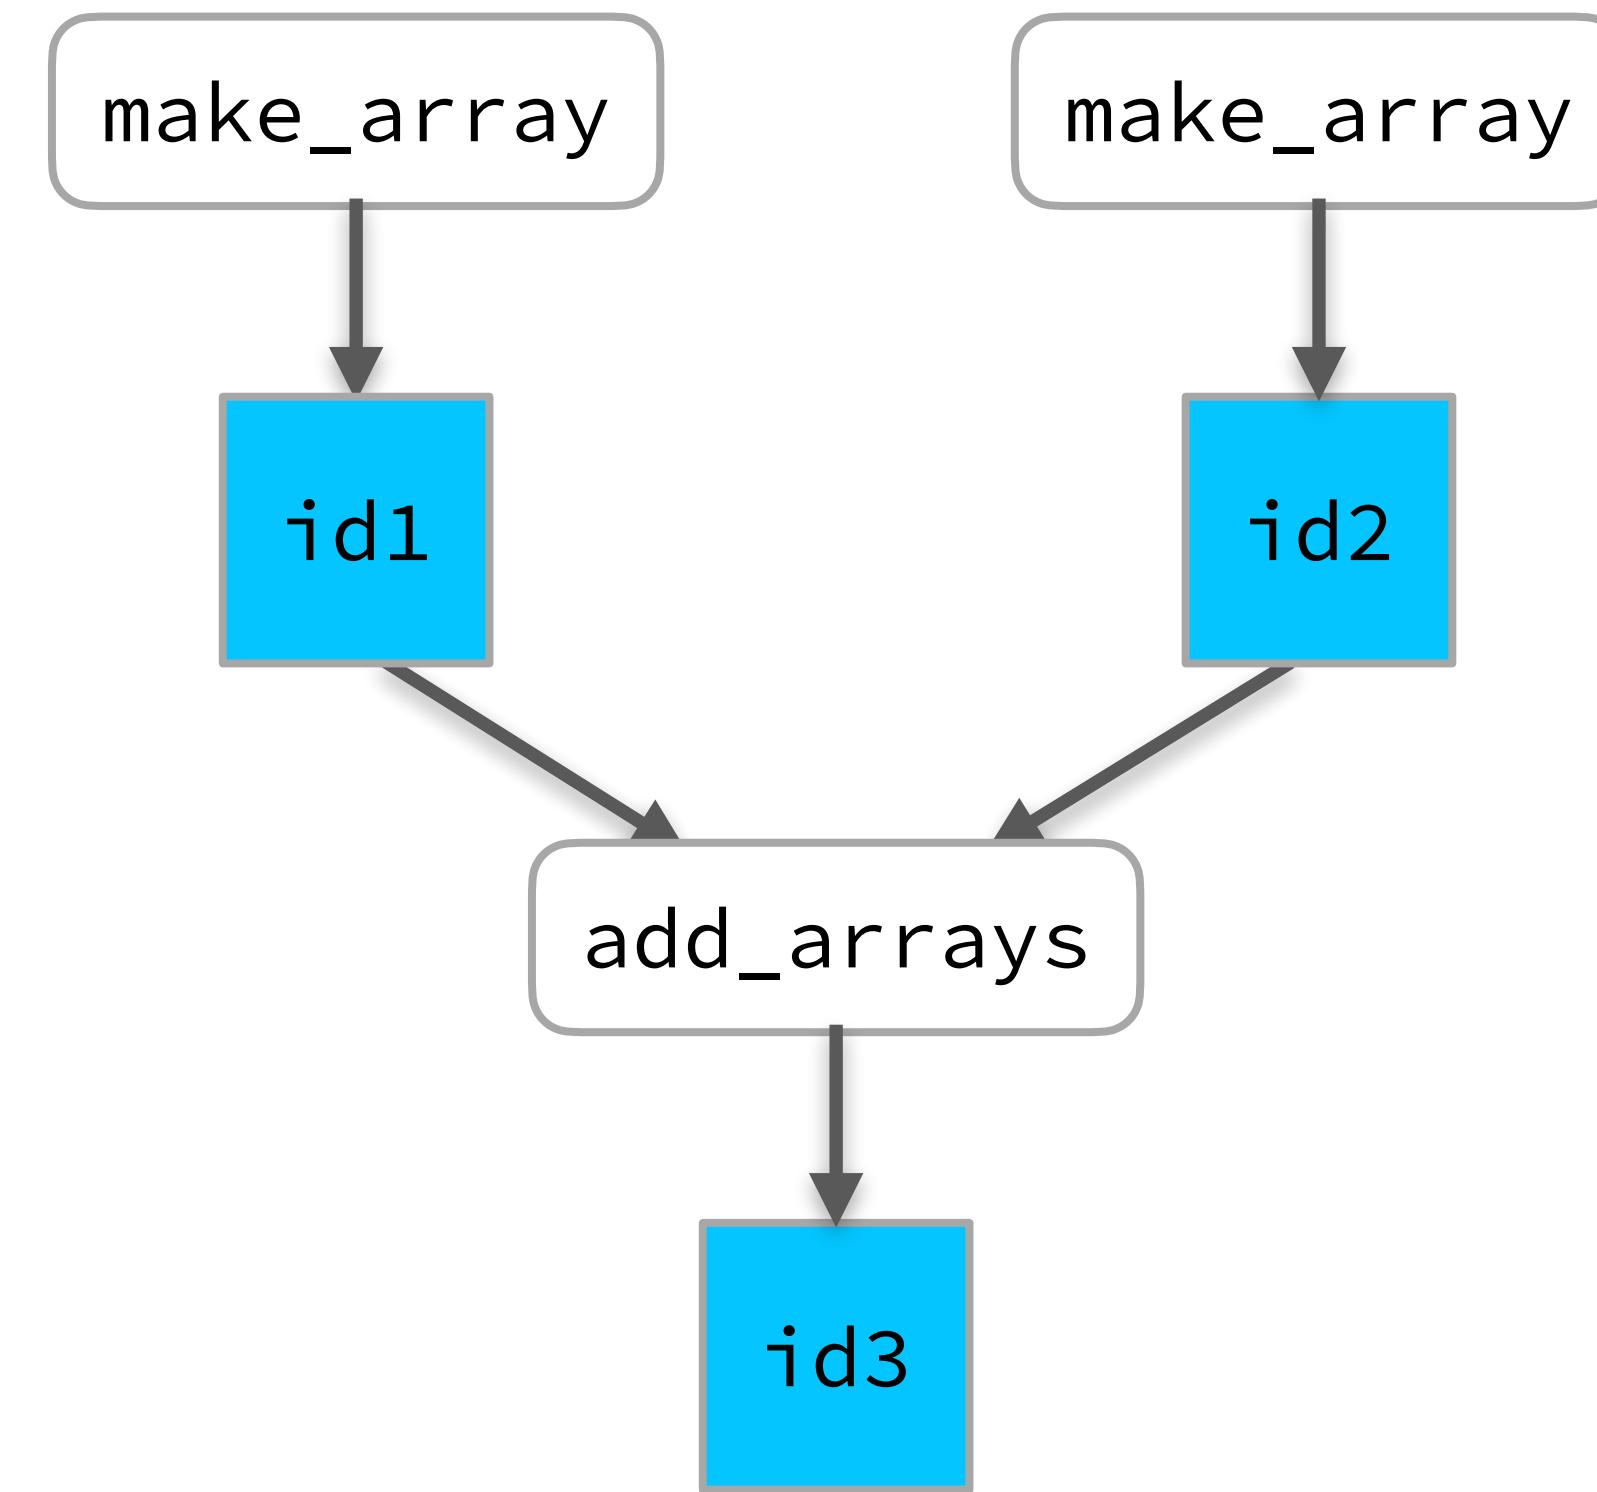
```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

Ray handles extracting the arrays from the object ids

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

```
id1 = make_array.remote(...)  
id2 = make_array.remote(...)  
id3 = add_arrays.remote(id1, id2)  
ray.get(id3)
```

Ray handles sequencing of async dependencies



What about
distributed
state?

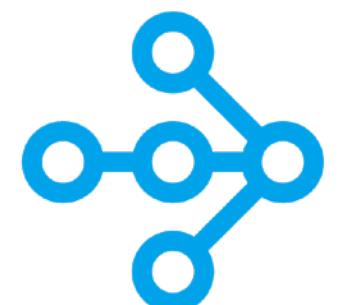
API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote
def make_array(...):
    a = ... # Construct a NumPy array
    return a
```

```
@ray.remote
def add_arrays(a, b):
    return np.add(a, b)
```

```
id1 = make_array.remote(...)
id2 = make_array.remote(...)
id3 = add_arrays.remote(id1, id2)
ray.get(id3)
```



API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

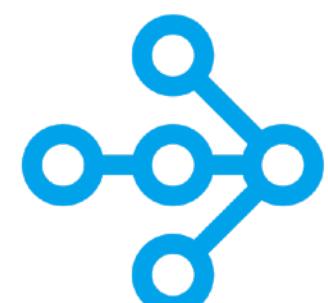
```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

```
id1 = make_array.remote(...)  
id2 = make_array.remote(...)  
id3 = add_arrays.remote(id1, id2)  
ray.get(id3)
```

Classes -> Actors

```
class Counter(object):  
    def __init__(self):  
        self.value = 0  
    def increment(self):  
        self.value += 1  
    return self.value
```

The Python
classes you
love...



API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a  
  
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)  
  
id1 = make_array.remote(...)  
id2 = make_array.remote(...)  
id3 = add_arrays.remote(id1, id2)  
ray.get(id3)
```

... now a remote
“actor”

Classes -> Actors

```
@ray.remote  
class Counter(object):  
    def __init__(self):  
        self.value = 0  
    def increment(self):  
        self.value += 1  
        return self.value  
    def get_count(self):  
        return self.value
```

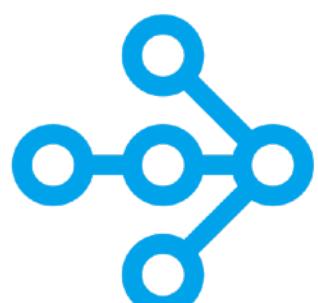
You need a
“getter” method
to read the state.



API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a  
  
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)  
  
id1 = make_array.remote(...)  
id2 = make_array.remote(...)  
id3 = add_arrays.remote(id1, id2)  
ray.get(id3)
```



Classes -> Actors

```
@ray.remote  
class Counter(object):  
    def __init__(self):  
        self.value = 0  
    def increment(self):  
        self.value += 1  
        return self.value  
    def get_count(self):  
        return self.value  
  
c = Counter.remote()  
id4 = c.increment.remote()  
id5 = c.increment.remote()  
ray.get([id4, id5]) # [1, 2]
```

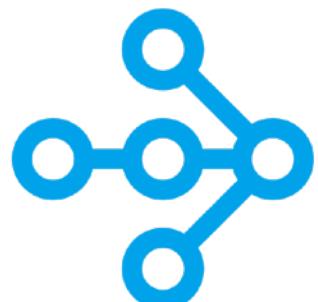
API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote
def make_array(...):
    a = ... # Construct a NumPy array
    return a

@ray.remote
def add_arrays(a, b):
    return np.add(a, b)

id1 = make_array.remote(...)
id2 = make_array.remote(...)
id3 = add_arrays.remote(id1, id2)
ray.get(id3)
```



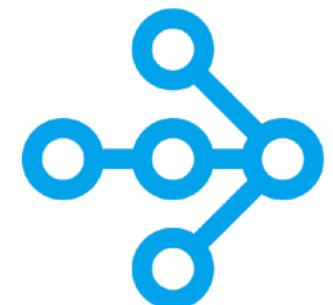
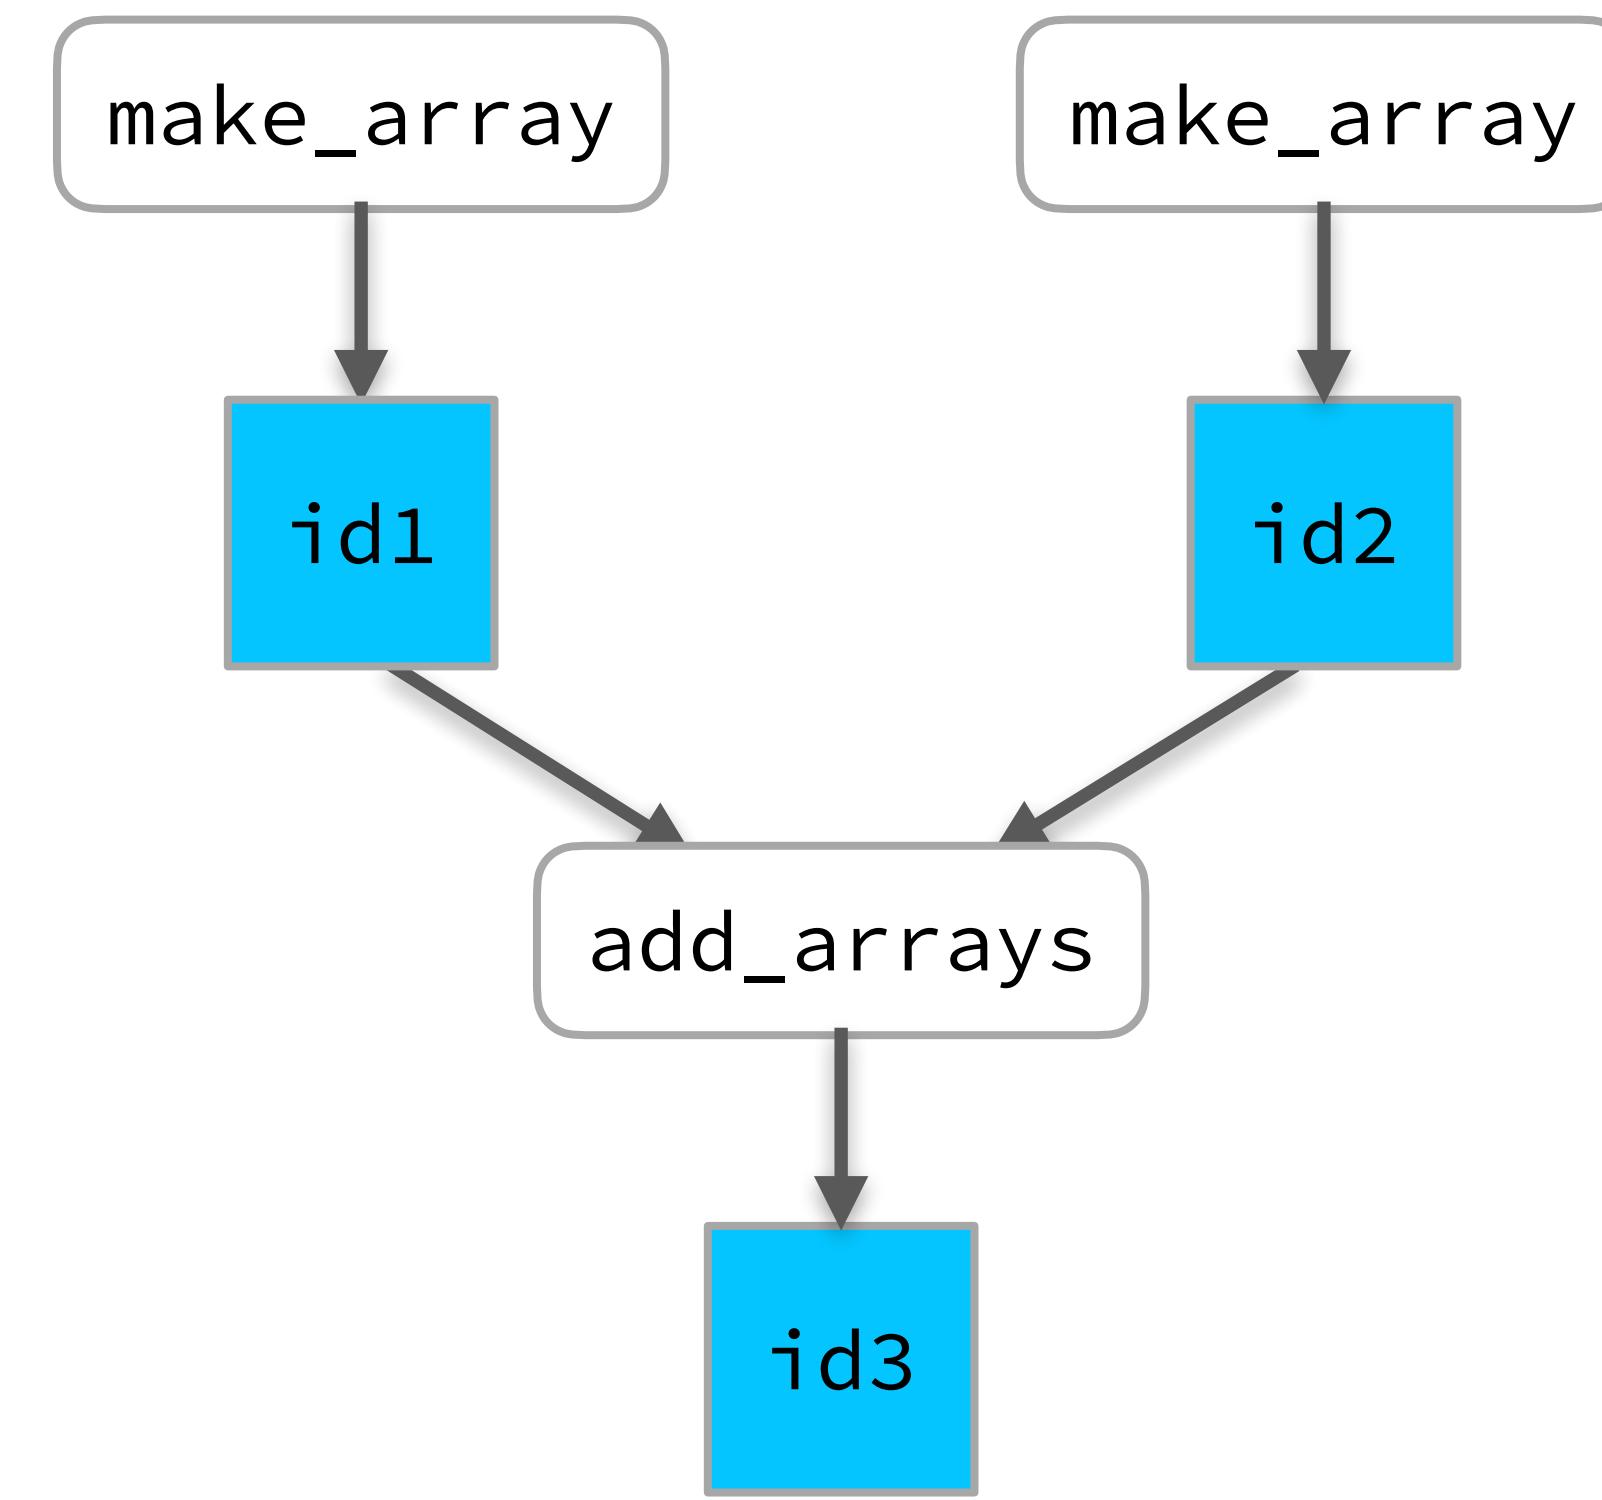
Classes -> Actors

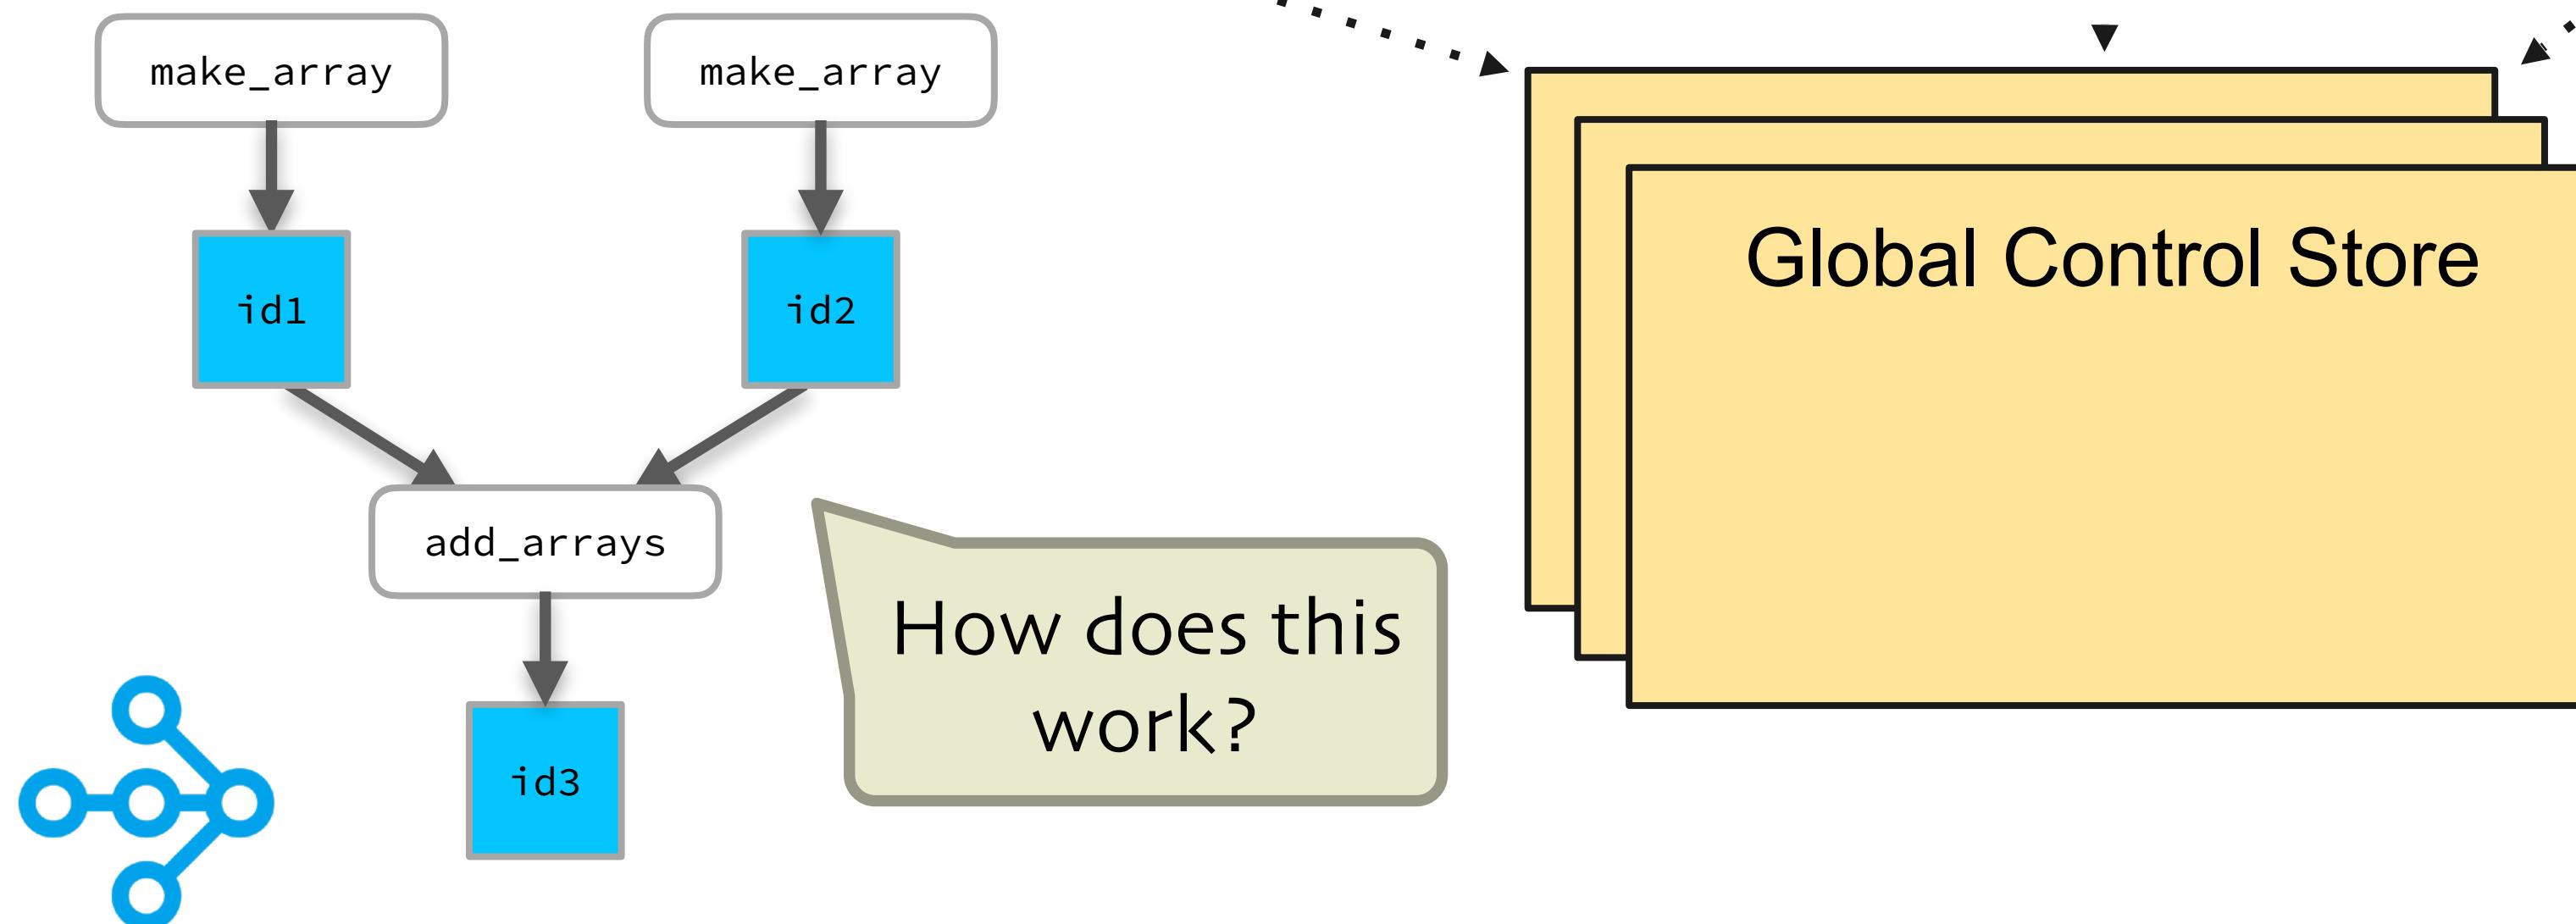
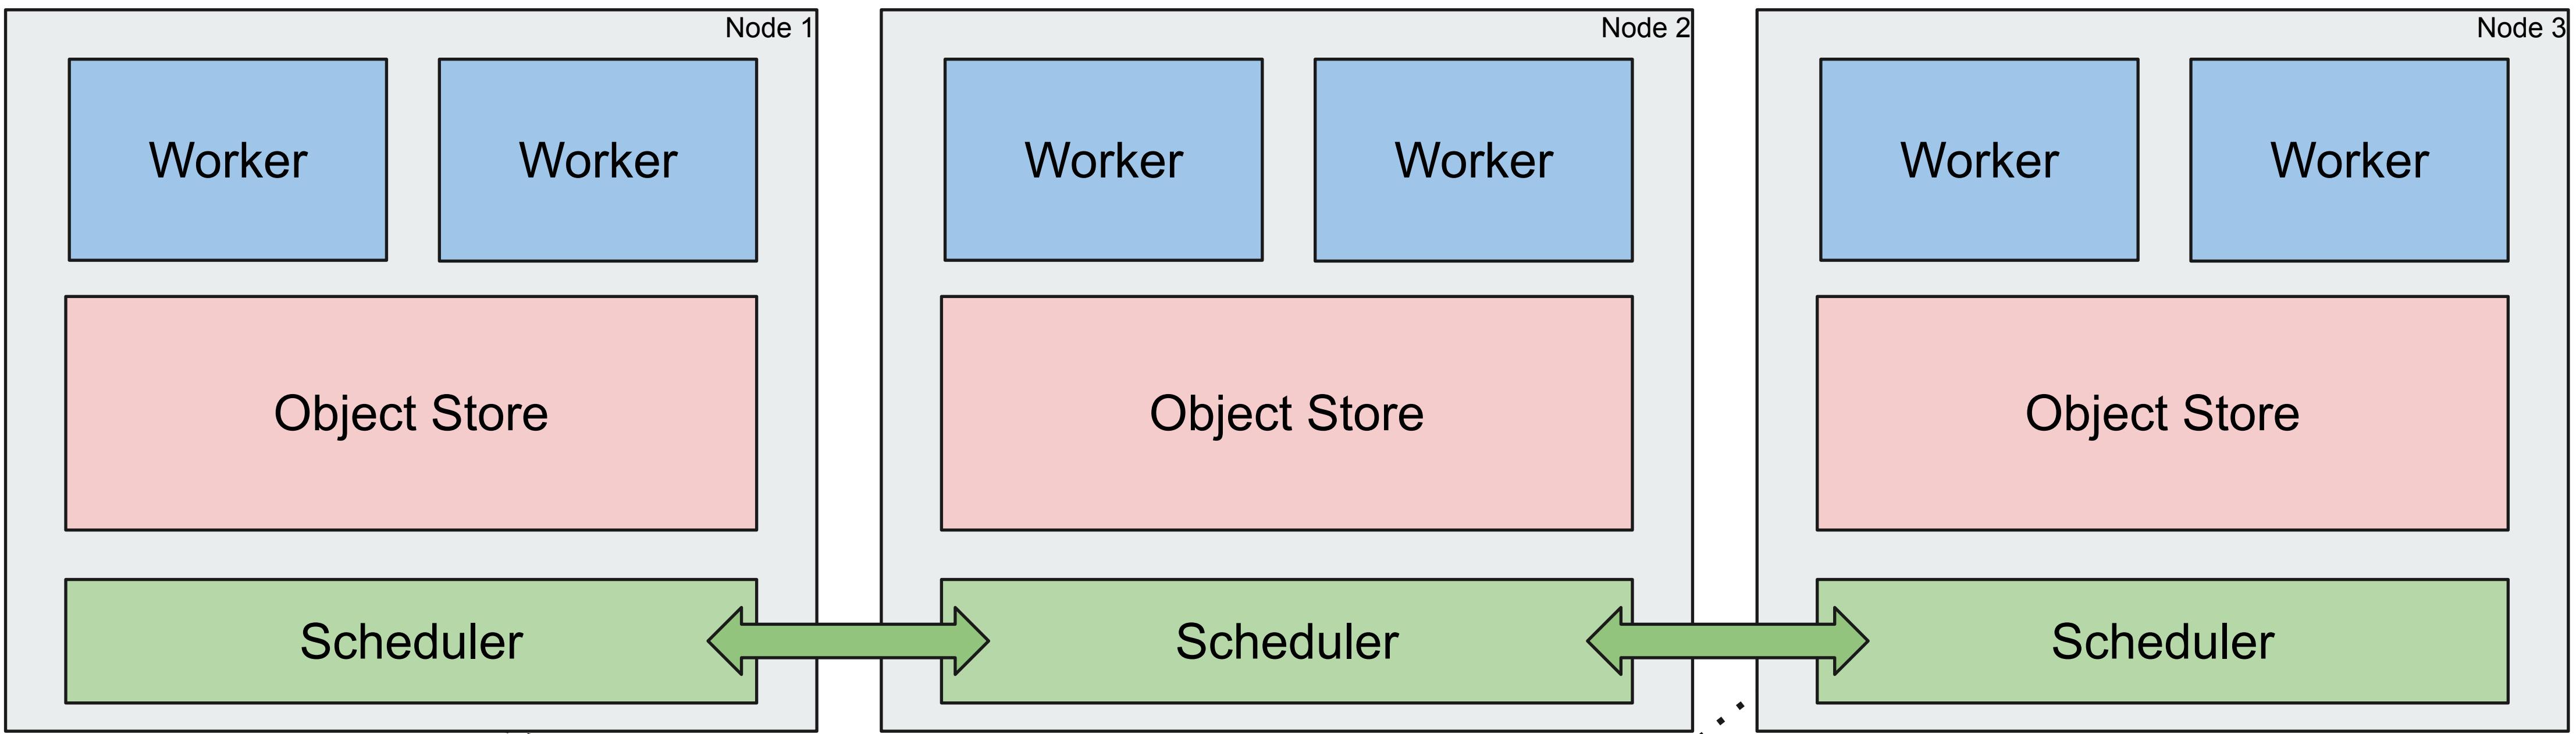
```
@ray.remote(num_gpus=1)
class Counter(object):
    def __init__(self):
        self.value = 0
    def increment(self):
        self.value += 1
        return self.value
    def get_count(self):
        return self.value

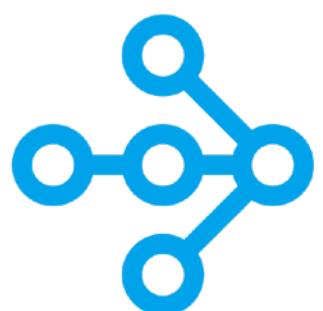
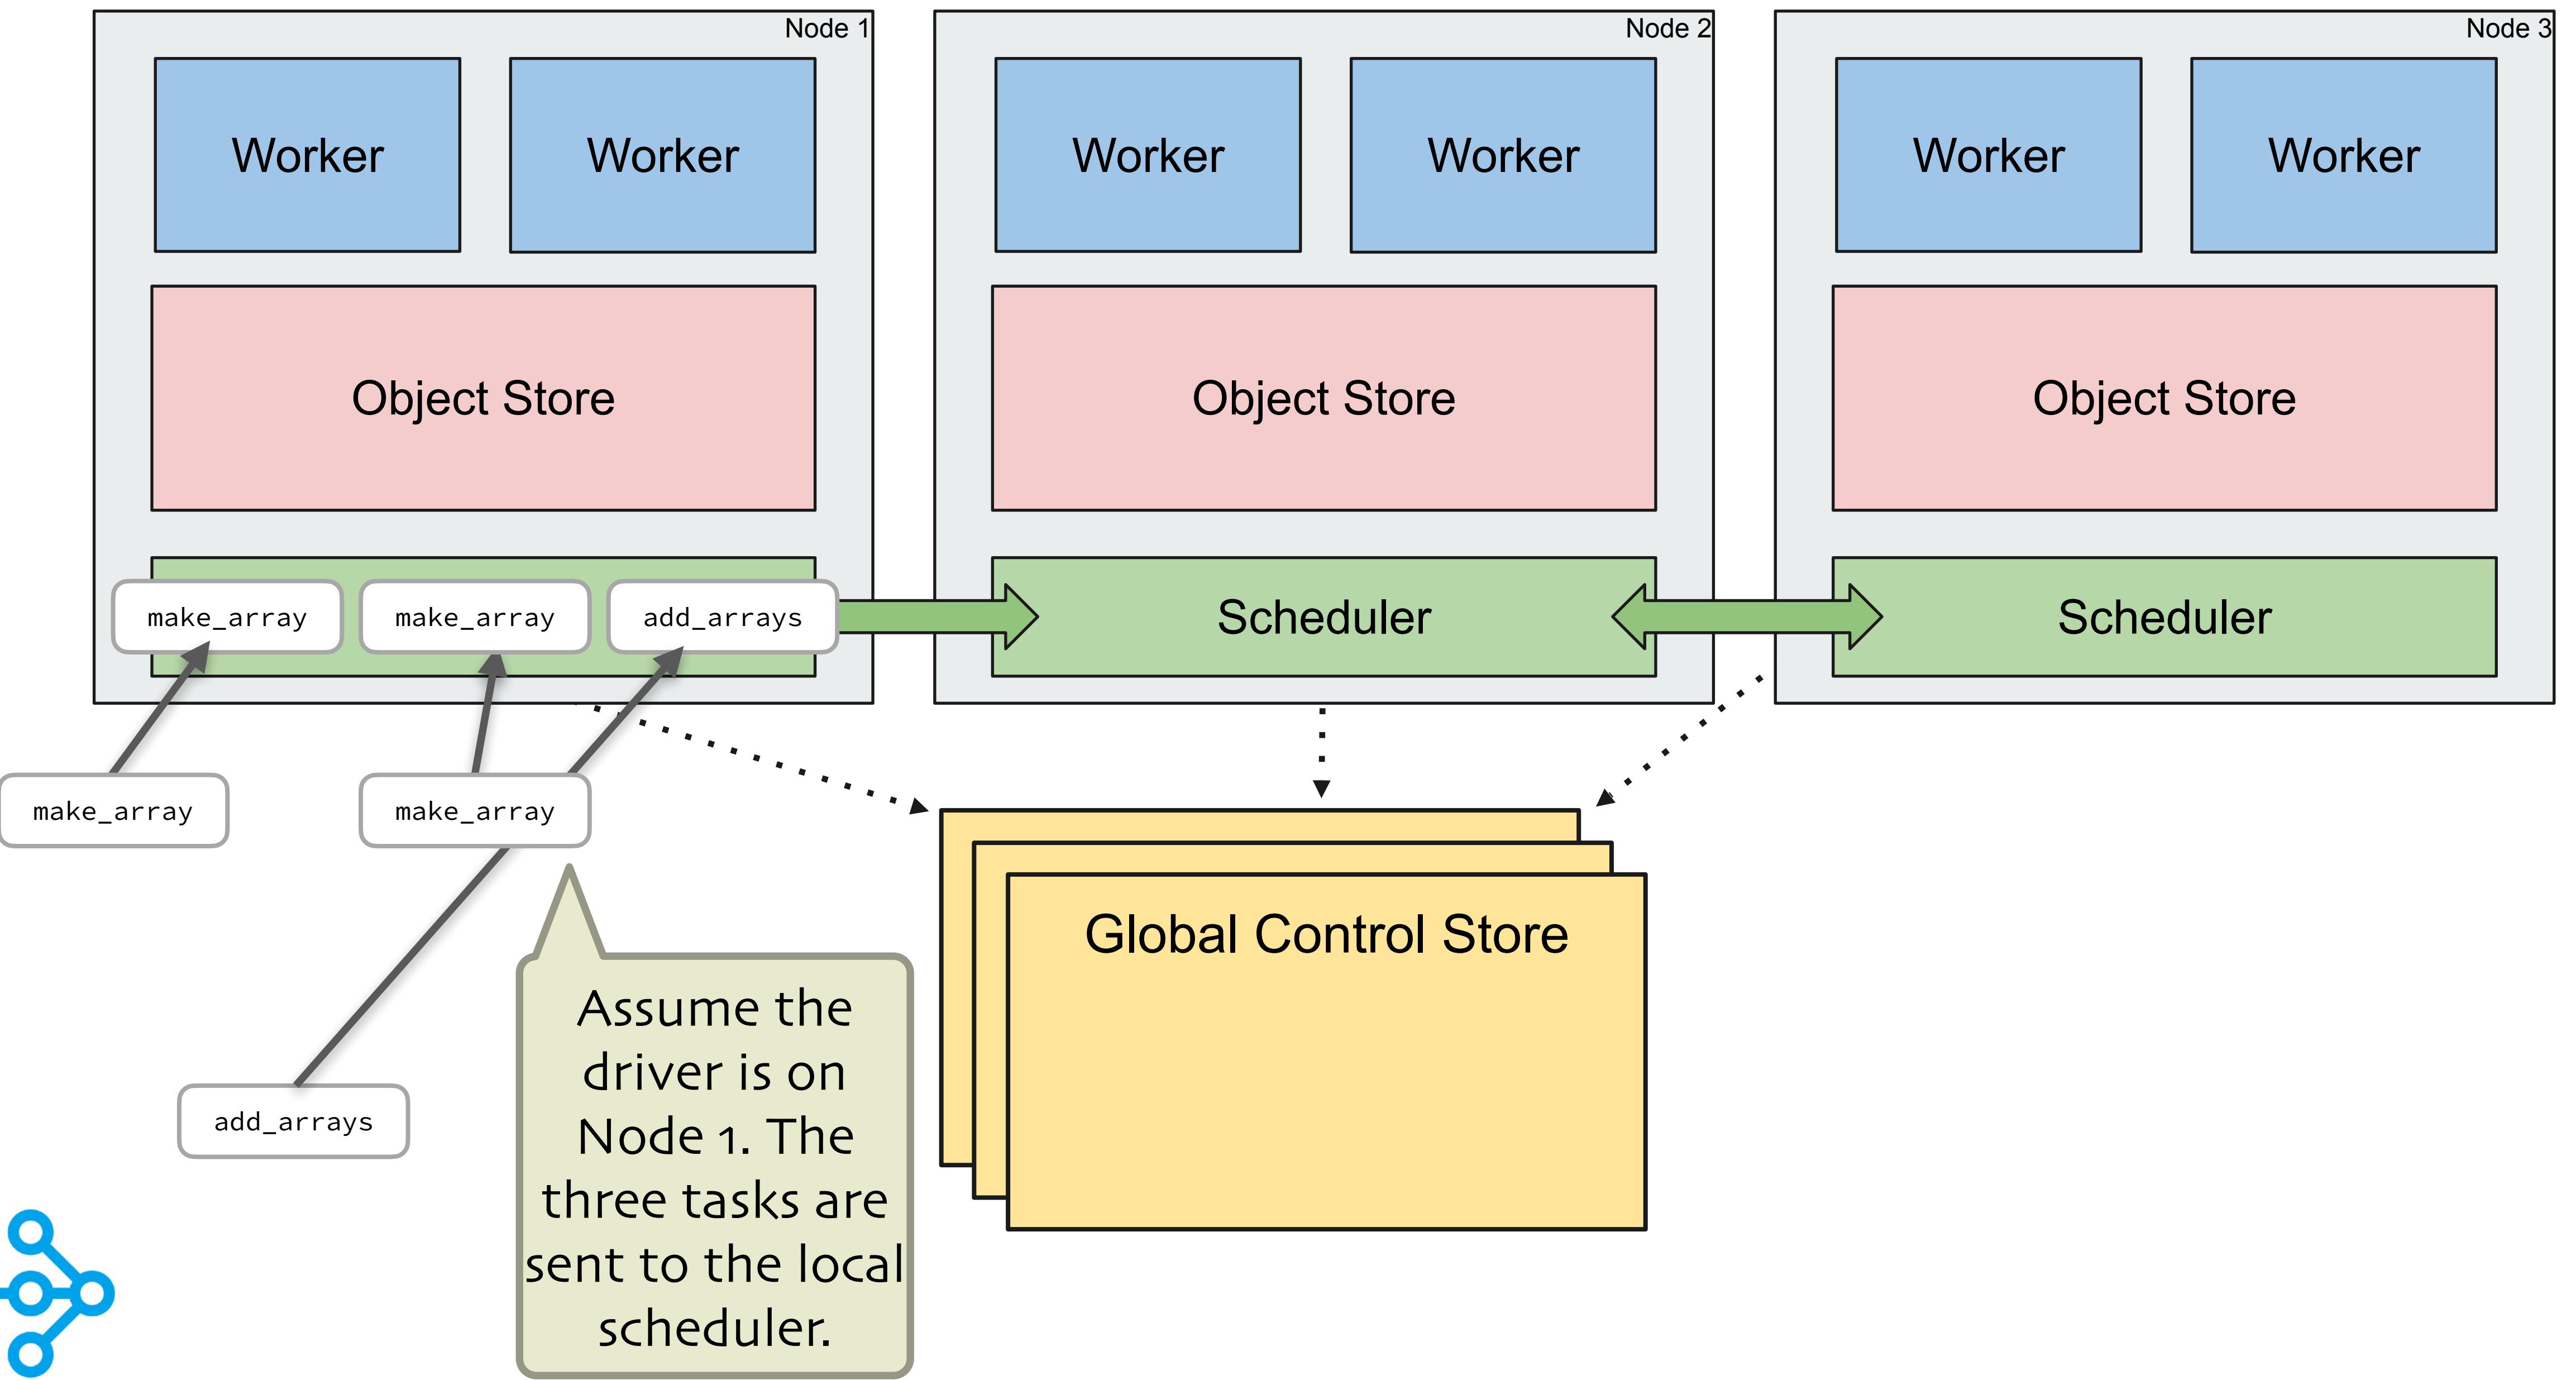
c = Counter.remote()
id4 = c.increment.remote()
id5 = c.increment.remote()
ray.get([id4, id5]) # [1, 2]
```

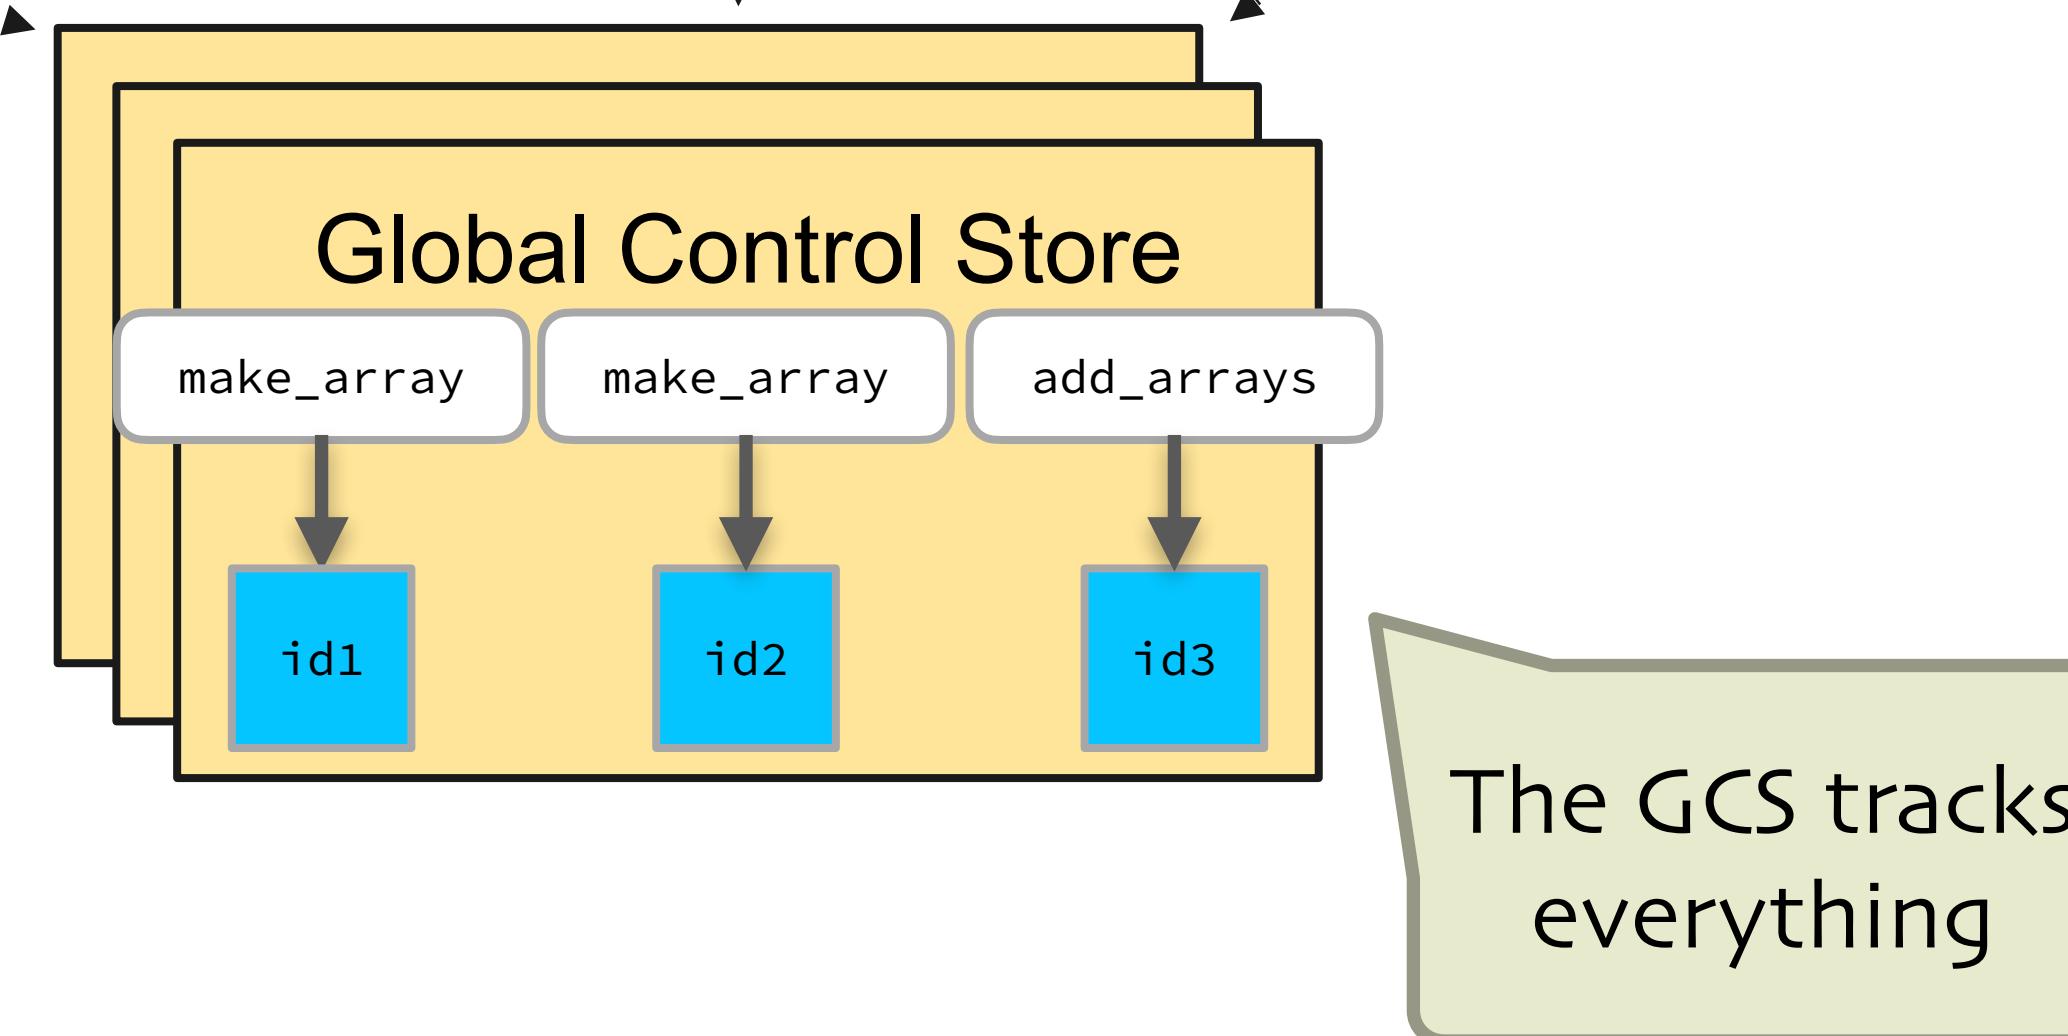
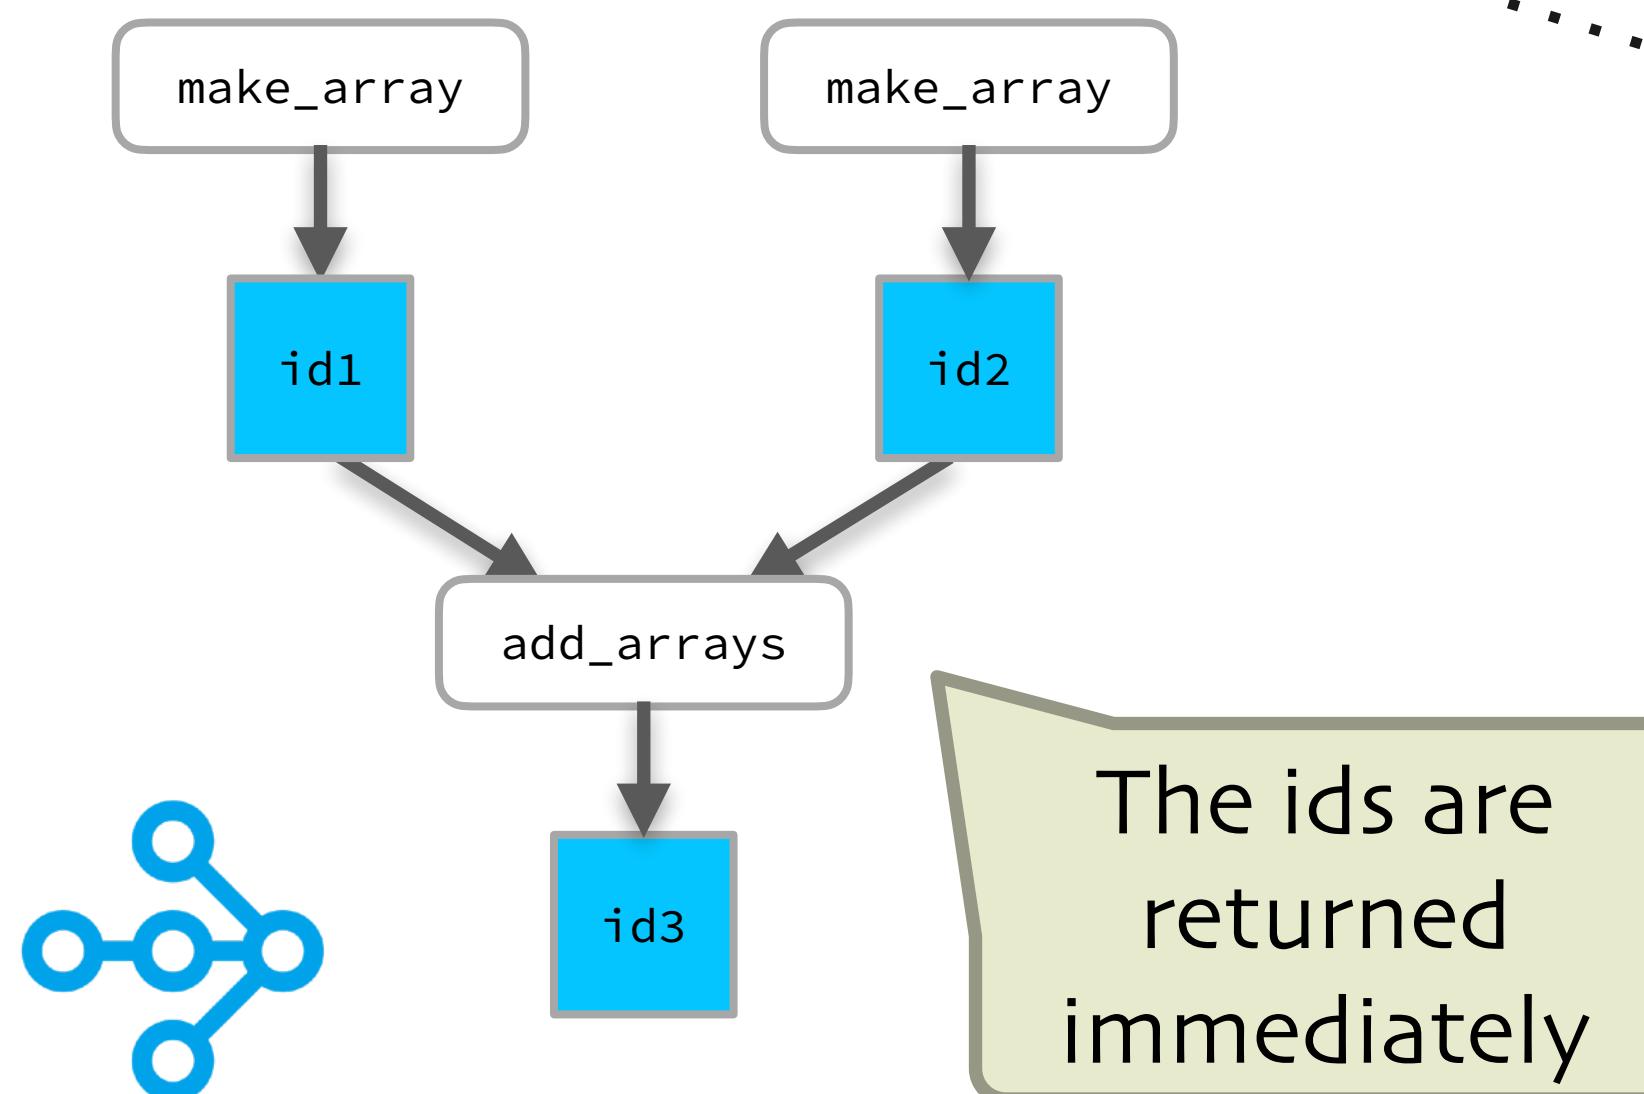
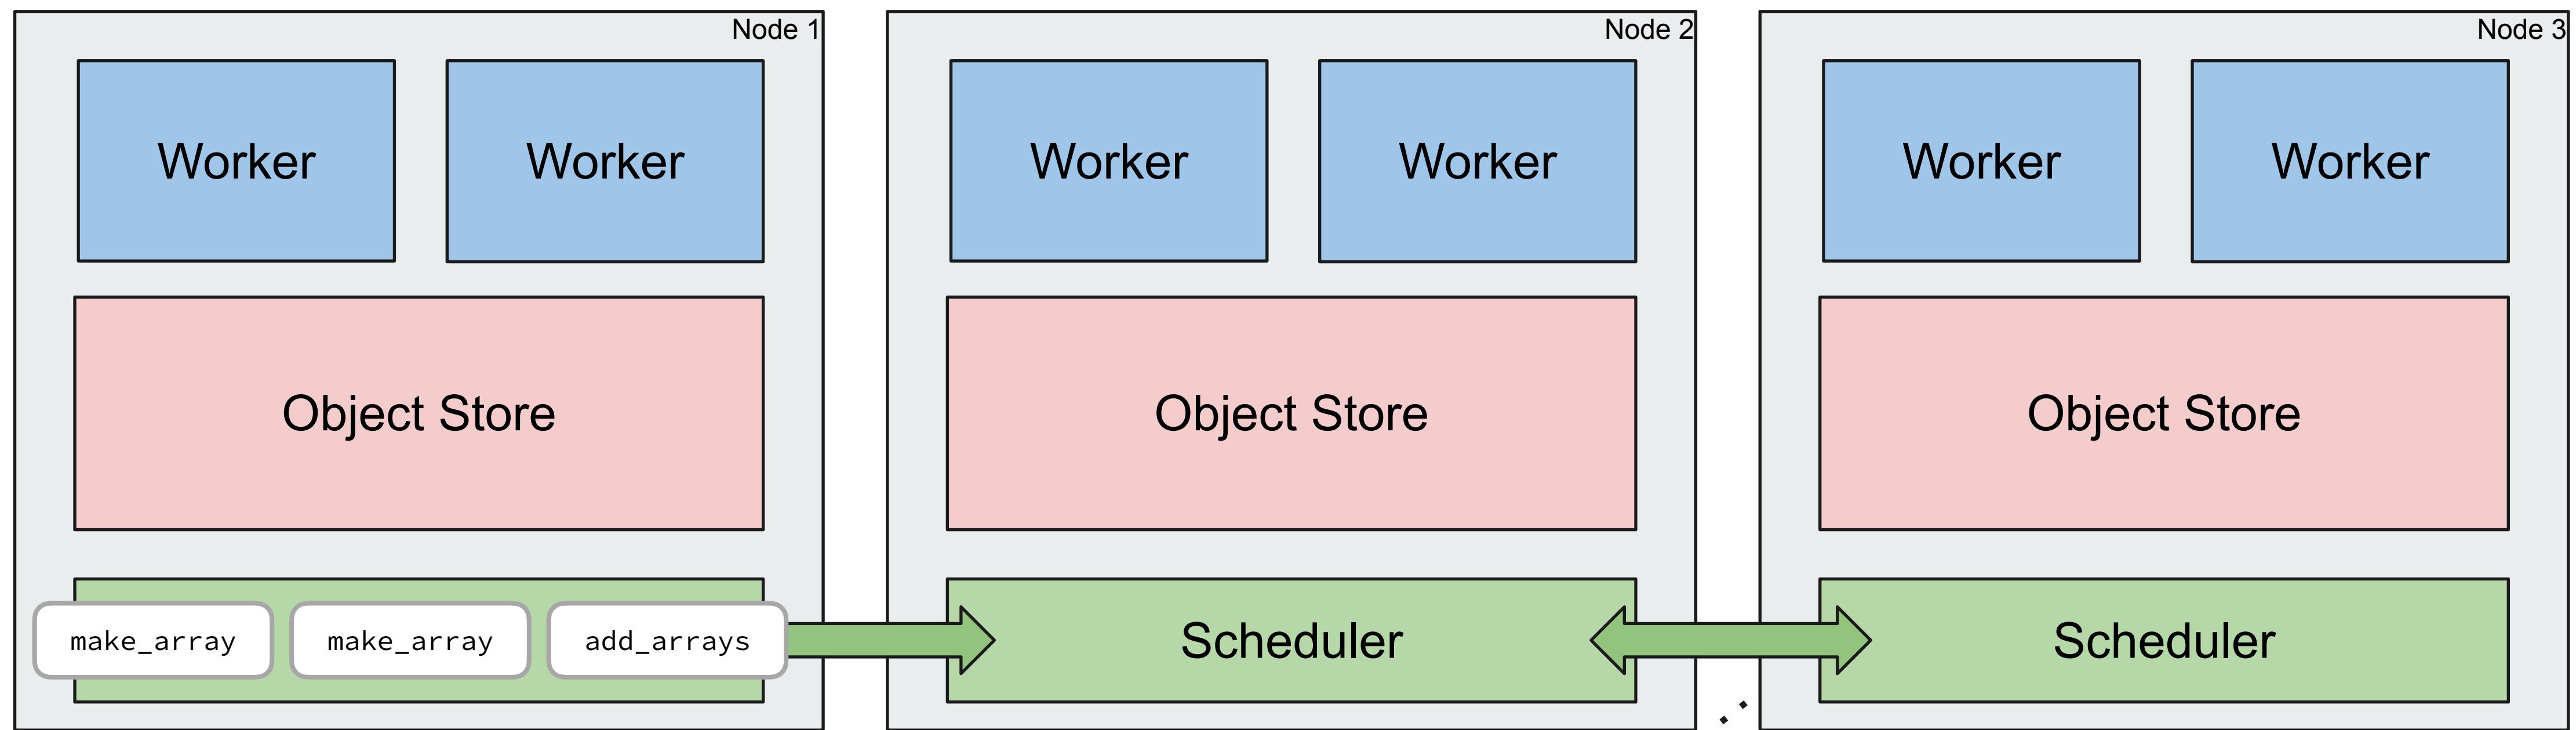
Optional configuration specifications

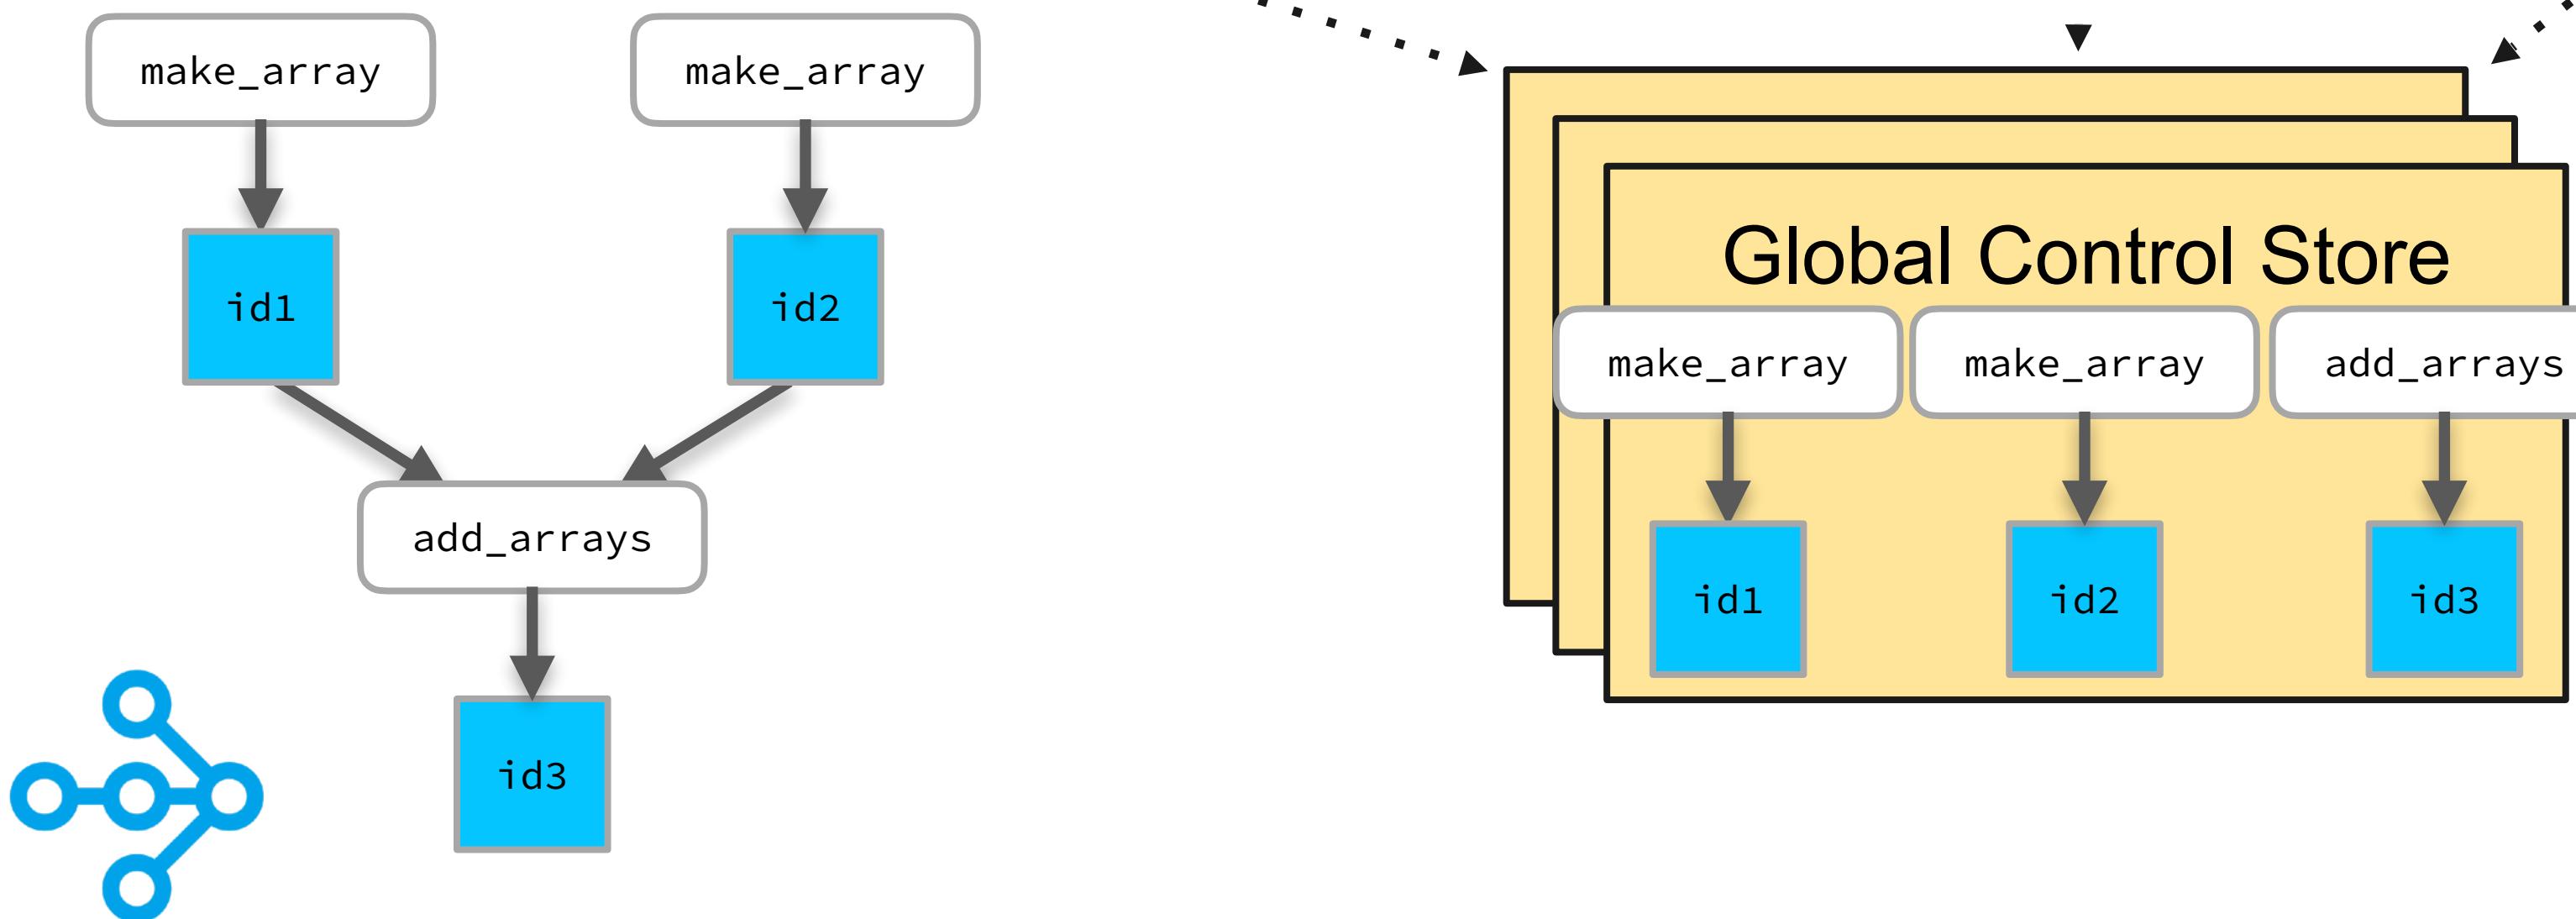
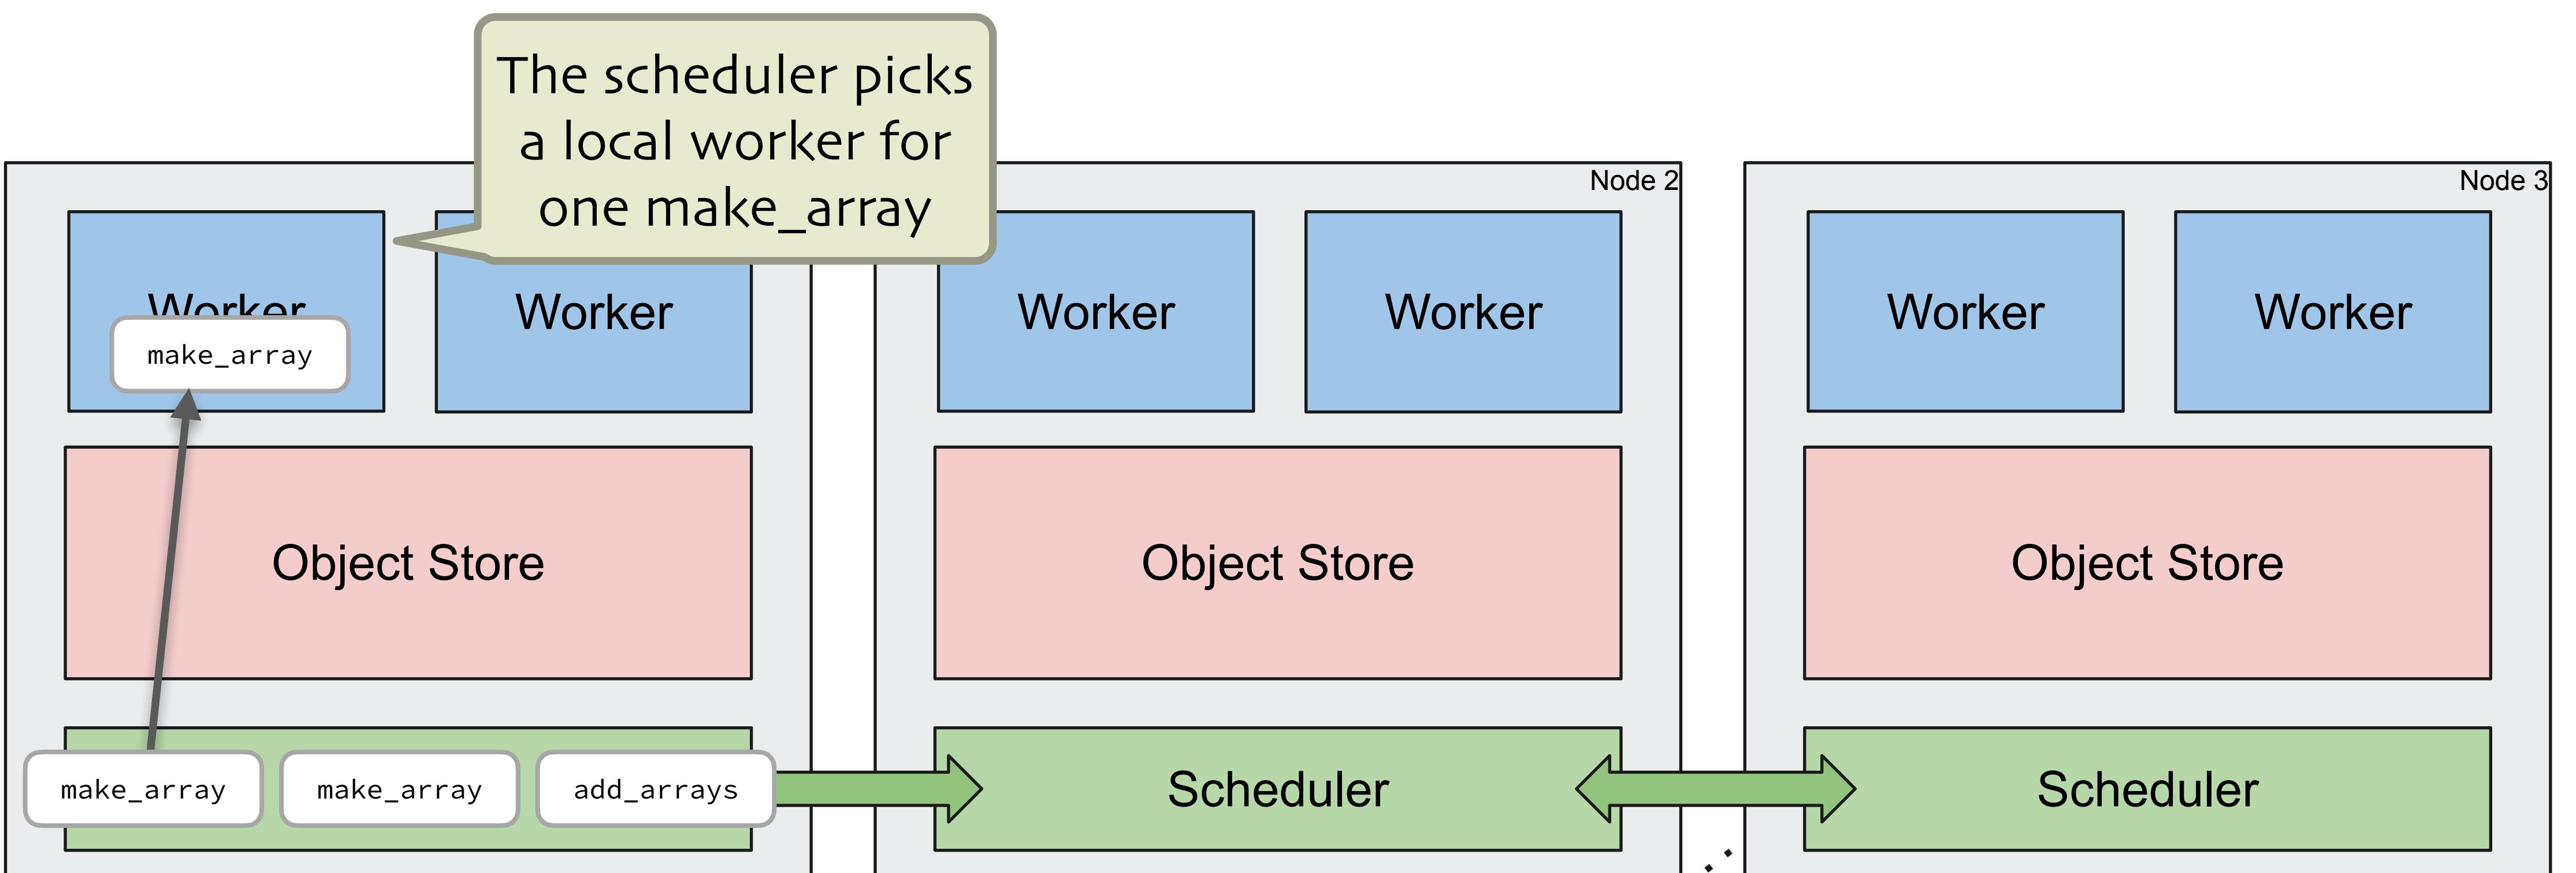
How Does This Work?

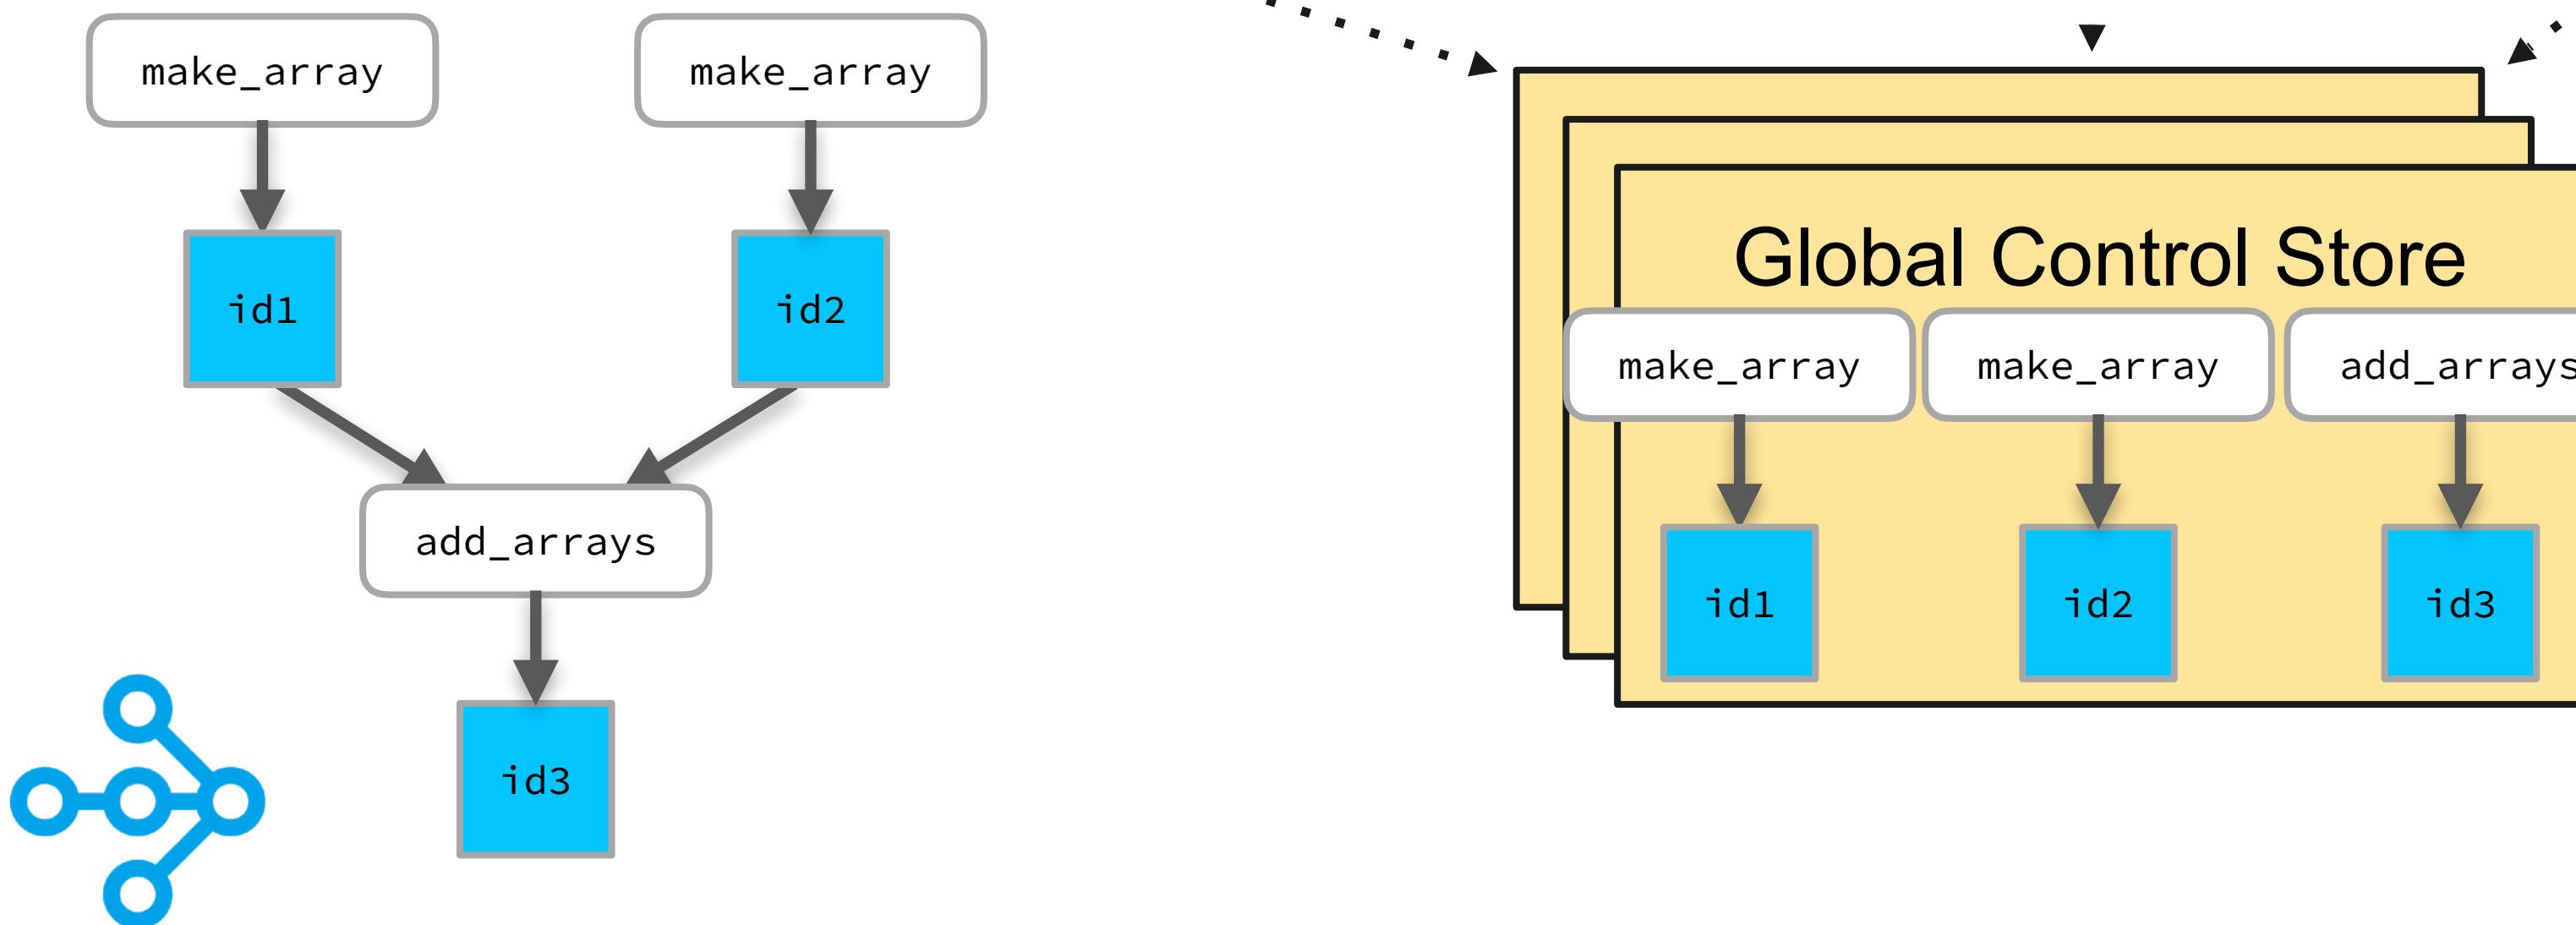
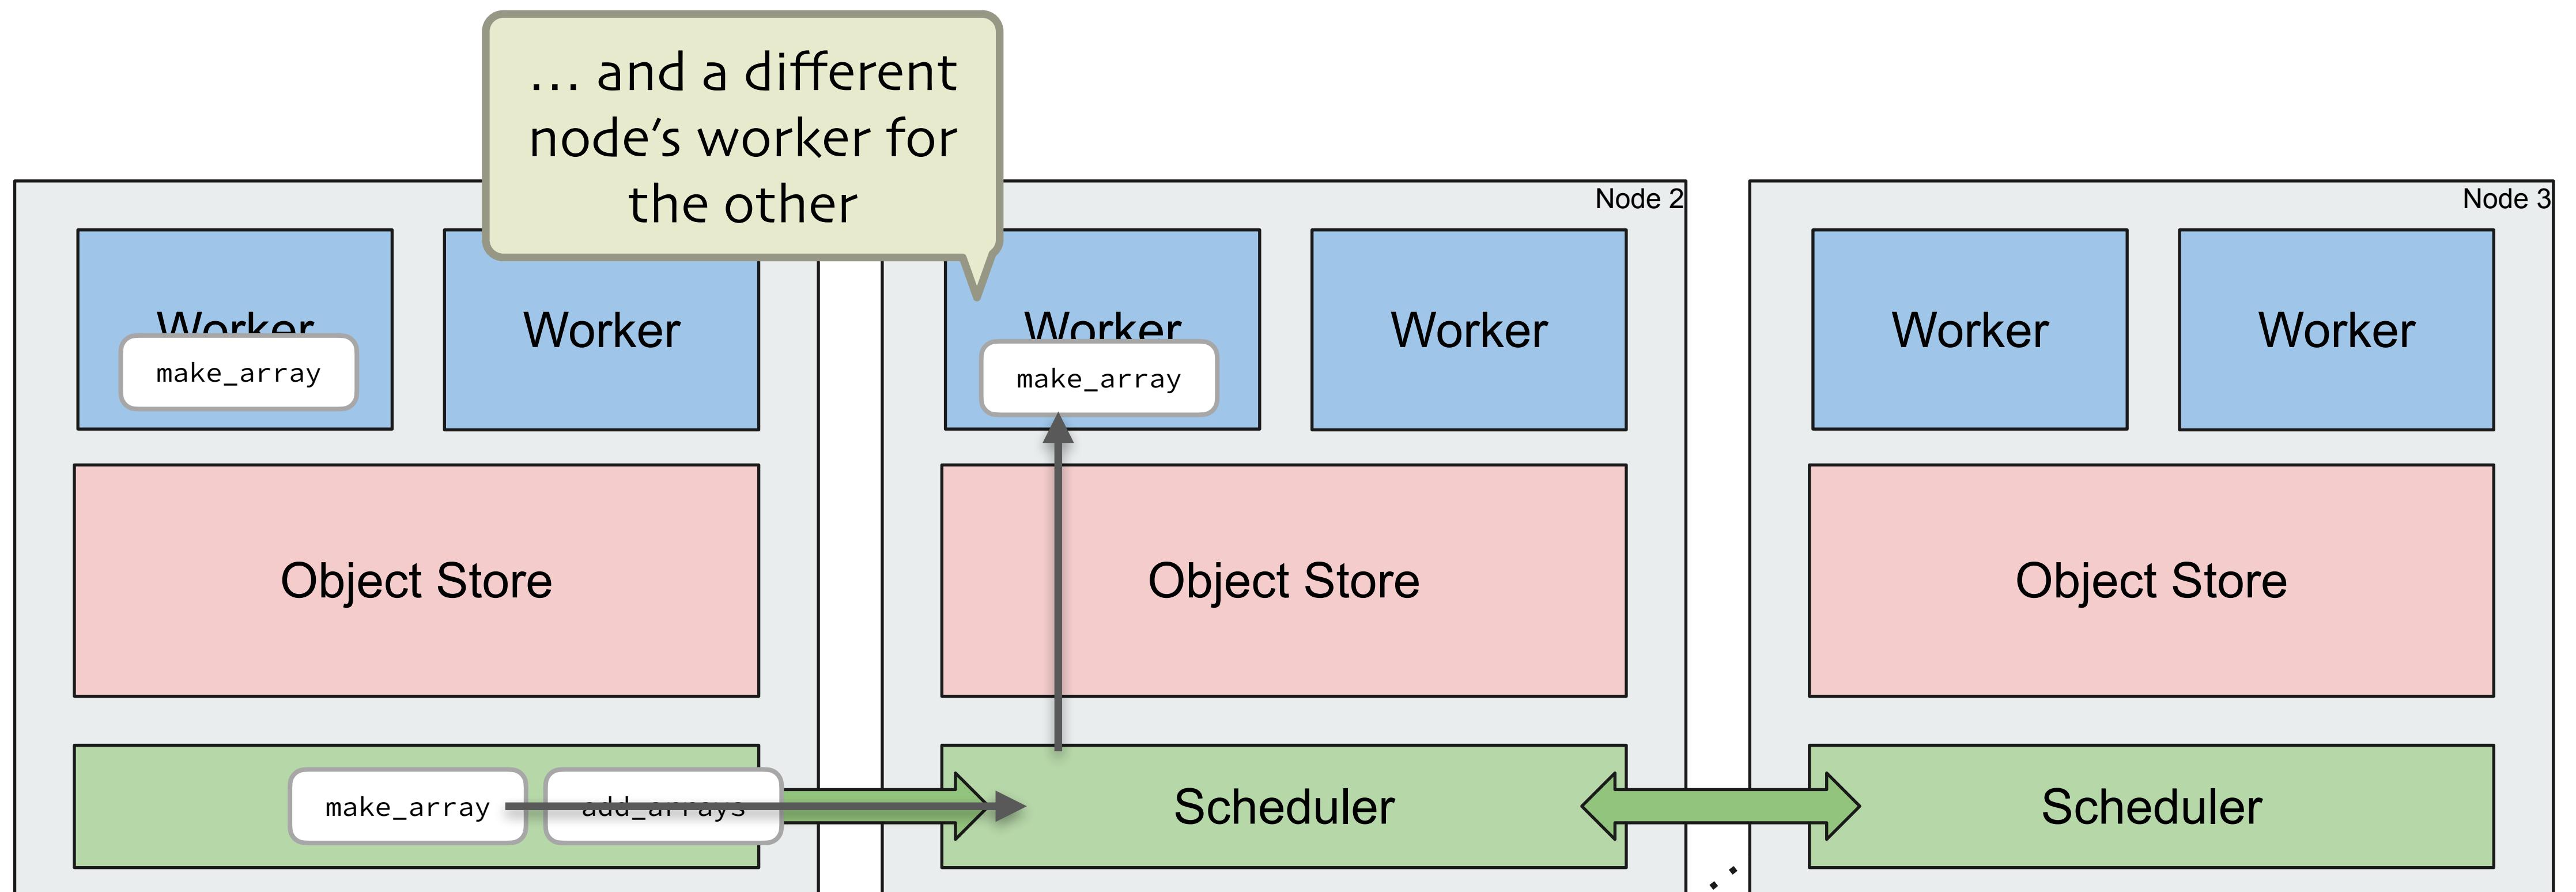






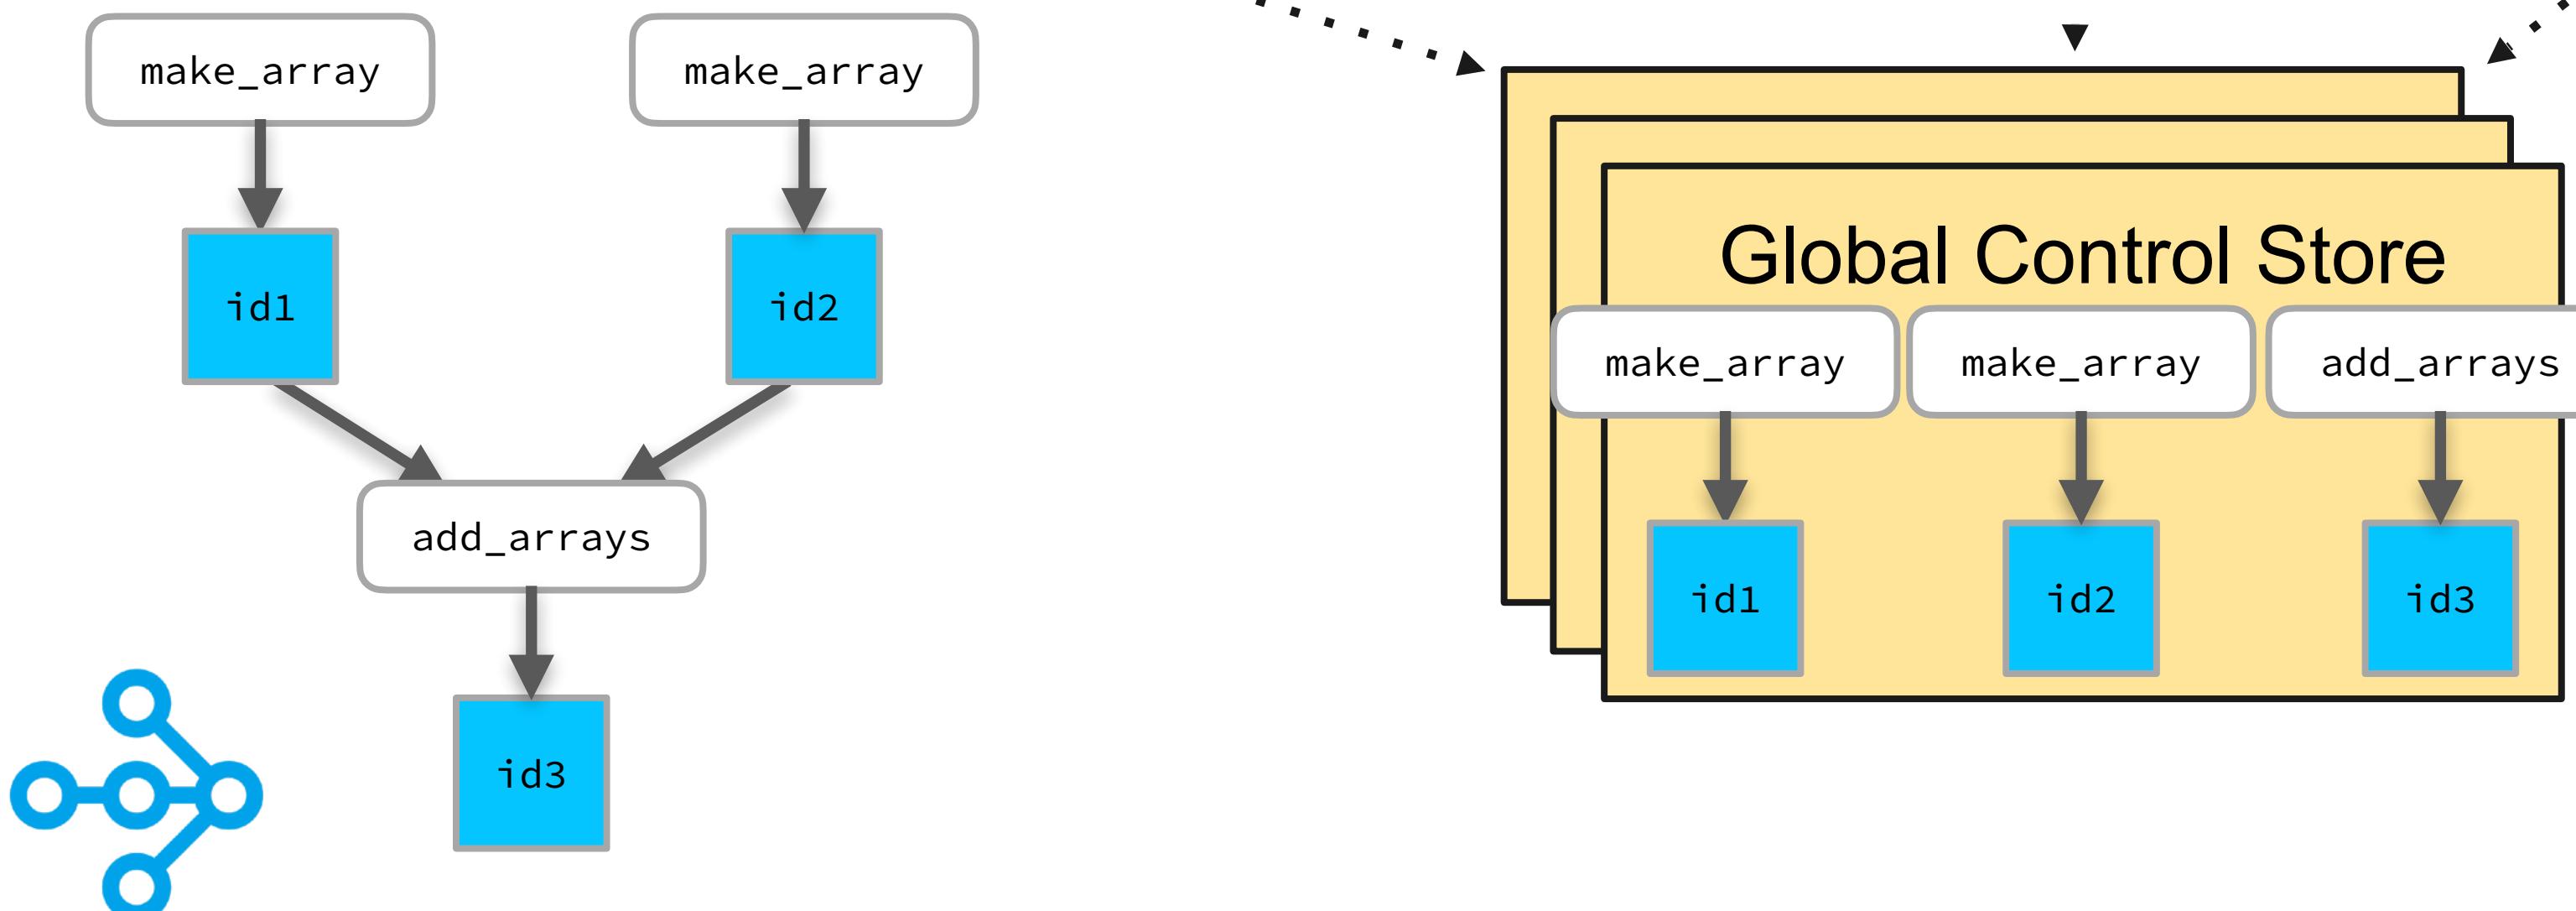
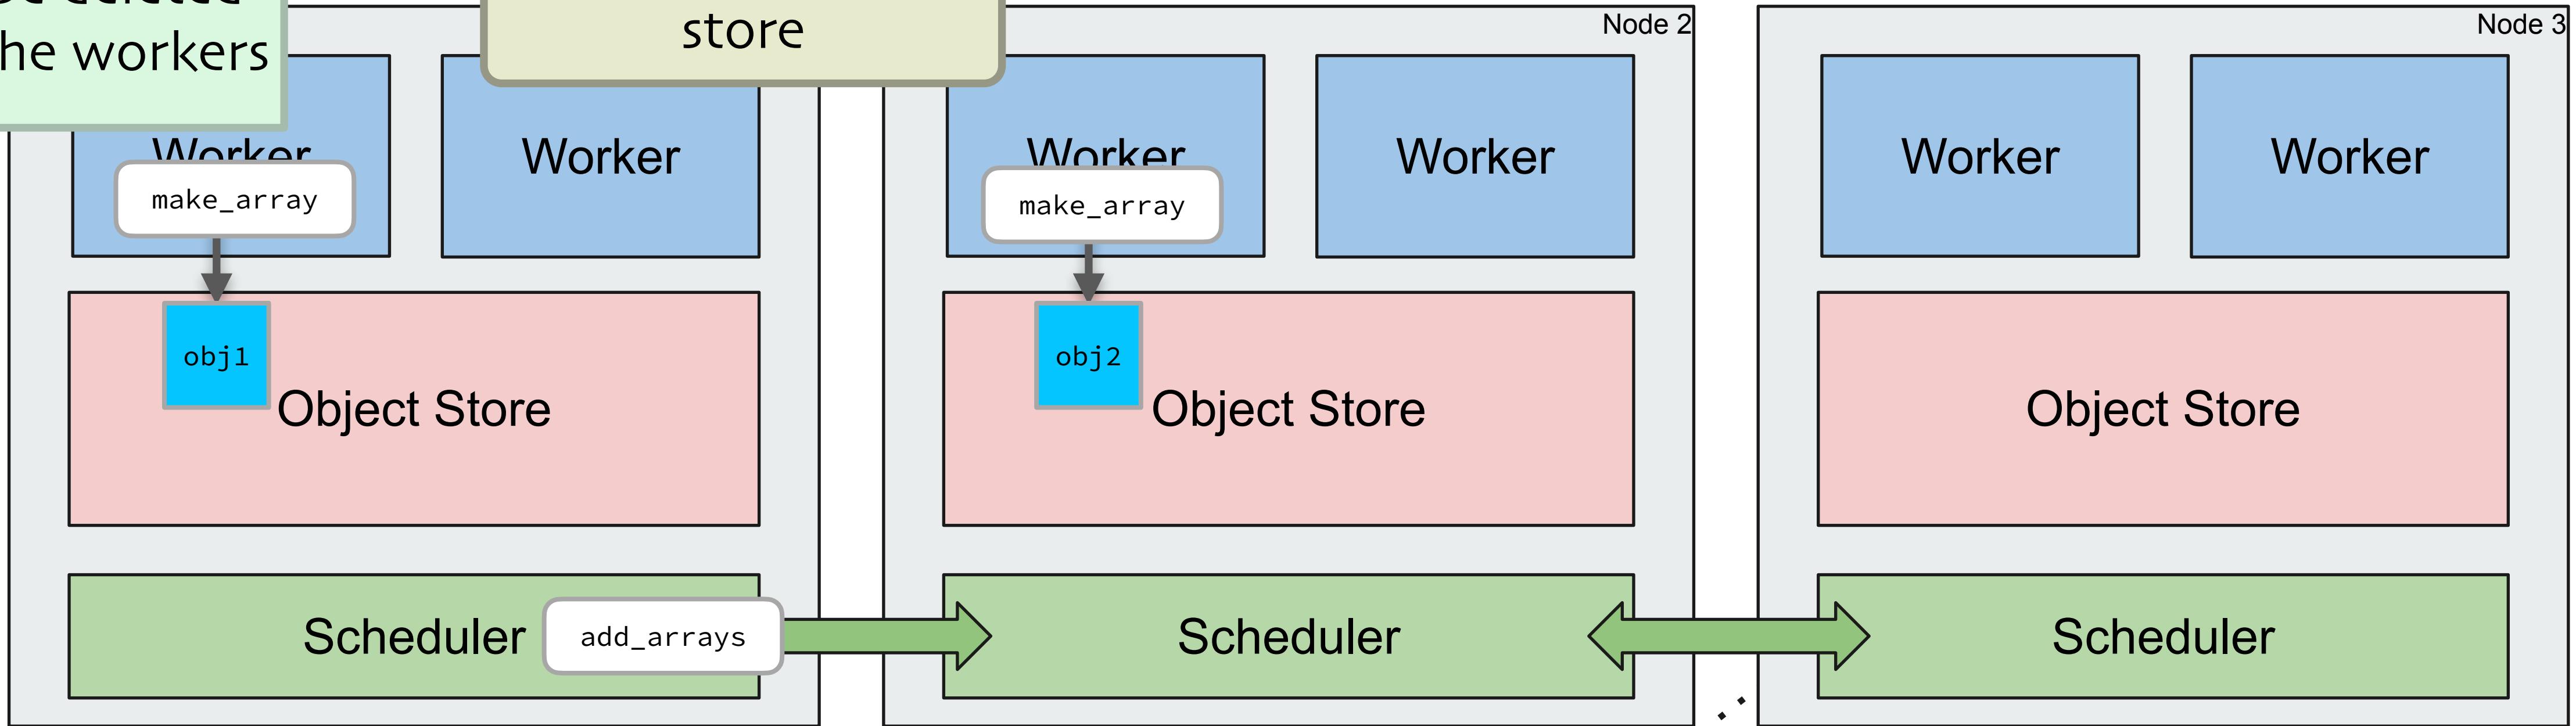


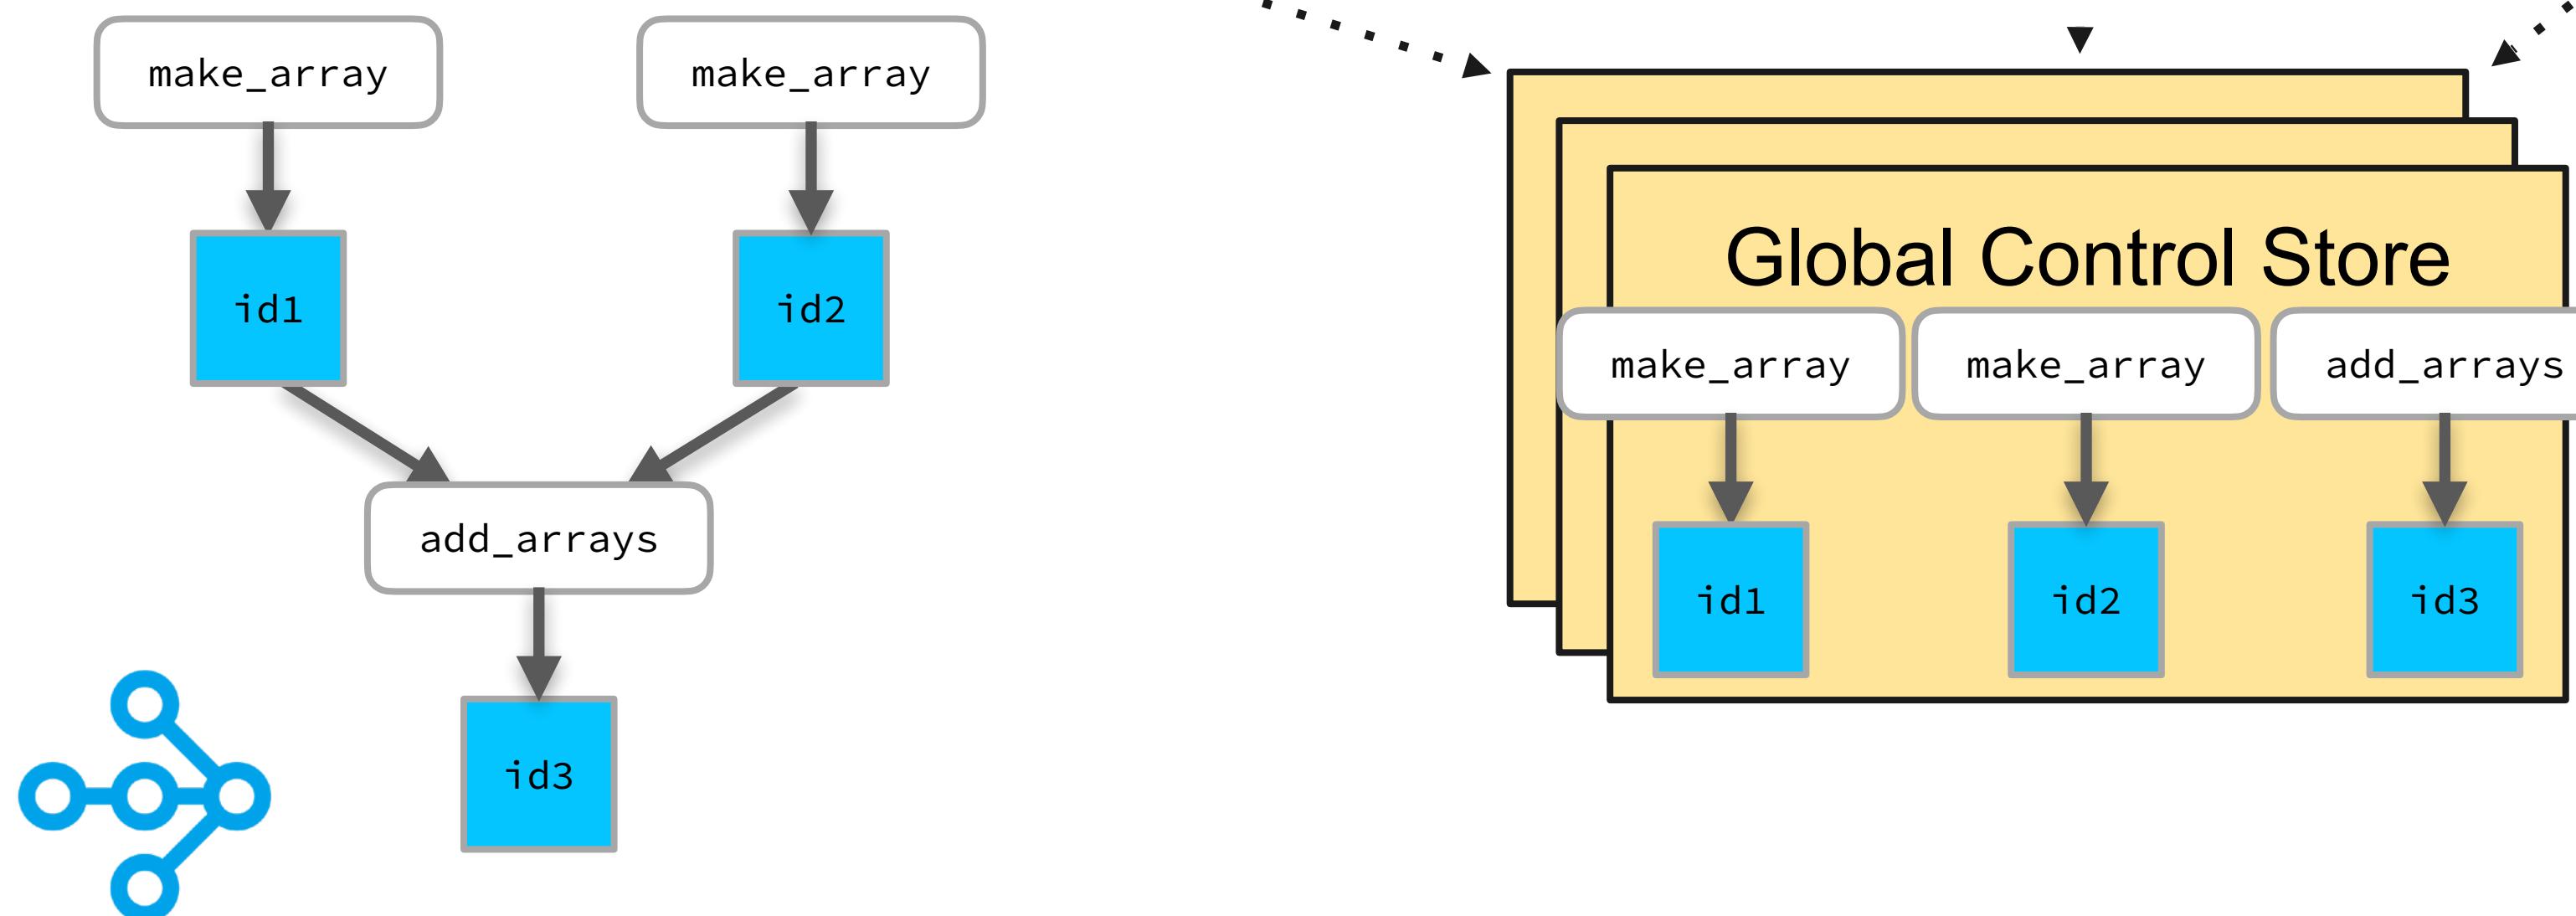
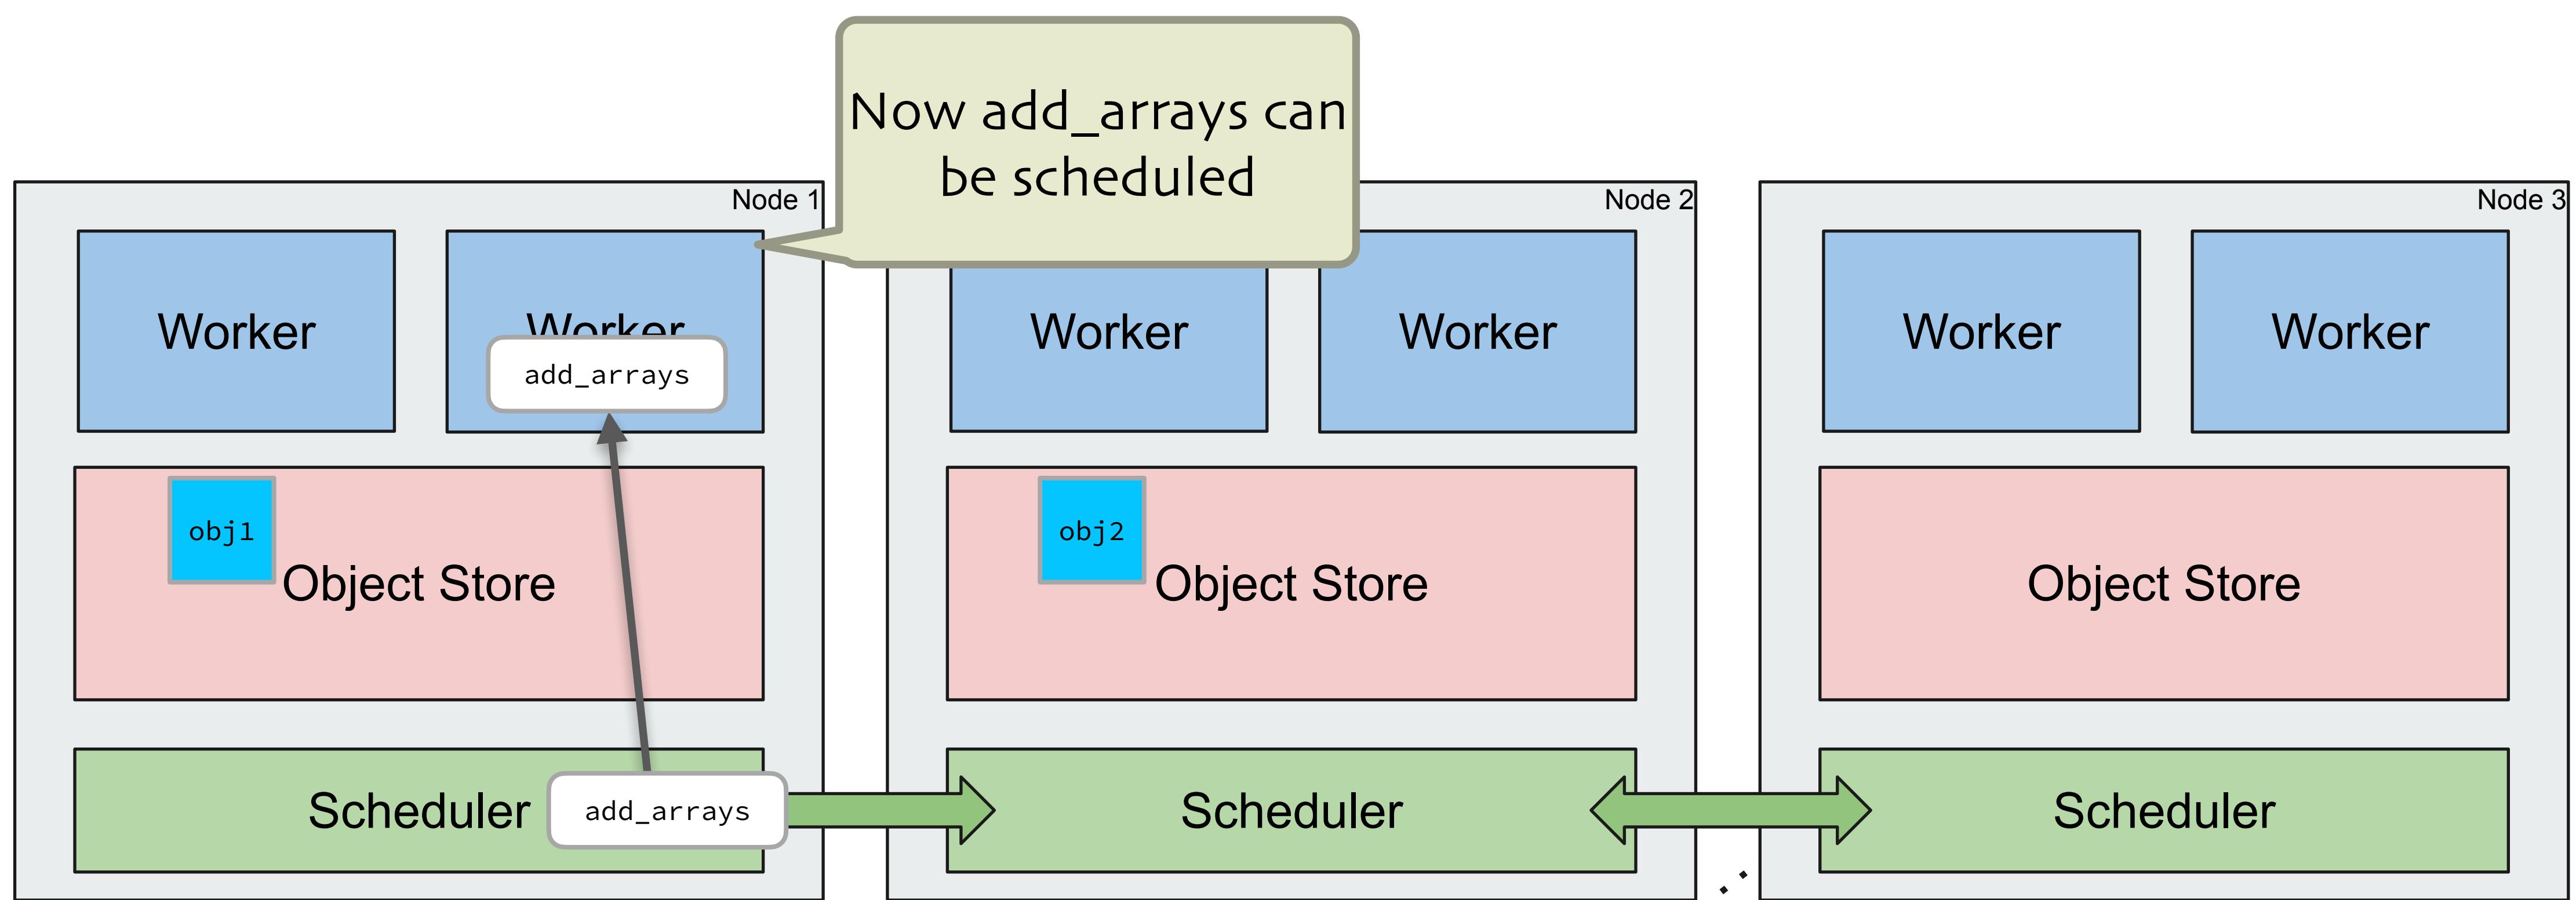




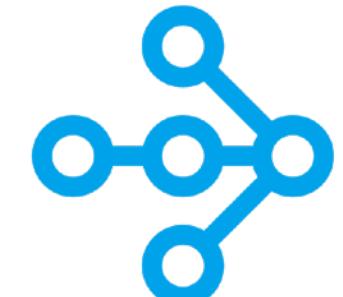
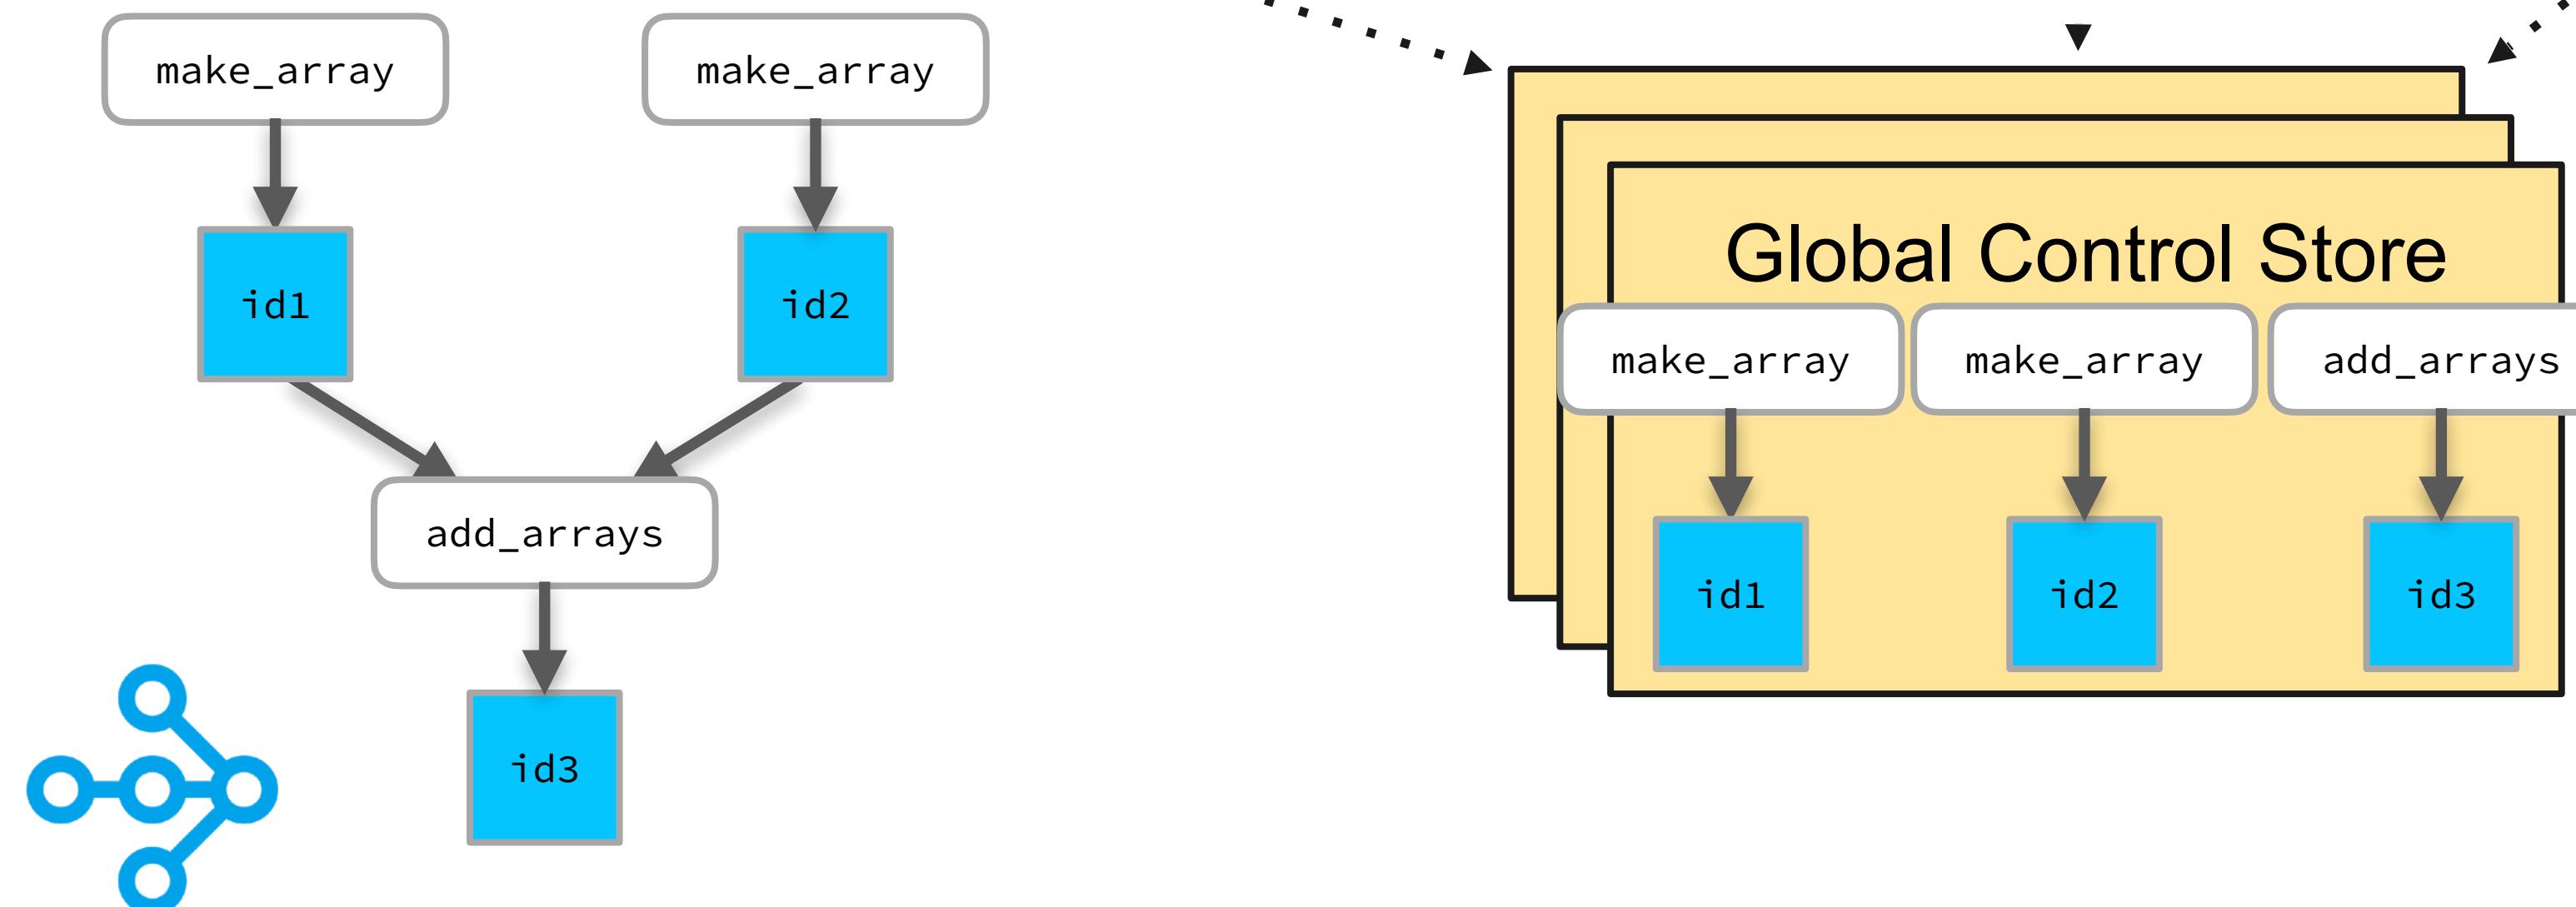
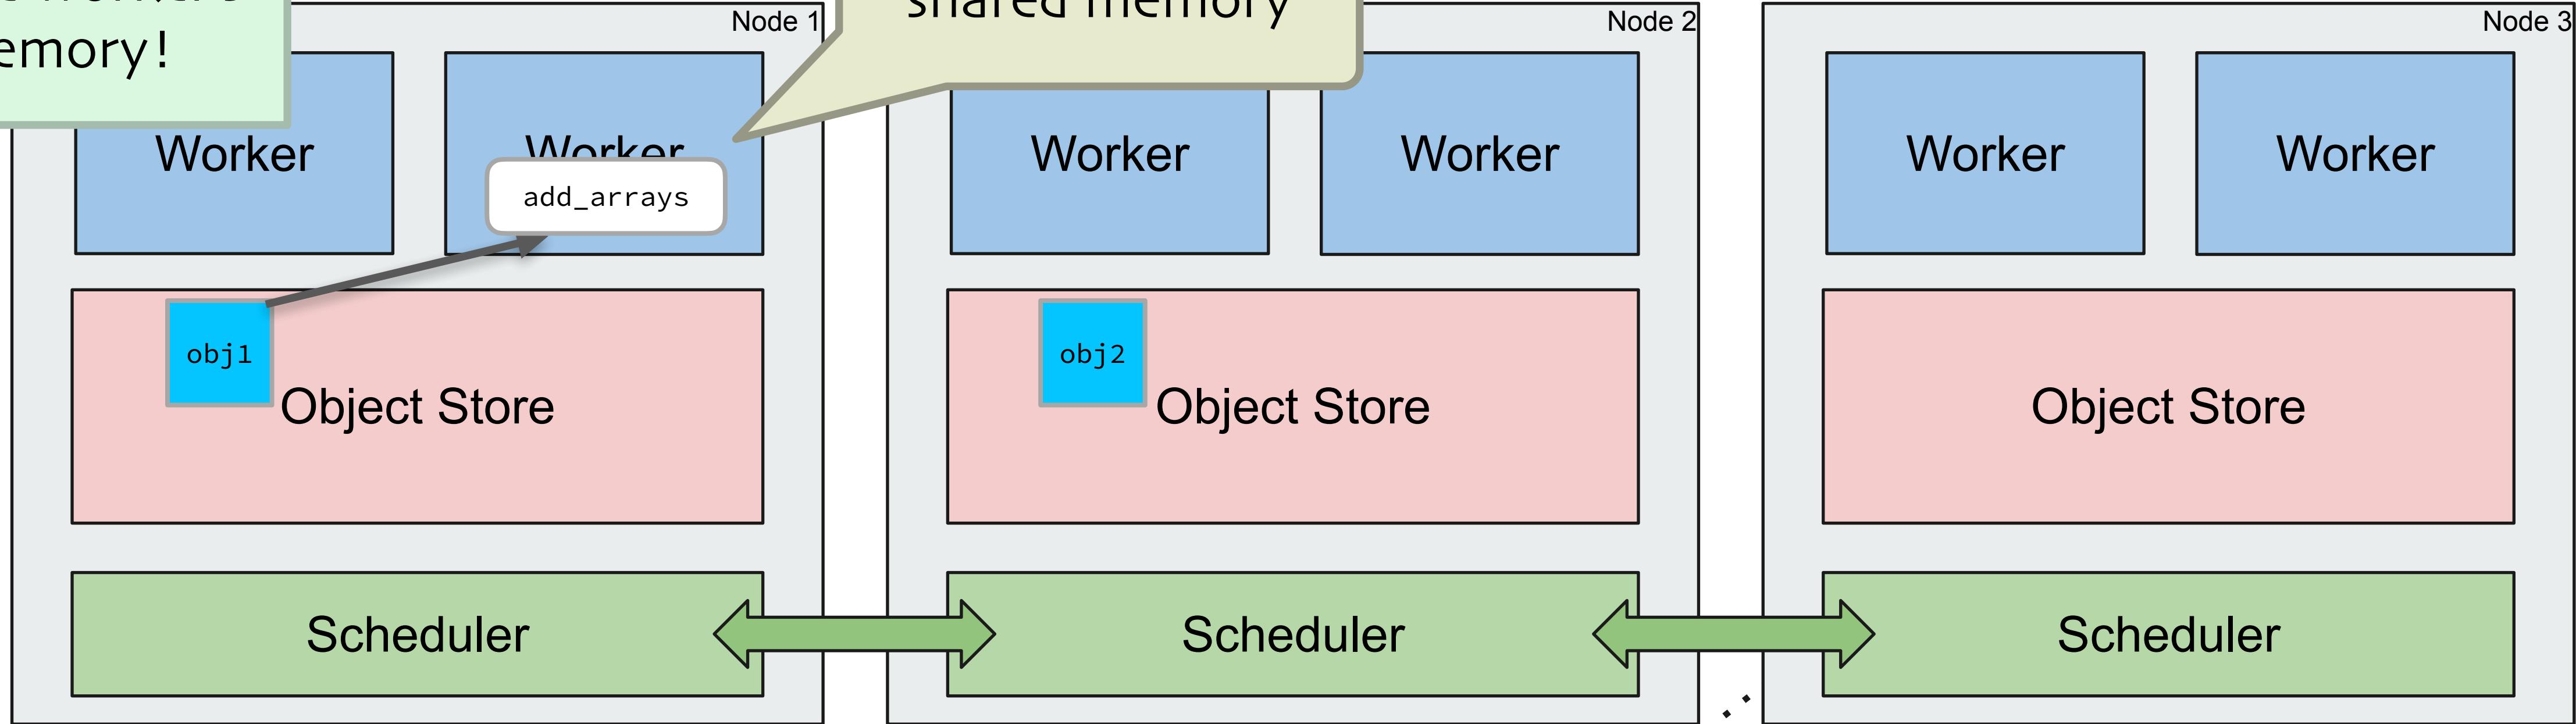
The tasks can now be deleted from the workers

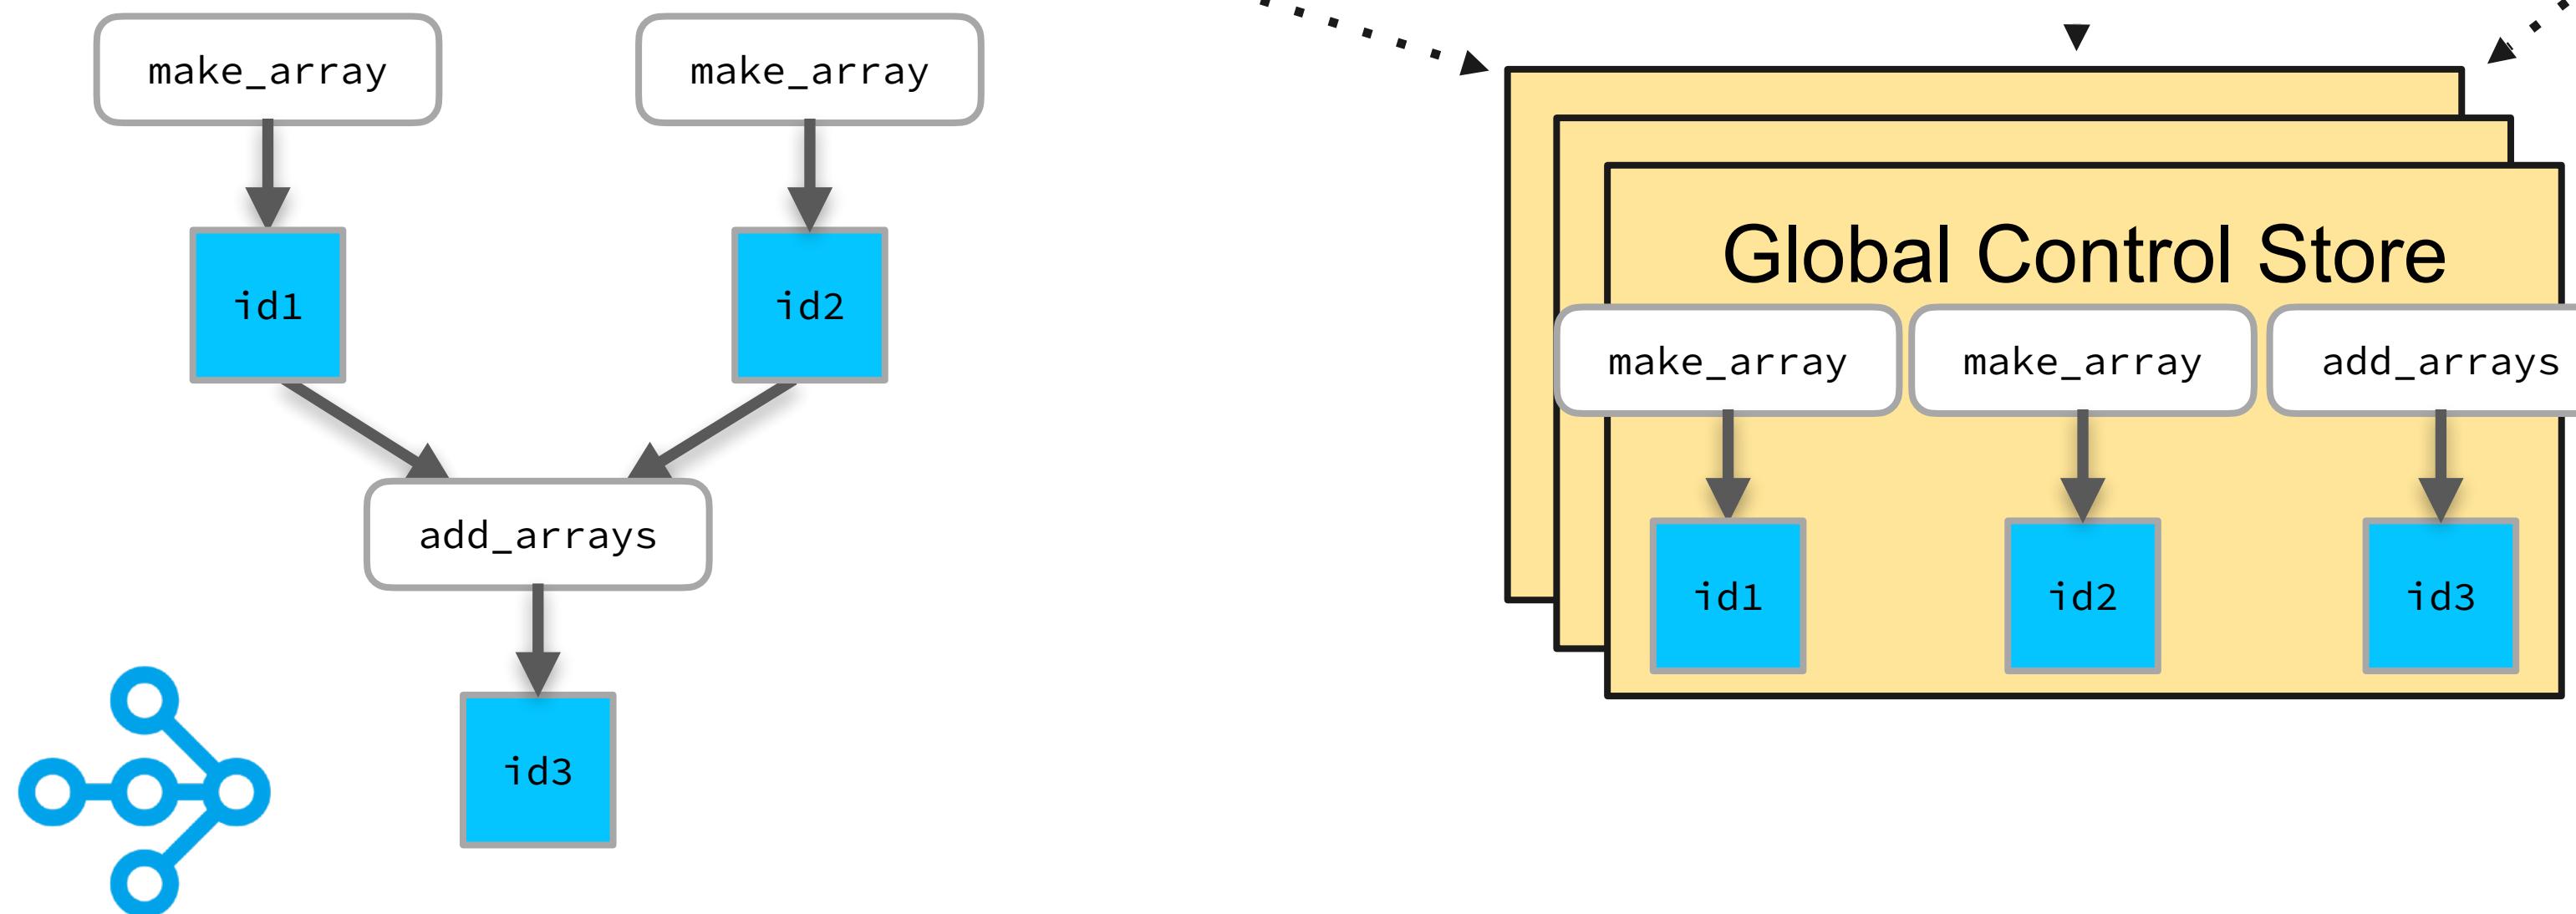
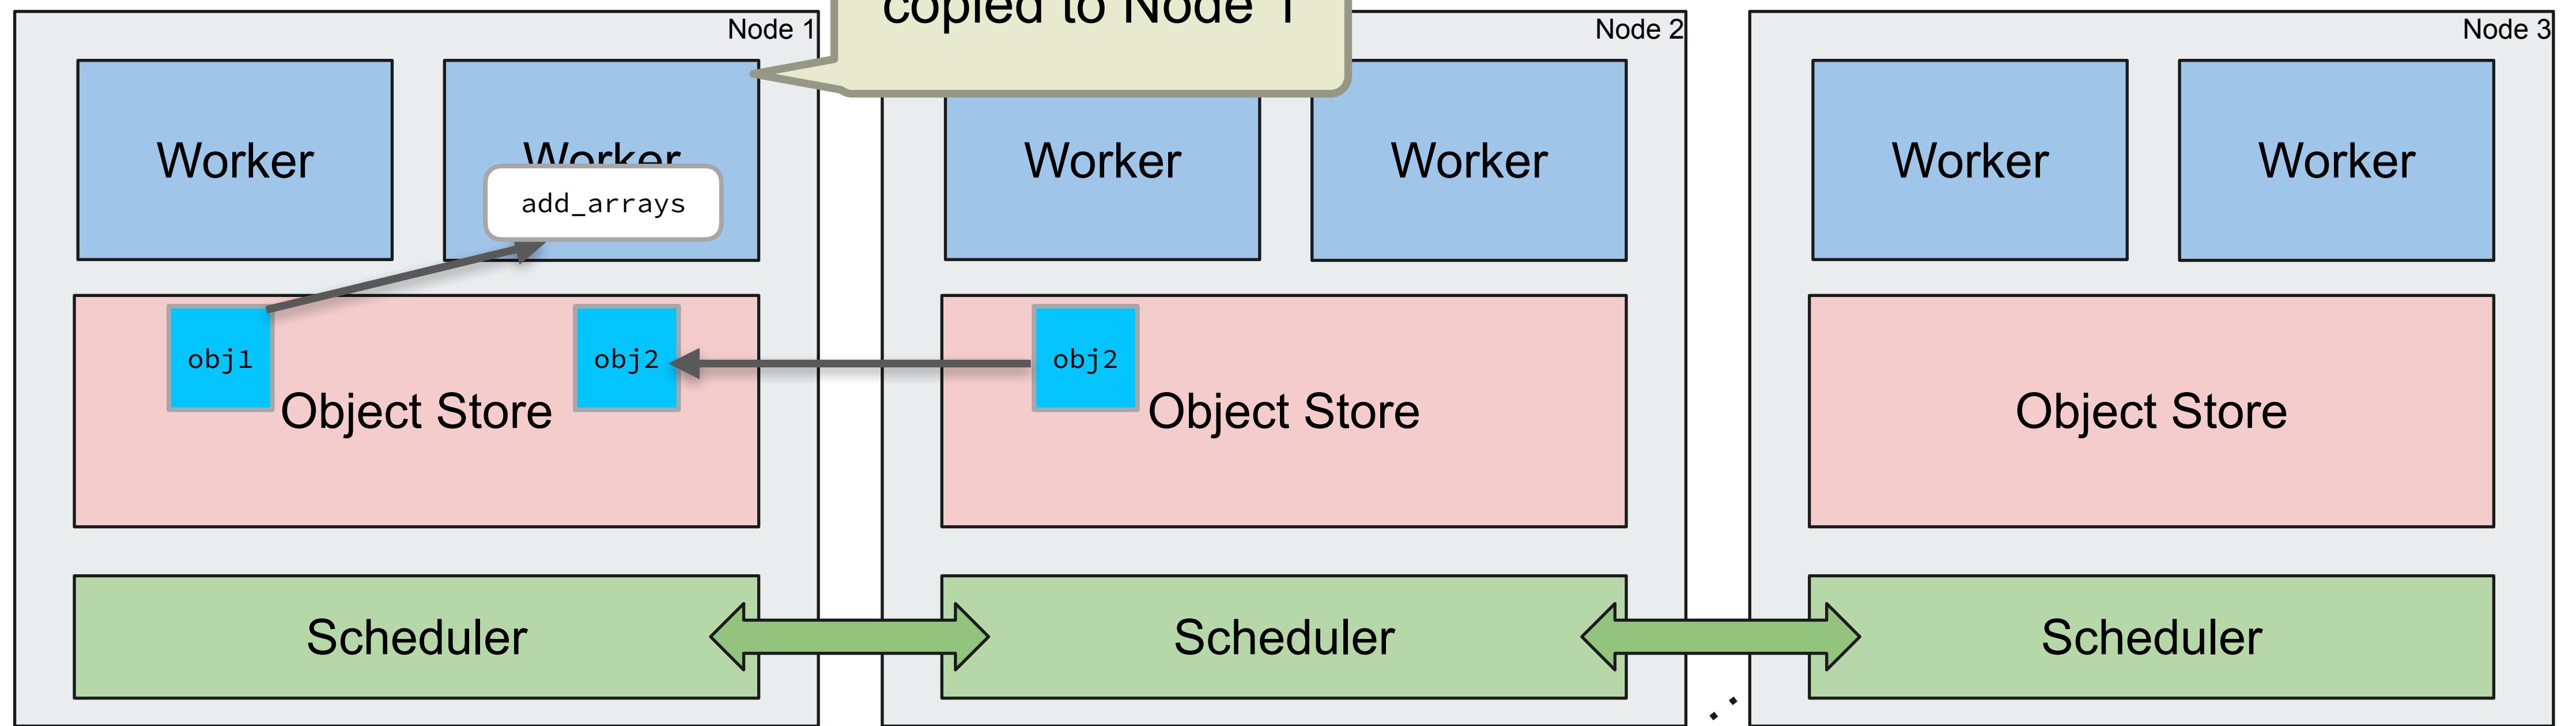
Each task writes its objects to the object store

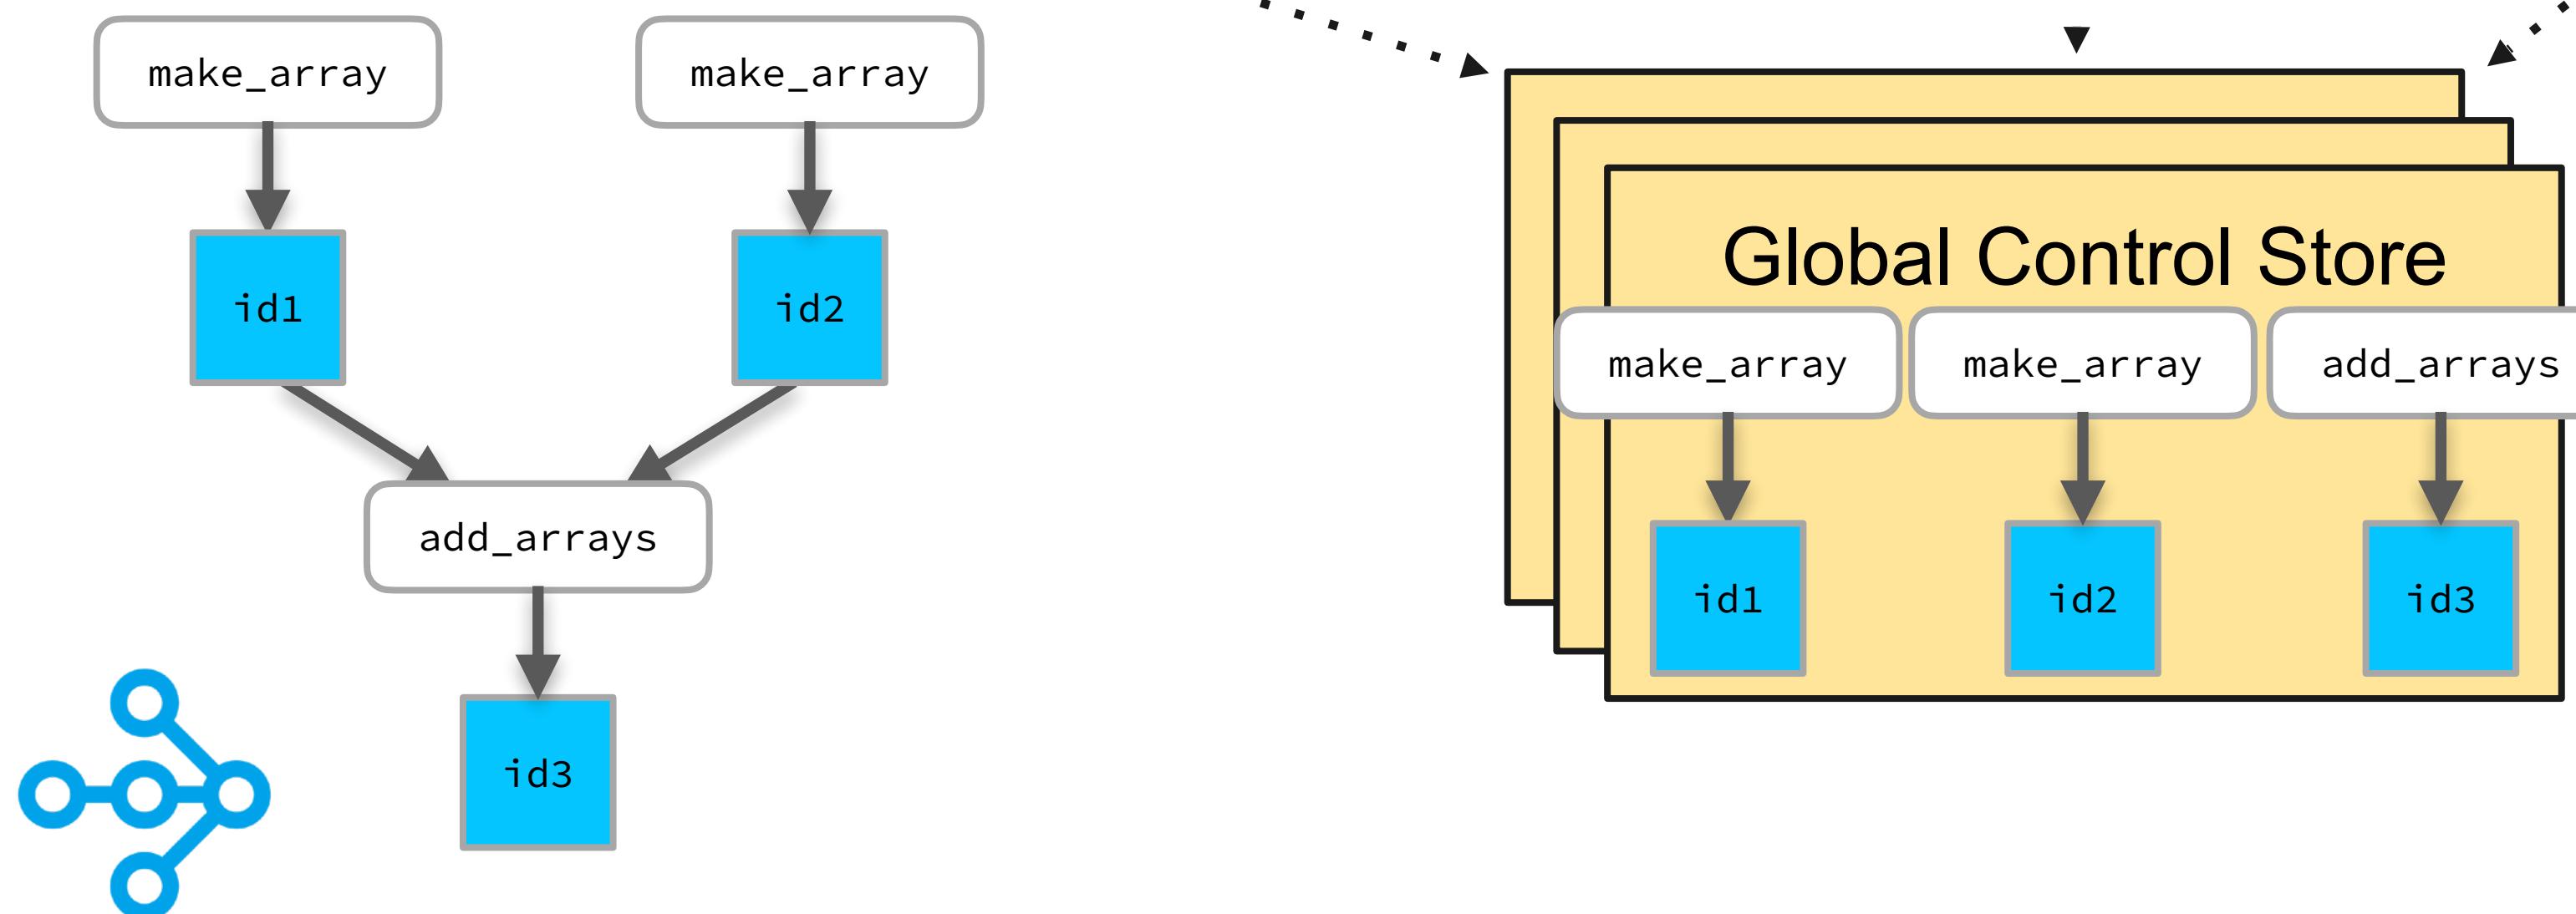
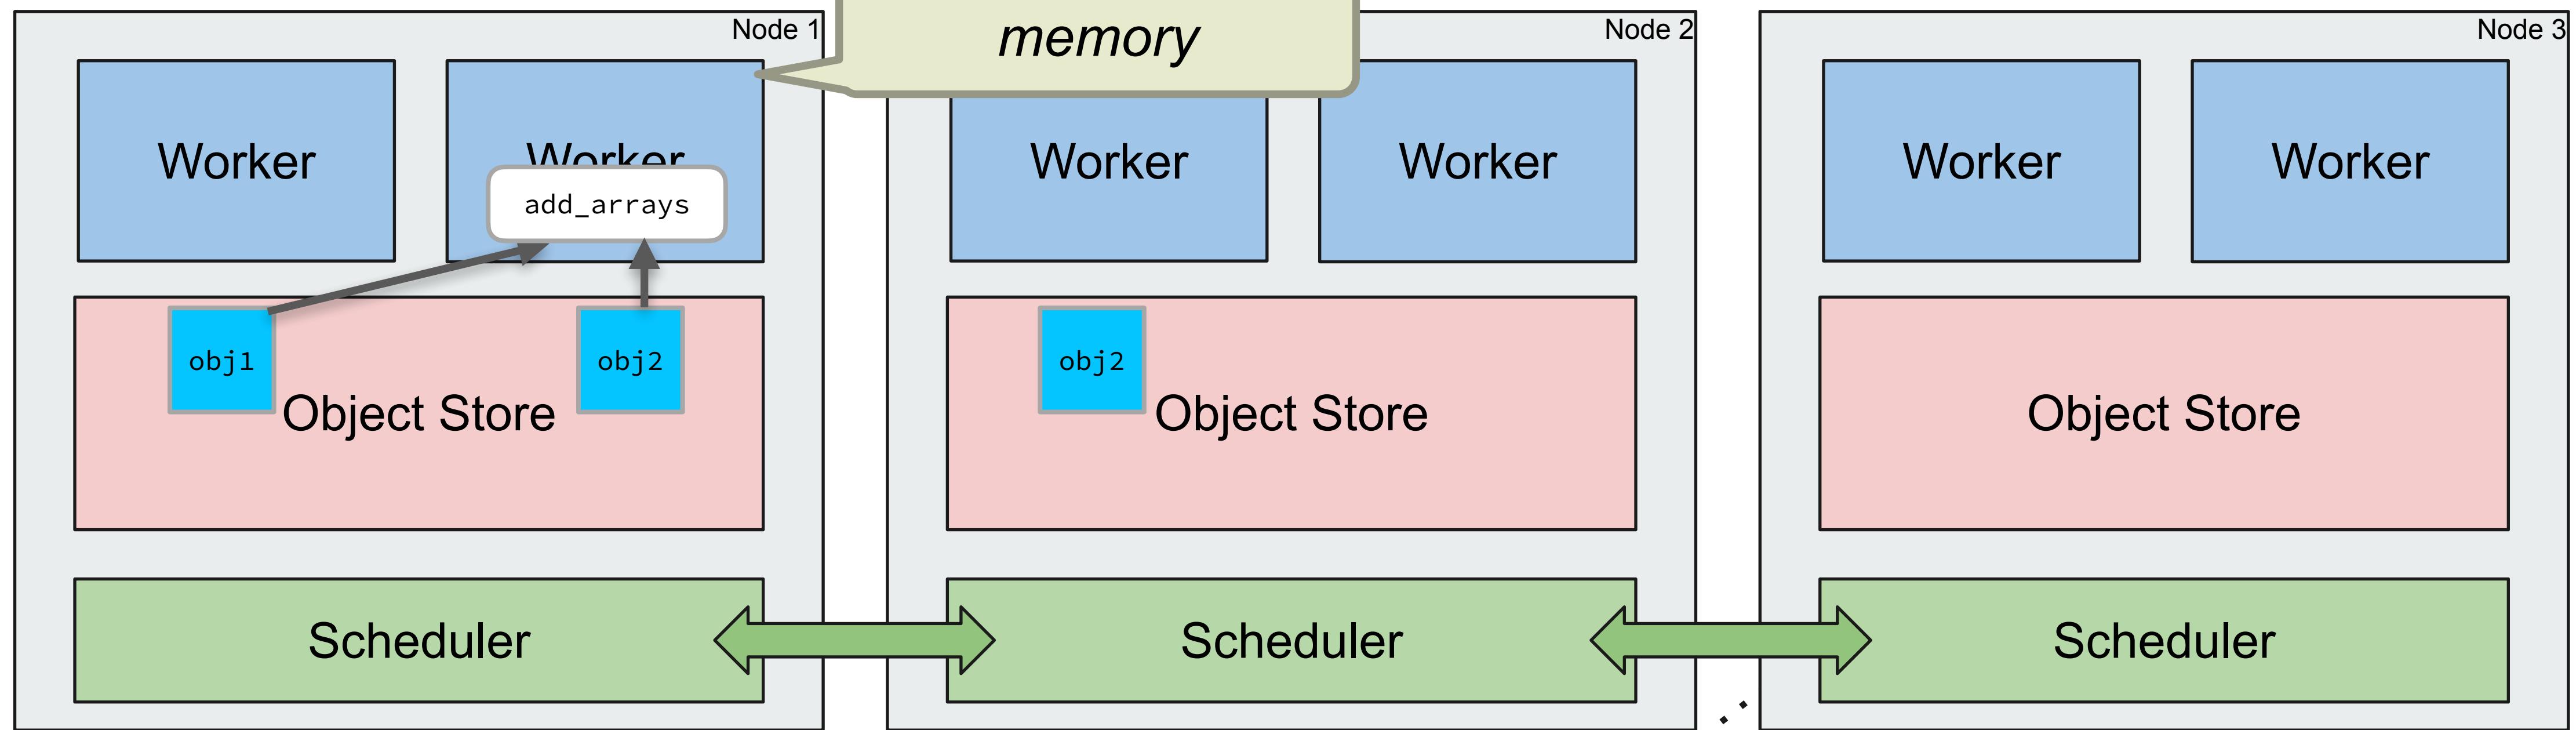


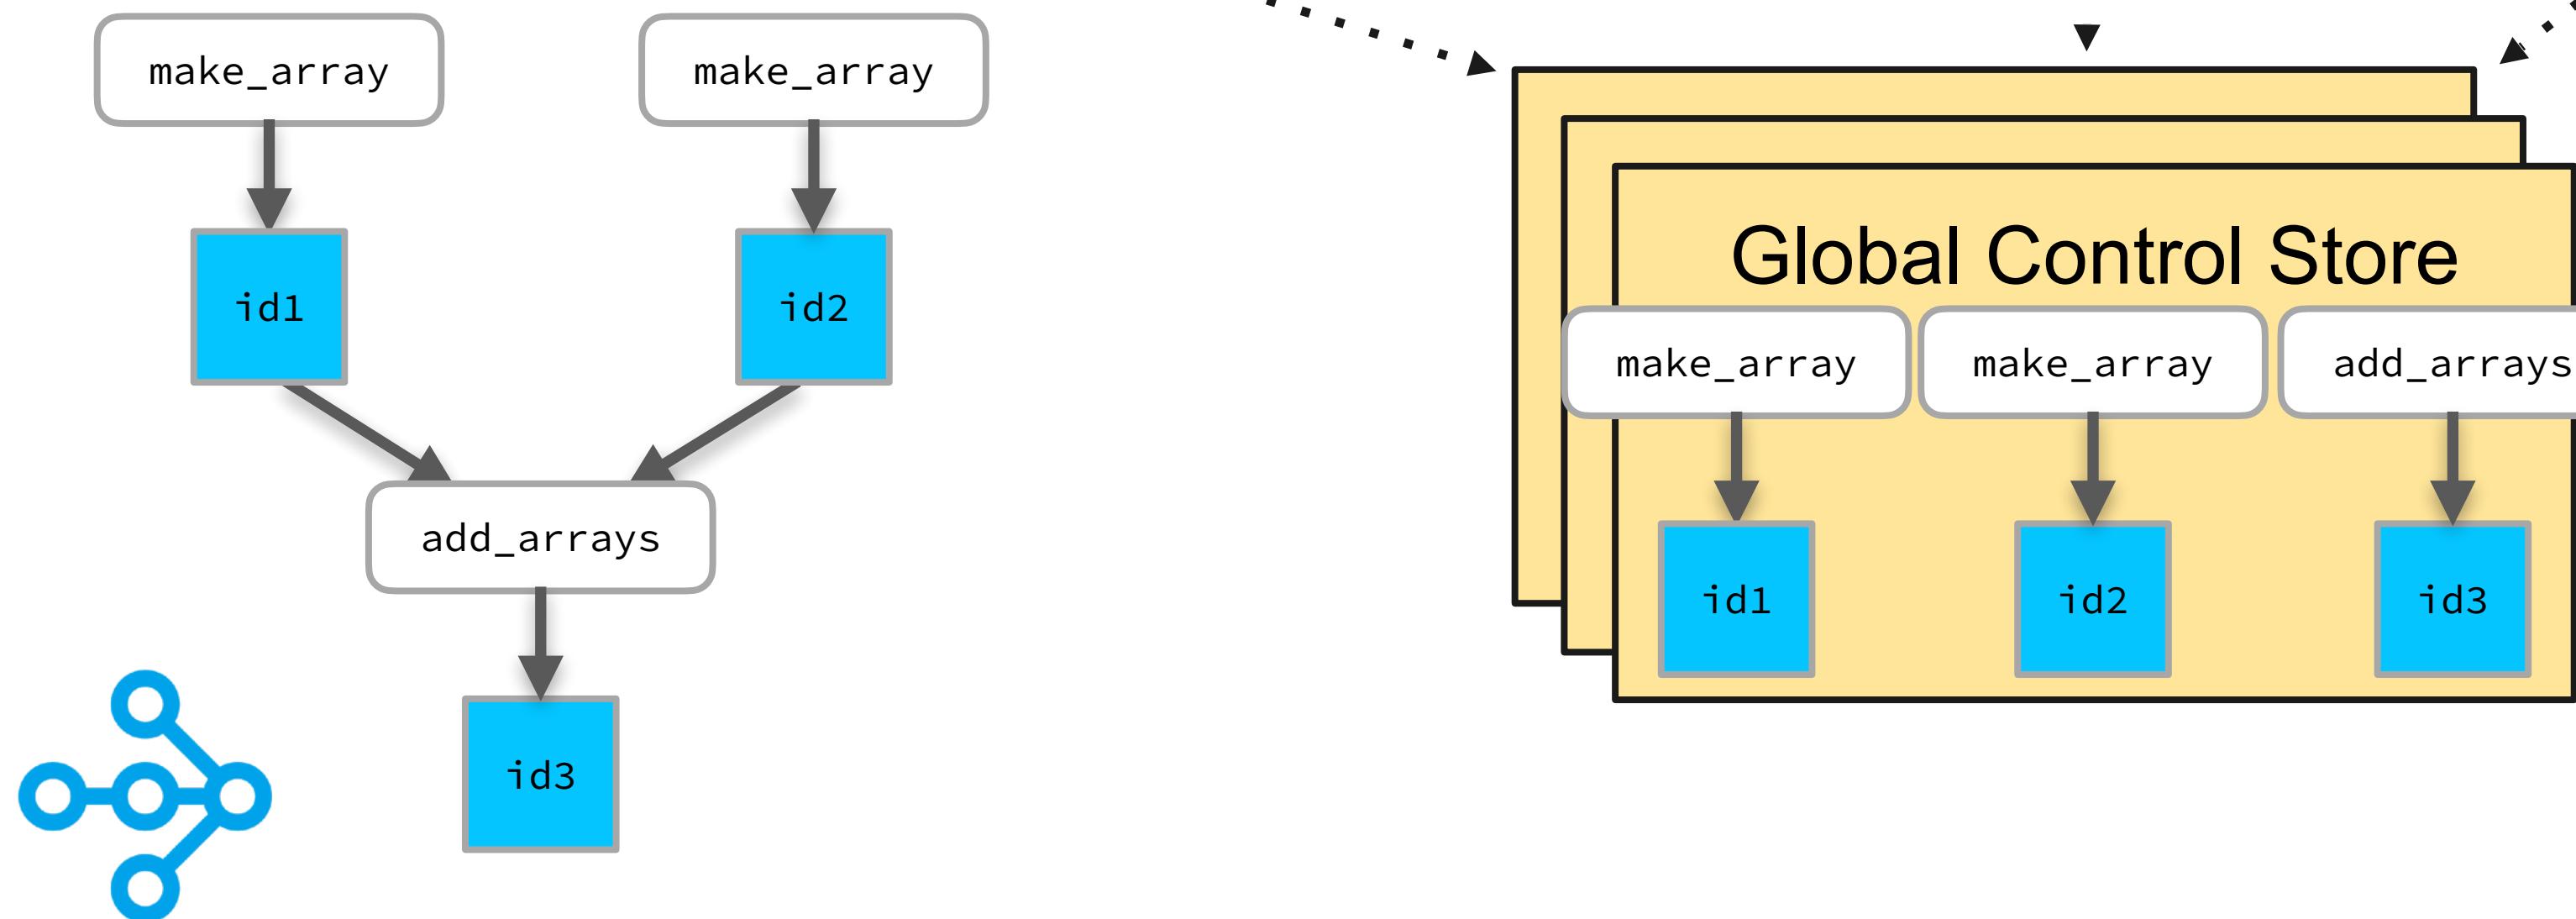
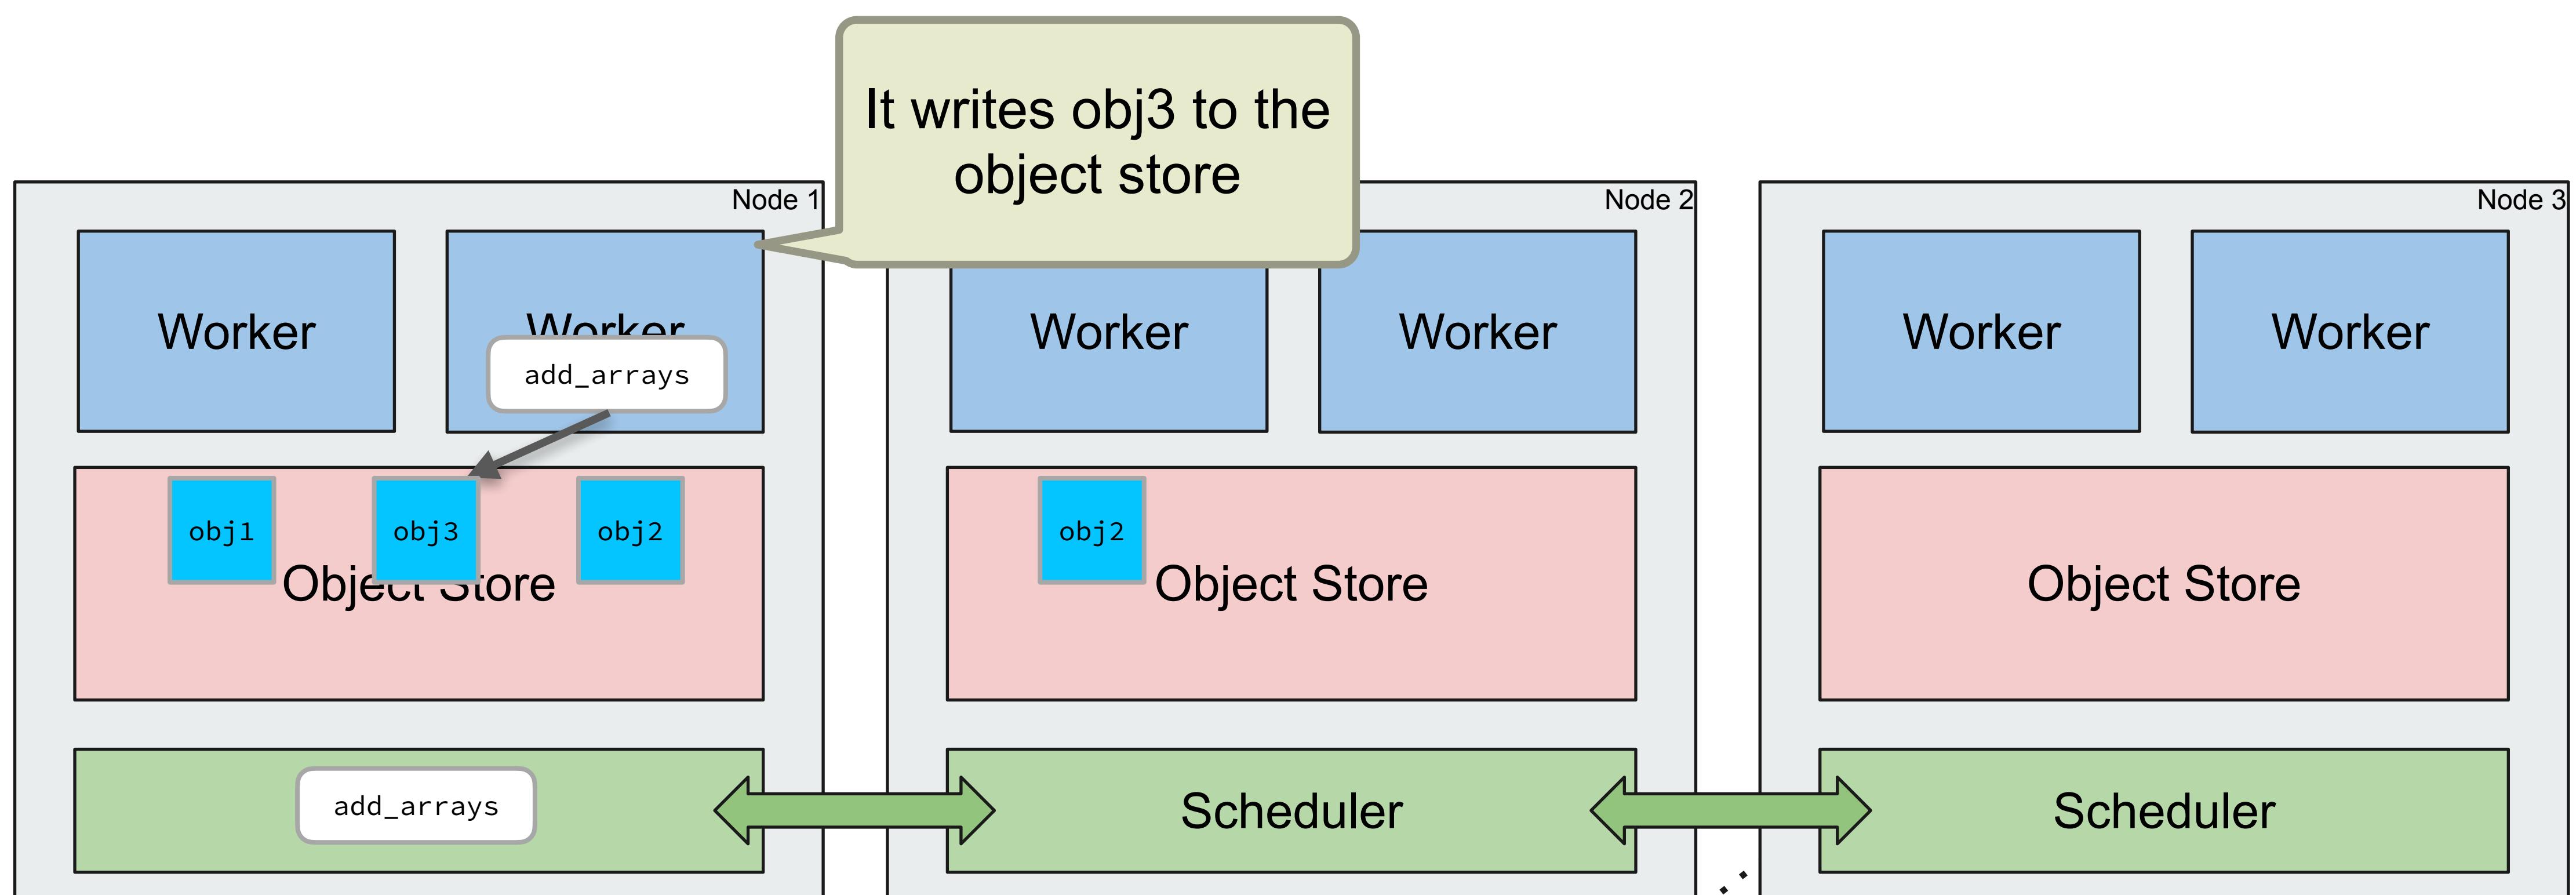


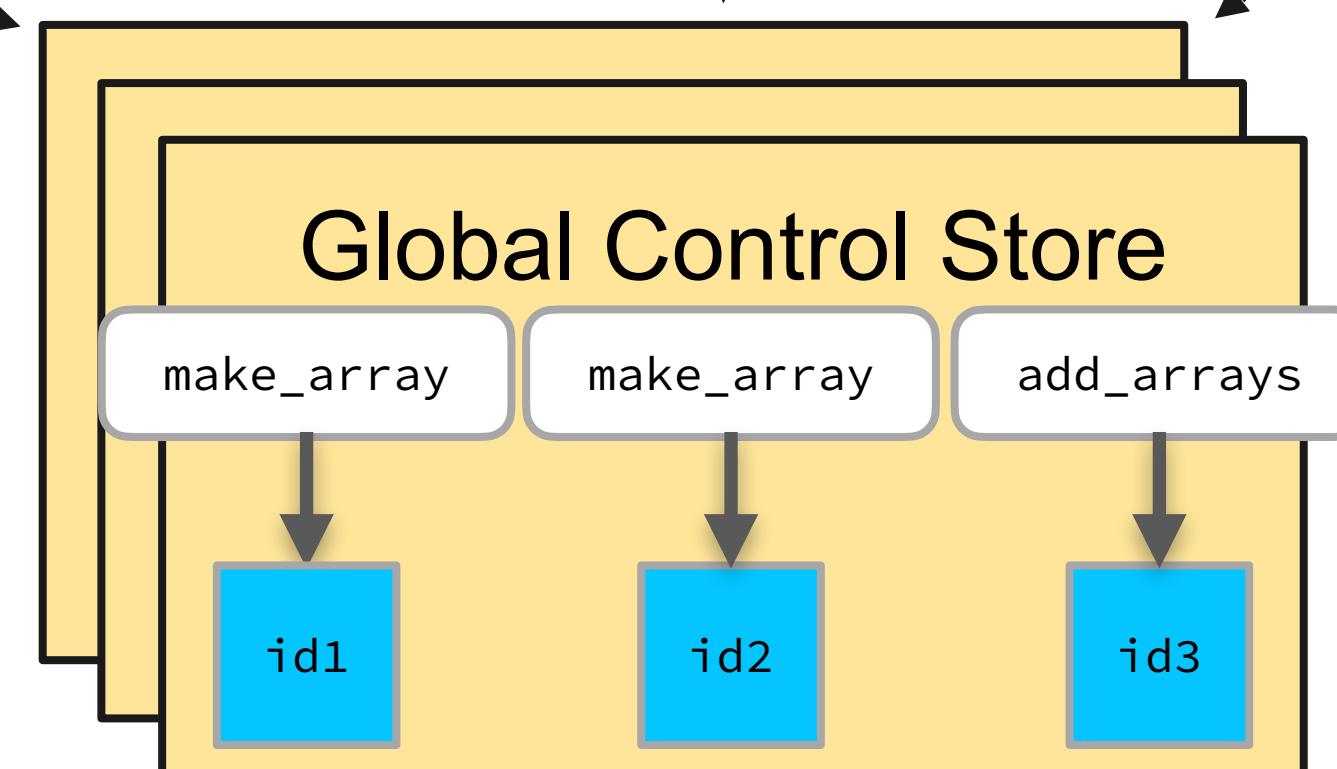
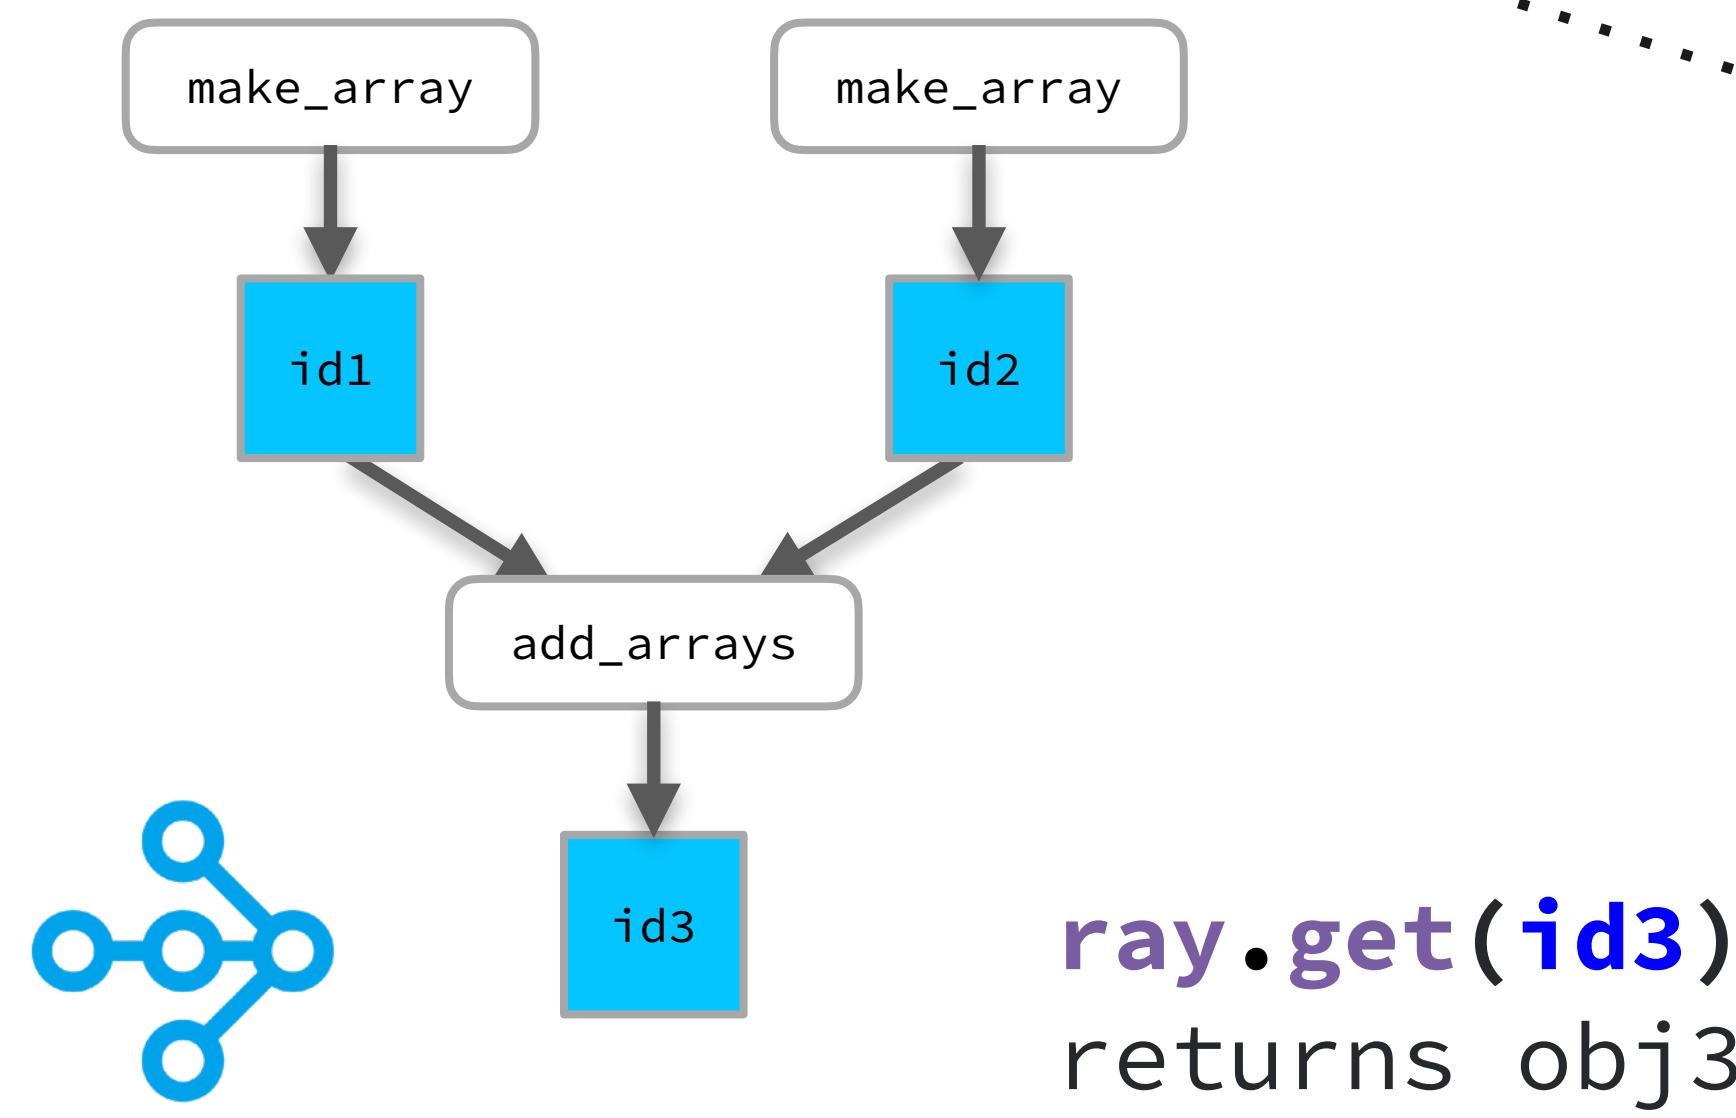
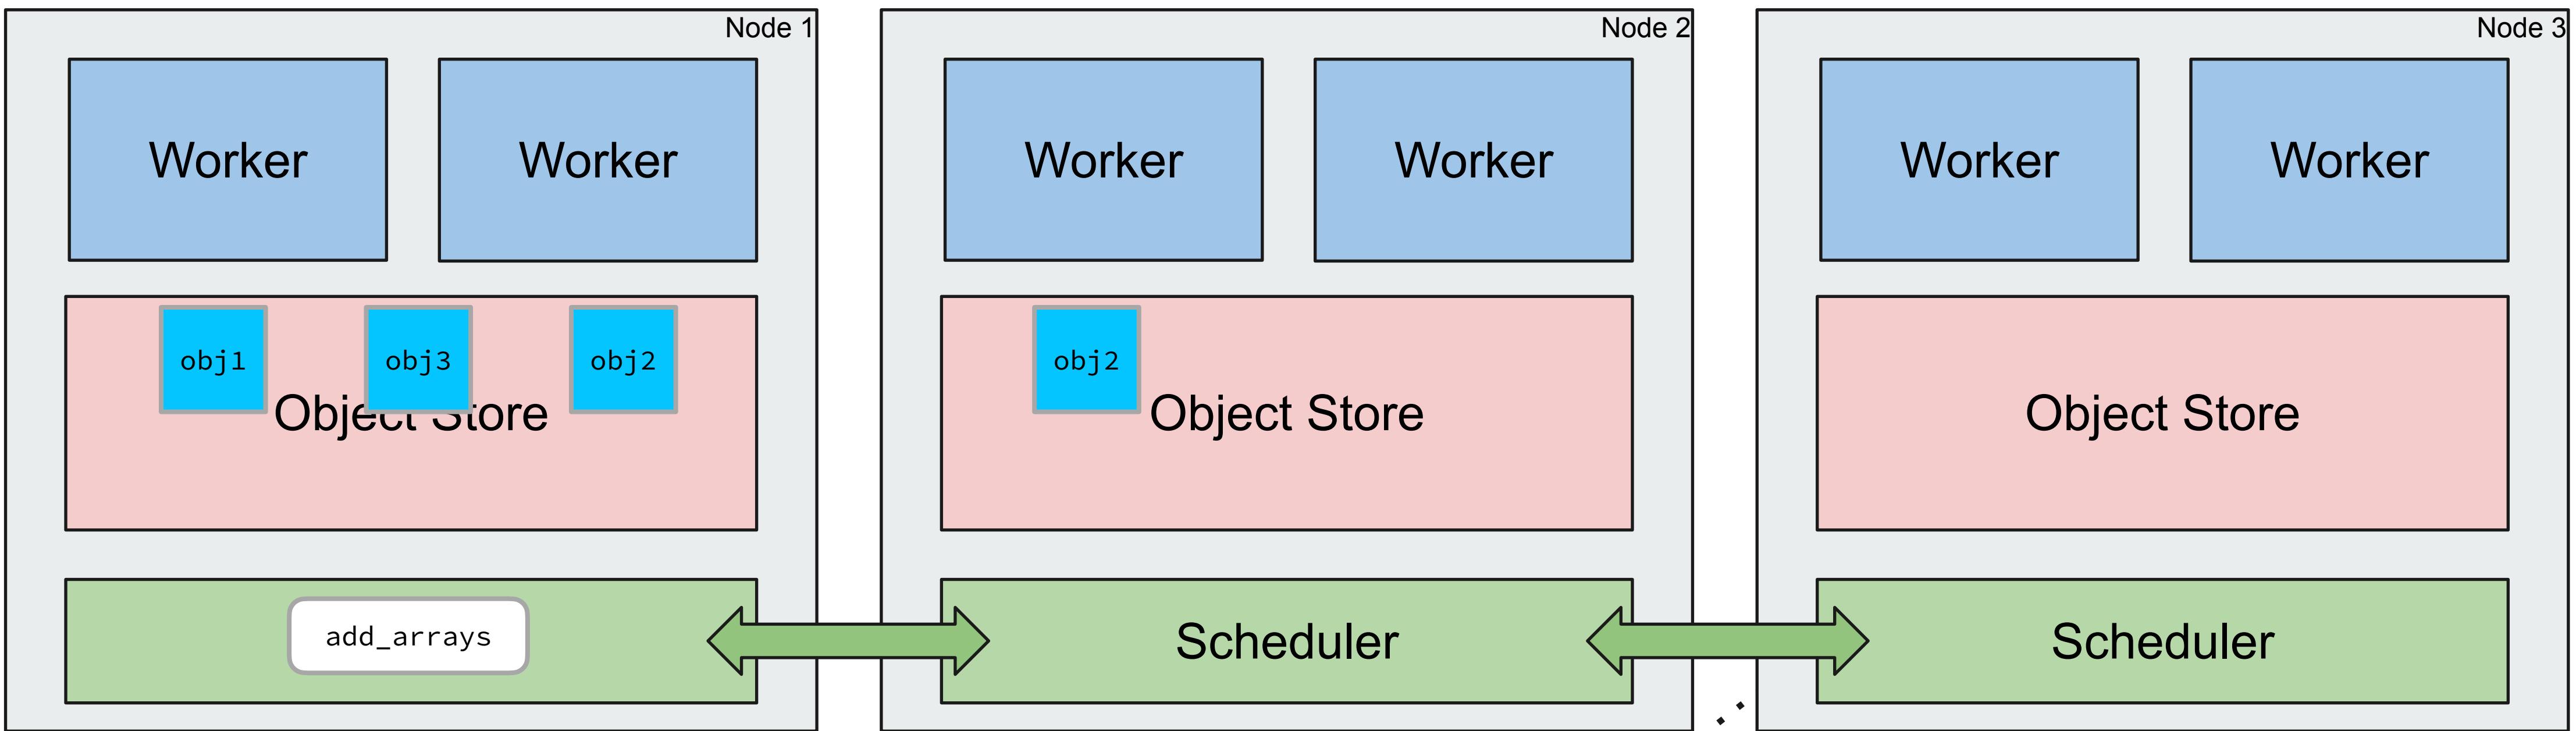
No need to copy
to the worker's
memory!











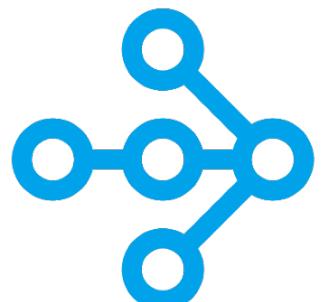
And we're done!



Ray Community

Community and Resources

- ray.io
- ray.readthedocs.io/en/latest/
- Tutorials (free): [Anyscale Academy](#)
- github.com/ray-project/ray.git
- Need help?
 - Ray Slack: [ray-distributed.slack.com](#)
 - [ray-dev](#) group





Migrating to Ray

If you're already using...

- `asyncio`
- `joblib`
- `multiprocessing.Pool`

For example, from this:

```
from multiprocessing.pool import Pool
```

To this:

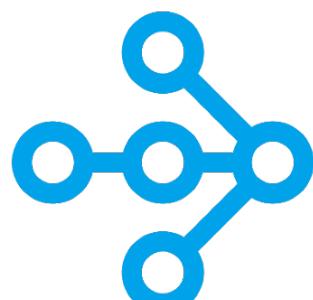
```
from ray.util.multiprocessing.pool import Pool
```

- Use Ray's implementations
 - Drop-in replacements
 - Change import statements
 - Break the one-node limitation!

See these blog posts:

<https://medium.com/distributed-computing-with-ray/how-to-scale-python-multiprocessing-to-a-cluster-with-one-line-of-code-d19f242f60ff>

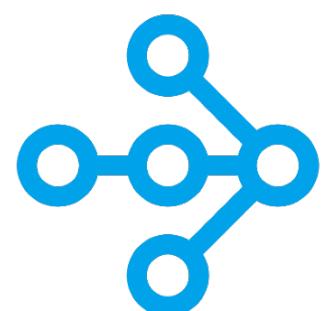
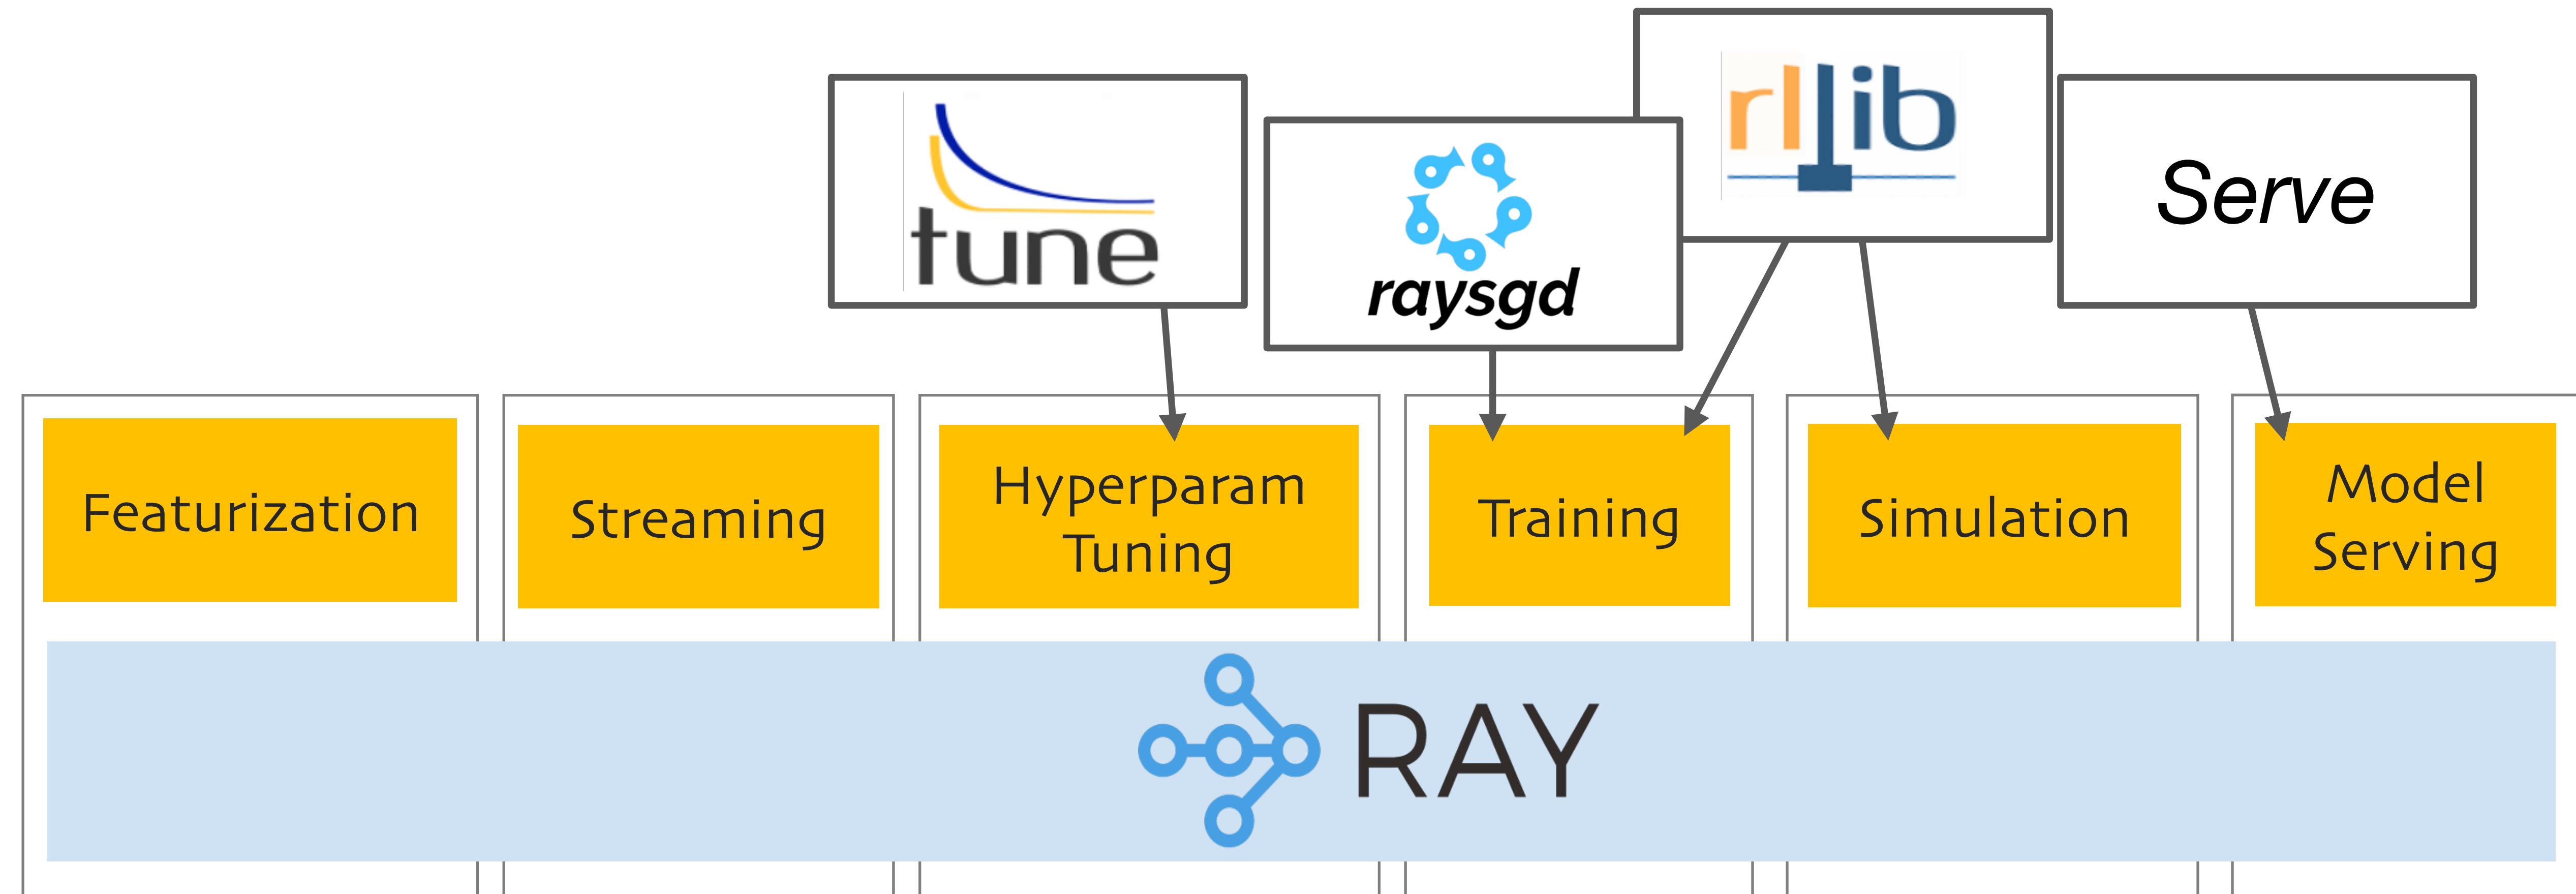
<https://medium.com/distributed-computing-with-ray/easy-distributed-scikit-learn-training-with-ray-54ff8b643b33>



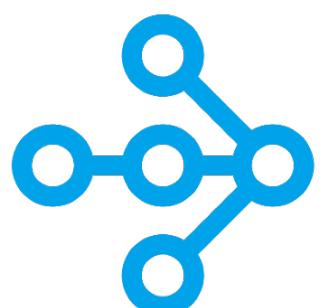
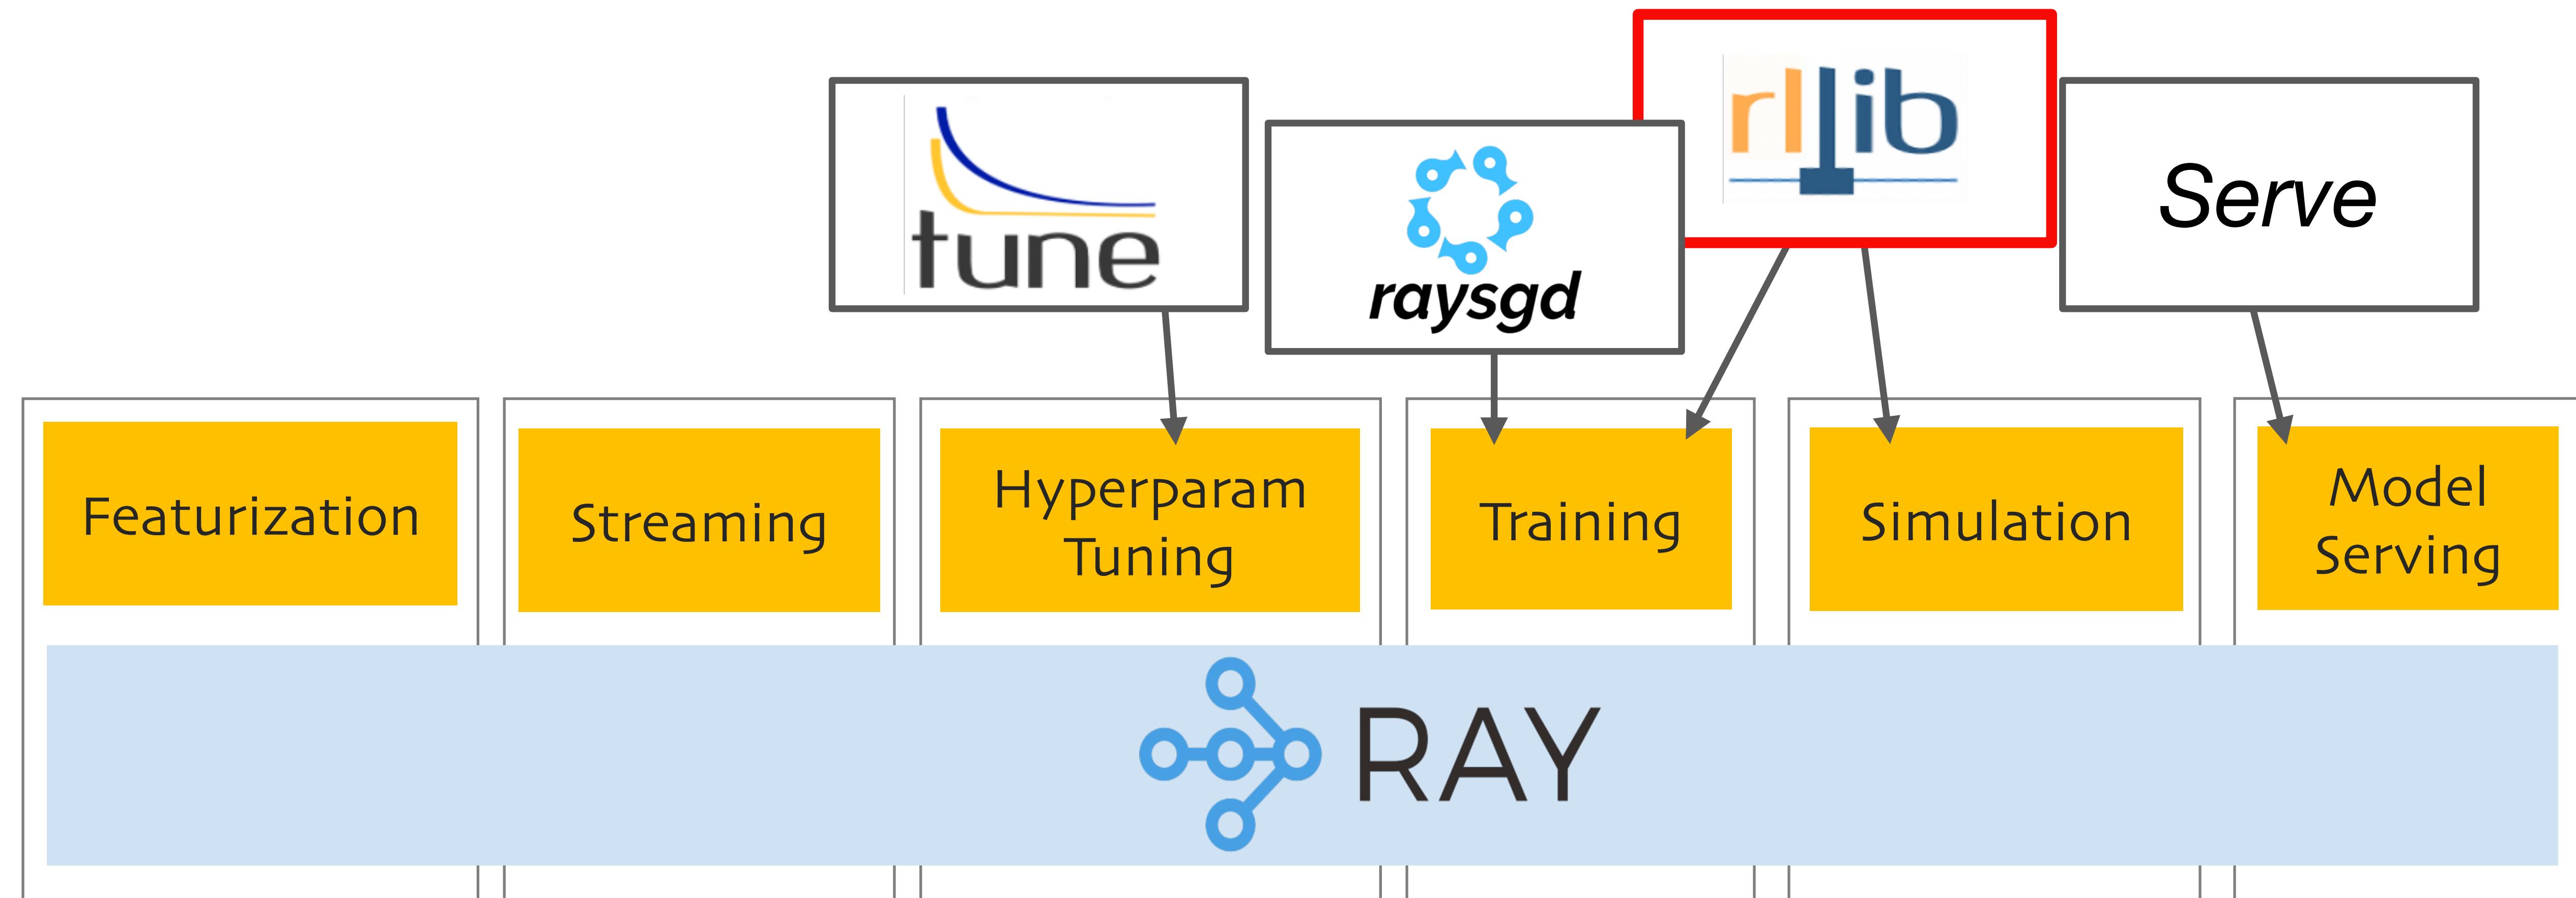


Machine Learning with Ray-based Libraries

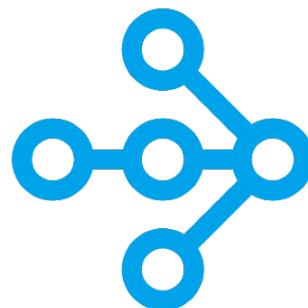
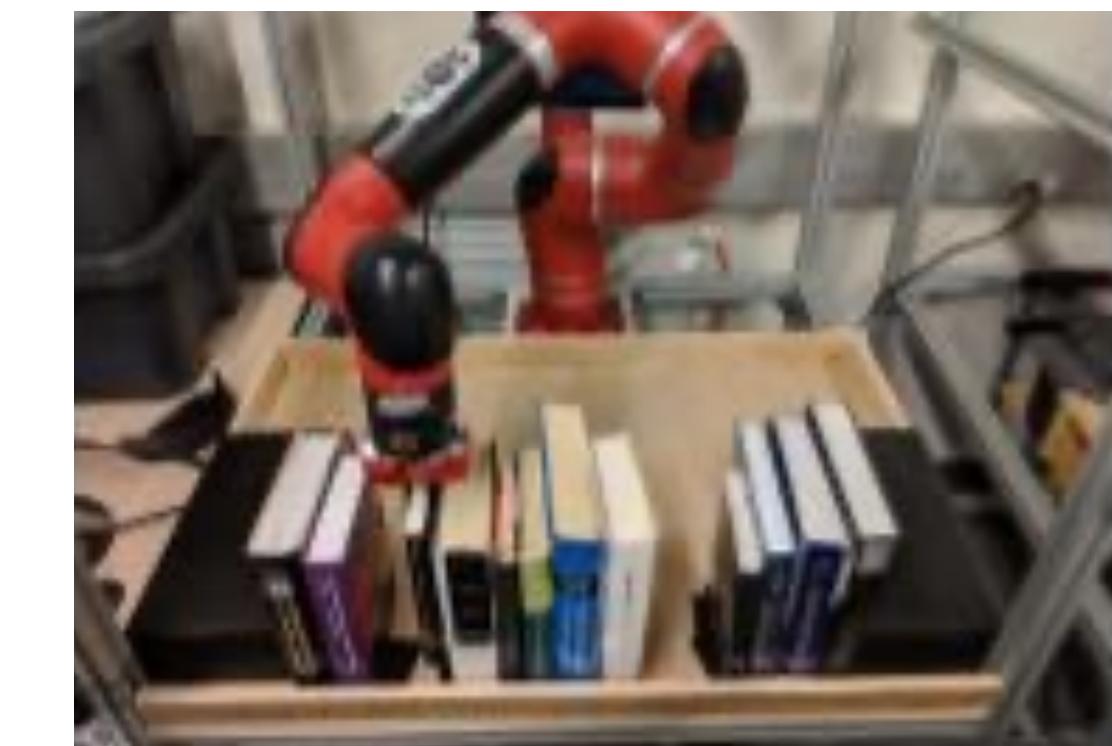
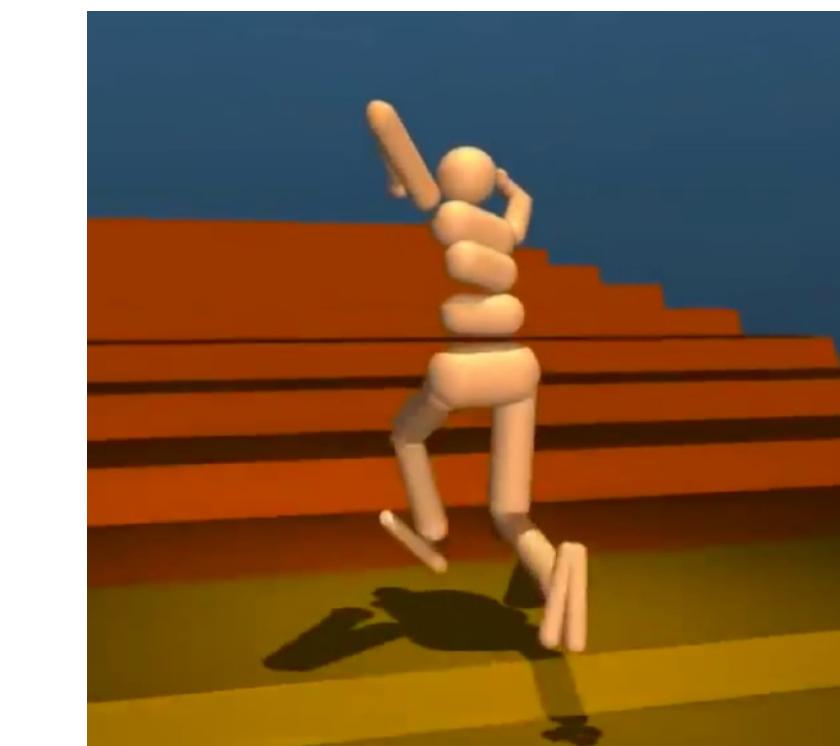
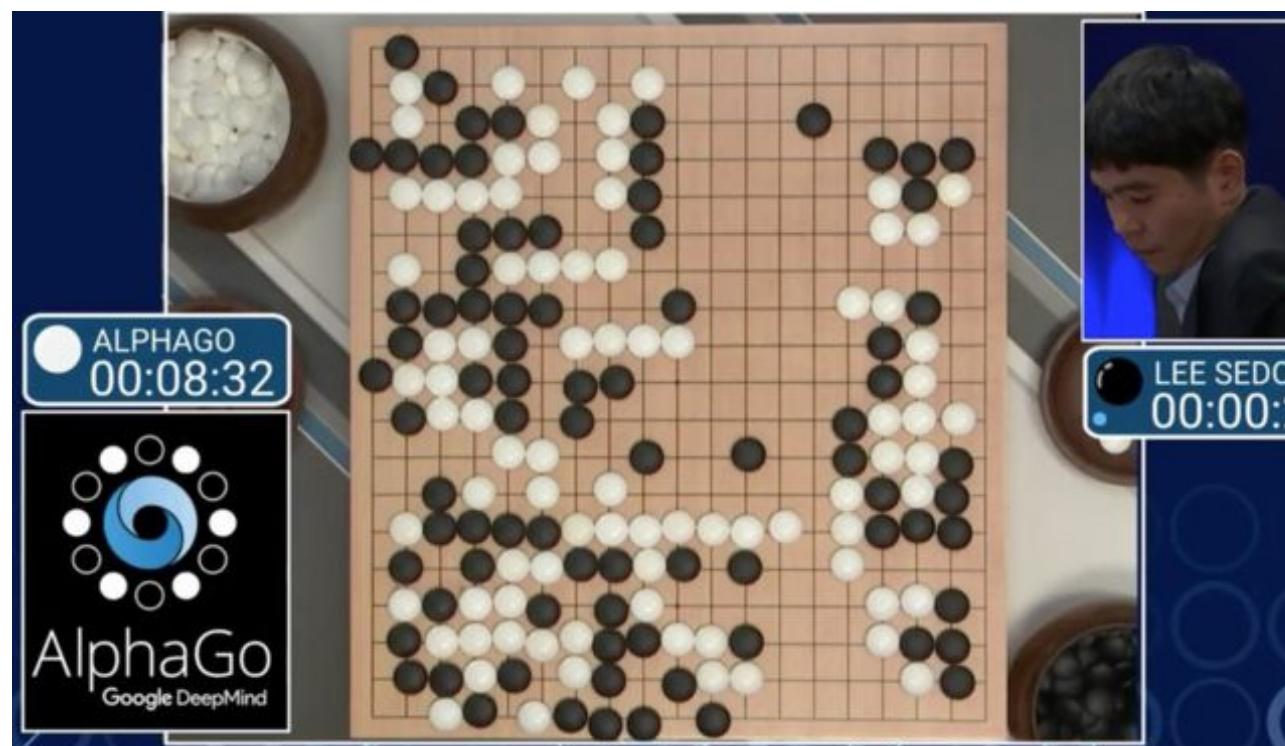
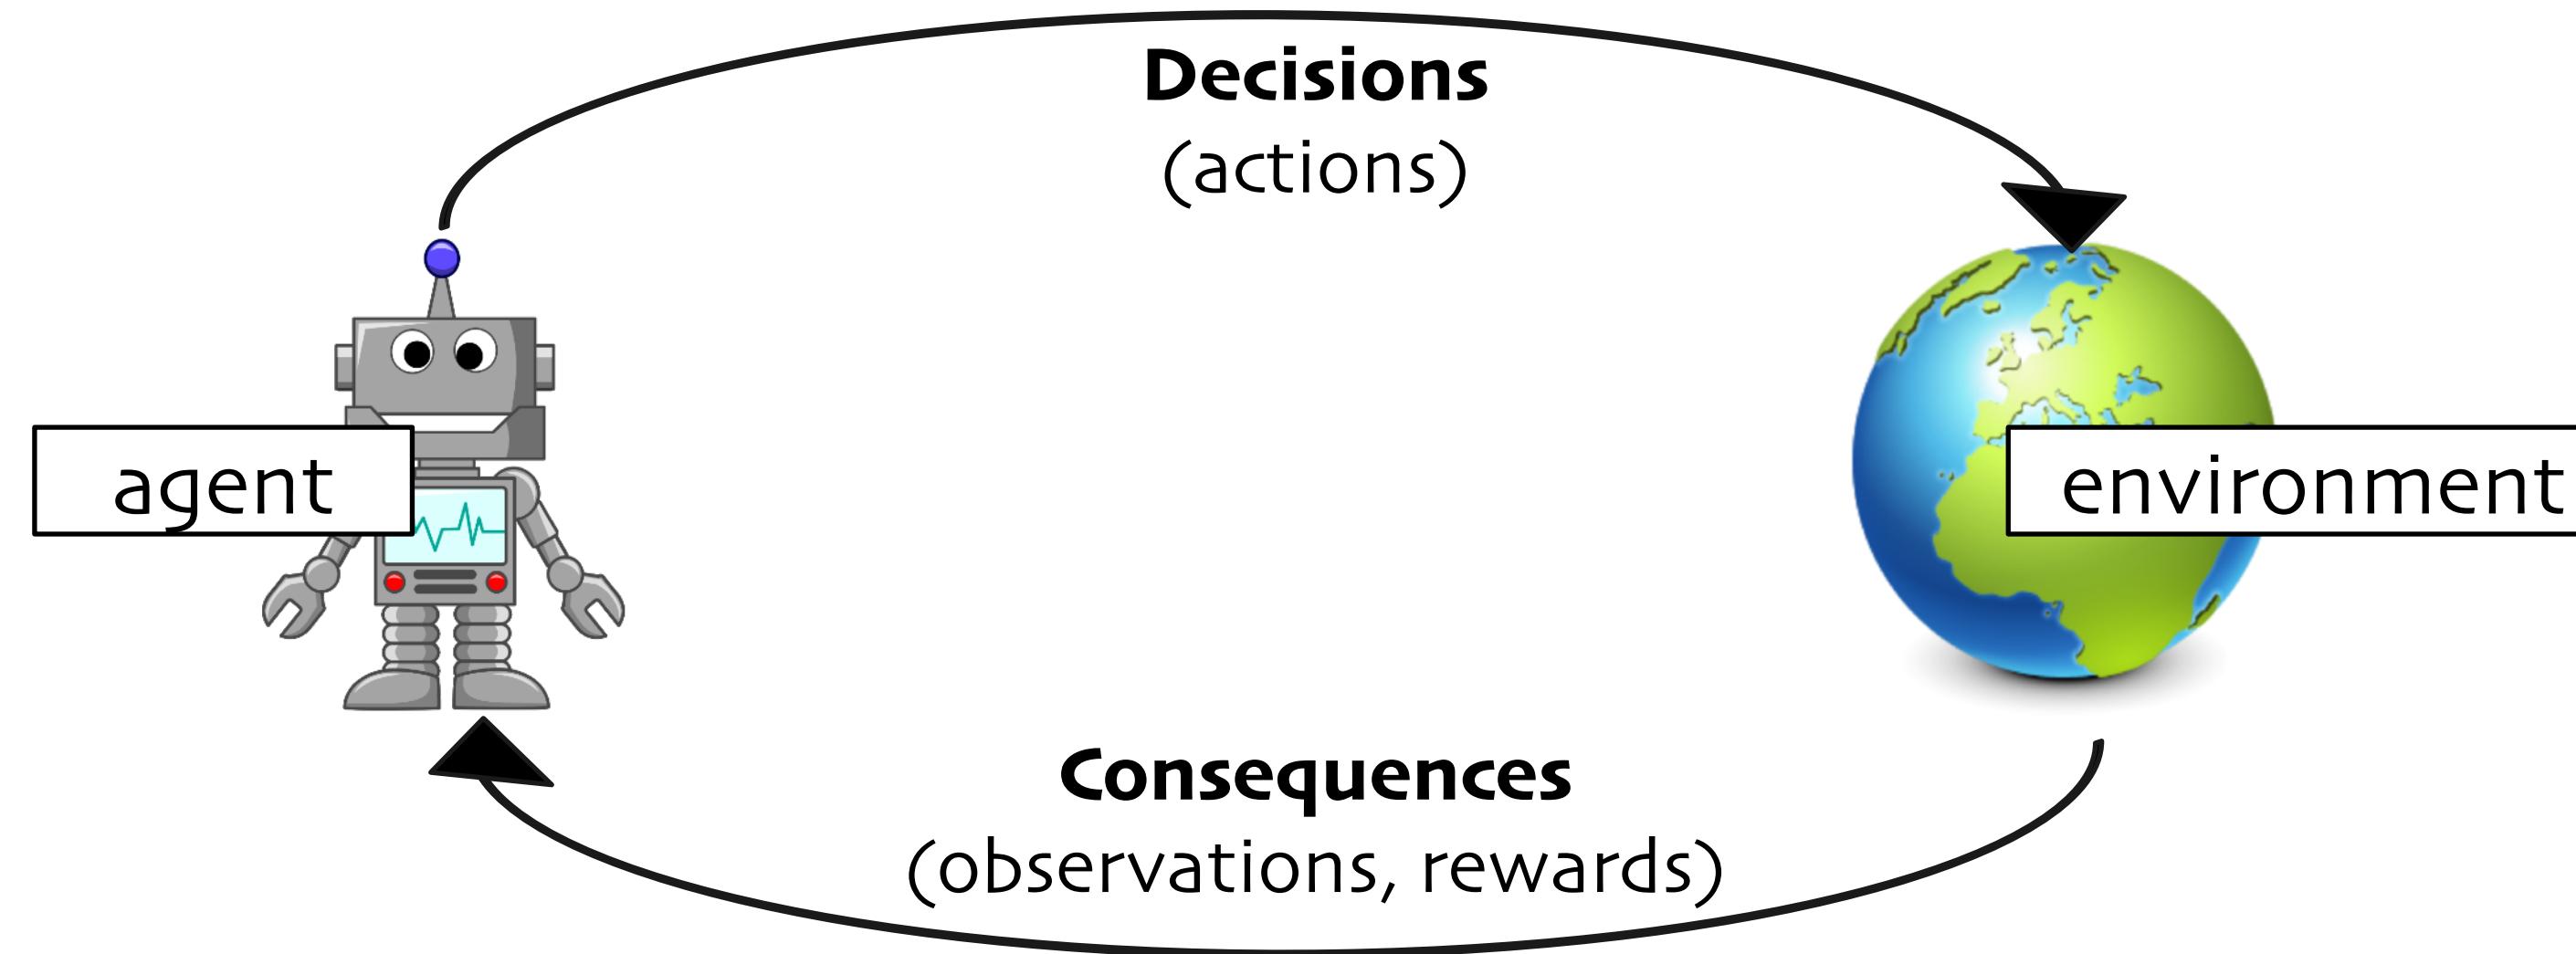
Ray Libraries



Reinforcement Learning - Ray RLLib



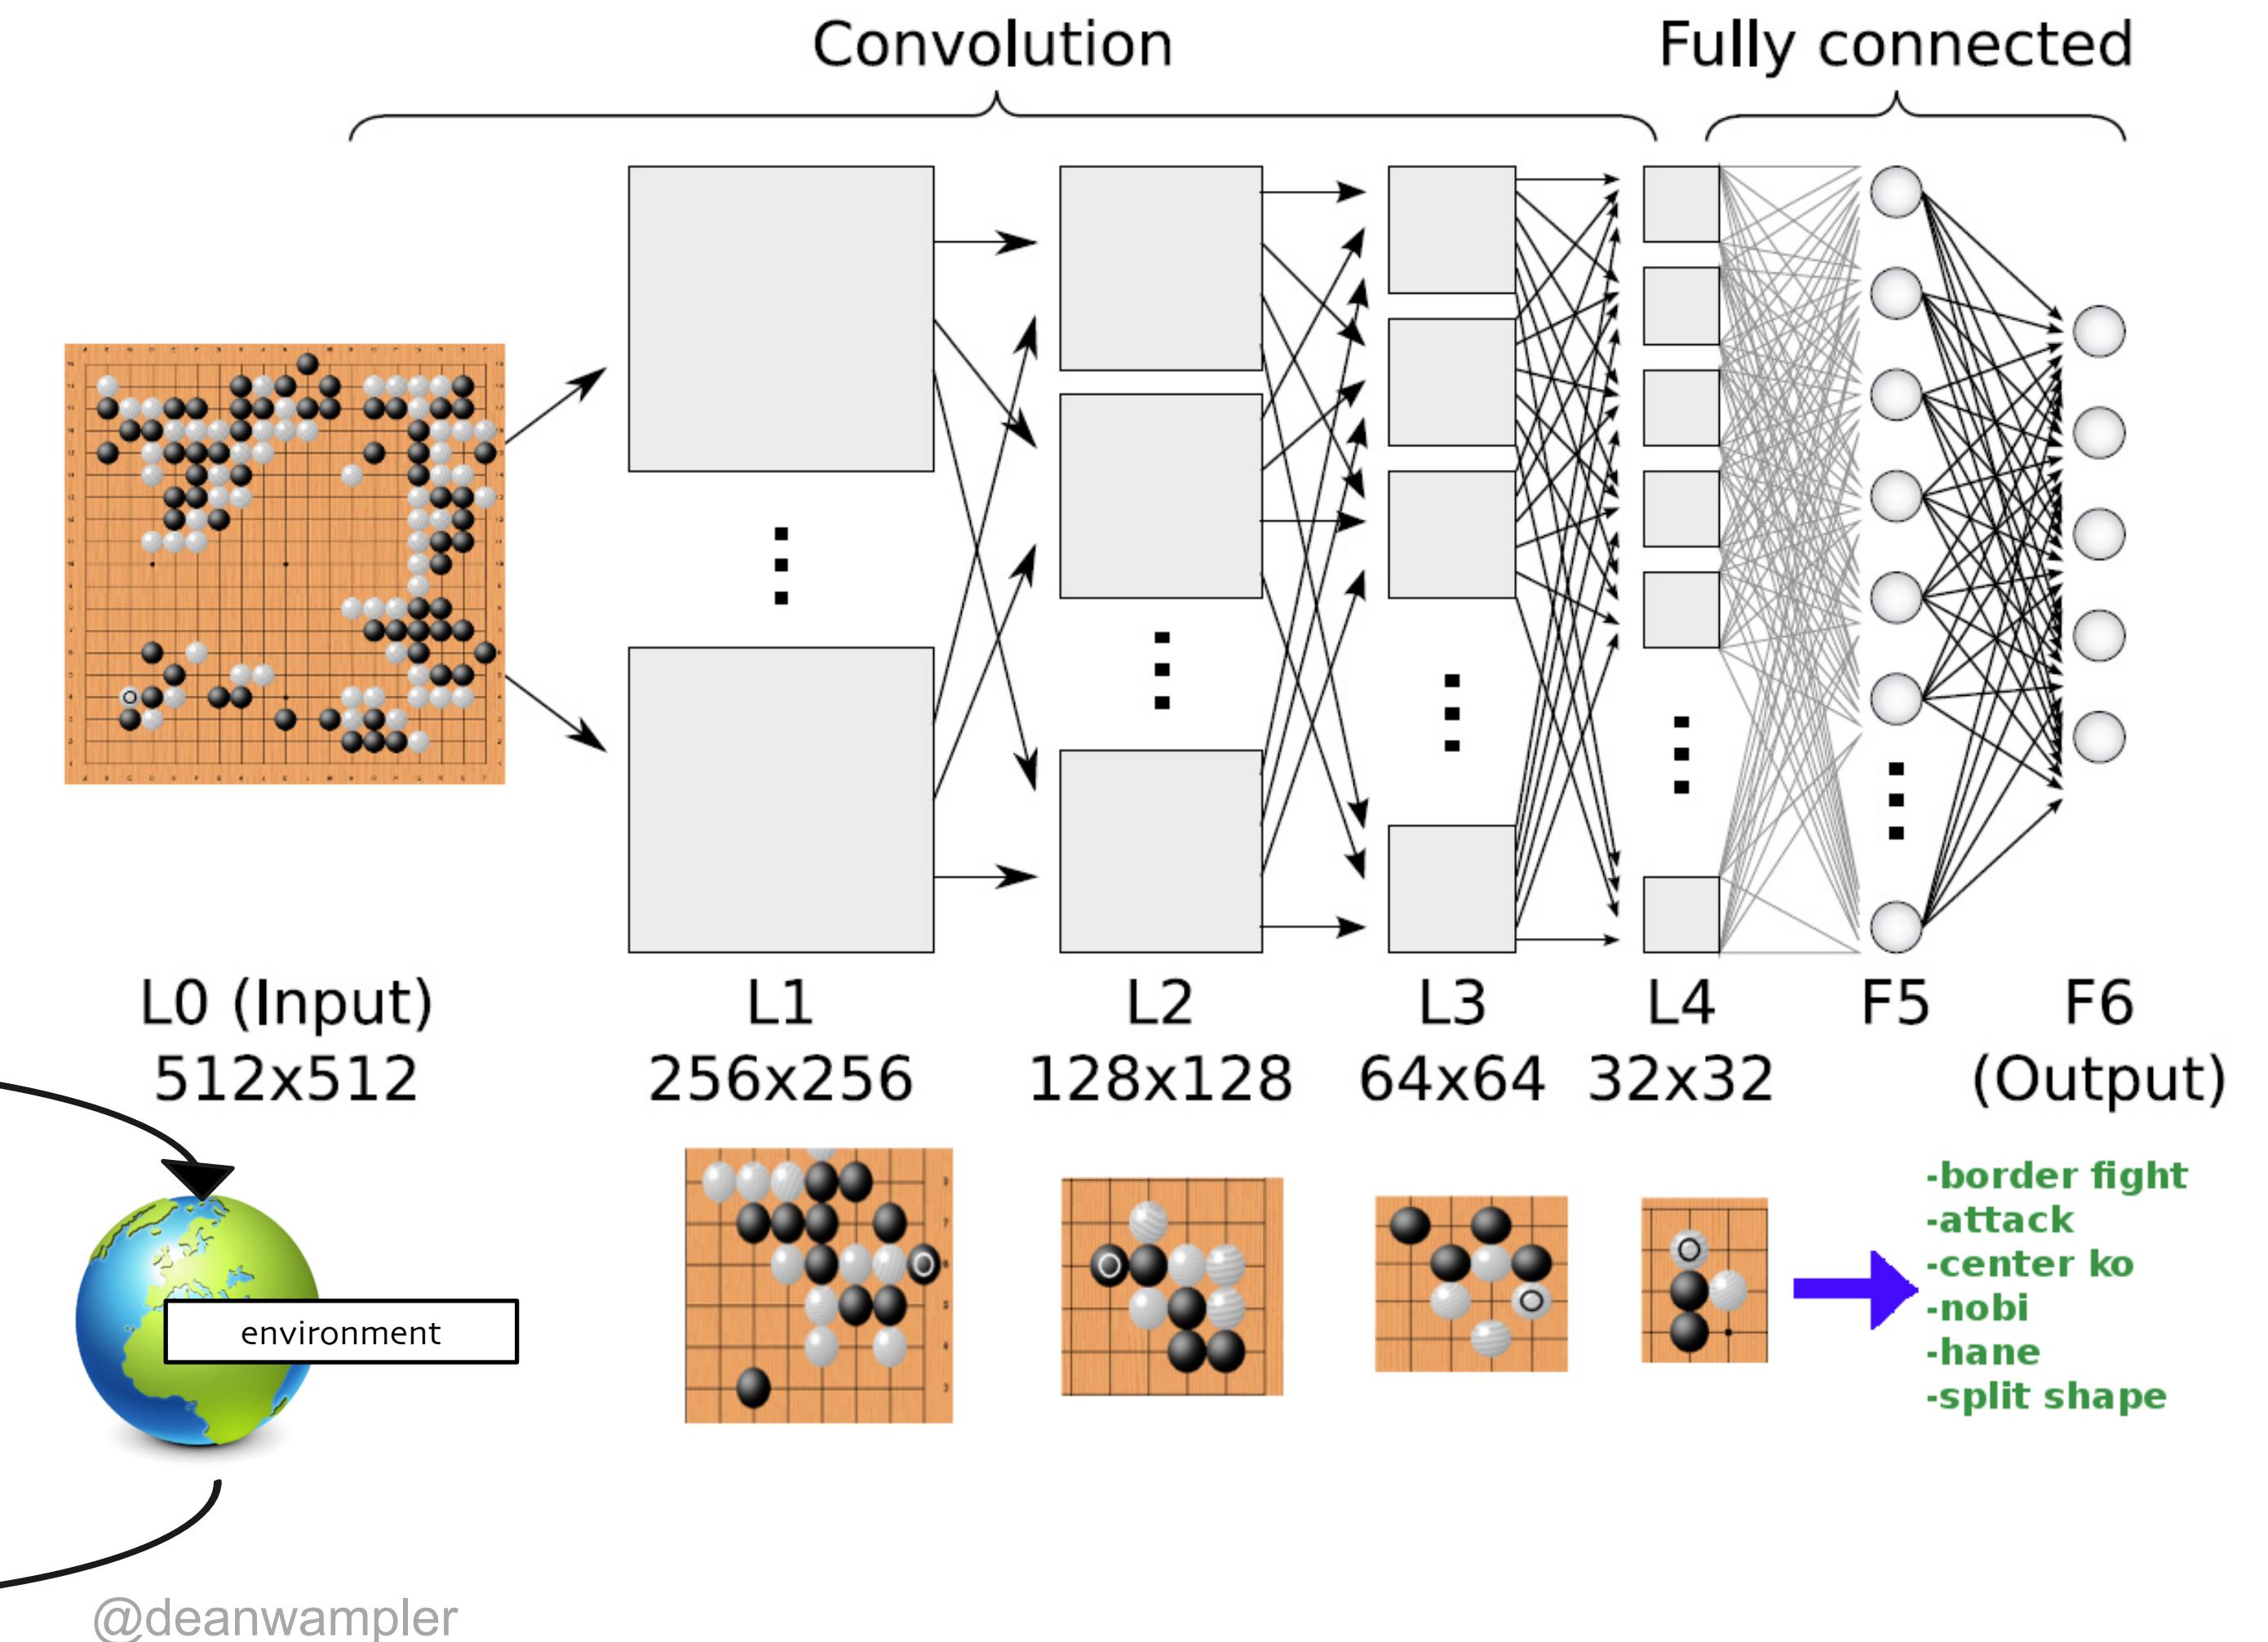
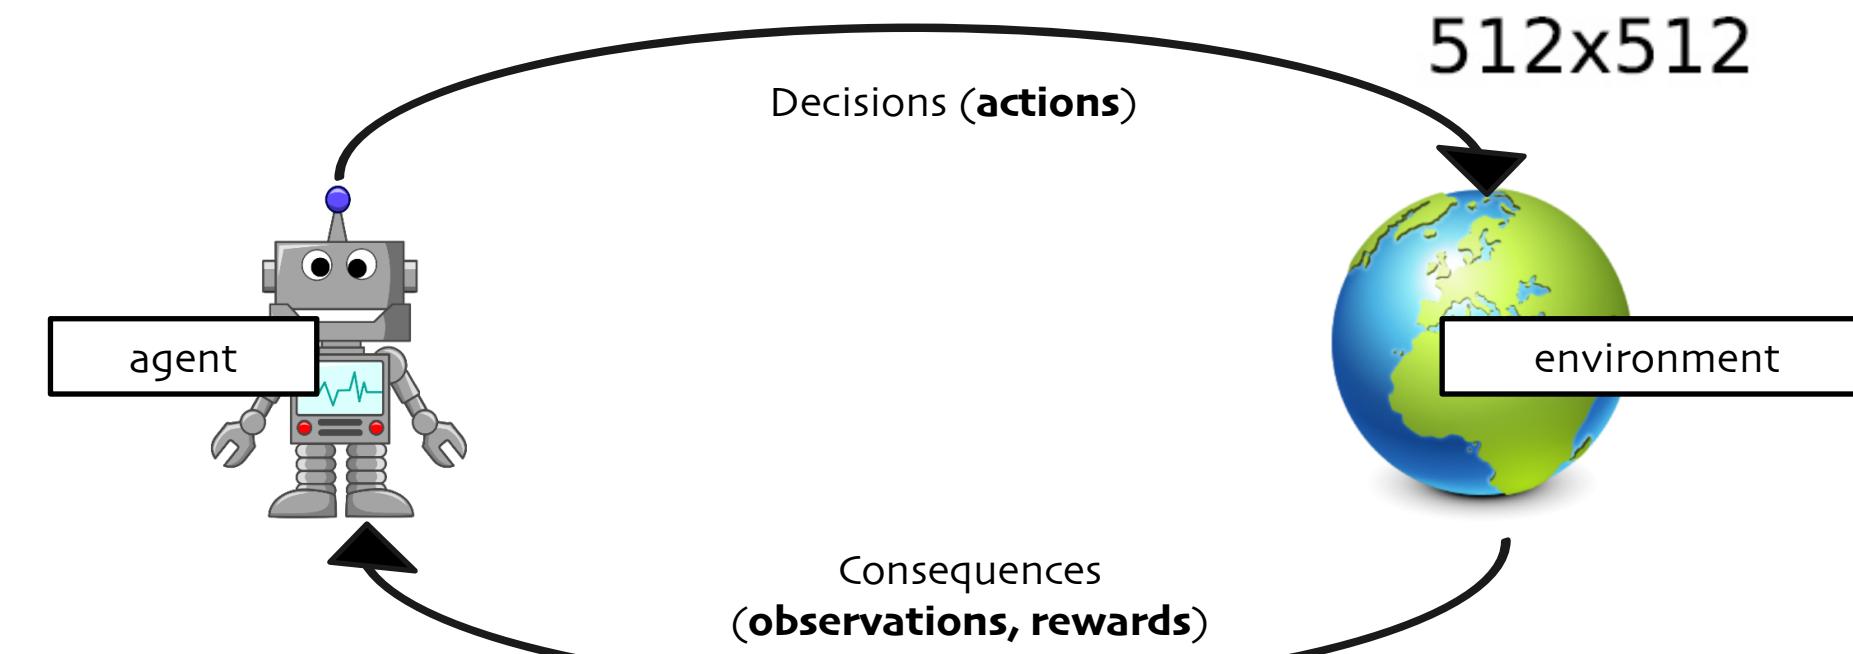
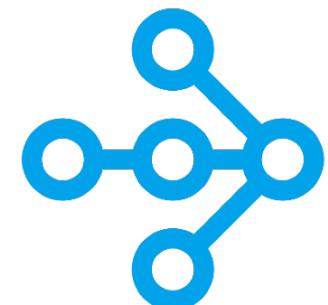
Background: Reinforcement Learning



Go as a Reinforcement Learning Problem

AlphaGo (Silver et al. 2016)

- **Observations:**
 - board state
- **Actions:**
 - where to place the stones
- **Rewards:**
 - 1 if win
 - 0 otherwise



Growing Number of RL Applications

Industrial
Processes

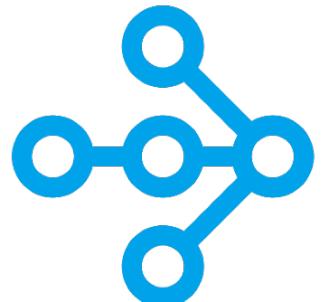
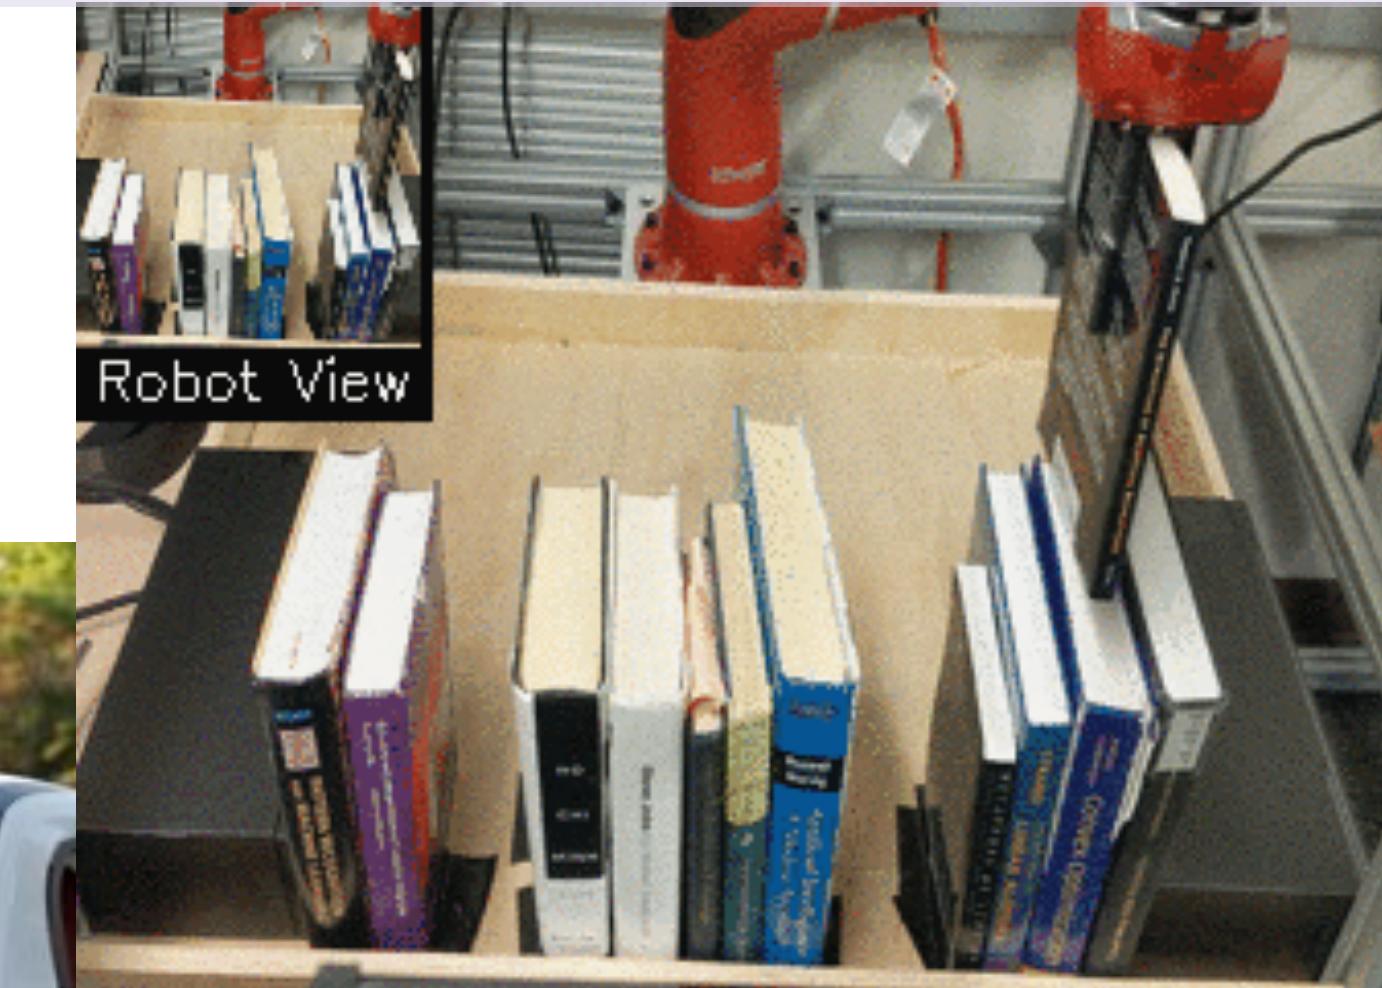
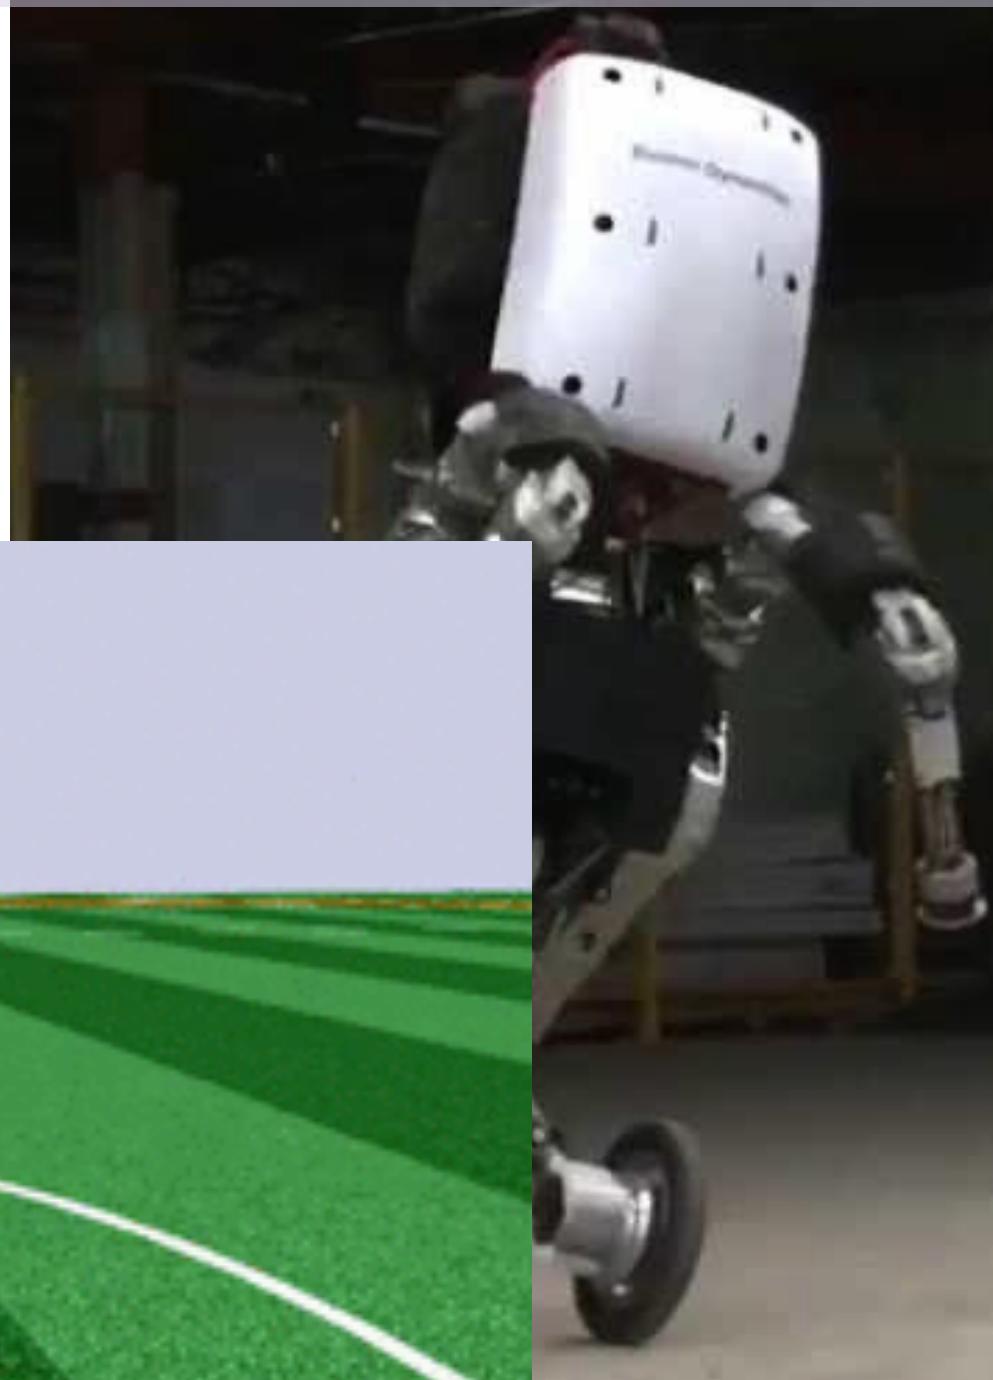
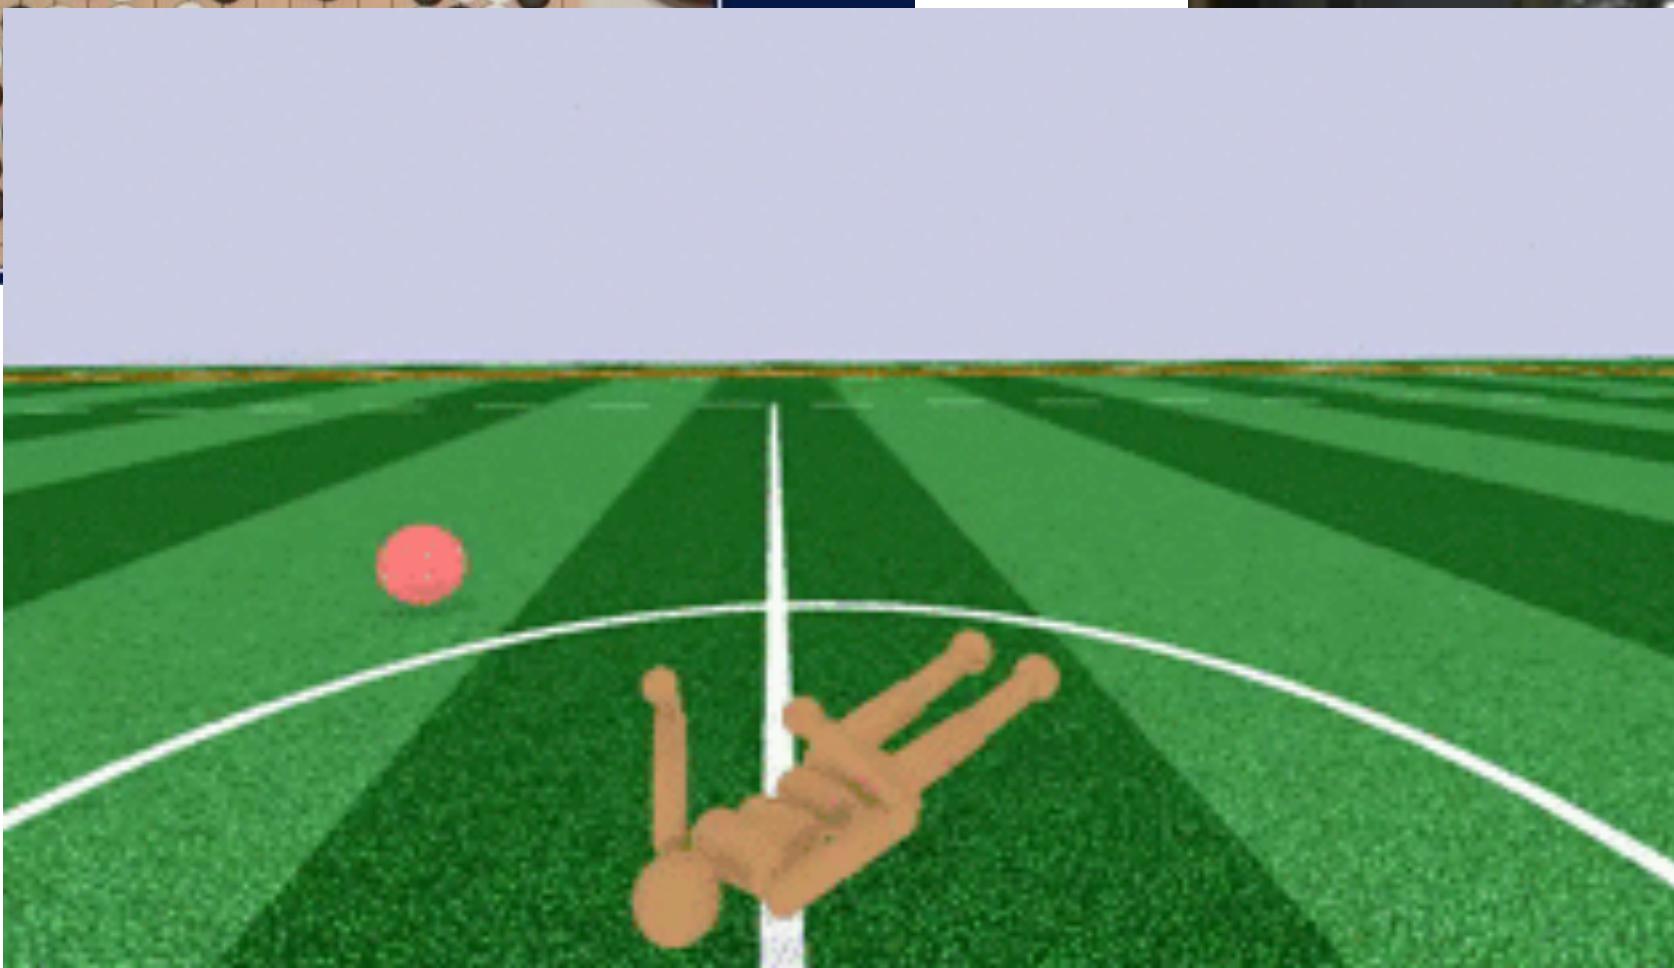
System
Optimization

Advertising

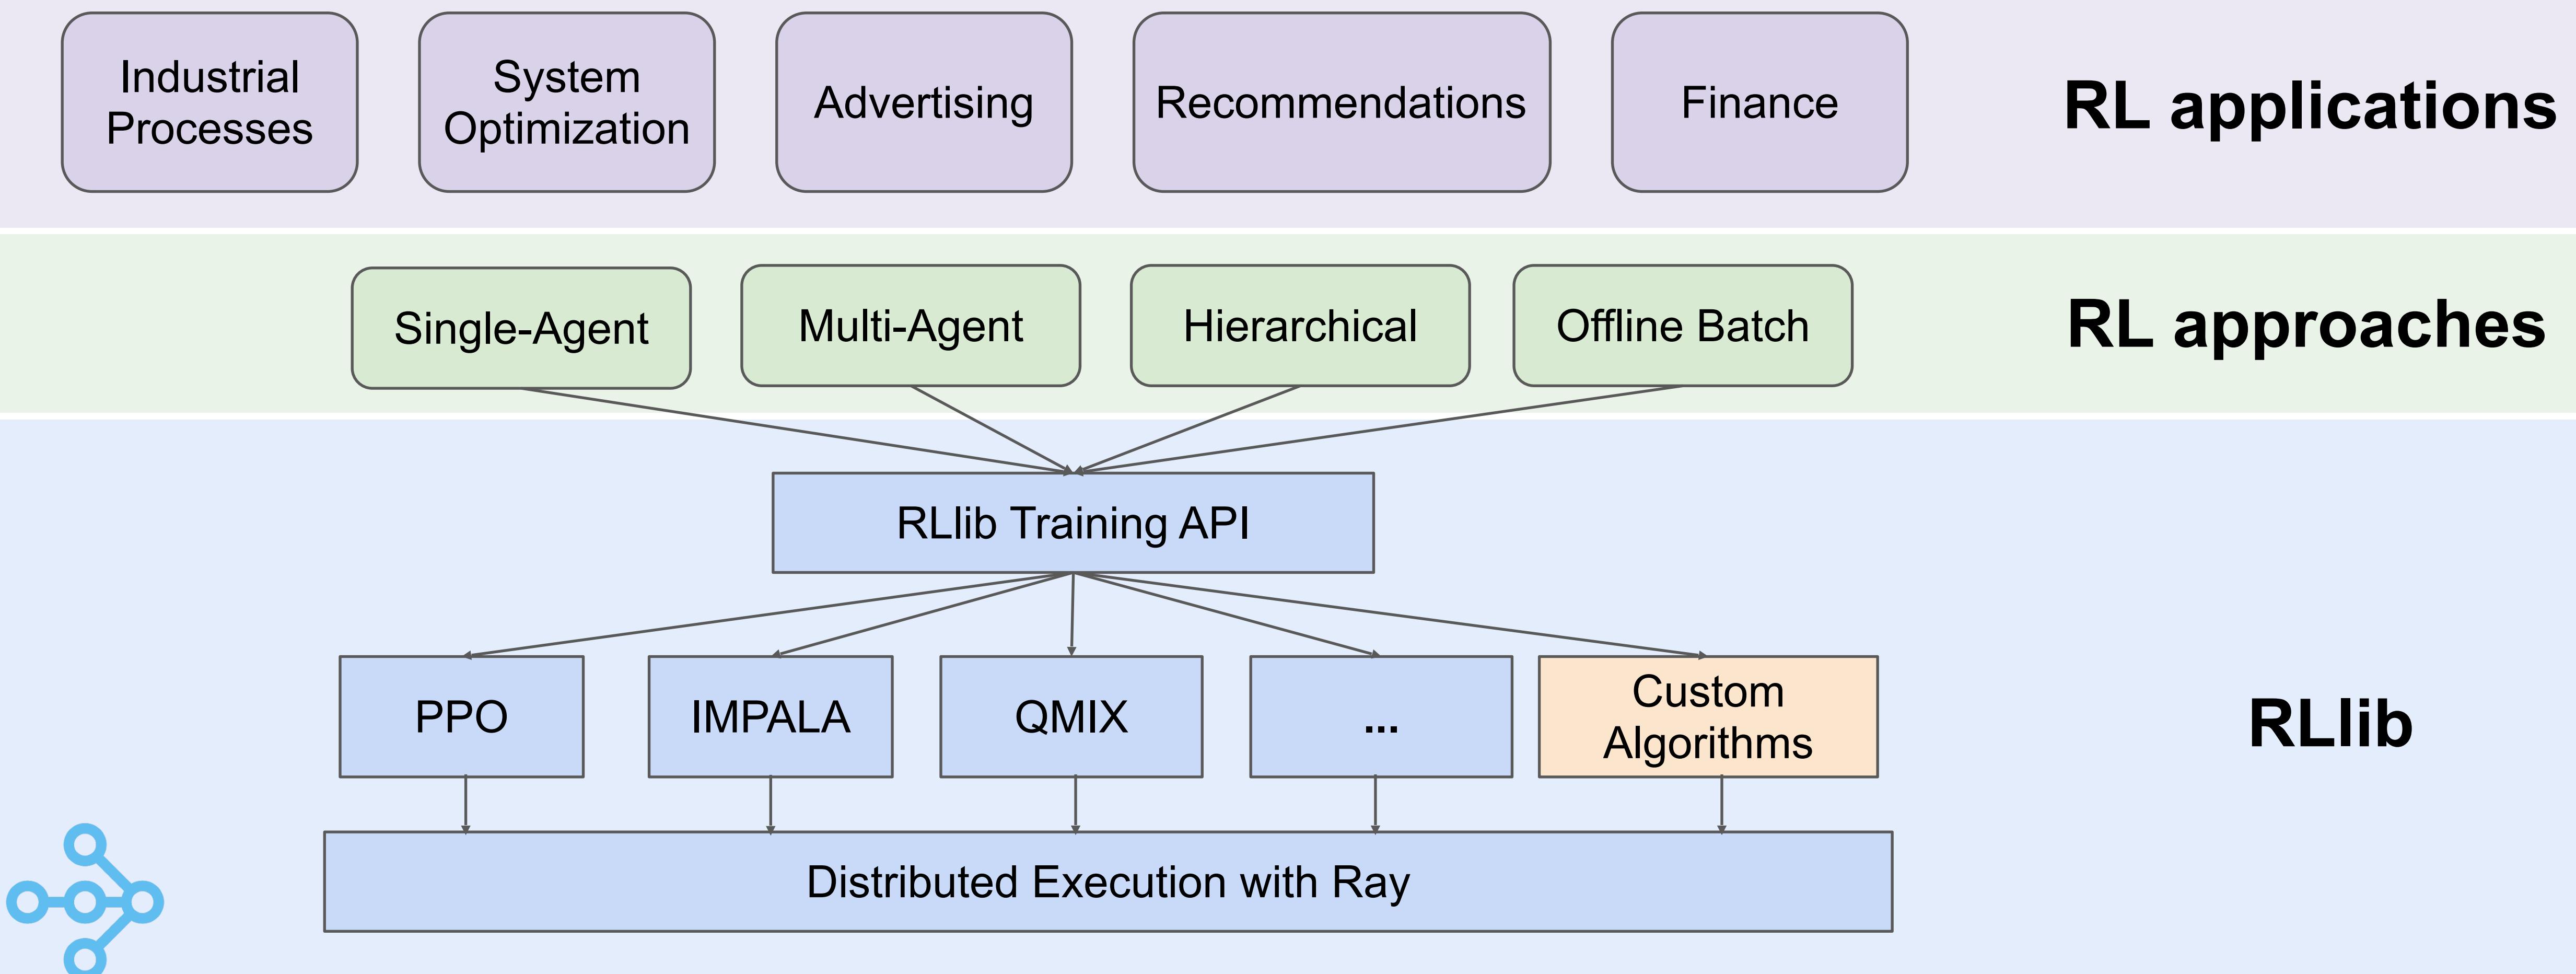
Recommendations

Finance

RL applications

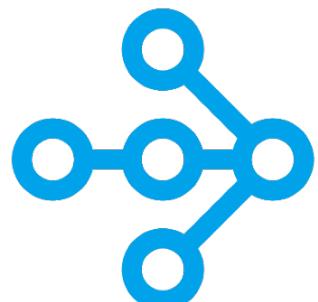


RLLib: A Scalable, Unified Library for RL



Broad Range of Scalable Algorithms

- High-throughput architectures
 - [Distributed Prioritized Experience Replay \(Ape-X\)](#)
 - [Importance Weighted Actor-Learner Architecture \(IMPALA\)](#)
 - [Asynchronous Proximal Policy Optimization \(APPO\)](#)
- Gradient-based
 - [Soft Actor-Critic \(SAC\)](#)
 - [Advantage Actor-Critic \(A₂C, A₃C\)](#)
 - [Deep Deterministic Policy Gradients \(DDPG, TD3\)](#)
 - [Deep Q Networks \(DQN, Rainbow, Parametric DQN\)](#)
 - [Policy Gradients](#)
 - [Proximal Policy Optimization \(PPO\)](#)
- gradient-free
 - [Augmented Random Search \(ARS\)](#)
 - [Evolution Strategies](#)
- Multi-agent specific
 - [QMIX Monotonic Value Factorisation \(QMIX, VDN, IQN\)](#)
- Offline
 - [Advantage Re-Weighted Imitation Learning \(MARWIL\)](#)



Amazon SageMaker RL

Reinforcement learning for every developer and data scientist



Amazon SageMaker RL

End-to-end examples for classic RL and real-world RL applications

Robotics

Industrial Control

HVAC

Autonomous Vehicles

Operations

Finance

Games

NLP

RL Environments to model real-world problems

AWS Simulation Environments

Amazon Sumerian

AWS RoboMaker

Open Source Environments

EnergyPlus

RoboSchool

PyBullet

...

Custom Environments

Bring Your Own

Commercial simulators

MATLAB & Simulink

Open AI Gym

RL Toolkits that provide RL agent algorithm implementations

RL-Coach

DQN

PPO

HER

Rainbow

...

RL-Ray RLLib

APEX

ES

IMPALA

A3C

...

Open AI Baselines

TRPO

GAIL

...

...

TensorFlow

MxNet

PyTorch

Chainer

Training Options

Single Machine / Distributed

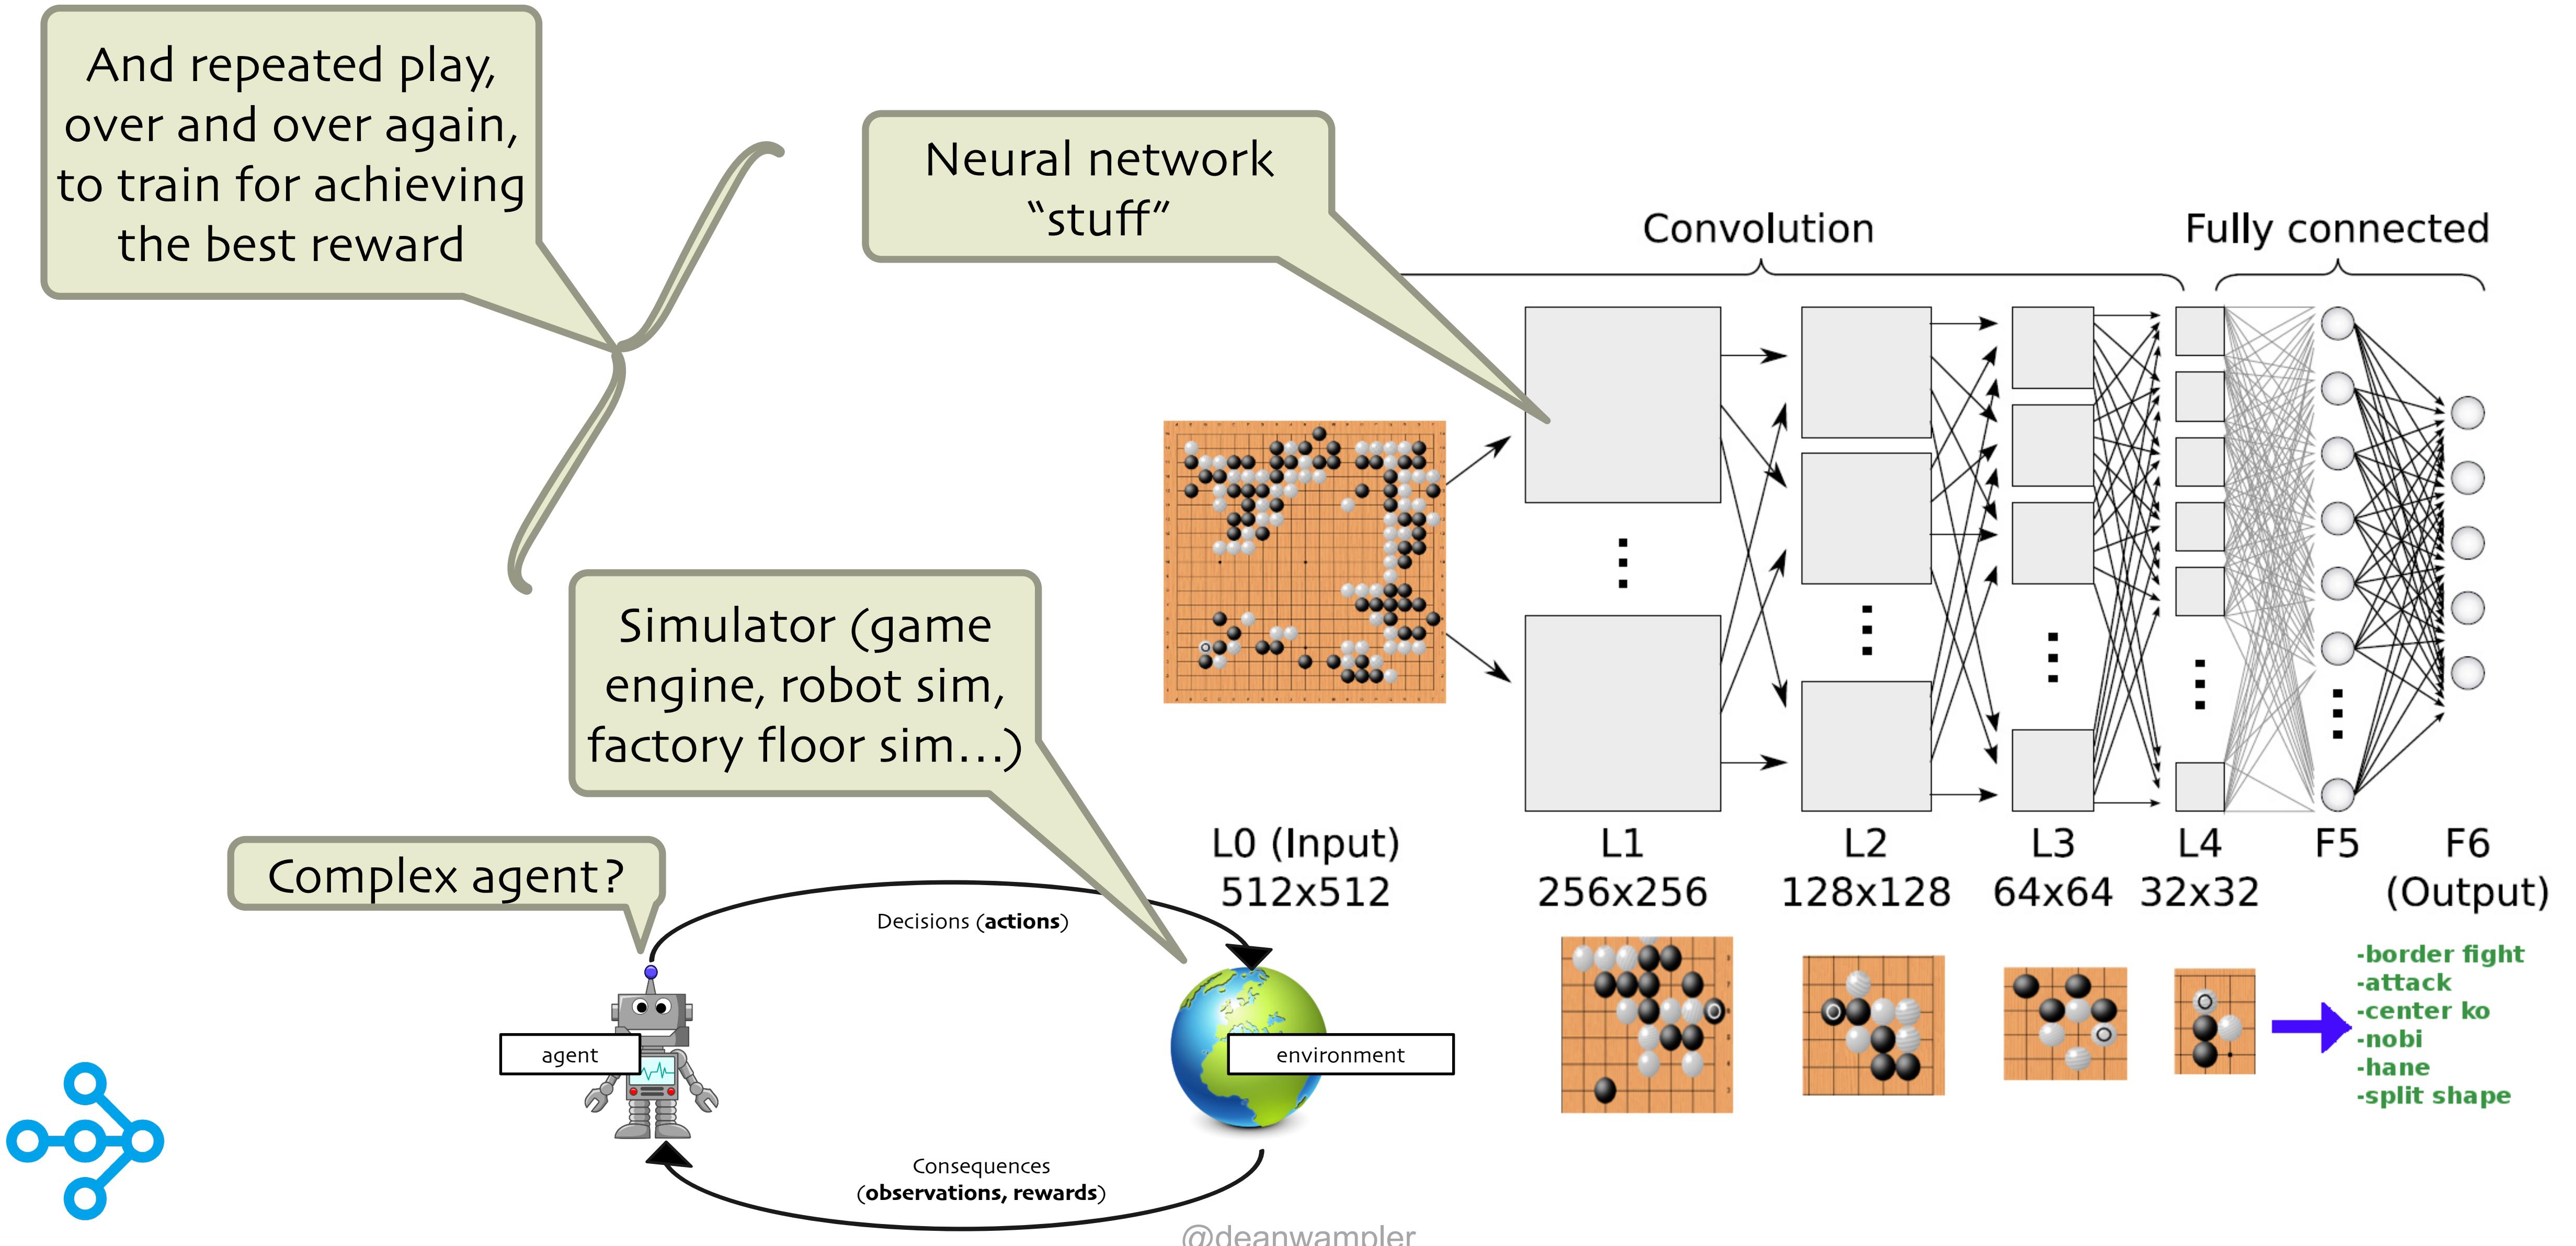
Local / Remote simulation

CPU / GPU Hardware

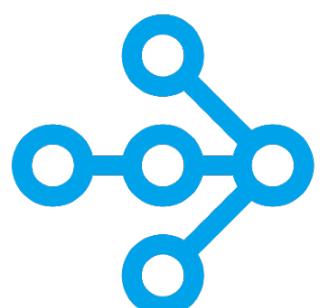
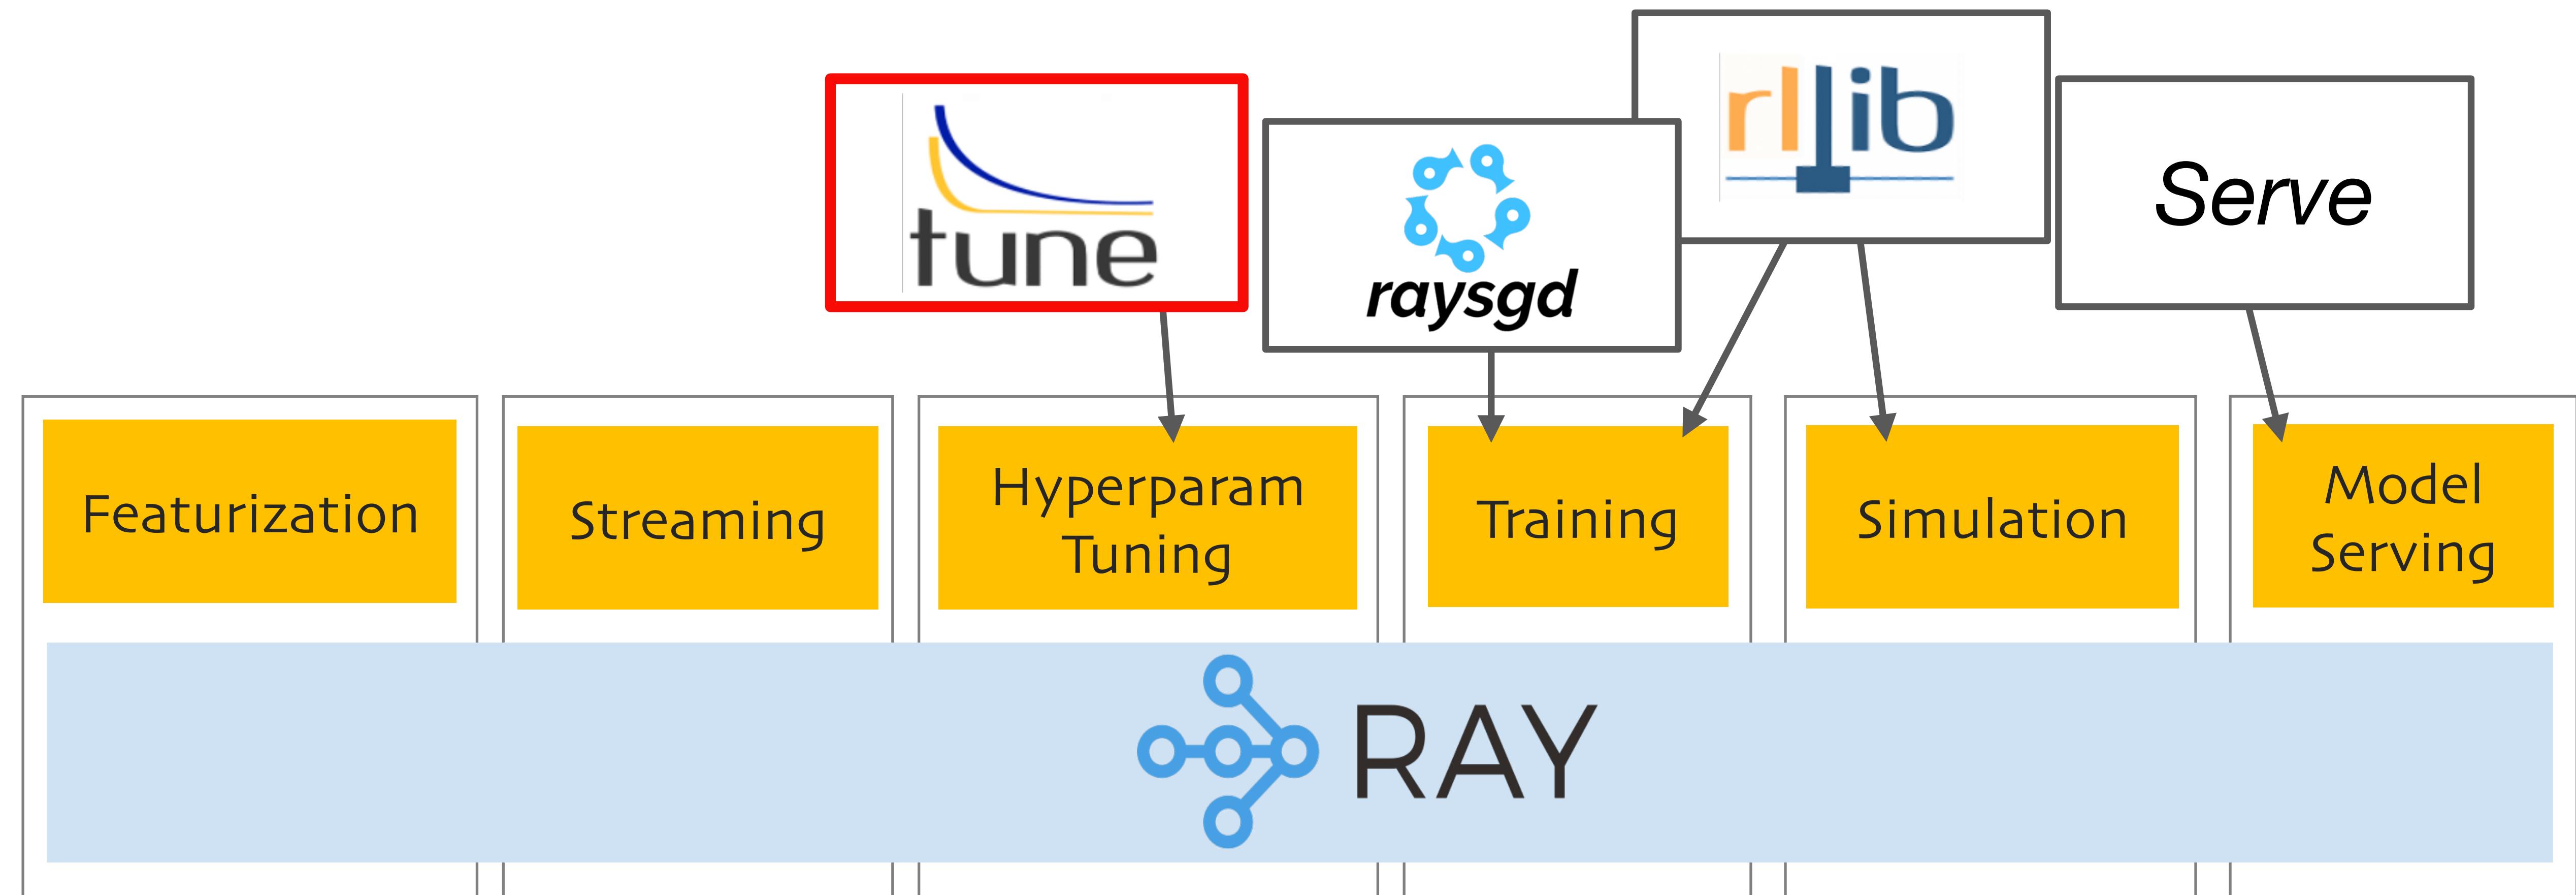
SageMaker supported

Customer BYO

Diverse Compute Requirements Motivated Creation of Ray!



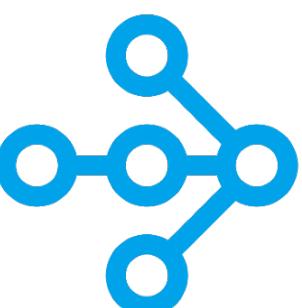
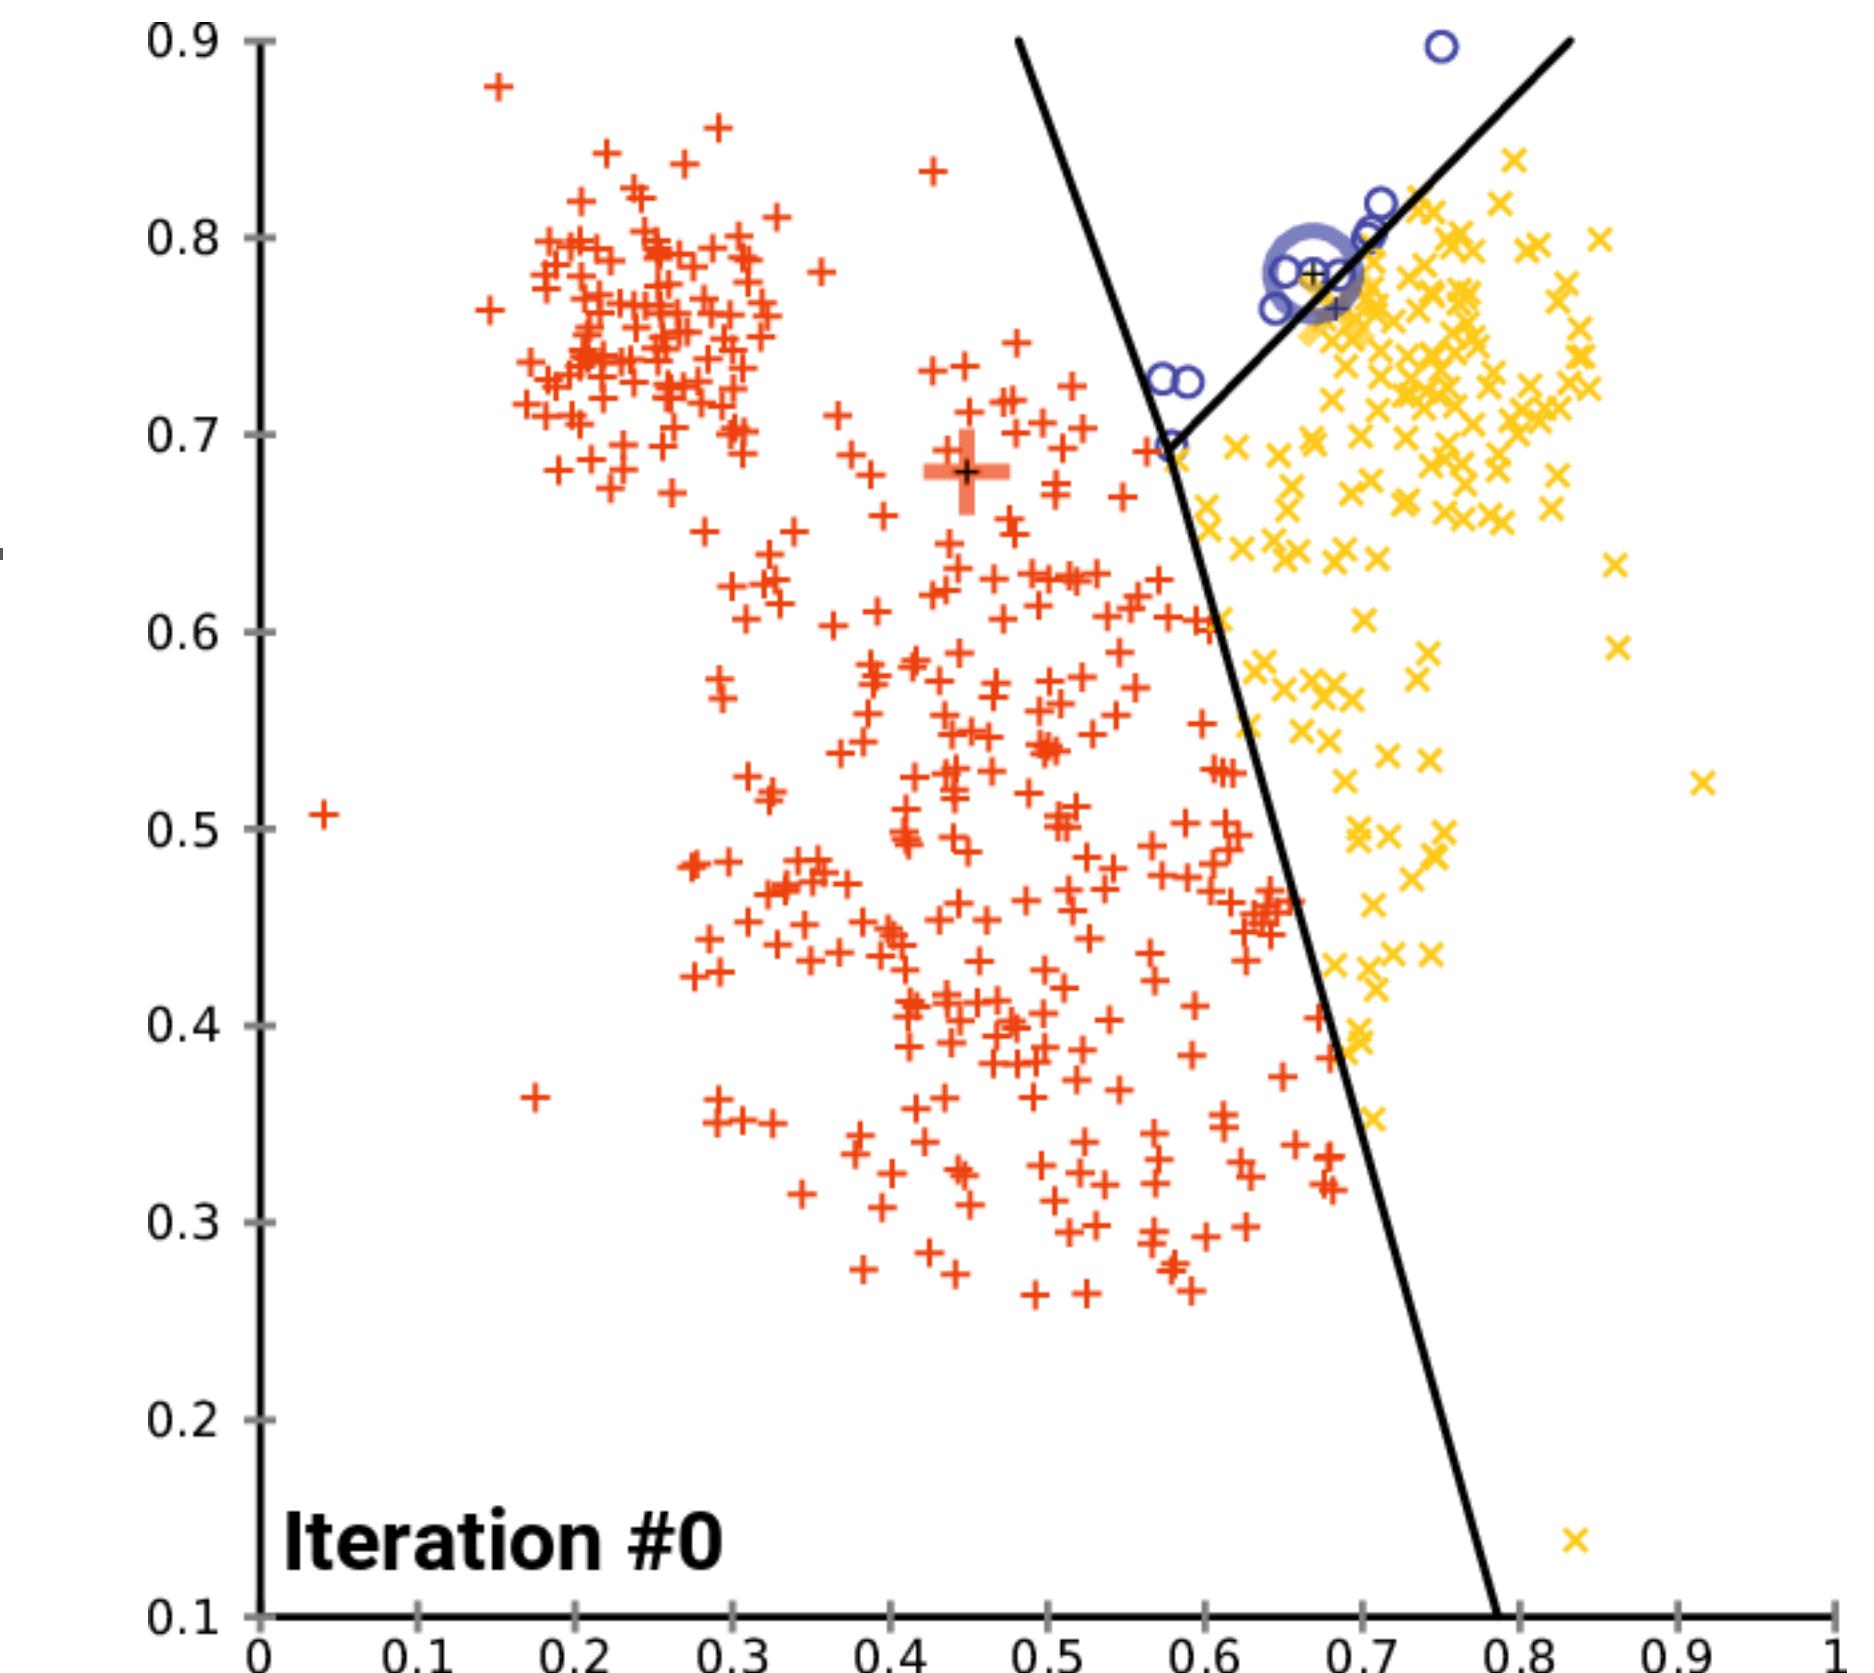
Hyperparameter Tuning - Ray Tune



What Is Hyperparameter Tuning?

Trivial example:

- What's the best value for "k" in k-means??
- k is a "hyperparameter"
- The resulting clusters are defined by "parameters"

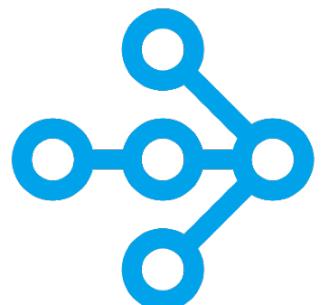
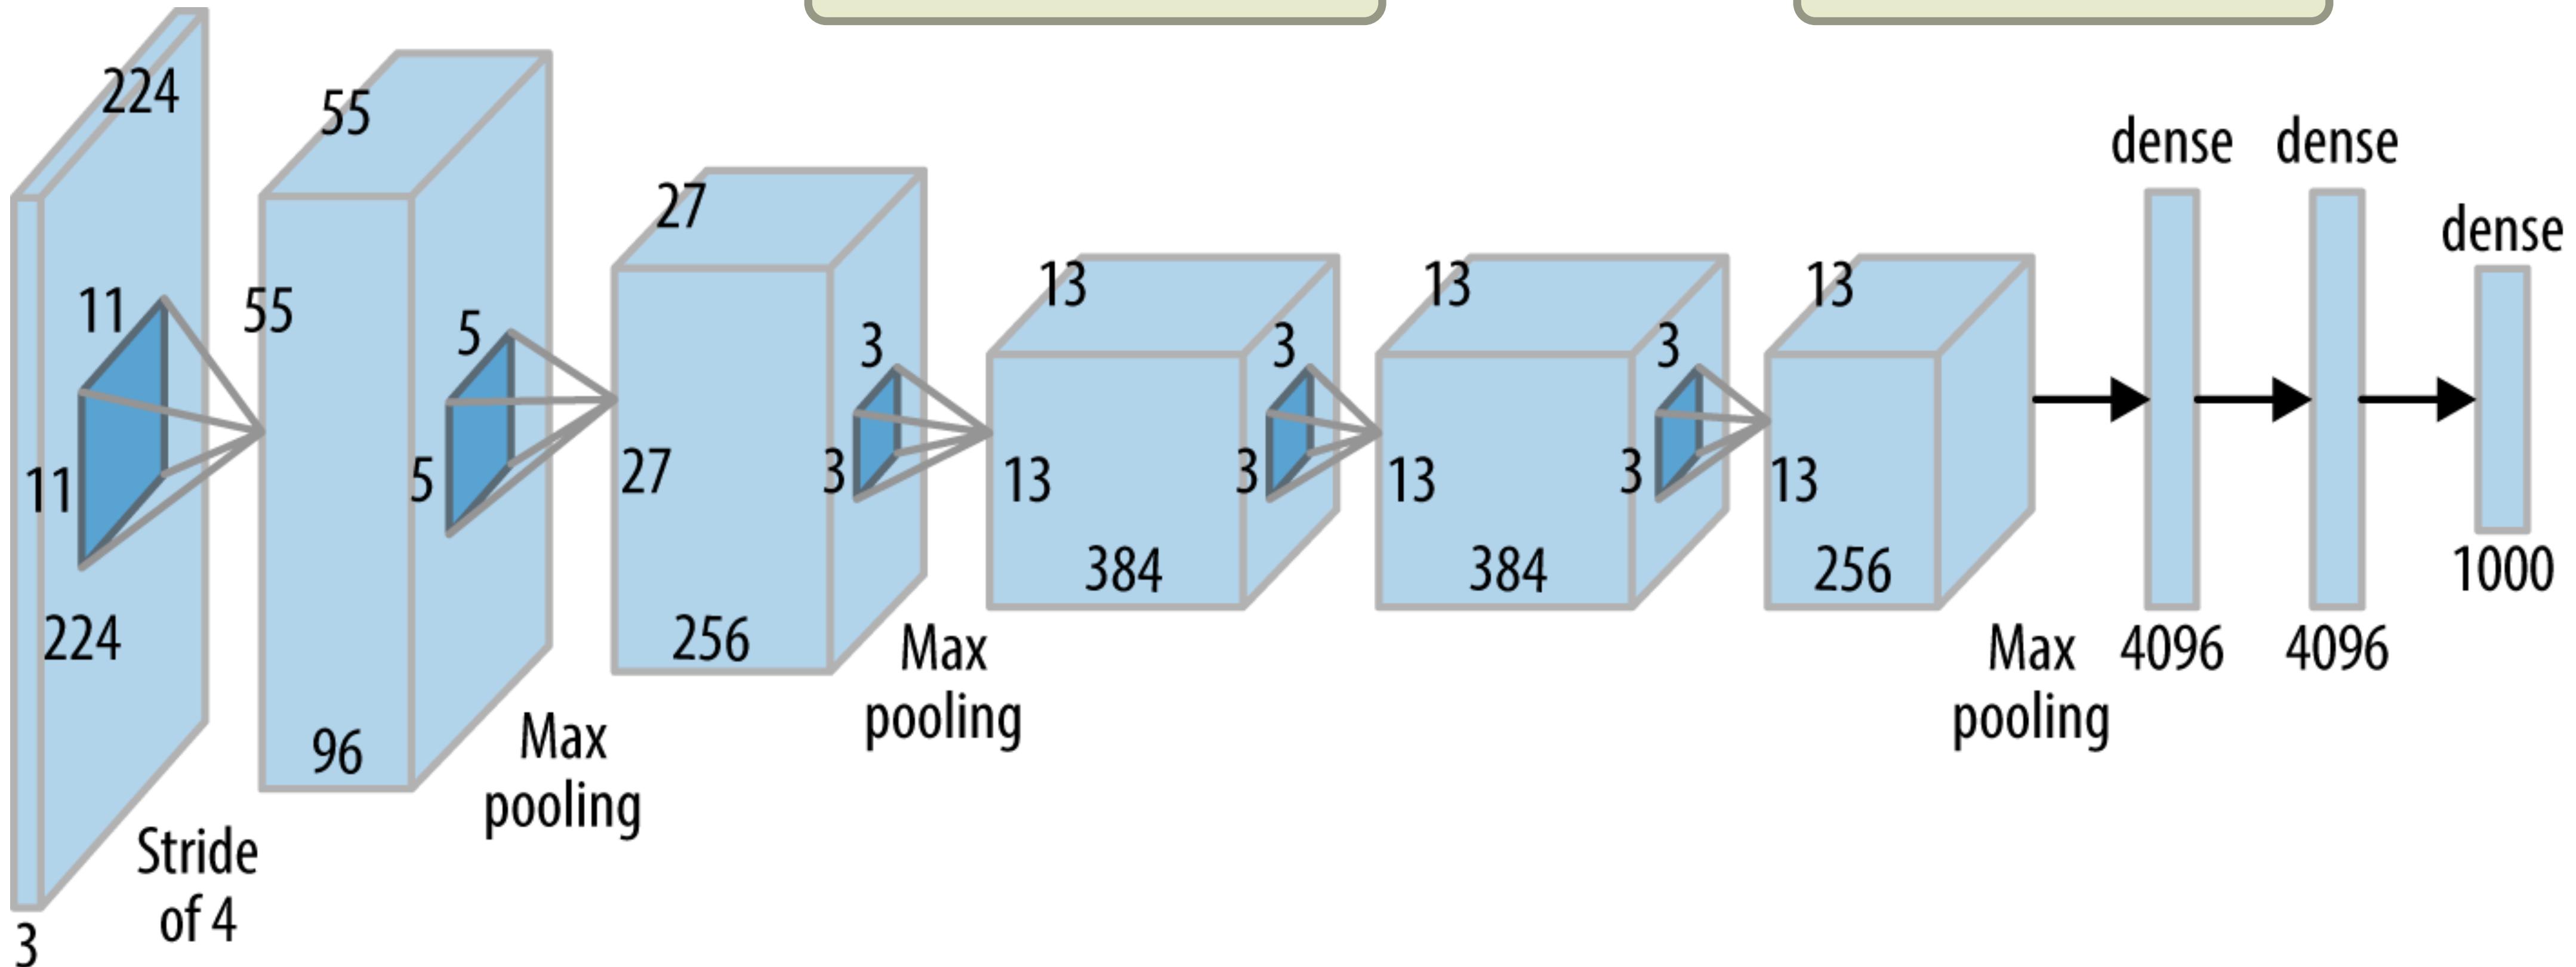


Source: https://commons.wikimedia.org/wiki/File:K-means_convergence.gif

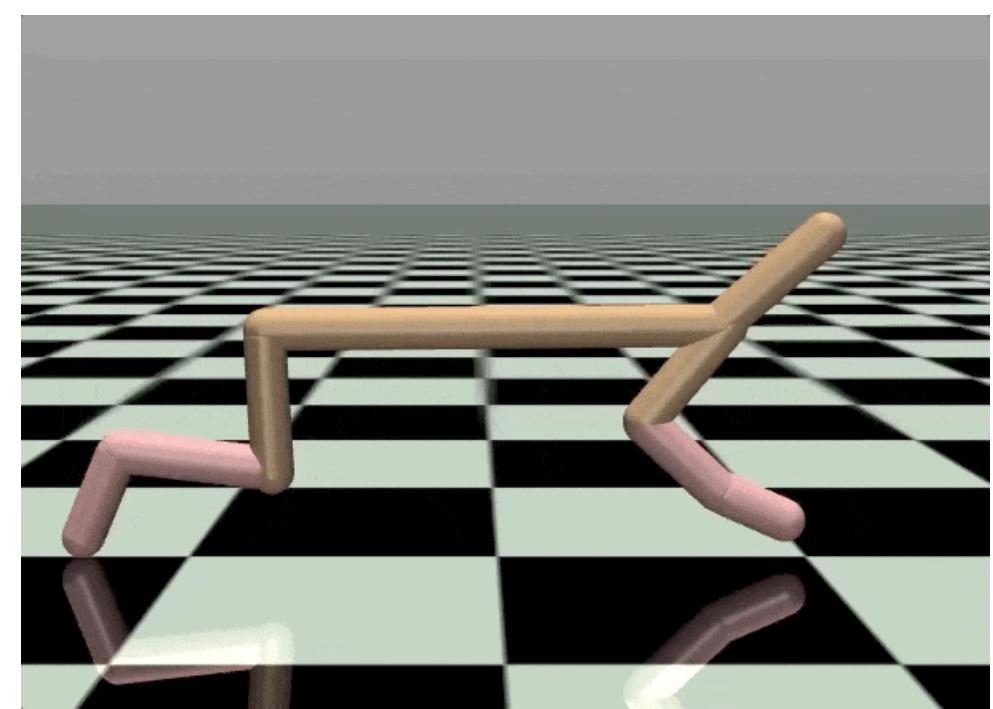
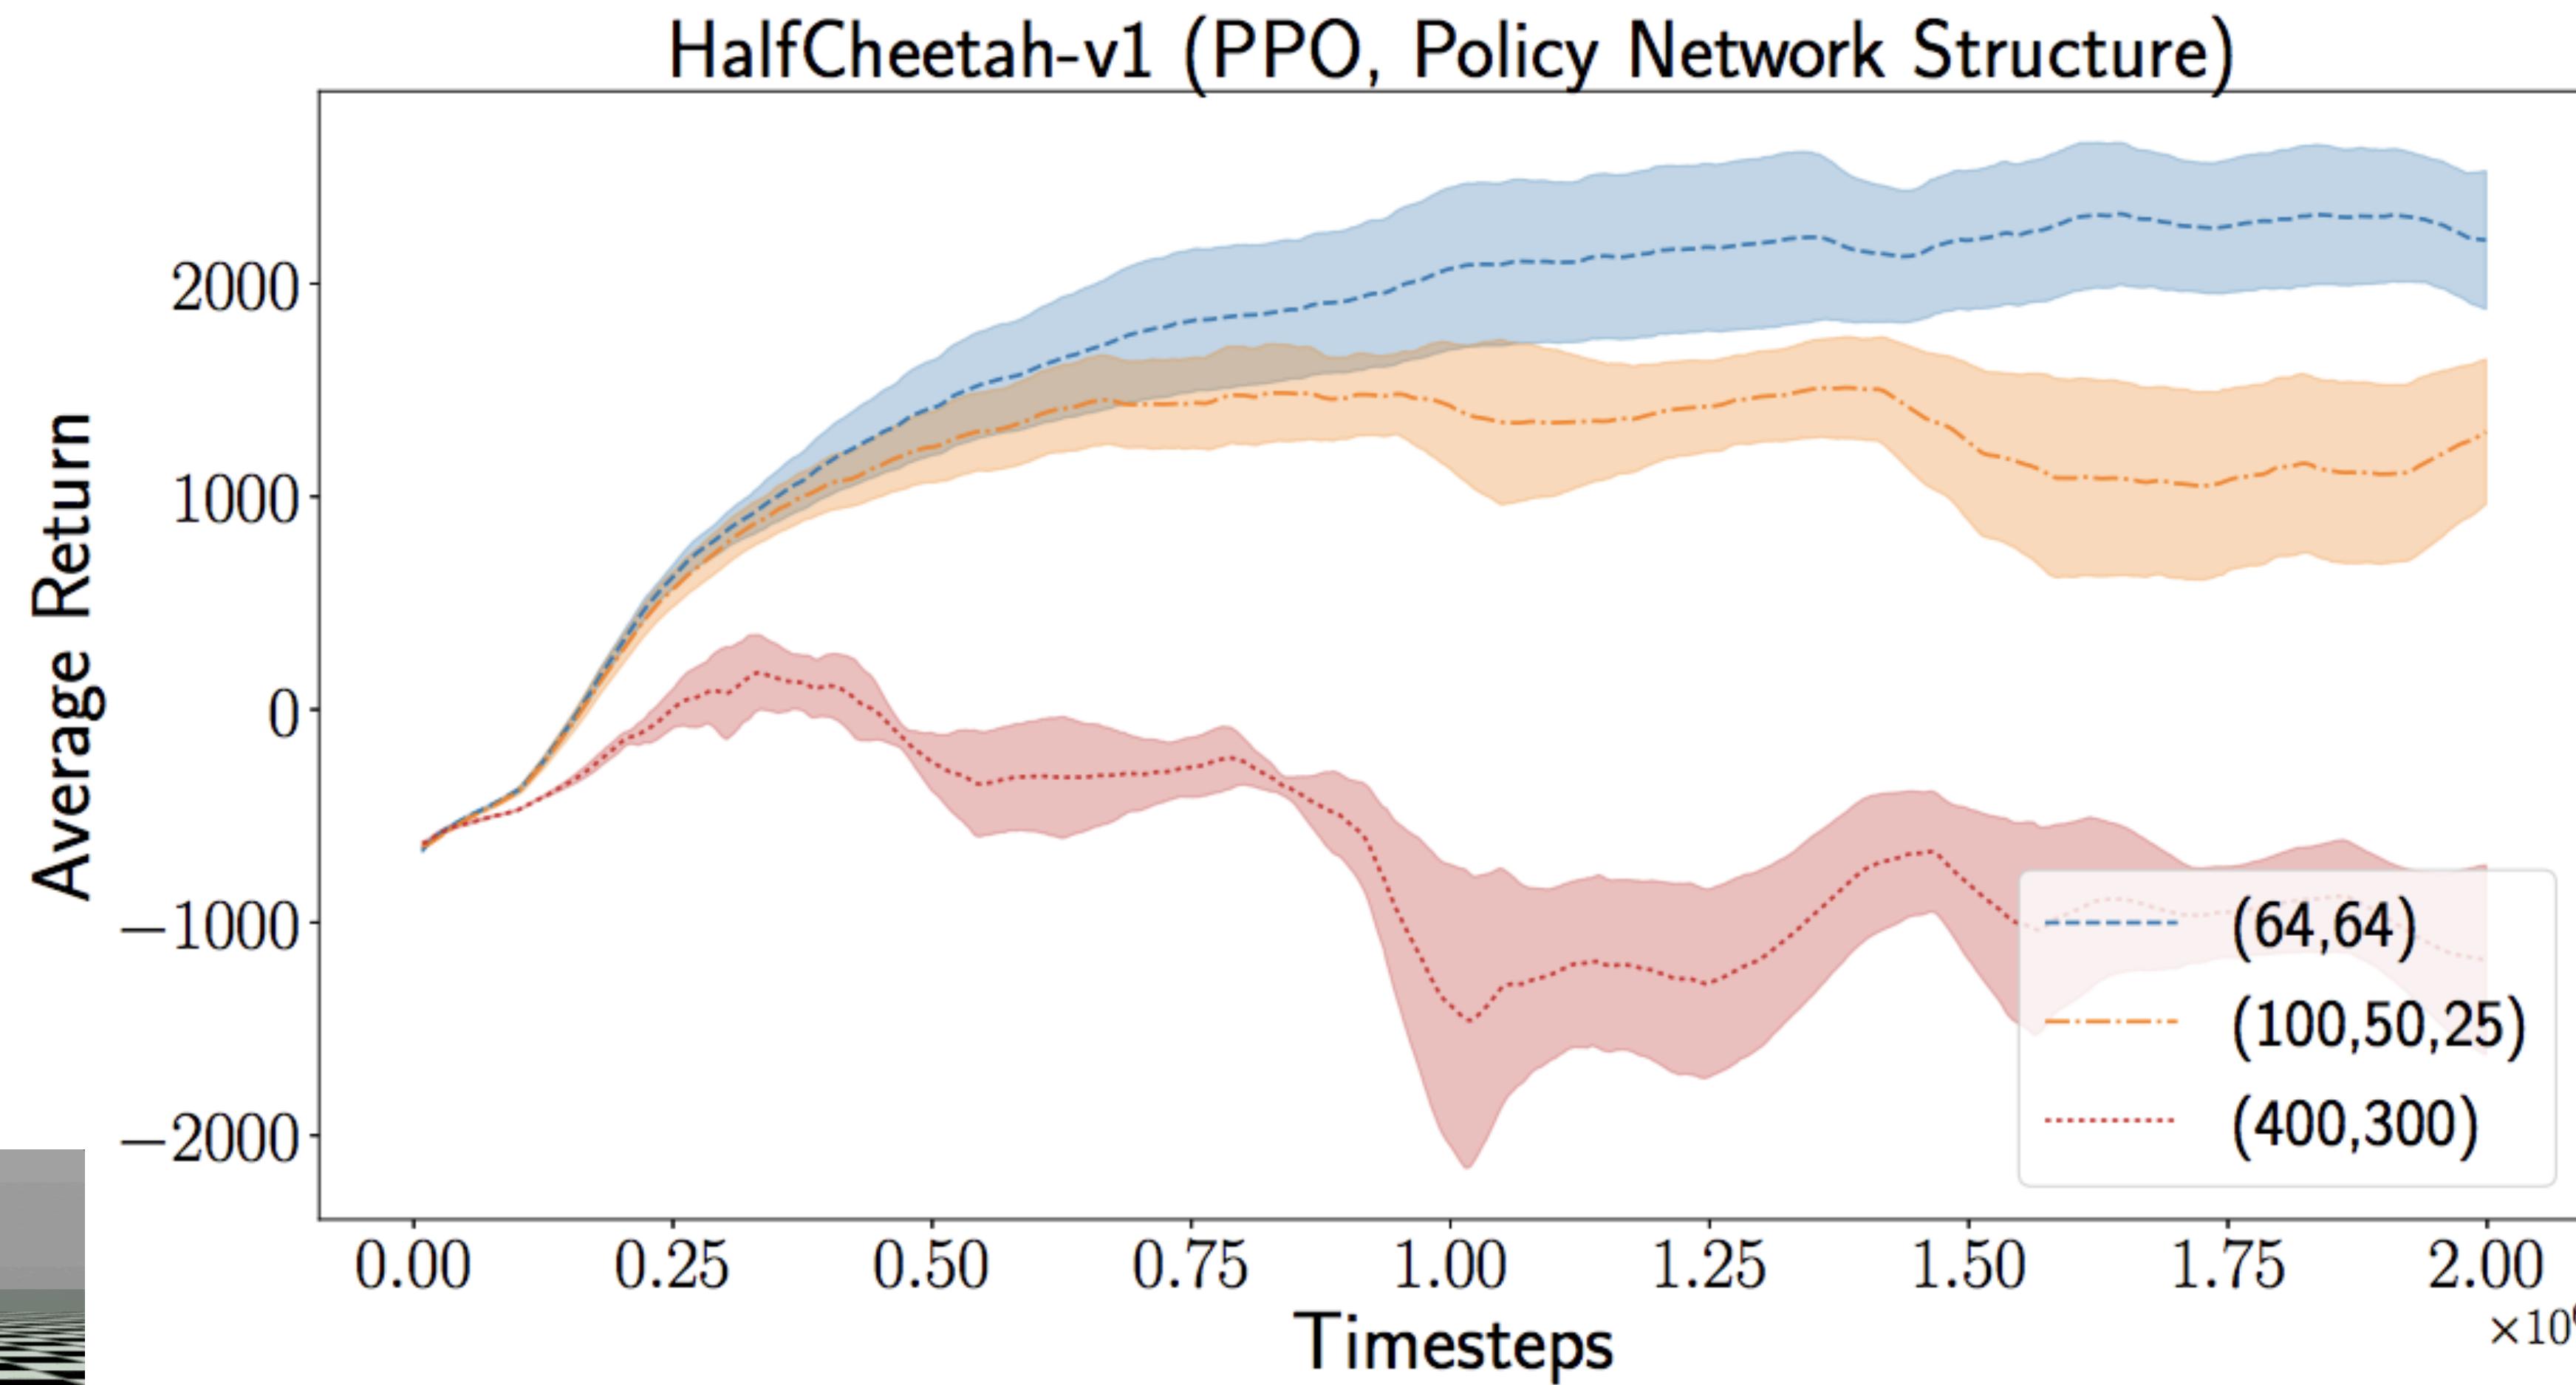
Nontrivial Example - Neural Networks

How many layers?
What kinds of layers?

Every number
shown is a
hyperparameter!



Hyperparameters Are Important for Performance

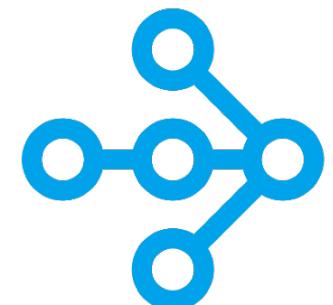
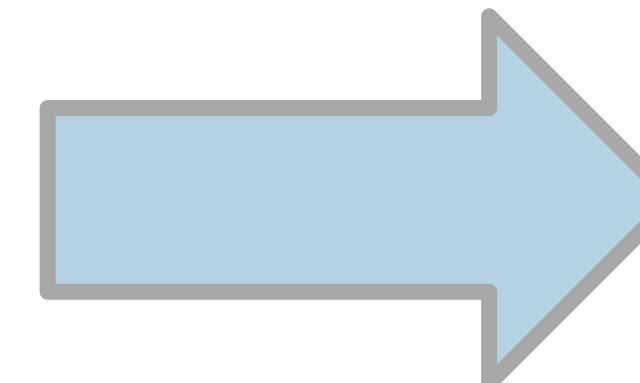


Why We Need a Framework for Tuning Hyperparameters

We want the best model

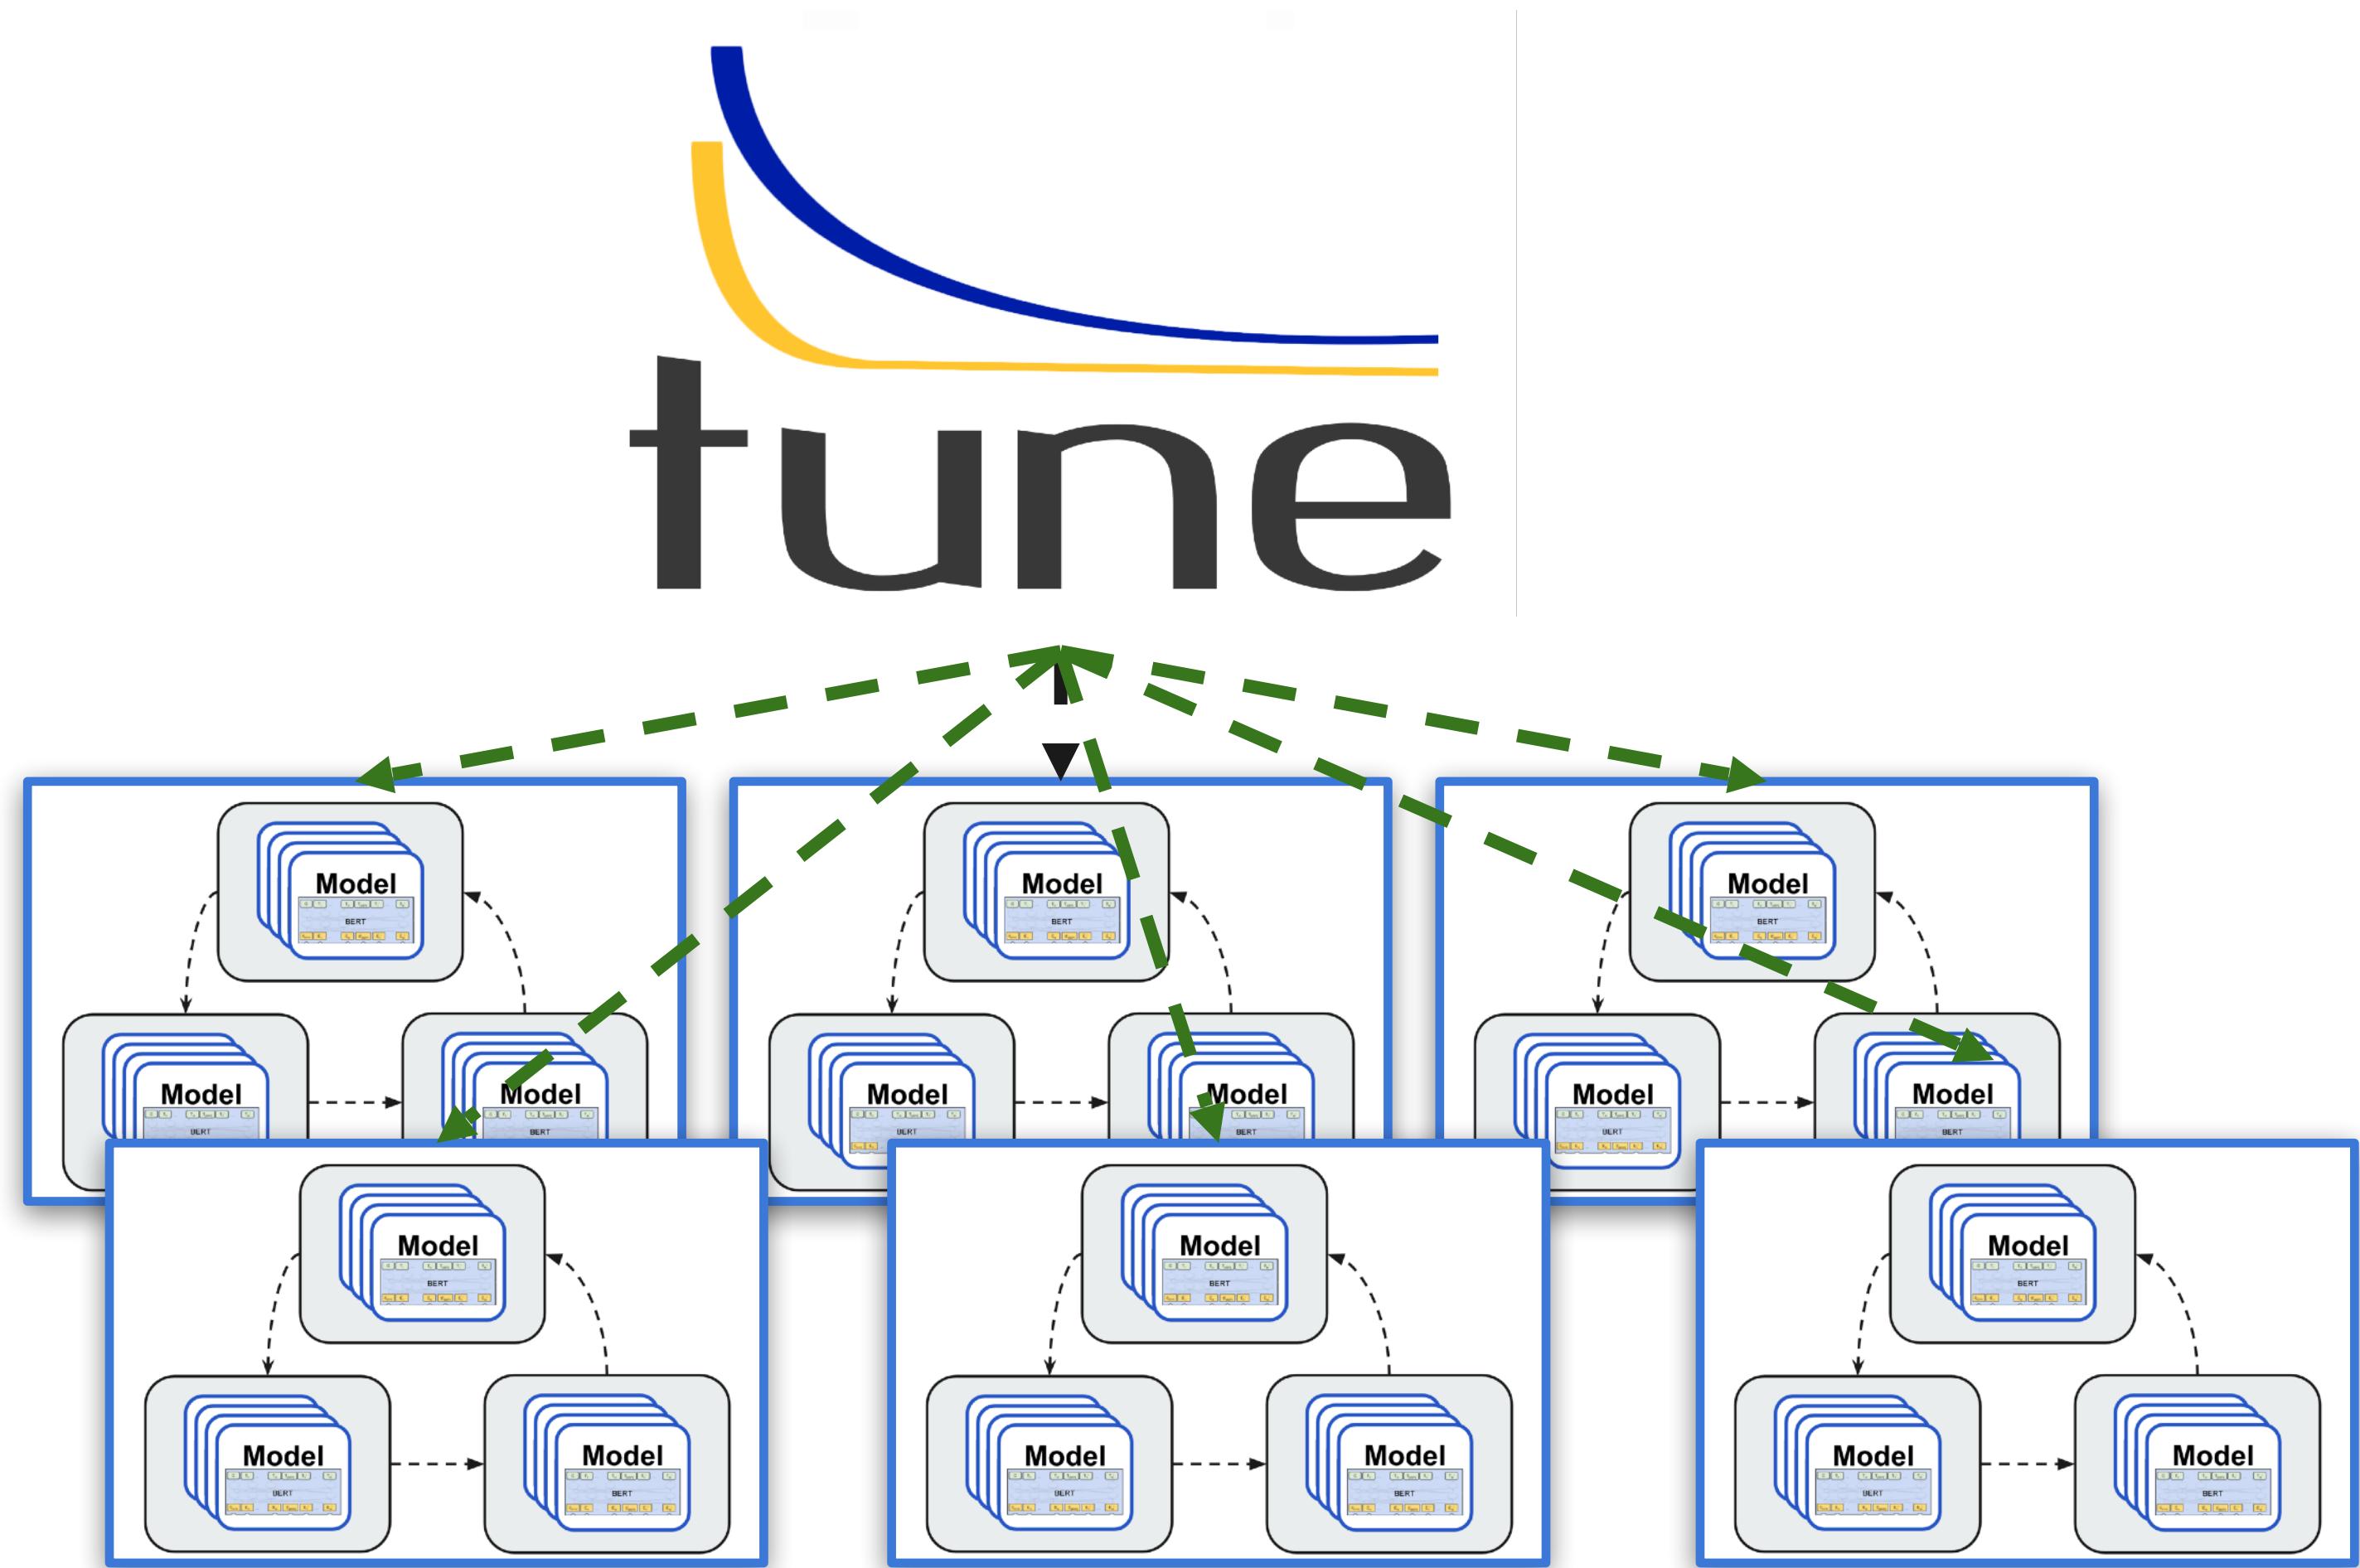
Resources are expensive

Model training is time-consuming

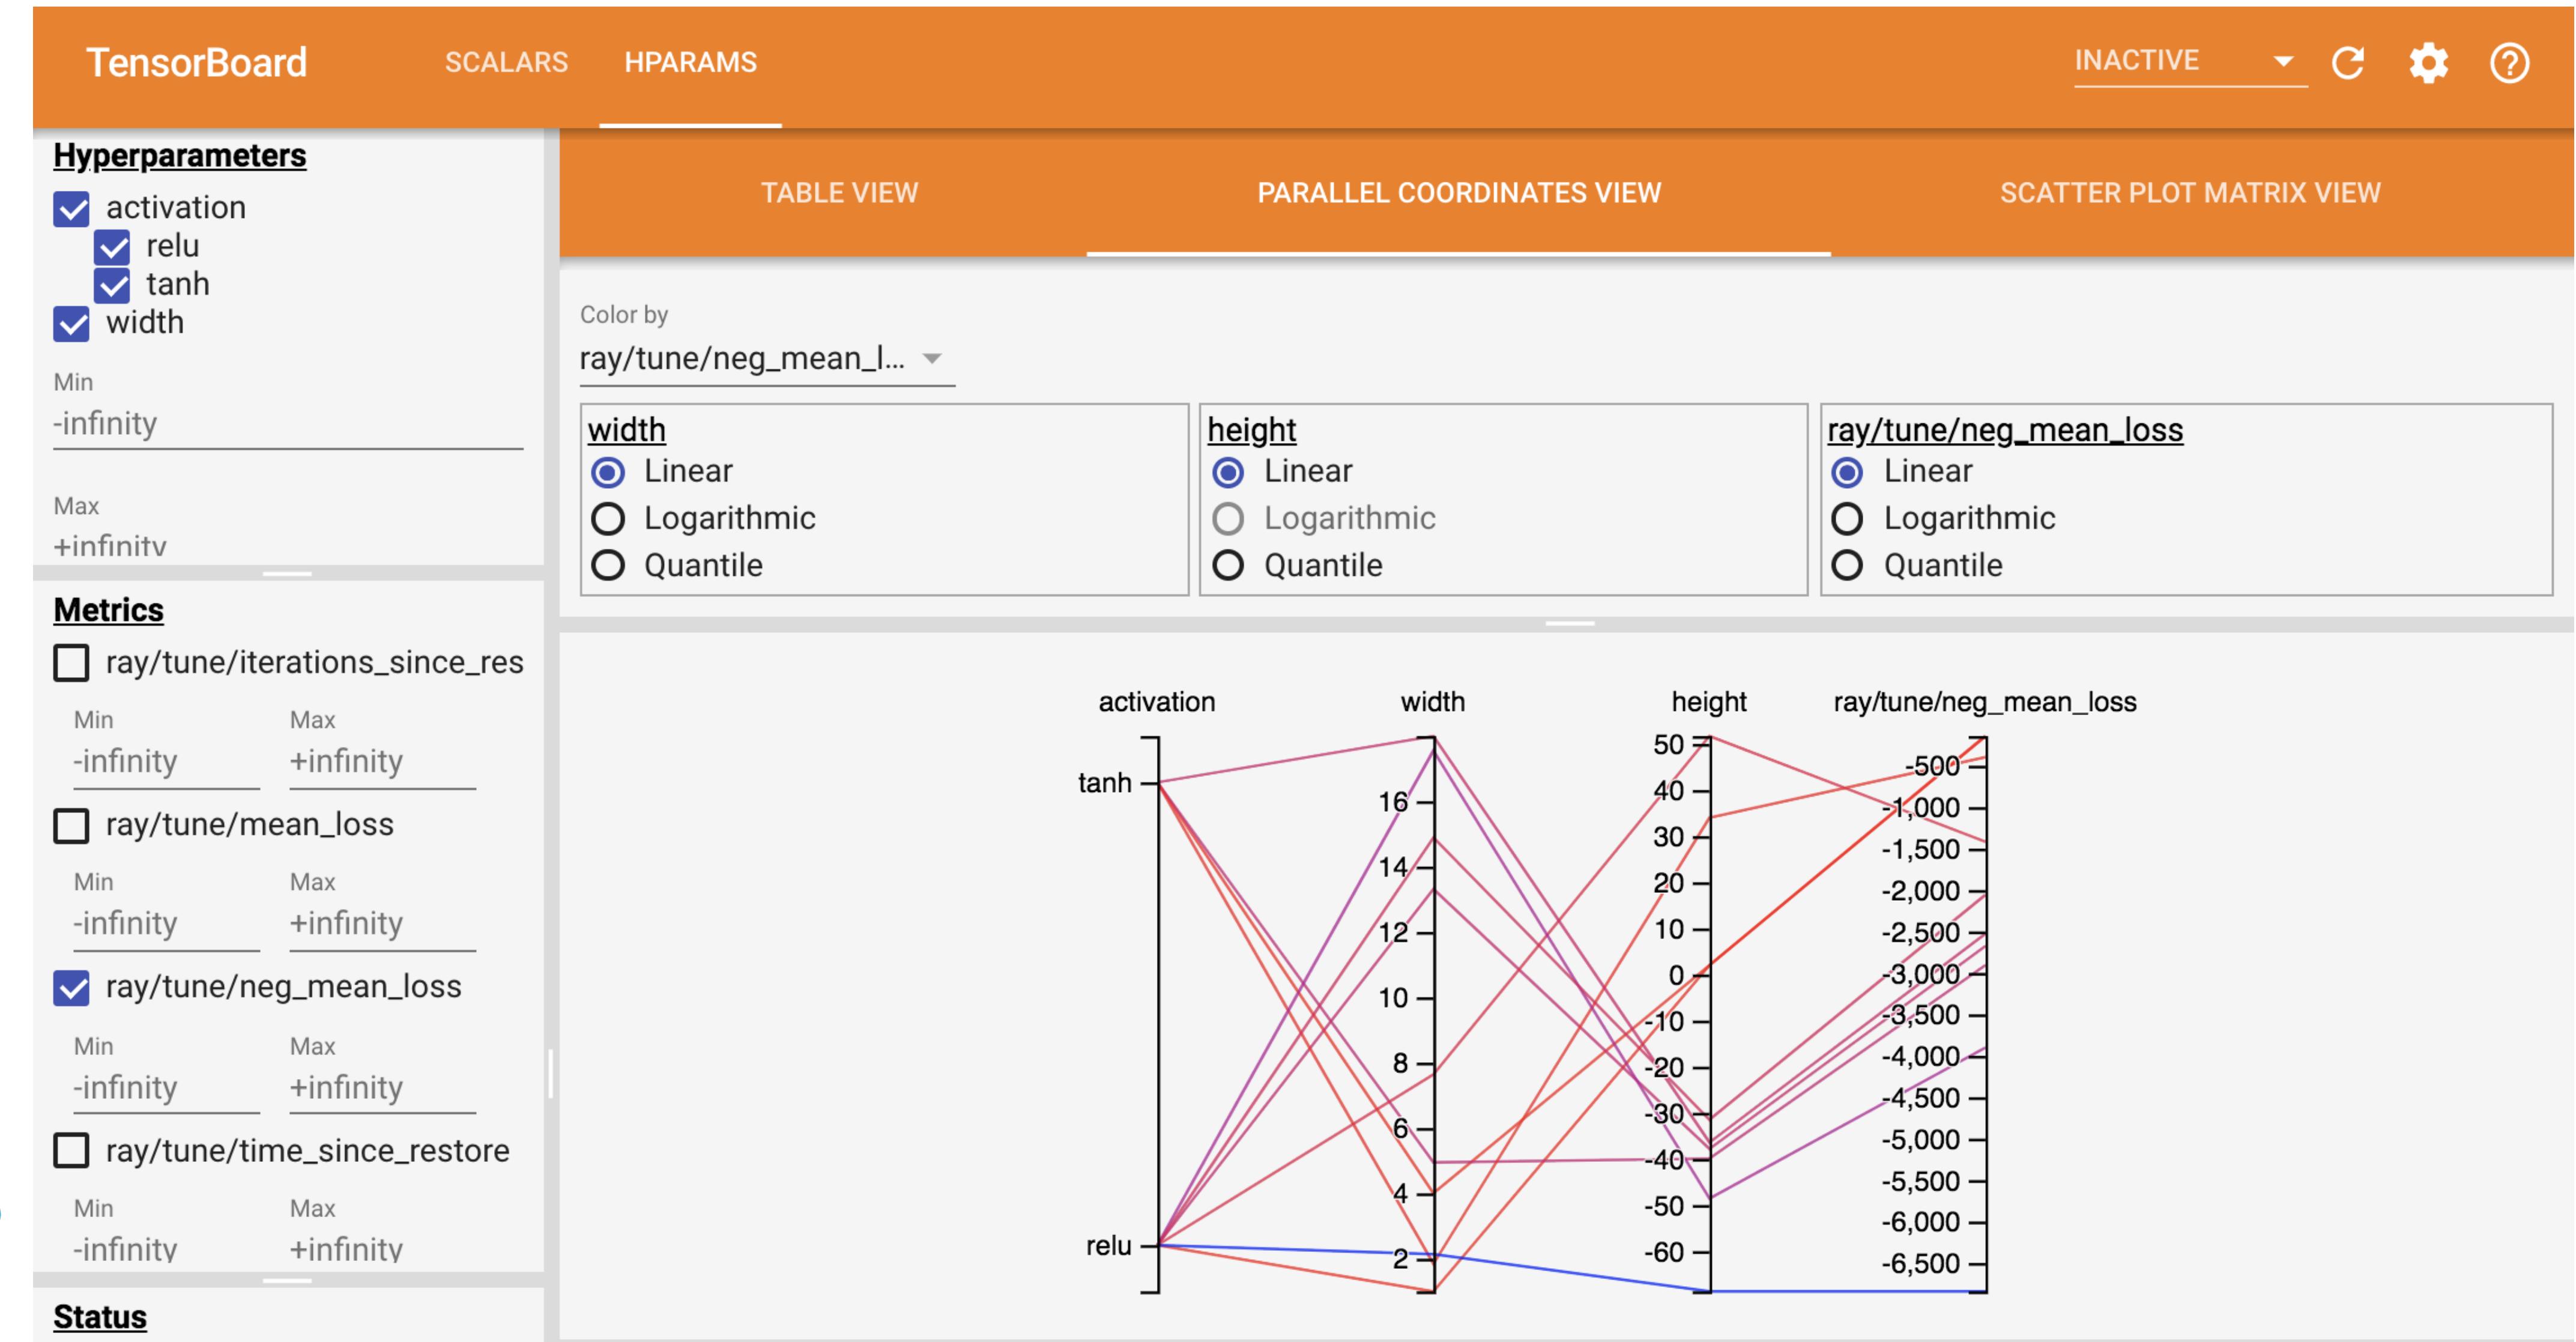


Tuning + Distributed Training

```
tune.run(PytorchTrainable,  
 config={  
     "model_creator": PretrainBERT,  
     "data_creator": create_data_loader,  
     "use_gpu": True,  
     "num_replicas": 8,  
     "lr": tune.uniform(0.001, 0.1)  
},  
    num_samples=100,  
    search_alg=BayesianOptimization()
```



Native Integration with TensorBoard HParams

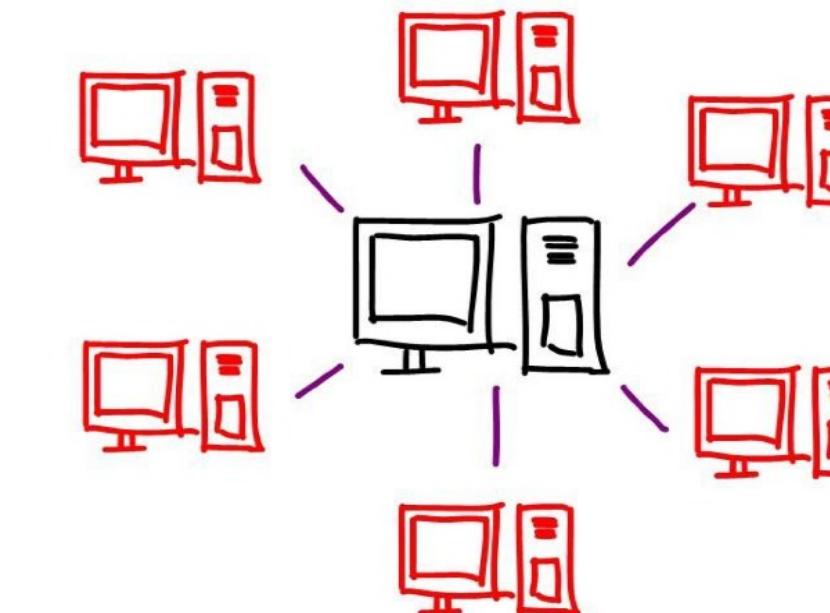


Tune is Built with Deep Learning as a Priority

Resource Aware
Scheduling



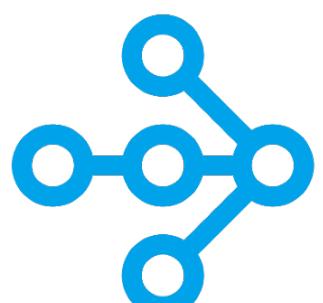
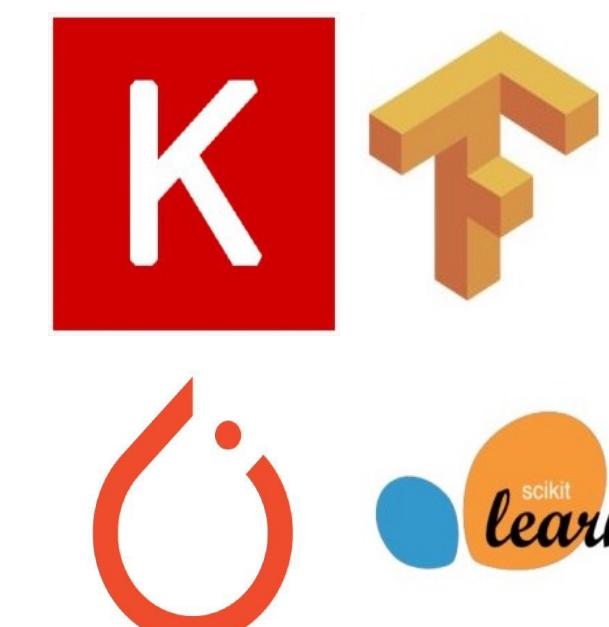
Seamless
Distributed Execution



Simple API for
new algorithms

```
class TrialScheduler:  
    def on_result(self, trial, result): ...  
    def choose_trial_to_run(self): ...
```

Framework Agnostic



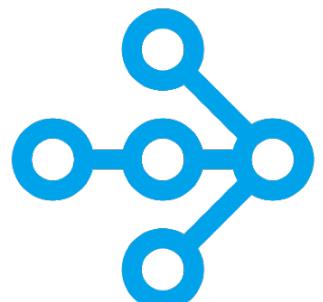
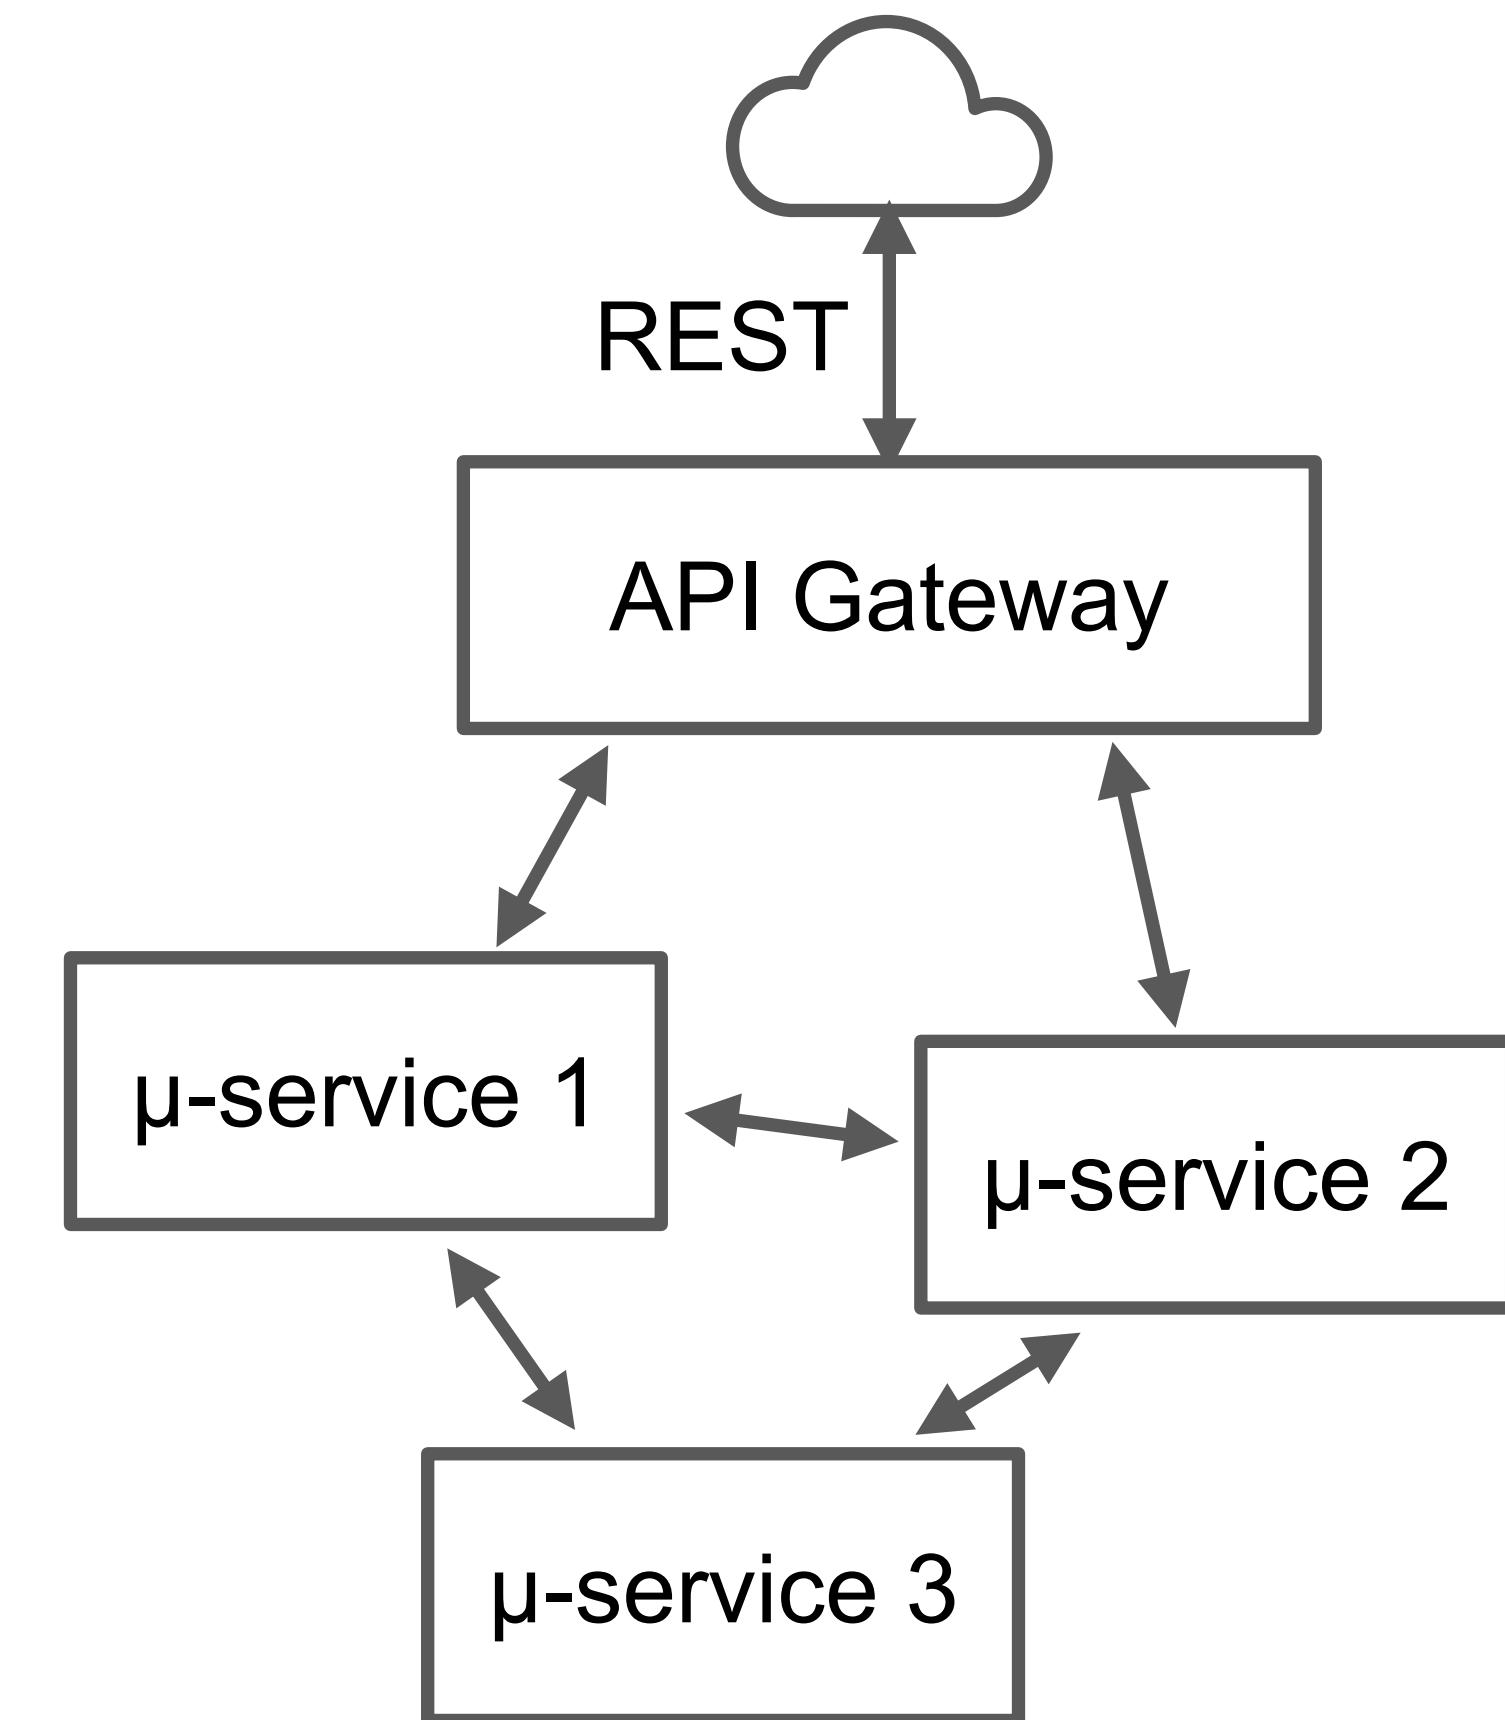
ray.readthedocs.io/en/latest/tune.html



what about Ray
for Microservices?

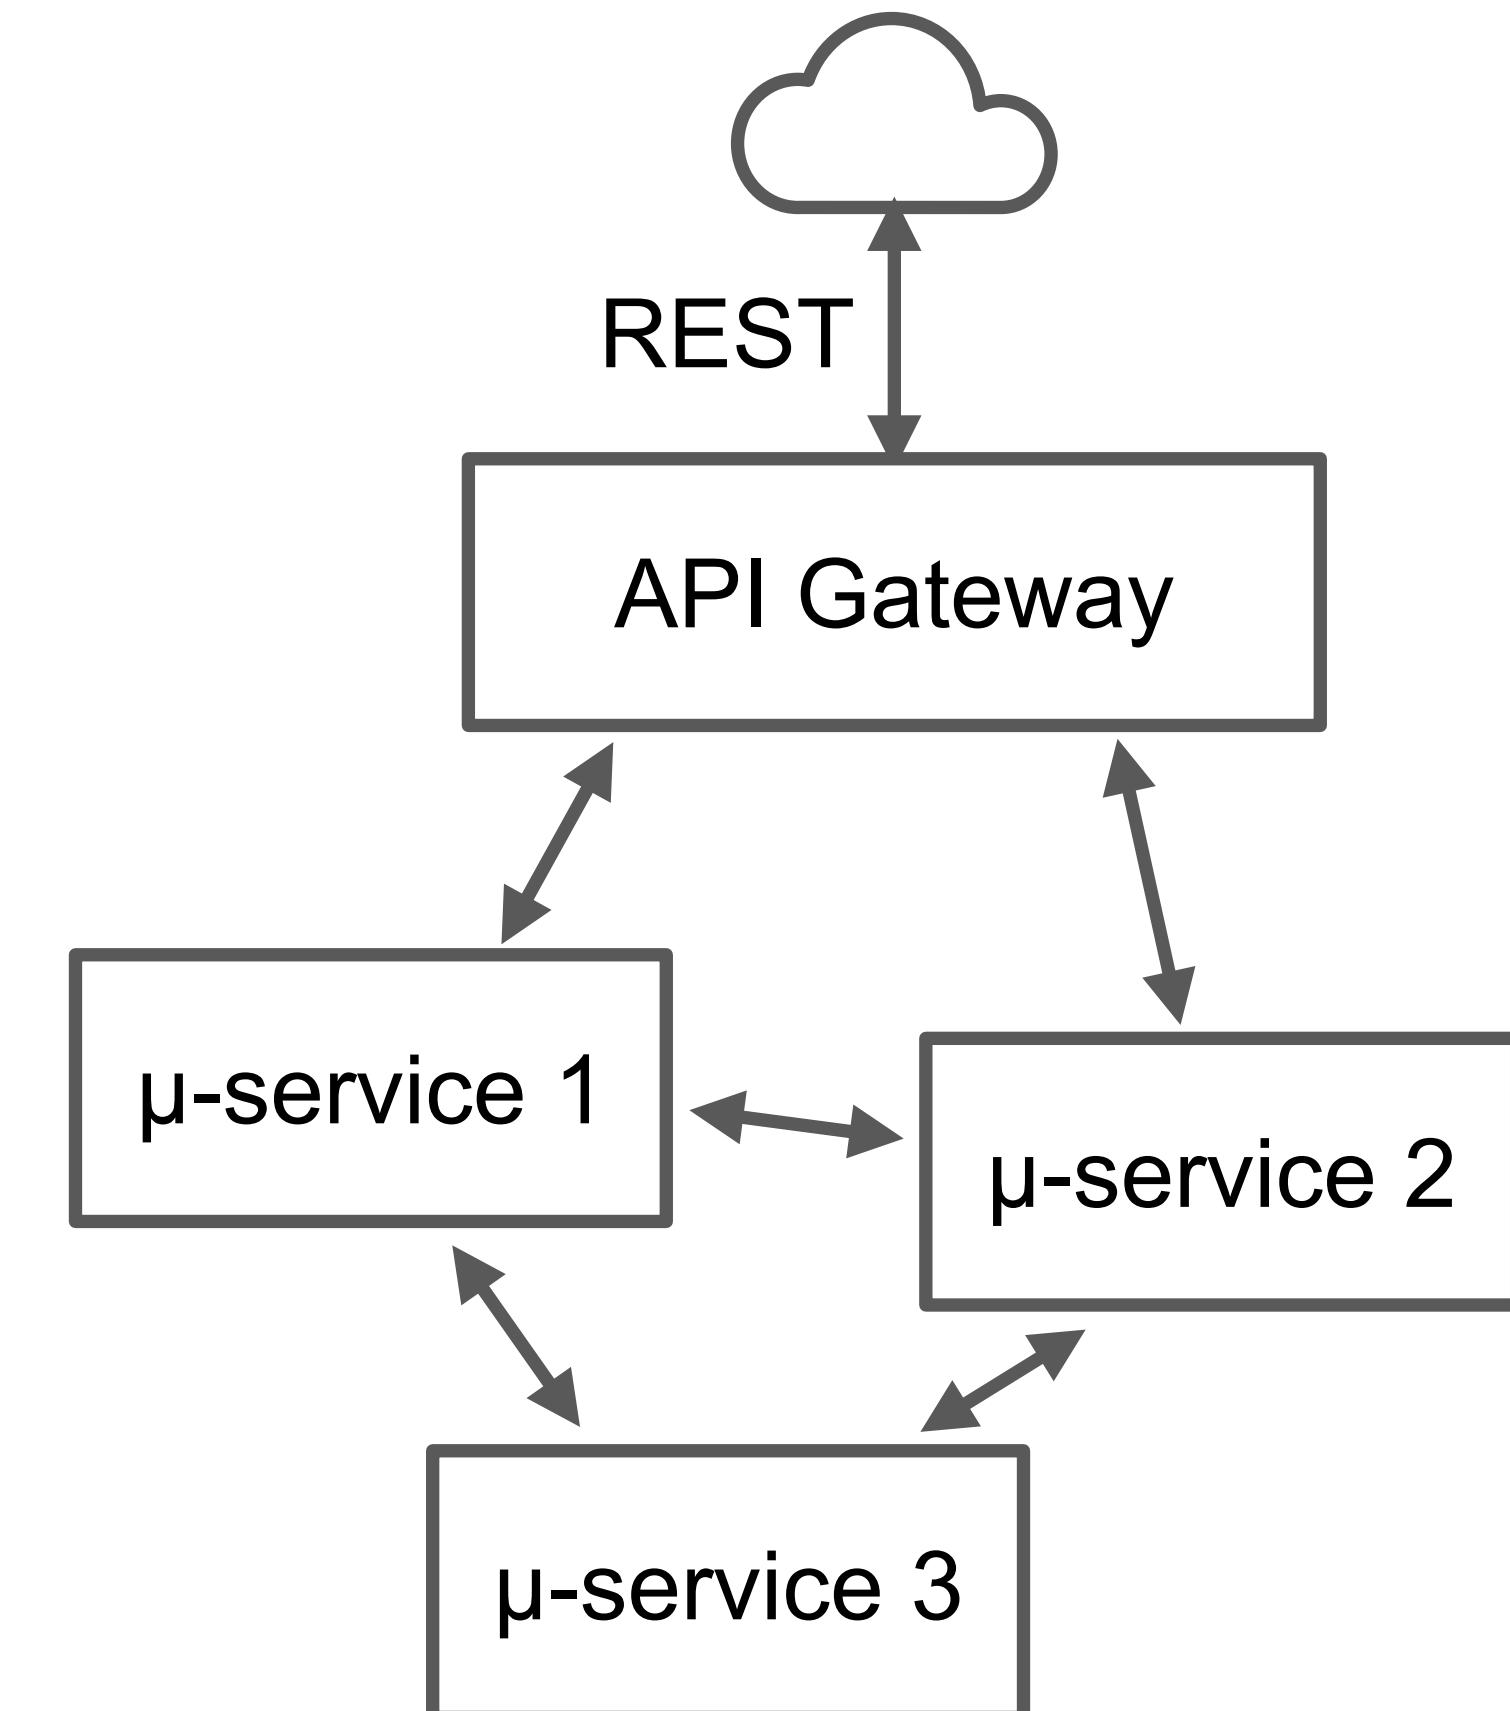
What Are Microservices?

- They partition the domain
 - Conway's Law - Embraced
 - Separate responsibilities
- Separate management

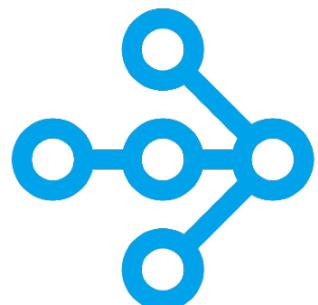


Conway's Law - Embraced

- “Any organization that designs a system will produce a design whose structure is a copy of the organization's communication structure”
- Let each team own and manage the services for its part of the domain

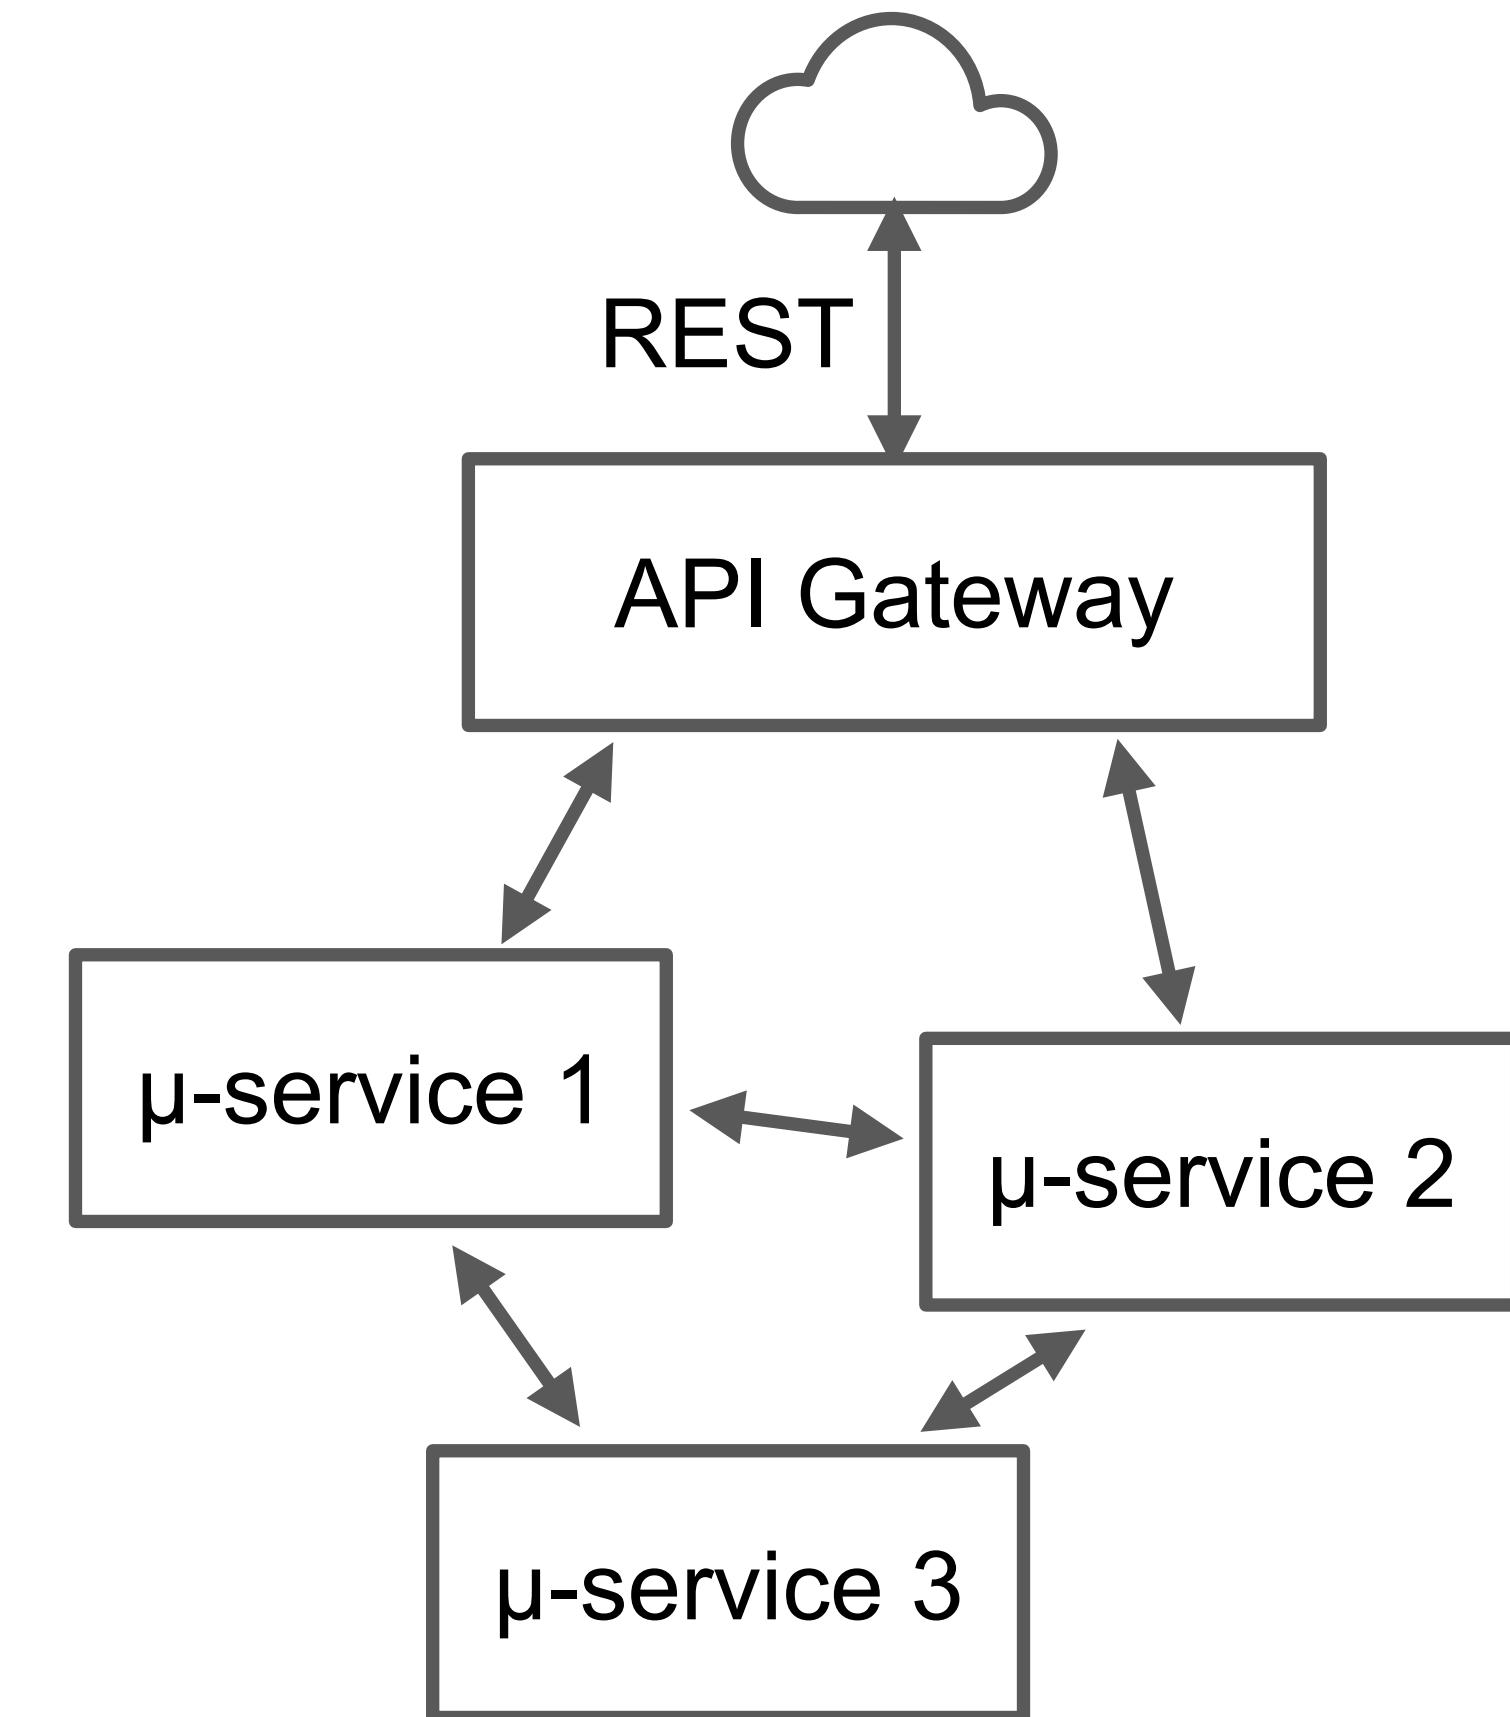


en.wikipedia.org/wiki/Conway's_law

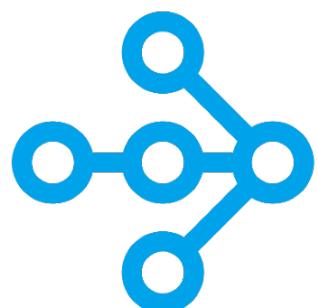


Separate Responsibilities

- Each microservice does “one thing”, a single responsibility with minimal coupling to the other microservices
- (Like, hopefully, the teams are organized, too...)

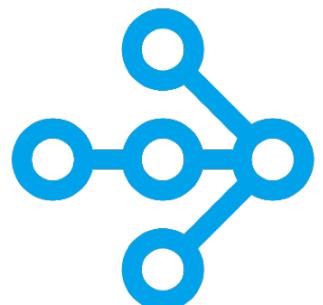
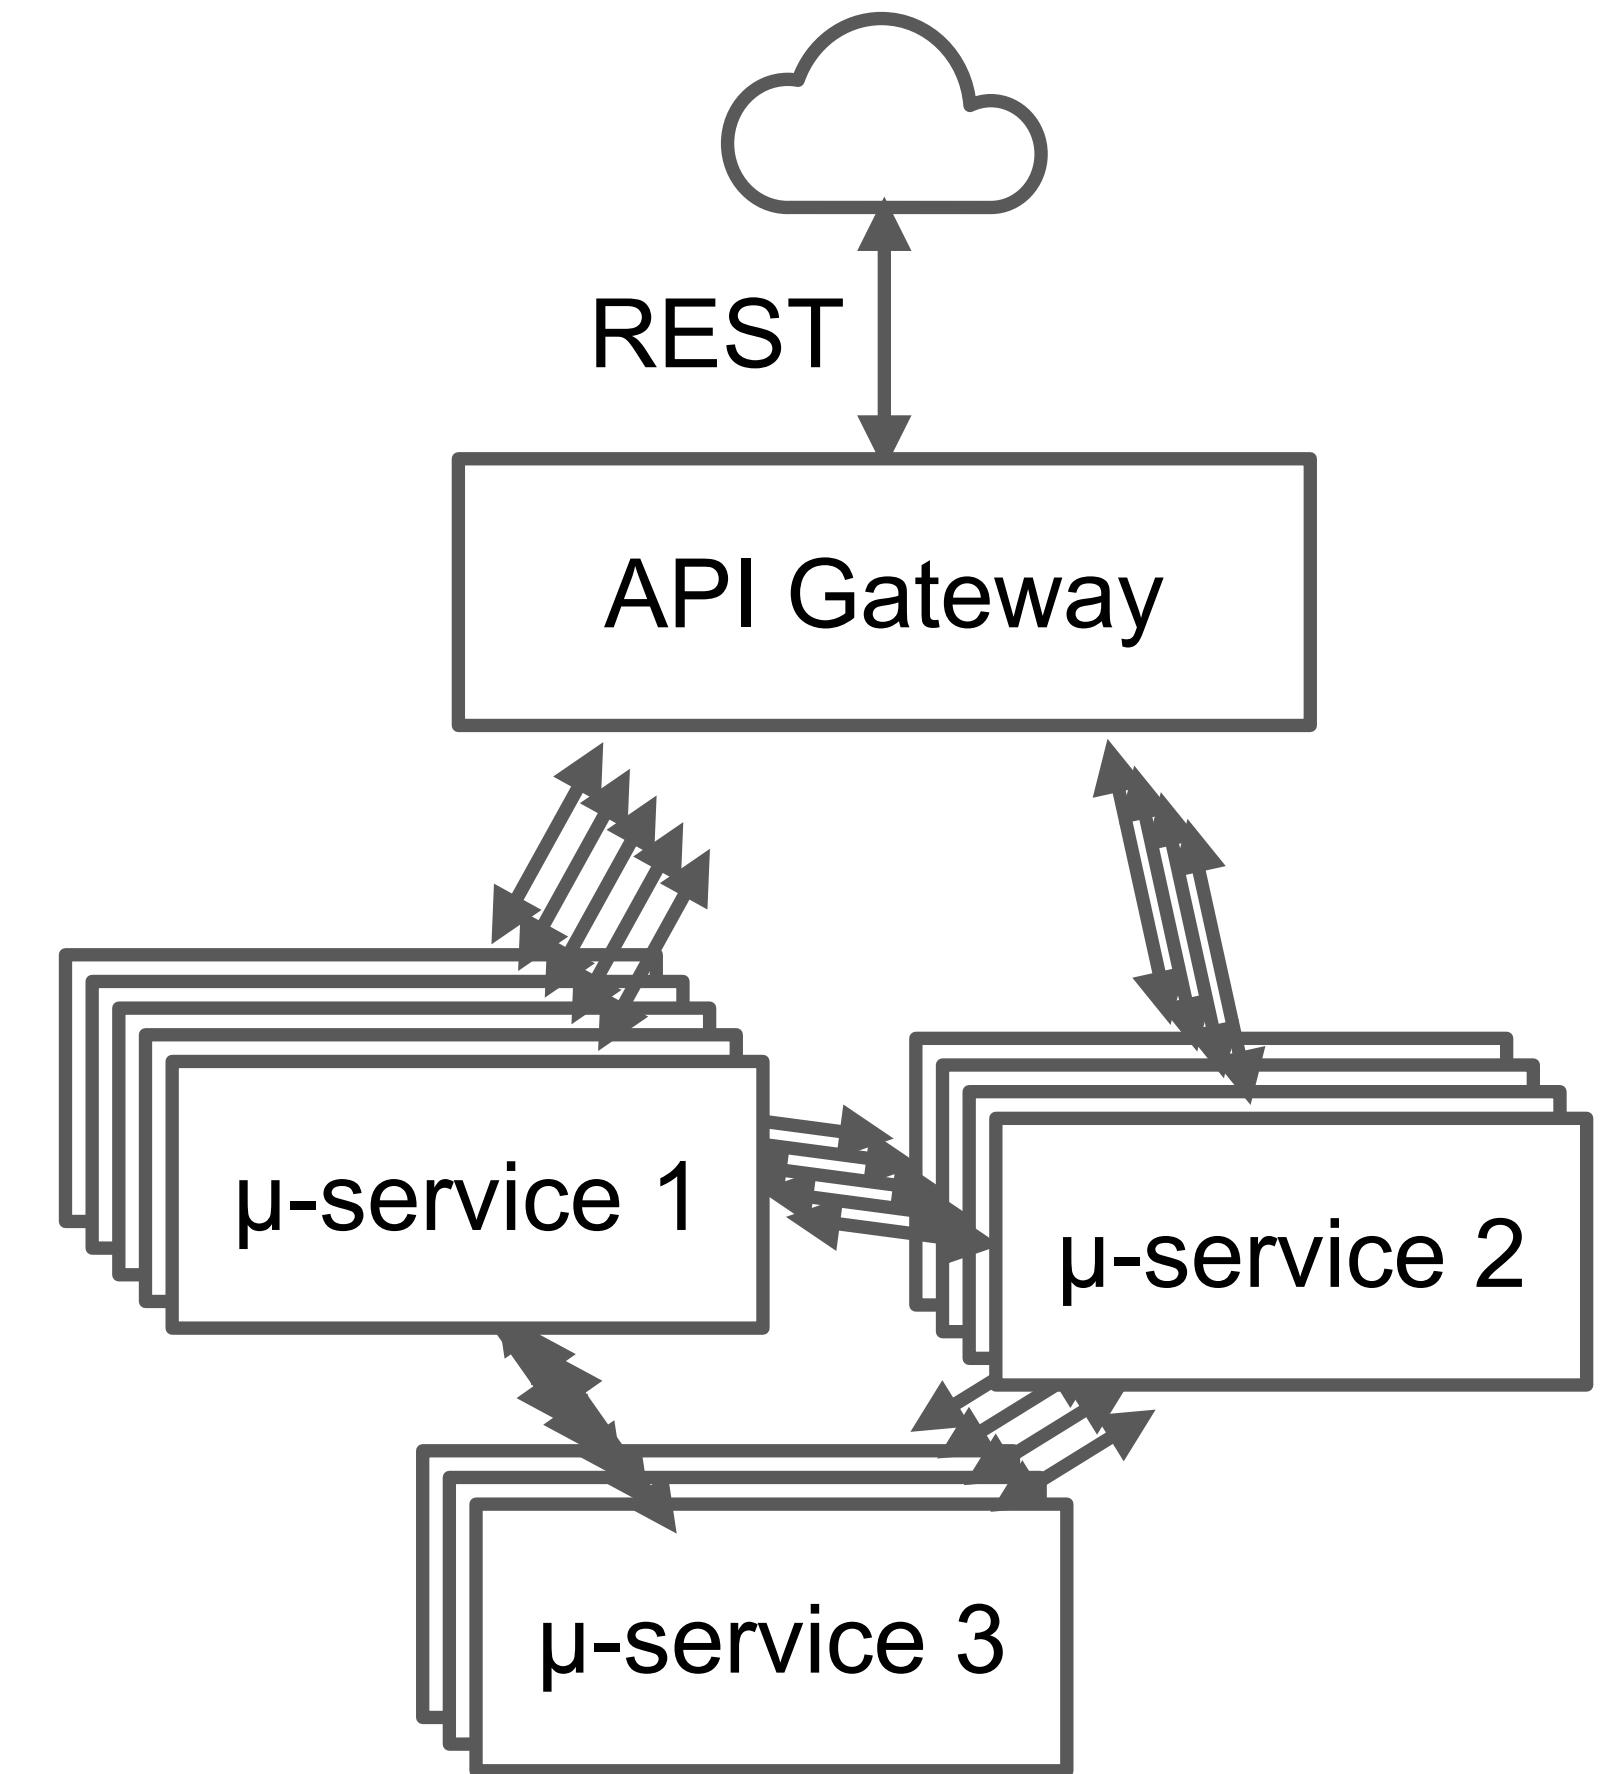


wikipedia.org/wiki/Single-responsibility_principle



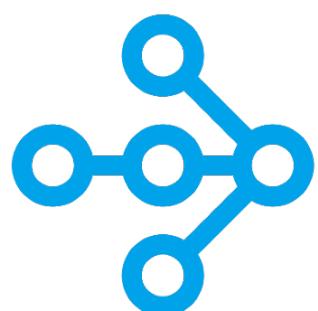
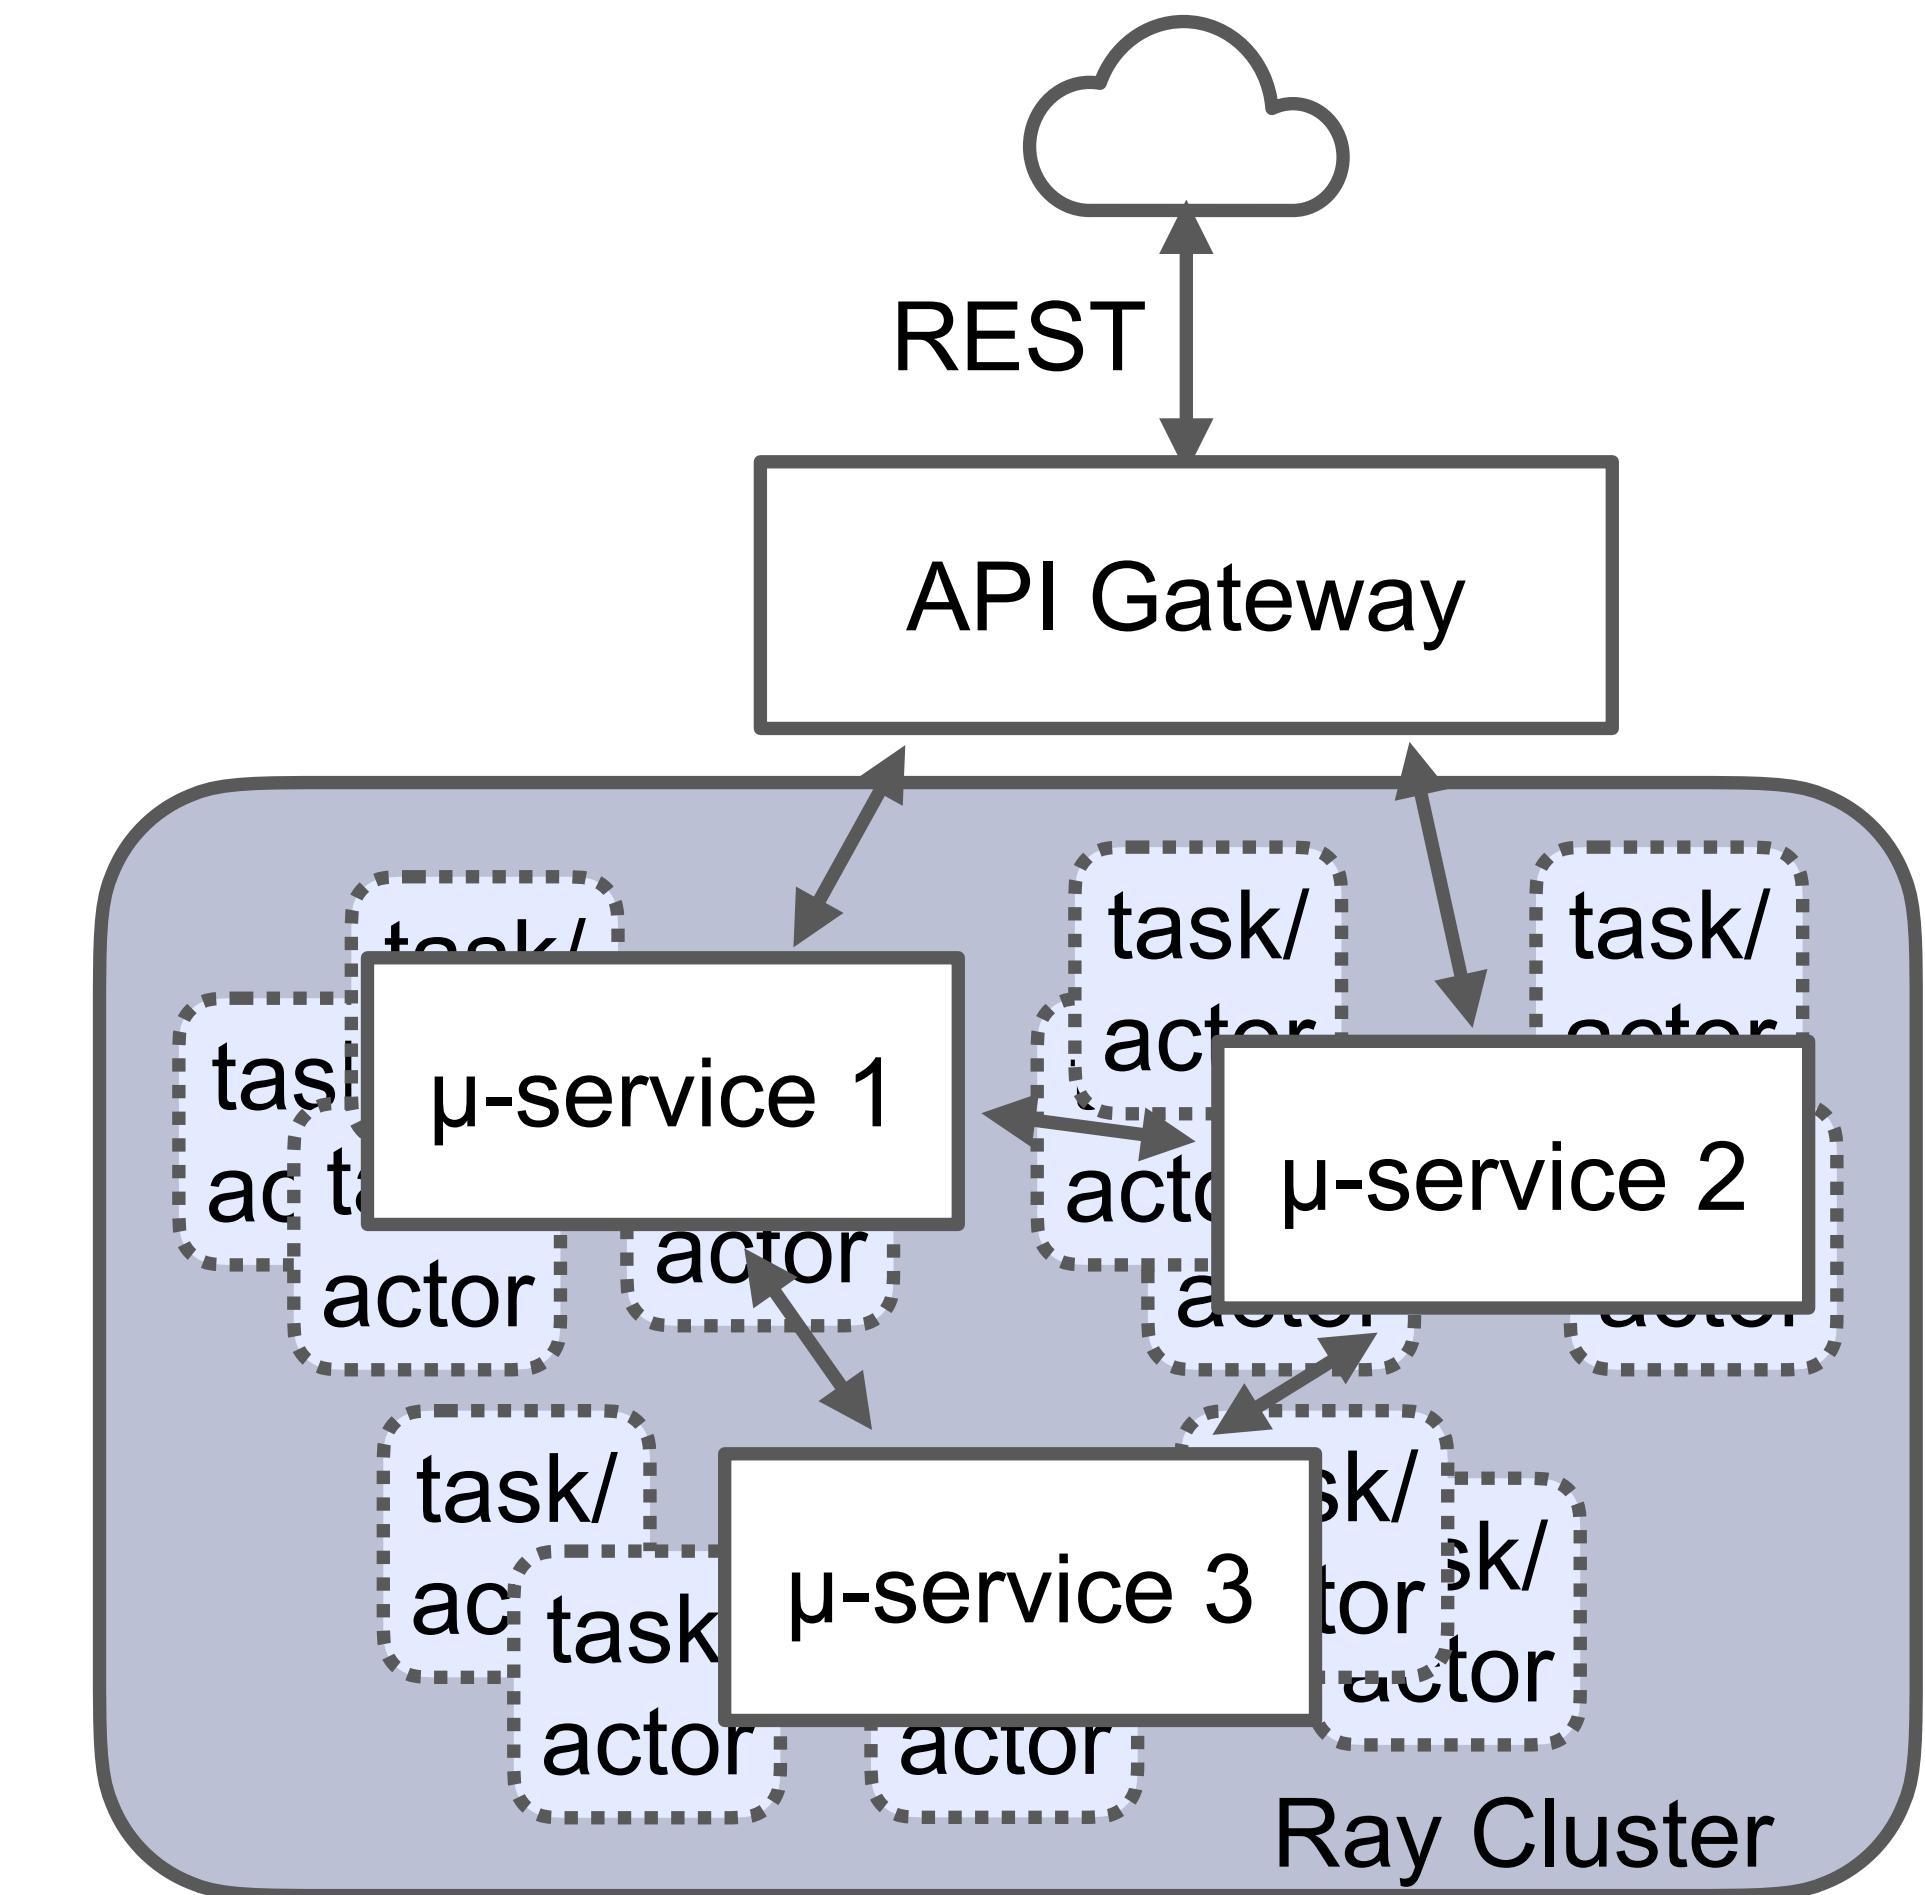
Separate Management

- Each team manages its own instances
- Each microservice has a different number of instances for scalability and resiliency
- But they have to be managed **explicitly**



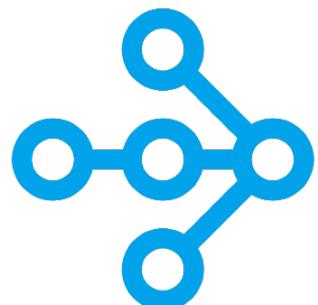
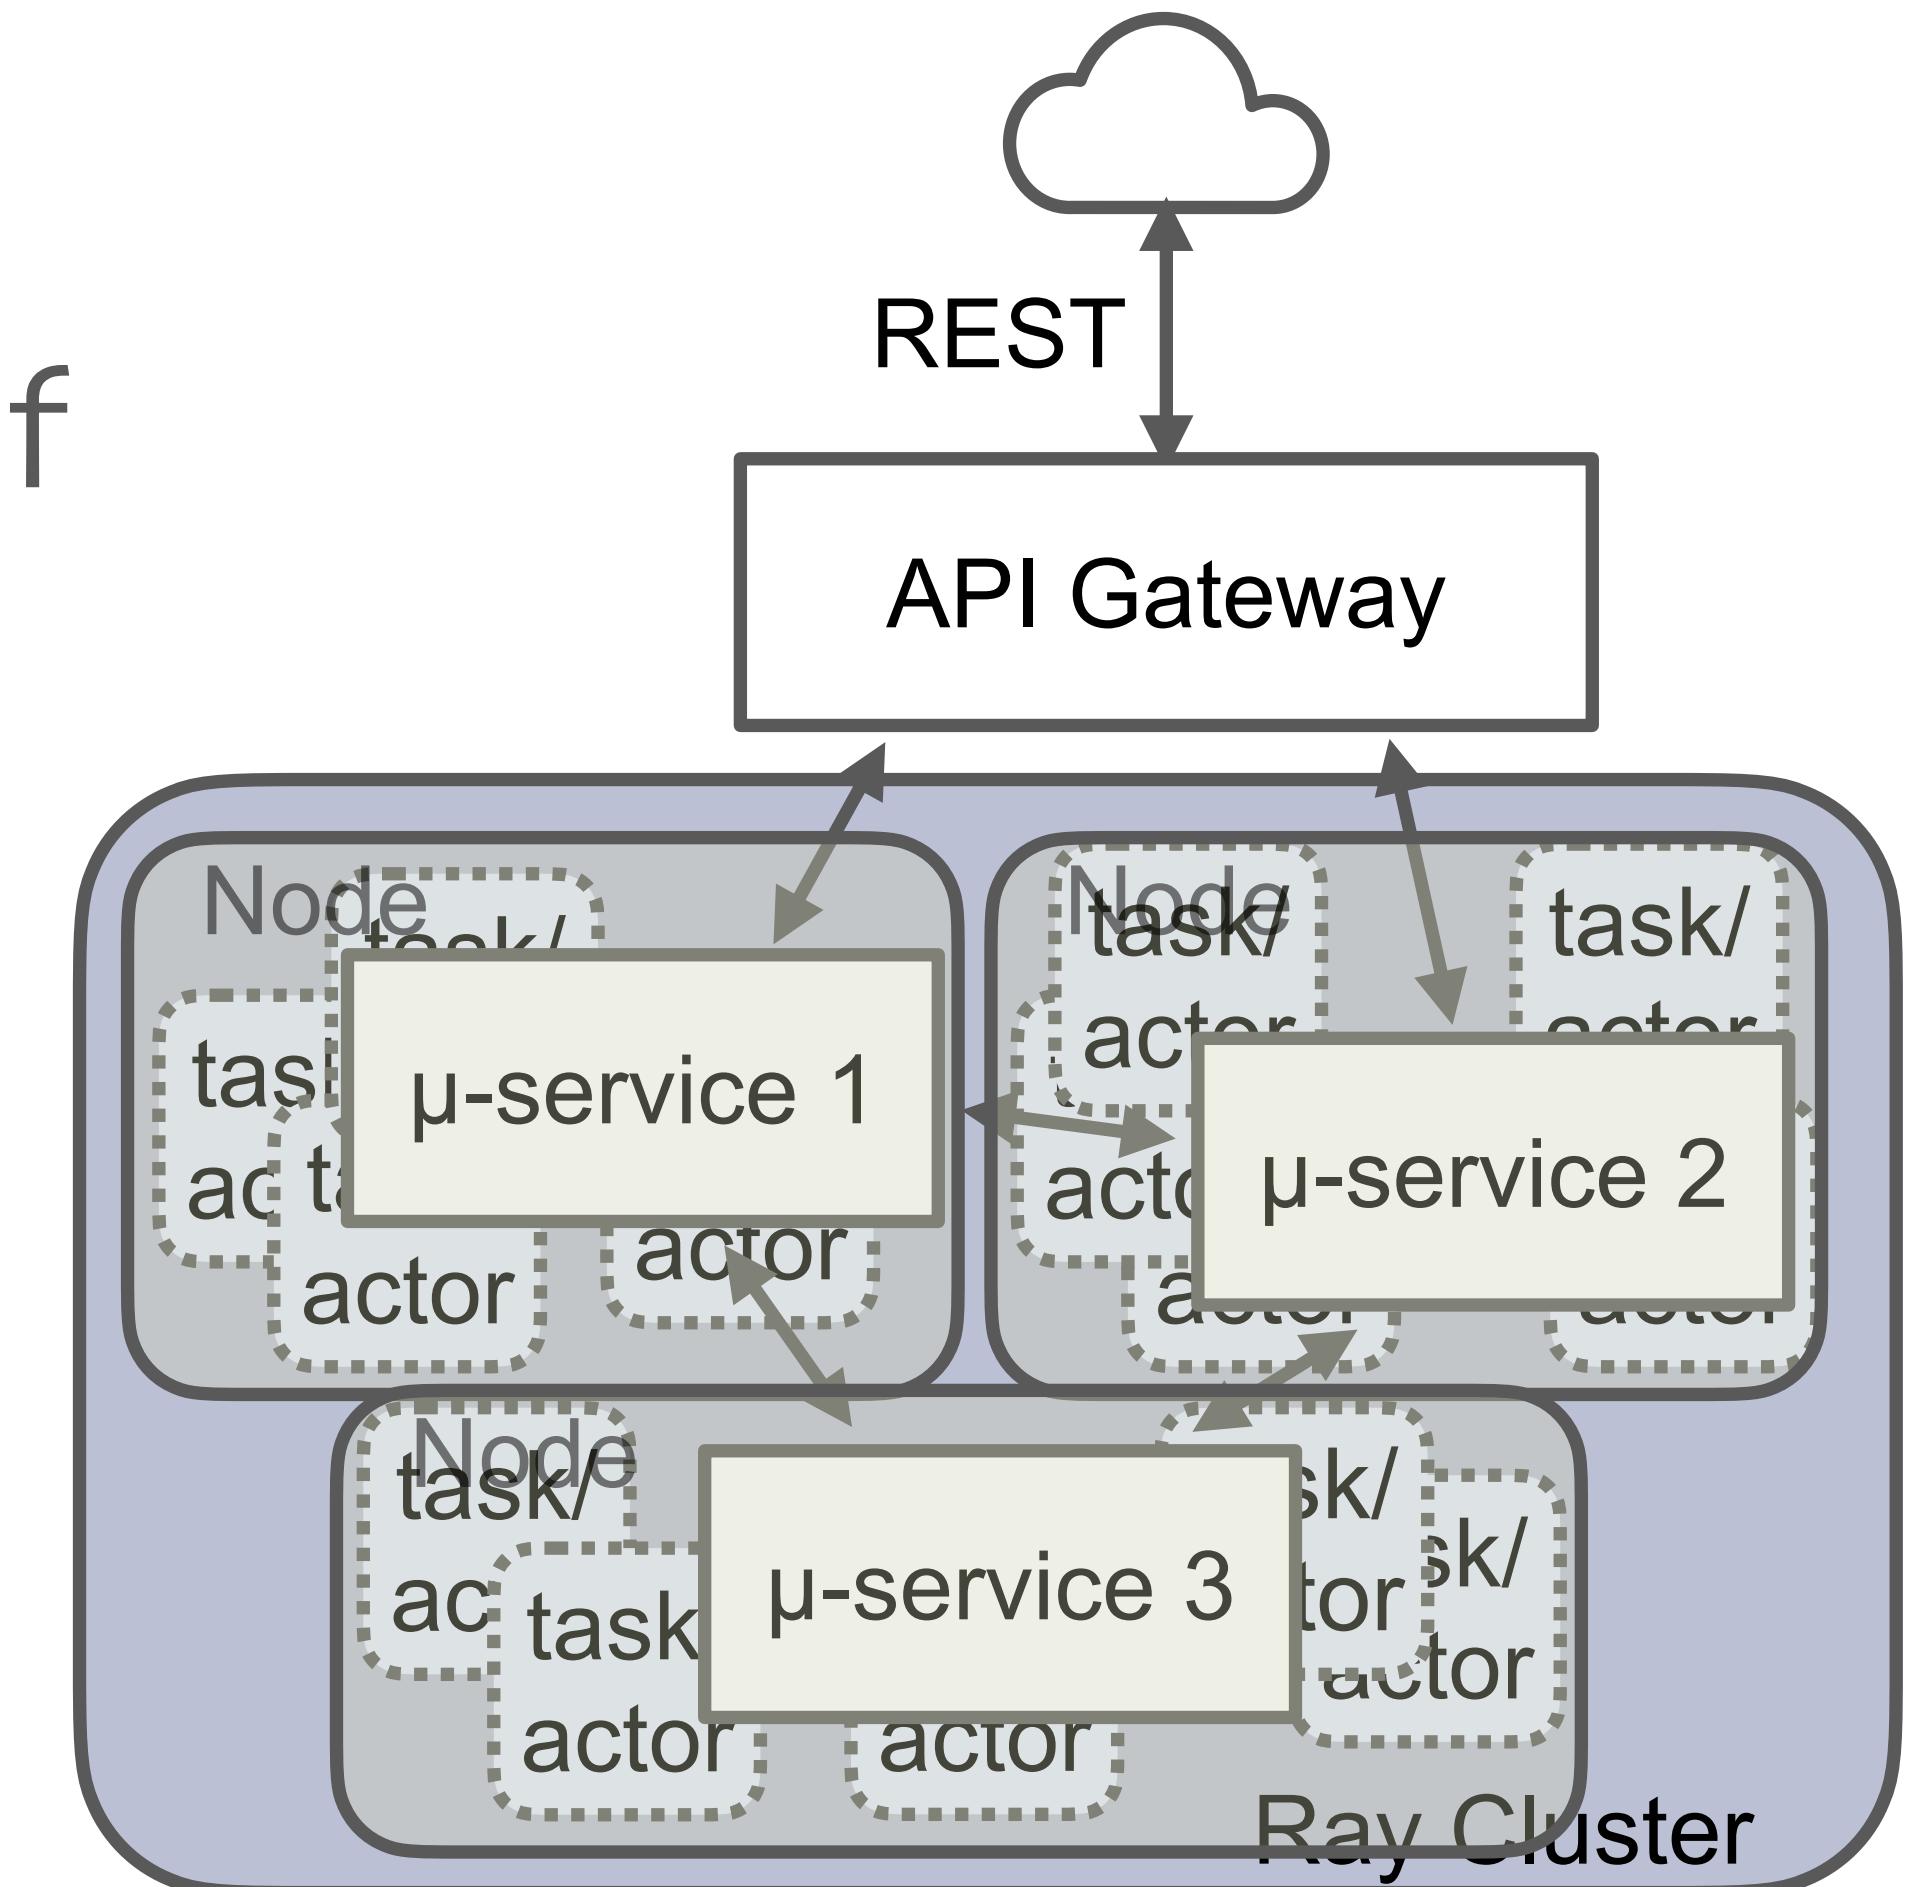
Management - Simplified

- With Ray, you have one “logical” instance to manage and Ray does the cluster-wide scaling for you.



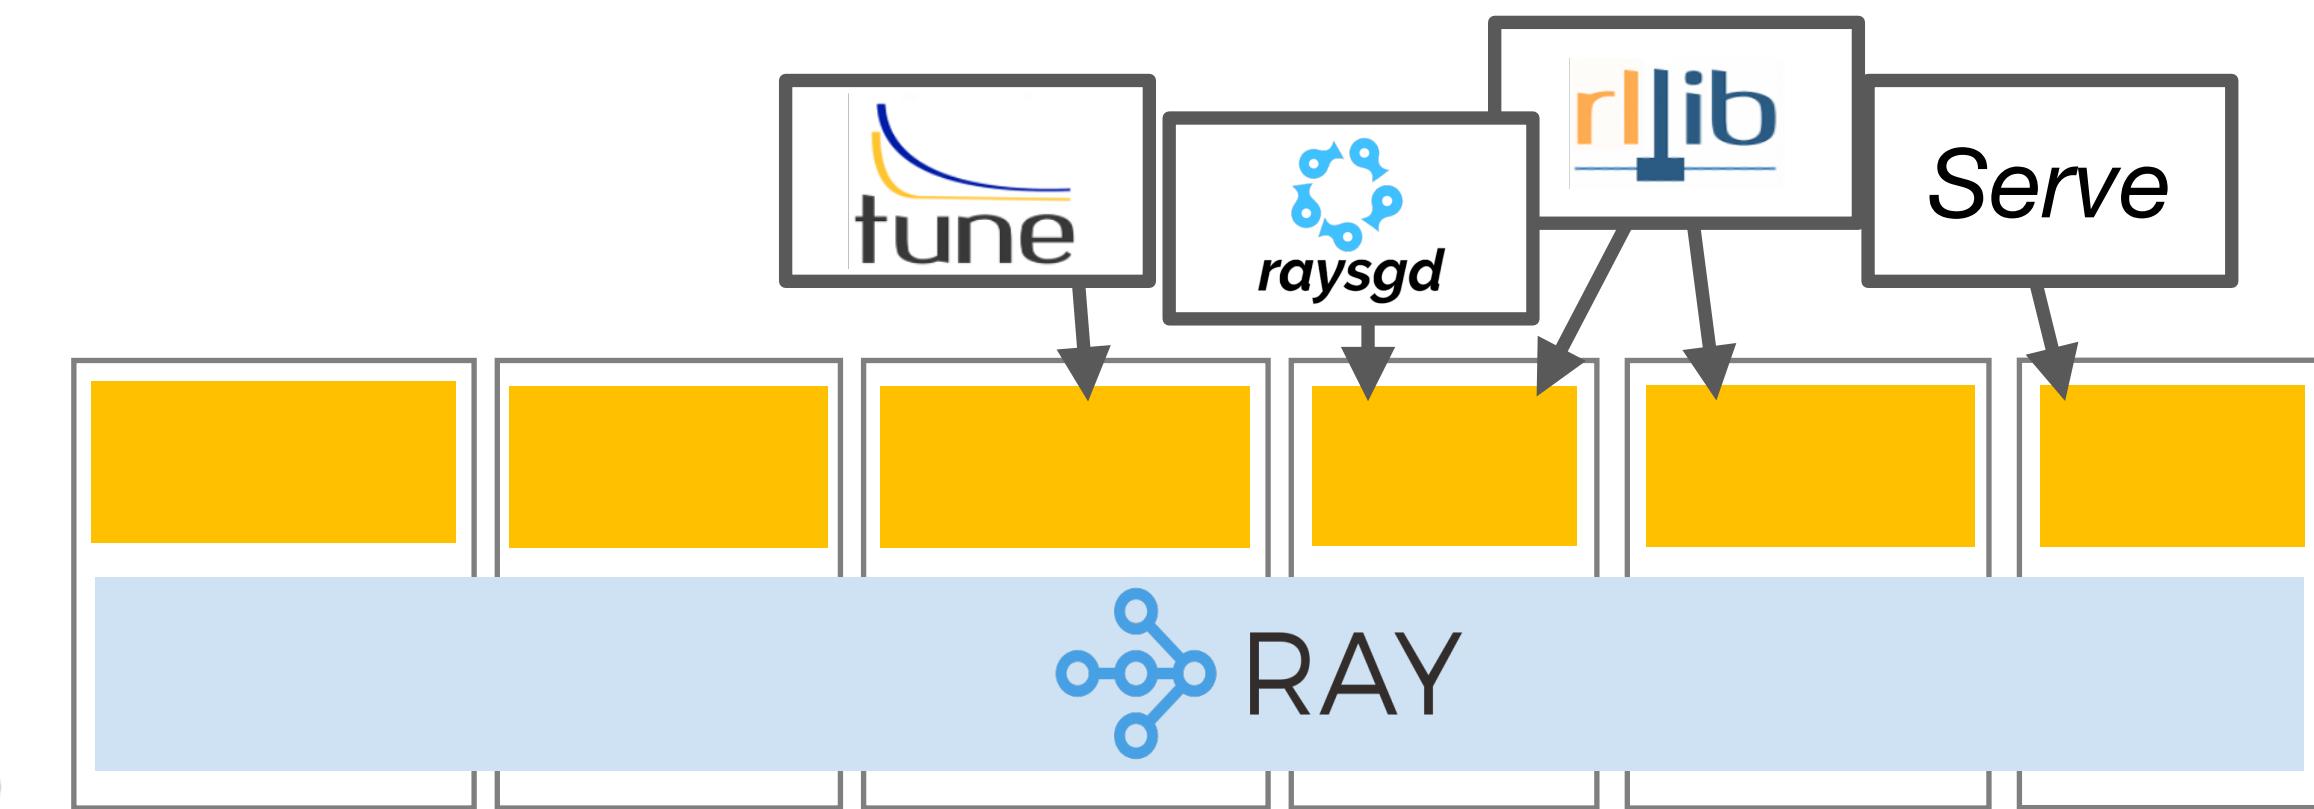
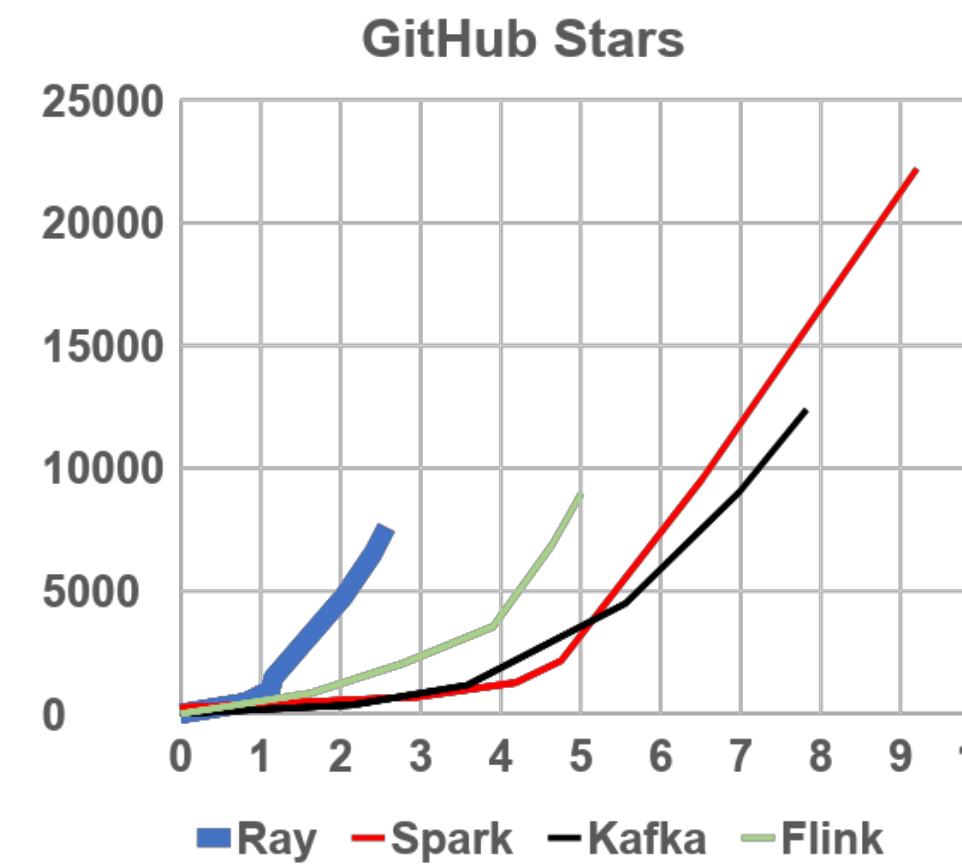
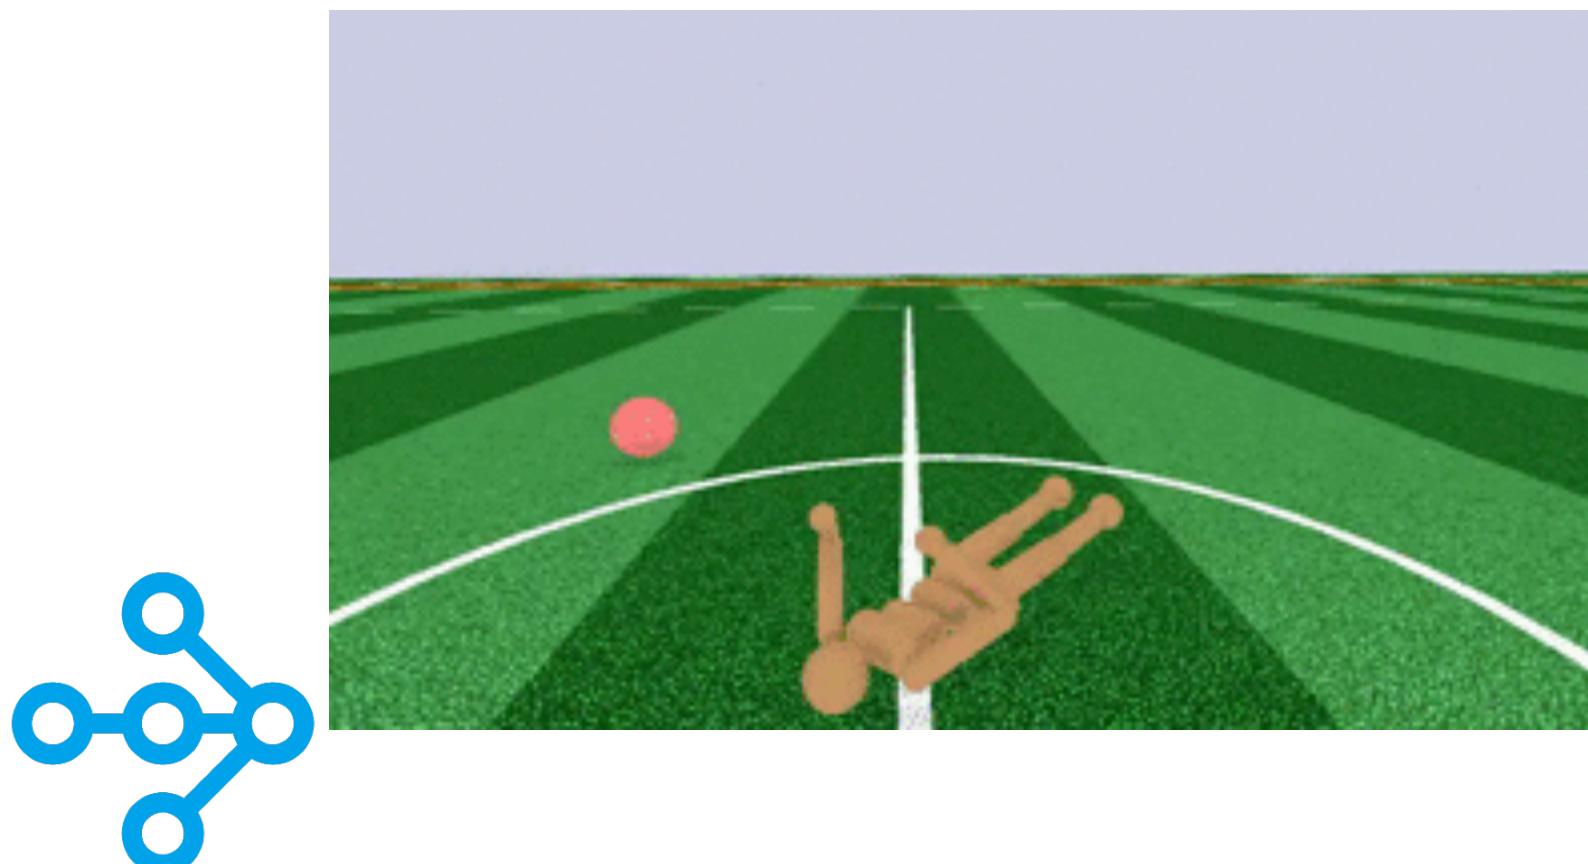
What about Kubernetes (and others...)?

- Ray scaling is very fine grained.
- It operates within the “nodes” of coarse-grained managers
- Containers, VMs, or physical machines



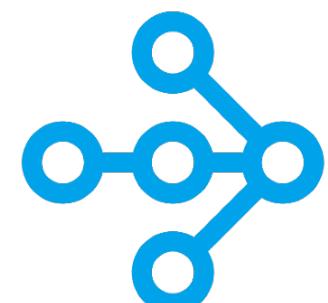
Conclusion

- Ray is the new state-of-the-art for distributed computing
 - The shortest path from your laptop to the cloud
 - Run complex distributed tasks on large clusters from simple code on your laptop



About Anyscale, Inc

- Spun out of U.C. Berkeley
- Making Ray the standard for distributed computing
- We are hiring!
- <https://anyscale.com>



Questions?

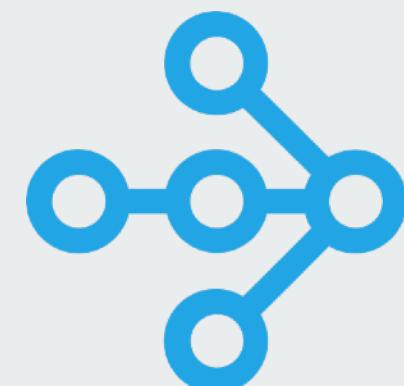
ray.io

anyscale.com - We're Hiring!

anyscale.com/events

raysummit.org

dean@anyscale.com



RAY SUMMIT
CONNECT



© 2019-2020, Anyscale.io

